Chapter 18 — Testing, Build and Tooling

Everything so far was compiled with a bare `javac` and run with `java`. Real projects add four things: a **build tool** that downloads **libraries** (other people's code, shipped as JARs) and runs every step (compile, test, package) in order, **automated tests** that prove the code still works after each change, **logging** that tells you what a running program did, and a **package** (a JAR, §18.12) you can ship. This chapter teaches all four on one small shop application, plus the JDBC (Java Database Connectivity) basics for talking to a database, and ends with a map of where to go next (Spring Boot, the most popular Java framework for services). **How the code was run.** All sources are in `code/`:

Folder	What	How it runs
`code/ch18/shop/`	the application: `Cart`, `Discounts`, `OrderService`, … and `ShopDemo` (has `main`)	the book's runner (`javac --release 25`, JDK 26)
`code/test/ch18/shop/`	JUnit tests (marked `// tools_run: skip` — they have no `main`)	`code/run_tests.sh`
`code/ch18/logging/`	`JulDemo`	the book's runner
`code/ch18/jdbc/`	`JdbcOrderRepository`, `JdbcDemo` (skip-marked: needs a database)	by hand against the SQL Handbook's private PostgreSQL (a popular open-source relational database) (18.11)
`code/build-files/`	`pom.xml`, `build.gradle.kts`, `settings.gradle.kts`	Maven 3.9.14 and Gradle 9.4.1 **offline**, from the local artifact cache (18.2, 18.3)
`code/package_jar.sh`	builds and runs a JAR with the `jar` tool	by hand (18.12)

Test libraries were downloaded from Maven Central (`repo1.maven.org`, SHA-1 checksum verified) into `ROOT/.java-books/lib/`: **JUnit 6.1.3** (`junit-platform-console-standalone`, the latest release at the time of writing), **Mockito 5.24.0** (`mockito-core`, `mockito-junit-jupiter`) with **Byte Buddy 1.18.14** (`byte-buddy`, `byte-buddy-agent`) and **Objenesis 3.3**, and **AssertJ 3.27.7**. The Maven and Gradle builds used the versions already in the local cache (JUnit 6.0.3, Mockito 5.23.0) because they ran without network. Every output below is real.

18.1 The shape of a real Java project

What it is

A Java project is more than `.java` files. It has **dependencies** (other people's JARs that your code needs), **tests**, **resources** (config files, templates), and **build outputs** (compiled classes, the final JAR). Both mainstream build tools agree on one **standard directory layout**, so any Java developer can open any project:

```
shop/
├── pom.xml                 (Maven)  or  build.gradle.kts + settings.gradle.kts (Gradle)
├── mvnw / gradlew          the build-tool "wrapper" script (commit it)
└── src/
    ├── main/
    │   ├── java/ch18/shop/Cart.java       production code (package = folder path)
    │   └── resources/                      copied as-is onto the class path (e.g. logback.xml)
    └── test/
        ├── java/ch18/shop/CartTest.java    tests: same package, separate folder
        └── resources/                      test-only resources
target/  (Maven)  or  build/  (Gradle)       generated — never commit
```

Tests live in the **same package** as the code they test but in a **separate source folder**. That gives them access to **package-private** members (§5.6), while the build keeps test classes out of the final JAR. This

chapter's `code/` folder mirrors it: `code/ch18/shop` is "main", `code/test/ch18/shop` is "test".

Coordinates and repositories

Every library is identified by three **coordinates**, written `groupId:artifactId:version` — for example `org.assertj:assertj-core:3.27.7`.

Term	Meaning
groupId	the organisation, reversed domain style (`org.junit.jupiter`)
artifactId	the library's name (`junit-jupiter-api`)
version	a release (`6.1.3`) or a **SNAPSHOT** (`1.0.0-SNAPSHOT`: work in progress, may change)
artifact	a file published under those coordinates — usually a JAR plus a `.pom` describing its own dependencies
repository	a server storing artifacts; the public one is **Maven Central** (`repo1.maven.org/maven2`, searchable at central.sonatype.com)
local repository	the cache on your disk (`~/.m2/repository` for Maven, `~/.gradle/caches` for Gradle) — downloaded once, reused by every project
transitive dependency	a dependency of a dependency, pulled in automatically (Mockito brings Byte Buddy and Objenesis)

A repository is laid out by coordinates: `org.assertj:assertj-core:3.27.7` lives at `repo1.maven.org/maven2/org/assertj/assertj-core/3.27.7/assertj-core-3.27.7.jar` — how this chapter fetched its jars by hand, and what a build tool does for every (transitive) dependency.

Why a build tool

Without one you'd download jars by hand, keep versions consistent, type long `-cp` lines and remember to test before packaging. A build tool turns that into one file and one command (`mvn verify`, `./gradlew build`) that behaves the same on every machine and on **CI** (*continuous integration*: a server that builds and tests every pushed change).

TL;DR

Standard layout (`src/main/java`, `src/test/java`), libraries identified by `group:artifact:version`, downloaded from Maven Central into a local cache. The build tool does the rest.

18.2 Maven

What it is

Apache Maven (2004) is a build tool driven by a **declarative** XML file, `pom.xml` (**P**roject **O**bject **M**odel). You describe *what* the project is — coordinates, dependencies, **plugin** settings (a plugin is an add-on JAR that does one kind of build work) — and Maven's fixed **lifecycle** decides *how* and *in which order* things happen. The current stable line is Maven 3.9 (3.9.16 at the time of writing); Maven 4.0 was still a release candidate.

The `pom.xml` — `code/build-files/maven/pom.xml`

Key parts (the full file adds pinned plugin versions):

```xml
<project ...>
  <modelVersion>4.0.0</modelVersion>
  <groupId>com.example</groupId>
  <artifactId>shop</artifactId>
  <version>1.0.0-SNAPSHOT</version>
  <packaging>jar</packaging>

  <properties>
    <maven.compiler.release>25</maven.compiler.release>   <!-- same as javac option: release 25 -->
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
  </properties>

  <dependencyManagement>                        <!-- versions only; adds nothing by itself -->
    <dependencies>
      <dependency>                              <!-- a BOM pins every JUnit module's version -->
        <groupId>org.junit</groupId> <artifactId>junit-bom</artifactId> <version>6.0.3</version>
        <type>pom</type> <scope>import</scope>
      </dependency>
      <dependency>                              <!-- one Byte Buddy for Mockito AND AssertJ -->
        <groupId>net.bytebuddy</groupId> <artifactId>byte-buddy</artifactId> <version>1.18.11</version>
      </dependency>
      <!-- ... byte-buddy-agent 1.18.11 ... -->
    </dependencies>
  </dependencyManagement>

  <dependencies>
    <dependency>
      <groupId>org.junit.jupiter</groupId> <artifactId>junit-jupiter</artifactId>
      <scope>test</scope>                       <!-- no version: taken from the BOM -->
    </dependency>
    <dependency>
      <groupId>org.mockito</groupId> <artifactId>mockito-junit-jupiter</artifactId>
      <version>5.23.0</version> <scope>test</scope>
    </dependency>
    <dependency>
      <groupId>org.assertj</groupId> <artifactId>assertj-core</artifactId>
      <version>3.27.7</version> <scope>test</scope>
    </dependency>
  </dependencies>

  <build>
    <plugins>
      <!-- maven-dependency-plugin 'properties' goal: exposes the mockito-core jar path as a property -->
      <plugin>
        <artifactId>maven-surefire-plugin</artifactId> <version>3.5.4</version>
        <configuration>
          <argLine>-javaagent:${org.mockito:mockito-core:jar} -Xshare:off</argLine>
          <excludedGroups>demo-failure</excludedGroups>        <!-- JUnit tags to skip -->
        </configuration>
      </plugin>
      <plugin>
        <artifactId>maven-jar-plugin</artifactId> <version>3.5.1</version>
        <configuration>
          <archive><manifest><mainClass>ch18.shop.ShopDemo</mainClass></manifest></archive>
        </configuration>
      </plugin>
    </plugins>
  </build>
</project>
```

What each part does:

- **Coordinates** of *this* project — what other projects would write to depend on it.
- `maven.compiler.release` — passes `--release 25` to `javac` (Chapter 1 §1.5).
- `dependencyManagement` — a table of versions; it *declares* nothing. A **BOM** (bill of materials) is a pom that is only such a table, imported with `<scope>import</scope>`, so all JUnit modules share one version.
- `dependencies` — what the project actually uses, each with a **scope**.

- `build/plugins` — configure the plugins that do the work (compile, test, jar).

Dependency scopes

Scope	On compile class path	On test class path	Packaged / at run time	Typical example
`compile` (default)	yes	yes	yes	Jackson, Guava
`provided`	yes	yes	no — the runtime supplies it	Servlet API inside an **app server** (e.g. Tomcat)
`runtime`	no	yes	yes	JDBC **driver** (the vendor library that talks to a specific database; code compiles against `java.sql` only)
`test`	no (main code)	yes	no	JUnit, Mockito, AssertJ
`import`	—	—	—	only for a BOM inside `dependencyManagement`

`system` also exists but is deprecated — don't use it.

Transitive dependencies and version conflicts

Maven builds a dependency **tree**. When two paths bring *different versions* of one library, Maven picks the **nearest** one in the tree; on a tie, the one declared **first**. This build hit a real conflict: `mockito-core` wants Byte Buddy 1.17.7 at depth 3, `assertj-core` wants 1.18.3 at depth 2 — so "nearest wins" picked AssertJ's version, and Mockito's `byte-buddy-agent` would have stayed at a different version. The fix is to **pin** the version in `dependencyManagement`. The real tree after pinning (`mvn -o dependency:tree`):

```
[INFO] com.example:shop:jar:1.0.0-SNAPSHOT
[INFO] +- org.junit.jupiter:junit-jupiter:jar:6.0.3:test
[INFO] |  +- org.junit.jupiter:junit-jupiter-api:jar:6.0.3:test
[INFO] |  |  +- org.opentest4j:opentest4j:jar:1.3.0:test
[INFO] |  |  +- org.junit.platform:junit-platform-commons:jar:6.0.3:test
[INFO] |  |  +- org.apiguardian:apiguardian-api:jar:1.1.2:test
[INFO] |  |  \- org.jspecify:jspecify:jar:1.0.0:test
[INFO] |  +- org.junit.jupiter:junit-jupiter-params:jar:6.0.3:test
[INFO] |  \- org.junit.jupiter:junit-jupiter-engine:jar:6.0.3:test
[INFO] |     \- org.junit.platform:junit-platform-engine:jar:6.0.3:test
[INFO] +- org.mockito:mockito-junit-jupiter:jar:5.23.0:test
[INFO] |  \- org.mockito:mockito-core:jar:5.23.0:test
[INFO] |     +- net.bytebuddy:byte-buddy-agent:jar:1.18.11:test
[INFO] |     \- org.objenesis:objenesis:jar:3.3:test
[INFO] \- org.assertj:assertj-core:jar:3.27.7:test
[INFO]    \- net.bytebuddy:byte-buddy:jar:1.18.11:test
```

Other conflict tools: `<exclusions>` on a dependency removes one transitive artifact; the **Maven Enforcer** plugin can fail the build on version-convergence problems (different versions of one library requested).

The lifecycle: phases, plugins and goals



- A **phase** is a step in the lifecycle. Running a phase runs **every phase before it**: `mvn package` = validate, compile, test-compile, test, then package.
- A **plugin** is a JAR that does work; each unit of work is a **goal**, written `plugin:goal` (`compiler:compile`, `surefire:test`). The `jar` packaging binds default goals to phases:

Phase	Default goal bound (packaging `jar`)	Result
process-resources	`resources:resources`	copy `src/main/resources` into `target/classes`
compile	`compiler:compile`	`target/classes`
test-compile	`compiler:testCompile`	`target/test-classes`
test	`surefire:test`	runs unit tests; fails the build on a failure
package	`jar:jar`	`target/shop-1.0.0-SNAPSHOT.jar`
integration-test / verify	`failsafe:*` if configured	integration tests (`*IT.java`) and checks
install	`install:install`	copy the JAR into `~/.m2/repository` for other local projects
deploy	`deploy:deploy`	upload to a remote repository (company Nexus/Artifactory, Maven Central)

Two other lifecycles exist: **clean** (`mvn clean` deletes `target/`) and **site** (project reports). Surefire runs test classes whose names match `Test*`, `*Test`, `*Tests` or `*TestCase` by default.

Everyday commands

Command	Does
`mvn clean verify`	full build from scratch with all checks — what CI usually runs
`mvn test`	compile and run unit tests
`mvn -Dtest=CartTest test`	run one test class (`CartTest#newCartIsEmpty` for one method)
`mvn package -DskipTests`	build the JAR without running tests (compiles them; `-Dname=value` sets a property)
`mvn dependency:tree`	print the resolved tree
`mvn -o …`	offline: use only the local repository
`./mvnw …`	the **Maven Wrapper**: downloads the exact Maven version the project pins (`mvn wrapper:wrapper` creates it); commit it so everyone and CI use the same Maven without installing it

A real build

The sources were copied into the standard layout next to this `pom.xml` and built **offline** (`mvn -o -B package`; `-B` = batch mode, no prompts or colours) from the local cache. Trimmed real output:

```
[INFO] --- dependency:3.10.0:properties (default) @ shop ---
[INFO] --- resources:3.3.1:resources (default-resources) @ shop ---
[INFO] --- compiler:3.15.0:compile (default-compile) @ shop ---
[INFO] Compiling 11 source files with javac [debug release 25] to target/classes
[INFO] --- resources:3.3.1:testResources (default-testResources) @ shop ---
[INFO] --- compiler:3.15.0:testCompile (default-testCompile) @ shop ---
[INFO] Compiling 6 source files with javac [debug release 25] to target/test-classes
[INFO] --- surefire:3.5.4:test (default-test) @ shop ---
[INFO] Using auto detected provider org.apache.maven.surefire.junitplatform.JUnitPlatformProvider
[INFO] Running ch18.shop.DiscountsTest
[INFO] Tests run: 27, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.105 s -- in ch18.shop.DiscountsTest
[INFO] Running ch18.shop.AssertjStyleTest
[INFO] Tests run: 6, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.395 s -- in ch18.shop.AssertjStyleTest
[INFO] Running ch18.shop.OrderServiceTest
[INFO] Tests run: 6, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.166 s -- in ch18.shop.OrderServiceTest
...
[INFO] Results:
[WARNING] Tests run: 48, Failures: 0, Errors: 0, Skipped: 1
[INFO] --- jar:3.5.1:jar (default-jar) @ shop ---
[INFO] Building jar: .../java18_mvn/target/shop-1.0.0-SNAPSHOT.jar
[INFO] BUILD SUCCESS
[INFO] Total time:  2.645 s
```

That is the lifecycle table in action. The real run taught two lessons: the first attempt failed because Maven 3.9.14's *default* resources plugin (3.4.0) wasn't cached — **pin every plugin version** so a Maven upgrade can't silently change your build — and the Byte Buddy conflict above. `java -jar target/shop-1.0.0-SNAPSHOT.jar` runs `ShopDemo` (output in 18.12).

Multi-module builds (in brief)

Big projects split into **modules** (`shop-domain`, `shop-web`, …) under a parent pom with `<packaging>pom</packaging>` and a `<modules>` list. The parent holds shared `dependencyManagement` and plugin settings; `mvn verify` at the root builds modules in dependency order.

Common mistakes

- Omitting plugin versions (build changes when Maven is upgraded) or using version ranges.
- Test libraries without `<scope>test</scope>` — they end up in the product.
- Fighting a conflict with random exclusions instead of reading `dependency:tree` and pinning in `dependencyManagement`.
- Running `mvn install` just to build — `verify` is enough unless another local project needs the JAR.

In the interview

30-second answer: "Maven is a declarative build tool. The `pom.xml` gives coordinates, dependencies with scopes, and plugin configuration. It resolves transitive dependencies from Maven Central into `~/.m2`, choosing the nearest version on conflicts, and runs a fixed lifecycle — compile, test, package, verify, install, deploy — where running a phase runs all earlier ones. Plugins' goals are bound to phases."

- *`dependencies` vs `dependencyManagement`?* → The first adds dependencies; the second only fixes versions (and scopes) for when they are used, including transitive ones — and overrides "nearest wins".
- *What is a BOM?* → A pom that only contains `dependencyManagement`; importing it aligns versions of a family of libraries (JUnit, Spring, Jackson).
- *`install` vs `deploy`?* → Local repository vs remote repository.

TL;DR

`pom.xml` = coordinates + dependencies (with scopes) + plugins. Fixed lifecycle; a phase runs all earlier phases. Pin plugin and conflicting library versions; read `dependency:tree`.

18.3 Gradle

What it is

Gradle (2008) is a build tool whose build file is **code** — a Kotlin (`build.gradle.kts`, the default for new builds) or Groovy (`build.gradle`) script. Instead of a fixed lifecycle it builds a **graph of tasks** and runs only what's needed, skipping **up-to-date** tasks (**incremental builds**: only redo work whose inputs changed — why a second `./gradlew build` with no changes is fast) and reusing outputs from a **build cache** (stored results of earlier runs, possibly shared between machines). It uses the same Maven repositories and coordinates. Gradle 9 (9.4.1 here) needs Java 17+ to run the build itself.

The build file — `code/build-files/gradle/build.gradle.kts`

```
plugins {
    application                                  // adds 'java' + the run task
}

group = "com.example"
version = "1.0.0-SNAPSHOT"

java {
    toolchain { languageVersion = JavaLanguageVersion.of(25) }   // build with a JDK 25
}

repositories {
    mavenCentral()
}

val mockitoAgent = configurations.create("mockitoAgent")   // a separate list holding just the agent jar

dependencies {
    testImplementation(platform("org.junit:junit-bom:6.0.3"))
    testImplementation("org.junit.jupiter:junit-jupiter")
    testRuntimeOnly("org.junit.platform:junit-platform-launcher")
    testImplementation("org.mockito:mockito-junit-jupiter:5.23.0")
    testImplementation("org.assertj:assertj-core:3.27.7")
    testImplementation("net.bytebuddy:byte-buddy:1.18.11")        // one Byte Buddy version for everybody
    testImplementation("net.bytebuddy:byte-buddy-agent:1.18.11")
    mockitoAgent("org.mockito:mockito-core:5.23.0") { isTransitive = false }
}

application { mainClass = "ch18.shop.ShopDemo" }

tasks.test {
    useJUnitPlatform { excludeTags("demo-failure") }
    jvmArgs("-javaagent:${mockitoAgent.asPath}", "-Xshare:off")
    testLogging { events("passed", "skipped", "failed") }
}

tasks.jar { manifest { attributes["Main-Class"] = "ch18.shop.ShopDemo" } }
```

`settings.gradle.kts` holds one line, `rootProject.name = "shop"` (multi-project builds list their subprojects there with `include(...)`).

- **Toolchain** — Gradle finds (or downloads) a JDK 25 to compile and test with, independent of the JDK running Gradle.
- **`platform(...)`** — imports a BOM, like Maven's `import` scope.
- **`testRuntimeOnly("org.junit.platform:junit-platform-launcher")`** — required since Gradle 9 when running on the JUnit Platform.
- **`useJUnitPlatform()`** — without it Gradle looks for JUnit 4 tests and finds none.

Configurations (Gradle's scopes)

Gradle configuration	Closest Maven scope	Meaning
`implementation`	compile	needed to compile and run, **not** exposed to consumers' compile class path
`api` (plugin `java-library`)	compile	part of your public API: consumers compile against it too
`compileOnly`	provided	compile time only (annotations, provided APIs)
`runtimeOnly`	runtime	run time only (JDBC driver, logging backend)
`testImplementation` / `testRuntimeOnly`	test	the same, for tests

`implementation` vs `api` is a real improvement over Maven: hiding implementation dependencies means fewer recompilations and less accidental **coupling**.

A real build

The same sources, built offline with Gradle 9.4.1 (`gradle --offline clean build --console=plain`; for the offline run `mavenCentral()` was replaced by `mavenLocal()` so Gradle could read the `~/.m2` cache):

```
> Task :clean
> Task :compileJava
...                                (classes, jar, startScripts, distTar, distZip, assemble, compileTestJava,
testClasses)
> Task :test
Cart > itemsViewIsReadOnly() PASSED
Cart > shippingIsFreeAbove50() SKIPPED
Cart > adding the same SKU twice merges the quantities PASSED
...
> Task :check
> Task :build
BUILD SUCCESSFUL in 5s
8 actionable tasks: 8 executed
```

`build` depends on `assemble` (jar + distributions from the `application` plugin) and `check` (tests). `gradle run` then ran `ShopDemo`. Useful commands: `./gradlew test --tests 'ch18.shop.CartTest'`, `./gradlew dependencies --configuration testRuntimeClasspath`, `./gradlew tasks`, `./gradlew build --scan` (online build report).

The Gradle Wrapper. `gradle wrapper` generates `gradlew`, `gradlew.bat` and `gradle/wrapper/gradle-wrapper.properties` (holding `distributionUrl`). Commit them and always run `./gradlew`. The user's OCP study repo is a real example: its `build.gradle` (Groovy DSL) pins a **Java 21 toolchain**, gives every chapter its own *source set* with its own compile task, and its wrapper pins Gradle 9.4.1.

Maven vs Gradle

	Maven	Gradle
Build file	XML, declarative	Kotlin/Groovy script
Model	fixed lifecycle of phases	task graph, incremental, build cache, daemon (a background process kept between builds)
Speed on big builds	fine	usually faster (up-to-date checks, caching)
Flexibility	limited — write a plugin for custom logic	very high (and easy to over-engineer)
Learning curve	gentle, very uniform across projects	steeper; builds differ more
Where you'll meet it	most enterprise/back-end Java, Spring examples	Android (mandatory), many newer projects, Kotlin

Both are good choices; follow the team. Spring Initializr (start.spring.io) generates either.

Common mistakes

- Using `gradle` from your PATH instead of `./gradlew`.
- Forgetting `useJUnitPlatform()` — "0 tests" and a green build.
- Putting everything in `api` "to be safe".
- Custom logic scattered in the build script; move it to a plugin in `buildSrc` or a convention plugin (a shared plugin applying a team's standard settings).

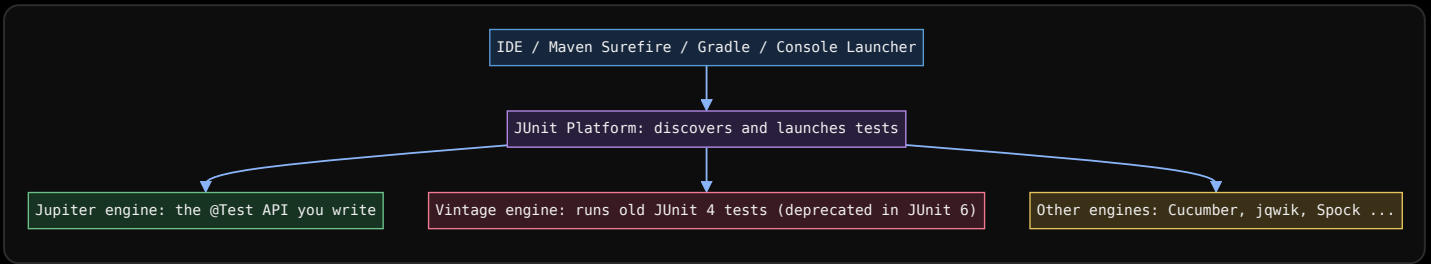
TL;DR

Gradle = scripted task graph with incremental builds; same repositories and coordinates as Maven. Use the wrapper, `implementation` by default, and remember `useJUnitPlatform()`.

18.4 JUnit: your first automated tests

What it is

An **automated test** runs code with known inputs and checks the result, reporting pass or fail with no human reading output. A **unit test** checks one small piece (a class or method) in isolation, in milliseconds. **JUnit** is Java's standard test **framework** (a library that calls *your* test methods). Its modern generation (JUnit 5, 2017; **JUnit 6**, released 30 September 2025) has three parts:



- **JUnit Platform** — the launcher that tools talk to.
- **JUnit Jupiter** — the programming model (`org.junit.jupiter.api.*`) and its engine.
- **JUnit Vintage** — an engine for legacy JUnit 4 tests.

JUnit 6 kept the same Jupiter API (these examples are identical in JUnit 5), moved the minimum to **Java 17**, gave Platform, Jupiter and Vintage **one shared version number** (6.x instead of 1.x/5.x), added JSpecify nullness annotations (marking whether a value may be `null`), a `--fail-fast` launcher option and test cancellation, and deprecated Vintage. Most JUnit 5 test code runs unchanged.

Why it matters

Tests let you **refactor** (restructure without changing behaviour), upgrade or fix without fear — hundreds of checks in a second. They document behaviour precisely ("SAVE25 needs a subtotal of at least 100"), catch **regressions** (things that used to work and broke) before users do, and push toward better design (hard-to-test code is usually too coupled).

The code under test

`Cart` holds `Item` records; adding the same SKU (*stock-keeping unit*, a product's id) twice merges quantities; money is `BigDecimal` (Chapter 7 §7.7).

```
public final class Cart {
    private final Map<String, Item> items = new LinkedHashMap<>();
    public void add(Item item) {
        items.merge(item.sku(), item,
                (old, extra) -> new Item(old.sku(), old.name(), old.unitPrice(), old.quantity() +
extra.quantity()));
    }
    public void remove(String sku) {
        if (items.remove(sku) == null) throw new NoSuchElementException("not in cart: " + sku);
    }
    public List<Item> items()   { return List.copyOf(items.values()); }
    public boolean isEmpty()    { return items.isEmpty(); }
    public int count()          { return items.values().stream().mapToInt(Item::quantity).sum(); }
    public BigDecimal subtotal() { /* sum of line totals, setScale(2, HALF_UP) */ }
}
```

The tests — `code/test/ch18/shop/CartTest.java`

```java
@Tag("fast")
@DisplayName("Cart")
class CartTest {

    private Cart cart;                                  // a NEW CartTest instance is created for every test

    @BeforeEach
    void setUp() { cart = new Cart(); }

    @Test
    void newCartIsEmpty() {
        assertTrue(cart.isEmpty());
        assertEquals(0, cart.count());
        assertEquals(new BigDecimal("0.00"), cart.subtotal());
    }

    @Test
    @DisplayName("adding the same SKU twice merges the quantities")
    void addSameSkuMerges() {
        cart.add(new Item("MUG", "Mug", new BigDecimal("12.50"), 1));
        cart.add(new Item("MUG", "Mug", new BigDecimal("12.50"), 2));

        assertAll("merged line",                        // reports EVERY failing check, not just the first
                () -> assertEquals(1, cart.items().size()),
                () -> assertEquals(3, cart.count()),
                () -> assertEquals(new BigDecimal("37.50"), cart.subtotal()));
    }

    @Test
    void removingUnknownSkuThrows() {
        NoSuchElementException e = assertThrows(NoSuchElementException.class, () -> cart.remove("NOPE"));
        assertEquals("not in cart: NOPE", e.getMessage());
    }

    @Test
    @Disabled("shipping costs not implemented yet")
    void shippingIsFreeAbove50() { fail("not written yet"); }

    @Nested
    @DisplayName("with a book and a mug")
    class WithTwoItems {
        @BeforeEach
        void fill() {                                   // runs AFTER the outer setUp()
            cart.add(new Item("BK-1", "Java book", new BigDecimal("45.00"), 2));
            cart.add(new Item("MUG", "Mug", new BigDecimal("12.50"), 1));
        }
        @Test void subtotalAddsLineTotals() { assertEquals(new BigDecimal("102.50"), cart.subtotal()); }
        @Test void removeDropsTheWholeLine() { cart.remove("BK-1"); assertEquals(1, cart.count()); }
    }
    // ... itemsViewIsReadOnly()
}
```

How it works:

1. `@Test` marks a test method. Test classes and methods don't need to be `public` (package-private is the convention); they must not be `private`.
2. The static import `org.junit.jupiter.api.Assertions.*` gives the `assertXxx` methods (an **assertion** fails the test if its check isn't true). Argument order is **(expected, actual)**, with an optional message **last**.
3. A failed assertion throws `AssertionFailedError`; any other uncaught exception also fails the test. A test that finishes normally passes.
4. `assertThrows` runs the lambda, checks the exception type and **returns** the exception so you can inspect it.
5. `@Nested` inner classes group tests sharing a situation; the outer `@BeforeEach` runs first.

The assertions you'll use

Assertion	Checks
`assertEquals(exp, act)` / `assertNotEquals`	`equals` (with a `delta` overload for doubles)
`assertTrue` / `assertFalse`	a condition
`assertNull` / `assertNotNull`	null-ness
`assertSame` / `assertNotSame`	same reference (`==`)
`assertArrayEquals`, `assertIterableEquals`	element by element
`assertLinesMatch`	lists of strings, regex lines allowed
`assertInstanceOf(Type.class, obj)`	type, returns the cast value
`assertThrows` / `assertThrowsExactly`	an exception of that type (or subtype / exactly that type)
`assertDoesNotThrow`	no exception; returns the result
`assertAll(...)`	runs every check, reports all failures together
`assertTimeout(Duration, ...)`	finishes in time (`assertTimeoutPreemptively` aborts it)
`fail("msg")`	fails immediately

Messages can be a `Supplier<String>` (`() -> "got " + r`), built only on failure.

Lifecycle — `LifecycleTest.java`

```
class LifecycleTest {
    static int instances;
    LifecycleTest()            { System.out.println("  constructor  (instance #" + ++instances + ")"); }

    @BeforeAll static void beforeAll() { System.out.println("@BeforeAll   (static, once)"); }
    @BeforeEach void beforeEach()      { System.out.println("  @BeforeEach"); }
    @Test void first()                 { System.out.println("    test: first"); }
    @Test void second()                { System.out.println("    test: second"); }
    @AfterEach void afterEach()        { System.out.println("  @AfterEach"); }
    @AfterAll static void afterAll()   { System.out.println("@AfterAll    (static, once)"); }
}
```

Output (from the run below):

```
@BeforeAll   (static, once)
  constructor  (instance #1)
  @BeforeEach
    test: second
  @AfterEach
  constructor  (instance #2)
  @BeforeEach
    test: first
  @AfterEach
@AfterAll    (static, once)
```

Three facts to remember:

- **A fresh instance per test method** (default `@TestInstance(Lifecycle.PER_METHOD)`), so fields can't leak state between tests. That is also why `@BeforeAll`/`@AfterAll` must be `static` — unless you annotate the class `@TestInstance(Lifecycle.PER_CLASS)`.
- **Method order is deterministic but deliberately not obvious** — here `second` ran before `first`. Tests must not depend on each other. If you truly need an order: `@TestMethodOrder(MethodOrderer.OrderAnnotation.class)` + `@Order(n)`.
- `@AfterEach` runs even when the test fails — use it for cleanup.

More annotations

Annotation	Purpose
`@DisplayName("…")`	readable name in reports
`@Disabled("reason")`	skip, with a reason shown in the report (`~` below)
`@Tag("fast")`	label for filtering (`--include-tag`, Maven `<groups>`, Gradle `includeTags`)
`@Timeout(2)`	fail if slower than 2 s
`@TempDir Path dir`	a fresh temporary directory, deleted afterwards (18.7)
`@ExtendWith(X.class)`	plug in an **extension** (hooks into the test lifecycle, e.g. Mockito, Spring) — replaces JUnit 4 runners and rules
`@EnabledOnOs`, `@EnabledIfEnvironmentVariable`, `@EnabledForJreRange` …	conditional execution

Assumptions (`assumeTrue(cond)`, `assumingThat`) stop a test as *aborted* rather than failed when a precondition doesn't hold (e.g. "only on CI"). Aborted = an assumption didn't hold; failed = an assertion or unexpected exception.

Running the tests from the command line

IDEs and build tools run JUnit for you. Without them, the **Console Launcher** (`junit-platform-console-standalone`, one JAR containing the Platform, Jupiter, Vintage and Suite engines) does it. The plain form — compile, then point the launcher at the classes:

```
$ javac --release 25 -cp '.java-books/lib/*' -d /tmp/java18t $(find ch18/shop test -name '*.java')
$ java -jar junit-platform-console-standalone-6.1.3.jar execute \
      --class-path /tmp/java18t:assertj-core-3.27.7.jar --scan-class-path \
      --include-classname '.*(Cart|Discounts)Test' --details=summary --disable-banner

Test run finished after 101 ms
...                                (13 containers found, started and successful)
[        34 tests found           ]
[         1 tests skipped         ]
[        33 tests successful      ]
[         0 tests failed          ]
```

`--scan-class-path` discovers every class whose name matches the default pattern (`Test*`, `*Test`, `*Tests`). The **exit code** is 0 only if nothing failed — that's what CI checks. Filter by tag with `--include-tag fast` (console), `<groups>fast</groups>` (Surefire) or `useJUnitPlatform { includeTags("fast") }` (Gradle).

For the Mockito tests the launcher must be started differently, because Mockito wants to be attached as a **Java agent** (a JAR the JVM loads before your code so it can modify classes as they load; 18.6) and the agent must be the *same* JAR the tests load. So `code/run_tests.sh` puts everything on one class path and starts the launcher by its main class:

```
LIB="../../../../.java-books/lib"; CP="$LIB/*"; OUT=$(mktemp -d)
javac --release 25 -cp "$CP" -d "$OUT" $(find ch18/shop test -name '*.java') || exit 1
if [ "${1:-}" = fail ]; then TAGS=(--include-tag demo-failure); else TAGS=(--exclude-tag demo-failure); fi
java -Xshare:off -javaagent:"$(ls "$LIB"/mockito-core-*.jar)" \
     -Djava.util.logging.SimpleFormatter.format='%4$s: %5$s%n' \
     -cp "$OUT:$CP" org.junit.platform.console.ConsoleLauncher execute \
     --scan-class-path "$OUT" "${TAGS[@]}" \
     --disable-banner --disable-ansi-colors --details=tree
```

`bash code/run_tests.sh` output (tree trimmed; notes in parentheses added):

```
WARNING: declined: customer=bob total=102.50
INFO: order placed: id=ORD-1 total=102.50 ref=PAY-9
INFO: order placed: id=ORD-1 total=76.88 ref=PAY-9
INFO: order placed: id=ORD-7 total=92.25 ref=ok
@BeforeAll   (static, once)
  ... (lifecycle lines shown above)
|
├─ JUnit Platform Suite ✔
├─ JUnit Jupiter ✔
|  ├─ AssertjStyleTest ✔          (6 tests, 18.7)
|  ├─ DiscountsTest ✔             (27 tests, 18.5)
|  ├─ Cart ✔
|  |  ├─ itemsViewIsReadOnly() ✔
|  |  ├─ shippingIsFreeAbove50() ↷ shipping costs not implemented yet
|  |  ├─ adding the same SKU twice merges the quantities ✔
|  |  ├─ newCartIsEmpty() ✔
|  |  ├─ removingUnknownSkuThrows() ✔
|  |  └─ with a book and a mug ✔
|  |     ├─ removeDropsTheWholeLine() ✔
|  |     └─ subtotalAddsLineTotals() ✔
|  ├─ OrderServiceTest ✔          (6 tests, 18.6)
|  └─ LifecycleTest ✔
|     ├─ second() ✔
|     └─ first() ✔
└─ JUnit Vintage ✔

Test run finished after 566 ms
...
[        48 tests found           ]
[         1 tests skipped         ]
[        47 tests started         ]
[        47 tests successful      ]
[         0 tests failed          ]
```

The `INFO:` / `WARNING:` lines at the top are the application's own logging (18.10) printed while `OrderServiceTest` ran.

Coming from JUnit 4

JUnit 4	Jupiter (JUnit 5/6)
`org.junit.Test`, public classes/methods	`org.junit.jupiter.api.Test`, package-private OK
`@Before` / `@After`	`@BeforeEach` / `@AfterEach`
`@BeforeClass` / `@AfterClass`	`@BeforeAll` / `@AfterAll`
`@Ignore`	`@Disabled`
`@Test(expected = X.class)`	`assertThrows(X.class, ...)`
`@RunWith`, `@Rule`	`@ExtendWith` extensions
`@Category`	`@Tag`
`assertEquals("msg", exp, act)` — message first	`assertEquals(exp, act, "msg")` — message last

Common mistakes

- Swapping expected and actual — failure messages then read backwards.
- Forgetting `@Test` (the method never runs, silently).
- Mixing JUnit 4's `org.junit.Test` with Jupiter assertions — the Jupiter engine ignores JUnit 4 tests.
- Tests that depend on each other's order or on shared static state.
- Printing results instead of asserting them — a test that can't fail is not a test.

In the interview

> **30-second answer:** "JUnit 5 and 6 split into the Platform, which launches tests for IDEs and build tools, and Jupiter, the API: `@Test`, `@BeforeEach` / `@AfterEach`, static `@BeforeAll` / `@AfterAll`, assertions like `assertEquals`, `assertThrows` and `assertAll`, plus `@Nested`, `@Tag`, `@DisplayName`, `@Disabled` and extensions via `@ExtendWith`. A new test instance is created per method so tests stay independent."

- *Why is `@BeforeAll` static?* → One instance per test method exists (2 tests → 2 `LifecycleTest` instances), so a once-per-class hook can't belong to an instance — unless `PER_CLASS` lifecycle is used.
- *`assertAll` vs several asserts?* → Several asserts stop at the first failure; `assertAll` reports all of them.
- *What changed in JUnit 6?* → Java 17 baseline, a single version for all modules, JSpecify annotations, Vintage deprecated; the Jupiter API is essentially the same.

TL;DR

`@Test` + `Assertions` + lifecycle hooks, one fresh instance per test, no reliance on order. Run from IDE, build tool, or the Console Launcher.

18.5 Parameterized, repeated and dynamic tests

The problem it solves

Rules like "a code is 4+ capital letters followed by 2 digits" need many inputs. Copy-pasting a test per input is noisy; a loop inside one test stops at the first failure and reports one anonymous result. **Parameterized tests** run the same method once per input, each reported separately. (They live in `junit-jupiter-params`, included in the `junit-jupiter` aggregate artifact.)

The rules under test — `Discounts.java`

```
public static boolean isValidCode(String code) {
    return code != null && code.matches("[A-Z]{4,}\\d{2}");
}
public static int percentFor(String code, BigDecimal subtotal) {
    if (!isValidCode(code)) return 0;
    return switch (code) {
        case "SAVE10" -> 10;
        case "SAVE25" -> subtotal.compareTo(BIG_ORDER) >= 0 ? 25 : 0;   // BIG_ORDER = 100.00
        default -> 0;
    };
}
public static BigDecimal apply(BigDecimal amount, int percent) {
    if (percent < 0 || percent > 100) throw new IllegalArgumentException("percent out of range: " + percent);
    return amount.multiply(BigDecimal.valueOf(100 - percent)).divide(BigDecimal.valueOf(100), 2,
RoundingMode.HALF_UP);
}
```

The tests — `DiscountsTest.java`

```java
@ParameterizedTest(name = "{0} is a valid code")
@ValueSource(strings = {"SAVE10", "SAVE25", "WELCOME99"})
void validCodes(String code) { assertTrue(Discounts.isValidCode(code)); }

@ParameterizedTest(name = "[{index}] {0} is rejected")
@NullAndEmptySource                                  // adds null and ""
@ValueSource(strings = {"save10", "SAV10", "SAVE1", "SAVE100"})
void invalidCodes(String code) { assertFalse(Discounts.isValidCode(code)); }

@ParameterizedTest(name = "{0} on {1} -> {2}%")
@CsvSource({
        "SAVE10,  20.00, 10",
        "SAVE25,  99.99,  0",     // boundary: just below 100
        "SAVE25, 100.00, 25",     // boundary: exactly 100
        "SAVE99, 500.00,  0",     // well-formed but unknown
})
void percentForCode(String code, BigDecimal subtotal, int expected) {   // strings converted to BigDecimal/int
    assertEquals(expected, Discounts.percentFor(code, subtotal));
}

static Stream<Arguments> roundingCases() {
    return Stream.of(
            Arguments.of("10.00", 10, "9.00"),
            Arguments.of("0.05", 50, "0.03"),              // 0.025 rounds half-up to 0.03
            Arguments.of("102.50", 25, "76.88"),           // 76.875 -> 76.88
            Arguments.of("19.99", 100, "0.00"));
}

@ParameterizedTest
@MethodSource("roundingCases")
void applyRoundsHalfUpToCents(String amount, int percent, String expected) {
    assertEquals(new BigDecimal(expected), Discounts.apply(new BigDecimal(amount), percent));
}

@RepeatedTest(3)
void applyIsDeterministic() { assertEquals(new BigDecimal("9.00"), Discounts.apply(BigDecimal.TEN, 10)); }

@TestFactory
Stream<DynamicTest> everyPercentKeepsTheResultWithinRange() {
    return Stream.of(0, 33, 50, 99, 100).map(p -> dynamicTest("percent " + p, () -> {
        BigDecimal r = Discounts.apply(new BigDecimal("80.00"), p);
        assertTrue(r.signum() >= 0 && r.compareTo(new BigDecimal("80.00")) <= 0, () -> "got " + r);
    }));
}
```

Output (part of the same run; some rows collapsed, notes in parentheses added):

```
|  ├─ DiscountsTest ✔
|  |  ├─ validCodes(String) ✔
|  |  |  ├─ "SAVE10" is a valid code ✔
|  |  |  ├─ "SAVE25" is a valid code ✔
|  |  |  └─ "WELCOME99" is a valid code ✔
|  |  ├─ applyRoundsHalfUpToCents(String, int, String) ✔   (4 rows)
|  |  ├─ invalidCodes(String) ✔
|  |  |  ├─ [1] null is rejected ✔
|  |  |  ├─ [2] "" is rejected ✔
|  |  |  ├─ [3] "save10" is rejected ✔
|  |  |  ├─ [4] "SAV10" is rejected ✔
|  |  |  ├─ [5] "SAVE1" is rejected ✔
|  |  |  └─ [6] "SAVE100" is rejected ✔
|  |  ├─ applyRejectsBadPercent(int) ✔                      (2 rows)
|  |  ├─ percentForCode(String, BigDecimal, int) ✔
|  |  |  ├─ "SAVE10" on "20.00" -> "10"% ✔
|  |  |  ├─ "SAVE25" on "99.99" -> "0"% ✔
|  |  |  ├─ "SAVE25" on "100.00" -> "25"% ✔
|  |  |  └─ "SAVE99" on "500.00" -> "0"% ✔
|  |  ├─ everyPercentKeepsTheResultWithinRange() ✔
|  |  |  ├─ percent 0 ✔
|  |  |  ├─ ... percent 33, 50, 99 ✔
|  |  |  └─ percent 100 ✔
|  |  └─ applyIsDeterministic() ✔
|  |     ├─ repetition 1 of 3 ✔
|  |     ├─ repetition 2 of 3 ✔
|  |     └─ repetition 3 of 3 ✔
```

Notice: the name pattern placeholders are `{0}`, `{1}` … and `{index}`; this JUnit version shows text arguments **in quotes**, so `null` (no quotes) and `""` (the empty string) are easy to tell apart. `@CsvSource` values start as text and are **converted** to the parameter types (`BigDecimal`, `int`) automatically.

Argument sources

Source	Supplies	Example use
`@ValueSource(strings/ints/…)`	one literal argument per run	valid/invalid inputs
`@NullSource`, `@EmptySource`, `@NullAndEmptySource`	`null`, empty string/collection/array	edge cases
`@EnumSource(Level.class)`	each enum constant (filterable with `names`)	behaviour per enum value
`@CsvSource({...})` / `@CsvFileSource(resources = "/x.csv")`	rows of several arguments	input → expected tables
`@MethodSource("factory")`	a `static` method in the test class (or `Class#method` elsewhere; non-static allowed with `PER_CLASS`) returning `Stream<Arguments>` or a collection	objects that can't be written as text
`@FieldSource` (5.11+)	a static field holding the arguments	shared fixed data
`@ArgumentsSource(MyProvider.class)`	your own provider	reusable generators

Related features: `@RepeatedTest(n)` runs a test n times (useful for suspected flakiness); `@TestFactory` returns **dynamic tests** built at run time, so the number of tests can depend on data (parameterized = fixed method fed by a source; dynamic = arbitrary tests from a factory). Since JUnit 5.13 a whole class can be parameterized with `@ParameterizedClass`.

Common mistakes

- Many inputs in one `@Test` with a loop — the first failure hides the rest.

- Hiding the *expected* value computation inside the test (re-implementing the logic). Expected values should be literal, worked out by hand — like `76.88` above.
- Forgetting boundaries: the `99.99` / `100.00` rows are the most valuable ones in that table.

TL;DR

One method, many inputs, each reported separately: `@ParameterizedTest` + a source. Put boundary values in the table.

18.6 Test doubles and Mockito

The problem it solves

`OrderService.checkout` charges a card and saves an order. A unit test must not charge real cards or need a real database — slow, flaky and expensive. So the service **receives its collaborators** through the constructor (**dependency injection**: an object is handed its dependencies instead of creating them; the *Dependency Inversion Principle*, LLD §4.10), and tests pass in stand-ins called **test doubles**.

```
public final class OrderService {
    private static final Logger LOG = Logger.getLogger(OrderService.class.getName());
    private final PaymentGateway gateway;          // interface: the card processor
    private final OrderRepository repository;      // interface: storage
    private final Clock clock;                     // time is a dependency too
    private final Supplier<String> ids;            // so is id generation

    public OrderService(PaymentGateway gateway, OrderRepository repository, Clock clock, Supplier<String> ids) {
... }

    public Order checkout(String customer, Cart cart, String discountCode) {
        if (cart.isEmpty()) throw new IllegalStateException("cart is empty");
        BigDecimal subtotal = cart.subtotal();
        int percent = Discounts.percentFor(discountCode, subtotal);
        BigDecimal total = Discounts.apply(subtotal, percent);

        PaymentGateway.PaymentResult result = gateway.charge(customer, total);
        if (!result.approved()) {
            LOG.warning(() -> "declined: customer=" + customer + " total=" + total);
            throw new PaymentDeclinedException(customer, total);
        }
        Order order = new Order(ids.get(), customer, total, clock.instant());
        repository.save(order);
        LOG.info(() -> "order placed: id=" + order.id() + " total=" + total + " ref=" + result.reference());
        return order;
    }
}
```

Injecting a `java.time.Clock` (Chapter 7 §7.10) and an id `Supplier` makes the result fully predictable: `Clock.fixed(...)` always returns the same instant.

The five kinds of test double

Double	What it is	In this chapter
Dummy	passed only to fill a parameter, never used	—
Stub	returns canned answers	the lambda `(c, amt) -> new PaymentResult(true, "ok")`
Fake	a real, working but simplified implementation	`InMemoryOrderRepository` (a `LinkedHashMap`)
Spy	a real object that also records calls	`spy(new InMemoryOrderRepository())`
Mock	programmed with expectations; you **verify** how it was called	`@Mock PaymentGateway gateway`

Rule of thumb: check **state** (the result, the fake's contents) when you can; **verify interactions** only when the interaction *is* the behaviour — "the card was charged the discounted amount", "nothing was saved after a decline". A longer discussion of doubles and test strategy is in LLD — Testing strategies.

Mockito in practice — `OrderServiceTest.java`

```java
@ExtendWith(MockitoExtension.class)                       // creates @Mock fields, checks for unused stubs
class OrderServiceTest {
    static final Instant NOW = Instant.parse("2025-09-16T10:00:00Z");

    @Mock PaymentGateway gateway;
    @Mock OrderRepository repository;
    @Captor ArgumentCaptor<Order> saved;
    OrderService service;
    Cart cart;

    @BeforeEach
    void setUp() {
        service = new OrderService(gateway, repository, Clock.fixed(NOW, ZoneOffset.UTC), () -> "ORD-1");
        cart = new Cart();
        cart.add(new Item("BK-1", "Java book", new BigDecimal("45.00"), 2));
        cart.add(new Item("MUG", "Mug", new BigDecimal("12.50"), 1));    // subtotal 102.50
    }

    @Test
    void approvedPaymentSavesTheOrder() {
        when(gateway.charge("ann", new BigDecimal("76.88"))).thenReturn(new PaymentResult(true, "PAY-9"));  // arrange

        Order order = service.checkout("ann", cart, "SAVE25");                                              // act

        assertEquals(new Order("ORD-1", "ann", new BigDecimal("76.88"), NOW), order);                       // assert
        verify(repository).save(saved.capture());         // called exactly once, capture the argument
        assertEquals("ORD-1", saved.getValue().id());
    }

    @Test
    void chargesBeforeSaving() {
        when(gateway.charge(anyString(), any())).thenReturn(new PaymentResult(true, "PAY-9"));
        service.checkout("ann", cart, null);
        InOrder inOrder = inOrder(gateway, repository);
        inOrder.verify(gateway).charge(eq("ann"), eq(new BigDecimal("102.50")));  // mixing matchers: all must be matchers
        inOrder.verify(repository).save(any(Order.class));
    }

    @Test
    void declinedPaymentThrowsAndSavesNothing() {
        when(gateway.charge(anyString(), any())).thenReturn(new PaymentResult(false, null));
        var e = assertThrows(PaymentDeclinedException.class, () -> service.checkout("bob", cart, null));
        assertEquals("payment of 102.50 declined for bob", e.getMessage());
        verifyNoInteractions(repository);
    }

    @Test
    void gatewayOutagePropagates() {
        when(gateway.charge(anyString(), any())).thenThrow(new IllegalStateException("gateway timeout"));
        assertThrows(IllegalStateException.class, () -> service.checkout("ann", cart, null));
        verify(repository, never()).save(any());
    }

    @Test
    void fakeRepositoryPlusSpy() {
        InMemoryOrderRepository fake = spy(new InMemoryOrderRepository());
        var svc = new OrderService((c, amt) -> new PaymentResult(true, "ok"),   // a stub written as a lambda
                fake, Clock.fixed(NOW, ZoneOffset.UTC), () -> "ORD-7");
        svc.checkout("cy", cart, "SAVE10");
        assertEquals(new BigDecimal("92.25"), fake.findById("ORD-7").orElseThrow().total());
        verify(fake, times(1)).save(any());
    }
    // ... emptyCartNeverReachesTheGateway()
}
```

All six passed (tree in 18.4). How Mockito works, step by step:

1. `@Mock` (or `mock(PaymentGateway.class)`) creates a generated subclass/**instrumented** class (bytecode rewritten at run time) whose methods record each call and return **defaults**: `null`, `0`, `false`, empty collections (an unstubbed `List<String> find()` returns an empty list), `Optional.empty()`.
2. `when(mock.method(args)).thenReturn(x)` **stubs** a call: when those arguments arrive, return `x`. Also `thenThrow(...)`, `thenAnswer(inv -> ...)`, and several values in a row (`thenReturn(a, b)`).
3. The code under test runs.
4. `verify(mock).method(args)` checks a call happened (exactly once by default). Variants: `times(n)`, `never()`, `atLeastOnce()`, `atMost(n)`, `verifyNoInteractions(...)`, `verifyNoMoreInteractions(...)`; `inOrder(...)` checks ordering.
5. **Argument matchers** — `any()`, `anyString()`, `eq(x)`, `argThat(pred)` — loosen matching. Rule: if one argument uses a matcher, *all* must (hence `eq("ann")` above); `verify(gateway).charge("ann", any())` throws `InvalidUseOfMatchersException`.
6. An `ArgumentCaptor` grabs what was passed so you can assert on it afterwards.

More Mockito you'll meet

Need	API
Stub a `void` method to throw	`doThrow(new X()).when(mock).voidMethod()`
Stub a spy without calling the real method	`doReturn(x).when(spy).method()`
Given/when/then wording (BDD — *behaviour-driven development* — style)	`BDDMockito.given(...).willReturn(...)`, `then(mock).should().method()`
Mock a `static` method	`try (MockedStatic<Util> m = mockStatic(Util.class)) { m.when(Util::now).thenReturn(...); }`
Allow an unused stub	`lenient().when(...)` or `@MockitoSettings(strictness = LENIENT)`
Mock without the extension	`mock(X.class)` or `MockitoAnnotations.openMocks(this)`

Mockito 5 notes. Mockito 5 needs Java 11+, and its default *inline* mock maker instruments classes at run time, so it can mock `final` classes, records and static methods. It does that through a Java agent. On JDK 21+ loading an agent *dynamically* (into an already-running JVM) prints a warning (JEP 451), and a future JDK will block it by default — without `-javaagent` the first run of this chapter printed:

```
Mockito is currently self-attaching to enable the inline-mock-maker. This will no longer work in future releases
of the JDK. Please add Mockito as an agent to your build as described in Mockito's documentation: ...
WARNING: A Java agent has been loaded dynamically (.../byte-buddy-agent-1.18.14.jar)
```

Hence `-javaagent:mockito-core.jar` in `run_tests.sh`, Surefire's `argLine`, and Gradle's `jvmArgs`. (`-Xshare:off` only silences a harmless class-data-sharing notice caused by the agent — Chapter 16 §16.2.)

When not to mock

- **Value objects** (defined only by their data) **and records** (`Item`, `Order`, `BigDecimal`) — just create real ones.
- **Types you don't own** (e.g. a JDBC `ResultSet`, an HTTP client): wrap them behind your own interface and mock that, or use an **integration test** (against real infrastructure such as a database, §18.9).
- **Everything** — tests that mock every collaborator end up restating the implementation ("calls A then B then C") and break on every refactor while catching few bugs. Prefer fakes for storage-like collaborators.

Common mistakes

- Stubbing with the wrong arguments — the stub never matches and the mock returns `null`. **Strict stubs** (the `MockitoExtension` default) turn that into an `UnnecessaryStubbingException`, as 18.8 shows.
- Verifying calls that aren't part of the behaviour (logging, getters).
- `when(spy.method())` on a spy calls the real method during stubbing — use `doReturn`.

- Mixing raw values and matchers in one call.

In the interview

> **30-second answer:** "A test double replaces a real collaborator: dummy, stub, fake, spy or mock. Mockito creates mocks, stubs them with `when(...).thenReturn(...)`, and verifies interactions with `verify`, argument matchers and captors. I inject collaborators through the constructor so tests can swap them, use fakes for things like repositories, and mock only true boundaries like payment gateways — checking state first and interactions only when they are the behaviour."

- *Mock vs stub?* → A stub provides answers; a mock is also verified for how it was called.
- *Mock vs spy?* → A mock is a hollow object with defaults; a spy wraps a real object and calls real methods unless stubbed.
- *How do you test code that uses `LocalDateTime.now()`?* → Inject a `Clock`, pass `Clock.fixed(...)` in tests.

TL;DR

Design for injection, then choose the lightest double: real object → fake → stub → mock. Mockito: `@Mock`, `when/thenReturn`, `verify`, matchers, captors; run it with `-javaagent`.

18.7 AssertJ fluent assertions

What it is

AssertJ is an assertion library: one entry point, `assertThat(actual)`, returns a type-specific object whose methods you chain — auto-complete shows only assertions that make sense for that type. Failure messages are detailed (see 18.8). It works with any test framework; Spring Boot's test **starter** (one dependency bundling related libraries) includes it. (Hamcrest — `assertThat(x, is(y))` — is the older matcher-style alternative.)

Code — `AssertjStyleTest.java`

```java
import static org.assertj.core.api.Assertions.*;

@Test
void collections() {
    assertThat(sampleCart().items())
            .hasSize(2)
            .extracting(Item::sku)
            .containsExactly("BK-1", "MUG");
}

@Test
void bigDecimalsCompareByValue() {
    assertThat(sampleCart().subtotal())
            .isEqualByComparingTo("102.5")              // ignores scale: 102.5 == 102.50
            .isGreaterThan(BigDecimal.ZERO);
}

@Test
void exceptions() {
    assertThatThrownBy(() -> new Item("X", "x", BigDecimal.ONE, 0))
            .isInstanceOf(IllegalArgumentException.class)
            .hasMessage("quantity must be positive: 0");
    assertThatIllegalArgumentException().isThrownBy(() -> Discounts.apply(BigDecimal.ONE, 200));
}

@Test
void recursiveComparisonIgnoringAField() {
    Order actual = new Order("ORD-42", "ann", new BigDecimal("9.99"), Instant.now());
    Order expected = new Order("ignored", "ann", new BigDecimal("9.99"), Instant.EPOCH);
    assertThat(actual).usingRecursiveComparison().ignoringFields("id", "placedAt").isEqualTo(expected);
}

@Test
void softAssertionsCollectAllFailures() {
    Cart cart = sampleCart();
    SoftAssertions.assertSoftly(soft -> {
        soft.assertThat(cart.count()).isEqualTo(3);
        soft.assertThat(cart.isEmpty()).isFalse();
        soft.assertThat(cart.items()).allSatisfy(i -> assertThat(i.quantity()).isPositive());
    });
}

@Test
void writesCsvIntoATemporaryDirectory(@TempDir Path dir) throws IOException {   // deleted after the test
    Order o = new Order("ORD-1", "ann", new BigDecimal("76.88"), Instant.parse("2025-09-16T10:00:00Z"));
    Path csv = OrderExporter.writeCsv(List.of(o), dir);
    assertThat(csv).exists().hasFileName("orders.csv");
    assertThat(Files.readAllLines(csv)).containsExactly(
            "id,customer,total,placed_at",
            "ORD-1,ann,76.88,2025-09-16T10:00:00Z");
}
```

All six passed (tree in 18.4). The last test uses JUnit's `@TempDir`: a real, empty directory per test, deleted afterwards — the right way to test file code (Chapter 13) without touching your disk.

JUnit assertions vs AssertJ

Check	JUnit Assertions	AssertJ
equality	`assertEquals(3, n)`	`assertThat(n).isEqualTo(3)`
list order and content	`assertEquals(List.of(...), list)`	`assertThat(list).containsExactly(...)` / `containsExactlyInAnyOrder` / `contains`
property of each element	loop	`extracting(Item::sku)`, `filteredOn(...)`, `allSatisfy(...)`
`BigDecimal` by value	`assertEquals(0, a.compareTo(b))`	`isEqualByComparingTo("102.5")`
exception and message	`assertThrows` then `assertEquals(msg, e.getMessage())`	`assertThatThrownBy(...).isInstanceOf(...).hasMessage(...)`
many checks, one report	`assertAll(...)`	`SoftAssertions.assertSoftly(...)`
object graph field by field	manual	`usingRecursiveComparison()`
strings	`assertTrue(s.startsWith("x"))`	`assertThat(s).startsWith("x").contains("y")`
files	manual	`assertThat(path).exists().hasContent(...)`
description	message argument	`.as("subtotal of %s", cart)` before the assertion

Common mistakes

- `assertThat(x);` with nothing chained — asserts nothing. (**Static analysis** tools — which inspect source for bugs without running it — and AssertJ's own `@CheckReturnValue` catch this.)
- Calling `.as("…")` **after** the assertion — too late; it must come first.
- `isEqualTo` on `BigDecimal` when scale doesn't matter — use `isEqualByComparingTo`.

TL;DR

`assertThat(actual).<fluent checks>` — readable, type-aware, great messages. Especially strong for collections, exceptions, `BigDecimal`, and object graphs.

18.8 Reading a failing test report

A test suite is only useful if its failures are quick to understand. `FailingDemoTest` fails **on purpose** — it carries `@Tag("demo-failure")`, which every normal run excludes. `bash code/run_tests.sh fail` runs only it:

```java
@Tag("demo-failure")
@ExtendWith(MockitoExtension.class)
class FailingDemoTest {
    @Mock PaymentGateway gateway;

    @Test
    void bigDecimalEqualsChecksScale() {
        // BigDecimal.equals compares value AND scale: 102.5 is not equal to 102.50
        assertEquals(new BigDecimal("102.5"), AssertjStyleTest.sampleCart().subtotal());
    }

    @Test
    void assertjListDiff() {
        List<String> skus = AssertjStyleTest.sampleCart().items().stream().map(Item::sku).toList();
        Assertions.assertThat(skus).containsExactly("MUG", "BK-1", "PEN");
    }

    @Test
    void groupedAssertionsReportEveryFailure() {
        Cart cart = AssertjStyleTest.sampleCart();
        assertAll("cart",
                () -> assertEquals(2, cart.count(), "count"),
                () -> assertEquals(2, cart.items().size(), "lines"),
                () -> assertEquals("100.00", cart.subtotal().toPlainString(), "subtotal"));
    }

    @Test
    void unusedStubIsReported() {
        // Strict stubs: a stubbing the code under test never uses is a smell (wrong argument? dead setup?)
        Mockito.when(gateway.charge("ann", BigDecimal.TEN)).thenReturn(new PaymentGateway.PaymentResult(true,
"x"));
        Assertions.assertThat(AssertjStyleTest.sampleCart().count()).isEqualTo(3);
    }
}
```

Real output (exit code 1; stack frames trimmed where marked):

```
|
├─ JUnit Platform Suite ✔
├─ JUnit Jupiter ✔
|  └─ FailingDemoTest ✔
|     ├─ bigDecimalEqualsChecksScale() ✘ expected: <102.5> but was: <102.50>
|     ├─ groupedAssertionsReportEveryFailure() ✘ cart (2 failures)
|     |         org.opentest4j.AssertionFailedError: count ==> expected: <2> but was: <3>
|     |         org.opentest4j.AssertionFailedError: subtotal ==> expected: <100.00> but was: <102.50>
|     ├─ unusedStubIsReported() ✘
|     |        Unnecessary stubbings detected.
|     |        Clean & maintainable test code requires zero unnecessary code.
|     |        Following stubbings are unnecessary (click to navigate to relevant line of code):
|     |          1. -> at ch18.shop.FailingDemoTest.unusedStubIsReported(FailingDemoTest.java:47)
|     |        Please remove unnecessary stubbings or use 'lenient' strictness. More info: javadoc for
UnnecessaryStubbingException class.
|     └─ assertjListDiff() ✘
|              Expecting actual:
|                ["BK-1", "MUG"]
|              to contain exactly (and in same order):
|                ["MUG", "BK-1", "PEN"]
|              but could not find the following elements:
|                ["PEN"]
└─ JUnit Vintage ✔

Failures (4):
  JUnit Jupiter:FailingDemoTest:bigDecimalEqualsChecksScale()
    MethodSource [className = 'ch18.shop.FailingDemoTest', methodName = 'bigDecimalEqualsChecksScale',
methodParameterTypes = '']
    => org.opentest4j.AssertionFailedError: expected: <102.5> but was: <102.50>
       org.junit.jupiter.api.Assertions.assertEquals(Assertions.java:1199)
       ch18.shop.FailingDemoTest.bigDecimalEqualsChecksScale(FailingDemoTest.java:26)
  ... (the other three repeat the messages above, each with its stack frames)

Test run finished after 67 ms
...
[         4 tests found           ]
[         0 tests successful      ]
[         4 tests failed          ]
```

How to read each one:

1. **`expected: <102.5> but was: <102.50>`** — looks equal, isn't: `BigDecimal.equals` compares value *and scale* (Chapter 7 §7.7). The last frame from *your* code (`FailingDemoTest.java:26`) is the line to open.
2. **`cart (2 failures)`** — `assertAll` ran all three checks and reports the two that failed (`count`, `subtotal`), each with its message label; the `lines` check passed.
3. **Unnecessary stubbings** — the test stubbed `charge("ann", 10)` but never called it. In real tests this usually means the stub's arguments don't match what the code passes, so the code got `null` instead. Mockito tells you the exact line.
4. **AssertJ's diff** prints actual, expected and exactly what's missing (and, when relevant, what's unexpected or out of order).

Build tools report the same information: Surefire writes XML/text reports under `target/surefire-reports/`, Gradle an HTML report at `build/reports/tests/test/index.html`; CI servers read the XML.

> **TL;DR**
>
> Read the first line (expected vs actual), then the first stack frame in your own code. Use `assertAll`/soft assertions and AssertJ to get complete, precise messages.

18.9 Test design

What makes a good unit test — FIRST

Letter	Property	How the tests here achieve it
Fast	milliseconds, so you run them constantly	no network, no DB, no sleeps: 47 tests in ~0.6 s
Independent	any order, any subset	fresh instance per test, `@BeforeEach` setup
Repeatable	same result everywhere, every time	fixed `Clock`, fixed ids, `@TempDir`
Self-validating	pass/fail without reading output	assertions, not `println`
Timely	written with (or before) the code	—

Structure: Arrange–Act–Assert

Every test has three blocks — **arrange** (inputs, doubles), **act** (one call to the code under test), **assert** (check outcome) — the same idea as BDD's *given–when–then*. `approvedPaymentSavesTheOrder` above is commented that way. One *behaviour* per test; several assertions about that one behaviour are fine.

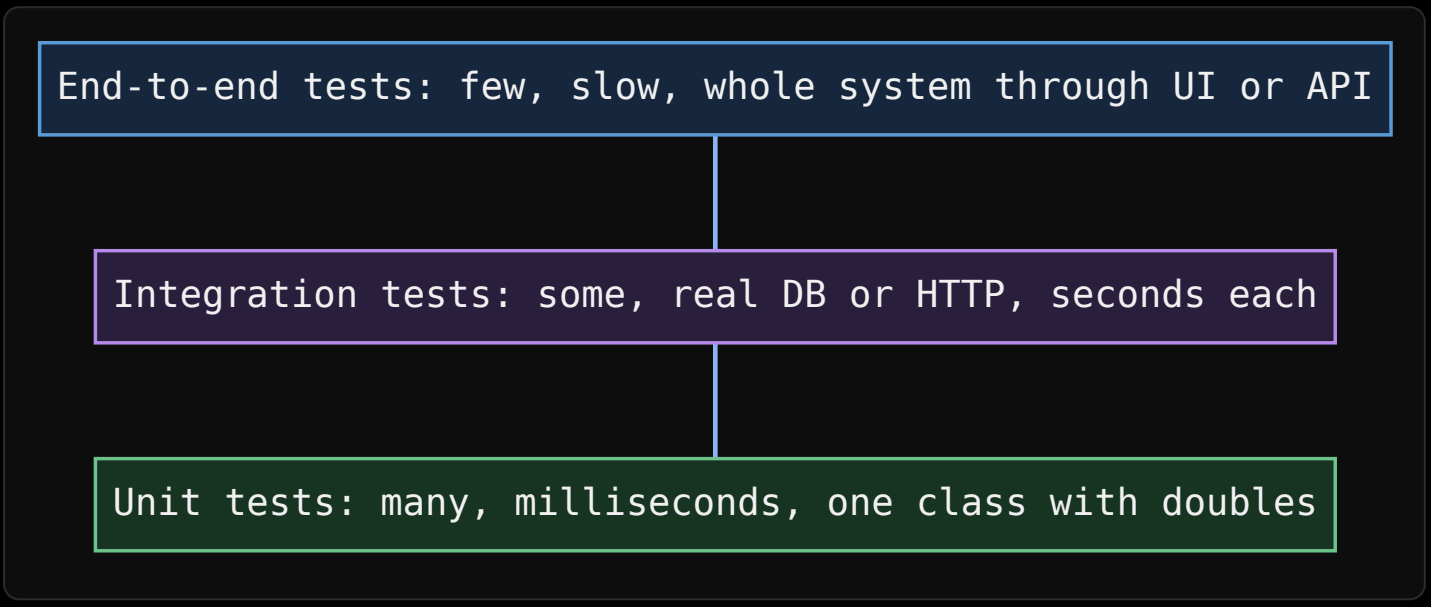
Naming

Name tests after behaviour, not methods: `declinedPaymentThrowsAndSavesNothing` beats `testCheckout2`. A failing test's name should already tell you what broke. `@DisplayName` allows full sentences.

Choosing inputs

- **Equivalence classes**: one representative per group that behaves the same (a valid code, a lowercase code, a too-short code).
- **Boundaries**: just below, at, just above each limit (`99.99`, `100.00` for "SAVE25 needs at least 100"; `-1`, `0`, `100`, `101` for percentages).
- **Special values**: `null`, empty, zero, negative, very large, duplicates, non-ASCII Unicode (accents, emoji).
- **Errors**: each exception path, including exceptions thrown by collaborators (`gatewayOutagePropagates`).

The test pyramid



Most tests should be fast unit tests; fewer **integration tests** exercise real infrastructure (a real database, real HTTP); a handful of **end-to-end** tests cover whole user journeys. **Testcontainers** (2.x at the time of writing) starts real databases/**brokers** (Kafka, RabbitMQ) in **Docker** containers for integration tests — the modern replacement for in-memory lookalikes such as H2, whose SQL **dialect**

differs from production (Testcontainers gives the same engine and dialect). Maven runs integration tests with the **Failsafe** plugin (`*IT.java`, `verify` phase); Gradle typically uses a separate test **source set** (a named group of source folders with its own compile task) or the JVM Test Suite plugin.

Coverage — useful, but not a goal

Code coverage measures which lines/**branches** (each path through an `if` / `switch`) the tests executed; **JaCoCo** is the standard Java tool (Maven plugin `jacoco-maven-plugin`, Gradle plugin `jacoco`), producing an HTML report. Low coverage reliably shows *untested* code; high coverage does **not** prove the tests check anything (90% line coverage only means 90% of lines ran; a test without assertions still "covers" lines). **Mutation testing** (PIT/pitest) is stronger: it makes small changes to your code (flip `>=` to `>`) and checks that some test fails — our `99.99` / `100.00` rows would kill that **mutant** (the altered copy).

Flaky tests

A **flaky** test sometimes passes and sometimes fails without code changes. Common causes and cures:

Cause	Cure
Current time, time zones	inject `Clock`; fixed zone
Randomness	inject `Random` with a seed, or a `Supplier`
Test order / shared static state	fresh state per test; no static mutables
`Thread.sleep` to "wait" for async work	wait for a condition (Awaitility library, latches — Chapter 14)
Real network or external services	doubles in unit tests; Testcontainers/WireMock (a fake HTTP server) in integration tests
Hash-order assumptions (`HashMap` iteration)	assert with `containsExactlyInAnyOrder` or sort first

TDD in one paragraph

Test-driven development repeats a short cycle: write a failing test (**red**), write the simplest code that passes (**green**), clean up (**refactor**) with the tests as a safety net. It keeps designs small and testable; many teams use it selectively (bug fixes: first write the test that reproduces the bug). See also LLD — refactoring with tests as a safety net.

In the interview

> **30-second answer:** "Good unit tests are fast, independent, repeatable and self-checking. I structure them arrange–act–assert, name them after behaviour, and pick inputs by equivalence classes and boundaries. I keep most tests at unit level with injected doubles, add integration tests against real infrastructure with Testcontainers, and treat coverage as a way to find untested code, not a target."

- *How do you test time-dependent code?* → Inject a `Clock`.
- *Is 100% coverage worth it?* → No; diminishing returns and it can't show assertion quality — mutation testing is a better signal.
- *What's a flaky test and how do you fix one?* → Nondeterministic outcome; remove the hidden input (time, randomness, order, network, sleeps).

TL;DR

FIRST, AAA, behaviour names, boundary inputs, mostly unit tests, real infrastructure for integration tests, no hidden time/random/order dependencies.

18.10 Logging

What it is

Logging records what a running program does, each record with a **level**, a **logger name** (usually the class) and a timestamp, sent to destinations (console, file, a central **log collector**) chosen by **configuration**, not code. `System.out.println` has no levels, no way to turn it down in production, no timestamps and no routing.

How a logging system works

1. Code asks for a **logger** named after its class: `Logger.getLogger(OrderService.class.getName())`.
2. Loggers form a **hierarchy** by name: `ch18.shop.OrderService` → `ch18.shop` → `ch18` → root. A logger without its own level **inherits** the nearest ancestor's.
3. Each call has a **level**; if it is below the logger's effective level, the call is discarded cheaply.
4. Accepted records go to **handlers/appenders** (console, file…), which have their own level and a **formatter/layout**.

Levels across the libraries

`java.util.logging`	SLF4J / Logback / Log4j 2	`System.Logger`	Use for
`SEVERE`	`ERROR`	`ERROR`	failures needing attention
`WARNING`	`WARN`	`WARNING`	unexpected but handled
`INFO`	`INFO`	`INFO`	important business/lifecycle events
`CONFIG`	—	—	configuration details (JUL only)
`FINE`	`DEBUG`	`DEBUG`	diagnostic detail
`FINER`, `FINEST`	`TRACE`	`TRACE`	very fine detail

`java.util.logging` (JUL), run — `code/ch18/logging/JulDemo.java`

JUL is built into the JDK (module `java.logging`), so it needs no dependency. By default the root logger is at `INFO` and a `ConsoleHandler` writes to **`System.err`** (settings in `$JAVA_HOME/conf/logging.properties`).

```java
// What you would normally put in a logging.properties file (-Djava.util.logging.config.file=...)
static final String CONFIG = """
        handlers = java.util.logging.ConsoleHandler
        .level = INFO
        java.util.logging.ConsoleHandler.level = ALL
        java.util.logging.SimpleFormatter.format = %4$-7s [%3$s] %5$s%6$s%n
        ch18.logging.level = FINE
        """;

public static void main(String[] args) throws IOException {
    System.setErr(System.out);                      // ConsoleHandler writes to System.err; merge for the demo
    LogManager.getLogManager().readConfiguration(
            new ByteArrayInputStream(CONFIG.getBytes(StandardCharsets.UTF_8)));
    // Normally a static final field; created AFTER the configuration here because the demo configures in code.
    Logger log = Logger.getLogger(JulDemo.class.getName());                    // "ch18.logging.JulDemo"

    System.out.println("-- 1. levels (ch18.logging is set to FINE, the root to INFO)");
    log.severe("disk full");
    log.warning("slow response: 1200 ms");
    log.info("server started on port 8080");
    log.config("pool size = 10");
    log.fine("cache miss for key 42");
    log.finer("not shown: FINER is below FINE");

    System.out.println("-- 2. another package inherits the root level INFO");
    Logger other = Logger.getLogger("com.acme.billing");
    other.info("billing ready");
    other.fine("not shown: com.acme.billing is at INFO");

    System.out.println("-- 3. parameters and lazy messages");
    log.log(Level.INFO, "user {0} bought {1} items", new Object[]{"ann", 3});   // MessageFormat style
    int[] calls = {0};
    Logger quiet = Logger.getLogger("com.acme.quiet");
    quiet.fine(() -> "expensive " + (++calls[0]));   // Supplier: not evaluated, FINE is disabled here
    quiet.fine("eager " + (++calls[0]));               // string built anyway, then thrown away
    System.out.println("message expressions evaluated: " + calls[0]);
    System.out.println("isLoggable(FINE) = " + quiet.isLoggable(Level.FINE));

    System.out.println("-- 4. logging an exception (message + stack trace)");
    try { Integer.parseInt("12x"); }
    catch (NumberFormatException e) { log.log(Level.WARNING, "bad quantity in request", e); }

    System.out.println("-- 5. a custom handler with its own format");
    Logger audit = Logger.getLogger("ch18.audit");
    audit.setUseParentHandlers(false);               // don't ALSO send to the root's ConsoleHandler
    Handler h = new StreamHandler(System.out, new java.util.logging.Formatter() {
        @Override public String format(LogRecord r) {
            return "AUDIT|" + r.getLevel() + "|" + r.getLoggerName() + "|" + formatMessage(r) + "\n";
        }
    });
    audit.addHandler(h);
    audit.info("ann changed her e-mail");
    h.flush();

    System.out.println("-- 6. System.Logger (Java 9): a facade that routes to JUL by default");
    System.Logger sys = System.getLogger("ch18.logging.system");
    sys.log(System.Logger.Level.INFO, "hello from System.Logger, {0}", "ann");
    sys.log(System.Logger.Level.DEBUG, "DEBUG maps to FINE -> shown, ch18.logging is FINE");
}
```

Output:

```
-- 1. levels (ch18.logging is set to FINE, the root to INFO)
SEVERE  [ch18.logging.JulDemo] disk full
WARNING [ch18.logging.JulDemo] slow response: 1200 ms
INFO    [ch18.logging.JulDemo] server started on port 8080
CONFIG  [ch18.logging.JulDemo] pool size = 10
FINE    [ch18.logging.JulDemo] cache miss for key 42
-- 2. another package inherits the root level INFO
INFO    [com.acme.billing] billing ready
-- 3. parameters and lazy messages
INFO    [ch18.logging.JulDemo] user ann bought 3 items
message expressions evaluated: 1
isLoggable(FINE) = false
-- 4. logging an exception (message + stack trace)
WARNING [ch18.logging.JulDemo] bad quantity in request
java.lang.NumberFormatException: For input string: "12x"
    at java.base/java.lang.NumberFormatException.forInputString(NumberFormatException.java:67)
    at java.base/java.lang.Integer.parseInt(Integer.java:527)
    at java.base/java.lang.Integer.parseInt(Integer.java:624)
    at ch18.logging.JulDemo.main(JulDemo.java:56)

-- 5. a custom handler with its own format
AUDIT|INFO|ch18.audit|ann changed her e-mail
-- 6. System.Logger (Java 9): a facade that routes to JUL by default
INFO    [ch18.logging.system] hello from System.Logger, ann
FINE    [ch18.logging.system] DEBUG maps to FINE -> shown, ch18.logging is FINE
```

What it shows:

- **Levels and inheritance**: `ch18.logging.level = FINE` applies to every logger below `ch18.logging`; `com.acme.billing` has no setting, so it inherits the root's `INFO`.
- **Two filters**: a record must pass the *logger's* level and then the *handler's* level (`ConsoleHandler.level = ALL` here; its default is `INFO`). JUL's default `ConsoleHandler` writes to `System.err`.
- **Formatter placeholders** in `SimpleFormatter.format` (a `String.format` pattern): `%1$` date/time, `%2$` source, `%3$` logger name, `%4$` level, `%5$` message, `%6$` exception with stack trace.
- **Lazy messages**: the `Supplier` overload skipped the string building; the plain overload built a string that was then discarded ("evaluated: 1"). Use suppliers (or `isLoggable` guards) for expensive messages.
- **Exceptions**: pass the exception as the last argument — you get the message *and* the stack trace (Chapter 8 §8.9).
- **Configuration timing**: levels from a configuration file are applied to loggers as they are created; configure logging at start-up (`-Djava.util.logging.config.file=logging.properties`) before your classes create loggers — that's why this demo, which configures in code, creates its logger afterwards.
- `System.Logger` (Java 9) is a tiny **facade** (a thin front-end API forwarding to whichever implementation is installed) in `java.base`; without other setup it forwards to JUL.

SLF4J + Logback: what most applications use (not run)

Real applications usually code against **SLF4J** (Simple Logging Facade for Java, 2.0.x) — an API only — and pick a **backend** (the library that actually writes the logs) at deployment: **Logback** (Spring Boot's default) or **Log4j 2**. Libraries log through SLF4J so the *application* decides where output goes. **Bridges** (adapter JARs) route JUL, Log4j 1, and Commons Logging calls into SLF4J so one configuration controls everything.

```java
// (SLF4J — not run: needs slf4j-api plus a backend such as logback-classic on the class path)
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.slf4j.MDC;

public final class OrderService {
    private static final Logger log = LoggerFactory.getLogger(OrderService.class);

    public Order checkout(String customer, Cart cart, String code) {
        MDC.put("customer", customer);                                  // added to every line on this thread
        try {
            log.debug("checkout started, {} items", cart.count());      // {} placeholders: no string built if DEBUG is off
            // ...
            log.info("order placed id={} total={}", order.id(), total);
            log.atInfo().addKeyValue("orderId", order.id()).log("order placed");  // SLF4J 2 fluent API: structured field
            return order;
        } catch (RuntimeException e) {
            log.error("checkout failed for {}", customer, e);           // exception LAST, no placeholder for it
            throw e;
        } finally {
            MDC.remove("customer");
        }
    }
}
```

```xml
<!-- src/main/resources/logback.xml (not run) -->
<configuration>
  <appender name="CONSOLE" class="ch.qos.logback.core.ConsoleAppender">
    <encoder><pattern>%d{HH:mm:ss.SSS} %-5level [%thread] %logger{36} %X{customer} - %msg%n</pattern></encoder>
  </appender>
  <logger name="com.example.shop" level="DEBUG"/>
  <root level="INFO"><appender-ref ref="CONSOLE"/></root>
</configuration>
```

- `{}` **placeholders** replace string concatenation and are cheap when the level is off.
- **MDC** (Mapped Diagnostic Context) attaches key/values (request id, user) to every log line of the current thread — `%X{customer}` prints it. (With virtual threads and thread pools, always clear it in `finally`.)
- **Structured logging** (JSON lines with fields) makes logs searchable in tools such as Elasticsearch/OpenSearch, Loki or Splunk.

Common mistakes

- Logging and rethrowing the same exception at every layer — one failure, five stack traces. Log once, where it's handled.
- Logging secrets or personal data (passwords, tokens, card numbers).
- String concatenation in hot debug logs (`log.debug("x=" + expensive())`).
- Several backends on the class path — SLF4J warns and picks one; keep exactly one.
- Treating the logging library as harmless: **Log4Shell** (CVE-2021-44228, December 2021) let attackers run code through message lookups in Log4j 2 versions before 2.15.0 — keep logging libraries patched like any other dependency.

In the interview

> **30-second answer:** "I code against SLF4J with a backend like Logback, one logger per class, levels ERROR to TRACE, parameterized `{}` messages, exceptions passed as the last argument, MDC for request ids, configuration outside the code. The JDK has java.util.logging and the `System.Logger` facade, fine for small tools."

- *Why a facade?* → Libraries shouldn't force a backend; the application chooses one, and bridges unify everything.

- *Why `{}` instead of concatenation?* → No string is built when the level is disabled.

TL;DR

Loggers per class, levels, hierarchy, handlers/appenders, config outside code. JUL is built in; SLF4J + Logback/Log4j 2 is the industry default.

18.11 JDBC basics

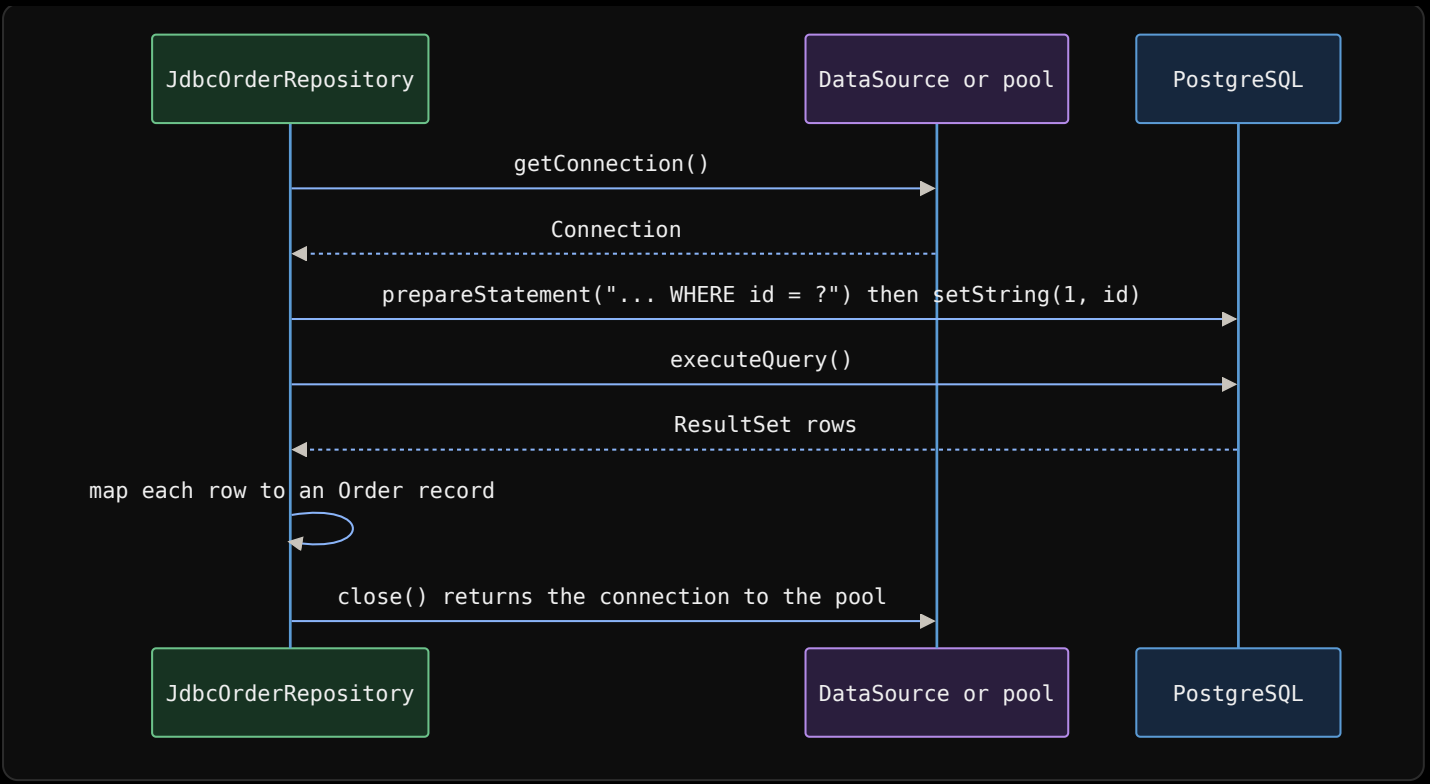
What it is

JDBC (Java Database Connectivity, packages `java.sql` and `javax.sql`) is the standard API for **relational databases** (tables of rows and columns, queried with SQL). Your code uses JDBC **interfaces**; a vendor's **driver** JAR (here PostgreSQL's `postgresql-42.7.13.jar`) implements them. Since JDBC 4.0 (Java 6) drivers register themselves through the **service loader** — the old `Class.forName("org.postgresql.Driver")` line is unnecessary. Every higher-level tool (Spring's `JdbcTemplate`, jOOQ, Hibernate/JPA — the Java Persistence API) is built on JDBC.

The SQL Handbook's Chapter 14 runs a longer JDBC program against the same private PostgreSQL (SQL injection demonstrated, batches, isolation levels); this section focuses on the Java side: resource handling, mapping rows to records, transactions, and hiding JDBC behind the `OrderRepository` interface so the rest of the app — and its tests — don't care.

The pieces

Type	Role
`DriverManager`	gets a `Connection` from a URL — simple, for scripts and tools
`DataSource`	a connection factory configured once — what applications inject; a **connection pool** (reusable open connections — HikariCP, Spring Boot's default) implements it
`Connection`	one session with the database; owns transactions
`PreparedStatement`	SQL with `?` placeholders, values bound by index (from **1**) with typed setters
`ResultSet`	a **cursor** over result rows (one row at a time): `next()`, then `getString("col")`, `getBigDecimal`, `getObject(col, LocalDate.class)`
`SQLException`	**checked** exception; `getSQLState()` gives a standard 5-character code (`23505` = unique violation)


JdbcOrderRepository
DataSource or pool
PostgreSQL
getConnection()
Connection
prepareStatement("... WHERE id = ?") then setString(1, id)
executeQuery()
ResultSet rows
map each row to an Order record
close() returns the connection to the pool
JdbcOrderRepository
DataSource or pool
PostgreSQL

The repository — `code/ch18/jdbc/JdbcOrderRepository.java`

```
public final class JdbcOrderRepository implements OrderRepository {
    private final DataSource dataSource;
    public JdbcOrderRepository(DataSource dataSource) { this.dataSource = dataSource; }

    @Override
    public void save(Order o) {
        String sql = "INSERT INTO orders (id, customer, total, placed_at) VALUES (?, ?, ?, ?)";
        try (Connection c = dataSource.getConnection();
             PreparedStatement ps = c.prepareStatement(sql)) {
            bind(ps, o);
            ps.executeUpdate();
        } catch (SQLException e) {
            throw new RepositoryException("save failed for " + o.id(), e);
        }
    }

    /** All or nothing: one transaction for the whole list. */
    public void saveAll(List<Order> orders) {
        String sql = "INSERT INTO orders (id, customer, total, placed_at) VALUES (?, ?, ?, ?)";
        try (Connection c = dataSource.getConnection()) {
            c.setAutoCommit(false);                       // start a transaction
            try (PreparedStatement ps = c.prepareStatement(sql)) {
                for (Order o : orders) { bind(ps, o); ps.addBatch(); }
                ps.executeBatch();
                c.commit();
            } catch (SQLException e) {
                c.rollback();                             // undo every insert of this call
                throw e;
            }
        } catch (SQLException e) {
            throw new RepositoryException("saveAll failed", e);
        }
    }

    private List<Order> query(String sql, String param) {
        try (Connection c = dataSource.getConnection();
             PreparedStatement ps = c.prepareStatement(sql)) {
            ps.setString(1, param);                       // bound parameter: never concatenate user input
            try (ResultSet rs = ps.executeQuery()) {
                List<Order> out = new ArrayList<>();
                while (rs.next()) out.add(map(rs));
                return out;
            }
        } catch (SQLException e) {
            throw new RepositoryException("query failed: " + sql, e);
        }
    }

    private static Order map(ResultSet rs) throws SQLException {       // row -> record
        return new Order(rs.getString("id"), rs.getString("customer"),
                rs.getBigDecimal("total"), rs.getTimestamp("placed_at").toInstant());
    }
    // ... findById / findByCustomer call query(...); bind(...) sets the four parameters;
    //     RepositoryException is an unchecked wrapper for SQLException
}
```

What to notice:

1. **try-with-resources** closes `ResultSet`, `PreparedStatement`, then `Connection`, even on exceptions (Chapter 8 §8.6). With a pool, `close()` returns the connection instead of disconnecting.
2. **Always `PreparedStatement` with `?`** for values — the driver sends them separately from the SQL text, which prevents **SQL injection** (input crafted to change the SQL command itself) and lets the database reuse the query **plan**.
3. **Transactions**: a connection starts in **auto-commit** mode (each statement is its own committed transaction). For several statements that must succeed together, `setAutoCommit(false)`, then `commit()` or `rollback()` (undo everything since the transaction began).

4. **Exception translation** (Chapter 8 §8.8): the checked `SQLException` is wrapped, with the cause kept, so the `OrderRepository` interface stays storage-neutral.
5. The class names only `java.sql` / `javax.sql` types — the driver is a **runtime** dependency (`runtime` scope / `runtimeOnly`).

Running it — `code/ch18/jdbc/JdbcDemo.java`

The demo creates a database `ch_18_jdbc` on the SQL Handbook's private server (port 5544, reachable only through a **Unix socket** — a local inter-process connection that looks like a file — hence the `junixsocket` socket-factory parameters in the URL; against a normal server the URL is just `jdbc:postgresql://localhost:5432/shop`), creates the table, and uses the repository through a `PGSimpleDataSource`:

```
PGSimpleDataSource ds = new PGSimpleDataSource();          // the driver's simple, non-pooled DataSource
ds.setURL(BASE + "ch_18_jdbc" + SOCKET);
ds.setUser("postgres");
JdbcOrderRepository repo = new JdbcOrderRepository(ds);
repo.save(new Order("ORD-1", "ann", new BigDecimal("76.88"), t));
repo.saveAll(List.of(/* ORD-2 for bob, ORD-3 for ann */));
System.out.println("findById(ORD-1)   = " + repo.findById("ORD-1").orElseThrow());
// ... findById("ORD-404"), findByCustomer("ann"), then saveAll(ORD-4, duplicate ORD-1) inside try/catch,
//     printing the SQLState of the cause, and finally whether ORD-4 exists
```

Commands (from the repository root) and real output:

```
$ bash tools_pg.sh start
postgres ready on port 5544 (socket /tmp)
$ javac --release 25 -cp '.java-books/lib/*' -d /tmp/java18o $(find java/chapters/18-testing-build-and-
tooling/code/ch18 -name '*.java')
$ java --enable-native-access=ALL-UNNAMED -cp "/tmp/java18o:.java-books/lib/*" ch18.jdbc.JdbcDemo
driver: PostgreSQL JDBC Driver 42.7.13
findById(ORD-1)   = Order[id=ORD-1, customer=ann, total=76.88, placedAt=2025-09-16T10:00:00Z]
findById(ORD-404) = Optional.empty
ann's orders      = [ORD-1, ORD-3]
saveAll failed: SQLState 23505 (BatchUpdateException)
ORD-4 present after rollback? false
```

(`--enable-native-access=ALL-UNNAMED` only silences the JDK's warning about the socket library loading native code — without it the program prints four `WARNING: A restricted method…` lines first.) The rollback worked: the valid `ORD-4` was undone together with the failing duplicate. `BatchUpdateException` is the `SQLException` subclass thrown by `executeBatch`.

Java ↔ SQL types

SQL (PostgreSQL)	Java	Getter / setter
`text`, `varchar`	`String`	`getString` / `setString`
`integer`, `bigint`	`int`, `long` (use `wasNull()` for SQL NULL) or `Integer` / `Long` via `getObject(col, Integer.class)`	`getInt`, `getLong`
`numeric`	`BigDecimal` — never `double` for money	`getBigDecimal` / `setBigDecimal`
`boolean`	`boolean`	`getBoolean`
`date`, `time`	`LocalDate`, `LocalTime` (JDBC 4.2)	`getObject(col, LocalDate.class)` / `setObject`
`timestamp`	`LocalDateTime`	`getObject(col, LocalDateTime.class)`
`timestamptz`	`OffsetDateTime` (pgjdbc does not accept `Instant` in `setObject`); `Timestamp.from(instant)` also works, as above	`getObject(col, OffsetDateTime.class)`

Testing database code

- **Unit tests** of services use the fake `InMemoryOrderRepository` — the reason the interface exists.
- **Integration tests** of `JdbcOrderRepository` itself need a real database: Testcontainers starts PostgreSQL in Docker, the test creates the **schema**, runs the repository, and checks rows. Mocking `Connection`/`ResultSet` tests nothing useful.

When not to write raw JDBC

For a handful of queries, raw JDBC is fine. For a real application, use a helper that removes the **boilerplate**: **Spring** `JdbcClient`/`JdbcTemplate` (SQL stays visible, mapping is a lambda), **jOOQ** (type-safe SQL builder), or **JPA/Hibernate** (an **ORM**, *object-relational mapper*: maps objects to tables and generates SQL — powerful, but hides queries; watch for the **N+1 problem**: one query for a list, then one extra query per item). The SQL book compares ORMs vs SQL in its Chapter 14.

Common mistakes

- Concatenating user input into SQL.
- Not closing connections — the pool runs dry and the app hangs.
- `getInt` on a NULL column returns `0`; check `wasNull()` or use `getObject(col, Integer.class)`.
- Using `double` for money columns.
- Forgetting `rollback()` (or `setAutoCommit(true)` before returning a connection to a pool that doesn't reset it).

In the interview

> **30-second answer:** "JDBC is the standard API; a driver implements it. Apps get connections from a pooled `DataSource`, use `PreparedStatement` with bound parameters to avoid SQL injection, map `ResultSet` rows to objects, close everything with try-with-resources, and control transactions with `setAutoCommit(false)`, `commit` and `rollback`. I hide it behind a repository interface and test it against a real database with Testcontainers."

- *`Statement` vs `PreparedStatement`?* → Prepared: parameters bound separately (safe, typed, plan reuse).
- *Why a connection pool?* → Opening a connection costs a network handshake, authentication and a server process; a pool reuses a bounded number of them.

TL;DR

`DataSource` → `Connection` → `PreparedStatement` → `ResultSet`, all in try-with-resources; bind parameters; explicit transactions for multi-statement work; hide JDBC behind an interface; test it against a real DB.

18.12 Packaging: JAR files

What it is

A **JAR** (Java ARchive) is a ZIP file of `.class` files and resources plus a **manifest**, `META-INF/MANIFEST.MF`, holding metadata (e.g. which class to start). A manifest with a `Main-Class` makes it an **executable JAR**, started with `java -jar`.

Building one by hand — `code/package_jar.sh`

This is what `jar:jar` (Maven) and the `jar` task (Gradle) do for you. Real run:

```
$ javac --release 25 -d classes $(find src -name '*.java')
$ jar --create --file shop.jar --main-class ch18.shop.ShopDemo -C classes .
$ jar --list --file shop.jar | head -5
META-INF/
META-INF/MANIFEST.MF
ch18/
ch18/shop/
ch18/shop/Cart.class
$ unzip -p shop.jar META-INF/MANIFEST.MF
Manifest-Version: 1.0
Created-By: 26 (Oracle Corporation)
Main-Class: ch18.shop.ShopDemo

$ java -jar shop.jar
items=3 subtotal=102.50
Order[id=ORD-1, customer=ann, total=76.88, placedAt=2025-09-16T10:00:00Z]
Sep 28, 2026 3:37:08 PM ch18.shop.OrderService checkout
WARNING: declined: customer=bob total=899.00
error: payment of 899.00 declined for bob
$ ls -l shop.jar
13629 bytes
```

- `-C classes .` means "change into `classes/` and add everything" — package folders become ZIP paths.
- `--main-class` (short form `-e`) writes the `Main-Class` manifest line; `Created-By` records the JDK that built it (26 here, compiled for 25).
- The `WARNING` lines come from JUL's *default* format (timestamp, source, level) because no custom configuration was loaded — compare 18.10.
- Short, older option spellings still work: `jar cfe shop.jar ch18.shop.ShopDemo -C classes .`, `jar tf shop.jar`.

Dependencies at run time

`java -jar` **ignores** `-cp` (so `java -cp lib/* -jar app.jar` won't find the libraries); a JAR that needs libraries must say where they are, or contain them:

Approach	How	Notes
Plain JAR + class path	`java -cp "shop.jar:lib/*" ch18.shop.ShopDemo`	simplest; ship a `lib/` folder (Gradle's `application` plugin builds exactly this into `distZip` with start scripts)
`Class-Path` manifest entry	`Class-Path: lib/postgresql-42.7.13.jar`	relative paths, space-separated
Fat / uber JAR	Maven Shade plugin, Gradle Shadow plugin	one file with all classes merged; watch for clashing `META-INF/services` files and signatures; still needs an installed JRE (unlike a `jlink` image)
Spring Boot JAR	`spring-boot-maven-plugin` / Gradle plugin	dependencies kept as nested JARs, loaded by Boot's own launcher
Custom runtime	`jlink` (Java 9)	a trimmed JDK containing only needed modules — see Chapter 16 §16.12
Native installer	`jpackage` (final in Java 16)	`.dmg` / `.pkg`, `.msi` / `.exe`, `.deb` / `.rpm` bundling a runtime
Container image	Dockerfile, Jib, Cloud Native Buildpacks	the common deployment unit for services
Native executable	GraalVM Native Image	**ahead-of-time** compiled to machine code, fast start-up; restrictions on reflection

Other JAR facts: a **multi-release JAR** (Java 9) holds version-specific classes under `META-INF/versions/N/`; a **modular JAR** contains `module-info.class` (Chapter 16, and the OCP modules chapter for exam depth); JARs can be **signed** with `jarsigner`. Libraries are published as a JAR + `.pom` (+ `-sources.jar` and `-javadoc.jar`) to a repository with `mvn deploy` or Gradle's `maven-publish` plugin.

TL;DR

JAR = ZIP + manifest. `jar --create --main-class`, run with `java -jar`. For dependencies: class path folder, fat JAR, Boot JAR, or container; `jlink` / `jpackage` to ship a runtime.

18.13 Other everyday tools

Need	Tool(s)	One-line note
IDE	IntelliJ IDEA, Eclipse, VS Code + Java extensions	import the Maven/Gradle build; don't hand-configure class paths (Chapter 1 §1.11)
Debugger	IDE debugger (**breakpoints** — where execution pauses to inspect variables — conditional breakpoints, watches, step into/over, evaluate expression); `jdb` on the command line	start the JVM with `-agentlib:jdwp=transport=dt_socket,server=y,suspend=n,address=*:5005` to attach remotely
Quick experiments	`jshell` (Java's REPL), single-file launch `java Foo.java`	Chapter 1 §1.9–1.10
Formatting	google-java-format, Palantir Java Format, Spotless (build plugin)	format automatically; stop debating style in reviews
Static analysis	compiler warnings (`-Xlint:all`), Error Prone, SpotBugs, PMD, Checkstyle, SonarQube	catch bugs and **smells** (patterns suggesting a design problem) before tests run
Null-safety	JSpecify annotations + NullAway	JUnit 6 and Spring Boot 4 both adopted JSpecify
Coverage / mutation	JaCoCo / PIT	18.9
Dependency hygiene	`dependency:tree`, `gradle dependencies`, Dependabot/Renovate (bots opening dependency-update pull requests), OWASP Dependency-Check	keep libraries current and free of known vulnerabilities
Documentation	`javadoc` (Markdown doc comments since Java 23 — Chapter 17)	generated HTML API docs
Profiling / monitoring	JFR, JMC, VisualVM, async-profiler	Chapter 16 §16.10
Benchmarks	JMH	Chapter 16 §16.9
Continuous integration	GitHub Actions, GitLab CI, Jenkins	run `./mvnw verify` / `./gradlew build` on every push

A typical CI job is only a few lines — check out the code, set up a JDK, run the wrapper:

```yaml
# .github/workflows/build.yml (GitHub Actions — not run)
name: build
on: [push, pull_request]
jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v5
      - uses: actions/setup-java@v5
        with: { distribution: temurin, java-version: '25', cache: maven }
      - run: ./mvnw -B verify
```

18.14 Where next: Spring Boot and friends

What it is

Most server-side Java uses a **framework** for dependency injection, configuration, web **endpoints** (URLs your service answers), data access, security, **observability** (logs, metrics, traces) and testing. The most widely used is **Spring Boot** (on the Spring Framework). **Spring Boot 4.0** was released on 20 November 2025 on Spring Framework 7 (4.1 followed in 2026); it runs on Java 17+ and supports Java 25.

What Spring Boot gives you

- **Dependency injection container**: you write classes like our `OrderService` with constructor parameters; Spring creates them (*beans*) and passes the right collaborators — exactly the design that made 18.6 testable, without the hand wiring in `ShopDemo`.
- **Starters**: one dependency per capability (`spring-boot-starter-webmvc`, `-jdbc`, `-data-jpa`, `-security`, `-actuator`) with compatible versions managed by Boot's BOM.
- **Auto-configuration**: sees a JDBC driver and a URL in `application.properties` → creates a HikariCP `DataSource`; sees the web starter → starts an **embedded** Tomcat (running inside your application, not as a separate server).
- **Executable JAR**: `./mvnw spring-boot:run` in development, `java -jar app.jar` in production.
- **Testing**: the test starter bundles JUnit Jupiter, Mockito and AssertJ (this chapter's tools); `@SpringBootTest` for full context tests, *slice* tests like `@WebMvcTest` and `@DataJdbcTest`, `@MockitoBean` to replace a bean with a mock, and Testcontainers integration.
- **Actuator**: **health checks**, **metrics** (via Micrometer), info endpoints for operations.

```java
// (Spring Boot — not run: needs the Spring Boot starters; shows the shape of an app)
@SpringBootApplication                              // component scan + auto-configuration
public class ShopApplication {
    public static void main(String[] args) { SpringApplication.run(ShopApplication.class, args); }
}

@RestController
@RequestMapping("/orders")
class OrderController {
    private final OrderService service;             // injected by Spring through the constructor
    OrderController(OrderService service) { this.service = service; }

    @GetMapping("/{id}")
    Order byId(@PathVariable String id) { return service.find(id).orElseThrow(); }
}
```

```properties
# src/main/resources/application.properties (not run)
spring.datasource.url=jdbc:postgresql://localhost:5432/shop
spring.datasource.username=shop
logging.level.com.example.shop=DEBUG
```

Generate a project at **start.spring.io** (Spring Initializr): choose Maven or Gradle, Java 25, and the starters you need.

Alternatives worth knowing

Framework	In one line
Quarkus	Kubernetes-oriented, fast start-up, strong GraalVM native-image support; Jakarta EE / MicroProfile APIs
Micronaut	compile-time dependency injection (no run-time reflection), small **footprint** (memory use)
Jakarta EE (formerly Java EE)	the standard specifications (Servlet, JPA, CDI, REST) implemented by application servers and runtimes such as WildFly, Open Liberty, Payara
Helidon, **Javalin**, **Vert.x**	lighter or **reactive** (non-blocking, event-driven) options for services

A learning path from here

1. Build a small **REST API** with Spring Boot (web + JDBC or JPA + validation), tested with this chapter's tools and Testcontainers.
2. Learn the design side: LLD Handbook — SOLID, patterns, and the case studies (parking lot, booking systems) — all in Java.
3. Learn the data side: SQL Handbook and *Designing Data-Intensive Applications* notes in `ddia/`.
4. Prepare for interviews: Chapter 19 — Java for coding interviews and the OCP Java 21 study repo for certification-level detail.

> **TL;DR**

Spring Boot = DI container + starters + auto-configuration + executable JAR + first-class testing. Everything in this chapter — constructor injection, JUnit/Mockito/AssertJ, SLF4J logging, JDBC/DataSource, Maven/Gradle, JARs — is what Boot builds on.

Chapter TL;DR

- **Layout & coordinates**: `src/main/java`, `src/test/java`; libraries are `group:artifact:version` from Maven Central, cached locally.
- **Maven**: `pom.xml`; scopes; nearest-wins conflicts fixed by `dependencyManagement`/BOMs; lifecycle validate → … → deploy. Pin plugin versions; use `mvnw`.
- **Gradle**: scripted incremental task graph; `implementation` / `api` / `runtimeOnly` / `testImplementation`; `useJUnitPlatform()`; `gradlew`, toolchains.
- **JUnit (5/6 Jupiter)**: `@Test`, `assertEquals(expected, actual)`, `assertThrows`, `assertAll`; new instance per test; `@Nested`, `@Tag`, `@TempDir`; parameterized tests. JUnit 6 = Java 17 baseline, unified versions, same API.
- **Mockito**: inject collaborators, mock boundaries: `when/thenReturn`, `verify`, matchers (all or none), captors, strict stubs; prefer fakes; attach with `-javaagent`.
- **AssertJ**: `assertThat(x).…` fluent, type-aware, precise failure messages; soft assertions; recursive comparison.
- **Test design**: FIRST, arrange–act–assert, behaviour names, boundaries, pyramid, Testcontainers for integration, coverage ≠ quality, remove time/random/order/network nondeterminism.
- **Logging**: logger per class, levels, hierarchy, handlers; JUL built in; SLF4J + Logback/Log4j 2 in apps; `{}` placeholders, exception last, MDC, no secrets.
- **JDBC**: `DataSource` → `Connection` → `PreparedStatement` → `ResultSet` in try-with-resources; bind parameters; explicit transactions; wrap behind a repository.
- **Packaging**: JAR = ZIP + manifest; `jar --create --main-class`, `java -jar`; fat JARs, Boot JARs, `jlink`, `jpackage`, containers.
- **Next**: Spring Boot builds on all of the above.

Next: Chapter 19 — Java for coding interviews.