LLD Handbook Study Edition — Vol 3: Case Studies II, Machine Coding, Revision and Practice

OOP · UML · SOLID · Clean Code · Refactoring · Patterns · Concurrency · 45 case studies in Java — every topic, readable length

6 parts · one place to study LLD: DesignGurus courses + classic books + GitHub problem sets, combined

what it is | problem | how it works | trade-offs | real world | interview | self-check | TL;DR

Personal study notes. Course diagrams belong to DesignGurus. Not for redistribution.

Contents

1. **Chapter 19 — Case Studies: Commerce & Marketplaces (Part 2 of Platforms) — Study Edition**
   - Design Amazon — Online Shopping System
   - Design an Online Stock Brokerage System
   - Design an Online Auction System
   - Design a Music Streaming Service (Spotify)
   - Design an Online Food Delivery Service (Swiggy) + FoodKart
   - Design a Grocery Store (POS & Checkout)
   - Chapter summary
2. **Chapter 20 — Case Studies: Pattern-Driven Components (Study Edition)**
   - 20.1 Design Splitwise
   - 20.2 Design a Notification Service
   - 20.3 Design an In-Memory Cache (and an LRU Cache from scratch)
   - 20.4 Design a Ride Sharing Service (Uber)
   - 20.5 Design a Logging Framework
   - 20.6 Design a Pub-Sub System
   - 20.7 Design a Rate Limiter
   - 20.8 Chapter summary
   - 20.9 Chapter self-check
   - 20.10 Chapter TL;DR
3. **Chapter 21 — Case Studies: Machine-Coding Round Problems (Study Edition)**
   - 21.0 How the machine-coding round works
   - 21.1 Design a Task Management System
   - 21.2 Design a Digital Wallet Service
   - 21.3 Design an Event Calendar (Flipkart)
   - 21.4 Design a LeetCode-like Coding Platform — "Coding Blox" (Flipkart)
   - 21.5 Design a Property Listing Site — "PropertyHunt" (ClearTrip)
   - 21.6 Design a Loan Ledger — "Ledger Co" (Navi)
   - 21.7 Design an Inventory Management System
   - 21.8 Design a Unix File Search (`find`)
   - Chapter summary
4. **Chapter 22 — Revision: Cheat Sheets & Glossary**
   - 22.1 How to revise
   - 22.2 The LLD interview answer framework (one page)
   - 22.3 OOP & relationships cheat sheet
   - 22.4 SOLID one-pager
   - 22.5 Other design principles and Clean Code rules
   - 22.6 Code smells → refactorings (all 24)
   - 22.7 Design patterns: all 23 GoF + extras
   - 22.8 Concurrency cheat sheet
   - 22.9 Case-study map
   - 22.10 Machine-coding and edge-case checklists
   - 22.11 Glossary
   - Self-check
   - TL;DR
5. **Chapter 23 — Interview Question Bank**
   - How to use this bank
   - Part A — "Design X" (LLD and machine coding) questions
   - Everyday objects
   - Booking and reservations
   - Games
   - Platforms and marketplaces
   - Components and infrastructure
   - Machine-coding round favourites
   - Part B — Concept questions
   - OOP basics
   - Relationships and UML
   - SOLID
   - Design principles and clean code
   - Creational patterns
   - Structural patterns
   - Behavioral patterns
   - Concurrency
   - Java specifics for LLD rounds
   - Part C — How to run a 60–90 minute machine-coding round
   - Sources of the question lists
6. **Chapter 24 — Mock Interview Drills (Study Edition)**
   - 24.0 How to use these drills
   - 24.1 Drill — Parking Lot
   - 24.2 Drill — Elevator System
   - 24.3 Drill — Library Management System
   - 24.4 Drill — Amazon Locker Service
   - 24.5 Drill — Vending Machine
   - 24.6 Drill — Online Blackjack Game
   - 24.7 Drill — Ride Sharing (practice session)
   - 24.8 Drill — Meeting Scheduler
   - 24.9 Drill — Movie Ticket Booking System
   - 24.10 Drill — Car Rental System
   - 24.11 Drill — ATM
   - 24.12 Drill — Chess Game
   - 24.13 Drill — Hotel Management System
   - 24.14 Drill — Food Delivery (practice session)
   - 24.15 Drill — Amazon Online Shopping System
   - 24.16 Drill — Stack Overflow
   - 24.17 Drill — Restaurant Management System
   - 24.18 Drill — Facebook
   - 24.19 Drill — Online Stock Brokerage System
   - 24.20 Drill — Jigsaw Puzzle
   - 24.21 Drill — Airline Management System
   - 24.22 Drill — ESPNcricinfo
   - 24.23 Drill — LinkedIn
   - Chapter summary

Chapter 19 — Case Studies: Commerce & Marketplaces (Part 2 of Platforms) — Study Edition

Where this fits. Chapter 18 designed platforms where people *talk*. This chapter designs platforms where people *buy, sell and trade*. Three pressures shape every design: **money must be exact, state must follow rules, and shared things must not be double-sold.** This is the study edition: every topic of the full chapter, without the padding.

#	Problem	The one idea that makes it work	Main patterns
19.1	Amazon — online shopping	An order is an immutable copy of the cart, moved through a **state machine**	State machine (enum), Strategy (payment), Observer, Facade
19.2	Stock brokerage + exchange matching engine	Four order types differ in only **two questions**; a book matches by **price, then time**	Template Method, Observer, Strategy (price policy)
19.3	Online auction	"Read highest bid, compare, replace" is **one atomic step**	Monitor lock, Observer, Specification
19.4	Music streaming (Spotify)	Play order is a **pluggable strategy**; plan rules live **in the plan**	Strategy, Observer, enum-with-behaviour
19.5	Food delivery (Swiggy) + FoodKart	Capacity and riders are **claimed atomically**; selection is a **strategy**	Strategy, Observer, State machine, CAS
19.6	Grocery store (POS & checkout)	A **pure pricing engine**: best offer per product, then a coupon spread over items so tax and refunds stay exact	Strategy (offers, tax, payment), Facade, Observer, enum state machine

Running the code: `cd chapters/19-case-platforms-2/code && javac --release 21 -d /tmp/lld_19 $(find . -name '*.java') && java -cp /tmp/lld_19 lld.ch19.amazon.Main` (or `brokerage`, `stockexchange`, `auction`, `music`, `fooddelivery`, `foodkart`, `grocery`). Java 21, standard library only; the stock-exchange `Main` also runs a small test suite. One file may hold several small package-private types.

Four rules that apply to every design in this chapter

1. **Money is `BigDecimal`, never `double`.** A `double` cannot store 0.1 exactly (`0.1 + 0.2` = `0.30000000000000004`). Build it from a `String` (`new BigDecimal("0.10")`) and compare with `compareTo`, not `equals` — `new BigDecimal("2.0").equals(new BigDecimal("2.00"))` is `false` because the scales differ.
2. **Status changes follow a state machine** — a fixed list of states plus allowed moves, kept *inside the enum*, so an illegal move ("cancel a delivered order") throws at once.
3. **Check-then-act on shared stock must be atomic.** "One left? Take it" is two steps; two threads can both see "yes". Do both under one lock, or with one compare-and-set (CAS: "set to B only if it is still A").
4. **Copy prices into orders.** Later catalog or menu price changes must not change old orders.

Design Amazon — Online Shopping System

Problem & clarifying questions

We design Amazon's **online retail part**: people list products, others find them, fill a cart, pay and track the parcel.

Question	Assumed answer
Buy without an account?	No. Browsing and carts are open; ordering needs a **registered member**
Who sells?	Members can list products (a marketplace)
Payment?	Credit card (clears at once) or bank transfer (clears later)
Cancel?	Yes, until shipped
Notifications?	On every order or shipment status change
Returns, wish lists, payouts, recommendations, multi-warehouse stock?	Out of scope — say so out loud

Requirements

Functional: (1) members add products to sell; (2) anyone searches by **name** or **category**; (3) anyone views, only members buy; (4) add/remove/change cart items; (5) check out the whole cart; (6) rate and review; (7) shipping address; (8) cancel an unshipped order; (9) notified on each order/shipment status change; (10) pay by card or bank transfer; (11) track shipment and order state.

Non-functional: two buyers never both get the last unit; exact `BigDecimal` money; orders and shipments keep a **history** of states; a new payment method or channel must not change checkout.

Use cases

Actors: **Admin** (block/unblock accounts, manage categories), **Guest** (search, cart, register), **Member** (guest + order + sell), **System** (notifications). In the use-case diagram, "Checkout" *includes* "Make payment" and "Add shipping address"; search by name/category *extend* "Search product".

Core entities

Class	Responsibility
`ProductCategory`	Named group; exists without products
`Product`	Price, seller, stock; **owns** reviews; guards stock (`reserve` / `release`)
`ProductReview`	Rating 1–5 + text; stores reviewer **by id** (no `Product ↔ Member` dependency)
`Search` / `Catalog`	Name-word and category indexes, both updated in `addProduct` ; points to, never owns, products
`ShoppingCart` / `Item`	Mutable cart owning its lines; `Item` is an immutable record (product id, name, qty, unit price)
`Order` / `OrderLog`	Immutable **copy** of cart lines + address + status on legal edges; one log entry per change
`Payment` (`CreditCardPayment` , `BankTransferPayment`)	Card clears now; transfer returns `PENDING`
`Shipment` / `ShipmentLog`	Parcel + tracking history; last entry = "where is it now"
Notification channel (`EmailChannel` , `SmsChannel`)	One delivery way (course: abstract `Notification` with subclasses; same idea)
`Account` , `Customer` → `Guest` , `Member` , `Admin`	`Customer` has only what *anyone* may do; ordering lives on `Member`
`OrderService`	Turns a cart into an order; drives the state machine

Class diagram

The course builds it in five stages (products → catalog → payment/shipment/notification → cart/order → accounts). Filled diamonds are composition (`Product` owns reviews, cart owns items, order owns log and

shipment); plain arrows are associations (order only *refers to* its payment); `Guest` and `Member` extend `Customer`.

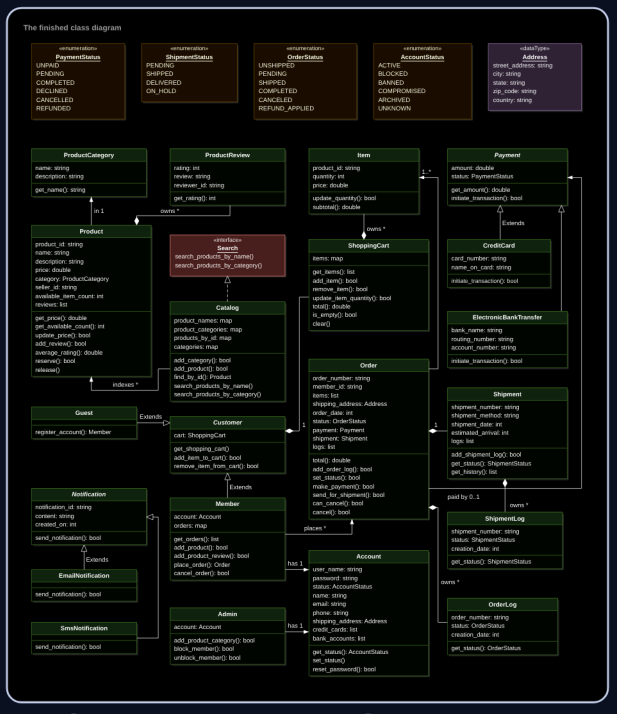
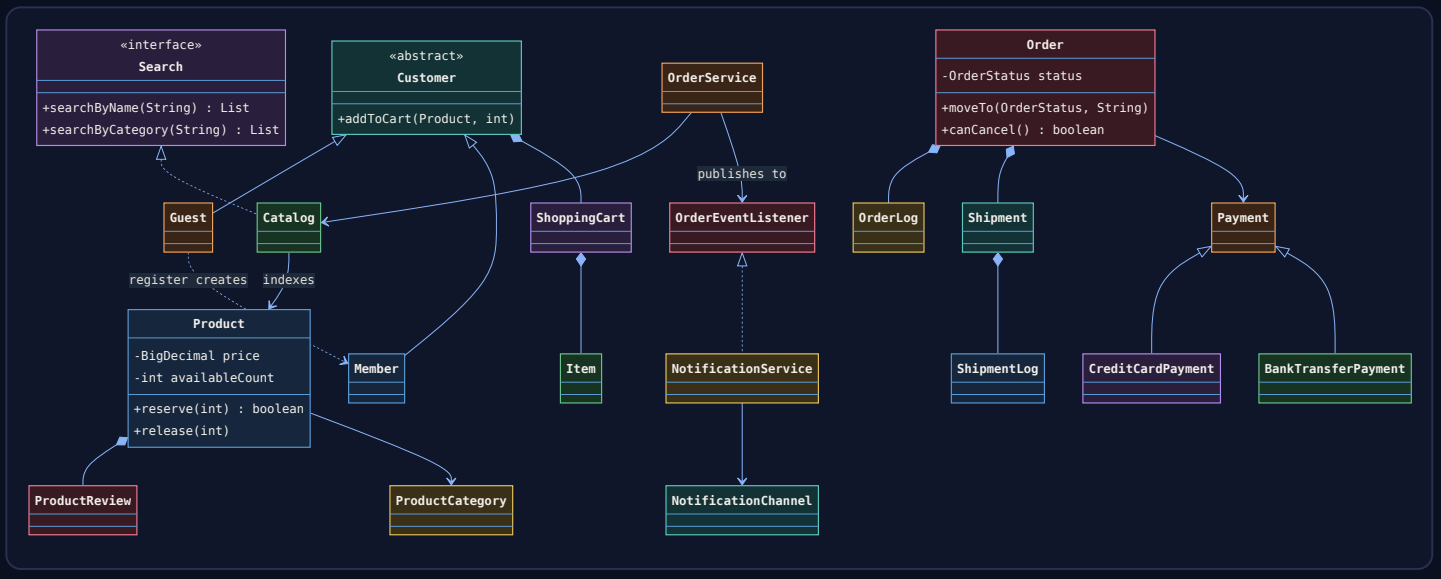


Figure — The finished class diagram of the online shopping system (course diagram)



Key flows

The course's activity diagram is the shopping loop (search, view, add, update cart, "more shopping?", checkout); its sequence diagrams cover search, add-to-cart, and **checkout**: get cart items → billing/shipping info → create order → verify items → pay → create shipment → return status (with error paths).

Checkout (`OrderService.placeOrder`) follows that order, but stock is **reserved** during verification, not just checked: member active and cart not empty → reserve each line (any failure releases what was reserved and fails) → new `Order` copying the lines → `payment.process()` : COMPLETED → `UNSHIPPED` + notify; PENDING (transfer) → notify PENDING; DECLINED → release stock, `CANCELED`.

Order state machine (`OrderStatus.next()`): PENDING → UNSHIPPED (paid) or CANCELED; UNSHIPPED → SHIPPED or CANCELED; SHIPPED → COMPLETED; CANCELED → REFUND_APPLIED. Requirement 8 is one line — `canCancel()` returns `status.canMoveTo(CANCELED)`; there is no `SHIPPED → CANCELED` edge.

Design decisions & patterns used

Decision	Why	Alternative rejected
Ordering on `Member`, not `Customer`	A guest *is a* customer; inheriting `placeOrder` it must reject breaks **LSP**	`Customer.placeOrder()` that throws for guests
`Guest.register()` returns a new `Member` carrying the cart	The cart survives registration	Mutating a guest (Java objects can't change class)
`Order` copies cart lines	A finished purchase must not change	Shared `Item` objects
Stock reserved at order time	No two buyers get the last unit (abandoned orders hold stock → timer release)	Check, pay, then decrement (race window)
State machine inside `OrderStatus`	One owner of the rules	Scattered `if`s
Sealed abstract `Payment`, a subclass per method (**Strategy**)	Methods *behave* differently; new method = new subclass (OCP); that is why it is not an enum	`enum PaymentType` + `switch`
Payment passed as `Function<BigDecimal, Payment>`	Total known only after building the order	Creating payment before the total
Observer for notifications	New channel = new `NotificationChannel`	Email/SMS code inside `OrderService`
`Shipment` keeps a log; reviews store `reviewerId`	History for tracking/disputes; no two-way dependency	Single status field; review holding `Member`
`OrderService` as **Facade**	Hides stock, payment, shipment, notification	Callers orchestrating five objects

Java implementation

Package `lld.ch19.amazon` : `Money.java` , `Catalog.java` , `Order.java` , `Payments.java` , `Accounts.java` , `OrderService.java` , `Main.java` .

```java
// File: lld/ch19/amazon/Order.java (excerpt)
enum OrderStatus {
    PENDING, UNSHIPPED, SHIPPED, COMPLETED, CANCELED, REFUND_APPLIED;

    private Set<OrderStatus> next() {
        return switch (this) {
            case PENDING -> EnumSet.of(UNSHIPPED, CANCELED);      // waiting for payment
            case UNSHIPPED -> EnumSet.of(SHIPPED, CANCELED);      // paid, not yet handed to carrier
            case SHIPPED -> EnumSet.of(COMPLETED);                // too late to cancel
            case CANCELED -> EnumSet.of(REFUND_APPLIED);          // money goes back if it was taken
            case COMPLETED, REFUND_APPLIED -> EnumSet.noneOf(OrderStatus.class);
        };
    }

    boolean canMoveTo(OrderStatus target) { return next().contains(target); }
}

// File: lld/ch19/amazon/Catalog.java (excerpt, in Product)
    synchronized boolean reserve(int quantity) {
        if (quantity <= 0 || quantity > availableCount) return false;
        availableCount -= quantity;
        return true;
    }

// File: lld/ch19/amazon/OrderService.java (excerpt)
    Order placeOrder(Member member, Address address, Function<BigDecimal, Payment> paymentMethod) {
        if (!member.canBuy()) throw new IllegalStateException("account " + member.account().status());
        List<Item> items = member.cart().snapshot();
        if (items.isEmpty()) throw new IllegalStateException("cart is empty");

        List<Item> reserved = new ArrayList<>();
        for (Item i : items) {
            Product p = catalog.find(i.productId()).orElseThrow();
            if (!p.reserve(i.quantity())) {
                reserved.forEach(r -> catalog.find(r.productId()).orElseThrow().release(r.quantity()));
                throw new IllegalStateException("out of stock: " + i.name());
            }
            reserved.add(i);
        }

        Order order = new Order("ORD-" + orderSeq.incrementAndGet(), member.id(), items, address);
        orders.put(order.orderNumber(), order);

        Payment payment = paymentMethod.apply(order.total());
        order.attachPayment(payment);
        switch (payment.process()) {
            case COMPLETED -> { transition(order, OrderStatus.UNSHIPPED, "paid by " + payment.describe());
member.cart().clear(); }
            case PENDING -> { publish(order, OrderStatus.PENDING); member.cart().clear(); }
            default -> {                                  // declined: give the stock back, keep the cart
                releaseStock(order);
                transition(order, OrderStatus.CANCELED, "payment declined");
            }
        }
        return order;
    }
```

Full runnable code: code/lld/ch19/amazon/ (see the Full Edition).

```
== 6. Illegal transition is rejected by the state machine ==
  ORD-1002: REFUND_APPLIED -> SHIPPED is not allowed
== 9. Ten members race for the last copy ==
  winners=1 losers=9 stock left=0
```

The demo also shows: name search matches every query word (`usb cable` finds "USB-C Cable"); the guest's cart survives registration; each status change goes to email + SMS; the cancel path `PENDING → UNSHIPPED → CANCELED → REFUND_APPLIED` returns stock; a blocked member is rejected; ten threads released together by a `CountDownLatch` produce exactly one winner.

Concurrency & edge cases

Risk	How the code prevents it
Last unit, two buyers	`Product.reserve` is `synchronized` (check + decrement in one step)
Multi-line cart, one line out of stock	`placeOrder` releases everything it reserved (all-or-nothing)
Cancel and ship at the same time	Both lock the `Order` around check-then-move
Payment declined	Stock released, order `CANCELED`, cart kept for retry
Bank transfer never clears	A timeout job cancels old `PENDING` orders
Price changes after adding to cart	`Item` stores unit price; production re-validates at checkout
Blocked member	`Member.canBuy()` checks `ACTIVE`
Deadlock	Each lock held briefly, never while waiting for another in opposite order

Real-world example

Amazon offers *Cancel items* only before the order enters shipping — our missing `SHIPPED → CANCELED` edge. Its 2007 **Dynamo** paper chose availability so "add to cart" always works during failures, merging conflicting cart versions later. At scale, cart, catalog and orders are separate services.

Extensions & follow-up questions

- **New payment (gift card, UPI)** — `GiftCardPayment extends Payment` (add to `permits`); `OrderService` untouched.
- **Last unit, distributed** — `UPDATE ... SET stock = stock - 1 WHERE id = ? AND stock >= 1`, or optimistic locking with a version column.
- **Abandoned orders** — reservation expiry; a job cancels `PENDING` orders after ~30 minutes.
- **Returns** — `COMPLETED → RETURN_REQUESTED → RETURNED → REFUND_APPLIED` (local enum change).
- **Multiple warehouses** — `Inventory` per warehouse; a Strategy picks one per line. **Split shipments** — one order, many `Shipment`s.
- **Search at scale** — Elasticsearch/OpenSearch fed by product events behind `Search`. **Payouts, recommendations, wish lists** — modules listening to order events.

In the interview

Plan: clarify (3 min) → actors/use cases (5) → class diagram, **saying the LSP sentence**: "Ordering is on `Member`, not `Customer`, otherwise every guest inherits a method it is not allowed to use" (10) → state machine, no cancel-after-ship edge (5) → code `placeOrder` and `cancel` (10) → concurrency and extensions (5).

> **30-second answer:** "Catalog with name and category indexes; a cart per customer; ordering only on `Member`; checkout reserves stock for every line atomically, copies the lines into an immutable order, then takes payment through a `Payment` strategy. The order moves through a state machine that forbids cancelling after shipping, and every status change is published to a notification observer."

- *Tax and discounts?* — A `PricingRule` chain (Chain of Responsibility or list of Strategies) applied when building the order; the order stores final line prices.
- *Double-click checkout?* — Client idempotency key; return the existing order for a repeated key.

Self-check

1. **Q:** Why can't `placeOrder` live on `Customer`? **A:** Guests extend `Customer` and must not order; inheriting a method you must refuse violates LSP.
2. **Q:** Two faults of "check stock, pay, then decrement"? **A:** A race lets two buyers pass the check; payment can succeed for goods no longer there.
3. **Q:** Which edges allow a cancel? **A:** `PENDING → CANCELED` and `UNSHIPPED → CANCELED`.
4. **Q:** Why `compareTo` on `BigDecimal`? **A:** `equals` also compares scale (`2.0` ≠ `2.00`).

TL;DR

- Guests browse and fill carts; **only members order** (LSP).
- Checkout = **reserve all lines atomically → copy into an immutable order → pay (Strategy) → publish (Observer)**.
- `OrderStatus` is a **state machine**; cancel after ship has no edge. Money in `BigDecimal`; histories in logs; stock guarded by a lock.

Design an Online Stock Brokerage System

Two halves, both asked in interviews:

- **Part A — the brokerage** (the course): members, money in/out, watchlists, four order types, lots and positions, statements.
- **Part B — the exchange's matching engine** (a Geektrust machine-coding problem kept in kumaransg/LLD as "Navi"): an order book matching by **price-time priority**, implemented and tested.

A **broker** holds your cash and shares and forwards orders; an **exchange** is where orders from all brokers meet and match. The code keeps them apart.

Problem & clarifying questions

An online brokerage lets members trade without a human broker, track transactions, see charts, and get alerts when a stock moves by a chosen amount — faster and cheaper, with real-time prices.

Question	Assumed answer
Order types?	Market, limit, stop-loss, stop-limit
Partial fills?	Yes, at several prices
Separate purchases tracked?	Yes — each is a **lot** (tax depends on which lot you sell)
Money in/out?	Check, wire, electronic bank transfer, each with a per-transfer limit
Reports?	Quarterly statements and a yearly tax statement
Notifications?	When an order executes
Margin, shorting, options, dividends/splits, streaming quotes, settlement delay, wash-sale rule?	Out of scope

Requirements

Functional (brokerage): (1) buy and sell; (2) **multiple watchlists**; (3) **market, limit, stop-loss, stop-limit** orders; (4) **multiple lots** per stock kept apart; (5) **quarterly** and **yearly tax** statements; (6) money via **check, wire, electronic transfer**; (7) notify on execution.

Functional (matching engine, Navi): (8) input lines `<id> <time> <symbol> <buy/sell> <price> <qty>` (e.g. `#1 09:45 BAC sell 240.12 100`) from a file given as argument, trades to console; (9) match each new order against the opposite side — lower sells and higher buys first, **earlier first** at the same price; (10) trade when buy ≥ sell, recorded at the **sell order's price**; (11) print `<buy_id> <price> <qty> <sell_id>`; unmatched quantity rests.

Non-functional: money never over-committed (ten $1,000 orders cannot all pass a $1,000 check); deterministic matching; engine safe under concurrency and conserves quantity.

Use cases

Admin (block/unblock), **Member** (search, trade, watchlists), **System** (order notifications, fetch quotes). Top cases: register/cancel membership, manage watchlists, search by symbol, place/cancel orders, deposit/withdraw.

Core entities

Class	Responsibility
`MoneyTransfer` → `CheckTransfer`, `WireTransfer`, `ElectronicBankTransfer`	Differ only in name and clearing days (5/1/3); the 100,000 limit lives once in the base
`Notifier` / `EmailNotifier`	"Order filled" messages (an `SmsNotifier` would be a second implementation)
`Order` (abstract) + `MarketOrder`, `LimitOrder`, `StopLossOrder`, `StopLimitOrder`	Shared fields; each subclass answers the **two questions**; `StopLimitOrder extends StopLossOrder`
`OrderPart`	One fill (e.g. 70 @ 10.00 + 30 @ 10.05)
`Stock`, `StockInventory`	Symbol + latest price; inventory is the **single source of truth** for prices
`StockExchange`, `OrderListener`	Holds open orders, fills on new orders/ticks, calls back `onFill` / `onCancelled` (never knows `Member`)
`StockLot`, `StockPosition`, `LotSale`	One purchase; all lots of a symbol, sold FIFO unless a lot is named; per-lot realised gain
`Watchlist`	Named set of **symbols** (never stale)
`Trade`, `Statement` → `QuarterlyStatement`, `TaxStatement`	Fills as the member sees them; statements share day-range filtering
`Member`, `Admin`	Cash, **reserved funds**, positions, watchlists, trades; implements `OrderListener`
(engine) `OrderBook`, `MatchingEngine`, `Trade`, `ExecutionPricePolicy`	Price levels with FIFO queues; routing and locking; trade price

Class diagram

Course stages: transfers and notifications → **the order hierarchy** (the part that earns most credit) → stock, inventory, exchange → lots and positions → watchlists, trades, statements → members and admins. `Member` extends `Account` and implements `OrderListener`; it owns positions, watchlists and trades.

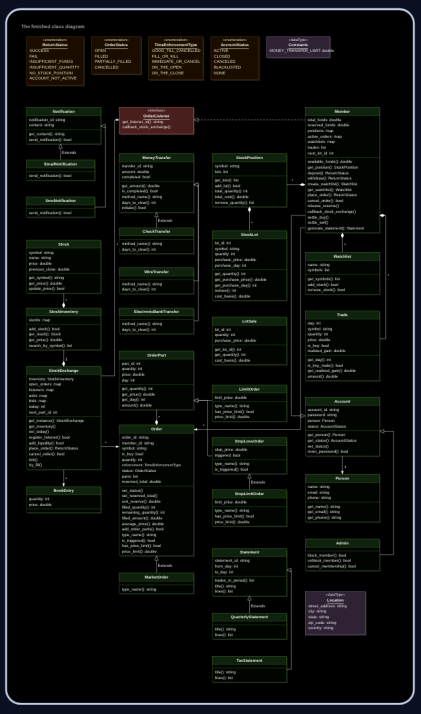
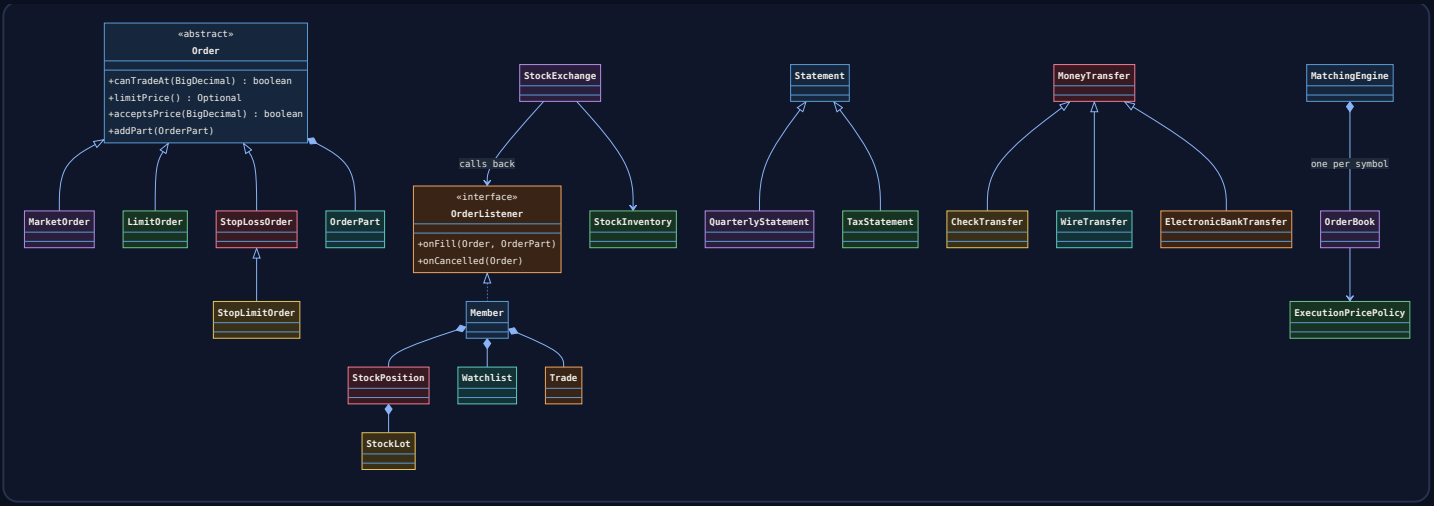


Figure — The finished class diagram of the online stock brokerage (course diagram)



Key flows

Buy order (course activity diagram): pick stock and type, enter quantity, limit and time enforcement → check funds → reserve money → send to exchange → "order placed". **Sell order** mirrors it, choosing **which lots** and checking owned shares instead of money.

The four order types share everything except two answers:

Type	Q1: may it trade at this price?	Q2: worst acceptable price	Plain words
Market	Always	None	"Now, at whatever the market gives"
Limit	Always	The limit	"Buy at 9.50 or less / sell at 10.90 or more"
Stop-loss	Only after price **reaches the stop**, then forever	None (becomes market)	"If it falls to 10.00, get me out"
Stop-limit	Only after the stop	The limit	"If it rises to 305, buy — but pay no more than 306"

A SELL stop triggers when price ≤ stop; a BUY stop when price ≥ stop; once triggered it "latches".

Fill: member reserves (funds or free shares) → exchange checks `canTradeAt`, fills level by level while `acceptsPrice` → `onFill` moves cash, releases the reservation, adds a lot or sells FIFO. Later `tick(symbol, price)` re-checks open orders (stops trigger here).

Matching engine — price-time priority. Per symbol, an `OrderBook` with two sides, each a `TreeMap` price → FIFO `ArrayDeque` (bids highest-first, asks lowest-first). For an incoming order: look at the best opposite price; stop if it does not **cross**; else fill the oldest order at that level for `min(remaining, resting)`; remove filled orders and empty levels; repeat; rest any remainder at the back of its own level.

Order	Book before	What happens
#1 sell 240.12 × 100	empty	rests
#2 sell 237.45 × 90	asks 240.12	rests; best ask 237.45
#3 buy 238.10 × 110	asks 237.45×90, 240.12×100	trade `#3 237.45 90 #2`; 20 left; 240.12 too high; rests 20 @ 238.10
#4 buy 237.80 × 10	best ask 240.12	no cross; rests
#5 buy 237.80 × 40	same	rests **behind #4**
#6 sell 236.00 × 50	bids 238.10×20, 237.80×(10, 40)	fills #3 20, #4 10, #5 20 — all at 236.00 (sell price)

Design decisions & patterns used

Decision	Why	Alternative rejected
One abstract `Order` + two abstract questions	Only "may I trade?" and "worst price?" differ. The course calls it plain polymorphism — not State (no order changes kind); its only named patterns are Observer and Singleton (Template Method/Strategy labels here are ours)	Four copied classes; one class with `type` + `switch`
`StopLimitOrder extends StopLossOrder`	Inherits the trigger, overrides only `limitPrice()`	Duplicated trigger logic
`OrderPart` per fill	Each fill has its own price and becomes a lot	Average price only (loses tax info)
Observer `OrderListener`	Exchange never imports `Member`	Exchange updating balances
Reserved funds / shares	Stops over-commitment; money actually moves only in `onFill`, called by the exchange (course's `callbackStockExchange`)	Balance check only at fill time
Lots, FIFO default	Tax depends on which shares were sold	One average-cost position
Template Method `Statement`	`render` final (shared filter); subclasses fill `body`	Two copies of the filter
`MoneyTransfer` base holds the limit	One rule, three methods	Limit copied per subclass
Singleton?	Course makes `StockExchange` a Singleton; we inject one instance (testable)	Global `getInstance()`
Engine: `TreeMap<price, ArrayDeque>` per side	Best price O(log n); FIFO gives time priority	One `PriorityQueue` (O(n) cancel, no visible levels)
Engine: lock per symbol	Symbols match in parallel; same symbol serialised	Global lock or no lock
Engine: `ExecutionPricePolicy` (**Strategy**)	Navi: sell price; real exchanges: resting order's price	Hard-coding one rule

Java implementation

Part A — Brokerage (`lld.ch19.brokerage`)

Files: `Orders.java`, `Market.java`, `Holdings.java`, `Accounts.java`, `Main.java`. The simulated exchange offers 70 shares per price level, 5 cents apart, 3 levels, so fills come in parts.

```java
// File: lld/ch19/brokerage/Orders.java (excerpt)
abstract class Order {
    // ... id, memberId, symbol, side, quantity, timeInForce, parts, status
    /** Question 1. Market and limit orders always say yes; stop orders wait for the trigger. */
    abstract boolean canTradeAt(BigDecimal marketPrice);

    /** Question 2. Empty means "any price" (market behaviour). */
    abstract Optional<BigDecimal> limitPrice();

    /** Is a fill at this price acceptable? BUY: price <= limit. SELL: price >= limit. */
    final boolean acceptsPrice(BigDecimal price) {
        return limitPrice().map(l -> side == Side.BUY ? price.compareTo(l) <= 0 : price.compareTo(l) >= 0)
                           .orElse(true);
    }
    // ...
}

final class MarketOrder extends Order {
    // ...
    @Override boolean canTradeAt(BigDecimal p) { return true; }
    @Override Optional<BigDecimal> limitPrice() { return Optional.empty(); }
    // ...
}
// ... LimitOrder: canTradeAt -> true, limitPrice -> Optional.of(limit)
class StopLossOrder extends Order {
    // ...
    @Override final synchronized boolean canTradeAt(BigDecimal p) {
        if (!triggered) {
            triggered = side() == Side.SELL ? p.compareTo(stopPrice) <= 0 : p.compareTo(stopPrice) >= 0;
        }
        return triggered;
    }
    @Override Optional<BigDecimal> limitPrice() { return Optional.empty(); }
    // ...
}

/** A stop order that becomes a LIMIT order (not a market order) once triggered. */
final class StopLimitOrder extends StopLossOrder {
    // ...
    @Override Optional<BigDecimal> limitPrice() { return Optional.of(limit); }
    // ...
}

// File: lld/ch19/brokerage/Accounts.java (excerpt, in Member)
    private synchronized ReturnStatus reserveFor(Order o) {
        if (status != AccountStatus.ACTIVE) return ReturnStatus.ACCOUNT_NOT_ACTIVE;
        if (o.side() == Side.BUY) {
            BigDecimal perShare = o.limitPrice().orElse(inventory.priceOf(o.symbol()));
            BigDecimal need = Money.times(perShare, o.quantity());
            if (need.compareTo(availableFunds()) > 0) return ReturnStatus.INSUFFICIENT_FUNDS;
            reservedFunds.put(o.id(), need);           // released part by part as fills arrive
            return ReturnStatus.SUCCESS;
        }
        // ... sells: free shares = position - shares promised to open sells of this symbol
    }
```

```
== 3. Market buy 100 AAPL fills in two parts (two lots) ==
    [email alice@mail.test] filled 70 AAPL @ 10.00 (O1 PARTIALLY_FILLED)
    [email alice@mail.test] filled 30 AAPL @ 10.05 (O1 FILLED)
Tax statement for Alice (days 1-365)
  AAPL L1: 70 sh bought @ 10.00 sold @ 11.00 gain $70.00
  AAPL L2: 30 sh bought @ 10.05 sold @ 10.95 gain $27.00
  AAPL L3: 20 sh bought @ 9.40 sold @ 10.95 gain $31.00
  AAPL L3: 30 sh bought @ 9.40 sold @ 9.95 gain $16.50
  realised gain: $144.50
```

Reading the output: a 200,000 check fails the 100,000 limit. The market order reserved $1,000.00 but cost $1,001.50 — a market order can spend more than reserved (the course's warning). A limit buy waits for the 9.40 tick; a 120-share sell goes FIFO across three lots; the stop does nothing at 10.50 and fires at 9.95; the stop-limit triggers at 310 but its 306 limit blocks the fill until 305.50; a second $2,000 order is refused by the reservation and cancelling returns the money; the tax statement shows gain **per lot**.

Full runnable code: code/lld/ch19/brokerage/ (see the Full Edition).

Part B — Matching engine (`lld.ch19.stockexchange`)

Files: `OrderBook.java`, `MatchingEngine.java` (routing, `synchronized (book)` per symbol, sequence taken inside the lock, trades published outside it), `MatchingEngineTests.java` (nine tests), `Main.java`.

```java
// File: lld/ch19/stockexchange/OrderBook.java (excerpt)
interface ExecutionPricePolicy {
    BigDecimal price(Order buy, Order sell, Order incoming);

    ExecutionPricePolicy SELL_PRICE = (buy, sell, incoming) -> sell.price();
    ExecutionPricePolicy RESTING_PRICE = (buy, sell, incoming) -> incoming == buy ? sell.price() : buy.price();
}

final class OrderBook {
    // ...
    private final NavigableMap<BigDecimal, Deque<Order>> bids = new TreeMap<>(Collections.reverseOrder());
    private final NavigableMap<BigDecimal, Deque<Order>> asks = new TreeMap<>();
    // ...
    /** Match the incoming order against the other side, then rest whatever is left. */
    List<Trade> add(Order incoming) {
        List<Trade> trades = new ArrayList<>();
        NavigableMap<BigDecimal, Deque<Order>> opposite = incoming.side() == Side.BUY ? asks : bids;

        while (!incoming.isDone() && !opposite.isEmpty()) {
            Map.Entry<BigDecimal, Deque<Order>> best = opposite.firstEntry();
            if (!crosses(incoming, best.getKey())) break;          // best price does not meet ours: stop
            Deque<Order> queue = best.getValue();
            Iterator<Order> it = queue.iterator();
            while (!incoming.isDone() && it.hasNext()) {
                Order restingOrder = it.next();                    // oldest first = time priority
                int qty = Math.min(incoming.remainingQty(), restingOrder.remainingQty());
                Order buy = incoming.side() == Side.BUY ? incoming : restingOrder;
                Order sell = incoming.side() == Side.SELL ? incoming : restingOrder;
                trades.add(new Trade(buy.id(), pricePolicy.price(buy, sell, incoming), qty, sell.id()));
                incoming.fill(qty);
                restingOrder.fill(qty);
                if (restingOrder.isDone()) {
                    it.remove();
                    resting.remove(restingOrder.id());
                }
            }
            if (queue.isEmpty()) opposite.pollFirstEntry();
        }
        if (!incoming.isDone()) rest(incoming);
        return trades;
    }
    // ... cancel, bestBid, bestAsk, crosses, rest
}
```

```
== Trades (execution price = sell order's price) ==
#3 237.45 90 #2
#3 236.00 20 #6
#4 236.00 10 #6
#5 236.00 20 #6
== Same input, real-exchange rule (price of the resting order) ==
#3 237.45 90 #2
#3 238.10 20 #6
#4 237.80 10 #6
#5 237.80 20 #6
  9/9 passed
```

The first block matches the problem's expected output exactly; the second prices each trade at the *resting* buyer's price. Tests cover price priority, time priority, no-cross, partial fill, cancel, symbol isolation, validation, and concurrency (8 threads × 500 orders: quantity conserved, book never crossed).

Full runnable code: code/lld/ch19/stockexchange/ (see the Full Edition).

Concurrency & edge cases

Risk	Where	Prevention
Over-spending	Brokerage	`reserveFor` (synchronized on member) subtracts all reservations first
Selling shares twice	Brokerage	Free shares = position − remaining open sells of that symbol
Member/exchange deadlock	Brokerage	Member reserves under its lock, submits outside it; only exchange → member calls
Market order costs more than reserved	Brokerage	Known gap; reserve a buffer (+5%) or a price band
Time in force not enforced	Brokerage	Enum has GTC, FOK, IOC, on-the-open, on-the-close but nothing reads it; belongs in the exchange's fill loop
Same-price orders "together"	Engine	Lock acquisition order = time priority; `sequence` taken inside the lock
Cancel racing a match	Engine	Same per-symbol lock: matched or cancelled, never both
Bad input (duplicate id, zero qty, negative price)	Engine	Validated before registering
Slow trade publishing	Engine	Published after releasing the lock
Floating-point prices	Both	`BigDecimal`; `TreeMap` compares with `compareTo`, so `237.8` = `237.80`

Real-world example

Zerodha offers **SL** (stop-loss limit) and **SL-M** (stop-loss market) — our `StopLimitOrder` and `StopLossOrder`; orders go to exchanges like NSE with price-time priority books. **LMAX** (Martin Fowler's "The LMAX Architecture") runs matching on a **single thread** fed by a ring buffer (the Disruptor) — "one serialised writer per book" taken to its end. US tax rules treat shares as sold FIFO unless you name lots, hence `StockPosition.sell`'s optional lot id.

Extensions & follow-up questions

- **Trailing stop** — new subclass moving its stop with the best price; only `canTradeAt` changes.
- **IOC / FOK** — IOC cancels the remainder after one pass; FOK checks available quantity first and cancels if short.
- **Price alerts** — a `PriceAlert` observer on ticks, firing once past the threshold.
- **Modify order** — reducing quantity keeps priority; price change or increase = cancel + new.
- **Faster engine** — single writer per symbol via a queue, pre-allocated objects, `long` price ticks.
- **Market orders in the engine** — limit at +∞ / 0, never rest the remainder. **Settlement** (T+1) — `unsettledCash`.

In the interview

Brokerage plan: clarify (3) → actors (3) → **key sentence**: *"The four order types differ in exactly two things: whether they may trade yet, and the worst price they accept"* (8) → lots and tax (4) → reserved funds (3) → code the two questions and `StockPosition.sell` (10) → Observer, statements, extensions (5). **Navi, 90 minutes:** parse lines → `OrderBook` with two `TreeMap`s of `ArrayDeque` → matching loop → print trades → tests (sample, price, time, partial, cancel). Pass the sample first, then refactor.

> **30-second answer:** "One abstract `Order` with four tiny subclasses answering 'may I trade now?' and 'what is my worst price?'. Fills are `OrderPart`s and each purchase becomes a lot, sold FIFO. Money is reserved at placement so a member cannot over-commit; the exchange calls back through an `OrderListener`. The matching engine keeps, per symbol, two `TreeMap`s of price levels with FIFO queues — best price first, oldest first — under one lock per symbol."

- *Book data structure and complexity?* — `TreeMap` levels: best level and new level O(log L); FIFO append O(1), O(1) per fill; cancel O(log L) + O(k) in the level (O(1) with a linked handle).

- *Trade price?* — Navi: the sell price. Real continuous trading: the resting order's price (incoming may get "price improvement").

Self-check

1. **Q:** SELL stop at 10.00; price 10.50 → 9.95 → 10.20. When does it trade? **A:** At 9.95, and it stays triggered.
2. **Q:** What does stop-limit add? **A:** A worst price after triggering — it becomes a limit order.
3. **Q:** Why did one market order create two lots? **A:** It filled at two prices; a lot is defined by its purchase price.
4. **Q:** Bids b1 then b2 at 50.00; a sell for 7 at 49. Trades? **A:** b1 gets 5, b2 gets 2 (time priority).
5. **Q:** What is a "crossed" book? **A:** Best bid ≥ best ask — those orders should already have traded.

TL;DR

- **Brokerage:** abstract `Order`, four subclasses answering *may I trade now?* and *worst price?*; fills are `OrderPart`s; purchases are **lots**; money **reserved** at placement; **Observer** callback; statements are a **Template Method**.
- **Matching engine:** per symbol, two `TreeMap<price, FIFO queue>`; best price, then oldest; rest the remainder; one lock per symbol; trade price is a pluggable policy.

Design an Online Auction System

Problem & clarifying questions

An online auction (like eBay) lets a **seller** list an item with a starting price and end time; **bidders** outbid each other; at the end the highest bidder wins — if the hidden **reserve price** was reached.

Question	Assumed answer
Listing contents?	Name, description, category, starting price, minimum increment, optional reserve, duration
Valid bid?	First ≥ starting price; later ≥ current highest + increment
Seller bids? Leader outbids self?	No and no
When does it end?	At end time; a late bid is refused even if the closing job hasn't run
Who is told what?	Outbid bidder immediately; winner and seller at close; seller if unsold
Proxy bidding, anti-sniping, payments?	Extensions

Requirements

Functional: (1) register and log in (login is the usual auth service, not coded); (2) create listings; (3) search by name, category, price range; (4) bid on **active** listings; (5) keep the **current highest bid** and notify; (6) close at end time — **highest bidder wins** if the reserve is met. **Non-functional:** (7) many simultaneous bids stay consistent (one true highest); (8) new bid rules, channels or search criteria without rewriting the core.

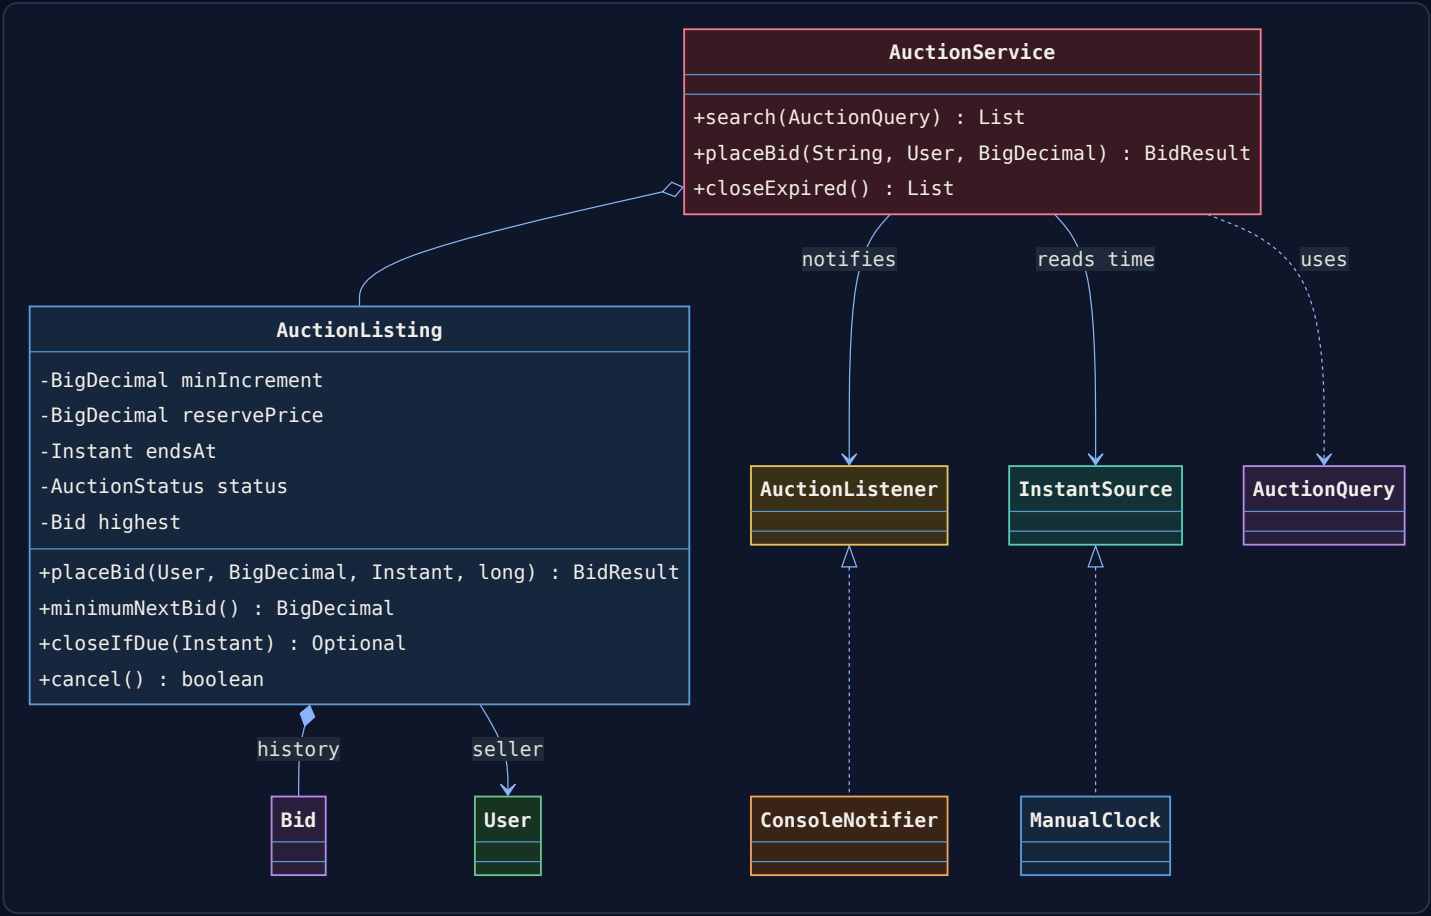
Use cases

Actor	Use cases
Seller	Create listing, cancel (only while nobody has bid), see result
Bidder	Search, view listing and history, bid, be told when outbid or won
System (scheduler)	Activate, close expired, declare winner, notify

Core entities

Class	Responsibility
`User`, `Bid`	Records; bid = bidder, `BigDecimal` amount, time, sequence
`AuctionStatus`	`SCHEDULED → ACTIVE → CLOSED`, or `CANCELLED`
`AuctionListing`	Item, seller, prices, times, history, highest bid; **owns the bidding rules** under its own lock
`BidOutcome` / `BidResult` / `CloseResult`	Why accepted/refused, next minimum, accepted bid, who was outbid; winner and reserve met
`AuctionQuery`	Fluent search; each criterion a `Predicate` (**Specification**)
`AuctionListener` / `ConsoleNotifier`	Observer: new high bid, outbid, won, unsold
`AuctionService`	Facade; calls listeners **outside** listing locks
`ManualClock`	An `InstantSource` moved forward in demos and tests

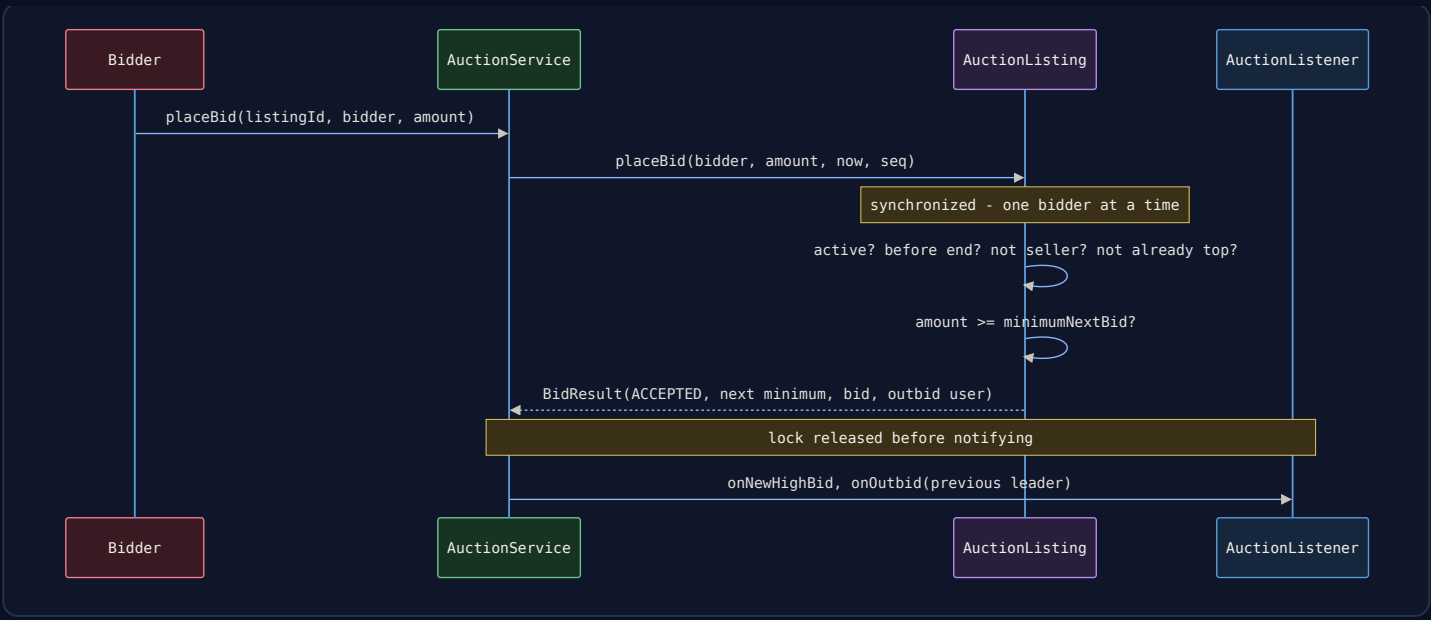
Class diagram



Key flows

Lifecycle: SCHEDULED → ACTIVE (start time) or CANCELLED (seller withdraws); ACTIVE → CANCELLED (withdraw with no bids) or CLOSED (end time). Activation is **lazy**: every method first calls `refresh(now)`, so no thread is needed to flip the state.

Placing a bid (the critical section):



Closing: a scheduler (a `ScheduledExecutorService` in production) calls `closeExpired()`; `closeIfDue(now)` atomically moves `ACTIVE → CLOSED` and returns a winner only if the reserve is met; notifications go out after locks are released.

Design decisions & patterns used

Decision	Why	Alternative rejected
Bid rules inside `AuctionListing` under its lock (a *monitor*: all methods `synchronized`)	Compare-and-replace is atomic; per-listing lock, so auctions don't wait for each other	Global lock; or `AtomicReference<Bid>` + CAS (valid, but seller/leader/time/increment rules are easier under one lock)
Refusals returned (`BidOutcome`), not thrown	A low bid is normal business; caller gets the next minimum	Exception per low bid
Observer called outside the lock	Slow notifications don't block bidders; `BidResult` carries a **snapshot** taken inside the lock	Notifying inside `synchronized` (slow, deadlock risk)
Minimum increment	No 1-cent wars; predictable next minimum	Any higher amount
Reserve price checked only at close	Seller protection without revealing it	Rejecting bids below reserve
Injected `InstantSource`	Testable time (`ManualClock`)	`Instant.now()` everywhere
Specification for search	Each criterion a `Predicate`, combined with `and`	An ever-growing `search(...)` parameter list
Not a Singleton	awesome-low-level-design uses a Singleton `AuctionSystem`; we pass one instance so tests get fresh ones	Global shared state

Java implementation

Package `lld.ch19.auction`: `Auction.java`, `AuctionService.java`, `Main.java`.

```java
// File: lld/ch19/auction/Auction.java (excerpt, in AuctionListing)
    synchronized BidResult placeBid(User bidder, BigDecimal rawAmount, Instant now, long seq) {
        refresh(now);
        BigDecimal amount = money(rawAmount);
        if (status != AuctionStatus.ACTIVE) return refuse(BidOutcome.AUCTION_NOT_ACTIVE);
        if (!now.isBefore(endsAt)) return refuse(BidOutcome.AUCTION_ENDED);  // closing job may not have run yet
        if (bidder.id().equals(seller.id())) return refuse(BidOutcome.SELLER_CANNOT_BID);
        if (highest != null && highest.bidder().id().equals(bidder.id())) return refuse(BidOutcome.ALREADY_HIGHEST);
        if (amount.compareTo(minimumNextBid()) < 0) return refuse(BidOutcome.TOO_LOW);

        Optional<User> outbid = Optional.ofNullable(highest).map(Bid::bidder);
        highest = new Bid(bidder, amount, now, seq);
        bids.add(highest);
        return new BidResult(BidOutcome.ACCEPTED, minimumNextBid(), Optional.of(highest), outbid);
    }

    /** First bid must reach the starting price; later bids must beat the leader by the increment. */
    synchronized BigDecimal minimumNextBid() {
        return highest == null ? startingPrice : highest.amount().add(minIncrement);
    }

    /** ACTIVE -> CLOSED once time is up. Returns empty if not yet due (or already closed). */
    synchronized Optional<CloseResult> closeIfDue(Instant now) {
        refresh(now);
        if (status != AuctionStatus.ACTIVE || now.isBefore(endsAt)) return Optional.empty();
        status = AuctionStatus.CLOSED;
        boolean reserveMet = highest != null && highest.amount().compareTo(reservePrice) >= 0;
        return Optional.of(new CloseResult(this, reserveMet ? Optional.of(highest) : Optional.empty(), reserveMet));
    }

// File: lld/ch19/auction/AuctionService.java (excerpt)
    BidResult placeBid(String listingId, User bidder, BigDecimal amount) {
        AuctionListing a = listings.get(listingId);
        if (a == null) throw new IllegalArgumentException("no listing " + listingId);
        BidResult r = a.placeBid(bidder, amount, clock.instant(), bidSeq.incrementAndGet());
        r.accepted().ifPresent(top -> {                // snapshot taken inside the lock
            listeners.forEach(l -> l.onNewHighBid(a, top));
            r.outbid().ifPresent(u -> listeners.forEach(l -> l.onOutbid(a, u, top.amount())));
        });
        return r;
    }
```

Full runnable code: code/lld/ch19/auction/ (see the Full Edition).

```
Ana bids 90 (below start) -> TOO_LOW (next minimum 100.00)
Ana bids 120 (already top) -> ALREADY_HIGHEST (next minimum 110.00)
Bob bids 105 (increment is 10) -> TOO_LOW (next minimum 110.00)
Sam bids on own item -> SELLER_CANNOT_BID (next minimum 120.00)
accepted=28 (varies run to run), highest=650.00 by bidder49, history strictly increasing by >= increment: true
late bid -> AUCTION_ENDED (next minimum 660.00)
```

Every refusal rule fires; outbid notices print before the bidder's own result line (the service notifies before returning). The watch closes unsold (60 < 500 reserve). Of 50 concurrent bids the accepted count varies, but the highest is always 650.00 and the history always rises by at least the increment — the invariant the lock protects.

Concurrency & edge cases

Case	Handling
Same amount at the same time	First through the lock wins; the second is `TOO_LOW`
Bid exactly at end time	`!now.isBefore(endsAt)` → `AUCTION_ENDED` (end is exclusive)
Close job runs late / races a bid	Late bids refused anyway; both `synchronized`, so the bid lands before or after the close
Slow or failing listener	Outside the lock; production publishes to a retrying queue
Seller cancels after bids	Refused
Money precision	Scale 2 `BigDecimal`, `compareTo`
Very hot auction	Single writer per auction (partition by id, e.g. Kafka), or `UPDATE ... SET high = ? WHERE id = ? AND high < ?`

Real-world example

eBay uses **proxy bidding** (it bids for you up to your maximum, by the increment), hidden **reserve prices**, and fixed end times — hence "sniping". Other sites use a **soft close** (late bids extend the end) to defeat it.

Extensions & follow-up questions

- **Proxy bidding** — per-bidder `maxAmount`; after each accepted bid, inside the lock, the strongest proxy responds with `min(its max, competitor + increment)`; ties to the earlier proxy (`Bid.seq`).
- **Anti-sniping** — soft close: a bid in the last N minutes extends `endsAt` (non-final, same lock).
- **Buy It Now** — fixed price closes the auction if paid before the first bid.
- **Dutch / sealed / Vickrey** — an `AuctionRules` strategy (`validate`, `winner`, `price`); Vickrey: highest wins, pays second-highest.
- **Payment after winning** — `AWAITING_PAYMENT`, 48 hours, else offer to the second bidder.
- **Scale the close job** — a priority queue by `endsAt`.

In the interview

Plan: bid rules (4 min) → entities and lifecycle (8) → code `placeBid` under the listing lock and **say why the lock is per listing** (12) → closing and notifications outside the lock (6) → concurrency and extensions (10).

> **30-second answer:** "Each listing owns its bid rules and highest bid, and `placeBid` is synchronized on the listing, so compare-and-replace is atomic and different auctions never contend. Refusals return a reason and the next minimum. The service notifies outbid users and, at close, the winner and seller, always after releasing the lock. Time comes from an injected clock, so closing and late bids are testable."

- *Why not `AtomicReference` + CAS?* — The rule also checks seller, leader and time; one small lock is simpler and still fast per listing.
- *Close on time across servers?* — A scheduler with a time-ordered index; `closeIfDue` is idempotent, so running twice is harmless.
- *Live price for thousands?* — Push `onNewHighBid` to pub/sub and WebSockets; viewers read a cache.

Self-check

1. **Q:** When is a bid `TOO_LOW`? **A:** Below the starting price with no bids, or below highest + increment.
2. **Q:** Why notify after the listing call returns? **A:** Short lock; no slowdown or deadlock from notification code.
3. **Q:** Bid at exactly the end instant? **A:** Refused — the end is exclusive.
4. **Q:** Where does the reserve matter? **A:** Only at close.
5. **Q:** What is the Specification pattern here? **A:** Search criteria as composable `Predicate`s.

TL;DR

- One **lock per listing**; all bid rules inside it; refusals are values, not exceptions.
- `SCHEDULED → ACTIVE → CLOSED | CANCELLED` ; activation lazy, closing idempotent.
- **Observer** after the lock; **injected clock**; **Specification** for search.

Design a Music Streaming Service (Spotify)

Problem & clarifying questions

Users search a big catalog, build playlists and play songs with play, pause, skip, previous, seek, shuffle and repeat. Free and paid plans differ; the service suggests music.

Question	Assumed answer
Search?	Songs, albums, artists
Playlists?	Users create, reorder, edit
Controls?	Play, pause, resume, skip, previous, seek; shuffle; repeat off / all / one
Plans?	FREE (ads, limited skips/hour, no downloads), PREMIUM (no ads, unlimited skips, offline)
Recommendations?	A history-based stub (real ones are ML)
Audio bytes, DRM, CDNs?	Out of scope — the player models *what* plays

Requirements

Functional: (1) browse/search songs, albums, artists; (2) playlists; (3) accounts (auth is the usual service); (4) play, pause, skip, seek (+ previous, shuffle, repeat); (5) recommendations from history; (6) FREE vs PREMIUM limits. **Non-functional:** (7) concurrent users on a read-mostly, thread-safe catalog; (8) large scale (indexes in production); (9) extensible — sharing, offline, new play orders and recommenders.

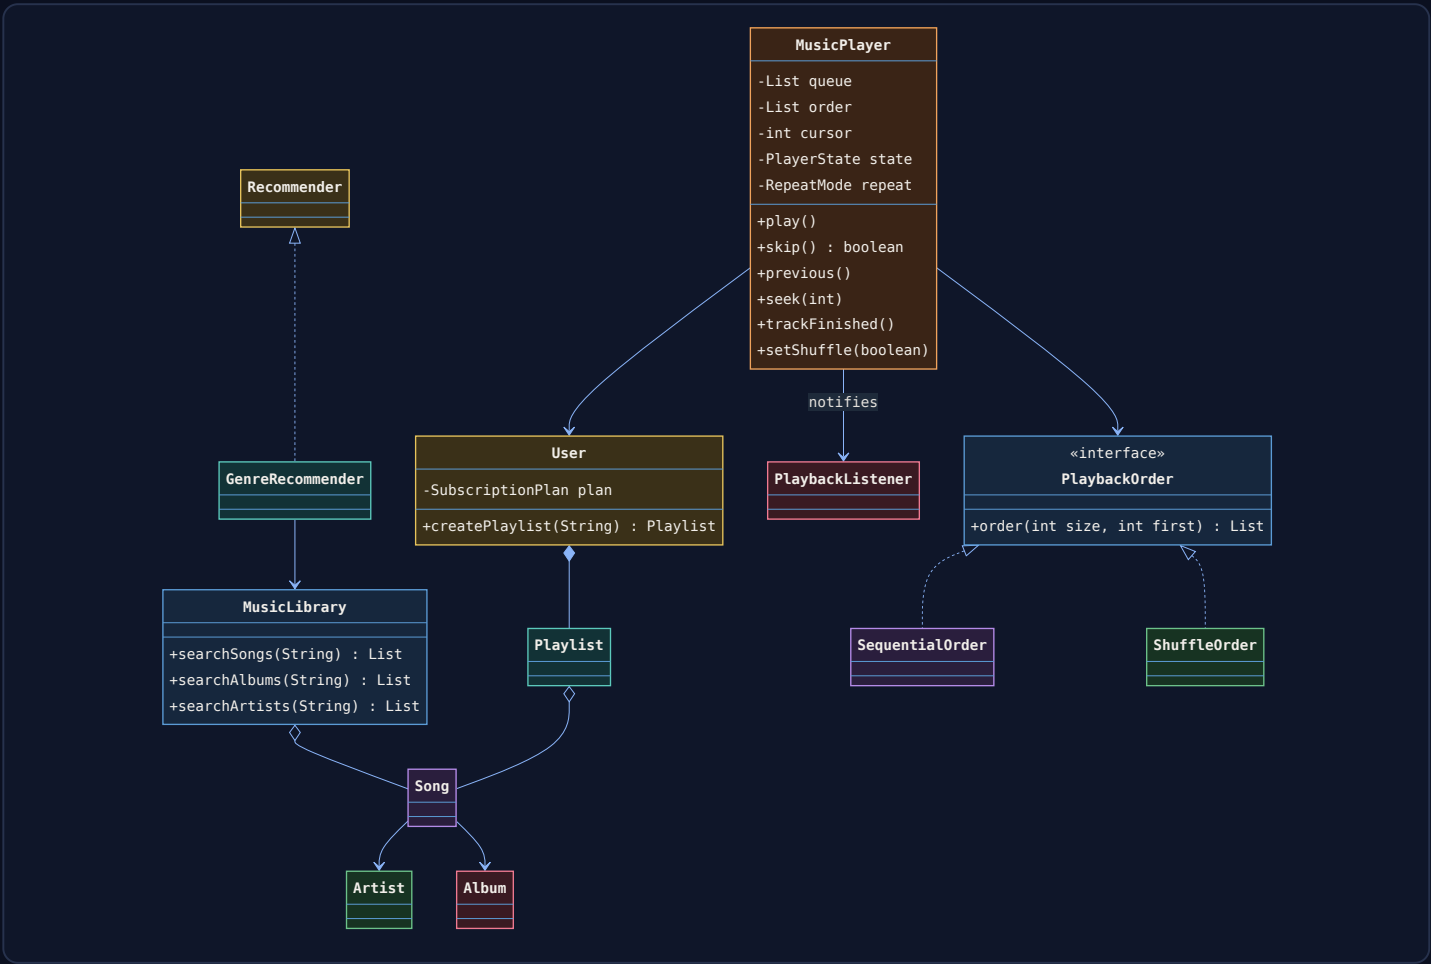
Use cases

Actor	Use cases
Listener	Search, controls, shuffle/repeat, playlists, upgrade plan
Artist / label	Add songs and albums
System	Record plays (history, royalties), ads for free users, recommend

Core entities

Class	Responsibility
`Artist` , `Album` , `Song` , `Genre`	Immutable catalog records
`MusicLibrary`	Thread-safe catalog with search by title, album, artist
`SubscriptionPlan`	Enum **carrying the rules**: price, skips/hour, ads, offline
`User` , `Playlist`	Plan, playlists, history; ordered song list (add, remove, move)
`PlaybackOrder` → `SequentialOrder` , `ShuffleOrder`	Strategy: a permutation starting from the current song
`RepeatMode` , `PlayerState`	`OFF/ALL/ONE` ; `STOPPED/PLAYING/PAUSED`
`MusicPlayer`	One session: queue, order, cursor, position, state, skip window, ad counter
`PlaybackListener`	Observer on "song started"
`Recommender` → `GenreRecommender`	Strategy: unplayed songs from favourite genres

Class diagram



Key flows

Player states: STOPPED → PLAYING (play) ⇄ PAUSED (pause/resume); PLAYING → PLAYING on skip or track end with more songs; PLAYING → STOPPED at end of queue with repeat OFF.

Song ends or Next pressed: natural end with repeat ONE → replay; else more songs in the order → cursor + 1; else repeat ALL/ONE → build a new order (reshuffle if shuffling) and play the first; else STOPPED. Pressing Next on repeat-ONE moves on (repeat-one only affects natural ends). **Skip** first checks the plan's rolling one-hour limit: a deque of skip timestamps, drop those older than an hour, compare with `skipsPerHour` .

Shuffle keeps the current song: a new permutation starts with the playing song and shuffles the rest (`Collections.shuffle(list, random)`); turning it off rebuilds the sequential order from the current song. The queue never changes — only the **order of indexes** — so nothing is lost.

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy for play order	Sequential, shuffle, "smart shuffle" interchangeable	`if (shuffle)` in skip, previous, wrap
Order = indexes into an unchanged queue	Shuffle off restores the original	Physically shuffling the queue
Seeded `Random`	Repeatable demos/tests	`Math.random()`
Plan rules in `SubscriptionPlan` enum	No `if (plan == FREE)` ; new plan = new constant	Boolean flags on `User`
Observer for plays	History, royalties, recommendations react to "song started"	Player calling each subsystem
Strategy for recommendations	Genre today, collaborative filtering tomorrow	One algorithm in `User`
One `MusicPlayer` per session, all `synchronized`	Phone + desktop may send commands at once; library uses concurrent maps	One shared player
State as enum, not State pattern	Three trivial states; GoF State pays off only for substantial per-state behaviour	Three state classes
No Singletons	awesome-low-level-design makes `MusicLibrary` , `UserManager` , `MusicRecommender` Singletons; we inject	Global `getInstance()`

Java implementation

Package `lld.ch19.music` : `Catalog.java` , `Player.java` , `Main.java` .

```java
// File: lld/ch19/music/Catalog.java (excerpt)
enum SubscriptionPlan {
    FREE(new BigDecimal("0.00"), 6, true, false),
    PREMIUM(new BigDecimal("10.99"), Integer.MAX_VALUE, false, true);
    // ... monthlyPrice, skipsPerHour, ads, offlineDownloads
}

// File: lld/ch19/music/Player.java (excerpt)
/** Strategy: in what order should the queue be played? */
interface PlaybackOrder {
    /** A permutation of 0..size-1 that starts with {@code first}. */
    List<Integer> order(int size, int first);
}

final class ShuffleOrder implements PlaybackOrder {
    private final Random random;
    ShuffleOrder(long seed) { this.random = new Random(seed); }   // seeded: repeatable demo and tests
    @Override public List<Integer> order(int size, int first) {
        List<Integer> rest = new ArrayList<>(IntStream.range(0, size).filter(i -> i != first).boxed().toList());
        Collections.shuffle(rest, random);
        rest.add(0, first);                    // the song you are hearing keeps playing
        return rest;
    }
}
// ... in MusicPlayer:
    synchronized boolean skip() {
        if (state == PlayerState.STOPPED) return false;          // nothing to skip
        Instant now = clock.instant();
        while (!recentSkips.isEmpty() && recentSkips.peekFirst().isBefore(now.minus(Duration.ofHours(1))))
recentSkips.pollFirst();
        if (recentSkips.size() >= user.plan().skipsPerHour) {
            System.out.println("  [player] skip limit reached on " + user.plan() + " plan");
            return false;
        }
        recentSkips.addLast(now);
        advance(false);
        return true;
    }

    synchronized void setShuffle(boolean on) {
        int now = order.isEmpty() ? 0 : order.get(cursor);
        playbackOrder = on ? new ShuffleOrder(42) : new SequentialOrder();
        order = playbackOrder.order(queue.size(), now);
        cursor = 0;
    }
```

Full runnable code: code/lld/ch19/music/ (see the Full Edition).

```
== 4. Shuffle keeps the current song, reorders the rest ==
  now Blue in Green - Miles Davis, up next [Love of My Life - Queen, So What - Miles Davis, Bohemian Rhapsody - Queen,
Hello - Adele]
  [player] skip limit reached on FREE plan
  recommended: [You're My Best Friend - Queen, All Blues - Miles Davis, When We Were Young - Adele]
```

Repeat-ONE replays "Bohemian Rhapsody" on natural end; on FREE an ad plays every three songs, repeat-ALL wraps and reshuffles, and the seventh skip in the hour is refused; PREMIUM has no ads and with repeat OFF the queue ends; the recommender suggests *unheard* songs from the most-played genres.

Concurrency & edge cases

Case	Handling
Two devices press "next" together	`synchronized` player: applied one after the other
Catalog updated during search	`ConcurrentHashMap`, weakly consistent iteration
Playlist edited while playing	`load` copies it (`List.copyOf`) — the queue is a snapshot
Seek out of range	Clamped to `[0, duration]`
Previous after 3 seconds	Restarts the song; within 3 s goes back one
Skip when stopped	Refused, not counted
Plan change mid-session	`plan` is `volatile`
Empty queue	`play()` throws `IllegalStateException`

Real-world example

Spotify's free tier has limited skips (e.g. six per hour on mobile) and ads; Premium removes both and allows downloads — our enum's rules. In 2014 Spotify engineers wrote that truly random shuffle "felt" non-random when an artist repeated, so they spread same-artist songs apart — a new `PlaybackOrder` in our design. Discover Weekly is the industrial `Recommender`.

Extensions & follow-up questions

- **Play next / add to queue** — a separate user queue consumed before the context order.
- **Offline** — `offlineDownloads`, a `DownloadManager`, and a local/remote `AudioSource` Strategy.
- **Collaborative playlists** — collaborators + edit log, changes pushed via Observer.
- **Crossfade / gapless** — a `TransitionStrategy` around `trackFinished`.
- **Royalties** — a `PlaybackListener` counting a play after a listening threshold.
- **Search at scale** — inverted index with prefix and fuzzy matching. **Smart shuffle** — a `PlaybackOrder` spacing out artists.

In the interview

Plan: controls, plans, scope (4 min) → catalog + search (6) → user, playlist, plan (5) → **player**: queue + order + cursor, state machine, shuffle/repeat (15) → Observer, recommender, concurrency, extensions (10).

> **30-second answer:** "Immutable catalog records in a thread-safe library. Each user has a player session holding the queue, a play order produced by a `PlaybackOrder` strategy, a cursor and a repeat mode. Shuffle rebuilds the order around the current song. Plan limits live in the plan enum. A playback listener records history, which feeds a pluggable recommender."

- *How does shuffle avoid repeats?* — It is a permutation (each index once), not "random next".
- *Previous in shuffle mode?* — Walk back along the same order list.

Self-check

1. **Q:** Why store the order as indexes? **A:** Shuffle off restores the original order.
2. **Q:** Repeat-ONE and Next? **A:** Moves on; replay only on natural end.
3. **Q:** How is "6 skips per hour" enforced? **A:** Deque of skip times; drop entries older than an hour; refuse at the limit.
4. **Q:** Why an enum, not the State pattern? **A:** The states have almost no distinct behaviour.

TL;DR

- Catalog = immutable records in a concurrent library.
- Player = **queue + order (Strategy) + cursor + repeat mode + state**; shuffle keeps the current song.
- Plan rules in the **enum**; plays via **Observer**; recommendations via **Strategy**.

Design an Online Food Delivery Service (Swiggy) + FoodKart

Two shapes:

- **Part A — Swiggy-style platform** (awesome-low-level-design): restaurants and menus, cart, tracked orders, payments, **rider assignment** by a pluggable strategy. We add **kitchen capacity** (max items cooking at once), the twist of a reported Flipkart food-ordering variant (not part of FoodKart).
- **Part B — FoodKart** (Flipkart machine coding, kumaransg/LLD): command-driven — users and pincodes, one-dish restaurants, ratings, `show_restaurant` by rating or price, `place_order` — plus a **restaurant selection strategy** filtered by capacity.

Problem & clarifying questions

Question	Assumed answer
Mix restaurants in one cart?	No — one restaurant per cart
Who confirms?	The restaurant; it refuses if the kitchen is full
When is capacity freed?	When food is ready, not when delivered
Rider choice?	A strategy: nearest, or fair share (fewest deliveries today)
No rider free?	Order waits in a FIFO queue; the next freed rider takes the oldest
Customer cancel?	Only before cooking (PENDING or CONFIRMED)
Payments?	UPI, card, cash on delivery (nothing to refund)
GPS, ETA, batching, surge?	Extensions

Requirements

Functional (Part A): (1) browse restaurants, view menus, order; (2) restaurants manage menus, prices, availability; (3) agents accept and fulfil; (4) track status with notifications; (5) several payment methods; (6) real-time notifications to customers, restaurants, agents.

Functional (Part B — FoodKart): (7) `register_user(name, gender, phone, pincode)`, phone unique; (8) `login_user(phone)` switches context; (9) `register_restaurant(name, pincodes "A/B", dish, price, quantity)` by the logged-in owner — **one** dish, **many** pincodes; (10) `update_quantity`, `update_location` — owner only; (11) `rate_restaurant(name, 1..5, optional comment)`, rating = average; (12) `show_restaurant(rating|price)` — **serviceable** (my pincode **and** quantity > 0), **descending**; (13) `place_order(name, qty)`, bonus order history; (14) *variant:* order without naming a restaurant, chosen by a **selection strategy** among those delivering here **with** capacity for the whole quantity.

Non-functional: no double-booked kitchen capacity or riders under concurrency; no negative stock, no rider on two orders; in-memory, demoable driver, graceful failures.

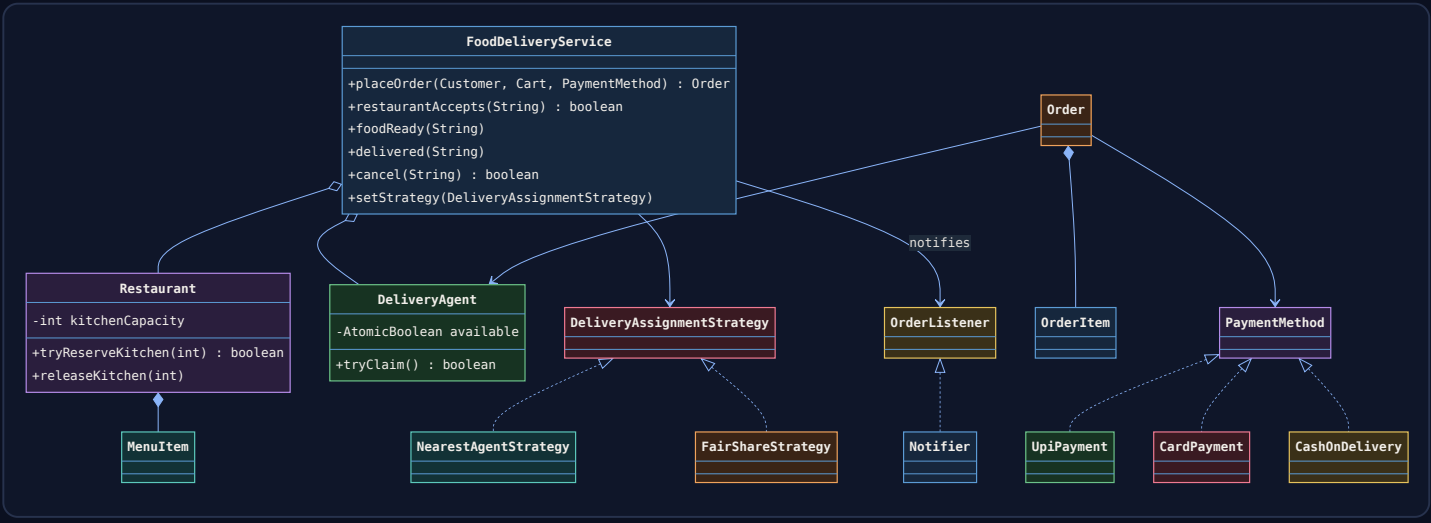
Use cases

Actor	Use cases
Customer	Browse nearby, cart, pay, order, track, cancel, rate
Restaurant	Manage menu, accept/refuse, preparing, ready
Delivery agent	Get assigned, pick up, deliver
System	Fees, rider assignment, waiting queue, notifications
Admin	Onboard restaurants and delivery agents
Payment gateway	Charge UPI/card (included in "place order"); refund extends "cancel" and a kitchen-full refusal when paid upfront

Core entities

Class	Responsibility
`Location`, `Customer`, `MenuItem`	Point + `distanceTo`; customer details; name, price, available flag
`Restaurant`	Menu, open flag, **kitchen capacity**, `tryReserveKitchen` / `releaseKitchen`
`Cart`	One restaurant + quantities; refuses a second restaurant
`OrderItem`	Snapshot of name, **unit price at order time**, quantity
`OrderStatus`	PENDING → CONFIRMED → PREPARING → READY_FOR_PICKUP → OUT_FOR_DELIVERY → DELIVERED, or CANCELLED
`Order`	Customer, restaurant, items, fee, payment, history, agent
`PaymentMethod` → `UpiPayment`, `CardPayment`, `CashOnDelivery`	Sealed; COD not paid upfront
`DeliveryAgent`	Location, `AtomicBoolean available`, deliveries today; `tryClaim()` is a CAS
`DeliveryAssignmentStrategy` → `NearestAgentStrategy`, `FairShareStrategy`	Ranks free riders
`OrderListener` / `Notifier`, `FoodDeliveryService`	Observer; facade with the waiting-for-rider queue
(FoodKart) `User`, `Restaurant`, `Review`, `FoodOrder`, `SortBy`, `RestaurantSelectionStrategy`, `FoodKartService`	`Restaurant.take(q)` is the atomic capacity check

Class diagram



FoodKart (Part B): `FoodKartService` (current user; `registerUser`, `loginUser`, `registerRestaurant`, `updateQuantity`, `updateLocation`, `rateRestaurant`, `showRestaurants(SortBy)`, `placeOrder(name, qty)`, `placeOrder(qty, strategy)`, `orderHistory`) holds `User`s, `Restaurant`s (pincodes, one dish, price, quantity, `take`, `serviceableFor`, `rating`; owns `Review`s) and `FoodOrder`s; `RestaurantSelectionStrategy` → `HighestRatingStrategy`, `LowestPriceStrategy`.

Key flows

Order states: PENDING → CONFIRMED (accepted, kitchen slots reserved) or CANCELLED (kitchen full / customer; refund if prepaid); CONFIRMED → PREPARING or CANCELLED (slots released); PREPARING → READY_FOR_PICKUP (slots released) → OUT_FOR_DELIVERY (rider assigned) → DELIVERED. Real apps split pickup into "rider assigned / arrived / picked up" — a local enum change.

Place order (activity): browse within radius → cart (a second restaurant's item is refused) → choose payment → restaurant open and items available? (no → reject) → copy prices, fee by distance → if paid upfront, payment must succeed → order PENDING, restaurant notified → kitchen has room? yes → CONFIRMED (slots reserved); no → CANCELLED and refunded if prepaid.

Food ready → rider: `foodReady` marks READY_FOR_PICKUP, releases slots, adds the order to the queue; while the queue is not empty: the strategy ranks free agents, `tryClaim()` (CAS true → false) on each in turn; claimed → poll, assign, OUT_FOR_DELIVERY, notify; none → stop, order waits. `delivered()` frees the rider and runs the same loop.

FoodKart `place_order(name, qty)` **:** logged in? → restaurant exists? → delivers to my pincode? → `take(qty)` atomically → record history → "Order Placed Successfully." else "Cannot place order". **Selection variant:** filter by pincode and `quantity ≥ qty` → strategy picks → `take(qty)`; if a racer took it, drop that restaurant and ask again (bounded retries).

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy for riders; `rank` returns a list	Swappable at run time (`setStrategy`); the list lets us try the next rider if the first was just taken	Returning one rider
CAS on the rider (`compareAndSet(true,false)`)	Exactly one claimer wins, even from outside the `synchronized` dispatcher (another dispatcher, a "rider accepts" button)	`isAvailable()` then set
Kitchen capacity reserved on accept, released on ready	"Cannot accept while the kitchen is full"	Unlimited orders
FIFO waiting queue	No free rider is normal at peak; fair	Failing the order
State machine in `OrderStatus`	Cancel after cooking impossible	Scattered `if` s
`OrderItem` copies price	Later menu changes can't alter orders	Referencing `MenuItem`
Sealed `PaymentMethod` records	Closed set; COD overrides `paidUpfront()`	`String paymentType`
Observer `OrderListener`	Customer, restaurant, rider apps react to the same events	Service printing itself
One restaurant per cart	Single kitchen, pickup, fee	Split orders
FoodKart: `SortBy` enum holds its comparator	`sorted(sortBy.comparator)` , descending, name tie-break	`if (sort.equals("rating"))`
FoodKart: owner recorded; `take(q)` synchronized	Only the owner updates; two users can't take the last plates	Anyone updates; get + set

Java implementation

Part A — Food delivery platform (`lld.ch19.fooddelivery`)

Files: `Model.java` , `FoodDeliveryService.java` , `Main.java` .

```java
// File: lld/ch19/fooddelivery/Model.java (excerpt)
    synchronized boolean tryReserveKitchen(int items) {
        if (itemsInKitchen + items > kitchenCapacity) return false;
        itemsInKitchen += items;
        return true;
    }
// ... in DeliveryAgent:
    boolean tryClaim() { return available.compareAndSet(true, false); }

    void completeDelivery() { deliveriesToday.incrementAndGet(); available.set(true); }

// File: lld/ch19/fooddelivery/FoodDeliveryService.java (excerpt)
interface DeliveryAssignmentStrategy {
    List<DeliveryAgent> rank(Order order, List<DeliveryAgent> free);
}
// ...
/** Fewest deliveries today first (spreads earnings), nearest as tie-breaker. */
final class FairShareStrategy implements DeliveryAssignmentStrategy {
    @Override public List<DeliveryAgent> rank(Order o, List<DeliveryAgent> free) {
        Location r = o.restaurant().location();
        return free.stream().sorted(Comparator.comparingInt(DeliveryAgent::deliveriesToday)
                .thenComparingDouble(a -> a.location().distanceTo(r))).toList();
    }
}
// ...
    private synchronized void dispatchWaiting() {
        Order o;
        while ((o = awaitingRider.peek()) != null) {
            List<DeliveryAgent> free = agents.stream().filter(DeliveryAgent::isAvailable).toList();
            Optional<DeliveryAgent> got = Optional.empty();
            for (DeliveryAgent a : strategy.rank(o, free)) if (a.tryClaim()) { got = Optional.of(a); break; }
            if (got.isEmpty()) {
                System.out.println("    (" + o.id() + " waiting for a free rider)");
                return;
            }
            awaitingRider.poll();
            o.assign(got.get());
            transition(o, OrderStatus.OUT_FOR_DELIVERY);
        }
    }
```

Full runnable code: code/lld/ch19/fooddelivery/ (see the Full Edition).

```
== 4. Kitchen capacity: 4 items max at Dosa Corner ==
    (Dosa Corner kitchen full: needs 2 item slots, free 1)
    [OD103 CANCELLED] -> customer Asha
    refund 111.20 to card ****4242
    (OD104 waiting for a free rider)
    [OD101 DELIVERED] -> customer Asha
    [OD104 OUT_FOR_DELIVERY] -> customer Vikram: Ravi is on the way
  o6 rider: Meena (fewer deliveries today)
```

The far grill (beyond 10 km) is hidden; a second restaurant in the cart is refused; fee = 20 + 8 × 1.4 km = 31.20 and a later dosa price change does not alter the order; a card order refused for a full kitchen is cancelled and refunded; an order out for delivery cannot be cancelled while a confirmed card order can (and is refunded); with both riders at one spot, fair share picks Meena (1 delivery) over Ravi (2).

Part B — FoodKart (`lld.ch19.foodkart`)

Files: `FoodKart.java`, `Main.java` (replays the sample test case, then extras).

```java
// File: lld/ch19/foodkart/FoodKart.java (excerpt)
    synchronized boolean take(int q) {
        if (q <= 0 || q > quantity) return false;
        quantity -= q;
        return true;
    }

    /** Serviceable = delivers to this pincode AND has non-zero quantity (from the problem statement). */
    boolean serviceableFor(String pincode) { return pincodes.contains(pincode) && quantity() > 0; }
// ...
enum SortBy {
    RATING(Comparator.comparingDouble(Restaurant::rating).reversed()),
    PRICE(Comparator.comparing(Restaurant::price).reversed());

    final Comparator<Restaurant> comparator;
    SortBy(Comparator<Restaurant> c) { comparator = c.thenComparing(Restaurant::name); }
}

/** Strategy: when the user does not name a restaurant, which one should we pick? */
interface RestaurantSelectionStrategy {
    Optional<Restaurant> select(List<Restaurant> canFulfil);
}
// ... HighestRatingStrategy, LowestPriceStrategy
    String placeOrder(int qty, RestaurantSelectionStrategy strategy) {
        User u = requireLogin();
        List<Restaurant> canFulfil = restaurants.values().stream()
                .filter(r -> r.pincodes().contains(u.pincode()) && r.quantity() >= qty).toList();
        for (int attempt = 0; attempt < 3; attempt++) {           // retry if a racer took the capacity
            Optional<Restaurant> pick = strategy.select(canFulfil);
            if (pick.isEmpty()) return "Cannot place order";
            if (pick.get().take(qty)) {
                record(u, pick.get(), qty);
                return "Order Placed Successfully from " + pick.get().name() + ".";
            }
            canFulfil = canFulfil.stream().filter(r -> r != pick.get()).toList();
        }
        return "Cannot place order";
    }
```

Full runnable code: code/lld/ch19/foodkart/ (see the Full Edition).

```
show_restaurant(price):
  Food Court-2, Burger
  Food Court-1, NI Thali
place_order(Food Court-1, 2): Order Placed Successfully.
place_order(Food Court-2, 7): Cannot place order
highest rating: Order Placed Successfully from Food Court-2.
lowest price  : Order Placed Successfully from Food Court-4.
```

The sample block reproduces every expected line (price-sorted Burger 120 before Thali 100; owner updates to `Burger - 8`). Extras show owner-only updates, rating validation, a comment-less rating and a non-serviceable pincode. The variant: the two strategies choose **different** restaurants, Food Court-1 (3 plates) is never picked for 4, and once nobody has 4 left the order is cleanly refused.

Concurrency & edge cases

Case	Handling
Two claims for one rider	`tryClaim()` CAS: one wins, the loser tries the next ranked rider
Two dispatch loops	`dispatchWaiting` is `synchronized` (FIFO kept)
Kitchen over-booked	`tryReserveKitchen` `synchronized`
Cancel races accept	Both lock the `Order`; accept re-checks PENDING; slots released only if reserved
Item sold out after adding	`placeOrder` re-checks `available`
Restaurant closes	Hidden in `browse`; `placeOrder` refuses
COD cancelled	No refund
FoodKart: no login / duplicates / rating out of 1–5	`requireLogin()` throws; `putIfAbsent` rejects; `Review` compact constructor rejects
FoodKart: capacity taken after selection	`take()` fails; restaurant dropped; strategy asked again

Real-world example

Swiggy and Zomato show the same states (placed, confirmed, preparing, rider assigned/picked up, delivered) and ask you to **reset your cart** when adding from a second restaurant. Rider dispatch uses distance, load and fairness and may batch two orders — variations of `DeliveryAssignmentStrategy`. FoodKart is a Flipkart 90-minute, in-memory, driver-class exercise rewarding clean separation over clever algorithms.

Extensions & follow-up questions

- **Batching** — a strategy accepting a rider already carrying one order when the second pickup is on the way (agent capacity 2).
- **ETA** — an `EtaEstimator`: prep time by load + travel time.
- **Surge / distance fees** — a `DeliveryFeePolicy` Strategy.
- **Rider rejects** — rider back to available, order to the queue head, rider excluded for it.
- **Multi-dish FoodKart** — `Map<String, MenuItem>`; `take` per item, all-or-nothing.
- **Order split across restaurants** — the strategy returns a set covering all items; reserve in each, roll back on failure.
- **Live tracking** — rider location events forwarded by an Observer over WebSocket.

In the interview

Design round (45 min): clarify (5) → entities + one-restaurant cart (8) → **state machine** (7) → **assignment strategy + CAS claim + waiting queue** (12) → payments/notifications (5) → concurrency and extensions (8). **FoodKart (90 min):** model classes → one service method per command → driver replaying the sample → `SortBy` → owner checks and validation → selection strategy → history; runnable at every step, exact output lines.

> **30-second answer:** "Restaurants own their menus and a kitchen capacity reserved on accept and released when food is ready. Orders copy item prices and move through a state machine from PENDING to DELIVERED. When food is ready the order joins a FIFO queue; a pluggable assignment strategy ranks free riders and we claim one with a compare-and-set, so a rider can never get two orders. Every status change is published to observers for the customer, restaurant and rider apps."

- *Why release capacity on ready?* — The kitchen is free once food leaves; delivery is the rider's constraint.
- *Why does the strategy return a list?* — The top rider may be claimed between ranking and claiming; fall through without re-ranking.
- *Scale across cities?* — Partition by zone; one dispatcher (single writer) per zone.

Self-check

1. **Q:** Which states allow a customer cancel? **A:** PENDING and CONFIRMED — after that the food cost is spent.

2. **Q:** What stops two orders getting one rider? **A:** `AtomicBoolean.compareAndSet(true, false)` in `tryClaim()`.
3. **Q:** FoodKart "serviceable"? **A:** Delivers to the user's pincode and quantity > 0.
4. **Q:** Why is Burger (120) before NI Thali (100) by price? **A:** The statement asks for descending order.
5. **Q:** Patterns used? **A:** Strategy (assignment, selection, payment), Observer, state machine, Facade.

TL;DR

- One-restaurant cart; order copies prices; **state machine** PENDING → … → DELIVERED.
- **Kitchen capacity** reserved on accept, released on ready; **riders** claimed with **CAS**, chosen by a swappable **strategy**, FIFO queue when none are free.
- FoodKart: pincode serviceability, one dish per restaurant, average rating, `SortBy` enum, atomic `take`, capacity-filtered **selection strategy**.

Design a Grocery Store (POS & Checkout)

A **point-of-sale (POS) system** runs a supermarket till: scan, running total, offers, payment in one or more ways, receipt, stock down. Three things make it different from the Amazon cart: **weighed items** (produce priced per kg from the scale), **interacting offers** (buy-2-get-1, "3 for ₹80", % off produce, a basket coupon — the combining rule decides price, tax and refunds), and **the till never blocks a physical sale** (if the system says 0 but the customer holds the carton, the count is wrong). Money: rupees in `BigDecimal` (`Rs` in the output).

Problem & clarifying questions

Question	Assumed answer
Item identification?	Barcode (EAN-13) for packaged goods; 4-digit **PLU** + scale weight for produce.
Tax?	Shelf prices tax-inclusive (India/UK/EU); support US-style tax-on-top too.
Offers and stacking?	One item offer per product — the **best** for the customer; then one basket coupon on what is still owed.
Payment?	Cash (change), card, UPI; **split tender**.
Changes of mind?	Void a line; back to scanning if nothing paid; cancel (reverse card/UPI).
Returns?	With the receipt; refund the **net paid**, never more units than bought.
Stock?	Down on sale, up on return; alerts; never refuse a sale because of the count.

Requirements

Functional: lookup by barcode/PLU, weighed pricing; sale with voidable lines; item offers + basket coupons with stacking rules; tax per category (inclusive/exclusive); split payment; immutable receipt; returns against a receipt; stock updates with low/negative alerts. **Non-functional:** exact money (`BigDecimal`, `HALF_UP`, parts add up); several registers share catalog/stock/receipts safely; deterministic pricing; auditable voids and scan-time prices.

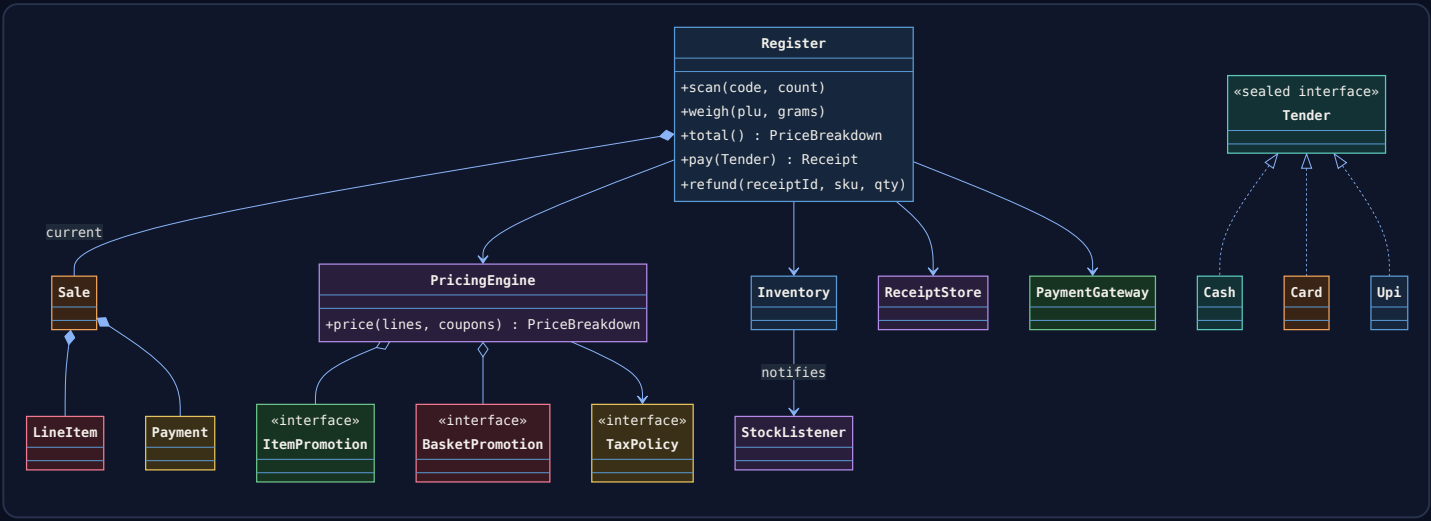
Use cases

Cashier / self-checkout (start, scan, weigh, void, coupon, total, pay, cancel), **Customer** (pay, return), **Supervisor** (approve voids/returns, reprice), **Payment network** (authorise, reverse), **Purchasing** (stock alerts).

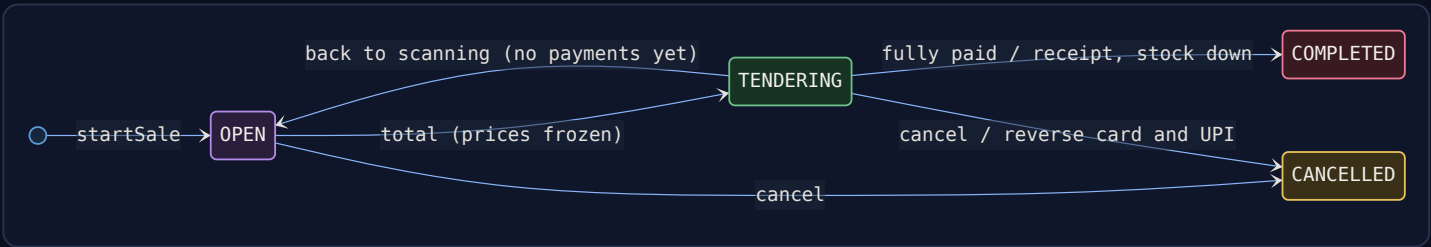
Core entities

Class	Responsibility
`Product`, `Catalog`	SKU, code, department, EACH/PER_KG, price, `TaxCategory`; lookup; reprice replaces the record.
`LineItem`, `Sale` + `SaleStatus`	One scan (price copied in); OPEN → TENDERING → COMPLETED / CANCELLED.
`ItemPromotion` (`BuyXGetYFree`, `MultiBuy`, `PercentOffDepartment`), `BasketPromotion` (`SpendThresholdCoupon`)	**Strategy** for offers.
`TaxPolicy` (`TaxInclusive`, `TaxExclusive`)	**Strategy** for tax.
`PricingEngine` → `PriceBreakdown`	Pure function: lines + coupons → discounts, net per SKU, tax, total.
`Tender` (sealed: `Cash`, `Card`, `Upi`), `Payment`, `PaymentGateway`	Tenders; applied amount; port to card/UPI networks.
`Receipt`, `ReceiptStore`	Immutable proof; atomic returns ledger.
`Inventory`, `StockListener`	Shared stock; **Observer** alerts.
`Register`	**Facade** for one lane.

Class diagram



Key flows



Pricing one basket (the hard method): skip voids → group by SKU → each SKU gets its **best** item offer → net per SKU → best basket coupon on the sum → **spread** the coupon over SKUs in proportion (remainder to the largest) → tax per category on discounted amounts → total (+ tax if added on top).

Split payment: `total()` freezes ₹588.23 → declined card throws, nothing recorded → UPI ₹400 → cash ₹500 → change ₹311.77 → receipt stored, stock reduced.

Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
Offers behind `ItemPromotion` / `BasketPromotion`	**Strategy**	New offer = new class	`switch` over offer types
Best offer per SKU, coupon after	Explicit stacking policy	Written once, explainable at the till	Apply all in list order (order-dependent, can go negative)
Pure `PricingEngine`	Stateless service	Same basket, same price; shared by all registers	Mutating discounts into lines while scanning
Freeze at `total()`	Snapshot	Price/offer changes mid-payment can't change the amount due	Re-price per payment
Spread the coupon over SKUs	Exact allocation	Tax per category and per-item refunds need it	One unallocated number
`TaxPolicy`	**Strategy**	Inclusive vs exclusive per store	Two engines
Sealed `Tender`, `Register` facade	Sealed types, **Facade**	Compiler-checked receipt switch; UI sees one API	Strings; UI calling every part
Stock may go negative; edge-triggered alerts	**Observer**	Shelf is the truth; one alert per crossing	Refusing the sale; alert spam

Java implementation

Package `lld.ch19.grocery`.

```
// File: lld/ch19/grocery/Pricing.java
    PriceBreakdown price(List<LineItem> lines, Set<String> coupons) {
        // ... group non-voided lines by SKU into SkuTotal(product, quantity, gross)
        for (SkuTotal t : bySku.values()) {
            Optional<Discount> best = itemOffers.stream()
                    .map(p -> p.evaluate(t.product(), t.quantity()))
                    .flatMap(Optional::stream)
                    .max(Comparator.comparing(Discount::amount));
            BigDecimal off = best.map(d -> d.amount().min(t.gross())).orElse(Money.ZERO);
            best.ifPresent(d -> discounts.add(new Discount(d.description(), d.sku(), off)));
            net.put(t.product().sku(), t.gross().subtract(off));
        }
        BigDecimal subtotal = sum(net.values());
        Optional<Discount> basket = basketOffers.stream()
                .map(b -> b.evaluate(subtotal, coupons))
                .flatMap(Optional::stream)
                .max(Comparator.comparing(Discount::amount));
        basket.ifPresent(d -> { discounts.add(d); spread(d.amount(), net); });
        // ... tax per TaxCategory on the discounted nets via taxPolicy; total = nets (+ tax if added on top)
    }
```

```
// File: lld/ch19/grocery/Register.java
    /** delta < 0 for a sale, > 0 for a return. merge() is atomic per key, so registers never lose updates. */
    void adjust(String sku, BigDecimal delta) {
        BigDecimal after = onHand.merge(sku, delta, BigDecimal::add);
        BigDecimal before = after.subtract(delta);
        BigDecimal point = reorderPoint.getOrDefault(sku, BigDecimal.ZERO);
        // edge-triggered: alert only when THIS change crosses the line, so each crossing alerts exactly once
        if (before.compareTo(point) > 0 && after.compareTo(point) <= 0)
            listeners.forEach(l -> l.stockAlert(sku, after, "LOW - reorder"));
        if (before.signum() >= 0 && after.signum() < 0)
            listeners.forEach(l -> l.stockAlert(sku, after, "NEGATIVE - recount the shelf"));
    }
```

Full runnable code: code/lld/ch19/grocery/ (see the Full Edition).

How it works: `scan` checks the EAN-13 check digit before lookup; `weigh` turns grams into exact kg (`movePointLeft(3)`). Cash above the amount due becomes change; card/UPI must be authorised and cannot exceed

the amount due. Returns: `ReceiptStore.claimReturn` uses `compute` (atomic check-and-add), and the refund is `share(after) − share(before)` of the SKU's net paid, so partial returns add up exactly.

Output (excerpt):

```
  error: Bad barcode 8901000000018: check digit wrong, rescan
   10% off PRODUCE: Bananas                          -Rs 7.47
   Buy 2 get 1 free: Chips 50g                      -Rs 20.00
   3 for Rs 80.00: Biscuits                         -Rs 10.00
   Coupon SAVE50                                    -Rs 50.00
 TOTAL                                             Rs 588.23
   incl. tax STANDARD @18%                          Rs 40.07
 Change                                            Rs 311.77
  return 1 of 3 chips -> refund Rs 12.29
  return 2 more chips -> refund Rs 24.58  (all three add up to the Rs 36.87 paid)
  inclusive total Rs 588.23 vs exclusive total Rs 644.73 (tax {EXEMPT=0.00, REDUCED=9.22, STANDARD=47.28})
  paneer on hand: -21  (sold 120 + 1 earlier; the till never refuses a real item)
  alert: PNR LOW - reorder (on hand 20)
  alert: PNR NEGATIVE - recount the shelf (on hand -1)
```

Chips qualified for both "buy 2 get 1" (₹20) and "3 for ₹50" (₹10): only the better applied. Returns refund the net paid, not the ₹20 shelf price. Four concurrent registers took paneer below zero without refusing a sale; each alert fired once.

Concurrency & edge cases

Case	Handling
Registers sell the same SKU at once	`ConcurrentHashMap.merge` (atomic per key)
Count says 0, item in hand	Sell; stock negative; one recount alert
Two desks refund the same line	`compute` in `claimReturn`; second over-claim refused
Price/offer change mid-sale	Line keeps scan price; breakdown frozen at `total()`
Card declined / cancel after card	Nothing recorded / non-cash payments reversed
Rounding in coupon split and refunds	Remainder to the largest SKU; refunds as differences of cumulative shares
One register, one cashier	`Register` / `Sale` single-threaded; shared parts are the thread-safe ones

Real-world example

Tills key produce by **PLU** codes (4011 = yellow bananas) and read an integrated scale; EAN-13 barcodes end in a check digit, so a smudged code is rejected instead of misread; tax-inclusive receipts (GST in India, VAT in the UK) show the tax contained, US receipts add sales tax — our two `TaxPolicy` strategies; cycle counts fix negative stock.

Extensions & follow-up questions

- **Loyalty / member prices** → a `CustomerContext` for the engine; points on the net paid.
- **Pre-packed weighed labels** → variable-measure barcodes decoded into a line.
- **Bundles (pasta + sauce)** → multi-SKU promotion; best-per-SKU becomes a small optimiser.
- **Refund to original tender** → keep tender splits on the receipt.
- **Offline mode** → queue authorisations under a floor limit; cache catalog and offers per register.
- **End of day** → `Shift` with float and expected cash.

In the interview

45 minutes: clarify (5) → entities + facade (8) → sale state machine (5) → **pricing engine and stacking** (12) → split payment + receipt (7) → returns, stock, concurrency (8).

30-second answer: "A Register facade drives one Sale through OPEN, TENDERING, COMPLETED or CANCELLED; lines copy the price at scan; produce is priced by weight. A pure pricing engine gives each SKU its single best offer, applies a basket coupon, spreads it over SKUs so tax and refunds are exact, and computes tax through an inclusive or exclusive policy. Prices freeze at total; tenders are sealed; cash

gives change; card and UPI go through a gateway. On completion an immutable receipt is stored and shared stock drops atomically with edge-triggered alerts — the till never refuses a real item."

- *Why spread the coupon?* Tax differs per category and returns must refund what was paid for that item.
- *Why stateless pricing?* Shared by all registers; reproducible for audits and returns.

Self-check

1. **Q:** Two offers match the chips — which applies? **A:** The larger discount only.
2. **Q:** Why is a one-chip refund ₹12.29? **A:** Net paid after the offer and the coupon share.
3. **Q:** What stops a double refund? **A:** `compute` makes check-and-add atomic.
4. **Q:** Why allow negative stock? **A:** The item is physically there; the count is wrong.

TL;DR

- **Register facade** → **Sale** state machine; scan-time prices; produce by **weight**.
- **Pure pricing engine**: best offer per SKU → coupon → spread → tax via **TaxPolicy**.
- Freeze at total; **split tender**; returns refund the **net paid**, claimed atomically; stock atomic, may go negative, **edge-triggered** alerts.

Chapter summary

Problem	Hard part	The move that solves it
Amazon	Guests vs members; last-item races; cancel rules	Ordering only on `Member` (LSP); atomic all-or-nothing reservation; order state machine
Brokerage	Four order types; tax lots; over-spending	Two abstract questions on `Order` ; `StockLot` + FIFO `LotSale` ; reserved funds
Stock exchange	Fair matching	Per-symbol `TreeMap<price, FIFO queue>` ; price then time; one lock per symbol
Auction	Concurrent bids	All bid rules in one per-listing critical section; notify after unlocking
Music streaming	Shuffle/repeat without losing order	Strategy-produced permutation of indexes; plan rules in the enum
Food delivery / FoodKart	Capacity and riders under concurrency	Atomic kitchen reservation; CAS-claimed riders; strategies for assignment and selection
Grocery POS	Weighed items, stacking offers, split tender, returns	Pure pricing engine (best offer per SKU, coupon spread, tax per category); frozen price; atomic return claims; edge-triggered stock alerts

Recurring patterns: Strategy (payments, play order, rider assignment, restaurant selection, trade price, grocery offers and tax), Observer (notifications, fills, plays), state machines in enums (orders, auctions, player), Template Method (statements), Facade (every `*Service`), and one concurrency rule — **make check-then-act atomic** (lock or compare-and-set).

Sources: DesignGurus *Grokking the Object-Oriented Design Interview* (Amazon, stock brokerage); awesome-low-level-design (auction, music streaming, food delivery); kumaransg/LLD machine-coding statements (Navi stock exchange, Flipkart FoodKart); grocery POS designed from first principles (GS1 EAN-13, IFPS PLU codes). All Java code is original.

Chapter 20 — Case Studies: Pattern-Driven Components (Study Edition)

Goal. These problems are not whole products but **building blocks** most systems contain: a ledger that splits money, a notification sender, a cache, a rider–driver matcher, a logger, a message broker and a rate limiter. Each is small enough to finish in an interview and each is **anchored on one or two design patterns**. Name the pattern and the change it absorbs, and half the question is answered.

Why components get their own chapter. A component has a small public surface (a few methods) and many hidden rules. The interviewer keeps asking "what if this rule changes?"; a good design answers "one new class, nothing else touched".

#	Component	Anchor pattern(s)	The change the pattern absorbs
20.1	Splitwise (expense sharing)	**Strategy** (split rules)	A new way to split a bill
20.2	Notification service	**Template Method** (retry loop), **Strategy** (send rules), **Observer** (outcomes), Simple Factory (channels)	A new channel, a new reason not to send, a new watcher of results
20.3	In-memory cache (+ LRU from scratch)	**Strategy** (eviction), **Observer** (removal listener), Builder	A new eviction rule
20.4	Ride sharing (Uber)	**Strategy** (matching, pricing, payment), **Decorator** (surge), State machine (trip)	A new matching rule, surge on/off, a new payment method
20.5	Logging framework	**Chain of Responsibility** (logger hierarchy), **Singleton** (LogManager), Strategy (formatter), Decorator (async appender)	A new sink, a new format, async writing
20.6	Pub-Sub system	**Observer / Publish-Subscribe**, append-only log with offsets	New subscribers, replay, slow consumers
20.7	Rate limiter	**Strategy** (algorithm), Simple Factory, Composite	A new algorithm, a new tier, combined limits

Each case follows the Chapters 14–21 template (problem → requirements → entities → diagrams → flows → decisions → code → edge cases → interview).

Running the code. Java 21, standard library only, under `code/lld/ch20/<name>/`, each package with a `Main`. From `chapters/20-case-components/code`:

```
javac --release 21 -d /tmp/lld_20 $(find . -name '*.java')
java -cp /tmp/lld_20 lld.ch20.splitwise.Main      # also: notification, cache, lrucache,
                                                  # ridesharing, carpool, logging, pubsub, ratelimiter
```

Money is always `BigDecimal`. A `double` cannot hold 0.1 exactly (`0.1 + 0.2` prints `0.30000000000000004`). Two rules: compare with `compareTo` (not `equals`, which also compares scale, so `2.0` ≠ `2.00`), and always give a rounding mode when dividing (`divide(x, 2, RoundingMode.HALF_EVEN)`), or 1/3 throws `ArithmeticException`.

Reading the course diagrams. Figures marked "(course diagram)" use one arrow legend throughout:

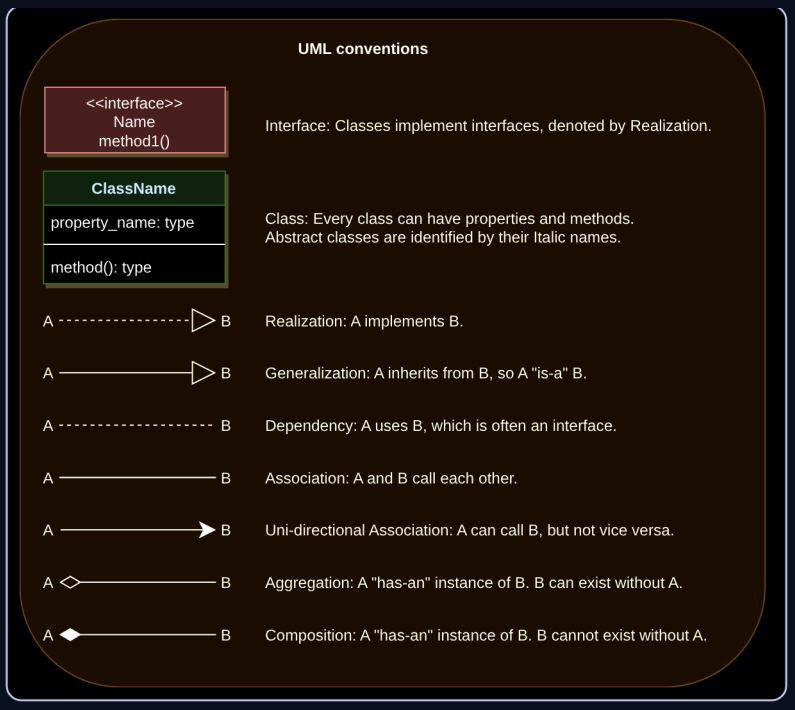


Figure — UML conventions used in the course class diagrams: interface, class, realization (implements), generalization (inherits), dependency, association, one-way association, aggregation and composition (course diagram)

The mermaid diagrams are drawn from **this chapter's Java code**.

20.1 Design Splitwise

Problem & clarifying questions

What it is. Splitwise remembers who paid for what when friends share costs and tells everyone who owes whom — the shared notebook for a weekend trip where Alice paid the hotel, Bob dinner, Carol the cab. The hard part is **not** the UI: it is deciding *which class owns the money arithmetic*, so the books balance after every operation.

Question	Assumed answer
Which split types?	Equal, exact, percentage. More later.
Must shares add up exactly?	Yes; otherwise the expense is rejected and no balance changes.
Rounding?	2-decimal money; no cent may be lost.
Pay only for yourself?	No: at least two people (payer counts) must be involved.
Settle up / simplify debts?	Yes / yes (on request, fewer transfers).
Currencies, receipts, undo?	Out of scope (say so).

Requirements

Functional: everything in the table, plus: users (name, email), groups, expense history per group, balance with one person and overall net. Invalid expense → **nothing changes**.

Non-functional: money correctness above all; a new split type changes no existing class (OCP); thread-safe.

Use cases

Actors: **Member** (adds expenses, settles up, asks what they owe), **Group creator**, **System** (computes shares, validates, updates the ledger, simplifies). *Add an expense* **includes** (always performs) *work out each share*, *check the shares add up*, *update the ledger*.

Core entities

Class	Responsibility
`User` (record)	Identity only: id, name, email.
`Group`	Name + member ids + expense ids. **Holds no balances.**
`SplitStrategy` + `EqualSplit`, `ExactSplit`, `PercentSplit`	Turn a total into one `Share` per person.
`Share` (record)	(person, amount) value object.
`Expense`	Description, total, payer, shares; **validates itself once** at creation.
`Ledger`	**One signed number per pair** of people; one mutator `adjust`.
`DebtSimplifier` / `Payment`	Greedy minimum-cash-flow → list of suggested transfers.
`Splitwise`	Facade; orders the steps, contains no arithmetic.
`Money`	2-decimal `BigDecimal` helpers.

Why the group holds no money: a person in two groups would have "Bob owes Alice" stored twice, and the copies would drift. One ledger = single source of truth. **`Group → User` is association, not composition:** deleting a group must not delete its people.

Class diagram

The course builds it in five stages, each leaving a working slice (copy this order in an interview): (1) `User`, `Group`; (2) `SplitStrategy` + rules + `Share`; (3) self-checking `Expense`; (4) `Ledger`; (5) `Splitwise` facade + `DebtSimplifier`.

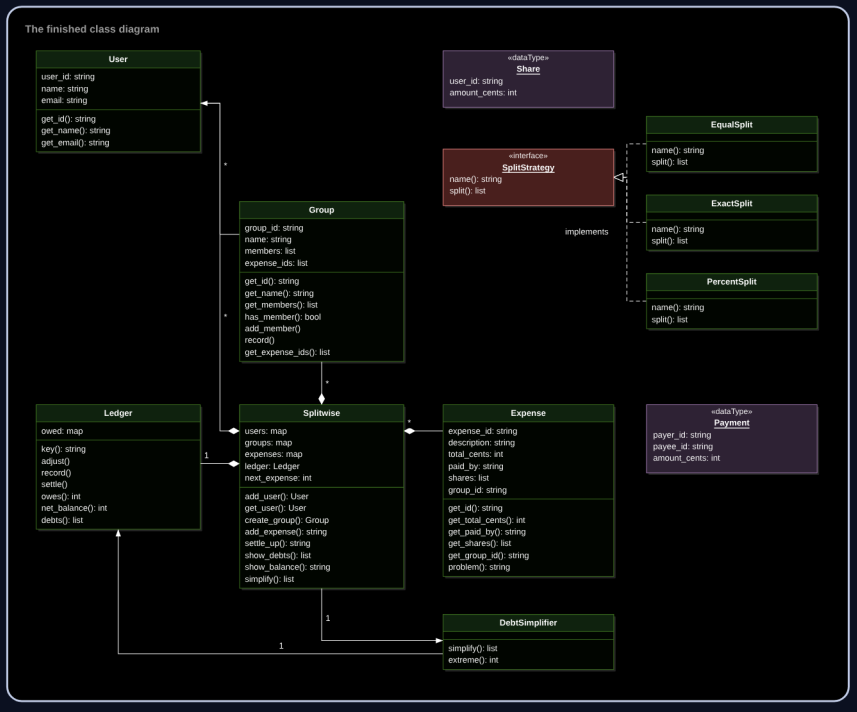
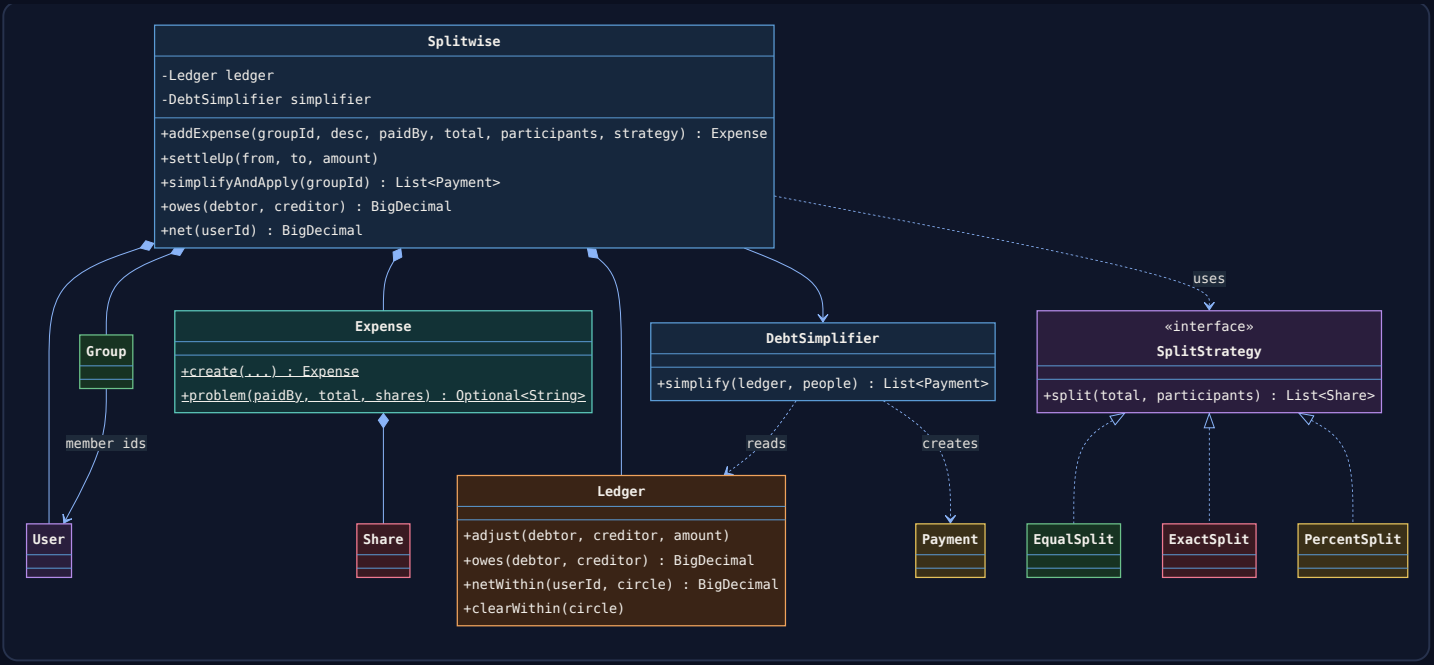


Figure — The finished Splitwise class diagram: service owning users, groups, expenses and the ledger, with the split strategies and the simplifier (course diagram)



Key flows

Adding an expense: strategy computes shares → `Expense.create` validates (≥ 2 people, sum = total) → only then, for each share, `ledger.adjust(sharer, payer, amount)`. If validation fails, an `IllegalArgumentException` with the reason is thrown and nothing is recorded.

Simplifying debts (greedy minimum cash flow):

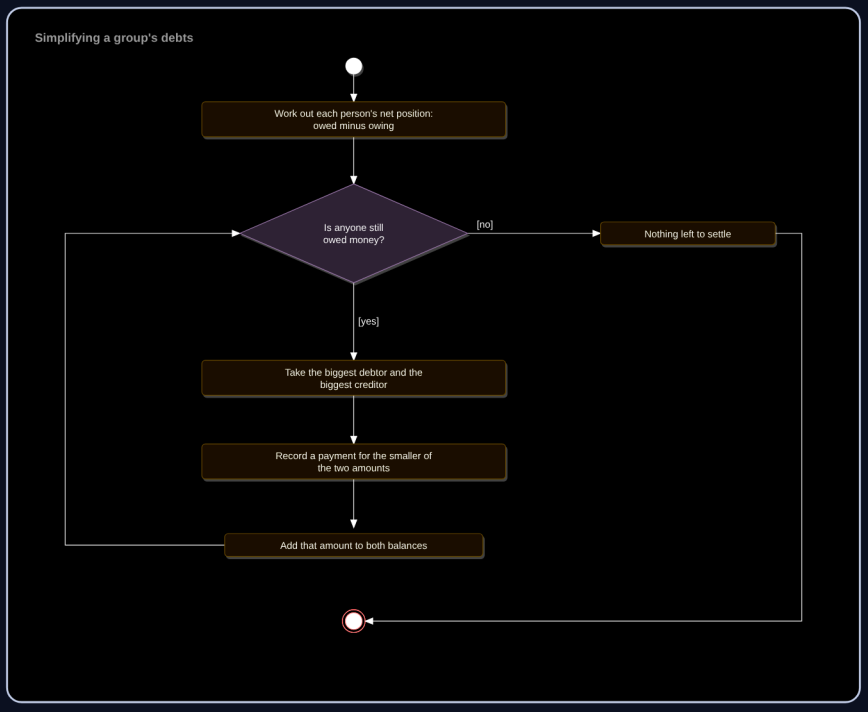


Figure — Simplifying debts: compute net positions, then the largest debtor repeatedly pays the largest creditor until every balance is zero (course diagram)

1. Net position per member = owed to them − they owe. Positive = creditor; nets sum to zero.
2. Creditors in one max-heap, debtors in another (ties by name → deterministic).
3. Pop biggest creditor C and debtor D; x = min(C, D); suggest "D pays C x".
4. Push back whichever side has a remainder; stop when a heap is empty.

Each round zeroes at least one person → **at most n − 1 payments**. It is **not always the true minimum** (that is NP-hard — it means finding the most zero-sum subgroups). Say so before you are asked.

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy for split rules	A new rule = a new class; all return `List<Share>`, so nothing downstream knows which ran.	`if (type == EQUAL) … else if …` in the service.
Expense validates in a **static factory**	Checked once; an invalid expense can never exist.	Checking in the ledger on every read.
`problem()` returns the **reason**	User sees *why* ("shares add up to 15.00 but the total is 30.00").	`boolean isValid()`.
One signed number per unordered pair	A debt can't disagree with itself; reverse debts cancel automatically.	Two maps "A owes B" / "B owes A".
One mutator `adjust`	Expenses, settle-ups, simplification share one cancel rule.	Separate `addDebt`, `settle`, `reduce`.
`DebtSimplifier` separate	SRP: keeping balances vs planning settlements.	Algorithm inside the ledger.
`BigDecimal`, spare cents handed out	100.00/3 → 33.34 + 33.33 + 33.33; percent split: last person absorbs remainder.	`double` (the course's whole-cents `long` is equally correct).
`Splitwise` is a **Facade**	One entry point that orders steps and decides nothing.	UI calling ledger/strategies directly.

Only one GoF pattern (Strategy) — fine: good designs use the pattern that absorbs the likely change, not a collection.

Java implementation

Key excerpts (the split strategy, the self-validating expense, the ledger's single mutator, and the facade's ordering):

```java
// File: code/lld/ch20/splitwise/SplitStrategy.java
interface SplitStrategy {
    List<Share> split(BigDecimal total, List<String> participants);
}

/** Equal split. Spare cents are handed out one at a time, in participant order. */
final class EqualSplit implements SplitStrategy {
    @Override
    public List<Share> split(BigDecimal total, List<String> participants) {
        if (participants.isEmpty()) throw new IllegalArgumentException("nobody to split between");
        BigDecimal n = BigDecimal.valueOf(participants.size());
        BigDecimal base = total.divide(n, 2, RoundingMode.DOWN);          // 100.00/3 -> 33.33
        BigDecimal leftover = total.subtract(base.multiply(n));           // 0.01
        int spareCents = leftover.divide(Money.CENT).intValueExact();     // 1
        List<Share> shares = new ArrayList<>();
        for (int i = 0; i < participants.size(); i++) {
            BigDecimal amount = i < spareCents ? base.add(Money.CENT) : base;
            shares.add(new Share(participants.get(i), amount));
        }
        return shares;
    }
}
// ... ExactSplit, PercentSplit (last person absorbs the remainder)

// File: code/lld/ch20/splitwise/Expense.java
    static Optional<String> problem(String paidBy, BigDecimal total, List<Share> shares) {
        if (total.signum() <= 0) return Optional.of("total must be positive");
        // ... negative / duplicate share checks
        Set<String> people = new HashSet<>(sharers);
        people.add(paidBy); // the payer counts as one of the two people even if they owe no share
        if (people.size() < 2) return Optional.of("an expense needs at least two people");
        BigDecimal sum = shares.stream().map(Share::amount).reduce(BigDecimal.ZERO, BigDecimal::add);
        if (sum.compareTo(total) != 0)
            return Optional.of("shares add up to " + sum + " but the total is " + total);
        return Optional.empty();
    }

// File: code/lld/ch20/splitwise/Ledger.java  (key (lo, hi), lo < hi; positive = "hi owes lo")
    /** The ONLY mutator: 'debtor' now owes 'creditor' 'amount' more. Opposite debts cancel. */
    void adjust(String debtor, String creditor, BigDecimal amount) {
        if (debtor.equals(creditor) || amount.signum() == 0) return;
        Pair key = Pair.of(debtor, creditor);
        BigDecimal delta = creditor.equals(key.lo()) ? amount : amount.negate();
        BigDecimal updated = balances.getOrDefault(key, Money.ZERO).add(delta);
        if (updated.signum() == 0) balances.remove(key); else balances.put(key, updated);
    }

// File: code/lld/ch20/splitwise/Splitwise.java
        List<Share> shares = strategy.split(total, participants);       // may throw
        Expense e = Expense.create("e" + ids.incrementAndGet(), description, paidBy, total, shares, groupId);
        for (Share s : e.shares()) ledger.adjust(s.userId(), paidBy, s.amount()); // payer's own share is a no-op
    // ... settleUp: refuse paying more than is owed, then ledger.adjust(to, from, amount) - the reverse debt cancels
```

Full runnable code: code/lld/ch20/splitwise/ (see the Full Edition).

Notes: `Money.of("10.005")` throws (`RoundingMode.UNNECESSARY` refuses silent rounding); `PercentSplit` rounds with `HALF_EVEN` (banker's rounding); `DebtSimplifier` uses two `PriorityQueue` s.

Output (excerpt):

```
Hotel shares: [alice=33.34, bob=33.33, carol=33.33]
  refused: shares add up to 15.00 but the total is 30.00
  suggested: dave pays alice 16.68
  suggested: carol pays alice 13.33
  suggested: bob pays alice 9.99
```

Four pairwise debts became three payments; every net (alice +40.00, bob −9.99, carol −13.33, dave −16.68) is unchanged.

Related machine-coding problem: The Ledger Co. (Navi)

The Geektrust/Navi round "Ledger Co." also keeps money records, but for **bank loans** (simple interest, monthly EMI, lump-sum prepayments). It is built in **Chapter 21 (The Ledger Co.)**. Both keep **one authoritative record** and compute every answer from it, and both use `BigDecimal` with a named rounding mode; the rules differ (*who owes whom* vs *how much is repaid by month N*).

Concurrency & edge cases

- **Concurrent expenses:** every public `Splitwise` method is `synchronized` (build → validate → adjust is atomic); at scale lock per group/pair or use a DB transaction.
- **Payer not a sharer** (Alice buys Bob's lunch) is valid; the payer's own share is a no-op `adjust`.
- **Overpaying on settle-up**, unknown users, duplicate/negative shares, zero totals: refused before anything changes.
- **Simplify touches other groups:** one number per *pair* means simplifying a group rewrites all debts between its members, even ones from another group. If that is unacceptable, keep per-group ledgers plus a derived global view.

Extensions & follow-up questions

- **Split by shares (2:1:1)?** New `ShareRatioSplit`; nothing else changes.
- **Delete/edit an expense?** Re-add the inverse, or better make the ledger an **append-only journal** (event sourcing): delete = append a reversing entry.
- **Is greedy optimal?** No. Nets A −6, B −5, C +2, D +4, E +5: greedy needs 4 payments; optimal 3 (B→E 5, A→D 4, A→C 2).
- **Is `Splitwise` doing too much?** At scale, yes: give users, groups, expenses and ledger their own repositories. Multiple currencies: `Money(amount, currency)` + one ledger per currency.

Real-world example

Splitwise really offers "simplify debts" for groups (the net-position idea). Bill-splitting in payment apps solves rounding the same way: one person absorbs the leftover cent.

In the interview

30-second answer: "Users and groups hold no money. Each split rule is a `SplitStrategy` returning shares. An `Expense` validates itself once: at least two people, shares add to the total. The `Ledger` keeps one signed number per pair and has a single `adjust`, so debts cancel automatically; settling up is an adjust the other way. Simplification: net positions + greedy biggest-debtor-pays-biggest-creditor, at most n − 1 payments, not always optimal."

Tip: say "one class per split rule" early; spend your time on the ledger.

Self-check

1. Why does `Group` hold no balances? → Each debt must live in one place; two copies drift.
2. Split 10.00 equally among 3. → 3.34, 3.33, 3.33.
3. Max payments from the greedy simplifier for n people? → n − 1.
4. Why `compareTo` for `BigDecimal`? → `equals` compares scale: `2.0` ≠ `2.00`.
5. Alice pays 20.00, Bob owes all of it — valid? → Yes, two people are involved.

TL;DR

- **Strategy** for split rules; **Expense validates itself** once.
- **Ledger = one signed number per pair**, one mutator; settle-up = reverse adjust.
- **Simplify** = net positions + greedy; ≤ n − 1 payments; optimal is NP-hard.
- **BigDecimal**, explicit rounding, no lost cents.

20.2 Design a Notification Service

Problem & clarifying questions

What it is. The component that sends emails, SMS and push alerts ("Your order has shipped", "482913 is your login code"). Analogy: a post-office counter that picks the delivery method, respects "do not disturb", retries if the courier fails, and keeps a receipt. **Why it is a design question:** it starts as one `sendEmail()` and grows fourteen `if`s (SMS, quiet hours, promo limits…). The exercise is *where each addition should have gone instead*.

Question	Assumed answer
Channels?	Email, SMS, push; more later.
Who picks the channel?	The user, **per kind** of notification, in preference order.
Wording?	Templates with placeholders: `Hi {name}, it arrives {eta}.`
Quiet hours?	Per user, may wrap midnight (22–07); urgent (OTP) ignores them.
Rate limits?	≤ N per user per kind per hour.
Failures?	Retry a channel a bounded number of times, then fall back to the next channel; all fail → FAILED.
Audit?	Every attempt recorded: SENT / FAILED / SUPPRESSED + reason.
Sync or async?	Caller must not wait for a slow gateway.

Requirements

Functional: the table above, plus: user has email, phone, device token; FAILED when no channel succeeds; **a new channel or a new reason to hold back needs no edit to the service.**

Non-functional: caller not blocked (async); thread-safe; testable without real time (injected clock); a broken listener never breaks delivery.

Use cases

Actors: **Product code** (asks), **Person** (channels, quiet hours), **Operator** (adds channels, watches outcomes), **Service** (renders, checks rules, tries channels, retries, records).

Core entities

Class	Responsibility
`User` / `Preferences`	Who someone is vs what they accept (ordered channels per kind, quiet hours with wrap logic in **one** place). Separate because they change for different reasons (SRP).
`Template` / `Message`	Subject + body with `{holes}` ; `render(values)` → identical `Message` for every channel.
`NotificationRequest`	kind, user, values, priority (NORMAL / URGENT).
`Gateway`	External provider (SMTP, SMS gateway, push service); swappable and fakeable.
`NotificationChannel` (abstract) + Email/Sms/Push	**Template Method**: final `deliver()` has the retry-with-backoff loop; subclasses give `addressOf(user)` .
`ChannelFactory`	**Simple Factory**: type → channel.
`RetryPolicy`	Max attempts + initial backoff (doubles).
`SendRule` + `QuietHoursRule` , `RateLimitRule`	**Strategy**: one reason not to send, returned as text.
`DeliveryRecord` / `DeliveryListener`	Outcome record; **Observer** (`LoggingListener` , `MetricsListener`).
`TimeSource`	Where "now" comes from (manual clock in tests).
`NotificationService`	Thin coordinator: rules → preferences → render → channels in order → `finish()` .

Class diagram

Course build stages: (1) `User` + `Preferences` ; (2) `Template` + `Message` + request; (3) channel hierarchy; (4) `SendRule` s; (5) service + listeners.

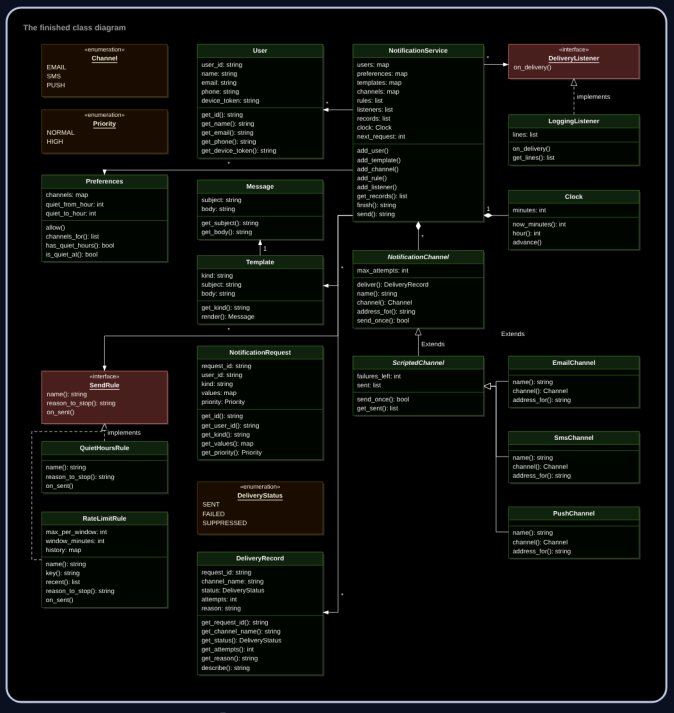
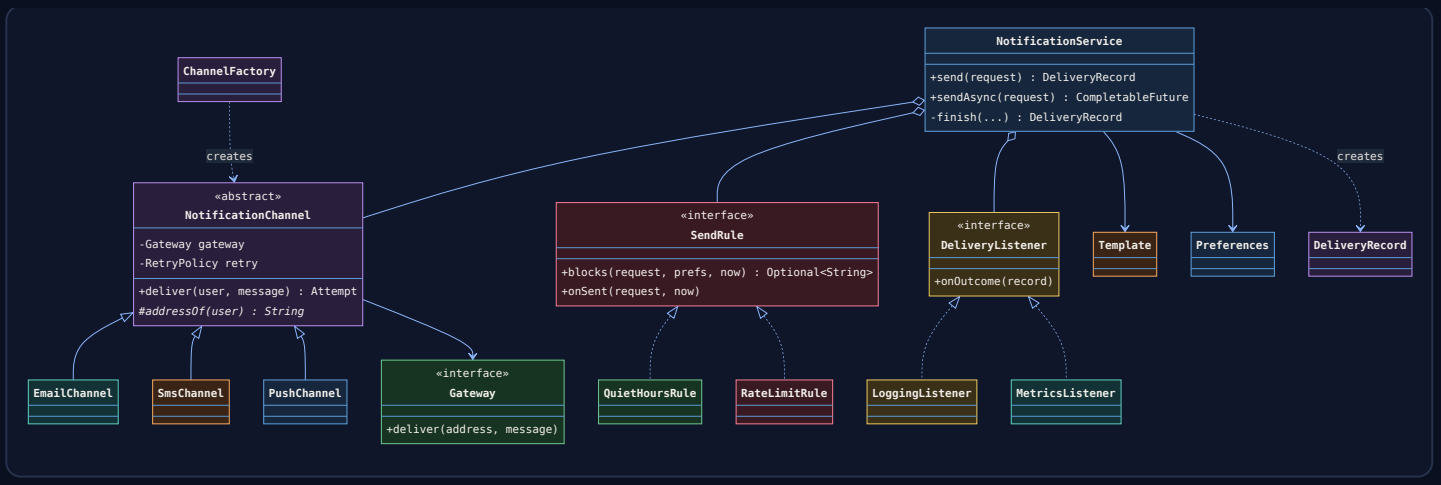


Figure — The finished notification service class diagram: the service in the middle, the channel hierarchy with the retry in the parent, the send rules, templates, preferences and the delivery listeners (course diagram)



Key flows

Sending one notification: ask every rule (any reason → SUPPRESSED); no accepted channel → SUPPRESSED; render the template **once**; try each accepted channel in the user's order (each with its own retry loop); first success → `onSent` to rules, SENT; all fail → FAILED. Every path exits through `finish()`.

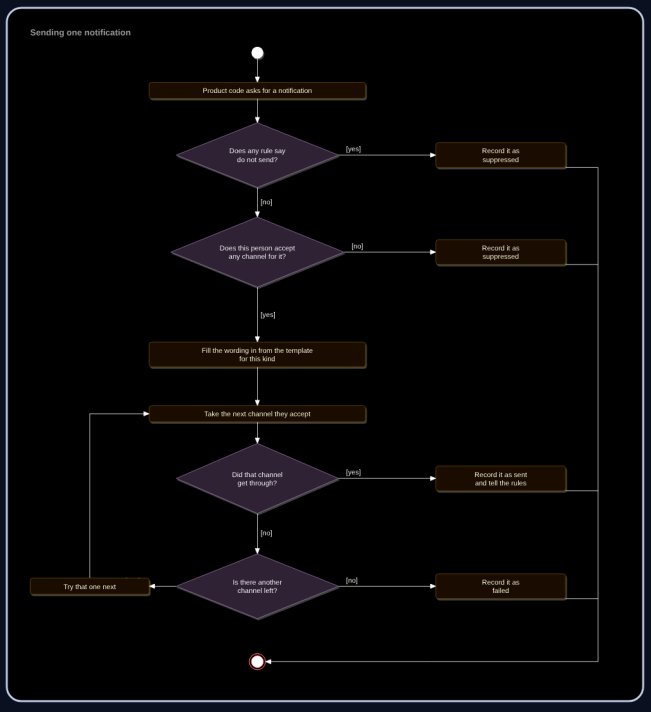


Figure — Sending one notification: ask every rule, render once, try each accepted channel in turn, and record sent / failed / suppressed (course diagram)

Design decisions & patterns used

Pattern / decision	Where	Why
Template Method	`final deliver()` ; subclasses fill `addressOf()`	Retry loop is identical for all channels; only a *step* varies. Strategy would replace the whole method.
Strategy	`SendRule`	Each "reason not to send" is a class; a new rule (unsubscribed, dedupe) = new class + one wiring line.
Observer	`DeliveryListener`	Logs, dashboards, alerts want outcomes; the service doesn't know who listens.
Simple Factory	`ChannelFactory`	One place maps type → class.
Adapter (target)	`Gateway`	One adapter per vendor SDK; tests use a fake.
Rule returns **reason text**	`SendRule.blocks`	Reason is recorded.
Urgency handled **inside** the rule	`QuietHoursRule`	Service needn't know which rules urgent ignores.
Rate limit counts only **successful** sends	`onSent` hook	Suppressed/failed messages don't use up the allowance.
Single exit `finish()`	Service	Every path stores a record and notifies listeners.
Render **once**	Service	Same text everywhere; wording change = one edit.

Three patterns, three *separate* axes of change: another channel (Template Method + Factory), another reason to hold back (Strategy), another interested party (Observer).

Java implementation

```java
// File: code/lld/ch20/notification/Channels.java
abstract class NotificationChannel {
    private final Gateway gateway;
    private final RetryPolicy retry;
    // ... constructor
    abstract ChannelType type();

    /** Hook: where this channel reaches the user, or null if the user has no such address. */
    protected abstract String addressOf(User user);

    final Attempt deliver(User user, Message message) {
        String address = addressOf(user);
        if (address == null || address.isBlank()) return new Attempt(false, 0, "no " + type() + " address");
        String lastError = null;
        long backoff = retry.initialBackoff().toMillis();
        for (int attempt = 1; attempt <= retry.maxAttempts(); attempt++) {
            try {
                gateway.deliver(address, message);
                return new Attempt(true, attempt, null);
            } catch (Exception e) {
                lastError = e.getMessage();
                if (attempt < retry.maxAttempts()) sleep(backoff);
                backoff *= 2;
            }
        }
        return new Attempt(false, retry.maxAttempts(), lastError);
    }
    // ...
}

final class SmsChannel extends NotificationChannel {
    SmsChannel(Gateway g, RetryPolicy r) { super(g, r); }
    ChannelType type() { return ChannelType.SMS; }
    protected String addressOf(User u) { return u.phone(); }
}

// File: code/lld/ch20/notification/NotificationService.java  (inside send)
        for (SendRule rule : rules) {
            Optional<String> reason = rule.blocks(req, prefs, now);
            if (reason.isPresent()) return finish(req, DeliveryStatus.SUPPRESSED, null, 0, reason.get());
        }
        // ... no accepted channel -> SUPPRESSED; render the template ONCE (missing value -> FAILED via finish)
        for (ChannelType type : accepted) {
            NotificationChannel channel = channels.get(type);
            if (channel == null) continue;
            Attempt a = channel.deliver(user, message);
            totalAttempts += a.attempts();
            if (a.success()) {
                rules.forEach(r -> r.onSent(req, now));
                return finish(req, DeliveryStatus.SENT, type, totalAttempts, null);
            }
            lastError = type + ": " + a.error();
        }
        return finish(req, DeliveryStatus.FAILED, null, totalAttempts, "all channels failed, last " + lastError);
```

Full runnable code: code/lld/ch20/notification/ (see the Full Edition).

Notes: `Preferences.isQuiet` handles wrap-around (start ≤ end → range; else "after start **or** before end"). A missing template value is an error, not a raw `{name}` in an inbox. Backoff doubles (the course retried immediately — a named weakness). `RateLimitRule` prunes old timestamps (the course's version leaked them). A throwing listener is caught in `finish()`. `sendAsync` uses a pool; production uses a durable queue.

Output (excerpt):

```
== 2. SMS fails 3 times -> retried, then falls back to EMAIL
    SMTP -> asha@mail.com : Order A-17 shipped | Hi Asha, it arrives Tuesday.
  [log] SENT       ORDER_SHIPPED user=u1 channel=EMAIL attempts=4
  [log] SUPPRESSED PROMO         user=u1 channel=- attempts=0 reason=quiet hours
  [log] SUPPRESSED PROMO         user=u1 channel=- attempts=0 reason=rate limit 2 per 60 min
```

Concurrency & edge cases

- Shared state: `ConcurrentHashMap` for users/prefs/templates, `CopyOnWriteArrayList` for listeners/history, rate-limit deque locked per key.
- **Check-then-act gap:** two concurrent sends can both pass `blocks()` before `onSent()`. Fix: reserve a slot atomically at check time (release on failure) or serialize per user.
- **Retries block worker threads** with `sleep`; real systems re-enqueue with a delay (retry topic).
- **Idempotency:** a retry after a timeout may deliver twice → pass an idempotency key; receivers tolerate duplicates.
- Missing address → 0-attempt failure, next channel. No template → FAILED; no accepted channel → SUPPRESSED.
- **Time zones:** quiet hours in the *user's* zone (`ZoneId` in preferences).

Extensions & follow-up questions

- **Add WhatsApp:** one subclass, one enum value, one `Gateway` adapter.
- **Too many messages:** a new `SendRule` (daily cap, dedupe), not a new `if`.
- **Campaign to a million users:** fan-out stage onto a queue + many workers calling `send`; rate-limit per provider.
- Priorities (separate queues), localization (templates by kind + locale), webhooks for receipts, a **circuit breaker** per provider.

Real-world example

Push goes through **FCM** (Android) and **APNs** (iOS) to a *device token*; email via SES/SendGrid, SMS via Twilio — each its own API, hence `Gateway`. **AWS SNS** fans one message out to many channels. LinkedIn's **Air Traffic Controller** decides whether/when/how to notify each member — `SendRule` at company scale. Phone "Do Not Disturb" = quiet hours with an urgent bypass.

In the interview

30-second answer: "Channels sit behind an abstract class whose final `deliver` holds retry-with-backoff (Template Method); each channel supplies only its address. Each do-not-send condition is a `SendRule` (Strategy) returning a reason. Preferences give an ordered channel list per kind; we render once and try channels in order, falling back on failure. Every outcome goes through one `finish` that records and notifies listeners (Observer). For scale, a queue sits in front."

Follow-ups: *Why not Strategy for channels?* → Channels share an algorithm and vary in one step: Template Method (injecting `RetryPolicy` + a `Sender` strategy is a valid composition alternative). *Duplicates on retry?* → Idempotency keys; at-least-once delivery.

Self-check

1. Which pattern puts the retry loop in one place? → Template Method (`final deliver()`).
2. Why `Optional<String>` from `blocks`? → The suppression reason is recorded.
3. Why not count suppressed sends against the rate limit? → The user never got them.
4. Quiet hours 22–07: is 06:00 quiet? 12:00? → Yes; no.
5. What must every exit of `send` do? → Go through `finish()`.

TL;DR

- **Template Method** channels (retry once in the parent); **Strategy** send rules returning reasons; **Observer** outcomes; **Factory** channels; **Adapter** providers.
- Render once, try channels in order, fall back, single `finish()` exit.
- Scale with a queue; watch idempotency and the rate-limit check-then-act gap.

20.3 Design an In-Memory Cache (and an LRU Cache from scratch)

Problem & clarifying questions

What it is. A small, fast store of recently used values, so you don't re-fetch them from a slow place (DB, another service). Analogy: the few books on your desk vs the library. When the desk is full, one book must go back — **which one** is the eviction policy, the rule most likely to change. Expect the question: *which data structure gives O(1) get and put?* (O(1) = time doesn't grow with the number of entries.)

Question	Assumed answer
Capacity in?	Number of entries.
Eviction?	LRU by default; swappable for FIFO or LFU without changing the cache.
TTL?	Optional per entry; an expired read is a miss and removes it.
Stats / notifications?	Hits, misses, evictions, expirations; subscribers learn about removals and the cause.
Thread-safe?	Yes (the course left it out; we add it).
Load on miss / write-back?	Out of scope (extensions).

Requirements

Functional: the table above, plus: `put` overwrites; a full cache evicts before inserting; O(1) get/put; invalidate one key or clear all; **every entry leaves through one code path**.

Non-functional: O(1); thread-safe; testable TTL (injected clock); listeners can't slow down or deadlock the cache.

Use cases

Actors: **Application** (read, store, remove), **Operator** (chooses rule, reads stats), **Cache** (evicts, expires, counts, notifies).

Core entities

Class	Responsibility
`Node<K,V>`	key, value, `expiresAtMillis`, `frequency`, own `prev` / `next`.
`DoublyLinkedList<K,V>`	Order, newest at front, two **sentinel** nodes (no null checks).
`EvictionPolicy<K,V>`	**Strategy**: `onInsert`, `onAccess`, `onRemove`, `victim()`. Owns the *order*.
`LruPolicy`, `FifoPolicy`, `LfuPolicy`	The three rules; LFU is O(1) with chained frequency buckets.
`EvictionPolicies` + `PolicyType`	**Simple Factory**.
`Cache<K,V>` + `Cache.Builder`	Map + policy + clock + stats + listeners; one lock; **one removal method**. **Builder** for config.
`CacheStats`, `RemovalReason`, `EvictionListener`	`LongAdder` counters; EVICTED/EXPIRED/REMOVED/CLEARED; **Observer**.
`ShardedCache<K,V>`	N caches chosen by key hash → less lock contention.

The core idea: two structures, each answering one question.

Question	Structure	Cost
"Where is key k?"	`HashMap<K, Node>`	O(1)
"Who leaves next?"	Doubly linked list by recency	O(1) at the ends
"Move this entry to the front"	Node has own links; the map hands it over	O(1), no search

A map can't name the oldest entry; a list can't find a key without walking. **Doubly** linked because unlinking a middle node needs its predecessor.

Class diagram

Course build order: (1) `Node` + list; (2) `EvictionPolicy` + LRU; (3) FIFO + LFU with no cache change; (4) `Cache` with clock, stats, TTL; (5) listeners. **Who owns what:** the cache owns the map (and nodes); the policy owns the *order*. Neither reaches into the other — that is why FIFO/LFU need no `Cache` change.

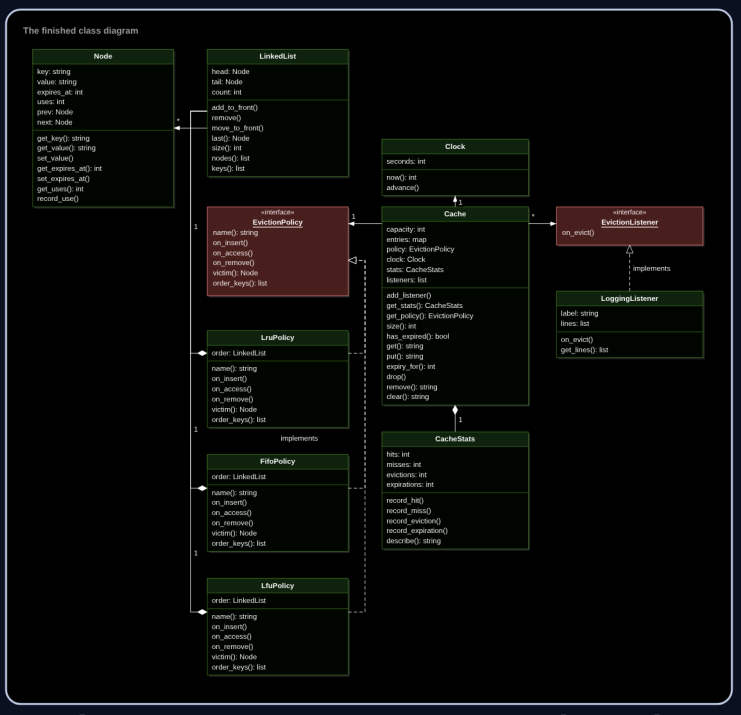
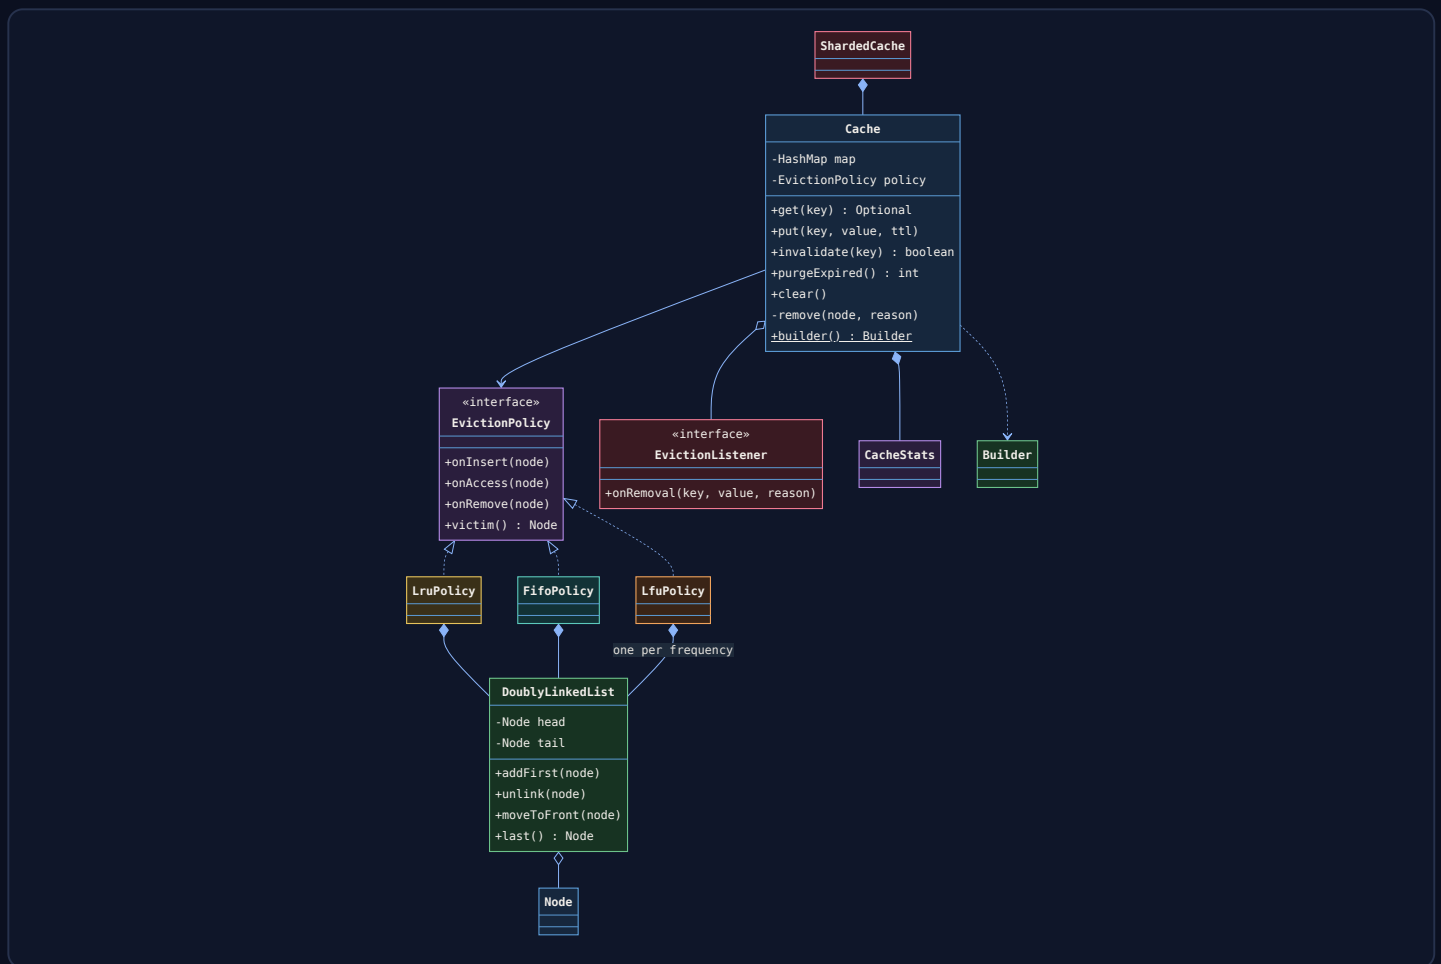


Figure — The finished cache class diagram: Cache with its map, policy, clock, stats and listeners; the three eviction policies behind one interface (course diagram)



Key flows

Read: absent → miss; present but expired → miss + removal (EXPIRED); else hit, `policy.onAccess`, return.
Store: replace if present (counts as a use); otherwise, only if full, ask `policy.victim()` and remove it through the single removal path, then insert. Listeners are called after the lock is released.

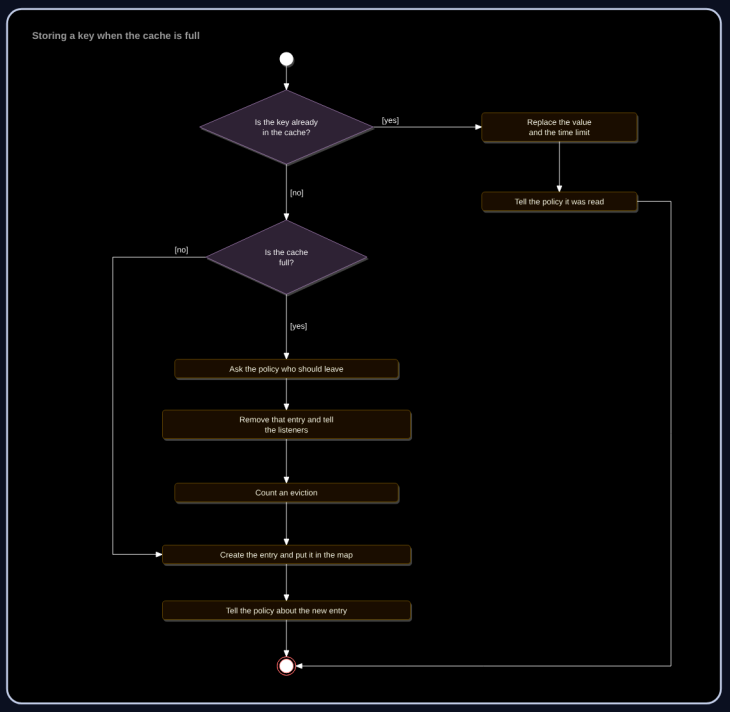


Figure — Storing a key: replace if present; otherwise, if full, ask the policy for a victim and remove it through the single removal path, then insert (course diagram)

How the three policies differ:

Event	LRU	FIFO	LFU (O(1) buckets)
`onInsert`	add to front	add to front	frequency = 1, bucket 1 (always first in chain)
`onAccess`	move to front	**nothing**	move bucket f → f+1 (created right after f); drop f if empty
`onRemove`	unlink	unlink	unlink; drop bucket if empty
`victim()`	back of list	back of list	back of the **first** bucket (LRU among least frequent)

FIFO is LRU with one empty method — evidence of good interface boundaries. The course's LFU scanned all entries (O(n)); ours keeps buckets in a **doubly linked chain in increasing frequency**, so the lowest bucket is always first and every operation is O(1) (Shah, Mitra & Matani, 2010).

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy for eviction	The rule most likely to change.	`switch (policyType)` in get/put.
Observer for removals, with reason	Logging, metrics, cleanup don't belong in the cache.	Hard-coded callbacks.
One removal method `remove(node, reason)`	Eviction, expiry, removal, clear all update map, policy, stats alike.	Four copies of clean-up code.
Sentinel nodes	Add/unlink are always the same pointer writes.	Null checks everywhere.
Injected clock (`LongSupplier`)	TTL tests run instantly.	`System.currentTimeMillis()` → tests must sleep.
Lazy expiry + optional sweeper	Dropped on read; `purgeExpired()` reclaims memory.	Timer per entry.
One `ReentrantLock` , listeners **after unlock**	LRU `get` writes (moves a node), so a read-write lock doesn't help; slow/re-entrant listeners can't block or deadlock.	`synchronized` around listener calls.
Builder	Defaults + validation once.	Telescoping constructors.
Sharding	N locks; threads on different shards never wait.	One global lock.

Java implementation

```java
// File: code/lld/ch20/cache/EvictionPolicy.java
interface EvictionPolicy<K, V> {
    void onInsert(Node<K, V> node);
    void onAccess(Node<K, V> node);
    void onRemove(Node<K, V> node);
    Node<K, V> victim();
}

final class LruPolicy<K, V> implements EvictionPolicy<K, V> {
    private final DoublyLinkedList<K, V> order = new DoublyLinkedList<>();
    public void onInsert(Node<K, V> n) { order.addFirst(n); }
    public void onAccess(Node<K, V> n) { order.moveToFront(n); }
    public void onRemove(Node<K, V> n) { order.unlink(n); }
    public Node<K, V> victim() { return order.last(); }
}
// ... FifoPolicy (onAccess does nothing), LfuPolicy (chained frequency buckets)

// File: code/lld/ch20/cache/Cache.java
    void put(K key, V value, Duration ttl) {
        List<Removal<K, V>> removed = new ArrayList<>(1);
        lock.lock();
        try {
            long expiresAt = ttl == null ? 0 : clockMillis.getAsLong() + ttl.toMillis();
            Node<K, V> existing = map.get(key);
            if (existing != null) {          // replace: counts as a use
                existing.value = value;
                existing.expiresAtMillis = expiresAt;
                policy.onAccess(existing);
                return;
            }
            if (map.size() >= capacity) removed.add(remove(policy.victim(), RemovalReason.EVICTED));
            Node<K, V> n = new Node<>(key, value, expiresAt);
            map.put(key, n);
            policy.onInsert(n);
        } finally {
            lock.unlock();
            notifyListeners(removed);
        }
    }
    // ... get, invalidate, purgeExpired, clear follow the same lock / collect / notify-after-unlock shape

    /** THE single exit path for every entry, whatever the reason. Caller holds the lock. */
    private Removal<K, V> remove(Node<K, V> n, RemovalReason reason) {
        map.remove(n.key);
        policy.onRemove(n);
        if (reason == RemovalReason.EVICTED) stats.evictions.increment();
        if (reason == RemovalReason.EXPIRED) stats.expirations.increment();
        return new Removal<>(n.key, n.value, reason);
    }
```

Full runnable code: code/lld/ch20/cache/ (see the Full Edition).

Output (excerpt) — same accesses (b read 3×, c once, a twice), three different victims:

```
== LRU (capacity 3): put a,b,c; get b,b,b,c,a,a; put d
  listener: b=2 left, EVICTED
== FIFO (capacity 3): put a,b,c; get b,b,b,c,a,a; put d
  listener: a=1 left, EVICTED
== LFU (capacity 3): put a,b,c; get b,b,b,c,a,a; put d
  listener: c=3 left, EVICTED
  stats: hits=2 misses=1 evictions=0 expirations=2 hitRate=67%
```

LRU Cache from scratch (the classic interview question)

Often asked directly: **"Implement an LRU cache with O(1) get and put"** (LeetCode 146). Same two structures, inline: HashMap for lookup; doubly linked list with `head` / `tail` sentinels (most recent after `head`). `get`: unlink + add at front. `put`: if present update + move to front; else if full remove `tail.prev` from list **and** map (why the node stores its key), then add. Java shortcut: `LinkedHashMap(accessOrder = true)` + `removeEldestEntry` — mention it, but write the manual version when asked "from scratch".

```java
// File: code/lld/ch20/lrucache/Main.java
        synchronized V get(K key) {
            Node n = index.get(key);
            if (n == null) return null;
            unlink(n);
            addFront(n);                // now the most recently used
            return n.value;
        }

        synchronized void put(K key, V value) {
            Node n = index.get(key);
            if (n != null) {            // update + mark as recently used
                n.value = value;
                unlink(n);
                addFront(n);
                return;
            }
            if (index.size() == capacity) {
                Node lru = tail.prev;    // least recently used lives just before the tail sentinel
                unlink(lru);
                index.remove(lru.key);
            }
            n = new Node(key, value);
            addFront(n);
            index.put(key, n);
        }
        // ... addFront / unlink: four and two pointer writes
    // ... jdkLru(capacity): new LinkedHashMap<>(16, 0.75f, true) overriding removeEldestEntry -> size() > capacity
```

Full runnable code: code/lld/ch20/lrucache/ (see the Full Edition). The demo prints LeetCode 146's expected values (1, miss, miss, 3, 4); this generic version returns `null` for a miss instead of `-1`.

Approach	get / put	Thread-safe	When to use
HashMap + own doubly linked list	O(1) / O(1)	Only with a lock	Interviews; custom behaviour
`LinkedHashMap(accessOrder=true)` + `removeEldestEntry`	O(1) / O(1)	No	Quick single-threaded caches
Caffeine / Guava	amortized O(1)	Yes	Production Java

Concurrency & edge cases

- **A read lock is not enough:** LRU `get` rewrites pointers, so it needs the exclusive lock.
- **Contention:** `ShardedCache` = 8 locks; trade-off: LRU order is only per shard.
- **Cache stampede** (many misses for one key all hit the DB): let the first thread load and others wait on its result (`computeIfAbsent` / a `CompletableFuture` per key) — Chapter 13's TTL cache does this with `ConcurrentHashMap.compute`.
- **TTL precisely:** expires *at* the deadline (`now >= expiresAt`); lazy — a dead entry holds a slot until read, swept, or evicted (then reported `EVICTED`). `put` without TTL on an existing key **clears** the old TTL.
- Capacity ≤ 0 refused by the builder; nulls unsupported (Caffeine too); use a monotonic clock (`nanoTime`) in production.

Extensions & follow-up questions

- **LFU:** one new policy; naive LFU is O(n), chained buckets O(1).
- **Size in bytes:** a `Weigher` strategy; evict until under max weight.
- **Load on miss:** *read-through* (`get(key, loader)`), *write-through* (cache + DB together), *write-back* (flush later; faster, may lose data), *cache-aside* (app does check → DB → put).
- **Distributed:** consistent hashing (high-level design). **Best policy?** Workload-dependent; Caffeine's *Window TinyLFU* mixes recency and frequency.

Real-world example

`LinkedHashMap` access order is the JDK's LRU. **Guava/Caffeine** expose the same knobs: `maximumSize`, `expireAfterWrite` / `expireAfterAccess`, `removalListener` (with causes like `EXPIRED`, `SIZE`, `EXPLICIT`, `REPLACED`), `recordStats()`. **Redis** picks eviction by config (`maxmemory-policy`: `allkeys-lru`, `allkeys-lfu`, `volatile-ttl`, `noeviction`…) and approximates LRU by sampling keys.

In the interview

30-second answer: "A HashMap key → node for O(1) lookup, and a doubly linked list with sentinels for O(1) ordering; the node carries its own links so we can unlink it directly. The eviction rule sits behind `EvictionPolicy` — LRU moves on access, FIFO doesn't, LFU uses frequency buckets for O(1). Every exit goes through one `remove` that updates stats; listeners are notified after the lock is released. One lock is correct; sharding reduces contention."

Follow-up: *Why does the node store the key?* → To delete it from the map on eviction.

Self-check

1. Which two structures give O(1) LRU? → HashMap (key → node) + doubly linked list.
2. What do sentinels buy? → No null checks; add/remove always the same writes.
3. FIFO vs LRU in code? → `onAccess` does nothing.
4. Why can't a read-write lock speed up LRU `get`? → `get` modifies the list.
5. Cache stampede fix? → One thread loads; others wait on the same future.

TL;DR

- **HashMap + doubly linked list (sentinels)** = O(1) get/put.
- **Strategy** eviction (LRU / FIFO / O(1) LFU); **Observer** removals; **Builder** config.
- **One removal path**; lazy TTL + sweeper; injected clock.
- One lock (gets are writes), listeners outside it, **shard** to scale; know `LinkedHashMap`, Caffeine, Redis policies.

20.4 Design a Ride Sharing Service (Uber)

Problem & clarifying questions

What it is. Connects a rider with a nearby driver — a radio taxi dispatcher who calls the nearest free cab, then the next if refused, and the meter shows the fare. Two decisions are the whole question: **which**

driver is asked (matching) and **what the trip costs** (pricing). Either should change with no edits elsewhere.

Question	Assumed answer
Vehicle types?	Economy, premium, van; rider asks for one.
Who gets a request?	Online, free drivers of the right type within range.
Offer one at a time or broadcast?	One at a time, nearest first; decline → next; none left → cancel.
Trip states?	Requested → accepted → started → ended; cancel only from requested/accepted. Illegal actions do nothing.
Fare / surge?	Base + per km + per minute, with a minimum, per type; surge multiplies the same fare.
Ratings / payment?	Each side rates once after a finished trip (average shown); charge rider's method (cash, wallet, card).
Map?	Grid with Manhattan distance (production: a routing engine giving road distance + duration).
Pool, scheduled rides, tracking, payouts?	Out of scope.

Requirements

Functional: the table above, plus: driver has one vehicle (plate, seats, type) and goes online/offline; rider cancels only before start; trip history for both sides.

Non-functional: a driver can **never** be on two trips (even under concurrency); rules swappable; deterministic tie-breaks.

Use cases

Actors: **Rider** (request, cancel, rate), **Driver** (online, accept/decline, start/end, rate), **System** (find, offer, price, surge, history).

Core entities

Class	Responsibility
`Location`, `VehicleType`, `Vehicle`	Grid point (Manhattan `distanceTo`), car type and details.
`Account` (abstract) → `Rider`, `Driver`	Shared id/name/phone, **rating as total + count**, finished trips. Rider + payment method; Driver + vehicle, location, online, **atomic busy flag** (`tryClaim`).
`TripStatus` / `Trip`	State machine table; trip holds **ranked candidate list + cursor**; every action checks state.
`PricingStrategy`, `StandardPricing` + `RateCard`	**Strategy**: base + km + min with floor, per type.
`SurgePricing`	**Decorator** wrapping any `PricingStrategy`.
`MatchingStrategy`	**Strategy**: `NearestDriverMatcher`, `BestRatedMatcher`.
`PaymentMethod`	**Strategy**: `CashPayment`, `WalletPayment`.
`TripListener`	**Observer** of trip events.
`RideSharingService`	Facade; owns accounts and trips; knows only the order of steps.

Rating as total + count: only the average is shown, and adding a rating is O(1). `Rider` / `Driver` extend `Account` because they pass the LSP "is-a" test.

Class diagram

Course stages: (1) accounts, vehicle, location; (2) `Trip` + `TripStatus`; (3) pricing + surge wrapper; (4) matching; (5) the service. In the course figure, the arrow from `SurgePricing` back to `PricingStrategy` ("has-a of my own type") is what makes it a Decorator rather than another Strategy.

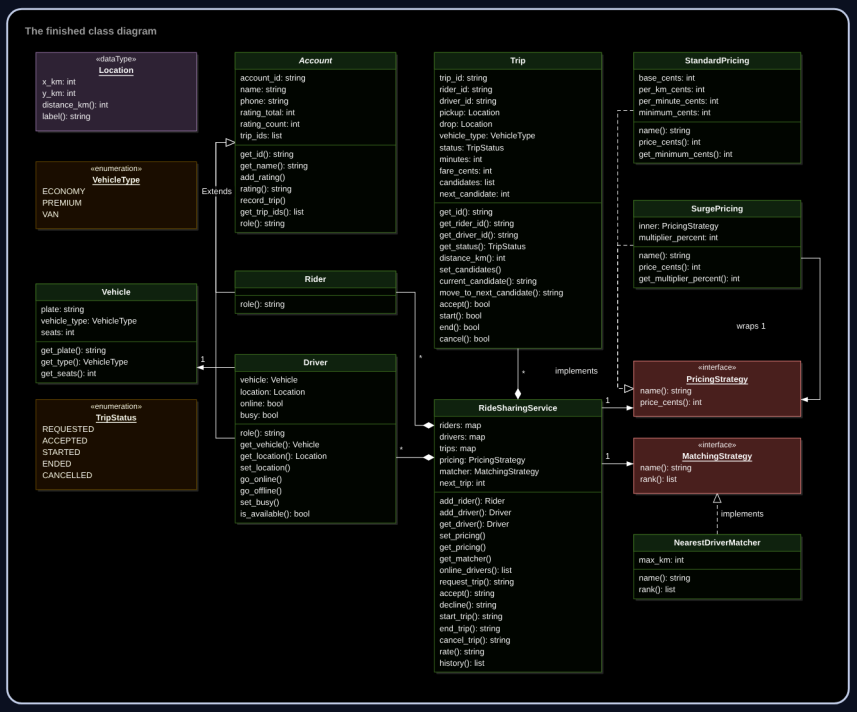
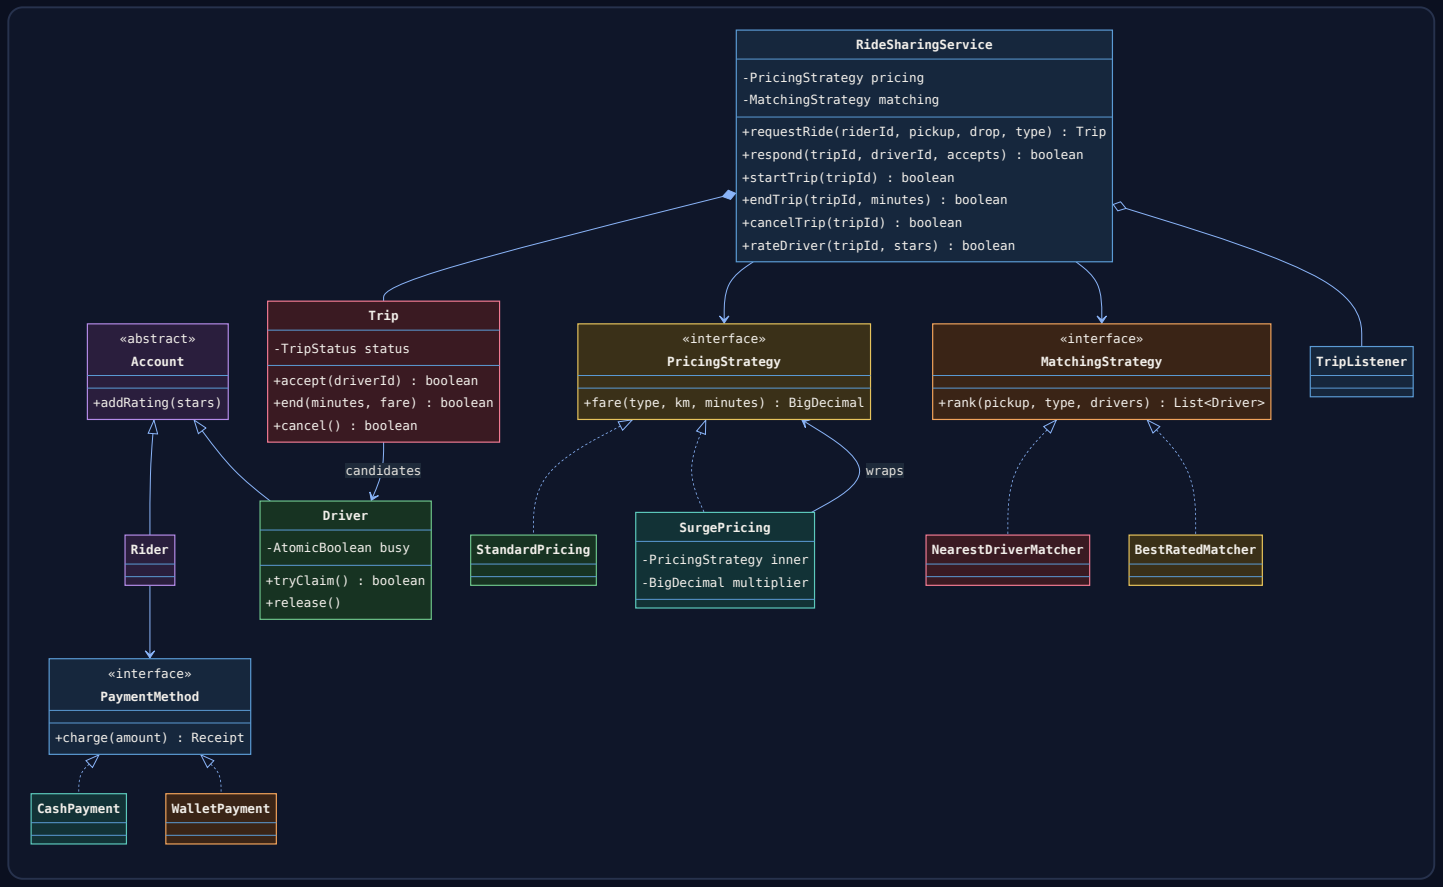
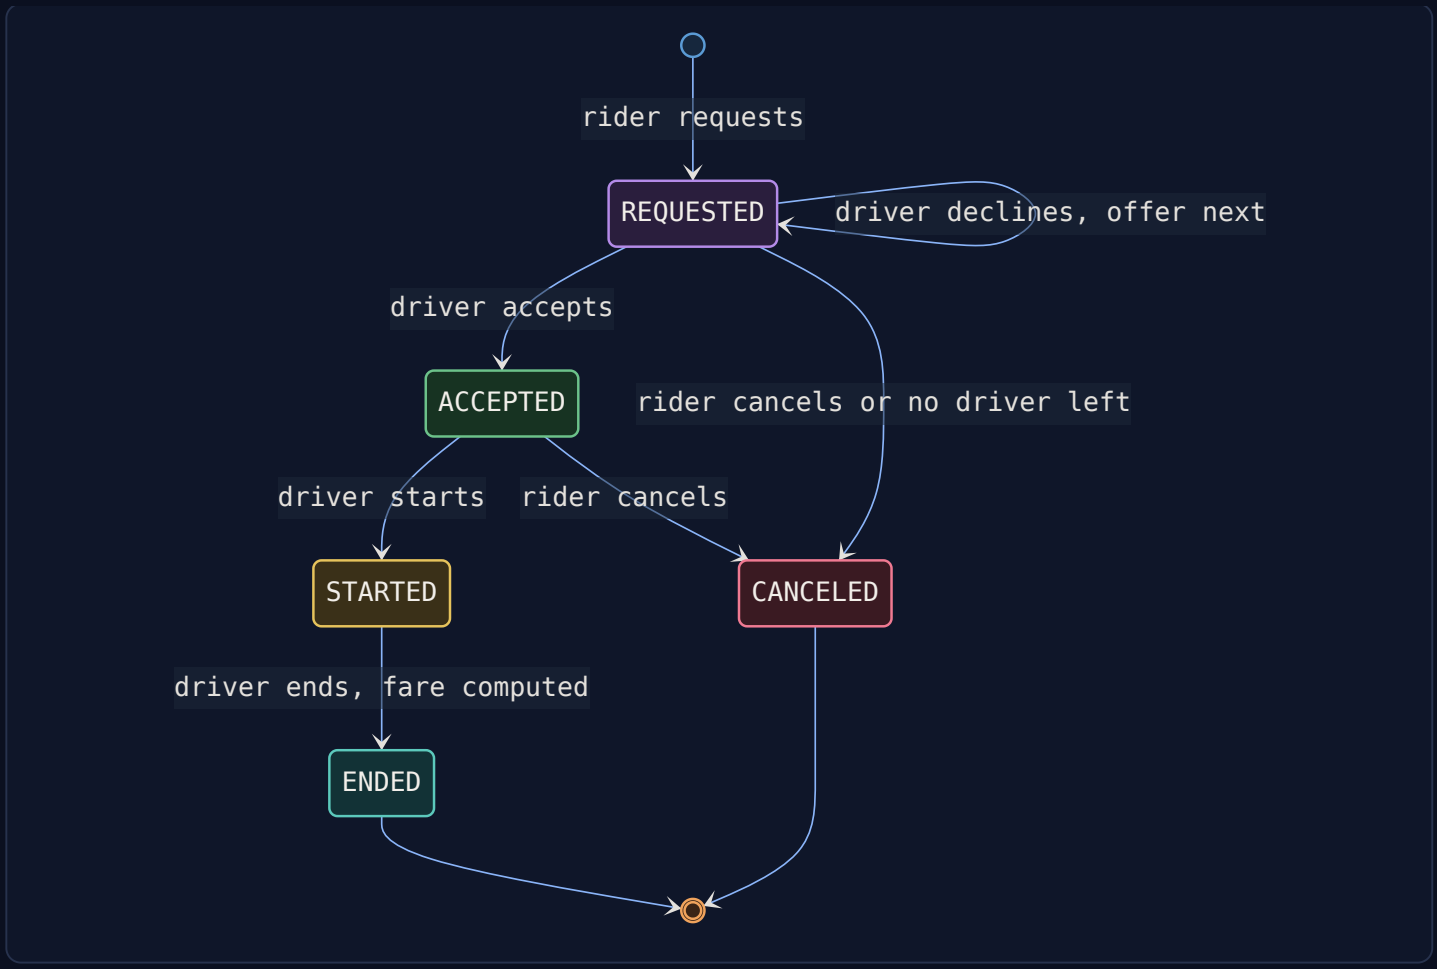


Figure — The finished ride sharing class diagram: the service owning riders, drivers and trips; PricingStrategy with StandardPricing and the SurgePricing wrapper; MatchingStrategy with the nearest-driver matcher (course diagram)



Key flows

Trip lifecycle (a state machine = fixed states + the only legal moves):



Finding a driver: rank free drivers of the right type nearest first; offer to the first; decline → move the cursor; nobody left → cancel. **Ending:** refuse if not started; else one `pricing.fare(...)` call (however many decorators wrap it), charge the rider, move the driver to the drop point, release.

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy for matching	Cities rank differently (nearest, best-rated…).	Sorting inside the service.
Matcher returns a **ranked list** stored on the trip	Decline = move cursor; no re-search; remembers who refused.	Returning one driver.
Strategy pricing + **Decorator** surge	Surge wraps *any* rule; off = remove the wrapper; decorators stack (surge + airport fee + discount).	`if (surge)` in pricing.
State machine via `TripStatus.canMoveTo`	One table; illegal moves return `false`. Full State pattern only pays when per-state behaviour is large.	Scattered `if (status == …)`.
Atomic claim (`AtomicBoolean.compareAndSet`)	Two trips racing for a driver: exactly one wins; loser treats it as decline.	Check `isAvailable()` then set: race.
Move driver to **drop point** at end	Next match must use the real position.	Forgetting (real bug class).
Observer for trip events	Apps, push, analytics listen.	Service calling them directly.
Tie-break by driver id; `BigDecimal` `HALF_UP`	Reproducible tests; exact fares (course uses whole cents and surge as a percentage, 150 = ×1.5).	`double`.

Java implementation

```java
// File: code/lld/ch20/ridesharing/Trip.java
enum TripStatus {
    REQUESTED, ACCEPTED, STARTED, ENDED, CANCELED;

    private static final Map<TripStatus, Set<TripStatus>> NEXT = Map.of(
            REQUESTED, EnumSet.of(ACCEPTED, CANCELED),
            ACCEPTED, EnumSet.of(STARTED, CANCELED),
            STARTED, EnumSet.of(ENDED),
            ENDED, EnumSet.noneOf(TripStatus.class),
            CANCELED, EnumSet.noneOf(TripStatus.class));

    boolean canMoveTo(TripStatus next) { return NEXT.get(this).contains(next); }
}

// File: code/lld/ch20/ridesharing/Strategies.java
/** Decorator: wraps ANY pricing rule and multiplies its answer. Remove the wrapper to switch surge off. */
final class SurgePricing implements PricingStrategy {
    private final PricingStrategy inner;
    private final BigDecimal multiplier;
    // ... constructor
    @Override
    public BigDecimal fare(VehicleType type, int km, int minutes) {
        return inner.fare(type, km, minutes).multiply(multiplier).setScale(2, RoundingMode.HALF_UP);
    }
}

// File: code/lld/ch20/ridesharing/Model.java  (Driver)
    boolean tryClaim() { return online && busy.compareAndSet(false, true); }

// File: code/lld/ch20/ridesharing/RideSharingService.java
    /** Driver answers the offer. A lost race for the driver counts as a decline. */
    boolean respond(String tripId, String driverId, boolean accepts) {
        Trip trip = trips.get(tripId);
        Driver current = trip.currentCandidate();
        if (current == null || !current.id().equals(driverId) || trip.status() != TripStatus.REQUESTED) return false;
        if (accepts && current.tryClaim()) {
            if (trip.accept(driverId)) {
                publish(trip, "accepted by " + driverId);
                return true;
            }
            current.release(); // trip was canceled meanwhile
            return false;
        }
        publish(trip, (accepts ? "could not claim " : "declined by ") + driverId);
        offerOrCancel(trip, trip.nextCandidate());
        return false;
    }
    // ... endTrip: price -> trip.end -> charge rider -> d.moveTo(drop) -> d.release() -> history
```

Full runnable code: code/lld/ch20/ridesharing/ (see the Full Edition).

`Trip` methods are `synchronized`, so a rider's cancel and a driver's accept can't interleave. `endTrip` relocates the driver **before** releasing them, so no matcher ever sees a free driver at the old spot.

Output (excerpt):

```
[T1] declined by D1
[T1] accepted by D2
[T1] ended: 7 km, 12 min, fare 13.90 (wallet debited 13.90, left 36.10)
[T2] ended: 1 km, 3 min, fare 7.50 (collect 7.50 in cash)
[T4] cancel refused in state STARTED
[T5] no driver left - canceled
trips accepted = 10 (never more than 10 drivers)
```

Trip 1: 2.50 + 7 × 1.20 + 12 × 0.25 = **13.90**. Trip 2: 4.45 → minimum 5.00 → surge × 1.5 = **7.50**. 40 concurrent requests vs 10 drivers → exactly 10 accepted.

Machine-coding variant: carpool / "offer and select a ride"

A Flipkart-style round (kumaransg/LLD collection) with **no dispatcher**: users *offer* spare seats on a route and others *select* a ride with a **selection strategy**.

- `add_user`, `add_vehicle` (several per user); `offer_ride(user, origin, seats, vehicle, destination)` **fails if that vehicle already has an active offer**.
- `select_ride(user, origin, destination, seats, strategy)` — seats 1 or 2; **Most Vacant** or **Preferred Vehicle** ("Activa").
- `end_ride`, `print_ride_stats` — taken/offered per user, counting only **ended** rides.

Design: `Ride` with an atomic `book`; `RideSelectionStrategy` (`MostVacant`, `PreferredVehicle`); `RideService` filters active rides by route and seats, excludes the user's own offers, and lets the strategy choose.

```
// File: code/lld/ch20/carpool/Main.java
    interface RideSelectionStrategy {
        Optional<Ride> select(List<Ride> candidates);
    }

    static final class MostVacant implements RideSelectionStrategy {
        public Optional<Ride> select(List<Ride> c) { return
c.stream().max(Comparator.comparingInt(Ride::availableSeats)); }
    }
    // ... PreferredVehicle(model): first ride whose vehicle model matches
    // ... selectRide: filter ACTIVE rides on the same route, not offered by this user, with enough seats;
    // ... strategy.select(matching), then ride.book(user, seats) (synchronized check-and-book)
```

Full runnable code: code/lld/ch20/carpool/ (see the Full Edition). The demo replays the problem's sample (answers 2(c), 2(b), none, 2(d), none; then the stats).

	Uber-style (main design)	Carpool variant
Who supplies the ride	Professional drivers	Any user, per offered ride
Who chooses	System offers to drivers	Rider picks with a selection strategy
Key pattern	Strategy + Decorator + state machine	Strategy (selection)
Key invariant	A driver on ≤ 1 trip	A vehicle has ≤ 1 active offer; seats never negative

Concurrency & edge cases

- **Double booking:** prevented by `tryClaim()` CAS (the 40/10 race proves it).
- **Cancel vs accept:** both via `synchronized` `Trip` methods; if cancel wins, `accept` fails and the driver is released.
- **Stale candidate list:** computed once; a claimed/offline driver fails `tryClaim` → next. A nearer driver freed later isn't noticed (real systems re-match).
- **Offer timeouts:** no answer in N seconds = decline (scheduled task per offer).
- No drivers → canceled immediately. Illegal transitions → `false`, state unchanged. Wallet too low → `Receipt(paid=false)`. One rating per side, 1–5 stars.

Extensions & follow-up questions

- **Rank by rating:** swap in `BestRatedMatcher`.
- **Airport fee, promo, tolls:** more decorators: `new PromoDiscount(new AirportFee(new SurgePricing(standard, 1.5)))`.
- **How is surge computed?** Demand (open requests) vs supply (free drivers) per area and window → multiplier; the pricing design doesn't change.
- **Nearby drivers fast:** geospatial index — cells (geohash, S2, Uber's H3 hexagons) → map cell → drivers; search pickup cell + neighbours.
- **Pool rides:** honest answer: `Trip` becomes many riders/stops and the matcher considers detours. Live tracking: latest location per driver, published via pub-sub (20.6). Scheduled rides: a scheduler starts matching shortly before.

Real-world example

Uber shows surge ("dynamic pricing") when demand exceeds supply and gives a driver a short window to accept before the next — the ranked-list-with-cursor idea. Uber open-sourced **H3**, a hexagonal grid used for pricing and dispatch. Ola, Lyft and Grab follow the same request → match → accept → trip → rate lifecycle.

In the interview

30-second answer: "Two rules behind interfaces: `MatchingStrategy` returns a ranked list of free drivers with the right car; `PricingStrategy` computes base + distance + time with a floor. Surge is a Decorator wrapping any pricing rule. The trip is a small state machine; every action checks the transition table and returns false if illegal. A decline moves a cursor along the stored ranking. Drivers are claimed with compare-and-set, so one driver never gets two trips."

Tip: name matching and pricing in the first minutes; spend real time on surge — *wrapping* the rule instead of adding a flag separates strong answers. *Strategy vs Decorator?* → Strategy replaces the rule; Decorator wraps it and stays in the call path (holds an `inner` of the same interface).

Self-check

1. Why a ranked list, not one driver? → Decline = next candidate, no re-search.
2. What makes `SurgePricing` a Decorator? → It is a `PricingStrategy` wrapping another and adjusting its result.
3. Can a rider cancel a STARTED trip? → No; `STARTED` only → `ENDED`.
4. Which Java class makes the claim atomic? → `AtomicBoolean.compareAndSet(false, true)`.
5. Carpool: when does `offer_ride` fail? → The vehicle already has an active offer.

TL;DR

- **Strategy** ×3 (matching, pricing, payment) + **Decorator** (surge) + **state machine** (trip).
- Ranked list + cursor; decline = next. **Atomic claim** prevents double booking; relocate driver at trip end.
- Ratings = total + count; `BigDecimal` fares with a minimum.
- Carpool: users offer seats; riders select with Most Vacant / Preferred Vehicle.

20.5 Design a Logging Framework

Problem & clarifying questions

What it is. Code writes messages ("payment 991 authorised") with a **level** (severity), and the framework sends them to **destinations** (console, file, log server) in a chosen **format** — a mini Log4j 2 / Logback /

`java.util.logging`. Analogy: a company mail room that routes notes by department and urgency stamp to inboxes, in the right envelope.

Question	Assumed answer
Levels?	TRACE < DEBUG < INFO < WARN < ERROR < FATAL (and OFF).
Per-module config?	Loggers named `app.payment`, `app.orders` form a **dot hierarchy**; a child inherits the nearest ancestor's level.
Destinations / formats?	Console, file, in-memory; pattern and JSON; both pluggable; each destination has its own minimum level.
Propagation?	Event goes to the logger's destinations then **up** to ancestors', unless additivity is off.
Async / thread-safe?	Optional async (app never waits for disk); yes.
Config files, rolling files, shipping?	Out of scope.

Requirements

Functional: the table above, plus: same name → same logger; exceptions attachable; lazy messages; async without changing appenders.

Non-functional: very cheap when disabled; thread-safe; logging never crashes the app; testable (injected clock).

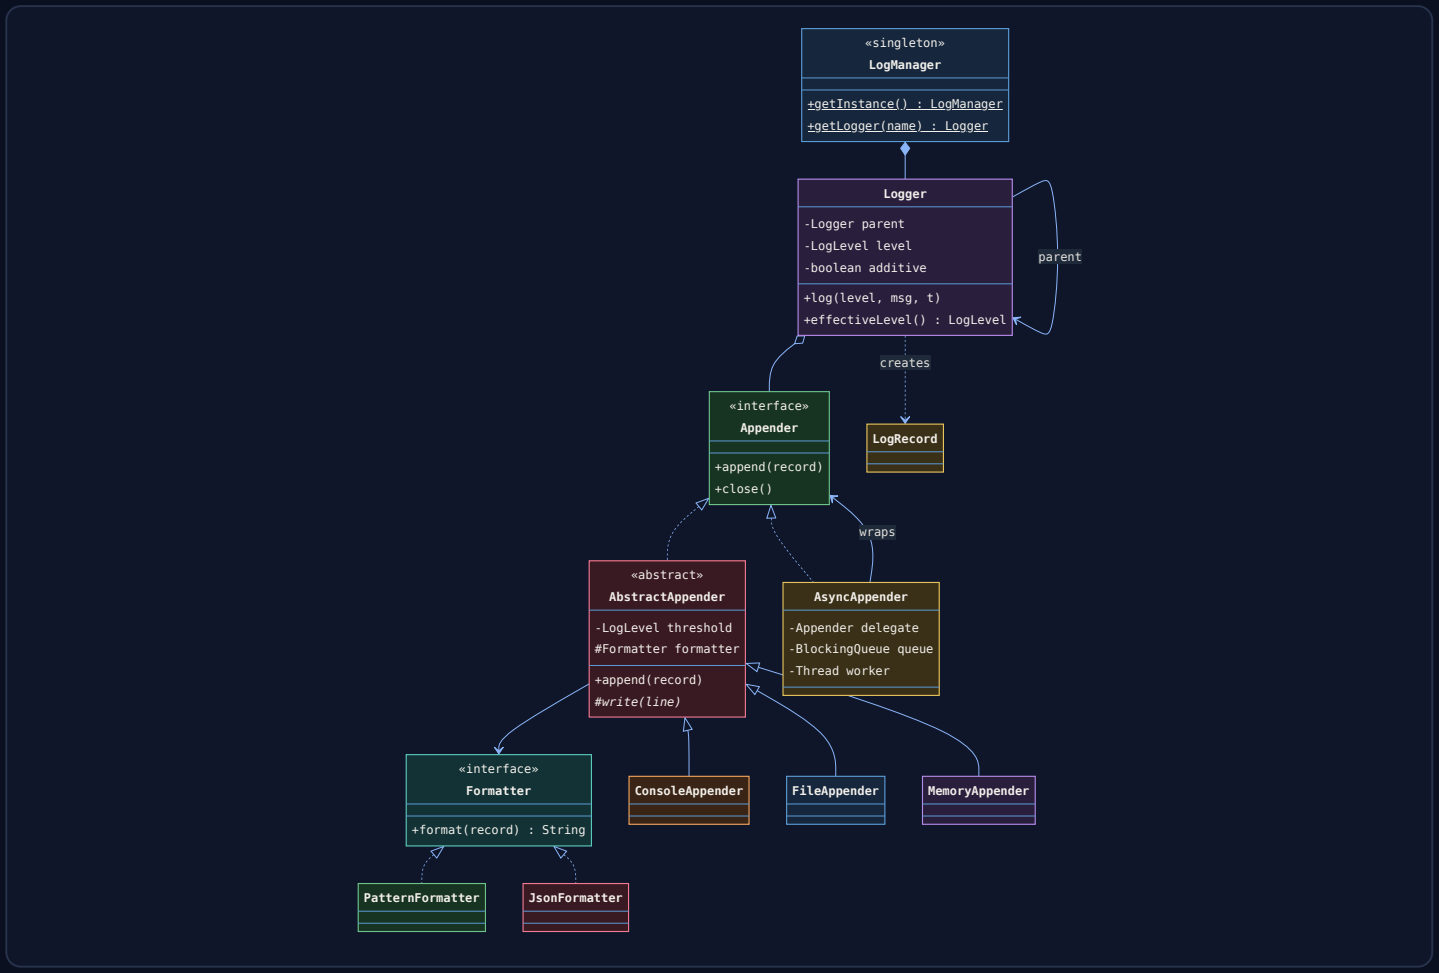
Use cases

Actors: **Developer code** (logs), **Operator** (configures levels, appenders, formats), **Framework** (filters, formats, routes, writes). Use cases: log a message, get a logger by name, set a level for a subtree, add a destination, choose a format, switch to async, stop propagation for a subtree.

Core entities

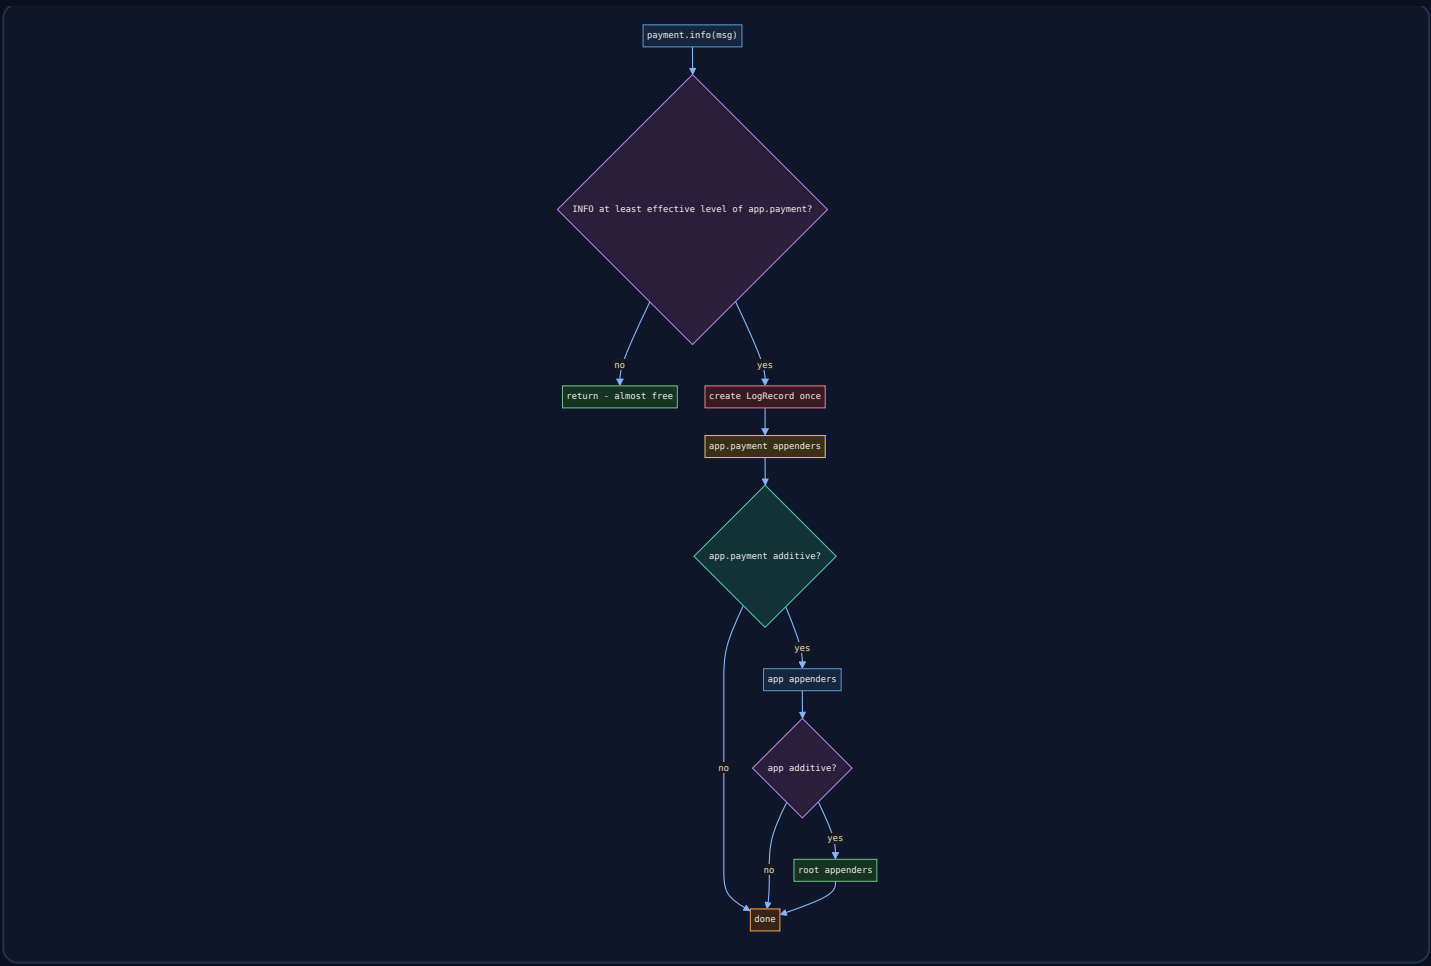
Class	Responsibility
`LogLevel`	Ordered severities; `isAtLeast(threshold)`.
`LogRecord`	One immutable event: time, level, logger, thread, message, error.
`Logger`	Named node: parent, optional level, appenders, additivity; `log()` walks up.
`LogManager`	**Singleton** registry; `getLogger(name)` creates parents first, one instance per name; holds the clock.
`Appender` / `AbstractAppender`	A sink; **Template Method** base checks threshold + formats, subclasses only `write(line)`.
`ConsoleAppender`, `FileAppender`, `MemoryAppender`	Concrete sinks.
`AsyncAppender`	**Decorator**: bounded queue + one writer thread around any appender.
`Formatter`	**Strategy**: `PatternFormatter`, `JsonFormatter`.

Class diagram



Key flows

Logging one event = Chain of Responsibility up the hierarchy (a request passes along a chain of handlers; each may handle it and/or pass it on):



The Log4j/Logback rule we copy: the **level check happens once, at the logger where the call was made**; parents' levels are not re-checked, only each **appender's own threshold** filters. So a DEBUG event from `app.payment` (set to DEBUG) still reaches root's console even though root is INFO.

Async: the app thread calls `append` → `queue.offer` (never blocks) → returns; a writer thread `take()`s and calls the slow delegate. `close()` enqueues a poison record, then `join()`s until the queue is drained.

Design decisions & patterns used

Pattern / decision	Where	Why
Chain of Responsibility	`Logger.log`: this → parent → … → root, stops at `additive == false`	Configure sinks once near the root; subtrees add or stop.
Singleton (holder idiom)	`LogManager`	One registry → same logger per name; lazy + thread-safe with no locking (JVM initializes the nested class once). Double-checked locking is easy to get wrong.
Static factory	`getLogger(name)`	Caches instances, creates parents first; `new Logger` would break the hierarchy.
Strategy	`Formatter`	Any sink × any format; no `JsonFileAppender` class explosion.
Template Method	final `append` → `write`	Threshold + formatting once.
Decorator	`AsyncAppender`	Async is an add-on, not `AsyncFileAppender`, `AsyncConsoleAppender` …
Level check first; lazy `Supplier<String>`	`isEnabled`, `debug(Supplier)`	Disabled logging = one comparison.
Appender failures caught	`Logger.log`	Logging never crashes the app.
Parent resolved **outside** `computeIfAbsent`	`LogManager.loggerFor`	A `ConcurrentHashMap` mapping function must not modify the same map (`IllegalStateException: Recursive update`).

Another common "CoR logger" answer chains *level handlers* (`ErrorHandler → WarnHandler → InfoHandler`). It works as a toy, but real frameworks put the chain on the **logger hierarchy**, because that is what operators configure. If the interviewer expects the level-handler version, explain both.

Java implementation

```
// File: code/lld/ch20/logging/Logger.java
    // ... effectiveLevel() walks up to the nearest logger with a level; isEnabled(); debug(Supplier) for lazy
messages
    void log(LogLevel lvl, String msg, Throwable t) {
        if (!isEnabled(lvl)) return;                    // cheap check first
        LogRecord r = new LogRecord(manager.now(), lvl, name, Thread.currentThread().getName(), msg, t);
        for (Logger l = this; l != null; l = l.parent) { // the chain
            for (Appender a : l.appenders) {
                try {
                    a.append(r);
                } catch (RuntimeException e) {
                    System.err.println("appender failed: " + e); // logging must never crash the app
                }
            }
            if (!l.additive) break;
        }
    }

/** Singleton registry of loggers (initialization-on-demand holder idiom: lazy and thread-safe). */
final class LogManager {
    // ...
    private static final class Holder { static final LogManager INSTANCE = new LogManager(); }

    static LogManager getInstance() { return Holder.INSTANCE; }

    private Logger loggerFor(String name) {
        if (name == null || name.isEmpty() || name.equals("root")) return root;
        Logger existing = loggers.get(name);
        if (existing != null) return existing;
        int dot = name.lastIndexOf('.');
        Logger parent = dot < 0 ? root : loggerFor(name.substring(0, dot)); // resolved OUTSIDE computeIfAbsent
        return loggers.computeIfAbsent(name, n -> new Logger(n, parent, this));
    }
    // ...
}

// File: code/lld/ch20/logging/Appenders.java
    @Override
    public final void append(LogRecord r) {
        if (r.level().isAtLeast(threshold)) write(formatter.format(r));
    }
    // ...
    @Override
    public void append(LogRecord r) {
        if (closed || !queue.offer(r)) dropped.incrementAndGet(); // full or closed: count it, never block
    }
```

Full runnable code: code/lld/ch20/logging/ (see the Full Edition).

`LogRecord` is created once and shared (immutable, thread-safe). Sinks synchronize `write` so lines never interleave. `AsyncAppender`'s promise is precise: `written + dropped` always equals the number of `append` calls; a throwing delegate is caught inside the writer loop so the writer never dies (which would leave `close()` waiting forever).

Output (excerpt):

```
2026-09-28 10:15:30.123 DEBUG [main] app.payment - card token=tok_42 (built lazily)
== 4. Additivity off: 'app.payment' events stop at its own appenders
  (nothing printed on console above) audit sink now has 2 lines
== 6. Async appender: 4 threads x 1000 events, then close() flushes
  written=4000 dropped=0 total=4000
```

Concurrency & edge cases

- Appenders in `CopyOnWriteArrayList`; level and additivity `volatile`; each sink's `write` synchronized.
- **Async queue full:** drop-and-count (never block). Alternatives: block (back-pressure, may stall the app) or drop low levels only — Logback's `AsyncAppender` by default discards TRACE/DEBUG/INFO at 80% full.
- **Shutdown:** without `close()` queued events are lost → JVM shutdown hook.

- **Recursive logging** from an appender can loop forever → thread-local guard.
- **Security:** never interpret messages — Log4Shell (CVE-2021-44228, Dec 2021) came from Log4j 2 evaluating `${jndi:...}` in logged strings. Mask card numbers, tokens, passwords.

Extensions & follow-up questions

- **Rolling files:** `RollingFileAppender` + `RollingPolicy` (size/time) Strategy.
- **Config file:** parse into `LogManager` calls; watch it for runtime changes.
- **Filters:** a `Filter` chain per appender — another Chain of Responsibility.
- **MDC / request id:** thread-local map copied into each record. Ship logs: network appender behind `AsyncAppender`, batching. Structured logging: key-value fields rendered as JSON.

Real-world example

SLF4J is the facade; **Logback** and **Log4j 2** implement exactly this model (dot-named hierarchy, inherited levels, appenders with layouts, *additivity*). Log4j 2 async loggers use the LMAX Disruptor ring buffer. `java.util.logging`: `Logger`, `Handler` (appender), `Formatter`, `Level`, `setUseParentHandlers(false)` (additivity off).

In the interview

> **30-second answer:** "Loggers form a dot hierarchy under a root managed by a singleton `LogManager`. A log call checks the inherited effective level, builds one immutable record, and passes it up the chain of loggers to their appenders — Chain of Responsibility — until additivity stops it. Each appender has its own threshold and a pluggable formatter (Strategy). `AsyncAppender` is a Decorator with a bounded queue and a writer thread."

Follow-ups: *Why is a disabled debug call cheap?* → One comparison before any string; `Supplier` for expensive ones. *Queued logs on crash?* → Lost unless flushed; close on shutdown.

Self-check

1. `app.payment` unset, `app` WARN, root INFO → effective level? → WARN.
2. Additivity=false? → Stops the event reaching ancestors' appenders.
3. Pattern making any appender async? → Decorator.
4. Why is `Formatter` separate from `Appender`? → Formats and sinks combine freely (Strategy).
5. How does `close()` lose nothing queued? → Poison pill + `join()`; refused records are counted as dropped.

TL;DR

> - Named **logger hierarchy**, inherited level, one cheap check up front.
> - **Chain of Responsibility** to the root (additivity stops it); **Singleton** `LogManager`.
> - **Strategy** formatters, **Template Method** appender base, **Decorator** async.
> - Immutable record, thread-safe sinks, never crash, never block on a full queue.

20.6 Design a Pub-Sub System

Problem & clarifying questions

What it is. Parts of a system talk **without knowing each other**: a *publisher* sends to a named *topic*; every *subscriber* of it receives the message. Analogy: a newspaper printed once, delivered to every subscriber,

new readers welcome any time. It is the **Observer pattern as infrastructure**: a **broker** in the middle decouples by topic, time and often machine. We design a simplified, Kafka-like, in-process broker.

Question	Assumed answer
Delivery to how many?	Every subscription gets every message (fan-out), in publish order.
Guarantee?	**At-least-once**: ack only after processing; failures retried; duplicates possible.
Slow subscriber?	Must not slow others or the publisher.
Keep messages?	Yes: append-only log with **offsets** → late subscribers and replay.
Poison message?	After N attempts → dead-letter topic.
Push or pull?	Push, one dispatcher per subscription.
Partitions, consumer groups, disk, retention?	Out of scope.

Requirements

Functional: create topics; publish (key, payload) → an **offset** (position in the log); subscribe from EARLIEST or LATEST; in-order delivery; commit only after the callback returns; retry, then dead-letter; **seek** to replay; unsubscribe; clean shutdown.

Non-functional: thread-safe publish; subscribers isolated; no busy-waiting; order preserved.

Use cases

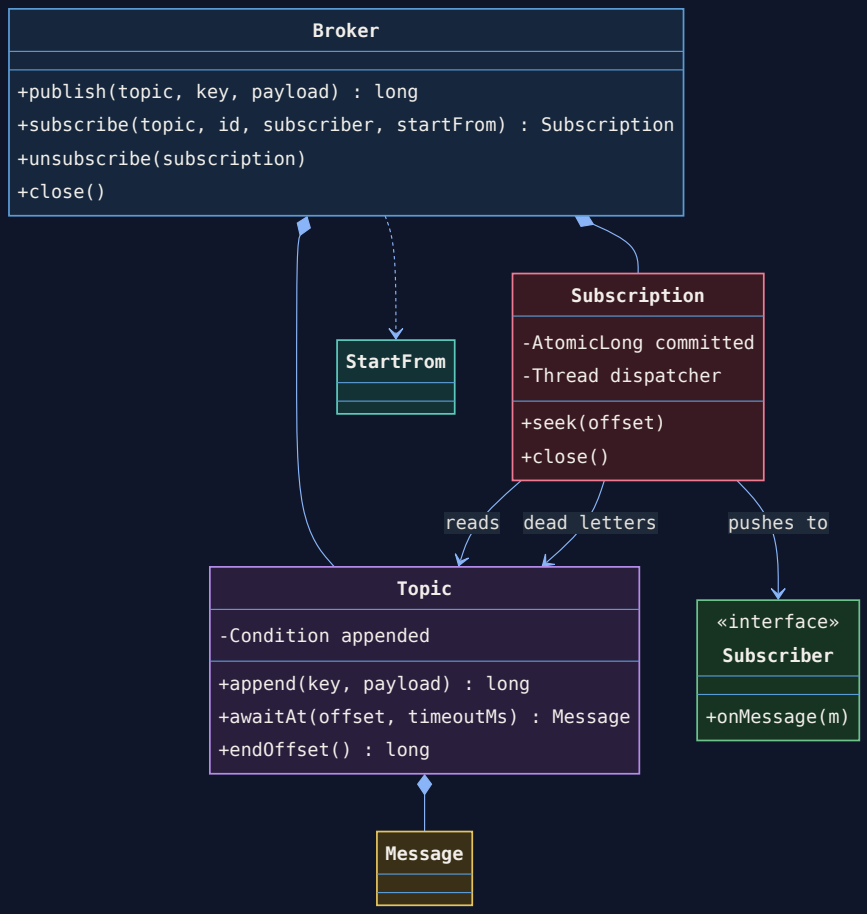
Actors: **Publisher**, **Subscriber** (sometimes fails), **Operator** (creates topics, replays, inspects dead letters), **Broker** (stores, dispatches, retries).

Core entities

Class	Responsibility
`Message`	topic, offset, key, payload; immutable.
`Topic`	Append-only `List<Message>` guarded by a `ReentrantLock` + `Condition` "appended".
`Subscriber`	Callback `onMessage`; throwing = negative ack.
`Subscription`	Committed offset (`AtomicLong`), **own dispatcher thread**, retry + dead-letter, `seek` .
`StartFrom`	EARLIEST / LATEST.
`Broker`	Registry, publish, subscribe, unsubscribe, DLQ, `awaitIdle` (tests), close.

Key idea: messages are not removed when consumed. Each subscription only remembers *how far it has read*. That gives late joiners, replay (`seek`) and independent progress for free — Kafka's model. (RabbitMQ-style queues delete on ack.)

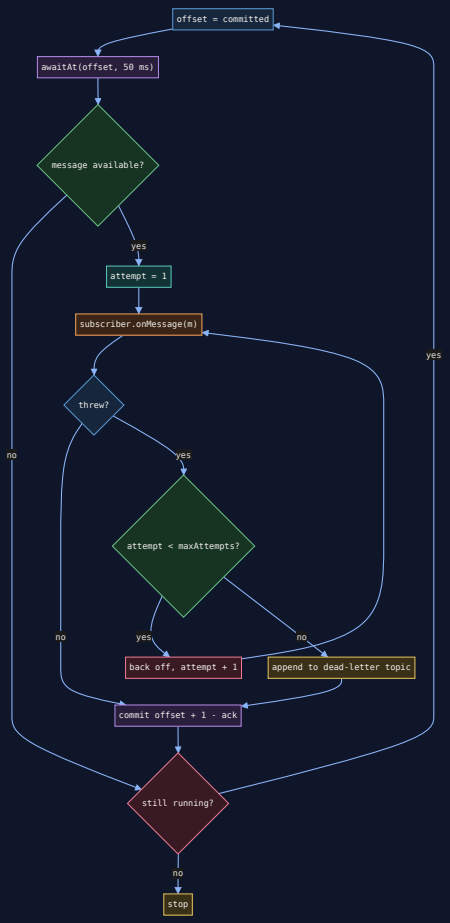
Class diagram



Key flows

Publish and fan-out: dispatchers sleep on the topic's `Condition`; `publish` appends (offset N) and `signalAll()`s; each dispatcher independently delivers message N and commits N+1 — a slow one delays only itself.

One subscription's dispatcher loop (at-least-once with dead letters):



Design decisions & patterns used

Decision	Why	Alternative rejected
Publish-Subscribe (Observer via broker)	Add subscribers with zero publisher changes.	Publisher calling consumers directly.
Append-only log + offsets	Replay, late joiners, independent progress; order = log order.	Delete on delivery; per-subscriber copies.
One dispatcher per subscription (virtual threads)	Slow subscriber delays only itself; one deliverer keeps order; virtual threads make it cheap.	One shared thread (slowest blocks all); pool per message (loses order).
`Lock` + `Condition` with timeout, in a `while`	No busy-waiting; timeout notices `close()` / `seek()` ; `while` handles spurious wake-ups.	`Thread.sleep` polling.
Ack after processing	At-least-once: a crash re-delivers.	Commit first = at-most-once.
Dead-letter topic	A poison message can't block forever.	Infinite retries.
`compareAndSet` on commit	A concurrent `seek()` wins.	Plain `set` .

Delivery guarantees:

Guarantee	How	Risk
At-most-once	Commit, then process	Crash loses the message
At-least-once (ours)	Process, then commit	Duplicates → subscribers must be **idempotent**
Exactly-once	At-least-once + dedup / transactions committing output and offset together	Complex; Kafka offers it within its own transactions

Java implementation

```
// File: code/lld/ch20/pubsub/Topic.java
    // ... append(): under the lock, add to the log and appended.signalAll()
    /** Returns the message at 'offset', waiting up to the timeout for it to be published. */
    Message awaitAt(long offset, long timeoutMs) throws InterruptedException {
        lock.lock();
        try {
            long nanos = TimeUnit.MILLISECONDS.toNanos(timeoutMs);
            while (offset >= log.size()) {         // loop: guards against spurious wake-ups
                if (nanos <= 0) return null;
                nanos = appended.awaitNanos(nanos);
            }
            return log.get((int) offset);
        } finally {
            lock.unlock();
        }
    }

// File: code/lld/ch20/pubsub/Broker.java  (Subscription)
    private void loop() {
        try {
            while (running) {
                long offset = committed.get();
                Message m = topic.awaitAt(offset, 50);
                if (m == null) continue;              // timed out: re-check 'running' and any seek()
                deliver(m);
                committed.compareAndSet(offset, offset + 1); // ack; a concurrent seek() wins
            }
        } catch (InterruptedException e) {
            Thread.currentThread().interrupt();
        }
    }

    private void deliver(Message m) throws InterruptedException {
        for (int attempt = 1; attempt <= maxAttempts; attempt++) {
            try {
                subscriber.onMessage(m);
                return;
            } catch (Exception e) {
                if (attempt == maxAttempts) {
                    deadLetters.append(id, m.topic() + "@" + m.offset() + " " + m.payload() + " -> " +
e.getMessage());
                    return;                           // skip it so one poison message cannot block the rest
                }
                Thread.sleep(5L * attempt);           // simple backoff
            }
        }
    }
    // ...
    void seek(long offset) { committed.set(offset); }
```

Full runnable code: code/lld/ch20/pubsub/ (see the Full Edition).

Output (excerpt):

```
inventory [0:order#1, 1:order#2, 3:order#3]  (offset 1 retried, offset 2 dead-lettered)
dead letter: orders@2 BAD-JSON -> cannot parse
live after replay [4:order#4, 3:order#3, 4:order#4] (order#4 seen twice: at-least-once)
```

Concurrency & edge cases

- Publishers: `append` under a lock → unique, gap-free offsets; messages immutable.
- **Unbounded log:** real brokers apply **retention** (by age/size); a lagging subscriber may then fall off the end → detect "offset out of range" and reset.
- **Subscriber blocks forever:** stalls only its own subscription; add a processing timeout. Java 21 caveat: a virtual thread blocking inside `synchronized` *pins* its carrier thread — prefer `ReentrantLock` in subscriber code, or platform threads.
- **Duplicates:** subscribers must be idempotent (remember processed ids).
- **Shutdown:** `close()` flips `running` and joins; the 50 ms timeout bounds it.

Extensions & follow-up questions

- **Consumer groups:** split a topic into **partitions**; same key → same partition (`hash(key) % n`) keeps per-key order; each group member owns partitions; offset per (group, partition). That is how Kafka scales.
- **Persist:** segment files + offset index; committed offsets in a special topic (`__consumer_offsets`).
- **Pull:** `poll(id, max)` + `commit(offset)` → back-pressure for free.
- Filtering by predicate, wildcard topics (`orders.*`), delayed retry topics (`orders.retry.5s`) instead of sleeping.

Real-world example

Apache Kafka is this model at scale (partitioned append-only logs, consumer offsets, consumer groups, per-partition order, retention by time/size). **Google Cloud Pub/Sub** and **AWS SNS + SQS** offer managed fan-out with acks and **dead-letter** queues; Spring Kafka's `DeadLetterPublishingRecoverer` is our `__dead_letters`. In-JVM: Guava `EventBus`, Spring `ApplicationEventPublisher`.

In the interview

> **30-second answer:** "Each topic is an append-only log; publishing appends and returns an offset. Each subscription keeps a committed offset and its own dispatcher thread that waits on a condition, calls the subscriber, and commits only after success — at-least-once, in order. Failures retry with backoff, then go to a dead-letter topic. Messages aren't deleted on read, so late subscribers start at the beginning and anyone can replay. To scale: partitions by key, consumer groups."

Follow-ups: *Slow consumer?* → One dispatcher per subscription; the log is shared, progress isn't. *Exactly-once?* → Hard; at-least-once + idempotent consumers. *Order?* → One writer assigns offsets, one dispatcher per subscription; with partitions, order is per key.

Self-check

1. What is an offset? → A message's position in the log; subscribers remember the next one.
2. Commit before or after processing for at-least-once? → After.
3. Why a dead-letter topic? → One poison message can't block a subscription forever.
4. Why does replay work here but not in a delete-on-ack queue? → Messages stay; seek moves the offset.
5. Why `while` around `Condition.await`? → Spurious wake-ups.

TL;DR

- **Pub-sub = Observer through a broker.**
- **Append-only log + per-subscription offsets** → fan-out, late joiners, replay.
- **One dispatcher per subscription** (isolation + order); **ack after processing** = at-least-once.
- Retry, then **dead-letter**; subscribers **idempotent**; scale with partitions and consumer groups.

20.7 Design a Rate Limiter

Problem & clarifying questions

What it is. Decides whether a request is allowed now, so one client can't overwhelm a service. Analogy: a turnstile allowing 10 people per minute per ticket type; over the limit: "come back in 6 seconds" — HTTP **429 Too Many Requests** with `Retry-After`. **Chapter 13 (the rate-limiter / token-bucket problem) already built the concurrency core** (a `synchronized` and a lock-free CAS token bucket, algorithm comparison). Here we

design the **objects around it**: pluggable algorithms, rules per API and tier, a factory, combined limits, a decision object.

Question	Assumed answer
Limit by?	(user, API); IP or API key is the same design, different key.
Different limits?	Per API and tier (FREE / PREMIUM), plus a default.
Algorithms?	Token bucket, fixed window, sliding log, sliding counter — per rule.
Several at once?	Yes: "2 per 10 s **and** 3 per minute" on `/login`.
Response?	allowed + remaining + retry-after.
Distributed?	Single process here; distributed in extensions.

Requirements

Functional: `check(userId, api)` → `Decision(allowed, remaining, retryAfter)`; rules per API + tier with a default; **one limiter per (user, api)**, created on first use; pluggable algorithms; combined rules must all allow.

Non-functional: ~O(1) per check; exactly N allowed under any concurrency; injectable clock; bounded memory.

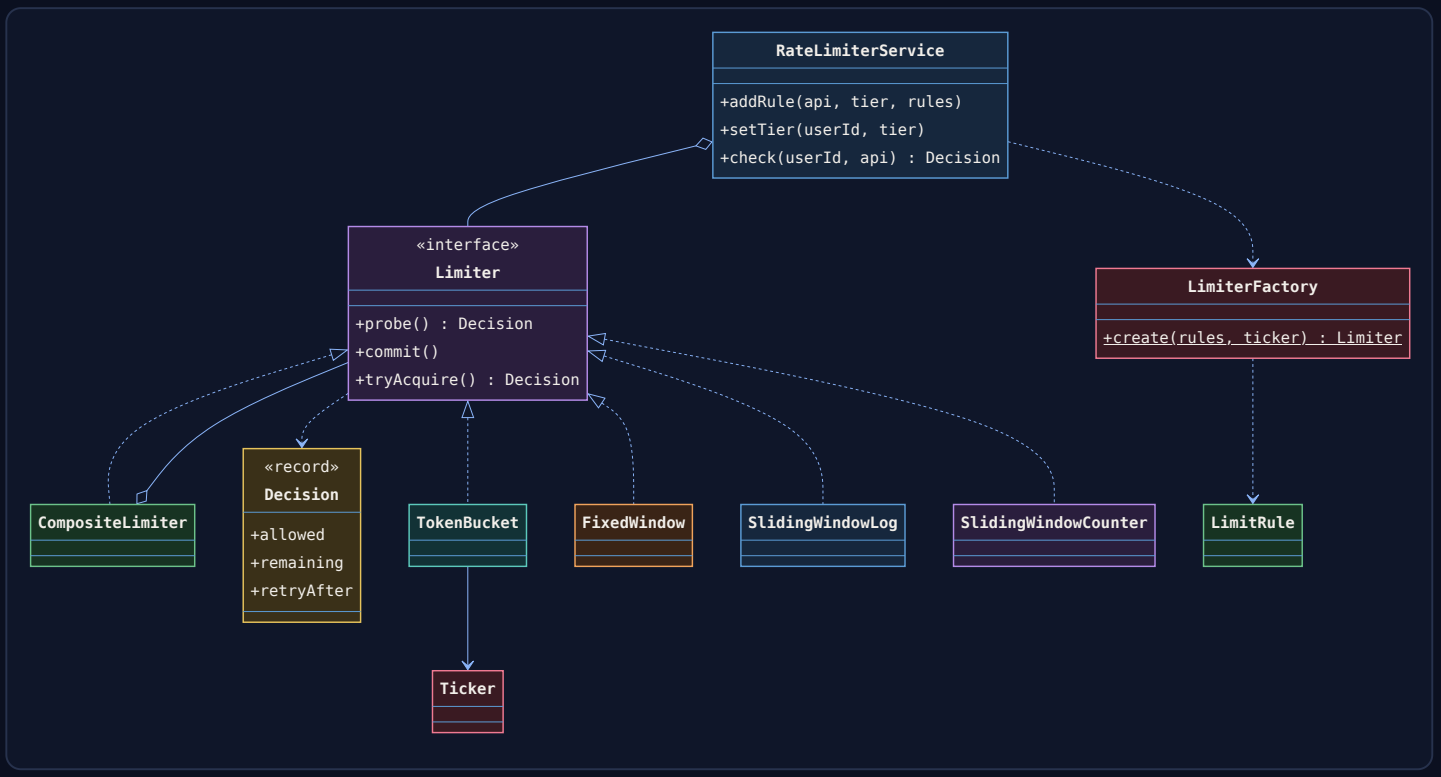
Use cases

Actors: **Client** (sends requests), **API gateway / filter** (calls `check` before the handler), **Operator** (configures rules and tiers).

Core entities

Class	Responsibility
`Limiter`	**Strategy**: `probe()` (would one more pass? consumes nothing), `commit()`, `tryAcquire()` = both atomically. One per key.
`TokenBucket`, `FixedWindow`, `SlidingWindowLog`, `SlidingWindowCounter`	The algorithms (`synchronized`).
`CompositeLimiter`	**Composite**: all children must allow; probe all, then commit all, under one lock.
`Decision`, `LimitRule`, `Algorithm`, `Tier`	Result; config (algorithm, limit, window); enums.
`LimiterFactory`	**Simple Factory**: rule(s) → limiter (single or composite).
`RateLimiterService`	**Facade**: resolves tier + rules, one limiter per (user, api) in a `ConcurrentHashMap`.
`Ticker`	Injected nanosecond clock.

Class diagram



Key flows

Gateway → `check("ann", "/search")` → `limiters.computeIfAbsent(key)` (first time: factory builds the limiter from the FREE-tier rules) → `tryAcquire()` → allowed: call the handler, 200; denied: **429 + Retry-After**.

The four algorithms (limit N per window W):

Algorithm	State per key	Bursts	Accuracy	Notes
Token bucket	tokens, last refill	Up to capacity	Good	Most common for APIs
Fixed window counter	window start, count	**Up to 2N across a boundary**	Poor at edges	Simplest
Sliding window log	one timestamp per request	None beyond N	Exact	O(N) memory per key
Sliding window counter	2 counters	Approximate	Good	previous × overlap + current
(Leaky bucket)	a queue	Constant output rate	—	Traffic shaping; not implemented

Design decisions & patterns used

Decision	Why	Alternative rejected
Strategy (`Limiter`)	Algorithms differ in accuracy/memory; new one = new class.	`switch` in `check` .
Simple Factory	One place turns config into objects.	`new TokenBucket` in the service.
Composite, probe all then commit all	"10/s and 1000/day" looks like one limiter; a request denied by one child uses up nothing in another.	Calling each child's `tryAcquire` in turn (a later denial leaves earlier children charged).
Facade service	Gateway asks one question.	Gateway juggling tiers/rules.
`computeIfAbsent` per key	Exactly one limiter even if the first two requests race.	`get` then `put` .
`Decision` with remaining/retryAfter	Maps to `X-RateLimit-Remaining` , `Retry-After` .	`boolean` .
Injected `Ticker` ; `synchronized` probe+commit	Tests never sleep; refill + check + take is one atomic step.	`nanoTime()` inside; unsynchronized fields.

Java implementation

```java
// File: code/lld/ch20/ratelimiter/Limiters.java
interface Limiter {
    Decision probe();

    void commit();

    /** Check and take in ONE atomic step (probe + commit under this limiter's lock). */
    default Decision tryAcquire() {
        synchronized (this) {
            Decision d = probe();
            if (d.allowed()) commit();
            return d;
        }
    }
}

/** Token bucket: bursts up to 'capacity', refills at 'perSecond'. Lazy refill, no background thread. */
final class TokenBucket implements Limiter {
    // ... capacity, refillPerNano, ticker, tokens, last
    @Override
    public synchronized Decision probe() {
        long now = ticker.nanos();
        tokens = Math.min(capacity, tokens + (now - last) * refillPerNano); // refilling is safe to do early
        last = now;
        if (tokens >= 1) return Decision.allow((long) (tokens - 1));
        long waitNanos = refillPerNano == 0 ? Long.MAX_VALUE : (long) Math.ceil((1 - tokens) / refillPerNano);
        return Decision.deny(Duration.ofNanos(waitNanos));
    }

    @Override
    public synchronized void commit() { tokens -= 1; }
}
// ... FixedWindow, SlidingWindowLog, SlidingWindowCounter

final class CompositeLimiter implements Limiter {
    private final List<Limiter> children;
    // ...
    @Override
    public synchronized Decision probe() {
        long remaining = Long.MAX_VALUE;
        for (Limiter l : children) {
            Decision d;
            synchronized (l) { d = l.probe(); }
            if (!d.allowed()) return d;           // first denial wins; nothing has been consumed
            remaining = Math.min(remaining, d.remaining());
        }
        return Decision.allow(remaining);
    }

    @Override
    public synchronized void commit() {
        for (Limiter l : children) synchronized (l) { l.commit(); }
    }
}

// File: code/lld/ch20/ratelimiter/RateLimiterService.java
    Decision check(String userId, String api) {
        // computeIfAbsent creates exactly ONE limiter per (user, api), even under a race
        Limiter limiter = limiters.computeIfAbsent(new Key(userId, api), k -> {
            Tier tier = tiers.getOrDefault(userId, Tier.FREE);
            List<LimitRule> r = rules.getOrDefault(api, Map.of()).getOrDefault(tier, defaultRules);
            return LimiterFactory.create(r, ticker);
        });
        return limiter.tryAcquire();
    }
```

Full runnable code: code/lld/ch20/ratelimiter/ (see the Full Edition).

Composite children are created for it alone by the factory and never shared, so the composite's lock makes probe-all + commit-all atomic.

Output (excerpt):

```
  t=0.0 ann /search #4 -> 429 Too Many Requests (retry after 1000 ms)
== 4. Boundary burst: limit 5 per 10 s, 5 calls at t=9.9 and 5 at t=10.1
  FIXED_WINDOW     allowed 10 of 10 within 0.2 s
  SLIDING_LOG      allowed 5 of 10 within 0.2 s
  SLIDING_COUNTER  allowed 5 of 10 within 0.2 s
  allowed = 50 (exactly 50, never more)
```

The classic point: **fixed window lets 10 through in 0.2 s** across a boundary; both sliding algorithms hold at 5. 640 concurrent attempts on a 50-token bucket → exactly 50.

Concurrency & edge cases

- Contention is per key (each limiter `synchronized`; map is `ConcurrentHashMap`), so different users never block each other.
- **Naive composite bug:** child 1 allows (and consumes), child 2 denies → child 1 is charged for a request that never ran. The demo's "3 per minute first" case shows the +11s call would wrongly be refused for ~49 s. Probe-then-commit fixes it (acquire-and-refund is the harder alternative).
- **Memory:** one limiter per key forever → evict idle limiters (expire-after-access cache, as in 20.3) with idle timeout ≥ the longest window — a fresh limiter then behaves identically.
- Monotonic clock; rule changes at runtime need clearing or versioning the limiter map.

Extensions & follow-up questions

- **Distributed:** shared state in Redis with an atomic Lua script (or `INCR` + `EXPIRE` for fixed windows); costs a network hop — mitigate with local token pre-allocation or sticky routing.
- **Leaky bucket:** bounded queue drained at a constant rate (protects fragile downstreams).
- Cost per request (`tryAcquire(permits)`), soft limits (warn at 80%), key by IP when logged out; runs in an API gateway or servlet filter / Spring `HandlerInterceptor`.

Real-world example

Stripe described combining a token-bucket request limiter, a concurrent-requests limiter and load shedders on Redis — the composite idea. **GitHub** returns `X-RateLimit-Limit/Remaining/Reset` (our `Decision`). HTTP 429 is RFC 6585. **NGINX** `limit_req` uses leaky bucket. Java: Guava `RateLimiter`, Resilience4j, **Bucket4j**. Cloudflare uses the sliding-window-counter approximation at the edge.

In the interview

30-second answer: "A `Limiter` interface with interchangeable algorithms — token bucket by default, sliding window when exactness matters. A factory builds limiters from rules chosen by API and tier; a composite combines '10 per second and 1000 per day'. The service keeps one limiter per (user, api) via `computeIfAbsent` and returns remaining and retry-after for the 429. Each limiter makes refill-check-take atomic; for many servers, move state into Redis with an atomic script."

Follow-ups: *Why is fixed window risky?* → Up to 2N across a boundary. *Token bucket vs sliding log?* → Bucket: bursts, O(1) memory; log: exact, stores every timestamp.

Self-check

1. Burst up to capacity, then a steady rate? → Token bucket.
2. Fixed window of 5 across a boundary in 0.2 s? → Up to 10.
3. Pattern making two rules look like one limiter? → Composite.
4. Why must refill + check + take be atomic? → Two threads could both take the last token.
5. Why probe every child before committing? → A request denied by one rule must not use a slot in another.

TL;DR

- **Strategy** algorithms, **Factory** from rules, **Composite** (probe all, commit all), **Facade** service.
- One limiter per (user, api) via `computeIfAbsent`; `Decision` → 200 or 429 + Retry-After.

- Token bucket (bursty, O(1)), fixed window (2× edge burst), sliding log (exact, O(N)), sliding counter (approx, O(1)).
- Atomic refill-check-take (Chapter 13); Redis + Lua when distributed.

20.8 Chapter summary

Component	Anchor pattern(s)	The one sentence to remember
Splitwise	Strategy	Split rules return shares; ledger = one signed number per pair; simplify with net positions + greedy (≤ n − 1).
Ledger Co. (full case in Chapter 21)	—	A loan ledger: same "one authoritative record + `BigDecimal`" idea, different rules.
Notification service	Template Method + Strategy + Observer	Retry once in the parent, one class per reason not to send, one exit that tells listeners.
In-memory cache	Strategy + Observer	HashMap + doubly linked list; policy owns the order; one removal path.
LRU cache	— (data structure)	Map to node, move to front on use, evict before tail; or `LinkedHashMap(accessOrder=true)`.
Ride sharing	Strategy + Decorator + state machine	Ranked list for declines, surge wraps pricing, atomic driver claim.
Carpool variant	Strategy	Offer seats per vehicle; select with Most Vacant / Preferred Vehicle.
Logging framework	Chain of Responsibility + Singleton + Decorator	Hierarchy with inherited levels; propagate up until additivity stops; async is a wrapper.
Pub-sub	Observer via broker	Append-only log + offsets; ack after processing; dead letters.
Rate limiter	Strategy + Factory + Composite	One limiter per key; swappable algorithms; 429 + Retry-After.

Common thread: the part most likely to change sits behind an interface, and the service only orders the steps. "Now add X" → a new class plus a one-line wiring change.

20.9 Chapter self-check

1. Two patterns in more than half the designs? → Strategy and Observer.
2. Strategy vs Decorator? → Strategy *replaces* a rule; Decorator *wraps* one and adds to its result.
3. Strategy vs Template Method? → Strategy swaps the whole algorithm (composition); Template Method fixes it and lets subclasses fill steps (inheritance).
4. Which designs used a "single exit", and why? → Notification (`finish`) and cache (`remove`): records, stats and listeners can't be skipped.
5. Three uses of CAS / `computeIfAbsent` for correctness? → Driver claim (`AtomicBoolean`), subscription commit (`AtomicLong.compareAndSet`), limiter creation (`computeIfAbsent`).

20.10 Chapter TL;DR

- Interviews probe the **rule most likely to change** — put it behind an interface.
- **Strategy** everywhere; **Observer** for "tell someone"; **Decorator** to add behaviour (surge, async); **Template Method** for a shared algorithm with one varying step; **Chain of Responsibility** for the logger hierarchy; **Composite** for combined limits.
- `BigDecimal` money; injected clocks; the smallest correct concurrency tool (one lock, CAS, `computeIfAbsent`, per-key locks, sharding).

- Volunteer the honest limits: greedy isn't optimal, LRU needs a write lock, at-least-once duplicates, fixed windows burst at edges.

Chapter 21 — Case Studies: Machine-Coding Round Problems (Study Edition)

> **Goal.** Chapters 14–20 taught whiteboard design. This chapter trains the **machine-coding round**: a written problem, 90–120 minutes, and **running code** the interviewer feeds their own inputs into. You get the scoring rubric, a skeleton to type from memory, a time plan, and eight real problems solved end to end in Java.

Each problem follows the usual template (problem → requirements → use cases → entities → diagram → flows → decisions → code → edge cases → extensions → interview → self-check → TL;DR). Code is Java 21, standard library only, one package per problem under `code/lld/ch21/<name>/`. From `chapters/21-case-machine-coding/code`:

```
javac --release 21 -d /tmp/lld_21 $(find . -name '*.java')
java -cp /tmp/lld_21 lld.ch21.taskmgmt.Main       # or wallet, calendar, codingblox, propertyhunt, ledger, inventory, filesearch
```

#	Problem	Source	The one idea that makes it work	Patterns
21.1	Task Management System	awesome-low-level-design	Status *state machine* + composable filters	State (enum), Specification, Strategy, Observer
21.2	Digital Wallet Service	awesome-low-level-design	`BigDecimal` money, idempotency keys, lock ordering	Strategy (offers), Facade
21.3	Event Calendar	Flipkart (kumaransg/LLD)	Half-open intervals + a boundary *sweep* for free slots	Strategy, Repository
21.4	Coding Blox	Flipkart (kumaransg/LLD)	Command dispatcher + level-based scoring	Command, Strategy, Repository
21.5	PropertyHunt	ClearTrip (kumaransg/LLD)	Optional criteria → list of predicates; unit normalisation	Builder, Specification, Strategy
21.6	Ledger Co (loans & EMIs)	Navi (kumaransg/LLD, Geektrust)	Derive balances from facts, don't store them	Strategy (interest), Command
21.7	Inventory Management	Common (Amazon/Flipkart-style)	Append-only movement ledger + per-SKU lock	Strategy (allocation), Observer (reorder)
21.8	Unix File Search (`find`)	Common OOD/machine-coding question	Specification filters + Composite AND/OR/NOT; Builder for criteria	Specification, Composite, Builder

(Food ordering "FoodKart" is solved in Chapter 19.)

21.0 How the machine-coding round works

What it is

An OO design interview where **working code is the answer**: a 1–2 page statement, your IDE, 90–120 minutes, no database/UI/framework — data in maps and lists, driven by `main`, a command list or tests. Unlike system design (servers, databases, queues) it never leaves one program; its vocabulary is **classes**. Analogy: whiteboard design *describes* a house; machine coding *builds a shed* while the inspector leans on every wall. Common at Flipkart, Swiggy, PhonePe, Navi, Razorpay, CRED, Meesho, Cleartrip.

How it is judged

Criterion	What is checked → how to score well
Working, demo-able code	Compiles, runs the sample, survives *their* inputs → Thin end-to-end path in the first hour; never finish with non-compiling code
Functional completeness	Every mandatory requirement; bonus after → Implement in the statement's order
Clean models	One job per entity, private fields, invariants inside the entity → Enums for fixed sets, records for values, classes for changing things
Separation of concerns	Parsing, logic, storage apart; no `System.out` in services → Driver → Service → Repository
Extensibility	"Add a new offer/sort/rule" = add a class → Strategy / Specification for "can be X or Y"
Edge cases	Unknown ids, duplicates, invalid input, wrong user → Custom exceptions; driver catches per command and **keeps going**
Readability	Names, small methods, no copy-paste → Nouns → classes, verbs → methods
Concurrency (when asked)	Wallets, inventory, booking → Say where the lock is and why; demo with a pool
Tests (bonus)	Extra credit → An `assert`-style self-check in `main` suffices

The five-step method, and how the clock changes it

Whiteboard method (Chapter 1): **Requirements** → **Actors & use cases** → **Classes** → **Relationships** → **Code** (roughly 5–8 / 3–5 / 8–10 / 8–10 / 12–15 min in a 45-minute round, code covering only 2–3 key methods). **Machine coding flips the proportions:** steps 1–4 take ~20 minutes, code the rest. A parking lot has ~12 classes but only 3 carry the main flows — build those, name the rest. Also judged: **scope**, **one job per class**, **relationships** (hold/own/extend), **code and model agree**, **a stated trade-off** — never UML notation. Top failures: **nothing runs at the end**, and a **god class** `Manager`.

A reusable project skeleton

```
src/
  model/          entities & value objects      (Task, Wallet, Loan, Property ...)   - enums, records, small classes
  repository/     storage contracts + in-memory (TaskRepository, InMemoryTaskRepository)
  service/        business rules, one per area  (TaskService, WalletService)          - no printing, no parsing
  strategy/       pluggable rules               (Offer, AllocationStrategy, SortBy)
  exception/      domain exceptions             (NotFoundException, InsufficientFundsException)
  driver/         Main + CommandProcessor       (parses input, calls services, prints output)
```

Layer	Knows about	Must NOT know about	Why
model	itself, other models	services, storage, console	Entities guard invariants (`Task.changeStatus` rejects `TODO → DONE`).
repository	model	services, console	Swap `HashMap` for a DB touching only this layer (DIP).
service	model, repository & strategy *interfaces*	concrete repositories, console	The part the interviewer reads most carefully.
strategy	model	services	Each "X or Y" rule → interface + small classes (OCP).
exception	—	—	One class per failure kind, readable message.
driver	services	repository internals	Parsing/printing replaceable; logic isn't.

(This book uses one package per problem; in the round, create the sub-packages.) Reusable blocks: a **generic in-memory repository** (Coding Blox), an `AtomicInteger` **ID generator**, and a **command processor** `Map<String, Consumer<String[]>>` (Command pattern, Chapter 10).

The 90-minute time plan

Minutes	Activity	Output
0–10	Read twice, ask questions, mark mandatory vs bonus, paste sample input into the driver	Numbered requirement list
10–20	Nouns → entities, verbs → methods; pick the 1–2 patterns the statement begs for	Packages, enums, records
20–60	Mandatory requirements **in order**, running after each	Sample input → sample output
60–75	Edge cases, validation, custom exceptions	Driver prints errors and continues
75–85	Bonus (only if all mandatory work); thread-safety for money/stock	Bonus commands
85–90	Clean names, delete dead code, final run, 2-minute walkthrough	Compiling code, demo, story

Run the program at least every 15 minutes — a compile error at minute 88 is a disaster, at minute 30 a 20-second fix.

What to say in the first two minutes

1. "Everything in memory behind a repository interface, so a DB can be added later."
2. "I'll drive it from `main` with the sample commands and print errors without crashing."
3. "Mandatory is A, B, C; bonus D, E. A–C first."
4. "Where does it need to be thread-safe?" (Wallets and inventory: always.)

Common mistakes in this round

Mistake	Do this instead
20 minutes on a perfect class diagram	Enums/records first; the diagram lives in names
One `Manager` with 15 maps	One service per area + repositories
`System.out` inside services	Services return data; the driver prints
`double` for money (`0.1 + 0.2 != 0.3`)	`BigDecimal` with fixed scale (21.2) or whole units (`long` paise)
Crashing on the first bad command	Catch domain exceptions in the driver loop
Bonus before mandatory	Mandatory in order; bonus last
Every class in full detail	Core flow first; stub or name the rest
Unseeded `Random` ; `LocalDate.now()`	Seed it / Strategy (21.4); inject a `Clock` (21.1, 21.2)

Real-world examples

Standard Spring Boot layering: controllers = driver, `@Service` = rules, `@Repository` = storage; in-memory repositories are the usual unit-test fakes.

In the interview

> **30-second answer:** "I read the statement twice, list mandatory and bonus requirements and clarify ambiguities. I set up model / repository / service / strategy / exception / driver packages, put the sample input in the driver, and implement requirements one by one, running after each. Services return data and throw domain exceptions; the driver prints and continues. Then edge cases, thread-safety where state is shared, and bonus features."

Why an interface for an in-memory repository? → Services depend on an abstraction; JDBC/Redis can replace it.

Self-check

- **Q:** The six packages? → **A:** model, repository, service, strategy, exception, driver.

- **Q:** Where does `System.out.println` live? → **A:** Only in the driver.
- **Q:** Minute 60, one mandatory requirement missing, bonus tempting? → **A:** Finish the mandatory one.
- **Q:** Why inject a `Clock`? → **A:** Time-based logic gives the same output every run.

TL;DR

Working code first, clean layering second, patterns where the statement asks for variation, graceful errors everywhere. **model** → **repository** → **service** → **strategy** → **exception** → **driver**. Run often.

21.1 Design a Task Management System

Problem & clarifying questions

The core of Jira/Asana/Trello: create, update, delete tasks (title, description, due date, priority, status), assign, move through statuses, comment, search/filter (priority, due date, assignee), reminders, task history, safe concurrent access.

Question	Assumed answer
Statuses and transitions?	TODO, IN_PROGRESS, BLOCKED, DONE, CANCELLED; only listed transitions (TODO can't jump to DONE; CANCELLED final; DONE can re-open).
Assignees?	One (nullable).
"Reminder" in a console app?	A method a scheduler calls: notify owners of open tasks due within N days or overdue.
Past due dates?	Not on creation.
Notifications?	A listener interface; the demo prints.
"Task history"?	Per-task audit log (built); a user's finished tasks is just a search: `assignedTo(u).and(withStatus(DONE))`.

Requirements

Functional: register users; create/update/delete tasks; assign; change status via allowed transitions; comment; search with any filter combination (assignee, status, min priority, due before, open only, title contains) + sort (due date, priority, id); reminders; per-task history. **Non-functional:** thread-safe; new filters/sorts/channels without editing code; deterministic (injected `Clock`).

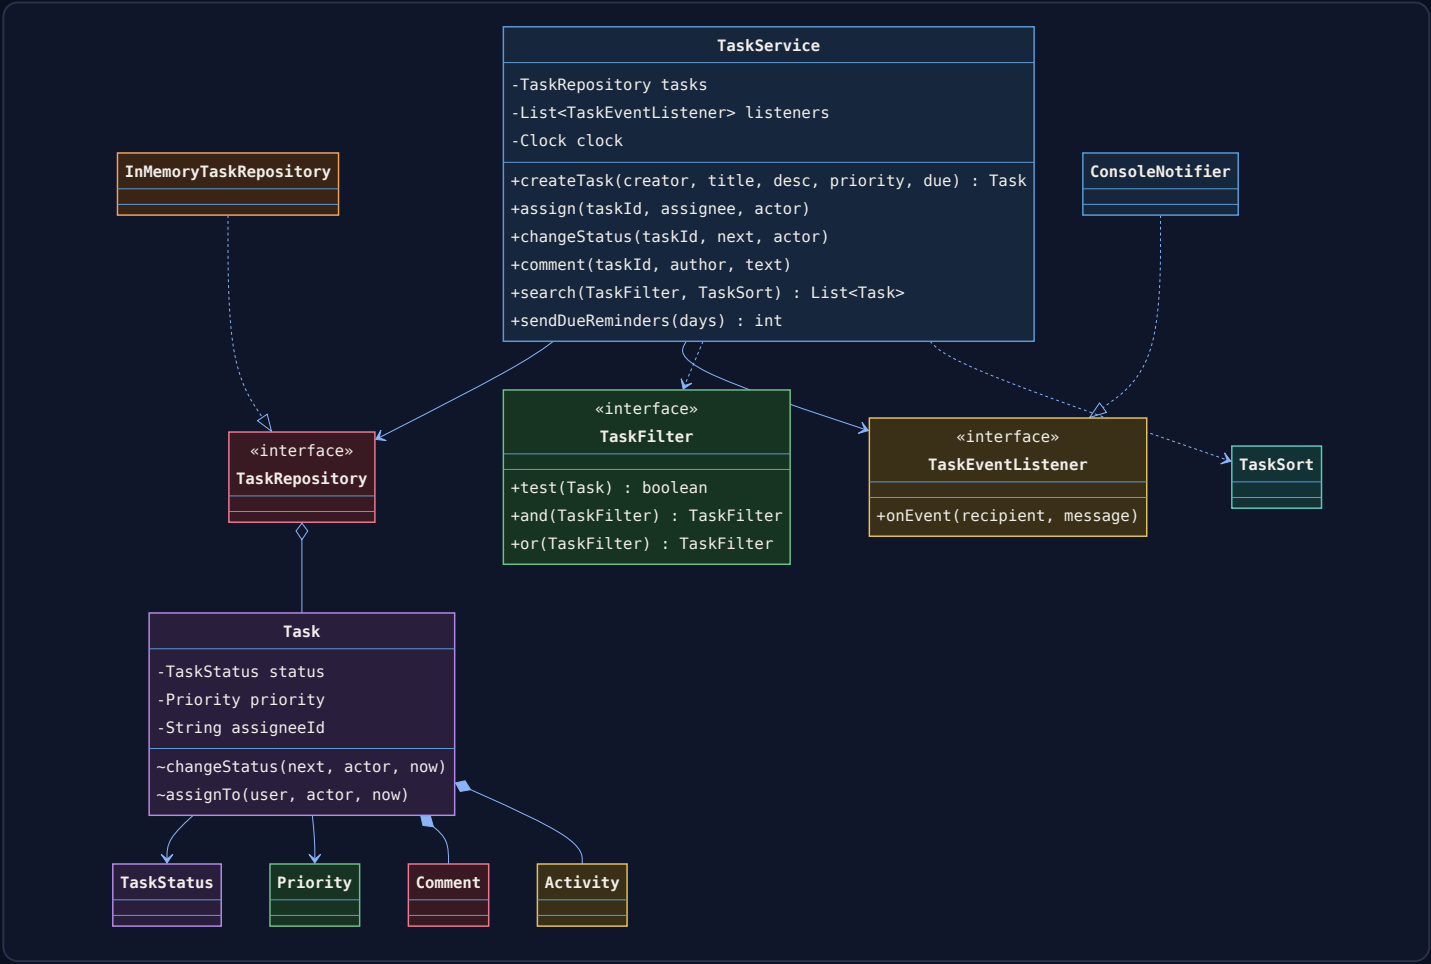
Use cases

Task creator (create, update, delete, assign, comment, search), **Assignee** (change status, comment, see own tasks), **Scheduler** (due-date reminders), **Notification channel** (receives events).

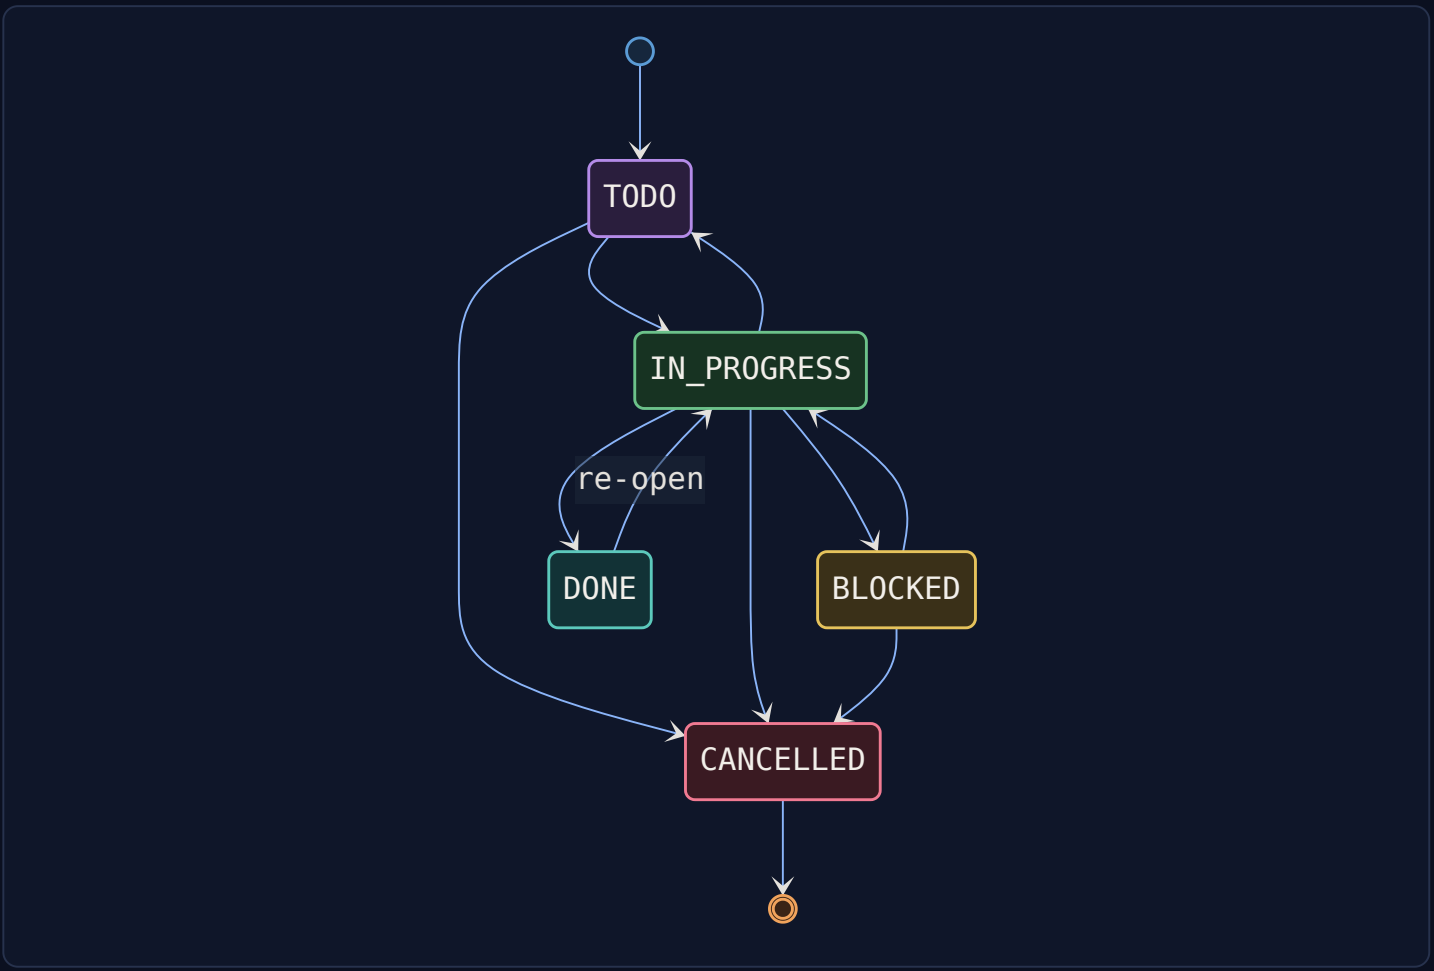
Core entities

Class	Responsibility
`Task`	Data; enforces legal status changes; records its own history; mutators `synchronized`.
`TaskStatus` (enum)	Statuses + allowed next statuses (tiny state machine).
`Priority` (enum)	LOW < MEDIUM < HIGH < CRITICAL (enum order compares).
`User`, `Comment`, `Activity`	Immutable records.
`TaskRepository` / `InMemoryTaskRepository`	Storage contract + `ConcurrentHashMap`.
`TaskFilter`	Composable predicate (**Specification**) with factories and `and` / `or`.
`TaskSort` (enum)	Named comparators (**Strategy**); `BY_ID` sorts `T2` before `T10`.
`TaskEventListener` / `ConsoleNotifier`	**Observer**.
`TaskService`	**Facade**: every use case, validation, clock, events.
`NotFoundException`, `InvalidTransitionException`	Clear failures.

Class diagram



Key flows



Change status: service finds the task → `task.changeStatus(next, actor, now)` → illegal: `InvalidTransitionException` (driver prints it); legal: log activity, set status, notify the creator.

Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
Transitions in the `TaskStatus` enum	State (enum)	One place for rules; `Task` asks `canMoveTo`.	Class per status — overkill with no per-state behaviour.
`TaskFilter extends Predicate<Task>` + `and/or`	Specification	Any combination; new filter = one factory method.	`search(assignee, status, priority, …)` with nullable params.
`TaskSort` enum of comparators	Strategy	Sort order is data.	`switch` on a string.
Listener list	Observer	Email/Slack/push without touching the service.	`System.out` in the service.
`TaskService` single entry	Facade	Simple driver API.	Driver using repositories.
Injected `Clock`	DI	Deterministic reminders/tests.	`LocalDate.now()` scattered.
`synchronized` `Task` mutators	Monitor per entity	Different tasks never block each other.	One global service lock.

Java implementation

```java
// File: lld/ch21/taskmgmt/TaskStatus.java
public enum TaskStatus {
    TODO, IN_PROGRESS, BLOCKED, DONE, CANCELLED;

    public Set<TaskStatus> allowedNext() {
        return switch (this) {
            case TODO -> EnumSet.of(IN_PROGRESS, CANCELLED);
            case IN_PROGRESS -> EnumSet.of(BLOCKED, DONE, TODO, CANCELLED);
            case BLOCKED -> EnumSet.of(IN_PROGRESS, CANCELLED);
            case DONE -> EnumSet.of(IN_PROGRESS);          // re-open
            case CANCELLED -> EnumSet.noneOf(TaskStatus.class);
        };
    }

    public boolean canMoveTo(TaskStatus next) { return allowedNext().contains(next); }

    public boolean isOpen() { return this != DONE && this != CANCELLED; }
}

// File: lld/ch21/taskmgmt/Task.java
    synchronized void changeStatus(TaskStatus next, String actor, LocalDateTime now) {
        if (!status.canMoveTo(next)) throw new InvalidTransitionException(id, status, next);
        log(now, actor, "status " + status + " -> " + next);
        status = next;
    }

// File: lld/ch21/taskmgmt/TaskFilter.java
@FunctionalInterface
public interface TaskFilter extends Predicate<Task> {

    static TaskFilter all() { return t -> true; }
    static TaskFilter assignedTo(String userId) { return t -> userId.equals(t.assigneeId()); }
    static TaskFilter withStatus(TaskStatus s) { return t -> t.status() == s; }
    static TaskFilter atLeast(Priority p) { return t -> t.priority().compareTo(p) >= 0; }
    // ... dueOnOrBefore, open, titleContains

    default TaskFilter and(TaskFilter other) { return t -> test(t) && other.test(t); }
    default TaskFilter or(TaskFilter other) { return t -> test(t) || other.test(t); }
}
```

Full runnable code: code/lld/ch21/taskmgmt/ (see the Full Edition).

Output (excerpt):

```
== reminders (due within 3 days)
  [notify carol] Reminder: T4 'Review PR #42' is due 2026-09-28
  [notify bob] Reminder: T3 'Upgrade to Java 21' is due 2026-10-01
== edge cases
  error: Task T2 cannot move from TODO to DONE
  error: Due date is in the past: 2026-09-27
  error: Cannot assign a CANCELLED task T4
```

Self-assignment sends no notification; reminders consider only open tasks; the illegal jump is rejected with a clear message.

Concurrency & edge cases

- Per-task `synchronized` mutators and getters; `ConcurrentHashMap` storage with `findAll()` returning a copy; listeners in `CopyOnWriteArrayList`; ids from `AtomicInteger`.
- Handled: unknown ids, blank title, past due date, illegal transition, assigning a closed task.
- **Not handled (say so):** lost updates on text fields (last writer wins) → optimistic locking with a `version` field.

Extensions & follow-up questions

- **Sub-tasks / dependencies:** `parentId` + `blockedBy`; forbid DONE while a blocker is open.
- **Custom workflows (Jira):** move the transition table into an injected `Workflow` (`Map<Status, Set<Status>>`) per project.

- Also: watchers (`Set<String>`), recurring tasks, `remindAt` + a `ScheduledExecutorService`, secondary indexes or Elasticsearch for search at scale.

Real-world example

Jira projects have *workflows* (statuses + allowed transitions that notify watchers) and JQL search like `assignee = currentUser() AND status != Done ORDER BY priority DESC` — compose filters, then sort, as `TaskFilter` + `TaskSort` do.

In the interview

Present: (1) statuses and transitions — shows you think in invariants; (2) `Task` and `TaskService` API; (3) filters as composable predicates; (4) observer; (5) run the demo. *Why not the full State pattern?* → Statuses have no behaviour besides legal transitions, so an enum table is simpler and still closed to illegal moves.

Self-check

- **Q:** Where is "TODO cannot go to DONE" coded? → **A:** `TaskStatus.allowedNext()`, enforced by `Task.changeStatus`.
- **Q:** Add a "created by X" filter? → **A:** `static TaskFilter createdBy(String u) { return t -> u.equals(t.createdBy()); }` — nothing else changes.
- **Q:** Why a lock per `Task`, not in the service? → **A:** Edits to different tasks never contend.

TL;DR

Task guards its own state machine and history; service = facade; filters = Specification, sorts = Strategy, notifications = Observer; thread-safety per task.

21.2 Design a Digital Wallet Service

Problem & clarifying questions

A wallet like Paytm, PhonePe or a PayPal balance: wallets, top-up, transfer, statement. The awesome-low-level-design statement also lists accounts, payment methods, external transfers, multi-currency, security and **concurrent transactions with consistency**. **Cashback** is added here (common in variants; shows Strategy and double-entry). The interviewer focuses on **money arithmetic**, **no double-spending under concurrency**, and **no double-charging on retries** (idempotency).

Question	Assumed answer
Money type?	`BigDecimal`, scale 2; > 2 decimals rejected, not rounded.
Negative balance?	No.
Client retries after a timeout?	Each request has an **idempotency key**; a repeat returns the original result and moves no money.
Offers?	"10% on first transfer, max 50" and "flat 25 on transfers ≥ 1000"; no stacking — best wins.
Cashback source?	A funded `PROMO` **wallet**; cashback = ordinary transfer PROMO → sender, so money is never created. Budget gone → no cashback, transfer stands.
Multi-currency?	Out of scope (extensions).

Requirements

Functional: create wallet; top up; atomic transfer; statement with running balance; best cashback from the PROMO wallet; idempotent top-up and transfer. **Non-functional:** no lost updates, negative balances or deadlocks; **total money conserved** (only top-ups add money); exact decimals; new offers without touching transfer code.

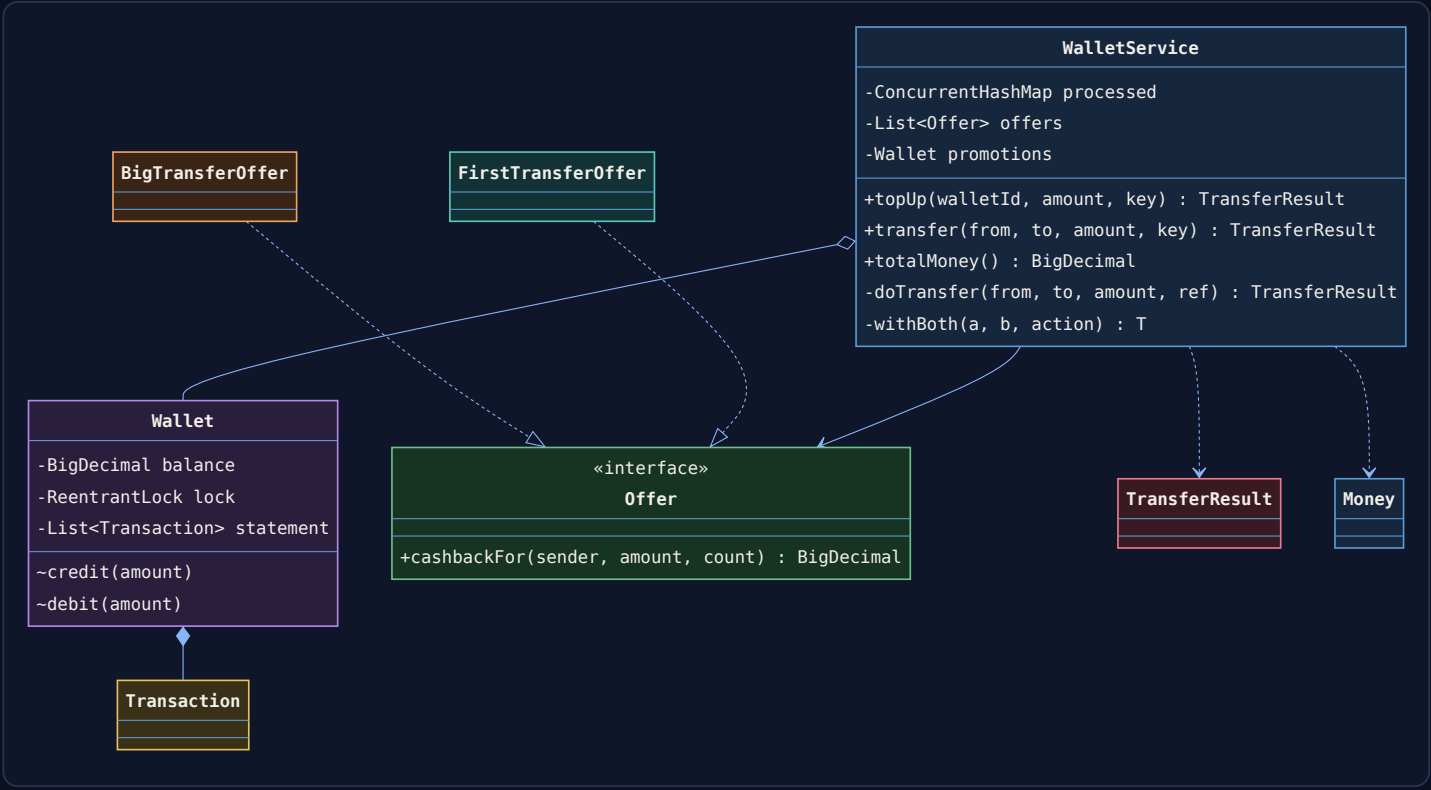
Use cases

User (create wallet, top up, transfer, statement), **Bank/card gateway** (top-up source), **Marketing** (defines offers, funds PROMO), **Client app** (retries with the same key).

Core entities

Class	Responsibility
`Wallet`	Balance + own statement + its `ReentrantLock`; `debit` refuses to go below zero.
`Transaction` (record)	Immutable line: type, amount, balance after, counterparty, reference.
`TxnType`	TOP_UP, TRANSFER_OUT/IN, CASHBACK_OUT (on PROMO)/IN.
`TransferResult` (record)	Stored for idempotent replays.
`Money`	Validates/normalises amounts.
`Offer` + `FirstTransferOffer`, `BigTransferOffer`	**Strategy**: cashback for a completed transfer.
`WalletService`	**Facade**: wallets, lock ordering, idempotency, offers.
`InsufficientFundsException`, `WalletNotFoundException`	Failures.

Class diagram



Key flows

Transfer: validate → `processed.computeIfAbsent(key, doTransfer)` (a known key returns the stored result) → lock both wallets **in id order** → debit, credit, record TRANSFER_OUT/IN → unlock → best offer → second locked step PROMO → sender (PROMO sorts before W1), recording CASHBACK_OUT/IN → store result.

Design decisions & patterns used

Decision	Why	Alternative rejected
`BigDecimal` scale 2, reject > 2 decimals	Exact money; no silent rounding of `1.005`.	`double` (wrong); `long` paise is also fine — say so.
Lock per wallet, acquired in id order	A→B and B→A would deadlock if each locked its source first; a global order makes a cycle impossible.	One global lock (serialises every transfer).
Debit before credit, under both locks	`debit` throws before any change → atomic.	Credit first + compensation.
Idempotency via `computeIfAbsent`	Check + run is one atomic step per key; if it throws, nothing is stored and the client may retry.	`containsKey` then `put` (race).
Cashback from a funded `PROMO` wallet	Double-entry: every credit has a debit; total changes only on top-ups; campaigns can't overspend.	Crediting from thin air.
Cashback as a second step, after the transfer locks	At most two locks at once; transfers don't queue on PROMO.	Locking three wallets per transfer.
`Offer` interface, best wins	New promotions = new classes; no-stacking in one method.	`if/else` on codes.

`computeIfAbsent` **detail:** the call is atomic, the function runs only if absent, an exception records nothing, and other updates to the *same map* may block while it runs — so keep the function short and never update the map inside it (ours only touches wallets).

Java implementation

```java
// File: lld/ch21/wallet/WalletService.java
    public TransferResult transfer(String fromId, String toId, BigDecimal amount, String idempotencyKey) {
        if (fromId.equals(toId)) throw new IllegalArgumentException("Cannot transfer to the same wallet");
        BigDecimal amt = Money.requirePositive(amount);
        Wallet from = get(fromId), to = get(toId);
        return processed.computeIfAbsent(idempotencyKey, key -> doTransfer(from, to, amt, key));
    }

    private TransferResult doTransfer(Wallet from, Wallet to, BigDecimal amt, String ref) {
        // Step 1: the transfer itself, under both wallet locks.
        int n = withBoth(from, to, () -> {
            from.debit(amt);                        // throws before any change if funds are short
            to.credit(amt);
            from.record(txn(ref, from, TxnType.TRANSFER_OUT, amt, to.id()));
            to.record(txn(ref, to, TxnType.TRANSFER_IN, amt, from.id()));
            return transferCounts.computeIfAbsent(from.id(), k -> new AtomicInteger()).incrementAndGet();
        });
        // Step 2: cashback is a second, ordinary transfer PROMO -> sender (same lock-ordering rule).
        BigDecimal offered = bestOffer(from, amt, n);
        // ... withBoth(promotions, from, ...): if the budget covers it, debit PROMO, credit sender,
        // ... record CASHBACK_OUT / CASHBACK_IN; otherwise cashback = 0
        return new TransferResult(ref, from.id(), to.id(), amt, cashback);
    }

    /** Locks two wallets, always the smaller id first, so two opposite operations can never deadlock. */
    private static <T> T withBoth(Wallet a, Wallet b, Supplier<T> action) {
        Wallet first = a.id().compareTo(b.id()) < 0 ? a : b;
        Wallet second = first == a ? b : a;
        first.lock().lock();
        try {
            second.lock().lock();
            try { return action.get(); } finally { second.lock().unlock(); }
        } finally { first.lock().unlock(); }
    }
```

Full runnable code: code/lld/ch21/wallet/ (see the Full Edition).

Output (excerpt):

```
  TransferResult[reference=req-3, from=W1, to=W2, amount=200.00, cashback=20.00]
  retry returns same result? true
  error: At most 2 decimal places: 1.005
  failed req-5 can be retried later: 10.00
== concurrency: 1000 opposite transfers on 8 threads
  total before=2400.50 after=2400.50 conserved=true
```

Retried `req-3` moved no money; failed `req-5` stored nothing, so its key worked later; the total was conserved with no deadlock.

Concurrency & edge cases

- Every read-modify-write of a balance holds the wallet's lock; all two-lock paths go through `withBoth`, so the order can't be forgotten; never more than two locks held.
- **Same key, different body** (client bug): today the stored result is returned; production also stores a request hash and rejects mismatches.
- Cashback runs after the transfer's locks are released, so briefly a transfer is visible without its cashback — fine for a reward.

Extensions & follow-up questions

- **Multi-currency:** `Currency` per wallet + `FxRateProvider`; store both amounts and the rate. **Payment methods / payouts:** `PaymentMethod.charge` before crediting; `PENDING_PAYOUT` + `PayoutGateway`. **Refunds:** post a compensating transaction, never edit history.
- **Distributed:** row locks (`SELECT ... FOR UPDATE`) or optimistic versions; idempotency keys under a unique constraint.
- **Slow transfers blocking the map:** `putIfAbsent(key, new CompletableFuture<>())`; winner runs outside the map, duplicates `join()`; remove the key on failure.
- **Limits (daily cap, KYC):** validators in a Chain of Responsibility before `doTransfer`.
- **Why not `AtomicReference<BigDecimal>` CAS?** A transfer changes two balances together; CAS on two variables isn't atomic.

Real-world example

Stripe's `Idempotency-Key` header returns the saved result of the first attempt on a retry instead of charging twice. Paytm/PhonePe cashback rules ("X% up to a cap on your first transaction") are `FirstTransferOffer`.

In the interview

Lead with the three risks (exact money, concurrency, retries — and "cashback comes out of a budget wallet"). Show `Money`, then `doTransfer` with lock ordering (draw the A→B / B→A deadlock), then idempotency; run the concurrent demo and point at `conserved=true`. Offers are a 2-minute Strategy mention.

Self-check

- **Q:** Why `new BigDecimal("0.1")` over `new BigDecimal(0.1)`? → **A:** The double constructor captures the binary approximation; the String one is exact.
- **Q:** W1→W2 and W2→W1 at once — why no deadlock? → **A:** Both lock W1 first.
- **Q:** A transfer fails for insufficient funds; reuse the key? → **A:** Yes — `computeIfAbsent` stores nothing when the function throws.
- **Q:** Where does "max 3 transfers per minute" go? → **A:** A validator before `doTransfer`.

TL;DR

`BigDecimal` money, lock both wallets in a global order, debit before credit, idempotency key → at most once, Strategy offers paid from a funded promotions wallet; prove it with a conserved-total concurrent demo.

21.3 Design an Event Calendar (Flipkart)

Problem & clarifying questions

Flipkart statement (kumaransg/LLD `Event_calendar_flipkart`): users with **working hours**; **teams** (a user in at most one); one-day **events** whose participants are users and/or teams; a user can be in **one event at a time**. For teams the creator says how many **representatives (N)** each must send; the system blocks N *available* members, else the event isn't created. Required: create user, team, event; **a user's events in a time range** (including as representative); **suggest free slots** for users/teams on a day. Bonus: **update event time**, **recurring events**. No DB/UI; "do not make assumptions, ask".

Question	Answer used
Does 14:00–15:00 clash with 15:00–16:00?	No: **half-open** `[start, end)`.
Must events fit every attendee's working hours?	Yes.
Which members represent a team?	First N free members in team order (pluggable strategy).
Direct attendee also in an invited team?	Not picked again as representative.
Free for a team with N?	At least N members working and not busy.
Recurring series with one clash?	Reject the whole series.

The sample is inconsistent. It says A + T1 (N = 1) are free "10AM to 3PM and 4PM to 7PM", but T1's members start at 11:00 and 11:30, so by the rules the answer is `11:00-15:00, 16:00-19:00` (what our program prints). Pointing this out politely is exactly what "ask, don't assume" tests. (It also says "Onboard 5 Users" but lists six.)

Requirements

Functional: create user (working hours); create team (one team per user); create event, rejected if anyone is outside hours, busy, or a team lacks N free members; a user's events in a range; free slots; bonus: atomic update time; recurring daily/weekly. **Non-functional:** "is everyone free?" + booking must be **atomic**; representative choice replaceable.

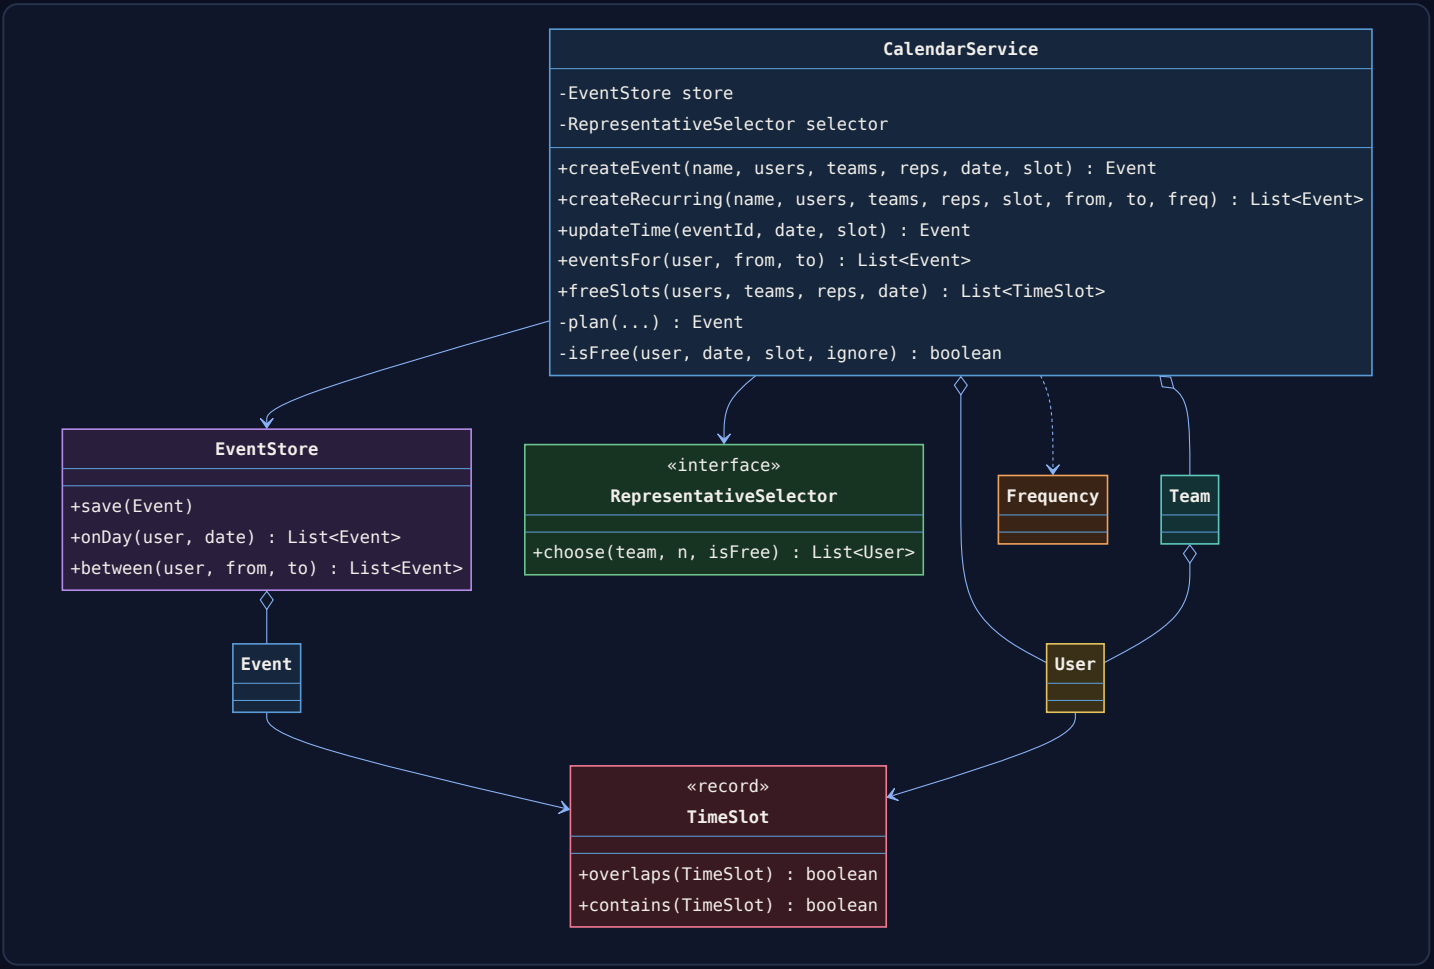
Use cases

Admin (onboard user, create team), **Organiser** (create/recurring/move event, find free slots), **Attendee** (own events in a range).

Core entities

Class	Responsibility
`TimeSlot` (record)	Half-open `[start, end)`; `overlaps`, `contains`; validates start < end.
`User`, `Team` (records)	Name + working hours; name + ordered members.
`Event` (record)	One occurrence: date, slot, request (users, teams, N), actual `attendees`, optional `seriesId`.
`Frequency`	DAILY / WEEKLY, knows the next date.
`EventStore`	Repository + index `user → date → events` (`TreeMap.subMap` range queries).
`RepresentativeSelector`	**Strategy**: pick N free members.
`CalendarService`	Use cases; planning separated from saving.
`CalendarException`	Readable rejection reason.

Class diagram



Key flows

Create event = plan, then save: reject if no participants; each direct user must be inside working hours with no overlapping event; for each team the selector picks N free members not already attending (fewer → reject "team has only k free"); then save and index per attendee.

Free-slot sweep: collect every interesting time that day — each person's working-hour start/end and each of their events' start/end — and sort. Nothing changes between two consecutive cuts, so test each piece once (all direct users free, each team has ≥ N free) and merge neighbouring good pieces. For A + T1 the cuts are 10:00, 11:00, 11:30, 15:00, 16:00, 18:30, 19:00, 19:30 → 11:00–15:00 and 16:00–19:00.

Design decisions & patterns used

Decision	Why	Alternative rejected
Half-open intervals	Back-to-back meetings don't clash; `overlaps` = `a.start < b.end && b.start < a.end`.	Closed intervals.
`plan()` separate from `save()`	Side-effect-free plan reused by create, recurring (plan all, then save all) and update (plan ignoring the old event, then swap) → all-or-nothing for free.	Save as you go + rollback.
`RepresentativeSelector`	Fairness / least-busy pluggable.	Hard-coded loop.
Index `user → TreeMap<date, events>`	Conflict checks read one user's one day; ranges use `subMap`.	Scanning all events.
Boundary sweep	Exact, any N, no minute loops.	Checking every 15/30 min (misses 11:30).
`synchronized` service methods	Check-then-book atomic.	Per-user locks (needs lock ordering; a scale-up).

Java implementation

```java
// File: lld/ch21/calendar/TimeSlot.java
/** Half-open interval [start, end) inside one day. 14:00-15:00 and 15:00-16:00 do not overlap. */
public record TimeSlot(LocalTime start, LocalTime end) {
    // ... compact constructor rejects start >= end
    public boolean overlaps(TimeSlot o) { return start.isBefore(o.end) && o.start.isBefore(end); }
    public boolean contains(TimeSlot o) { return !o.start.isBefore(start) && !o.end.isAfter(end); }
    // ...
}

// File: lld/ch21/calendar/CalendarService.java  (freeSlots: the sweep)
        TreeSet<LocalTime> cuts = new TreeSet<>();
        for (User u : everyone) {
            cuts.add(u.workingHours().start());
            cuts.add(u.workingHours().end());
            for (Event e : store.onDay(u.name(), date)) { cuts.add(e.slot().start()); cuts.add(e.slot().end()); }
        }
        List<TimeSlot> result = new ArrayList<>();
        LocalTime prev = null;
        for (LocalTime cut : cuts) {
            if (prev != null) {
                TimeSlot piece = new TimeSlot(prev, cut);
                boolean ok = direct.stream().allMatch(u -> isFree(u, date, piece, null))
                        && tms.stream().allMatch(t -> t.members().stream()
                                .filter(m -> isFree(m, date, piece, null)).count() >= reps);
                if (ok) {
                    int last = result.size() - 1;
                    if (last >= 0 && result.get(last).end().equals(piece.start()))
                        result.set(last, new TimeSlot(result.get(last).start(), piece.end()));
                    else result.add(piece);
                }
            }
            prev = cut;
        }
        return result;
    // ...
    private boolean isFree(User u, LocalDate date, TimeSlot slot, Event ignore) {
        if (!u.workingHours().contains(slot)) return false;
        return store.onDay(u.name(), date).stream()
                .filter(e -> ignore == null || !e.id().equals(ignore.id()))
                .noneMatch(e -> e.slot().overlaps(slot));
    }
```

Full runnable code: code/lld/ch21/calendar/ (see the Full Edition).

Output (excerpt):

```
  FAILED: Event2: C is not available 2026-09-29 14:00-15:00
  ok: Event3 2026-09-28 15:00-16:00 [C, E, B, D]
  FAILED: Event5: F is not available 2026-09-28 10:00-11:00
== 5. free slots today for A + T1 (1 representative)
  [11:00-15:00, 16:00-19:00]
== bonus: update event time
  FAILED: Event1: team T1 has only 1 free member(s), needs 2
  FAILED: Sync@2026-09-28: A is not available 2026-09-28 15:30-16:30
```

The refused update keeps the old booking; the "Sync" series is rejected as a whole.

Concurrency & edge cases

- All service methods `synchronized` → "is free?" + save is atomic.
- Handled: start ≥ end, unknown user/team, user already in a team, no participants, N < 1, any clash in a series, moving an event onto a slot overlapping *itself* (ignored).
- Out of scope: time zones (real calendars store UTC instants + the user's zone), multi-day events (split into per-day slots).

Extensions & follow-up questions

- **Invites:** `InviteStatus {PENDING, ACCEPTED, DECLINED}`; declined attendees don't block.

- **Fair representatives:** `LeastBusySelector`.
- **Recurring at scale:** store one rule (iCalendar RRULE) and expand lazily per query.
- **Rooms:** a room is a "user" with building hours. **Big teams:** count arrays or a segment tree over minutes.

Real-world example

Google Calendar's "Find a time" and Outlook's Scheduling Assistant intersect attendees' busy blocks the same way, and store recurring meetings as an RRULE plus exceptions.

In the interview

Draw a timeline: free time = working time minus busy time, intersected. One sentence on half-open intervals. Code `TimeSlot`, `plan()`, `createEvent` first (first three requirements), then `eventsFor`, then the sweep. Mention the sample inconsistency. Point out that the bonus features fall out of `plan()`.

Self-check

- **Q:** Why does `plan()` take an `ignore` event? → **A:** When moving an event, its own booking mustn't conflict with the new time.
- **Q:** Why is all-or-nothing recurrence easy? → **A:** Plan every occurrence first, save only if all succeed.
- **Q:** Half-open overlap test? → **A:** `a.start < b.end && b.start < a.end`.
- **Q:** Why a `TreeMap` per user? → **A:** Range query = `subMap(from, to)`.

TL;DR

Half-open `TimeSlot`s, per-user day index, side-effect-free `plan()` reused by create/recurring/update, pluggable representative selector, boundary sweep for free slots — and ask about the contradictory sample.

21.4 Design a LeetCode-like Coding Platform — "Coding Blox" (Flipkart)

Problem & clarifying questions

Flipkart SDE-2 statement (kumaransg/LLD `leetcode-lld-flipkart-coding-blox`). Users sign up, create contests and attend others'. Questions and contests have a level (LOW/MEDIUM/HIGH); a contest uses **all questions of its level**. Command-driven:

Command	Meaning
`CreateUser <name>`	Default score **1500**.
`CreateQuestion <level> <score>` / `ListQuestion [level]`	Ids auto-increment from 1.
`CreateContest <name> <level> <creator>` / `ListContest [level]`	Creator auto-attends.
`AttendContest <contestId> <user>`	Register.
`RunContest <contestId> <creator>`	Creator only. Simulate solved questions (random), sum `contestPoints`, then: LOW → `score + (points − 50)`, MEDIUM → `score + (points − 30)`, HIGH → `score + points`.
`LeaderBoard <asc/desc>`	Users sorted by score.
Bonus `ContestHistory` / `WithdrawContest`	Per-contest results; leave before the run (not the creator).

Expectations: your own sample data, driver or tests, modular OO with interfaces, graceful edge cases, DRY.

Question	Answer used
How many questions solved?	Random 1..all, behind a **seeded** `SolveSimulator`.
Run twice / attend after run?	No (errors) — ratings never applied twice.
Tie-break / empty level?	By name / can't create a contest with no questions of its level.

Requirements

Functional: all commands above including both bonus ones. **Non-functional:** commands from a built-in list; errors printed, processing continues; scoring rule and solving logic replaceable; rating updates safe if two contests finish at once.

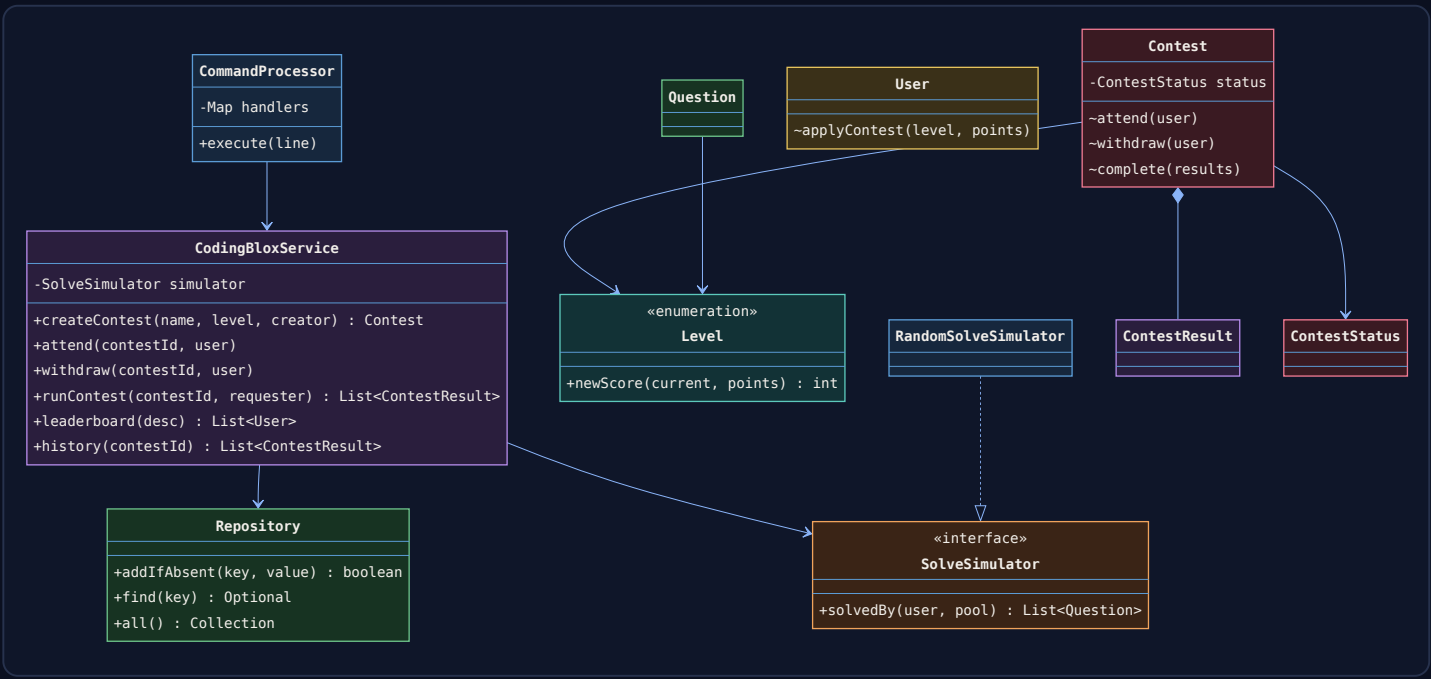
Use cases

User (sign up, attend/withdraw, leaderboard, history), **Contest creator** (create, run), **Admin / problem setter** (questions), **Judge** (simulated: who solved what).

Core entities

Class	Responsibility
`Level` (enum)	LOW/MEDIUM/HIGH with penalty; `newScore(current, points)` is the rating rule.
`User`	Name + score; `applyContest` updates atomically.
`Question`, `ContestResult` (records)	id/level/score; user/points/solved ids.
`Contest` + `ContestStatus`	Participants, OPEN/COMPLETED, results; guards "open only" rules.
`SolveSimulator` / `RandomSolveSimulator`	**Strategy** for "who solved what".
`Repository<K,V>`	Generic in-memory store reused three times.
`CodingBloxService`	Business rules per command.
`CommandProcessor`	Parses a line, dispatches via a map, prints.
`BloxException`	Readable failure.

Class diagram



Key flows

RunContest: requester must be the creator → pool = questions of the contest level → for each participant, `simulator.solvedBy` and sum points → `contest.complete(results)` (throws if already COMPLETED) → only then apply ratings. Contest lifecycle: **OPEN** (attend/withdraw allowed) → **COMPLETED** on `RunContest` by the creator.

Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
Command name → handler map	Command	New command = one `put`; parsing/printing out of the service.	Parsing in the service / giant `switch`.
Scoring rule on `Level`	Strategy (per-constant data)	A new level brings its own rule.	`if (level == LOW)…` in `runContest`.
`SolveSimulator` interface	Strategy	Seeded for demo, scripted for tests, real judge later.	`new Random()` in the service.
`Contest.complete` checks status	Entity guards invariant	"Run once" enforced where state lives.	Check in the driver.
Generic `Repository<K,V>`	Repository	Three stores, one class (DRY).	Three copy-pasted maps.
Results → `complete` → ratings	Safe ordering	If `complete` throws, no rating changed.	Updating ratings in the loop.

Java implementation

```
// File: lld/ch21/codingblox/Level.java
public enum Level {
    LOW(50), MEDIUM(30), HIGH(0);

    private final int penalty;
    Level(int penalty) { this.penalty = penalty; }

    /** newScore = currentScore + (contestPoints - penalty), exactly as the statement defines. */
    public int newScore(int current, int contestPoints) { return current + contestPoints - penalty; }
    // ... parse(String)
}

// File: lld/ch21/codingblox/CodingBloxService.java
    public List<ContestResult> runContest(int contestId, String requester) {
        Contest c = contest(contestId);
        if (!c.creator().equals(requester)) throw new BloxException(requester + " did not create contest " +
contestId);
        List<Question> pool = listQuestions(c.level());
        List<ContestResult> results = new ArrayList<>();
        for (String name : c.participants()) {
            List<Question> solved = simulator.solvedBy(name, pool);
            int points = solved.stream().mapToInt(Question::score).sum();
            results.add(new ContestResult(name, points, solved.stream().map(Question::id).toList()));
        }
        c.complete(results);                              // throws if already run: scores never double-apply
        for (ContestResult r : results) user(r.user()).applyContest(c.level(), r.points());
        return results;
    }

// File: lld/ch21/codingblox/CommandProcessor.java
        handlers.put("runcontest", a -> svc.runContest(Integer.parseInt(arg(a, 1)), arg(a, 2))
                .forEach(r -> out("  " + r.user() + " solved " + r.solvedIds() + " -> " + r.points() + " pts")));
        // ... one entry per command
```

Full runnable code: code/lld/ch21/codingblox/ (see the Full Edition).

Output (excerpt):

```
> RunContest 1 Ross
  Ross solved [2] -> 20 pts
  Monica solved [1, 2] -> 30 pts
  Joey solved [1, 2] -> 30 pts
> RunContest 1 Ross
ERROR: Contest 1 has already run
> LeaderBoard score desc
  Monica : 1540
  Chandler : 1510
  Joey : 1480
  Ross : 1470
```

Contest 1 is LOW: Ross 1500 + (20 − 50) = 1470; Joey/Monica 1480. Contest 2 is MEDIUM: Monica 1480 + (90 − 30) = 1540, Chandler 1500 + (40 − 30) = 1510.

Concurrency & edge cases

- `ConcurrentHashMap` repositories; `synchronized` `Contest` methods; `User.applyContest` under the user's lock, so two contests finishing together can't lose an update.
- **Known gap (say it):** a user attending between the participant copy and `complete` gets no result → add a `RUNNING` status switched atomically with handing back the participant list.
- Handled: duplicates, unknown ids/levels, attend/withdraw after the run, creator withdraw, non-creator run, bad args, unknown command.

Extensions & follow-up questions

- Fixed number of questions: a `QuestionSelectionStrategy` at creation. Real submissions: a judge + `Submission(user, question, verdict, time)`, ICPC-style scoring.
- Elo-style ratings: a `RatingStrategy` over all participants' results. Scheduled contests: `SCHEDULED → RUNNING → COMPLETED` + scheduler.
- Leaderboard at scale: `TreeSet` by (score, name) or a Redis sorted set updated on each change.

Real-world example

LeetCode weekly contests recompute each participant's rating from their relative rank after the contest; HackerRank and Codeforces share the shape: a problem pool, registration, ratings only after finalisation.

In the interview

Paste the sample commands first; build CreateUser → CreateQuestion → CreateContest → AttendContest → RunContest → LeaderBoard, running after each. Name the three moves: command map, scoring on the enum, seeded randomness behind an interface. Bonus commands are tiny because `Contest` stores results and guards status.

Self-check

- **Q:** Why update ratings only after `complete`? → **A:** A second run throws before any rating changes.
- **Q:** Add `VERY_HIGH` with +20 bonus? → **A:** An enum constant with its own rule (e.g. negative penalty).
- **Q:** Deterministic `RunContest` in tests? → **A:** Inject a scripted or seeded `SolveSimulator`.
- **Q:** Where is "creator always attends" enforced? → **A:** The `Contest` constructor.

TL;DR

Command map in the driver, rules in the service, invariants in `Contest`, scoring on `Level`, randomness behind a seeded `SolveSimulator`; complete first, then update ratings.

21.5 Design a Property Listing Site — "PropertyHunt" (ClearTrip)

Problem & clarifying questions

ClearTrip statement (kumaransg/LLD `PropertyHunt`; the longer screenshot version is followed). Users **register**, **list** properties to sell or rent, **view own listings**, **shortlist** others' properties, **view shortlisted** (sold ones show as SOLD), **search** available properties by any combination of **location** (case-

insensitive), **price range**, **type** (sell/rent), **size range**, **rooms** (1/2/3BHK), **sort by** price or size; **mark sold** (owner only); **view sold**. Ranges `X-Y` are inclusive; a single `X` means "at least X" for size and "at most X" for price. **Bonus:** sq ft / sq m / sq yd; **nearby** locations shown last; prices like `40k`, `60L`, `2Cr`; comma-separated values for location and rooms.

Question	Answer used
Input format?	One command per line, fields split by `
Shortlist my own property?	No (the text readme's sample does it — a conflict to ask about).
Shortlist a sold property?	No; but one sold *after* shortlisting stays, shown as SOLD.
Internal units?	Price in whole rupees (`long`), size in sq ft (`double`).
Ordering with a location?	Exact matches, then nearby, then the chosen sort.
Mark sold records buyer/price?	Status only; a `Sale` record is an extension.

The sample `bellandur` search returns a *sarjapur* property — consistent only because sarjapur declares `nearby bellandur|…` (the text readme drops that field).

Requirements

Functional: RegisterUser, ListProperty, Search, Shortlist, ViewShortlisted, ViewListed, MarkSold, ViewSold + all four bonus items. **Non-functional:** filters combine freely; new filters/units without rewriting search; clear errors; concurrent-safe.

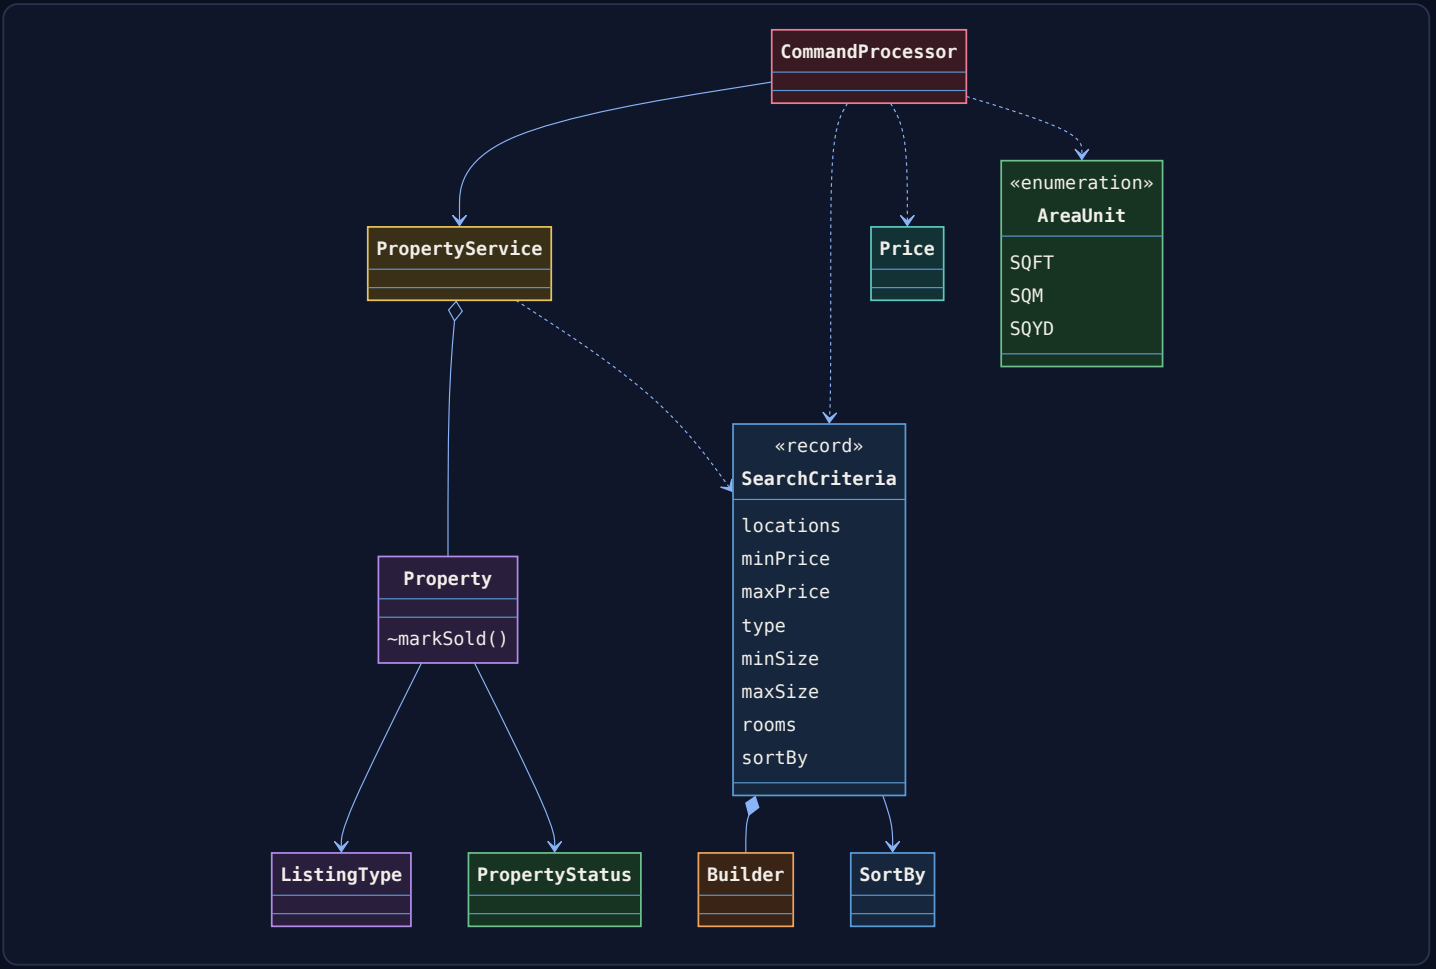
Use cases

Seller / landlord (list, view listed, mark sold, view sold), **Buyer / tenant** (search, shortlist, view shortlisted). Roles are per property, not per account.

Core entities

Class	Responsibility
`Property`	Listing data, owner, nearby locations, status; `markSold` guards double sale.
`ListingType` , `PropertyStatus` , `SortBy`	SELL/RENT, AVAILABLE/SOLD, sort comparators (**Strategy**).
`Price`	Parse `40k/60L/2Cr/1.5Cr` ↔ rupees.
`AreaUnit`	SQFT/SQM/SQYD → sq ft; parses "70 sqm".
`SearchCriteria` + `Builder`	All-optional search parameters (**Builder**).
`PropertyService`	Use cases; owner-vs-buyer rules.
`CommandProcessor` , `HuntException`	Parses `

Class diagram



Key flows

Search: parse fields → normalise with `Price.parse` / `AreaUnit.parseSize` → `SearchCriteria.Builder` → service builds a **list of predicates** (AVAILABLE, location or nearby, price, type, size, rooms) → reduce with `Predicate.and` → sort by (location rank, price or size, id) → driver prints the table. Property lifecycle: AVAILABLE → SOLD (owner only).

Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
All-optional criteria with a Builder	Builder	Eight optional fields; eight nulls are unreadable.	Telescoping constructors.
Each criterion → `Predicate`, reduced with `and`	Specification	Any combination; new filter = one line.	`if` per combination.
`SortBy` enum of comparators	Strategy	"Sort by rooms" = one constant.	`switch` in search.
Normalise at the boundary (rupees, sq ft)	—	Business code compares plain numbers.	Comparing strings like "90L".
Location rank 0/1/2	—	One function filters *and* orders (exact, nearby, excluded).	Two searches merged.
Shortlist = set of ids per user	—	Status read live, so SOLD appears automatically.	Copied objects go stale.

Java implementation

```java
// File: lld/ch21/propertyhunt/PropertyService.java
    public List<Property> search(String user, SearchCriteria c) {
        requireUser(user);
        Set<String> locs = c.locations() == null ? null : lower(c.locations());
        List<Predicate<Property>> filters = new ArrayList<>();
        filters.add(p -> p.status() == PropertyStatus.AVAILABLE);
        if (locs != null) filters.add(p -> locationRank(p, locs) < 2);
        if (c.minPrice() != null) filters.add(p -> p.price() >= c.minPrice());
        if (c.maxPrice() != null) filters.add(p -> p.price() <= c.maxPrice());
        if (c.type() != null) filters.add(p -> p.type() == c.type());
        if (c.minSize() != null) filters.add(p -> p.sizeSqft() >= c.minSize());
        if (c.maxSize() != null) filters.add(p -> p.sizeSqft() <= c.maxSize());
        if (c.rooms() != null) filters.add(p -> c.rooms().contains(p.rooms()));
        Predicate<Property> all = filters.stream().reduce(x -> true, Predicate::and);

        Comparator<Property> order = locs == null ? c.sortBy().comparator
                : Comparator.<Property>comparingInt(p -> locationRank(p, locs)).thenComparing(c.sortBy().comparator);
        return properties.values().stream().filter(all).sorted(order).toList();
    }
    // ... shortlist(): rejects own, sold and duplicate properties; stores only the id
```

Full runnable code: code/lld/ch21/propertyhunt/ (see the Full Edition).

Output (excerpt):

```
-> Search|Jack|bellandur|null|null|null|null|price
<- 1   A 3BHK for sale      bellandur  90L    1800sqft  3BHK  Sale
<- 2   A 2BHK for rent      sarjapur   40k    1500sqft  2BHK  Rent
-> ViewShortlisted|Jill
<- 1   A 3BHK for sale      bellandur  90L    1800sqft  3BHK  Sale  SOLD
<- 2   A 2BHK for rent      sarjapur   40k    1500sqft  2BHK  Rent  AVAILABLE
-> ListProperty|Jill|Bad|x|abc|sell|100 sqft|1
ERROR: Bad price: abc
```

Nearby sarjapur comes after exact bellandur; a sold property leaves searches but stays (SOLD) in shortlists.

Concurrency & edge cases

- `ConcurrentHashMap`s, concurrent shortlist sets, `synchronized` `markSold` (no double sale).
- A search racing a sale may still show it — fine for listings; a purchase flow re-checks under a lock.
- Handled: unregistered user, unknown property, own/duplicate/sold shortlist, non-owner mark sold, bad price/size/unit/type.

Extensions & follow-up questions

- Descending / multi-key sort: direction flag + `.reversed()`, or chained `thenComparing`.
- Fuzzy locations: a `LocationResolver`. "Within 5 km": lat/long + haversine; at scale a geo-index (geohash, R-tree) or a search engine.
- Millions of listings: secondary indexes or Elasticsearch/OpenSearch. Price-drop alerts: Observer on `Property`.

Real-world example

99acres, MagicBricks and NoBroker filter by locality, budget, BHK, buy/rent and area, suggest nearby localities, and quote prices in lakhs and crores — hence parsing `60L` and `2Cr`.

In the interview

Start with model and units (`Price.parse("90L")`, `AreaUnit`), then register → list → search with *one* filter → add predicates → sort. Call out the Builder + predicate list as the extensibility story and "normalise at the boundary".

Self-check

- **Q:** Buyer searches in sq m, seller listed in sq yd? → **A:** Both converted to sq ft at parse time.

- **Q:** Why shortlist ids, not objects? → **A:** Status is read live.
- **Q:** A single number in the price field means? → **A:** "At most X" (size: "at least X").
- **Q:** Add a "furnished" filter? → **A:** A `Property` field, an optional criterion, one predicate line.

TL;DR

Normalise price and size at the edges, turn each optional criterion into a predicate and AND them, sort by (location rank, key, id), keep shortlists as ids.

21.6 Design a Loan Ledger — "Ledger Co" (Navi)

Problem & clarifying questions

The Geektrust "Ledger Co" challenge (kumaransg/LLD `ledger_company_navi`, used by Navi). Banks lend; borrowers repay in **EMIs** (Equated Monthly Instalments, same amount each month) and may pay a **lump sum** any time.

Command	Meaning
`LOAN BANK BORROWER P N R`	Principal `P`, `N` years, `R` % simple interest per year.
`PAYMENT BANK BORROWER LUMP_SUM EMI_NO`	Lump sum paid right after EMI number `EMI_NO`.
`BALANCE BANK BORROWER EMI_NO`	Print `BANK BORROWER AMOUNT_PAID EMIS_LEFT` after EMI `EMI_NO`.

Rules: `I = P × N × R / 100`; `A = P + I`; EMI = `A / (N × 12)` **rounded up** (86.67 → 87, 100.10 → 101); paid after EMI k = `k × EMI` + every lump sum with EMI number ≤ k; EMIs left = remaining ÷ EMI **rounded up**; the last EMI shrinks to what is owed. A lump sum reduces the *number* of EMIs, never the amount. Only `BALANCE` prints.

Question	Answer used
Paid can exceed `A`?	No — capped at `A` (EMIs left = 0).
Whole-number interest?	Assumed rounded up (say so); official samples are whole either way.
One loan per bank-borrower?	Yes; a second `LOAN` is an error.
`EMI_NO` beyond tenure?	Error; `0` is valid.

Requirements

Functional: LOAN, PAYMENT, BALANCE as above; reproduce the official samples. **Non-functional:** input from a list/file; a bad line doesn't stop the rest; interest rule replaceable; exact `BigDecimal` arithmetic.

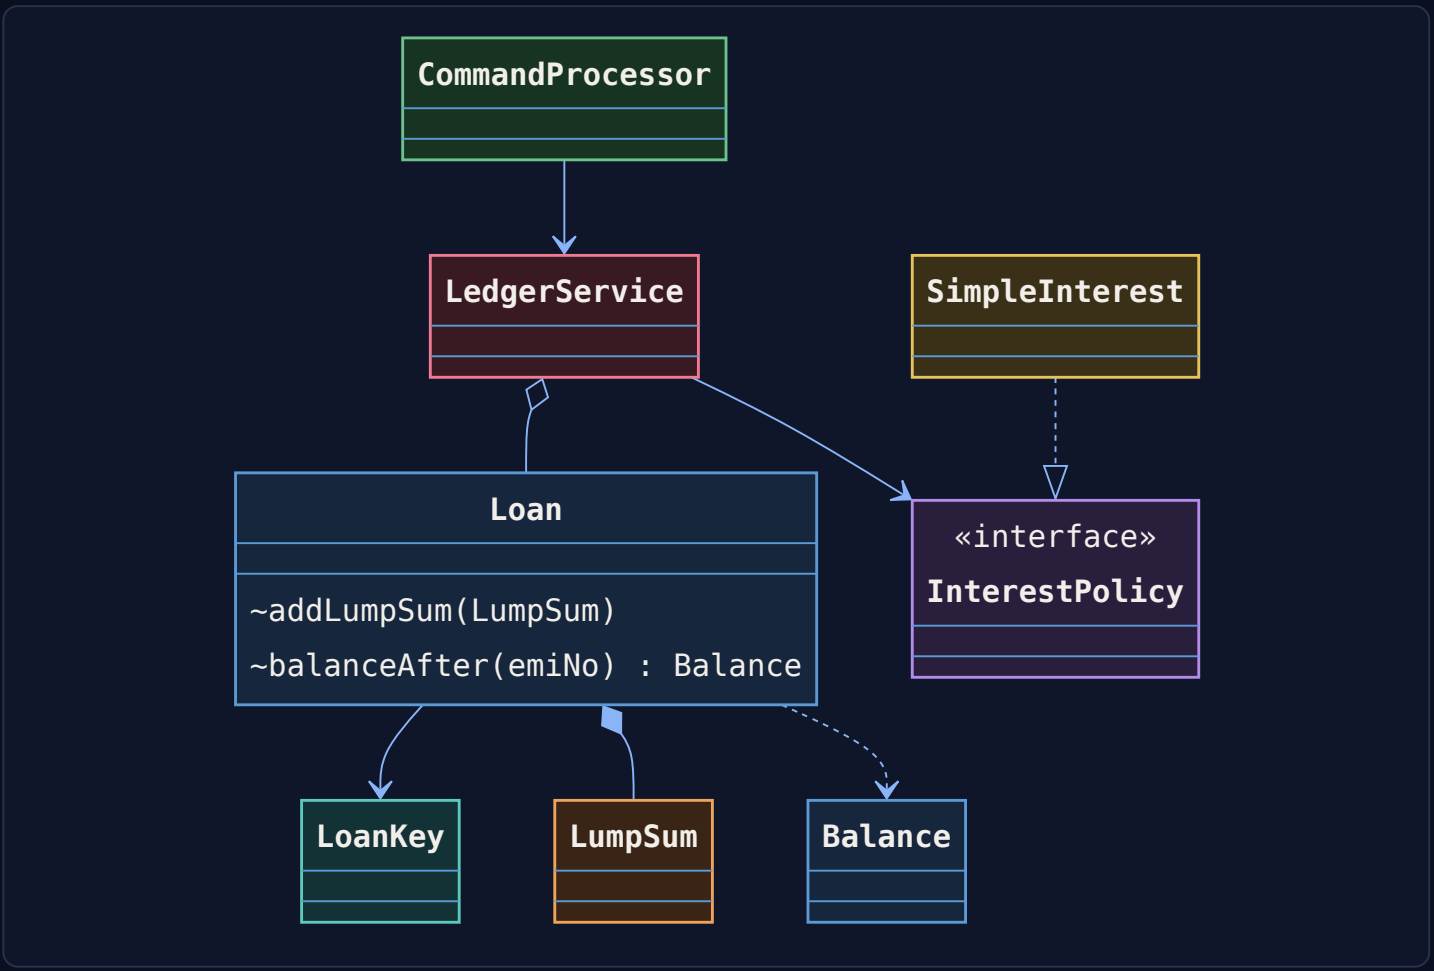
Use cases

Bank (records a loan), **Borrower** (pays lump sums), **Anyone** (balance after EMI k).

Core entities

Class	Responsibility
`LoanKey` (record)	(bank, borrower) — a perfect map key (records give `equals/hashCode`).
`Loan`	Computes `A` and EMI once; stores lump sums; derives a `Balance` for any EMI.
`LumpSum`, `Balance` (records)	amount + after-EMI; output line.
`InterestPolicy` / `SimpleInterest`	**Strategy** for total interest.
`LedgerService`, `CommandProcessor`, `LedgerException`	Loans store; parsing and validation; readable errors.

Class diagram



Key flows

Worked example (Sample 2, Shelly): P = 15000, N = 2, R = 9 → I = 2700, A = 17700, EMI = ⌈737.5⌉ = 738. After 12 EMIs + a 7000 lump sum: paid 15856, remaining 1844, EMIs left ⌈2.499⌉ = 3 → `UON Shelly 15856 3`.

Design decisions & patterns used

Decision	Why	Alternative rejected
Derive balance from facts (A, EMI, lump sums) per query	Queries may ask about *any* EMI, even before a later payment.	A mutable `remaining` field — wrong for earlier EMIs.
`InterestPolicy` Strategy	Reducing-balance or compound later without touching `Loan`.	Formula hard-coded.
`LoanKey` record as key	Value semantics for free.	`"IDIDI#Dale"` strings.
`BigDecimal` + `RoundingMode.CEILING`	Exact "always up".	`Math.ceil` on doubles (4.000000001 → 5).
Driver returns `List<String>`	Easy to compare with expected output.	Printing in the service.

Java implementation

```java
// File: lld/ch21/ledger/Loan.java
    Loan(LoanKey key, BigDecimal principal, int years, BigDecimal rate, InterestPolicy policy) {
        if (principal.signum() <= 0 || years <= 0 || rate.signum() < 0)
            throw new LedgerException("Invalid loan terms for " + key);
        this.key = key;
        this.tenureMonths = years * 12;
        this.totalRepayable = principal.add(policy.interest(principal, BigDecimal.valueOf(years), rate));
        this.emi = totalRepayable.divide(BigDecimal.valueOf(tenureMonths), 0, RoundingMode.CEILING);
    }
    // ... addLumpSum validates amount > 0 and EMI number in range

    /** Amount paid after `emiNo` EMIs (lump sums made at or before that EMI count), capped at A. */
    synchronized Balance balanceAfter(int emiNo) {
        if (emiNo < 0 || emiNo > tenureMonths) throw new LedgerException("EMI number out of range: " + emiNo);
        BigDecimal lumps = lumpSums.stream().filter(l -> l.afterEmi() <= emiNo)
                .map(LumpSum::amount).reduce(BigDecimal.ZERO, BigDecimal::add);
        BigDecimal paid = emi.multiply(BigDecimal.valueOf(emiNo)).add(lumps).min(totalRepayable);
        BigDecimal remaining = totalRepayable.subtract(paid);
        int left = remaining.divide(emi, 0, RoundingMode.CEILING).intValueExact();
        return new Balance(key.bank(), key.borrower(), paid, left);
    }
```

Full runnable code: code/lld/ch21/ledger/ (see the Full Edition).

Output (excerpt) — the Sample 1 and 2 blocks match the official Geektrust expected outputs:

```
== Sample 2
IDIDI Dale 1326 9
IDIDI Dale 3652 4
UON Shelly 15856 3
MBI Harry 9044 10
== Edge cases
SBI Asha 1120 0
ERROR EMI number out of range: 13  [BALANCE SBI Asha 13]
ERROR Not a number: ten  [PAYMENT SBI Asha ten 1]
```

Asha's 2000 lump sum overpays a 1120 loan, so paid is capped at 1120 and 0 EMIs remain.

Concurrency & edge cases

- `ConcurrentHashMap` + `putIfAbsent` makes "one loan per pair" atomic; `synchronized` `Loan` methods keep PAYMENT and BALANCE apart.
- Handled: duplicate/unknown loan, EMI out of range, invalid terms, overpayment (capped), bad arguments, unknown command.

Extensions & follow-up questions

- **Reducing-balance EMI:** new policy, `EMI = P·r·(1+r)^n / ((1+r)^n − 1)` ; `Loan` API unchanged.
- **Lump sum that reduces the EMI instead:** a `PrepaymentPolicy` (REDUCE_TENURE / REDUCE_EMI).
- **Late payments:** record dated payment events and derive EMIs paid. **File input:** `Files.readAllLines` into the same `run`.

Real-world example

Navi, banks and NBFCs show an EMI schedule and allow part-prepayment, usually choosing between fewer EMIs (this puzzle's rule) or a smaller EMI. Real loans use reducing-balance interest, not flat simple interest — worth raising.

In the interview

Show the three formulas and compute one sample by hand (proves you understood the rounding); code `Loan.balanceAfter` (the heart), the service, the command loop; run both samples. Emphasise "state = facts, balances = derived".

Self-check

- **Q:** Why ceil the EMI? → **A:** A floor leaves a shortfall; rounding up guarantees the EMIs cover `A`.

- **Q:** Does `BALANCE … 5` include a PAYMENT after EMI 10? → **A:** No; only `afterEmi ≤ 5`.
- **Q:** Why compute the balance rather than store it? → **A:** Queries arrive for any EMI in any order.
- **Q:** Compound interest? → **A:** A new `InterestPolicy` injected into `LedgerService`.

TL;DR

EMI = ⌈(P + PNR/100) / 12N⌉. Store loan facts and lump sums; derive paid and EMIs-left on demand with ceiling rounding; interest is a Strategy; the driver is a tiny command loop.

21.7 Design an Inventory Management System

Problem & clarifying questions

A common e-commerce / quick-commerce question: track **SKUs** (stock keeping units — one id per sellable item) across **warehouses**, record every **stock movement** (received, shipped, transferred, adjusted after a physical count), fulfil orders by choosing warehouses, and raise **reorder alerts**.

Question	Answer used
Stock per warehouse?	Yes: (SKU, warehouse) → quantity.
Split orders?	Yes; warehouse order is a strategy (nearest city first / most stock first).
Partial shipments?	No — all-or-nothing.
Reorder alert when?	*Total* stock ≤ reorder point; once per dip, re-armed on recovery.
Reorder how much?	`max − current`.
Negative stock?	Never.

Requirements

Functional: add products (reorder point, max level) and warehouses; receive; ship by strategy; transfer; cycle-count adjust; stock per warehouse and total; SKU movement history; reorder alerts. **Non-functional:** **no overselling** under concurrency; different products in parallel; full audit trail.

Use cases

Warehouse staff (receive, transfer, adjust), **Order service** (ship), **Purchasing** (reorder alerts), **Manager** (stock and history).

Core entities

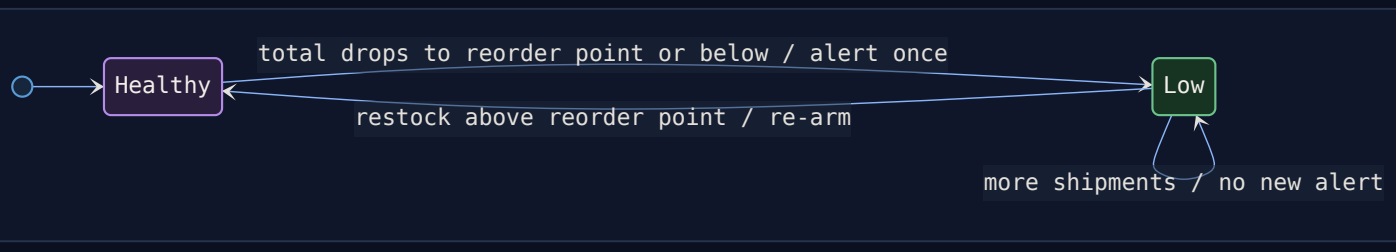
Class	Responsibility
`Product`, `Warehouse` (records)	SKU, name, reorder point, max level; id, city.
`StockMovement` + `MovementType`	Append-only ledger line: INBOUND, OUTBOUND, TRANSFER, ADJUSTMENT.
`Allocation` (record)	"take q from warehouse w".
`AllocationStrategy` + `NearestCityFirst`, `MostStockFirst`	**Strategy**: order to drain warehouses.
`ReorderListener`	**Observer** for low stock.
`InventoryService`, `InventoryException`	Levels, per-SKU locks, movements, alerts; readable errors.

Class diagram



Key flows

Ship: lock the SKU → total < qty? reject, nothing changed → strategy orders a *snapshot* of the levels → build the full plan → apply it, writing OUTBOUND movements → `checkReorder` → unlock.



Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
One lock per SKU	Lock striping by key	Covers all of a SKU's warehouses with no lock-ordering puzzle; different SKUs never contend.	Lock per (SKU, warehouse) — needs ordering; one global lock.
Plan, then apply	—	Check total and build the full plan before changing anything → all-or-nothing.	Deduct and roll back.
`AllocationStrategy`	Strategy	Speed vs fewer split parcels as small classes.	`if (mode == …)` .
Append-only movement ledger	Event log / audit	Every change explainable; levels rebuildable; adjustments record the delta.	Quantities with no history.
Once-per-dip `ReorderListener`	Observer + edge-triggering	Notified once on crossing, not every sale.	Alert on every movement (spam).
Strategy gets a snapshot	Defensive copy	Can't modify live stock.	Passing the live map.

Java implementation

```
// File: lld/ch21/inventory/InventoryService.java
    /** All-or-nothing: either the whole quantity is allocated across warehouses, or nothing moves. */
    public List<Allocation> ship(String sku, int qty, AllocationStrategy strategy, String orderRef) {
        positive(qty);
        return withSku(sku, () -> {
            Map<String, Integer> levels = stock.get(sku);
            int total = levels.values().stream().mapToInt(Integer::intValue).sum();
            if (total < qty) throw new InventoryException("Only " + total + " of " + sku + " available, order needs "
+ qty);
            List<Allocation> plan = new ArrayList<>();
            int remaining = qty;
            for (String wh : strategy.order(Map.copyOf(levels), warehouses)) {
                int take = Math.min(remaining, levels.get(wh));
                if (take > 0) { plan.add(new Allocation(wh, take)); remaining -= take; }
                if (remaining == 0) break;
            }
            for (Allocation a : plan) {
                change(sku, a.warehouseId(), -a.quantity());
                record(MovementType.OUTBOUND, sku, a.warehouseId(), null, -a.quantity(), orderRef);
            }
            return plan;
        });
    }
    // ...
    /** Fires once when stock crosses down to the reorder point; re-arms when stock recovers. */
    private void checkReorder(String sku) {
        Product p = products.get(sku);
        int total = totalUnlocked(sku);
        if (total <= p.reorderPoint()) {
            if (alerted.add(sku)) listeners.forEach(l -> l.lowStock(p, total, p.maxLevel() - total));
        } else {
            alerted.remove(sku);
        }
    }
    // ... withSku(sku, action): look up the SKU's ReentrantLock, lock, run, unlock
```

Full runnable code: code/lld/ch21/inventory/ (see the Full Edition).

Output (excerpt):

```
== ship 32 phones to a Bengaluru customer (nearest city first)
  [Allocation[warehouseId=BLR-1, quantity=30], Allocation[warehouseId=BLR-2, quantity=2]]
== ship 25 more (most stock first) -> crosses reorder point
  [ALERT] PHONE low: 18 left, reorder 82 to reach 100
  ship 1 more (no duplicate alert): [Allocation[warehouseId=BLR-1, quantity=1]]
== concurrency: 200 one-unit orders race for 50 cases
  shipped=50 rejected=150 left=0
```

The alert fires once at 18 (≤ 20), not again at 17; exactly 50 of 200 concurrent orders ship — no overselling.

Concurrency & edge cases

- Check → plan → apply under the SKU's `ReentrantLock`; different SKUs in parallel.
- **Ledger** = `ConcurrentLinkedQueue` (lock-free appends, order kept, no `ConcurrentModificationException`); a `CopyOnWriteArrayList` would copy the whole log on every movement. **Listeners** = `CopyOnWriteArrayList` (read often, changed rarely); they run under the SKU lock — hand slow ones (email) to an executor.
- Handled: unknown SKU/warehouse, bad quantities, same-warehouse transfer, orders larger than stock, zero-change adjustments (not logged).

Extensions & follow-up questions

- **Reservations (cart holds):** `onHand` and `reserved`; available = onHand − reserved; holds expire; shipping converts a reservation.
- **Batches / expiry:** FEFO (first-expiry-first-out) is just another `AllocationStrategy`.
- **Auto purchase orders:** a listener creating a pending PO for `max − current`.

- **Database / many services:** optimistic updates (`UPDATE ... SET qty = qty - ? WHERE qty >= ? AND version = ?`); ledger as an event stream (Kafka).

Real-world example

Amazon and Flipkart choose the fulfilment centre by proximity and stock and may split an order (`AllocationStrategy`); Shopify has multi-location fulfilment priorities; quick-commerce dark stores (Blinkit, Zepto, Instamart) depend on accurate per-store stock and fast reorder alerts.

In the interview

Start with the data shape: `stock[sku][warehouse] = qty` + an append-only movement list. Code receive and ship (plan-then-apply, per-SKU lock) first — most of the marks; then transfer/adjust, the once-per-dip observer, allocation strategies; run the 200-thread demo.

Self-check

- **Q:** Why check the total before planning? → **A:** Reject short orders before touching any warehouse.
- **Q:** Why is a per-SKU lock enough for multi-warehouse orders? → **A:** An order has one SKU; its lock covers all its warehouses.
- **Q:** Add FEFO? → **A:** Batch-level stock + a new `AllocationStrategy` ; `ship` unchanged.
- **Q:** Why log adjustments as movements? → **A:** Shrinkage is auditable and levels rebuildable.

TL;DR

Stock per (SKU, warehouse), every change a ledger movement, a lock per SKU prevents overselling, allocation is a Strategy, low-stock alerts are an edge-triggered Observer.

21.8 Design a Unix File Search (`find`)

Problem & clarifying questions

Build a small `find` (e.g. `find /var/log -type f -size +10M -mtime +7`): a file-system model, **filters**, a way to **combine** them with AND / OR / NOT, and a clean API for search options. The trap to avoid: one method with ten optional parameters. Terms: a **predicate** is a yes/no test on one entry; a **glob** is a wildcard name pattern (`*` , `?` , `[0-9]`); **depth** = levels below the start (start = 0).

Question	Answer used
Real disk?	In-memory tree (deterministic); `java.nio` adapter is an extension.
Tests?	name / iname glob, type, size, mtime, owner, empty.
Combining?	AND, OR, NOT, parentheses; precedence NOT > AND > OR; side-by-side tests = AND.
Size semantics?	GNU: round **up** to the unit, then compare — `-size -1M` matches only empty files. Files only.
Options?	min/max depth, limit, **prune** (skip e.g. `node_modules`).
API?	Fluent **Builder** + a `find` command-line parser.
Order / huge trees?	Pre-order DFS, children by name; stream results; no recursion; parallel variant.

Requirements

Functional: model files/dirs (name, owner, mtime, size); search with any AND/OR/NOT combination; depth, limit, prune; parse `find` lines; clear errors. **Non-functional:** Open/Closed (new test = one class); streaming, deep trees; searching while files are added must not throw.

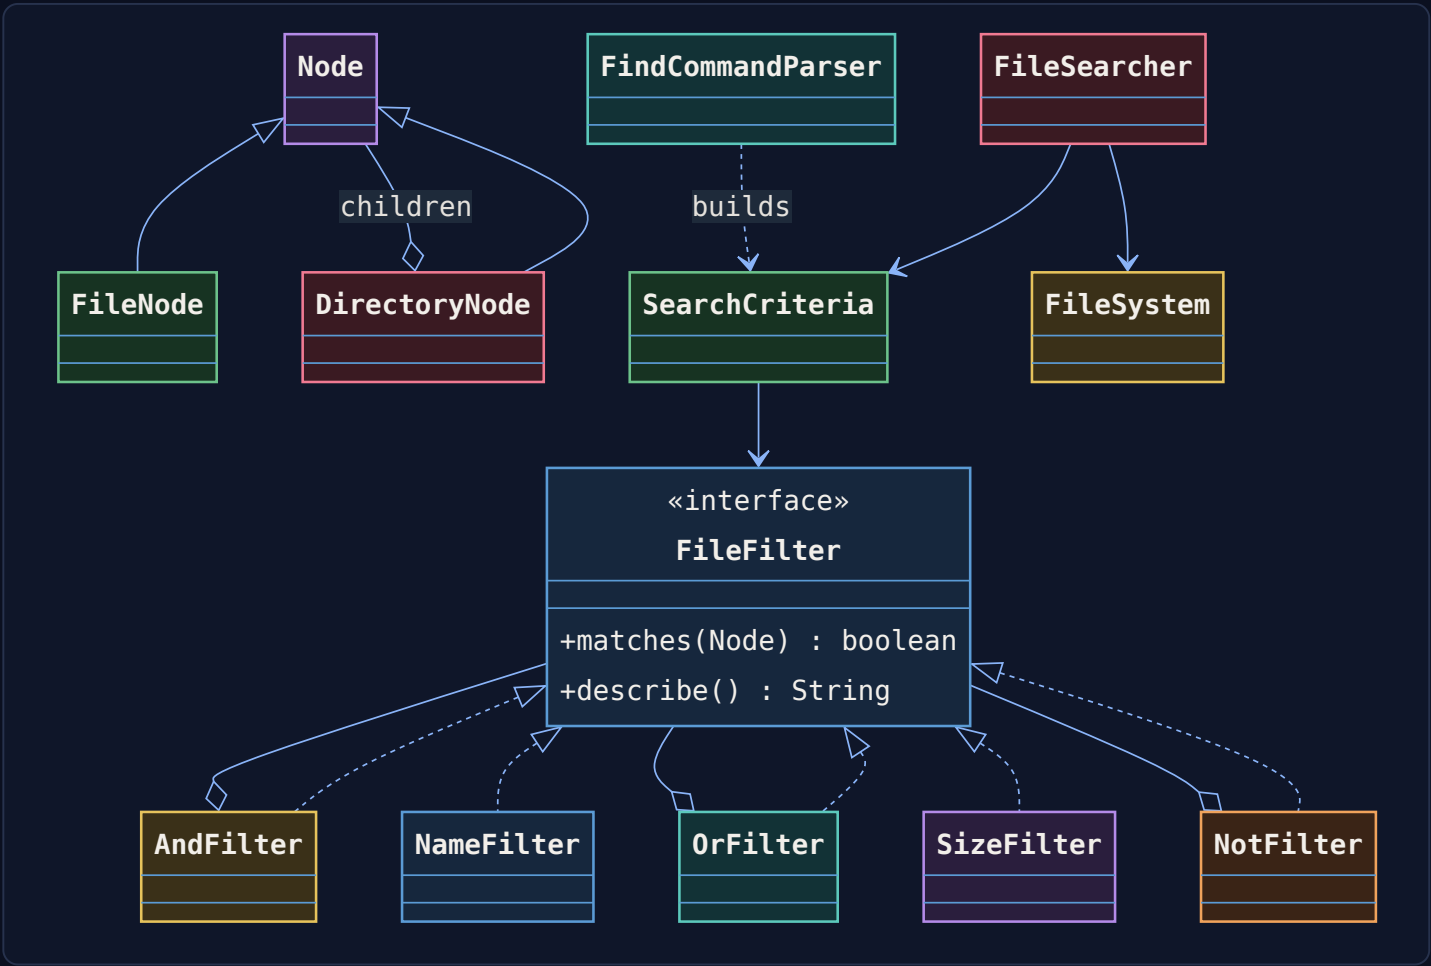
Use cases

Developer (Builder, search, stream), **Shell user** (type a command, get paths or an error), **Ops script** (big old logs, empty directories).

Core entities

Class	Responsibility
`Node` (sealed) → `FileNode`, `DirectoryNode`	**Composite** file tree; `path()`, `size()`, `totalSize()`; children in a sorted `ConcurrentSkipListMap`.
`FileSystem`	Root, `mkdirs`, `createFile`, `resolve`.
`FileFilter`	**Specification**: `matches`, `describe`, default `and` / `or` / `negate`.
`AndFilter`, `OrFilter`, `NotFilter`	**Composite** filter tree.
`Filters`	Leaf tests (name, iname, type, size, mtime, owner, empty, all).
`SearchCriteria` + `Builder`	Immutable, validated search description.
`FileSearcher`	Iterative DFS, streams to a `Consumer`; fork/join variant.
`FindCommandParser`	Recursive descent → `SearchCriteria`.

Class diagram



Two Composites: the **file tree** and the **filter tree**; the searcher walks one and asks the other one question per entry.

Key flows

`find /home ( -name '*.jpg' -o -name '*.png' ) ! -size +1M` parses to **AND(OR(name *.jpg, name *.png), NOT(size > 1M))**.

Search loop: push (start, 0) → pop → if depth ≥ minDepth and filter matches, emit → if directory, depth < maxDepth and not pruned, push children in reverse name order → repeat until empty or limit.

Precedence: `-o` < AND (`-a` or nothing) < `!` < `( )`. So `-name a -o -name b -type f` = `a OR (b AND file)`.

Design decisions & patterns used

Decision	Pattern	Why	Alternative rejected
One small class per test	**Specification**	New test = one record; loop unchanged	`search(name, ext, minSize, …)` with `if`s
And/Or/Not hold filters	**Composite**	Arbitrary nesting; printable, inspectable, parser output	Lambdas only (can't print; cf. §21.1)
`Node` / `FileNode` / `DirectoryNode`	**Composite**	Uniform walk and `totalSize()`	`instanceof` everywhere
`SearchCriteria.Builder`	**Builder**	Named optional settings, validated once, immutable	Telescoping constructors
Recursive-descent parser	Interpreter-style	One method per precedence level	String splitting
Explicit stack, `Consumer` output	Streaming	No stack overflow; constant memory	Recursion; always a full list
Short-circuit AND/OR	—	Cheap tests first	Evaluate every test
`ConcurrentSkipListMap` children	Concurrent collection	Sorted, safe to iterate during writes	`TreeMap` (CME)

Java implementation

```java
// File: lld/ch21/filesearch/FileFilter.java
public interface FileFilter {
    boolean matches(Node n);
    String describe();
    default FileFilter and(FileFilter other) { return AndFilter.of(this, other); }
    default FileFilter or(FileFilter other) { return OrFilter.of(this, other); }
    default FileFilter negate() { return new NotFilter(this); }
}
```

```java
// File: lld/ch21/filesearch/Filters.java  (SizeFilter.matches)
            if (node.type() != FileType.FILE) return false;
            long units = (node.size() + unitBytes() - 1) / unitBytes();      // round up
            return switch (cmp) { case '+' -> units > n; case '-' -> units < n; default -> units == n; };
```

```java
// File: lld/ch21/filesearch/FileSearcher.java
    public int search(SearchCriteria c, Consumer<Node> out) {
        Deque<Frame> stack = new ArrayDeque<>();
        stack.push(new Frame(fs.resolve(c.startPath()), 0));
        int found = 0;
        while (!stack.isEmpty() && found < c.limit()) {
            Frame f = stack.pop();
            Node n = f.node();
            if (f.depth() >= c.minDepth() && c.filter().matches(n)) { out.accept(n); found++; }
            if (n instanceof DirectoryNode d && f.depth() < c.maxDepth() && !c.pruneNames().contains(d.name())) {
                List<Node> kids = new ArrayList<>(d.children());
                Collections.reverse(kids);                    // push in reverse so we pop in name order
                for (Node k : kids) stack.push(new Frame(k, f.depth() + 1));
            }
        }
        return found;
    }
```

```java
// File: lld/ch21/filesearch/FindCommandParser.java
    private FileFilter parseOr() {
        FileFilter left = parseAnd();
        while (peekIs("-o", "-or")) { pos++; left = left.or(parseAnd()); }
        return left;
    }
    // parseAnd: unary ( [-a] unary )* ; parseUnary: ! unary | ( expr ) | -name G | -size N | ...
```

Full runnable code: code/lld/ch21/filesearch/ (see the Full Edition).

How it works: globs compile once to a regex on the last path component; `-size` accepts `[+-]N[cbkMG]` (no suffix = 512-byte blocks); `-mtime` counts whole days against an injected `now`; each Builder `where` is ANDed; the parser skips `-maxdepth` / `-mindepth` at every level; `searchParallel` forks one task per sub-directory and joins in name order — identical results.

Output (excerpt):

```
$ find /home ( -name '*.jpg' -o -name '*.png' ) ! -size +1M
  parsed: from /home where ((name='*.jpg' OR name='*.png') AND NOT size+1M)
   /home/asha/photos/c.png
   /home/asha/photos/d.jpg
$ find /home -size -1M
  parsed: from /home where size-1M
   /home/ravi/todo.md
$ from /home/asha/projects where (name='*.java' OR name='*.js') prune=[node_modules]
   /home/asha/projects/app/Main.java
   /home/asha/projects/app/Util.java
   find / ( -name a                  -> Missing ')'
   find / -maxdepth 1 -mindepth 3    -> mindepth 3 > maxdepth 1
   sequential=1600 parallel=1600 same order=true first=/data/d000/sub0/f060.csv last=/data/d199/sub4/f099.csv
```

`-size -1M` found only the empty file (the rounding gotcha); pruning skipped `node_modules`; 20,000 generated files gave the same 1,600 matches sequentially and in parallel.

Concurrency & edge cases

- Weakly consistent `ConcurrentSkipListMap` iterators: no exception when files are added mid-search.
- Parallel walk only reads; each task builds its own list — no locks.
- Real disks: don't follow symlinks by default (GNU `find` needs `-L`), or track visited inodes.
- Handled: missing start path, `mindepth > maxdepth`, unknown predicate, missing argument, bad `-size` / `-mtime` / `-type`, unbalanced `(`, unclosed `[` in a glob, limit mid-walk.

Extensions & follow-up questions

- **Real disk** → adapter over `java.nio` / `Files.walkFileTree` (`SKIP_SUBTREE` for prune).
- **Actions** (`-delete`, `-exec`) → a `Consumer` per action; delete in post-order.
- **Content search** → `ContainsTextFilter`, placed last in the AND.
- **Repeated searches** → an index like `locate`, possibly stale.
- **More tests** → one record + one parser `case`.

Real-world example

GNU findutils' `find` uses this precedence, short-circuits, rounds `-size` up and does not follow symlinks by default; Java's `Files.find(start, maxDepth, predicate)` splits walking from testing the same way; IDE "Find in files" builds the same filter tree from checkboxes.

In the interview

First 10 minutes: `Node` model + `FileFilter` with two leaves → demo. Then And/Or/Not, the Builder with depths, the iterative walk; the parser last. Say it: "a new filter is one class; the searcher never changes."

Self-check

- **Q:** `-name a -o -name b -type f`? → **A:** `a OR (b AND file)`.
- **Q:** Why does `-size -1M` match only empty files? → **A:** Rounded up, any non-empty file is ≥ 1 MiB unit.
- **Q:** Prune vs filtering a directory out? → **A:** Prune never enters it; a filter only hides the entry.
- **Q:** Why records instead of lambdas for And/Or/Not? → **A:** Printable, inspectable, what the parser builds.

TL;DR

Two Composites (files, filters); Specification leaves combined by AND/OR/NOT; a Builder for criteria; an iterative streaming DFS; a recursive-descent parser with `!` > AND > OR and GNU's rounded-up `-size`.

Chapter summary

Problem	Core entities	Key pattern(s)	Concurrency point	Trap to avoid
Task Management	Task, TaskStatus, TaskService	Specification filters, enum state machine, Observer	per-task `synchronized`	TODO → DONE jumps; filters as many parameters
Digital Wallet	Wallet, Transaction, Offer	Strategy offers, idempotency, funded cashback	lock both wallets in id order	`double` money; check-then-act on keys; cashback from nothing
Event Calendar	TimeSlot, Event, EventStore	Strategy for representatives, plan-then-save	atomic check-and-book	closed intervals; trusting the sample blindly
Coding Blox	Contest, Level, User	Command map, Strategy (solver, scoring)	atomic rating update	unseeded randomness; running a contest twice
PropertyHunt	Property, SearchCriteria	Builder + predicate list, Strategy sort	synchronized mark-sold	comparing "90L" strings; stale shortlist copies
Ledger Co	Loan, LumpSum, Balance	Strategy interest, derived state	synchronized lump sums	storing a running balance; floor instead of ceil
Inventory	StockMovement, AllocationStrategy	Strategy allocation, Observer alerts	per-SKU lock	partial shipments; alert spam
Unix find	Node (Composite), FileFilter, SearchCriteria	Specification + Composite filters, Builder, recursive-descent parser	weakly consistent children map; read-only parallel walk	parameter-list search methods; `-o` vs AND precedence

The pattern behind all eight: entities guard their own rules, services orchestrate, strategies hold the "X or Y" rules, repositories hide storage, the driver parses and prints, and every error is a clear message — not a crash.

Chapter 22 — Revision: Cheat Sheets & Glossary

What this chapter is. The whole handbook on a few pages: the answer framework (Ch 1); OOP, relationships and UML (Ch 2–3); SOLID, other principles and clean code (Ch 4–6); all 24 smells (Ch 7); all 23 GoF patterns + extras and a "which pattern?" table (Ch 8–11); the concurrency toolbox (Ch 12–13); a map of every case study (Ch 14–21); machine-coding and edge-case checklists; and a glossary. Nothing here is new and there is no code — each row tells you where the full explanation lives.

How to read the references. "Ch 12 · Deadlock" means the *Deadlock* topic in Chapter 12; "Ch 16 · Car Rental" means the Car Rental case study in Chapter 16. When a row or heading already names the topic, only the chapter is given ("Ch 2"). Numbers such as "22.7" refer to sections of this chapter.

22.1 How to revise

In 15 minutes (the morning of the interview)

Minutes	Read	Why
0–3	22.2 (the framework)	You will use it in the first minute of the interview
3–6	22.4 (SOLID one-pager) and the relationship table in 22.3	The two things asked about most often by name
6–10	"Which pattern?" table in 22.7	Patterns should come from a signal, not from memory
10–13	Concurrency cheat sheet and deadlock checklist in 22.8	Every booking/inventory problem has a race
13–15	The row(s) of 22.9 for the problems you expect	One sentence per problem: entities, pattern, the atomic step

The night before (about 2 hours)

1. **Read 22.2–22.8 slowly** (45 min), saying *one handbook example* per row; if you can't, open the referenced topic.
2. **Three unpractised case studies** from 22.9 (45 min): requirements, classes, class diagram and the one hard method, without looking; then compare with the chapter's TL;DR.
3. **Scan the glossary** (22.11); read up only the terms you can't explain in one sentence (20 min).
4. **Stop.** Sleep matters more than one more pattern.

Common mistakes when revising

- Re-reading chapters instead of **recalling** first and then checking. Recall is what the interview tests.
- Memorising all 23 patterns equally. Six come up in almost every round (22.7): Strategy, Factory, Observer, State, Singleton, Builder.
- Skipping the "when *not* to use it" column. Interviewers ask "would you always do that?" (Ch 4 · SOLID in interviews).

22.2 The LLD interview answer framework (one page)

Source: Chapter 1 (how the interview runs, the answer framework, requirements, classes, SOLID check, patterns, interview code, communication tips, common mistakes).

TL;DR

Clarify → requirements → use cases → entities → relationships → class diagram → key flows → patterns/SOLID → code → edge cases/concurrency → extensions. Say each step out loud as you enter it.

The eleven steps

#	Step	Ask yourself	Output
1	Clarify	Who uses it? Single site or many?	Shared understanding
2	Requirements	Functional? Non-functional (thread-safe, extensible, in-memory)? **Out of scope?**	8–12 numbered lines + out-of-scope list
3	Use cases	What can each actor do?	Actor → use-case list
4	Entities	Which nouns hold state or behaviour? Which are just values?	Classes, one sentence each
5	Relationships	Uses / knows / has / owns / is-a? How many?	Typed lines with multiplicity
6	Class diagram	Fields, key methods, interfaces, enums	UML class diagram
7	Key flows	Steps of the 2–3 main use cases; lifecycle of the core object	Sequence / activity / state diagram
8	Patterns & SOLID	A rule that varies? A lifecycle? Many listeners? Two reasons to change?	1–2 named patterns + one rejected
9	Code	Interfaces, enums, models, services, driver	Running code for the main flows
10	Edge cases & concurrency	Bad input? Full? Two users, one seat?	Validation, exceptions, locks
11	Extensions	"What if X is added?"	Short answers, maybe one class

Time budgets

45-minute OOD round		**90-minute machine-coding round**	
Clarify requirements	5–8 min	Read, requirements, assumptions	5–10 min
Actors and use cases	3–5 min	Model: entities, enums, service interfaces	10–15 min
Find the classes	8–10 min	Core flow end to end	30–40 min
Connect the classes	8–10 min	Remaining features, validation, edge cases	15–20 min
Write the code	12–15 min	Polish, one extension point, demo	5–10 min
		Buffer	5 min (never zero)

Rule of thumb: roughly **a quarter of the hour goes before you name a class**, and in machine coding you **get something running by the halfway mark** (Ch 1 · The interview).

Finding classes (Ch 1 · Finding the classes)

- **Nouns → candidate classes; verbs → methods.** A verb with no class to live on means a class is missing.
- Cut candidates with five tests: is it a **value**? the **same thing** under another name? does it have **changing state**? does anything **call** it? does it differ in **behaviour or only data**?
- **Different data → field or enum. Different behaviour → class.**
- One sentence per class, **no "and"**. Put each method where its data is (Information Expert).
- Beginners most often forget the **record-keeping class** (ticket, loan, booking) and the **rule class** (rate, strategy).

What the interviewer judges (Ch 1 · The interview)

1. Did you **bound** the problem? 2. Does each class have **one job** you can say in a sentence? 3. Are the **relationships** right? 4. Does the **code match the diagram**? 5. Can you explain a **trade-off** (the choice and the rejected option)?

Machine coding adds: does it **run**, is it **modular**, **readable**, does it **handle bad input**, and (senior roles) is shared state **thread-safe**.

Code habits interviewers notice (Ch 1 · Interview code)

Habit	Instead of
Pass time in (`long nowMinute` , `Instant now` , injected `Clock`)	Reading the clock inside logic
Money as `BigDecimal` from strings, or whole cents in a `long`	`double` (0.1 + 0.2 ≠ 0.3)
`IllegalArgumentException` / `IllegalStateException` / a custom exception	Silent failure or error codes
`Optional` when "not found" is normal	Returning `null`
`private final` fields	Public mutable fields
Domain names (`ParkingTicket`)	`Manager` , `Helper` , `Data2`

Skeleton: **Models** → **Repositories** → **Strategies** → **Services** → **Driver** (Ch 1 · Interview code).

In the interview

30-second answer (from Ch 1 · Common mistakes): "I clarify and write 8 to 12 requirements with an out-of-scope list, list actors and use cases, pick classes from the nouns and give each one responsibility, connect them and draw a class diagram, sketch the key flow, run a quick SOLID check and pick one or two patterns where the signal is clear, then code the two or three core flows, handle edge cases and concurrency, and finish with extensions and trade-offs."

Common mistakes (Ch 1)

Mistake	Fix
Coding in minute two	8–13 minutes on requirements and actors first
No out-of-scope list	Write it down and agree it
One subclass per value (`CompactSpot` , `LargeSpot`)	Different data → field/enum
God class (`ParkingLotManager`)	One sentence per class; split on "and"
Inheritance for reuse ("Car extends Engine")	Composition
Pattern stuffing	One or two patterns, each with its signal
All twelve classes at equal depth	Two or three flows fully
Ignoring two users on one resource	Lock or atomic operation on the shared state
Non-running code in machine coding	Thin end-to-end path running by halfway

---

22.3 OOP & relationships cheat sheet

Source: Chapter 2 (OOP) and Chapter 3 (relationships and UML).

The four pillars (Ch 2)

Pillar	One line	Java tools	Smell when missing
Encapsulation (Ch 2)	Data is private; only the owning class's methods change it, and they enforce the rules	`private` fields, no setters, validation in constructors	Public fields; setters that break invariants
Abstraction (Ch 2)	Expose only what callers need; hide how	Interfaces, abstract classes, small public API	Callers reach into internals
Inheritance (Ch 2)	Child reuses and specialises parent — **only when "child is a parent" is truly correct**	`extends`, `super`, `@Override`	Subclass refuses or throws on parent methods
Polymorphism (Ch 2)	One call, different behaviour per object; removes type switches	Overriding + dynamic dispatch	`switch` / `instanceof` chains on type

Encapsulation vs abstraction: hiding *data* behind rule-enforcing methods vs hiding *complexity* behind a minimal set of operations (Ch 2 self-check).

Other OOP facts worth one line each

Topic	Remember	Where
Overloading vs overriding	Compile time, by declared parameter types vs run time, by the actual object. Fields and `static` methods are not polymorphic (statics are hidden, not overridden)	Ch 2 · Compile-time vs runtime polymorphism
`this`, `static`, `final`	Current object; one per class; cannot be reassigned/overridden/extended	Ch 2 · Classes & keywords
Interfaces	Contracts; can have `default`, `static`, `private` methods; a functional interface has one abstract method (lambdas)	Ch 2
Enums	Fixed sets with fields and behaviour ("which spot types fit a truck" lives on `VehicleType`)	Ch 2 · Enums with fields and behaviour
Composition over inheritance	Avoids fragile base classes; allows run-time swapping	Ch 2, Ch 5 · Composition over inheritance
`equals` / `hashCode`	Equal objects must have equal hash codes; override both or hash collections break	Ch 2 · Object identity vs equality
Immutability & records	Final class, private final fields, no setters, defensive copies. Records are only **shallowly** immutable — copy collections in the compact constructor	Ch 2 · Immutability & records
Access modifiers	Default to `private`; package-private = same package only	Ch 2 · Access modifiers and packages
Generics	Type-safe reusable infrastructure; **PECS:** producer `extends`, consumer `super`	Ch 2 · Generics basics for LLD
Coupling & cohesion	Low coupling (between classes), high cohesion (inside a class)	Ch 2 · Coupling and cohesion
OOA → OOD → OOP	What the problem is → how, with objects → code. Connections that carry data become objects	Ch 2 · Object-oriented analysis and design

Abstract class vs interface (Ch 2 · Abstract classes vs interfaces)

	Abstract class	Interface
Instantiable?	No	No
Instance fields (state)	Yes	No (only `public static final` constants)
Constructors	Yes	No
Method bodies	Any	`default`, `static`, `private` only
How many per class	**One** (`extends`)	**Many** (`implements`)
Relationship expressed	"is a" (shared identity and state)	"can do" (capability / role)
Typical use	Shared state + partial implementation for a close family (Template Method)	Contracts, roles, strategies, test seams

Rule of thumb: start with an interface; add an abstract class underneath when implementations share real state or a fixed algorithm (JDK: `List` + `AbstractList` + `ArrayList`).

The six relationships (Ch 3)

Weakest to strongest coupling: **dependency, association, aggregation, composition, realization, generalization** (Ch 3 · self-check).

Relationship	Phrase	UML line	Java shape	Lifetime	Example
Dependency (Ch 3)	A *uses* B	dashed, open arrow → B	B as parameter / local / return / static call	none	`CheckoutService` uses `PaymentGateway`
Association (Ch 3)	A *knows* B	solid line (arrow for one-way)	field of type B, created elsewhere	independent	`Order` → `Customer`
Aggregation (Ch 3)	A *has* B, parts live on	solid, hollow ◇ at A (the whole)	field/collection, parts passed in	B outlives A	`Airline` ◇— `Aircraft`
Composition (Ch 3)	A *owns* B	solid, filled ◆ at A (the whole)	A creates B, hides it, B not shared	B dies with A	`ParkingFloor` ◆— `ParkingSpot`
Realization (Ch 3)	A *implements* B	dashed, hollow △ at B	`implements`	—	`EmailNotifier` ⇢ `Notifier`
Generalization (Ch 3)	A *is a* B	solid, hollow △ at B	`extends`	—	`CardPayment` → `Payment`

Choosing, in order (Ch 3 · Choosing the relationship): interface A promises to fulfil → realization. Substitutable with no refused methods → generalization. Not kept in a field → dependency. Part that dies with the whole → composition. Part that outlives it → aggregation. Otherwise → association.

UML arrow legend (Mermaid)



Mermaid syntax: `A ..> B` dependency · `A --> B` association (navigable A→B; `A -- B` unspecified) · `A o-- B` aggregation and `A *-- B` composition (A is the whole) · `B <|.. A` realization · `B <|-- A` generalization · `A "1" *-- "1..*" B` adds multiplicities.

Multiplicity → code (Ch 3 · TL;DR): `1` required field; `0..1` `Optional` /nullable; `1..*` non-empty collection checked in the constructor; exact numbers checked in constructors. Visibility: `+` public, `-` private, `#` protected, `~` package.

UML diagrams in the interview (Ch 3 · UML diagrams)

Diagram	Use it for	Budget in 45 min
Use case (Ch 3 · Use-case diagrams)	Scope: actors + what each does. «include» = always; «extend» = sometimes	Usually spoken, 3–5 min
Class (Ch 3 · Class diagrams)	**The answer.** 10–15 boxes, every line typed, with multiplicity	Always drawn, 10–15 min
Sequence (Ch 3 · Sequence diagrams)	Call order between objects (3–4 objects, 6–8 messages)	One behaviour diagram only
Activity (Ch 3 · Activity diagrams)	A flow that branches; fork/join for parallel steps	"
State machine (Ch 3 · State machine diagrams)	One object's lifecycle; transition label `event [guard] / action`	"

UML has 14 diagram types: 7 structural, 7 behavioural (Ch 3 · What UML is).

Common mistakes

- Diamond at the wrong end (it goes at the **whole**); triangle at the wrong end (it goes at the **parent/interface**) (Ch 3 · Choosing the relationship).

- Lines without type or multiplicity (Ch 3 · Choosing the relationship).
- Inheritance used to share data; subclasses that only pass a constant to `super` → one class with an enum field (Ch 3 · self-check).

22.4 SOLID one-pager

Source: Chapter 4 (the five principles, SOLID and patterns, SOLID in the interview, over-engineering); smells from Chapter 7 (smell lookup table).

TL;DR

SOLID = five names for the ways code becomes painful to change. **Use them to diagnose, not as a checklist**, apply a fix when a real change asks for it, and always say what the fix costs.

Principle	One-line rule	Smell that signals a violation	Fix	Patterns that help
S — Single Responsibility (Ch 4 · Single Responsibility Principle)	A class should have one reason to change (one actor asks for its changes)	The responsibility sentence needs an "and"; *Divergent Change*, *Large Class*	Split by actor (Extract Class); each new class needs a name that stands alone	Command, Proxy, Decorator (each concern in its own class)
O — Open/Closed (Ch 4 · Open/Closed Principle)	Add new behaviour by adding code, not by editing working code	A `switch` / `if` chain over a type; *Repeated Switches*, *Shotgun Surgery*	Put the part that *has* varied behind an interface; add variants as classes (Replace Conditional with Polymorphism)	Strategy, Template Method, Decorator, Factory Method, Observer, Command, Visitor
L — Liskov Substitution (Ch 4 · Liskov Substitution Principle)	A subtype must be usable anywhere its parent is expected	Subclass throws, does nothing, or refuses a parent method; *Refused Bequest*	Better parent, a narrower interface, or composition (Replace Subclass with Delegate)	Composite, Null Object
I — Interface Segregation (Ch 4 · Interface Segregation Principle)	No client should depend on methods it does not use	Empty method bodies or `UnsupportedOperationException` in implementers; fat interfaces	Split into **role interfaces**, one per kind of caller	Adapter
D — Dependency Inversion (Ch 4 · Dependency Inversion Principle)	Depend on abstractions, not concrete details; the policy owns the abstraction	A high-level class does `new` of (or names) a low-level one; cannot test without the real DB/provider	Inject an interface owned by the high-level class; `main` wires the details	Strategy, Factory Method / Abstract Factory, Adapter, Observer

LSP in detail (Ch 4 · Liskov Substitution Principle, TL;DR): a subclass must accept at least as much, promise at least as much, keep invariants and history, and throw nothing new. Square/Rectangle is worse than `Penguin.fly()` throwing because it fails **silently** (wrong area, no error).

Five-minute pass over your diagram (Ch 4 · SOLID in interviews, Ch 1 · SOLID check)

Run S → O → L → I → D over the diagram you just drew. **Two findings is a good pass; five means rework; zero means you did not look.** S and O earn the most credit: O is what interviewers probe with a late requirement ("what if trucks are charged differently?" → "add a class").

Related ideas (Ch 4 · Cohesion, Coupling, SoC, SOLID Q&A)

Idea	Question it asks
SRP	How many actors can ask this class to change?
Cohesion	Do the parts of this class belong together?
Coupling	How much does A depend on B?
Separation of Concerns	Is each area of work in its own module?
SRP vs ISP	SRP splits **classes** (guards the worker); ISP splits **interfaces** (guards the caller)
DIP vs DI vs IoC	DIP = principle (depend on abstractions); DI = technique (pass dependencies in); IoC = the broader idea that a framework calls your code; IoC containers perform DI

Common mistakes: over-engineering (Ch 4)

Warning sign	Principle misapplied	Do instead
`XHelper`, `XUtils`, `XManager2` after a split	SRP	Split by actor; names must stand alone
Six small classes that always change in the same commit	SRP over-split	Merge them back
Interface + factory + registry for one implementation	OCP (*Speculative Generality*)	Write the `if`; refactor at the second or third case (rule of three)
A subclass that "works" by silently doing nothing	LSP	Better parent or composition
One interface per method	ISP over-split	Group by caller into role interfaces
An interface in front of stable `String`-like classes	DIP	Invert only volatile dependencies
A DI container for a 200-line program	DIP	Wire by hand in `main`

YAGNI is the counterweight: SOLID says *what* a good seam looks like; YAGNI says *when* to cut it.

In the interview

Say **observation → cost → fix → price of the fix**, and describe before you label: "`Reporting` implements nine methods but calls two — an interface segregation problem" (Ch 4 · SOLID in interviews).

22.5 Other design principles and Clean Code rules

Other principles (Chapter 5)

Principle	One-line idea	Smell it fixes	Key Java tool	Where
DRY	Each piece of knowledge has one authoritative home	Duplicated rules drifting apart	Value objects, constants, methods	Ch 5 · DRY
KISS	Choose the simplest design that works	Over-engineering	Plain code, `switch` expressions	Ch 5 · KISS
YAGNI	Don't build features before they are needed	Speculative Generality	Delete unused code	Ch 5 · YAGNI
Law of Demeter	Only talk to your immediate friends	Train wrecks `a.getB().getC().do()`	Intent-revealing delegate methods	Ch 5 · Law of Demeter
Tell, Don't Ask	Tell objects what to do; don't pull their data out and decide for them	Anemic model, Feature Envy	Command methods on the owner	Ch 5 · Tell, Don't Ask
Composition over inheritance	Prefer "has-a" to "is-a" for reuse	Fragile base class	Forwarding / decorators	Ch 5 · Composition over inheritance
Program to an interface	Depend on supertypes, choose the concrete class once at the edge	Hard-wired concrete dependencies	Interfaces + DI	Ch 5 · Program to an interface
Encapsulate what varies	Isolate the part that changes behind an interface	Growing type switches	Strategy, lambdas	Ch 5 · Encapsulate what varies
Loose coupling	Interacting objects know little about each other	God orchestrators	Observer, events, DI	Ch 5 · Loose coupling
Hollywood Principle / IoC	"Don't call us, we'll call you" — the framework owns the flow	Duplicated skeletons	Template Method, callbacks, DI	Ch 5 · Hollywood Principle
Separation of Concerns	Each part handles one concern	Spaghetti code	Layers, MVC	Ch 5 · Separation of Concerns, Ch 4
Least Astonishment	Do what the name promises and nothing more	Hidden side effects	Honest names, `Optional`	Ch 5 · Least Astonishment
Fail fast	Detect problems early and loudly	Distant `NullPointerException` s	Guard clauses, `requireNonNull`	Ch 5 · Fail fast
Immutability	Objects never change after creation	Aliasing, races	Records, `List.copyOf`	Ch 5 · Immutability
Design by Contract	Preconditions (caller), postconditions (method), invariants (always)	Magic error values	Checks at method entry/exit	Ch 5 · Design by Contract
GRASP	Nine guidelines for "which class should do this?"	God classes	Expert, Creator, Controller…	Ch 5 · GRASP
Domain-Driven Design	Model the business: bounded contexts, entities vs value objects, small aggregates, domain events, repositories per root	Anemic model, big ball of mud, primitive obsession	Records as value objects, aggregate roots, events + outbox	Ch 5 · DDD

GRASP in one line each (Ch 5 · GRASP):

Guideline	Rule
Information Expert	The class with the data does the work
Creator	The container (or the one with the initial data) creates
Controller	A thin non-UI object receives each system event and delegates
Low Coupling	Choose the design with fewer, more stable dependencies
High Cohesion	Each class does one related set of things
Polymorphism	Let each type implement the varying behaviour; don't switch on types
Pure Fabrication	Invent a helper class (e.g. a repository) when no domain class is a clean home
Indirection	Put a middleman between things that shouldn't know each other
Protected Variations	Wrap what's unstable behind a stable interface

Tensions to mention (Ch 5 summary): DRY vs loose coupling (sharing code couples its users); YAGNI vs encapsulate-what-varies (abstract only real or clearly predicted variation); Law of Demeter vs cohesion (too many delegates → Middle Man); Information Expert vs SoC (domain objects compute, fabrications persist).

Clean Code rules checklist (Chapter 6)

Area	Rule of thumb	Where
Names	Reveal intent; pronounceable; searchable; one word per concept; code reads like the requirement	Ch 6 · Meaningful names
Functions	Small; do one thing; one level of abstraction; top-down (stepdown); ≤ 2–3 arguments; **no flag arguments**; no hidden side effects; commands *or* queries (CQS); exceptions, not error codes	Ch 6 · Functions
Comments	Say it in code; comment only the *why*, warnings, legal text, public contracts; delete commented-out code	Ch 6 · Comments
Formatting	Newspaper order (headline first); related things close; one automated team style	Ch 6 · Formatting
Objects vs data structures	Objects hide data and expose behaviour; data structures expose data; no hybrids; no train wrecks	Ch 6 · Objects and data structures
Error handling	Unchecked exceptions with context; wrap vendor exceptions; empty collections / `Optional` / special case instead of `null`; never pass `null`; always clean up	Ch 6 · Error handling
Boundaries	Keep third-party code behind your own interfaces; learn it with tests	Ch 6 · Boundaries
Unit tests	Three laws of TDD; one concept per test; Given–When–Then; **F.I.R.S.T.** (Fast, Independent, Repeatable, Self-validating, Timely)	Ch 6 · Unit tests
Classes	Small by responsibility; cohesive; open for extension; depend on abstractions	Ch 6 · Classes
Systems	Build objects in one place (composition root) and inject them; cross-cutting concerns in proxies/aspects	Ch 6 · Systems
Simple design	Beck: passes the tests → reveals intention → no duplication → fewest elements (*Clean Code* swaps the middle two)	Ch 6 · Simple design
Concurrency	Separate concurrency code; share as little as possible; use the library; keep locks small; plan shutdown	Ch 6 · Concurrency
Smells	The C/E/F/G/J/N/T heuristics catalogue as a review vocabulary	Ch 6 · Smells and heuristics
Code review	Correctness → design → readability → tests → security → performance → operability; automate style; small PRs	Ch 6 · Code review checklist
Error strategy	Classify, translate at boundaries, retry only transient + idempotent with backoff, one global handler that logs once	Ch 6 · Error-handling strategy
Testing strategy	Many unit, some integration, few E2E; right test double; verify state over interactions; don't mock what you don't own	Ch 6 · Testing strategies

Boy Scout Rule: leave the code a little cleaner than you found it (Ch 6 intro).

22.6 Code smells → refactorings (all 24)

Source: Chapter 7 (code smells, refactoring catalogue, smell lookup table). ***Bold*** *= the usual first move.*

#	Smell	Signal (one line)	Group	Refactorings that cure it
1	Mysterious Name	Name doesn't say what it does/holds (`data` , `proc()` , `Manager`)	2nd-ed extra	**Change Function Declaration**, Rename Variable, Rename Field
2	Duplicated Code	Same structure in more than one place	Dispensables	**Extract Function**, Slide Statements, Pull Up Method
3	Long Function	You want a comment to explain a block	Bloaters	**Extract Function**, Replace Temp with Query, Introduce Parameter Object, Preserve Whole Object, Replace Function with Command, Decompose Conditional, Replace Conditional with Polymorphism, Split Loop
4	Long Parameter List	Many parameters, same types, boolean flags	Bloaters	Replace Parameter with Query, **Preserve Whole Object**, **Introduce Parameter Object**, Remove Flag Argument, Combine Functions into Class
5	Global Data	Readable and writable from anywhere (incl. mutable singletons)	2nd-ed extra	**Encapsulate Variable** (then shrink its scope)
6	Mutable Data	Shared data changed after creation	2nd-ed extra	**Encapsulate Variable**, Split Variable, Slide Statements, Extract Function, Separate Query from Modifier, Remove Setting Method, Replace Derived Variable with Query, Combine Functions into Class/Transform, Change Reference to Value
7	Divergent Change	One class changed for many different reasons (SRP)	Change Preventers	**Split Phase**, Move Function, Extract Function, **Extract Class**
8	Shotgun Surgery	One change forces many small edits in many classes	Change Preventers	**Move Function**, **Move Field**, Combine Functions into Class/Transform, Split Phase, Inline Function, Inline Class
9	Feature Envy	Method uses another class's data more than its own	Couplers	**Move Function**, Extract Function
10	Data Clumps	Same group of values always travel together (`startDate, endDate`)	Bloaters	**Extract Class**, **Introduce Parameter Object**, Preserve Whole Object
11	Primitive Obsession	Money as `double` , status as `String`	Bloaters	**Replace Primitive with Object**, Replace Type Code with Subclasses, Replace Conditional with Polymorphism, Extract Class, Introduce Parameter Object
12	Repeated Switches	The *same* switch on a type code in several places	OO Abusers	**Replace Conditional with Polymorphism**
13	Loops	Classic loop that filters/maps/aggregates	2nd-ed extra	**Replace Loop with Pipeline**
14	Lazy Element	A class/function that doesn't pull its weight	Dispensables	**Inline Function**, **Inline Class**, Collapse Hierarchy
15	Speculative Generality	Hooks "we might need someday"; abstract class with one subclass	Dispensables	Collapse Hierarchy, Inline Function, Inline Class, Change Function Declaration, Remove Dead Code
16	Temporary Field	Field set only in some situations, `null` otherwise	OO Abusers	**Extract Class**, Move Function, Introduce Special Case

#	Smell	Signal (one line)	Group	Refactorings that cure it
17	Message Chains	`order.getCustomer().getAddress().getCity()`	Couplers	**Hide Delegate**, Extract Function, Move Function
18	Middle Man	Most methods just forward to another class	Couplers	**Remove Middle Man**, Inline Function, Replace Superclass/Subclass with Delegate
19	Insider Trading	Two modules poke at each other's internals	Couplers	**Move Function**, **Move Field**, Hide Delegate, Replace Subclass/Superclass with Delegate
20	Large Class	Too many fields/methods — a God Class	Bloaters	**Extract Class**, Extract Superclass, Replace Type Code with Subclasses
21	Alternative Classes with Different Interfaces	Same job, different method names (`sendMail` vs `dispatch`)	OO Abusers	**Change Function Declaration**, Move Function, Extract Superclass
22	Data Class	Only fields and getters/setters; others manipulate it	Dispensables	**Encapsulate Record**, Remove Setting Method, Move Function, Extract Function, Split Phase
23	Refused Bequest	Subclass doesn't want inherited methods, or throws on them (LSP)	OO Abusers	Push Down Method/Field, **Replace Subclass with Delegate**, **Replace Superclass with Delegate**
24	Comments	Comment used as deodorant for unclear code	Dispensables	**Extract Function**, Change Function Declaration, Introduce Assertion

Smell family → principle violated (Ch 7 · Smell lookup table)

Smell family	Principle
Divergent Change, Large Class	SRP
Repeated Switches, Shotgun Surgery	OCP
Refused Bequest	LSP
Long Parameter List with flags, fat interfaces	ISP / clean function design
Global Data, Mutable Data	Encapsulation
Hard-wired `new` of concrete dependencies	DIP
Duplicated Code	DRY
Speculative Generality	YAGNI
Message Chains, Insider Trading	Law of Demeter
Feature Envy, Data Class	Tell, Don't Ask

Refactoring toward patterns (Ch 7)

Smell → pattern, in short (the full symptom table is in 22.7): algorithm-choosing switch → **Strategy**; methods checking and changing a status field → **State**; `if (x == null)` everywhere → **Null Object / Special Case**; sibling subclasses with the same skeleton → **Template Method**; many confusing constructors → **Factory Method** / creation methods; object built by many ordered calls → **Builder**; undo/queueing, or a huge function with shared locals → **Command**; logging/caching mixed into a core class → **Decorator**; polling or hard-coded notification → **Observer**; two classes, same job, different interfaces → **Adapter**; subclass explosion from feature combinations → **Strategy / Decorator / Bridge**; a pattern that no longer earns its keep → **remove it** (Inline Class, Collapse Hierarchy).

Refactoring = changing structure without changing behaviour, in small steps, with tests as the safety net. Wear one of **two hats** at a time: adding function or refactoring (Ch 7 · Refactoring fundamentals).

22.7 Design patterns: all 23 GoF + extras

Source: Chapters 8–11 (catalogue in Ch 8 · The 23 GoF patterns; comparison tables at the end of Chapters 8, 9, 10 and 11).

Count: 5 creational + 7 structural + 11 behavioral = **23**. **Class-scoped** (inheritance) patterns: Factory Method, class Adapter, Template Method, Interpreter; almost everything else is object-scoped (composition) (Ch 8 · The 23 GoF patterns).

Creational (Chapter 8)

Pattern	Intent (one line)	Use when	Java / JDK example	Commonly confused with	Where
Singleton	Exactly one instance with a global access point	There really must be exactly one (config, pool, registry)	`Runtime.getRuntime()`; Spring singleton-scoped beans	A global variable; a DI singleton scope (preferred for testing)	Ch 8
Factory Method	Subclasses decide which concrete class to create via an overridable method	Shared procedure, varying product; others will add types	`Collection.iterator()`	Simple Factory (one method with a switch)	Ch 8
Abstract Factory	Create *families* of related objects without naming concrete classes	Products must match; swap a whole set at once	`DocumentBuilderFactory`; JDBC `Connection`	Factory Method (one product by inheritance vs a family by composition)	Ch 8
Builder	Build a complex object step by step; one validating `build()`	Many optional fields, cross-field validation, immutable result	`StringBuilder`, `HttpClient.newBuilder()`	A plain constructor (fine for a handful of required fields); telescoping constructors	Ch 8
Prototype	Create new objects by copying a configured instance	Setup is expensive, or only the interface is known	Copy constructors, `ArrayList.clone()`	Memento (a snapshot to restore, not a copy to use)	Ch 8

Singleton notes (Ch 8 · TL;DR): eager, holder-class or **`enum`** (most robust: blocks reflection and serialization duplicates); double-checked locking needs a **`volatile`** field. It is global state that hurts testing — prefer injecting one shared instance. Many case studies reject it (Ch 14 self-check).

Structural (Chapter 9)

Pattern	Intent (one line)	Use when	Java / JDK example	Commonly confused with	Where
Adapter	Convert one interface into another that clients expect	"The vendor API doesn't match ours" (after the fact)	`Arrays.asList`, `InputStreamReader`	Bridge (same shape, but Bridge is designed up front); Facade	Ch 9 · Adapter Pattern
Bridge	Split abstraction from implementation so both vary independently	Class names with two nouns; two dimensions multiplying classes	JDBC `Connection` + drivers	Adapter; Strategy (same shape)	Ch 9 · Bridge Pattern
Composite	Treat single objects and trees uniformly	Part-whole trees (files, menus, UI)	`java.awt.Container` is a `Component` holding `Component`s	Decorator (also wraps a component, but adds behaviour to one)	Ch 9 · Composite Pattern
Decorator	Add responsibilities dynamically by wrapping; same interface, stackable	"Sometimes logging, sometimes caching, sometimes both"	`BufferedInputStream`, `Collections.synchronizedList`	Proxy (same shape; Proxy *controls access*); Chain of Responsibility (in CoR someone can stop the flow)	Ch 9 · Decorator Pattern
Facade	One simple entry point to a complex subsystem	Every caller repeats the same multi-step sequence	SLF4J `LoggerFactory`, service layers	Adapter (converts one object's interface); Mediator	Ch 9 · Facade Pattern
Flyweight	Share intrinsic (immutable) state across huge numbers of objects	Millions of near-identical objects	`Integer.valueOf`, `String.intern()`	Object Pool (reuses objects one borrower at a time; flyweights are shared at once)	Ch 9 · Flyweight Pattern
Proxy	A stand-in that controls access (lazy, remote, protective, caching)	It's expensive, remote or restricted	`java.lang.reflect.Proxy`, Spring `@Transactional`, Hibernate lazy load	Decorator	Ch 9 · Proxy Pattern

Memory hooks (Ch 9 · Structural comparison): Adapter *make it fit* · Decorator *add to it* · Proxy *guard it* · Facade *simplify it* · Bridge *split it* · Composite *nest it* · Flyweight *share it*.

Behavioral (Chapters 10–11)

Pattern	Intent (one line)	Use when	Java / JDK example	Commonly confused with	Where
Chain of Responsibility	Pass a request along handlers until one handles (or stops) it	Several checks, any may reject; handler not known in advance	Servlet `Filter` chain; `java.util.logging` handlers	Decorator (every layer always runs)	Ch 10
Command	Turn a request into an object: queue, log, retry, undo	Undo/redo, job queues, scheduling	`Runnable` / `Callable`; Swing `UndoManager`	Strategy (Strategy = *how*; Command = *what*)	Ch 10
Iterator	Traverse a collection without exposing how it is stored	Walk a collection/tree/paged API	`Iterator`, `Spliterator`, JDBC `ResultSet`	Visitor (what to do at each node)	Ch 10
Mediator	Centralise communication between peers in one coordinator	Many objects interact in tangled ways; order of reactions matters	Chat servers, UI controllers; `Executor` / `Timer` are a loose fit (Ch 10)	Observer (broadcast, no coordination); Facade	Ch 10
Memento	Capture and restore state without breaking encapsulation	Undo, checkpoints, rollback	A `record` snapshot; DB `ROLLBACK`, Git commits (real-world)	Command undo (store actions, not states); Prototype	Ch 10
Interpreter	One class per grammar rule; evaluate sentences of a small language	User-written rules/expressions	`java.util.regex.Pattern`, SpEL	Composite (same tree); Visitor (many operations over the AST)	Ch 10
Observer	Notify many dependents automatically when one object changes	An open-ended set of independent listeners	`PropertyChangeListener`, `java.util.concurrent.Flow`, Spring `@EventListener`	Mediator; pub-sub (Observer through a broker, Ch 20 · Pub-sub)	Ch 11
State	Behaviour changes with internal state, as if the class changed	Many states/transitions; methods repeat a `switch` on status	TCP connection (GoF's example), order lifecycles. `Thread.State` is only an enum of states, not the pattern	Strategy (same diagram; states know each other and advance themselves)	Ch 11
Strategy	Make a family of algorithms interchangeable behind one interface	Several ways to do one job, chosen at run time	`Comparator`; `RejectedExecutionHandler` policies	State; Template Method; Decorator (Strategy *replaces* a rule, Decorator *wraps* one)	Ch 11
Template Method	Fixed algorithm skeleton in a base class; subclasses fill in steps	Long fixed flow, a few varying steps	`AbstractList`, `InputStream.read(byte[],…)`, `HttpServlet.service()`	Strategy (Template = inheritance, one step varies;	Ch 11

Pattern	Intent (one line)	Use when	Java / JDK example	Commonly confused with	Where
				Strategy = composition, whole algorithm varies)	
Visitor	Add operations to a stable set of classes without changing them (double dispatch)	Many new operations over a fixed type set (ASTs)	`javax.lang.model.element.ElementVisitor`; `FileVisitor` is named after it but is really a tree-walk callback	Iterator; in Java 21, **sealed types + pattern-matching** `switch` usually replace it	Ch 11

Beyond GoF: extras you will meet

Pattern	Intent (one line)	Use when	Java / real example	Commonly confused with	Where
Simple Factory (idiom)	One `create(type)` method maps a type value to a class	Small, stable set of types chosen by config	`Calendar.getInstance()`, `EnumSet.noneOf`	Factory Method	Ch 8
Object Pool	Borrow and return a bounded set of expensive objects	DB connections, limited resources	HikariCP, `ThreadPoolExecutor`	Flyweight	Ch 8
Null Object	A do-nothing implementation instead of `null`	`null` checks everywhere for an optional collaborator	`Collections.emptyList()`, `OutputStream.nullOutputStream()`	`Optional` (for lookups)	Ch 11 · Beyond GoF
Dependency Injection	Receive collaborators (as interfaces) through the constructor	Classes create their own collaborators; hard to test	Spring constructor injection; a hand-wired `main`	Service Locator (anti-pattern: the class asks for its dependencies)	Ch 11 · Beyond GoF
Repository	Collection-like interface over storage	Business logic mixed with storage code	Spring Data `CrudRepository`	DAO (table-oriented)	Ch 11 · Beyond GoF
Specification	A named, combinable business rule (`isSatisfiedBy`, and/or/not)	Same rule in queries and validation; finder-method explosion	`Predicate`; Spring Data JPA `Specification`	—	Ch 11 · Beyond GoF
MVC	Separate model, view and controller (Observer + Strategy + Composite)	UI, input handling and rules tangled	Spring MVC; Swing `JTable` + `TableModel`	MVP, MVVM (relatives)	Ch 11 · Beyond GoF
Event Sourcing (lite)	Store the events; derive current state by replaying them	Full history, audit, time travel	Append-only logs	Plain audit log	Ch 11 · Beyond GoF
Architectural styles (layered, client-server, modular monolith vs microservices, event-driven, pipe-and-filter, hexagonal/clean)	Fix the shape of the *whole system*: parts, dependency direction, communication	Deciding where classes live and answering "would you split this?"	Spring layered apps; Kafka consumers; Unix pipes; ports and adapters	Design patterns (a few classes inside one part); MVC (only the UI edge)	Ch 11 · Architectural patterns

"Which pattern?" by problem symptom

Combined from Ch 8 · How to pick a pattern, Creational patterns compared, Ch 9 · Structural comparison, Ch 11 · Choosing a behavioral pattern and Ch 7 · Refactoring toward patterns.

If you see / need…	Use	Not to be confused with
One operation, several interchangeable rules chosen at run time (pricing, matching, routing)	**Strategy**	State, Template Method
The class to create depends on a type value; caller shouldn't know the concrete class	**Factory** (Simple Factory / Factory Method)	Abstract Factory
Products must always match as a family	**Abstract Factory**	Factory Method
Constructor has ~6+ arguments, many optional	**Builder**	—
Another one of *this*, cheaply or without knowing its class	**Prototype**	Memento
Exactly one instance must exist	**Singleton** (or DI singleton scope)	A convenient global handle
Expensive objects created and destroyed repeatedly	**Object Pool**	Flyweight
Existing class has the wrong interface	**Adapter**	Facade, Bridge
Two dimensions multiplying classes (`WindowsButton` , `MacButton` , …)	**Bridge**	Adapter
Tree of part-whole objects treated the same	**Composite**	Decorator
Add behaviour in combinations without editing the class	**Decorator**	Proxy
Every caller repeats the same five steps into a subsystem	**Facade**	Mediator
Millions of near-identical objects	**Flyweight**	Object Pool
Lazy load, access control, remote, caching (same interface)	**Proxy**	Decorator
A request must pass several checks; any may reject it	**Chain of Responsibility**	Decorator (all layers run)
Undo/redo, queue, schedule, retry or log *requested actions*	**Command**	Strategy
Evaluate user-written rules in a small language	**Interpreter** (+ Specification)	Visitor
Walk a collection/tree/paged API without exposing structure	**Iterator**	Visitor
Many objects depend on each other and the **order** of reactions matters	**Mediator**	Observer
Roll back to an earlier state without exposing internals	**Memento**	Command undo
One change must reach an open-ended set of listeners	**Observer** (or a broker)	Mediator
Three or more methods repeat one `switch` over a status; legal actions depend on it	**State** (or enum + guarded method for 3–4 states)	Strategy
Several classes share a long fixed flow, differ in a few steps	**Template Method**	Strategy
Many new operations over a stable set of types	**Visitor**, or sealed types + pattern matching	Iterator
`null` checks everywhere	**Null Object** / `Optional`	—
Classes create their own collaborators	**Dependency Injection**	Service Locator
Business logic mixed with storage	**Repository**	DAO
Same rule duplicated in queries and validation	**Specification**	—

Three questions for the famous look-alikes (Ch 11 · Choosing a behavioral pattern):

1. Do the alternatives reference each other? Yes → **State**. No → **Strategy** (or Bridge if chosen once at wiring).
2. Can one participant stop the flow? Yes → **Chain of Responsibility**. Everyone always runs → **Decorator**. Everyone is told, nobody stops anyone → **Observer**.

3. Does a central object decide what happens next? Yes → **Mediator**. Each listener decides for itself → **Observer**.

The six that come up in almost every round (Ch 1 · Patterns in this round, Ch 8 · How to pick a pattern)

Pattern	Signal	Do not use when
Strategy	One operation, several swappable rules	Only one rule exists and nothing hints at a second
Factory	Creation depends on a type value	Only one concrete class exists
Observer	Several listeners the source should not know	A single listener the source already references
State	Behaviour depends on status; many transitions	3–4 states handled by one guarded method
Singleton	Genuinely exactly one	You just want a global handle
Builder	Many optional fields; long constructor	A handful of required fields

Common mistakes (Ch 8 · Anti-patterns)

Anti-pattern	Fix
God Object (one class does everything)	Split by responsibility; push behaviour to the data owners
Golden Hammer (favourite pattern everywhere)	Start from the problem's signal
Pattern overuse ("patternitis") — factories for one class, single-implementation interfaces	YAGNI: add the pattern when the second variant arrives
Singletonitis (singletons as globals)	Dependency injection
Cargo-cult patterns (structure without intent)	Name the property the pattern gives you, or drop it

In the interview

Use **one or two** patterns in a 45-minute design, placed where their signal appears, and **name one you rejected** and why (Ch 8 · How to pick a pattern, self-check).

22.8 Concurrency cheat sheet

Source: Chapter 12 (concepts) and Chapter 13 (classic problems).

TL;DR (Ch 12)

Races come from **shared mutable state**. Make the critical section atomic, or remove sharing/mutability. Visibility needs **happens-before** edges. Prevent deadlock with **lock ordering** or **tryLock + back-off**. In LLD: name the contended resource, keep **check + claim atomic**, lock fine-grained and in order, and say what changes with many servers.

Primitive | use for | gotcha

Primitive	Use for	Gotcha	Where
`synchronized`	Making a multi-step update atomic (mutex + visibility)	Don't lock on a changing reference, a boxed value or a string literal; no slow I/O while holding it; readers must use the same lock	Ch 12 · Mutex and synchronized
`ReentrantLock`	Same as `synchronized` plus try/timed/fair locking and `Condition`s	Call `lock()` **just before** `try`, `unlock()` in `finally`; fairness is noticeably slower	Ch 12 · ReentrantLock and reentrancy
`tryLock(timeout)`	Give up if the lock is busy; deadlock prevention with back-off	Never `unlock()` when `tryLock` returned `false`; don't retry in a tight loop (livelock)	Ch 12 · Try-lock & timed locking
`ReentrantReadWriteLock`	Many readers, few writers	Cannot upgrade read → write (self-deadlock); don't modify under the read lock; writer starvation in non-fair mode	Ch 12 · Read-write locks
`StampedLock`	Read-mostly data with optimistic reads	**Not reentrant**; `validate` the stamp before using values read optimistically	Ch 12 · Read-write locks
`volatile`	Visibility of a flag or a published immutable snapshot	Gives visibility, **not** atomicity: `count++` is still a race	Ch 12 · Java Memory Model
`AtomicInteger` / `AtomicLong` / CAS	Lock-free counters and single-variable claims	Separate `get()` then `set()` is still a race — use `compareAndSet` / `updateAndGet`; no side effects in the lambda (may run several times); ABA → `AtomicStampedReference`	Ch 12 · CAS & atomics
`LongAdder`	Very hot counters (metrics)	Use instead of `AtomicLong` when contention is high	Ch 12 · CAS & atomics
`Semaphore(N)`	Limit concurrent access to N; signalling between threads	Release in `finally`; don't release when `acquire` never succeeded	Ch 12 · Semaphores
`wait/notifyAll`, `Condition`	Waiting for a condition on shared state	Always `while`, never `if` (spurious wakeups); hold the monitor; `notify()` can wake the wrong thread; state change without notify = lost wakeup	Ch 12 · Condition variables
`CountDownLatch`	Wait for N events once	One-shot; `countDown()` in `finally`, `await(timeout)`	Ch 12 · CountDownLatch and CyclicBarrier
`CyclicBarrier` / `Phaser`	N threads meet at each phase (reusable)	Fewer threads than parties → everyone waits forever	Ch 12 · CountDownLatch and CyclicBarrier
`ConcurrentHashMap`	Shared map	`containsKey` then `put` is a race — use `putIfAbsent` / `computeIfAbsent` / `merge`; no slow work or recursive updates inside `compute`; no `null` keys/values	Ch 12 · Concurrent collections
`CopyOnWriteArrayList`	Read-mostly lists (listener lists)	Every write copies the array	Ch 12 · Concurrent

Primitive	Use for	Gotcha	Where
			collections
`BlockingQueue` (bounded)	Handing work between threads (producer–consumer)	Unbounded queue → no back-pressure → memory blow-up; one poison pill per consumer	Ch 12 · Producer-consumer
`ExecutorService` / `ThreadPoolExecutor`	Running tasks on a bounded pool	Call `shutdown()` ; `newFixedThreadPool` has an **unbounded** queue; exceptions in `submit`ted tasks are lost unless you call `get()` ; tasks waiting on tasks in the same pool → thread-starvation deadlock	Ch 12 · Creating threads, Thread pool pattern
Virtual threads	Many blocking I/O tasks	Pinning (Java 21–23: blocking inside `synchronized`); heavy `ThreadLocal`s waste memory	Ch 12 · Creating threads, Immutability & confinement
`ForkJoinPool` / parallel streams	Splitting CPU work	Use a threshold: tiny tasks cost more to schedule than to run; subtasks must touch disjoint data	Ch 12 · Choosing the tool, Ch 13 · Parallel merge sort
`CompletableFuture`	Combining async results (`thenApply` , `thenCombine` , `allOf`)	Handle errors with `exceptionally` rather than calling `get()` at every step	Ch 12 · Creating threads
`ThreadLocal`	Per-thread state (thread confinement)	Clear it (`remove()`) in pooled threads, or values leak into the next task	Ch 12 · Immutability & confinement
Immutability / confinement	Avoiding synchronization entirely	`final List` is not immutable if the list is mutable and exposed	Ch 12 · Immutability & confinement

Which `java.util.concurrent` class for which job (Ch 12 · Choosing the tool)

Counter → `AtomicInteger` / `AtomicLong` (`LongAdder` if very hot) · stop a loop from another thread → `volatile boolean` or `interrupt()` · atomic multi-step update → `synchronized` / `ReentrantLock` · give up if busy → `tryLock(timeout)` · many readers, few writers → `ReentrantReadWriteLock` , `StampedLock` or an immutable snapshot in a `volatile` field · limit access to N → `Semaphore(N)` · wait for a condition → `Lock` + `Condition` (or `wait/notifyAll`) in a `while` loop · wait for N events once → `CountDownLatch` · N threads meet each phase → `CyclicBarrier` / `Phaser` · hand work between threads → bounded `BlockingQueue` · shared map → `ConcurrentHashMap` with `putIfAbsent` / `compute` / `merge` · read-mostly list → `CopyOnWriteArrayList` · run tasks → bounded `ExecutorService` , virtual threads for blocking I/O · split CPU work → `ForkJoinPool` /parallel streams · combine async results → `CompletableFuture` · avoid synchronization → immutability, thread confinement.

Deadlock checklist (Ch 12 · Deadlock, Livelock, Starvation)

A deadlock needs **all four Coffman conditions**; break any one and it cannot happen.

#	Condition	How to break it
1	Mutual exclusion	Lock-free/immutable data, shareable resources (read locks)
2	Hold and wait	Acquire everything at once, or `tryLock` and release everything on failure
3	No preemption	Timeouts (`tryLock(timeout)`), leases that expire
4	Circular wait	**Global lock ordering** (always lock the lower id first)

- ☐ Do I ever hold two locks? If yes, is there **one global order** (account id, seat id), with a tie-breaker for equal keys?

- ⬜ Do I call a **listener, callback or another object's synchronized method while holding my lock**? Move it outside the lock.
- ⬜ Do tasks in a bounded pool **wait for other tasks in the same pool**? (thread-starvation deadlock)
- ⬜ Are retries randomised (**jitter**) so two threads don't back off in lockstep? (livelock, Ch 12)
- ⬜ Could one thread wait forever while others progress? Consider fairness (Ch 12 · Starvation).
- ⬜ Detection: a thread dump (`jstack` / `jcmd`) prints "Found one Java-level deadlock"; `ThreadMXBean.findDeadlockedThreads()` .

Deadlock vs livelock vs starvation: nobody progresses and all wait / everybody runs but nobody progresses / some progress while one never does (Ch 12 · self-check).

The LLD concurrency recipe (Ch 12 · Concurrency in LLD interviews)

1. **Name the shared mutable state** (free spots, seat → owner, stock).
2. **State the invariant** ("a seat has at most one booking", "stock ≥ 0").
3. **Find the critical operation** — usually a check-then-act (*find free → assign*).
4. **Pick the smallest tool that makes it atomic:** one item → `queue.poll()` , `putIfAbsent` , `compareAndSet` ; several items all-or-nothing → per-item locks **in sorted order**; across servers → DB constraints/transactions, optimistic versioning, or a distributed lock.
5. **Choose granularity:** per floor / show / SKU, not one global lock.
6. **Handle time:** holds with expiry, timeouts, idempotent retries.

Classic problems (Ch 13 · summary)

Problem	Key primitive	Invariant to state
FooBar (Ch 13 · Print FooBar Alternately)	2 semaphores	Exactly one permit exists across both
Zero Even Odd (Ch 13 · Print Zero Even Odd)	3 semaphores	Zero dispatches; one permit circulates
FizzBuzz (Ch 13 · Fizz Buzz Multithreaded)	monitor + `notifyAll`	Predicates partition 1..n
H2O (Ch 13 · Building H2O)	2 semaphores + barrier	≤ 2 H and ≤ 1 O per round
Dining Philosophers (Ch 13)	ordered locks / n−1 semaphore	No circular wait
TTL cache (Ch 13 · Thread-safe cache with TTL)	`ConcurrentHashMap` + `compute` + conditional remove	Never return or keep expired data
Striped map (Ch 13 · Striped concurrent map)	lock array	Each bucket belongs to exactly one stripe
Blocking queue (Ch 13 · Thread-safe bounded Blocking Queue)	lock + 2 conditions	0 ≤ count ≤ capacity; wait in loops
Bloom filter (Ch 13 · Concurrent Bloom Filter)	atomic OR	Bits only go 0 → 1; no false negatives
Merge sort (Ch 13 · Parallel merge sort)	Fork/Join	Subtasks touch disjoint ranges
Token-bucket rate limiter (Ch 13)	lock or CAS on immutable state	Refill + take is one atomic step

Common mistakes (Chapter 12)

- A global `synchronized` on `ParkingLot.park()` — correct, but serialises every gate.
- Checking availability with one call and booking with another (check-then-act).
- Locking seats in the order the user clicked them (deadlock between two multi-seat bookings).
- Forgetting that in-memory locks don't work across multiple app servers.
- Swallowing `InterruptedException` — rethrow, or restore the flag with `Thread.currentThread().interrupt()` (Ch 12 · Thread states).

- Calling `run()` instead of `start()` (Ch 12 · Thread states).

22.9 Case-study map

Source: Chapters 14–21. "Concurrency concern" is the race each chapter guards against, and "atomic step" is how.

Everyday objects (Chapter 14)

Problem	Key entities	Patterns used	Concurrency concern	Where
Parking lot	`ParkingLot`, `ParkingFloor`, `ParkingSpot`, `Vehicle`, `Ticket`, `DisplayBoard`, `SpotType` / `VehicleType` enums	Strategy (pricing, spot-finding), Factory (electric spots), Observer (boards); Singleton rejected	Two gates take one spot → atomic `pollFirst` + CAS; per-ticket lock for payment; `putIfAbsent` for plates	Ch 14
Vending machine	`VendingMachine`, `Slot`, `Inventory`, `CashBox`, `Denomination`, `VendingState` + 5 states	**State** (the whole design)	`synchronized` machine; change-making by bounded coin-change DP, not greedy	Ch 14
Coffee machine	`CoffeeMachine`, `Recipe`, `RecipeBook`, `IngredientStore`, `OrderResult` (sealed)	Observer (low stock), recipes as data	Two orders take the last milk → one `ReentrantLock` around check + deduct; notify outside the lock	Ch 14 · Coffee Vending Machine
ATM	`Atm`, `AtmState` + 4 states, `Transaction` (sealed), `NoteHandler`, `CashDispenser`, `BankService`, `Account`	State (session), Command (transactions), Chain of Responsibility (note cassettes)	Overdraw from two ATMs → per-account lock; transfers lock in id order	Ch 14
Elevator	`Elevator`, `StopList` (two sorted sets), `Door`, `Dispatcher`, `DispatchStrategy`, `Building`	Strategy (dispatch) + modes; LOOK scheduling	Per-car lock, synchronized dispatcher; tick-driven `step()`	Ch 14 · Elevator System
Traffic signal	`Intersection`, `TrafficLight`, `Phase`, `SignalState` + 5 states, `TimingPolicy`, `SafetyMonitor`	State (timed), Strategy (timing), Observer (lights)	Synchronized tick; safe with a timer thread plus emergency threads	Ch 14 · Traffic Signal Control System
Shipping locker (Amazon Locker)	`LockerStation`, `Compartment`, `Booking`, `CompartmentSelector`, `PickupResult` (sealed), `LockerListener`, `LockerNetwork`	Strategy (door choice), Observer (SMS), enum state machines, sealed results	Two carriers, one door → station lock around choose-and-mark; lazy expiry of holds and pickup windows; hashed single-use pickup and return drop-off codes, keypad lockout	Ch 14 · Shipping Locker System

Booking & inventory (Chapters 15–16)

Problem	Key entities	Patterns used	Concurrency concern	Where
Library	`Book` (title) vs `BookItem` (copy), `Member`, `BookLending`, `LendingService`, `Catalog`	Strategy (fines), notifier interface; FIFO reservation queue per title	Last copy lent twice → lock member then copy (fixed order)	Ch 15 · Library
Movie tickets (BookMyShow)	`Show`, `ShowSeat`, `Booking`, `SeatLockProvider`, `BookingService`, `CancellationPolicy`	Strategy (refund rule), gateway behind an interface	Book `ShowSeat`s; all-or-nothing hold with TTL → pay → atomic commit; refund if hold lost	Ch 15 · BookMyShow
Concert tickets	`SeatedInventory`, `GaPool`, `SeatState` (sealed), `Hold`, `VirtualWaitingRoom`, `PurchaseLimiter`	Sealed seat states, waitlist	Extreme contention → CAS per seat, CAS counter for GA, FIFO waiting room, per-user limits, TTL holds with a reaper	Ch 15 · Concert tickets
Hotel	`Hotel`, `Room`, `RoomBooking`, `Invoice`, `RoomCharge`, `Payment` subclasses, `FrontDesk`	Template Method (`RoomCharge`)	Room × night reserved atomically in `synchronized Room.tryReserve`	Ch 15 · Hotel
Car rental	`Vehicle`, `DateRange`, `Reservation`, `Bill` / `BillItem`, `ReservationAddOn`, `PricingStrategy`	Template Method (add-ons), Strategy + Decorator (pricing), state machine	Availability is about **dates** → `Vehicle.tryBook(range)` synchronized	Ch 16 · Car Rental
Restaurant	`Menu` / `MenuSection` / `MenuItem`, `Table`, `Reservation`, `Order`, `Meal` / `MealItem`, `Kitchen`, `Check`	Observer (order ready), Strategy (chef assignment, payment), state machine	Two receptionists, one table → `Table.tryHold(time)` synchronized; listeners after the lock	Ch 16 · Restaurant
Airline	`Flight` vs `FlightInstance`, `Seat` vs `FlightSeat`, `FlightReservation`, `Itinerary`, `FareStrategy`	Strategy + Decorator (fares), Facade	Selling one seat twice → `FlightInstance.holdSeats` all-or-nothing under the instance lock	Ch 16 · Airline
Course registration	`Course` vs `Section`, `Student`, `RegistrationRule`s, `Registrar`	Chain of rules / Specification	Overselling a section → section lock inside a student lock (fixed order)	Ch 16 · Course Registration
Meeting scheduler	`Interval` (half-open), `Room` (owns a `TreeMap` calendar), `Meeting`, `Invitation`, `RoomSelectionStrategy`, `MeetingObserver`, `Recurrence`	Strategy (room ranking), Observer (invites, RSVPs)	Two organizers, overlapping times → check (`lowerEntry`) and insert under the room lock; rank then try each room; lock order meeting → room	Ch 16 · Meeting Scheduler

Games (Chapter 17)

Problem	Key entities	Patterns used	Concurrency concern	Where
Blackjack + deck	`Card`, `Deck<T>`, `Shoe`, `Hand`, `BlackjackScoring`, `PlayerStrategy`, `DealerPolicy`	Strategy (player, dealer); injected randomness	One table = one thread; shared balances need an atomic wallet service	Ch 17 · Blackjack
Chess	`Piece` (sealed) + subclasses, `SlidingPiece`, `Board`, `Move`, `Rules`, `Game`	Polymorphic pieces, Factory Method, Command + Memento (undo)	`synchronized` game so a move and a resignation don't interleave	Ch 17 · Chess
Tic-tac-toe	`Board` (row/col/diagonal counters), `Game`, `Player`, `MoveStrategy`	Strategy (players)	Many games in a `ConcurrentHashMap`; per-game lock	Ch 17 · Tic-Tac-Toe
Snake & ladder	`Board`, `Jump`, `Dice`, `Rules`, `OverflowRule`, `WinStrategy`, `Game`	Strategy (dice, overflow, win)	Each game its own `Random` and its own thread of turns	Ch 17 · Snake and Ladder
2048	`LineMerger`, `Board`, `TileSpawner`, `Game`, `Direction`	Strategy (spawner); one pure merge function	Single-threaded; `synchronized move` if UI and input threads both call it	Ch 17 · 2048
Jigsaw puzzle	`Side` (FLAT/TAB/BLANK + shape code), `Piece` (immutable), `Board` + `Placement` (piece, rotation), `PuzzleValidator`, `SolverStrategy`	Strategy (brute force vs edge-first with a side index); one placement check for players and solvers	Co-op players → `synchronized` game; solvers carry an attempt budget	Ch 17 · Jigsaw

Platforms (Chapters 18–19)

Problem	Key entities	Patterns used	Concurrency concern	Where
Stack Overflow	`Post` (sealed) → `Question` / `Answer` / `Comment`, `Member`, `Tag`, `Bounty`, `QuestionIndex`	Strategy (reputation), Observer (badges), Facade; inverted index	Two votes at once → `synchronized` post; `AtomicInteger` reputation	Ch 18
Cricinfo	`Match` → `Innings` → `Over` → `Ball`, `Team`, `Playing11`, `StatQuery`	Template Method (stat queries), Observer (live pages)	Two scorers → `synchronized` match; observers in `CopyOnWriteArrayList`	Ch 18
LinkedIn	`Member`, `Profile`, `SocialContent`, `ConnectionInvitation`, `ConnectionService`, `SearchIndex`, `FeedService`	State (invitation), Strategy (feed ranking, notifications), Facade	Symmetric connections → all graph changes synchronized on the service	Ch 18
Facebook	`Member`, `Audience` (sealed), `PrivacyList`, `FriendService`, `VisibilityService`, `NewsFeedService`	Strategy (feed ranker), Observer (notifications)	Accept can't leave a one-way friendship → synchronized service, concurrent adjacency sets	Ch 18
Amazon	`Product`, `ShoppingCart`, `Order`, `OrderLog`, `Payment`, `Shipment`, `Customer` → `Guest` / `Member`	Strategy (payment), Observer (notifications), order state machine, Facade	Last unit, two buyers → `synchronized Product.reserve`; all-or-nothing multi-line reservation	Ch 19 · Amazon
Stock brokerage + exchange	`Order` → Market/Limit/StopLoss/StopLimit, `OrderPart`, `StockLot`, `OrderBook`, `MatchingEngine`	Observer (fills), Template Method (statements), trade-price policy	Over-spending → reserved funds; one lock per symbol; price-then-time matching	Ch 19 · Stock Brokerage
Auction	`AuctionListing`, `Bid`, `BidResult`, `AuctionService`, `AuctionListener`	Observer, Specification (search), status state machine	Concurrent bids → all bid rules in one per-listing critical section; notify after unlocking	Ch 19 · Auction
Music streaming (Spotify)	`Song` / `Album` / `Artist`, `MusicLibrary`, `Playlist`, `MusicPlayer`, `PlaybackOrder`, `Recommender`	Strategy (play order, recommendations), Observer (plays), plan rules in an enum	Two devices press "next" → `synchronized` player methods	Ch 19 · Music Streaming
Food delivery (Swiggy) + FoodKart	`Restaurant`, `Cart`, `Order`, `DeliveryAgent`, `DeliveryAssignmentStrategy`, `RestaurantSelectionStrategy`	Strategy (assignment, selection, payment), Observer, order state machine	Kitchen capacity reserved atomically; riders claimed with CAS	Ch 19 · Food Delivery
Grocery store (POS)	`Register`, `Sale`, `LineItem`, `PricingEngine`, `ItemPromotion` / `BasketPromotion`, `TaxPolicy`, `Tender` (sealed), `Receipt`, `Inventory`	Strategy (offers, tax, payment gateway), Facade (register), Observer (stock), sale state machine	Registers share stock → atomic `merge`, may go negative, edge-triggered alerts; returns claimed with `compute`	Ch 19 · Grocery Store

Pattern-driven components (Chapter 20)

Problem	Key entities	Patterns used	Concurrency concern	Where
Splitwise	`Expense`, `SplitStrategy` (Equal/Exact/Percent), `Share`, `Ledger`, `DebtSimplifier`	Strategy (splits)	One signed number per pair, one mutator; `BigDecimal`, no lost cents	Ch 20
Notification service	`NotificationChannel` (abstract), channels, `SendRule`, `ChannelFactory`, `DeliveryListener`	Template Method (channels), Strategy (send rules), Observer, Simple Factory, Adapter	Rate-limit check-then-act gap; idempotency	Ch 20
In-memory cache / LRU	`Cache<K,V>`, `Node`, `DoublyLinkedList`, `EvictionPolicy`, `Cache.Builder`, `ShardedCache`	Strategy (eviction), Observer (removals), Builder	One lock (a `get` is a write in LRU); listeners outside the lock; shard to scale	Ch 20
Ride sharing (Uber)	`Rider`, `Driver`, `Trip`, `MatchingStrategy`, `PricingStrategy`, `SurgePricing`	Strategy ×3, Decorator (surge), trip state machine	Double-booking a driver → atomic claim (`AtomicBoolean`)	Ch 20 · Ride sharing
Logging framework	`Logger`, `LogManager`, `LogRecord`, `Appender`, `AsyncAppender`, `Formatter`	Chain of Responsibility (hierarchy), Singleton (`LogManager`), Strategy (formatters), Template Method (appender), Decorator (async)	Thread-safe sinks; never block on a full queue; never crash the app	Ch 20
Pub-sub	`Broker`, `Topic`, `Subscription`, `Subscriber`, `Message`	Observer through a broker	One dispatcher per subscription; ack after processing (at-least-once) → idempotent subscribers	Ch 20 · Pub-sub
Rate limiter	`Limiter` (TokenBucket, FixedWindow, SlidingWindowLog/Counter), `CompositeLimiter`, `LimiterFactory`	Strategy, Simple Factory, Composite, Facade	Atomic refill-check-take; one limiter per key via `computeIfAbsent`; Redis + Lua across servers	Ch 20

Machine-coding problems (Chapter 21)

Problem	Key entities	Patterns used	Concurrency concern	Where
Task management	`Task`, `TaskStatus`, `TaskService`, `TaskRepository`, `TaskFilter`	Specification (filters), Strategy (sorts), Observer, enum state machine	Per-task `synchronized`	Ch 21 · Task management
Digital wallet	`Wallet`, `Transaction`, `Money`, `Offer`, `WalletService`	Strategy (offers), idempotency key	Lock both wallets in a global (id) order; debit before credit; total conserved	Ch 21 · Digital wallet
Event calendar (Flipkart)	`TimeSlot` (half-open), `Event`, `EventStore`, `RepresentativeSelector`, `CalendarService`	Strategy (representative), plan-then-save	Atomic check-and-book	Ch 21
Coding Blox (Flipkart)	`Contest`, `Level`, `User`, `Question`, `SolveSimulator`, `CommandProcessor`	Command map, Strategy (solver, scoring)	Atomic rating update; don't run a contest twice	Ch 21
PropertyHunt (ClearTrip)	`Property`, `SearchCriteria` (+ Builder), `PropertyService`	Builder + predicate list, Strategy (sort)	Synchronized mark-sold	Ch 21
Ledger Co (Navi)	`Loan`, `LumpSum`, `Balance`, `InterestPolicy`, `LedgerService`	Strategy (interest), derived state, command loop	Synchronized lump sums	Ch 21
Inventory management	`Product`, `Warehouse`, `StockMovement`, `AllocationStrategy`, `ReorderListener`	Strategy (allocation), Observer (edge-triggered alerts)	Overselling → lock per SKU	Ch 21 · Inventory
Unix file search (`find`)	`Node` → `FileNode` / `DirectoryNode`, `FileFilter`, `AndFilter` / `OrFilter` / `NotFilter`, `SearchCriteria` (+ Builder), `FileSearcher`, `FindCommandParser`	Specification + Composite (filters and file tree), Builder, recursive-descent parser	Weakly consistent children map; read-only fork/join walk	Ch 21 · Unix File Search

TL;DR — what every case study teaches

- **Name the real unit of inventory** (copy, show seat, room-night, vehicle-day) (Ch 15 summary).
- **Every "if free then take" is one atomic step** — a monitor, a CAS, or a conditional `UPDATE` (Ch 15 summary).
- **Put rules on the object that owns the data**; the service only orders the steps (Ch 16, Ch 20 summaries).
- **Strategy** for the rule most likely to change; **Observer** for "tell someone when it happens"; **State** (or an enum state machine) for lifecycles.
- **Inject** time and randomness; money in `BigDecimal` or integer cents; call listeners **outside** locks.

22.10 Machine-coding and edge-case checklists

Source: Ch 21 · Machine-coding round (and Ch 1 · Interview code).

Machine-coding checklist

Before coding (0–20 min)

- [ ] Read the statement twice; mark **mandatory vs bonus**; ask clarifying questions.
- [ ] Put the statement's sample input into the driver as a `List<String>` of commands.

- ☐ Ask: "Where does it need to be thread-safe?" (wallets and inventory: always yes).
- ☐ Nouns → entities, verbs → service methods. Pick the 1–2 patterns the statement begs for ("can be sorted by X or Y" → Strategy).
- ☐ Create the six packages: **model, repository, service, strategy, exception, driver**.

While coding (20–75 min)

- ☐ Implement mandatory requirements **in the statement's order**; run the driver after each one (at least every 15 minutes).
- ☐ Entities guard their own invariants (e.g. `Task.changeStatus` rejects `TODO → DONE`).
- ☐ Services return data and throw domain exceptions; **no** `System.out` **in services**.
- ☐ Repository = interface + `ConcurrentHashMap` implementation; IDs from an `AtomicInteger`.
- ☐ Driver = `Map<String, Consumer<String[]>>` command processor; catch domain exceptions per command, **print and continue**.

Before you stop (75–90 min)

- ☐ Edge cases and validation done (see below); custom exceptions with readable messages.
- ☐ Thread-safety where money or stock is shared; say where the lock is and why.
- ☐ Bonus only after every mandatory feature runs.
- ☐ Clean names, delete dead code, run everything once more, prepare a 2-minute walkthrough.

Rubric the interviewer uses (Ch 21 · Machine-coding round)

(1) **Working, demo-able code** — thin end-to-end path in the first hour, never finish non-compiling; (2) **functional completeness** — mandatory first, in order; (3) **clean models** — enums for fixed sets, records for values, classes for changing things; (4) **separation of concerns** — Driver → Service → Repository; (5) **extensibility** — new rule = new class; (6) **edge cases** — custom exceptions, the driver keeps going; (7) **readability** — names from the statement's nouns and verbs; (8) **concurrency (when asked)** — say where the lock is, demo with a pool; (9) **tests (bonus)** — an `assert`-style self-check in `main` is enough.

Edge-case checklist (Ch 1 · Interview code, Ch 21 · Machine-coding round, case studies)

Category	Questions to ask your code
Input	`null`, empty, negative, zero, duplicate IDs, unknown IDs, invalid numbers
Ownership	Wrong user acting on someone else's object; a player moving out of turn
Capacity	Full lot, zero seats left, empty playlist, max loans reached, party larger than any table
State	Paying twice, exiting without entering, cancelling a delivered order, moving after game over, running a contest twice
Time	Exit before entry, holds/bookings expiring, part hours (round up?), an event exactly at the end instant (half-open intervals), late returns
Money	`BigDecimal` scale and rounding; no lost cents in splits; compare with `compareTo`, not `equals` (Ch 4 money note)
Concurrency	Two threads booking the same seat/spot/rider; multi-item bookings locking in sorted order; retries that must be idempotent
Failure	Payment declined after a hold; hold lost before payment → compensating refund; listener throws
Randomness	Seeded `Random` or a scripted strategy so the demo is repeatable

Common mistakes (Ch 21 · Machine-coding round)

Besides the Chapter 1 mistakes in 22.2: 20 minutes on a perfect diagram (enums/records first); one `Manager` with 15 maps (a service per area + repositories); `System.out` in services (the driver prints); `double` money (`BigDecimal` or whole units in a `long`); crashing on a bad command (catch per command); bonus before mandatory; unseeded `Random` (seed it or use a Strategy); `LocalDate.now()` in logic (inject a `Clock`).

---|---| | 20 minutes on a perfect class diagram | Enums/records first; the diagram lives in names | | One `Manager` with 15 maps | One service per area + repositories | | `System.out.println` inside a service | Services return data; the driver prints | | `double` for money | `BigDecimal` with a fixed scale, or whole units in a `long` | | Crashing on the first bad command | Catch domain exceptions in the driver loop | | Bonus before mandatory | Mandatory in order; bonus last | | `Random` without a seed | Seed it, or put randomness behind a Strategy | | `LocalDate.now()` in logic | Inject a `java.time.Clock` |

22.11 Glossary

Alphabetical. "Where" gives the chapter and topic where the term is explained in full. **Not repeated here** (see the tables above): the 23 GoF patterns and the extras (Simple Factory, Object Pool, Null Object, Dependency Injection, Repository, Specification, MVC, Event Sourcing) → 22.7; the five SOLID principles → 22.4; DRY, KISS, YAGNI, Law of Demeter, Tell Don't Ask, Program to an interface, Encapsulate what varies, Loose coupling, Hollywood Principle, Separation of Concerns, Least Astonishment, Fail fast, Design by Contract and the nine GRASP guidelines → 22.5; the 24 code smells → 22.6; the four pillars, overloading/overriding, interfaces vs abstract classes, the six relationships and the other one-line OOP facts → 22.3; the `java.util.concurrent` primitives, Coffman conditions, deadlock, livelock and starvation → 22.8.

Term	Plain meaning	Where
ABA problem	A value changes A → B → A, so a compare-and-swap cannot notice the change; fix with a version stamp (`AtomicStampedReference`)	Ch 12 · CAS & atomics
Activity diagram	UML flowchart of a process, with decisions, swimlanes and fork/join	Ch 3 · Activity diagrams
Actor	Someone or something that uses the system (customer, admin, the system itself)	Ch 1 · Requirements
Aggregate (DDD)	A cluster of objects changed as one unit; only its root is reachable from outside, and the root guards the invariants	Ch 5 · DDD
Anemic domain model	Classes with data but no behaviour; the logic lives in services that pull the data out	Ch 5 · Tell, Don't Ask
Anti-corruption layer	Translation layer at a bounded context's border that keeps a foreign or legacy model out of yours	Ch 5 · DDD
Anti-pattern	A common "solution" that reliably causes more problems than it solves	Ch 8 · Anti-patterns
Append-only log	Messages are only ever added at the end; readers track their own offset	Ch 20 · Pub-sub
Association class	A class that holds data about a link itself (e.g. `Enrollment{grade}`)	Ch 3 · Class diagrams
At-least-once delivery	A message may be delivered more than once, never lost; consumers must be idempotent	Ch 20 · Pub-sub
Atomic variable	`AtomicInteger` / `AtomicLong` / `AtomicReference` : lock-free single-variable updates using CAS	Ch 12 · CAS & atomics
Back-pressure	Slowing producers when consumers can't keep up, e.g. with a bounded queue	Ch 12 · Producer-consumer
`BigDecimal`	Exact decimal type used for money; build from strings, compare with `compareTo`	Ch 1 · Interview code, Ch 4 (money note)
Bloom filter	Probabilistic set: may say "maybe present" wrongly, never says "absent" wrongly	Ch 13 · Concurrent Bloom Filter
Bounded context	Boundary inside which one domain model and one vocabulary apply consistently (often one team or service)	Ch 5 · DDD
Boy Scout Rule	Leave the code a little cleaner than you found it	Ch 6 (intro)
Caretaker / Originator	Memento roles: the originator makes and restores snapshots; the caretaker stores them without looking inside	Ch 10 · Memento
CAS (compare-and-swap)	Atomic "set to new value only if it still equals the expected value"	Ch 12 · CAS & atomics
Characterisation test	A test that records what existing code does now, before you refactor it	Ch 7 · Full source code and output
Check-then-act	Read a condition, then act on it, as two steps: a race unless made atomic	Ch 12 · Race conditions, Concurrency in LLD interviews
Circular wait	T1 waits for T2's lock while T2 waits for T1's; one of the four deadlock conditions	Ch 12 · Deadlock
Class	A blueprint for objects: fields plus methods	Ch 2 · Classes & keywords
Class diagram	UML diagram of classes, fields, methods and relationships — "the answer" in an LLD round	Ch 3 · Class diagrams

Term	Plain meaning	Where
Class explosion	Subclasses multiplying across two dimensions (`WindowsButton` , `MacButton` , …)	Ch 9 · Bridge Pattern
Clean code	Code whose intent is obvious to the next reader and which is easy to change safely	Ch 6 (intro)
Code review	A peer reads a change before merge: correctness first, then design, readability, tests, security, performance	Ch 6 · Code review checklist
Code smell	A surface symptom that usually signals a deeper design problem; not a bug by itself	Ch 7 · Code smells
Command–Query Separation (CQS)	A method either changes state or returns a value, not both	Ch 6 · Functions
Compensation	Undoing a completed step with a new action (e.g. a refund) because money can't be rolled back	Ch 15 · BookMyShow
Composition root	The one place (usually `main`) where objects are created and wired together	Ch 6 · Systems
Concurrency vs parallelism	Many tasks in progress (structure) vs many executing at the same instant (hardware)	Ch 12 · Concurrency vs parallelism
Condition variable	Lets a thread wait until shared state meets a condition (`wait/notify` , `Condition`)	Ch 12 · Condition variables
Constructor chaining	One constructor calling another with `this(...)`	Ch 2 · Classes & keywords
Critical section	Code that touches shared state and must run in one thread at a time	Ch 12 · Race conditions
Dead-letter queue	Where messages go after retries are exhausted, for later inspection	Ch 20 · Pub-sub
Deep vs shallow copy	Copying nested objects too vs sharing them with the original	Ch 8 · Prototype
Derive, don't store	Compute counts and totals from the underlying records so they can't drift	Ch 18 · summary
Design pattern	A named, reusable solution to a recurring design problem	Ch 8 · What is a design pattern?
Director	Optional Builder role that runs a fixed sequence of building steps	Ch 8 · Builder
Domain event	Past-tense record of a business fact (`OrderPlaced`) that other parts of the system react to	Ch 5 · DDD
Domain-Driven Design (DDD)	Modelling software on the business: bounded contexts, entities, value objects, aggregates, domain events, repositories	Ch 5 · DDD
Double dispatch	Choosing a method by both the element's and the visitor's type (Visitor's `accept`)	Ch 11 · Visitor
Double-checked locking	Lazy singleton that checks, locks, checks again; the field must be `volatile`	Ch 8 · Singleton
Dynamic dispatch	The JVM picks an overridden method from the object's actual class at run time	Ch 2 · Polymorphism, Compile-time vs runtime polymorphism
Entity (DDD)	Object defined by an identity that survives changes; two entities are equal when their ids are equal	Ch 5 · DDD
Expression problem	Visitor makes adding operations cheap but adding types expensive (and the reverse for plain polymorphism)	Ch 11 · Visitor

Term	Plain meaning	Where
Fairness	Serving waiting threads roughly in order so none starves; costs throughput	Ch 12 · Starvation
F.I.R.S.T.	Clean tests are Fast, Independent, Repeatable, Self-validating, Timely	Ch 6 · Unit tests
Flag argument	A boolean parameter that makes a function do two things; split it instead	Ch 6 · Functions
Fork/Join	Split a task recursively into subtasks run in parallel, then combine results	Ch 12 · Choosing the tool, Ch 13 · Parallel merge sort
Fragile base class	A change in a parent class silently breaks subclasses	Ch 5 · Composition over inheritance
Functional interface	An interface with one abstract method, usable with lambdas	Ch 2 · Interfaces
Functional requirement	A behaviour the system must have ("a customer takes a ticket")	Ch 1 · Requirements
Glob	Filename wildcard pattern: `*` any characters, `?` one character, `[0-9]` one character from a set	Ch 21 · Unix File Search
God class / God Object	One class that knows and does almost everything	Ch 8 · Anti-patterns, Ch 7 · Code smells
Golden Hammer	Applying a favourite tool or pattern to every problem	Ch 8 · Anti-patterns
GRASP	Larman's nine guidelines for assigning responsibilities to classes	Ch 5 · GRASP
Guard clause	An early check-and-return/throw at the top of a method instead of nested `if` s	Ch 7 · Refactoring catalogue, Ch 5 · Fail fast
Happens-before	Java Memory Model rule guaranteeing one action's effects are visible to another	Ch 12 · Java Memory Model
HLD vs LLD	Many services and machines vs classes inside one application	Ch 1 · The interview
Hold with TTL	A temporary claim (e.g. seats during payment) that expires automatically	Ch 15 · BookMyShow
Idempotency	Running an operation twice has the same effect as once; an idempotency key enforces it	Ch 21 · Digital wallet, Ch 20 · Pub-sub
Immutability	Objects that never change after creation; safe to share between threads	Ch 2 · Immutability & records, Ch 5 · Immutability
Intrinsic vs extrinsic state	Flyweight's shared, immutable part vs the per-use part passed in by the caller	Ch 9 · Flyweight Pattern
Invariant	A fact about an object that must always be true ("balance is never negative")	Ch 5 · Design by Contract
Inversion of Control (IoC)	The framework controls the flow and calls your code; DI is one way to achieve it	Ch 5 · Hollywood Principle, Ch 4 · SOLID Q&A
Inverted index	A map from each word/tag to the list of items containing it, for search	Ch 18 · Stack Overflow
Invoker / Receiver	Command roles: the invoker triggers commands; the receiver does the real work	Ch 10 · Command
Java Memory Model (JMM)	Rules for when one thread's writes become visible to another	Ch 12 · Java Memory Model

Term	Plain meaning	Where
Lock ordering	Always acquire locks in one global order to prevent circular wait	Ch 12 · Deadlock
Lock striping	Guarding a map with an array of locks, each covering some buckets	Ch 13 · Striped concurrent map
LOOK scheduling	Elevator keeps moving in one direction while stops remain ahead, then reverses	Ch 14 · Elevator System
Lost wakeup	A signal sent before the waiter started waiting, so the waiter sleeps forever	Ch 12 · Condition variables
LRU cache	Evicts the least recently used entry; hash map + doubly linked list gives O(1)	Ch 20 · Cache / LRU
Machine-coding round	LLD round where working, runnable in-memory code is the answer (90–120 min)	Ch 1 · The interview, Ch 21 · Machine-coding round
Matching engine	Per-symbol order book matching buy and sell orders by best price, then oldest first	Ch 19 · Stock Brokerage
Multiplicity	How many objects take part at each end of a relationship (`1` , `0..1` , `*` , `1..*`)	Ch 3 · Association
Mutex	A lock that lets only one thread into a critical section	Ch 12 · Mutex and synchronized
Non-functional requirement	A quality: thread safety, extensibility, in-memory storage, latency	Ch 1 · Requirements
Object	One instance of a class with real values	Ch 2 · Classes & keywords
Optimistic locking	Read with a version; the write succeeds only if the version is unchanged	Ch 12 · Read-write locks, Concurrency in LLD interviews
`Optional`	A container that may be empty; return it when "not found" is normal	Ch 1 · Interview code
Out-of-scope list	Written list of what you will not build; shrinks the problem	Ch 1 · Requirements
Pinning	A virtual thread stuck on its carrier thread while it blocks	Ch 12 · Creating threads
PLU code	Four-digit price look-up code for loose produce, keyed at the till together with the scale weight	Ch 19 · Grocery Store
Poison pill	A special message telling a consumer to stop; one per consumer	Ch 12 · Producer-consumer
Producer–consumer	Threads hand work through a (bounded) queue	Ch 12 · Producer-consumer
Pub-sub	Observer through a broker: publishers and subscribers don't know each other	Ch 20 · Pub-sub
Qualified association	An association looked up by a key (`Library [isbn] → Book` , i.e. a `Map`)	Ch 3 · Class diagrams
Race condition	The result depends on the timing of threads touching shared mutable state	Ch 12 · Race conditions
Rate limiter	Caps how many requests a key may make per time window (token bucket, windows)	Ch 13 · Rate limiter, Ch 20 · Rate Limiter
Record	Compact Java class for immutable data; only shallowly immutable	Ch 2 · Immutability & records
Reentrancy	A thread may re-acquire a lock it already holds	Ch 12 · ReentrantLock and reentrancy

Term	Plain meaning	Where
Refactoring	Changing internal structure without changing observable behaviour	Ch 7 · Refactoring fundamentals
Requirement	One thing the system must do, written so you can check it	Ch 1 · Requirements
Role interface	A small interface serving one kind of caller (the ISP fix)	Ch 4 · Interface Segregation Principle
Rule of three	Refactor to an abstraction at the second or third case, not the first	Ch 4 · Over-engineering
Sealed type	A type whose permitted subtypes are fixed, enabling exhaustive `switch`	Ch 11 · Visitor
Sequence diagram	UML diagram of messages between objects over time	Ch 3 · Sequence diagrams
Service Locator	Anti-pattern where a class asks a registry for its dependencies, hiding them	Ch 11 · Beyond GoF
Sharding	Splitting data across independent parts (each with its own lock) to scale	Ch 20 · Cache / LRU
Signaling	One thread tells another that something has happened (semaphore, latch)	Ch 12 · Signaling pattern
Sliding window	Rate-limit algorithm counting requests in a moving time window (log or counter)	Ch 20 · Rate Limiter
SOLID	SRP, OCP, LSP, ISP, DIP: five guidelines that keep one change from spreading	Ch 4 · What SOLID is
Split tender	Paying one bill with several payment methods (e.g. part UPI, part cash)	Ch 19 · Grocery Store
Spurious wakeup	A waiting thread wakes with no signal; why `wait()` sits in a `while` loop	Ch 12 · Condition variables
State machine diagram	UML diagram of one object's states and transitions (`event [guard] / action`)	Ch 3 · State machine diagrams
`static`	Belongs to the class, one per class, not to each object	Ch 2 · Classes & keywords
Test double	Anything standing in for a real dependency in a test: dummy, stub, spy, mock, fake	Ch 6 · Testing strategies
Test-Driven Development (TDD)	Failing test → make it pass → refactor, in very short cycles	Ch 6 · Unit tests
Thread	An independent path of execution inside a process, sharing its memory	Ch 12 · Processes vs threads
Thread confinement	Keeping data inside one thread so it never needs locking	Ch 12 · Immutability & confinement
Thread pool	A fixed set of reusable threads fed from a task queue	Ch 12 · Thread pool pattern
Thread states	NEW, RUNNABLE, BLOCKED, WAITING, TIMED_WAITING, TERMINATED	Ch 12 · Thread states
Thread-starvation deadlock	Tasks in a bounded pool wait for tasks queued behind them in the same pool	Ch 12 · Deadlock, Thread pool pattern
Token bucket	Rate-limit algorithm: tokens refill at a fixed rate; each request takes one; allows bursts	Ch 13 · Rate limiter
Two hats	While coding you are either adding function or refactoring, never both at once	Ch 7 · Refactoring fundamentals

Term	Plain meaning	Where
Ubiquitous language	One vocabulary shared by domain experts and the code (`place()` , not `setStatus(2)`)	Ch 5 · DDD
UML	Unified Modeling Language: 14 diagram types, 7 structural and 7 behavioural	Ch 3 · What UML is
Use case	One thing an actor can do with the system	Ch 1 · Requirements, Ch 3 · Use-case diagrams
Use-case diagram	UML diagram of actors, use cases and the system boundary («include», «extend»)	Ch 3 · Use-case diagrams
Value object	A small immutable object defined by its values (money, date range), usually a record	Ch 2 · Immutability & records, Ch 5 · DRY, Ch 5 · DDD
Virtual waiting room	FIFO queue that admits buyers gradually during a ticket rush	Ch 15 · Concert tickets

---

Self-check

- **Q:** Name the eleven framework steps in order. **A:** Clarify, requirements, use cases, entities, relationships, class diagram, key flows, patterns & SOLID, code, edge cases & concurrency, extensions.
- **Q:** Where does the diamond go, and what does filled mean? **A:** At the whole; filled = composition (the part dies with the whole).
- **Q:** A `switch (vehicleType)` computes fees. Which principle, smell and pattern? **A:** OCP; Repeated Switches (if it repeats); Strategy.
- **Q:** Decorator vs Proxy vs Chain of Responsibility? **A:** Decorator adds behaviour and every layer runs; Proxy controls access; in CoR one handler can stop the flow.
- **Q:** State vs Strategy? **A:** States know each other and advance themselves; strategies are chosen from outside.
- **Q:** Four Coffman conditions, and the usual fix? **A:** Mutual exclusion, hold and wait, no preemption, circular wait; global lock ordering.
- **Q:** Why is `volatile` not enough for `count++` ? **A:** Visibility, not atomicity; use `AtomicInteger` or a lock.
- **Q:** The six machine-coding packages? **A:** model, repository, service, strategy, exception, driver.
- **Q:** What makes every booking problem safe? **A:** Naming the real unit of inventory and making "if free then take" one atomic step (plus holds that expire).

TL;DR

- **Framework:** clarify → requirements → use cases → entities → relationships → diagram → flows → patterns/SOLID → code → edge cases/concurrency → extensions.
- **Classes:** different data → field; different behaviour → class; one sentence per class; composition over inheritance.
- **SOLID** diagnoses change pain; always state the cost of the fix, and know when not to apply it.
- **Patterns** come from signals: Strategy (rules vary), Factory (type value → class), Observer (many listeners), State (lifecycle), Singleton (exactly one), Builder (many optional fields). Use one or two, name a rejected one.
- **Concurrency:** find the shared state, state the invariant, make check-then-act atomic with the smallest tool, lock in a fixed order, notify outside locks.
- **Machine coding:** run early and often, mandatory first, services return data, the driver prints and keeps going.

Chapter 23 — Interview Question Bank

How to use this bank

This chapter collects the low-level design (LLD) and machine-coding questions that come up again and again in interviews, worded the way they usually appear on LeetCode Discuss interview-experience posts and in the popular public question lists (see the sources at the end). Use it as a checklist so that nothing surprises you.

- ⋆ marks a question that is asked very often. Do those first.
- **Part A** is the "Design X" prompts (parking lot, BookMyShow, Splitwise …). When this book has a full case study, the entry is a short reminder of what to say plus a link to the full design. When it doesn't, the entry is a complete outline you can build from.
- **Part B** is the concept questions (OOP, UML, SOLID, patterns, concurrency, Java specifics). Each has a short answer you can say out loud and a link to the full explanation.
- **Part C** is a minute-by-minute plan for a 60–90 minute machine-coding round.

How to drill: read the question, cover the answer, say your answer out loud (or sketch the classes), then compare. Mark the ones you stumbled on and come back to them the next day.

Part A — "Design X" (LLD and machine coding) questions

These are the "Design X" prompts that show up again and again in object-oriented design (OOD) rounds and in machine-coding rounds. An OOD round is a whiteboard-style talk about classes; a machine-coding round is 90–150 minutes of writing working code. ⋆ marks prompts asked very often, and every machine-coding favourite also gets a ⋆. When this book has a full case study, the entry here is a short reminder of what to say, followed by a link. When the book has no case study, the entry is a complete outline you can build from.

Everyday objects

Q1. ⋆ Design a Parking Lot

Asked at: Flipkart (SDE2 LLD round)

- Entities: `ParkingLot` → `Level` → `ParkingSpot` (sized: small, compact, large), `Vehicle`, `Ticket`, `EntryGate` / `ExitGate`, `FeeCalculator`.
- Patterns: Strategy for spot allocation and for pricing; Singleton (or one injected instance) for the lot; Factory for vehicles.
- What interviewers probe: two cars racing for the last free spot (claim the spot atomically), hourly vs flat pricing, and adding EV charging spots without editing old code. → *Full design:* Ch 14 — 14.1 Design a Parking Lot

Q2. ⋆ Design an Elevator / Lift System

- Entities: `ElevatorController`, `Elevator` (current floor, direction, door state), `Request` (hall call or car call), `Dispatcher`.
- Patterns: State (moving up, moving down, idle, doors open), Strategy for dispatching (nearest car, SCAN / LOOK).
- What interviewers probe: the scheduling algorithm (why plain first-come-first-served makes cars zig-zag), keeping requests consistent when many buttons are pressed at once, and overload or maintenance modes. → *Full design:* Ch 14 — 14.5 Design an Elevator System

Q3. ⋆ Design a Vending Machine

- Entities: `VendingMachine`, `Inventory`, `Product`, `Coin` / `Note`, `ChangeDispenser`.
- Pattern: State (idle, has money, dispensing, sold out) is the whole point of this question.
- What interviewers probe: making change with a limited coin stock, refunding on cancel, and what happens when a product sells out while money is inserted. → *Full design:* Ch 14 — 14.2 Design a Vending Machine

Q4. Design a Coffee Vending Machine

- Entities: `CoffeeMachine`, `Recipe` (ingredient → quantity), `IngredientStore`, `Beverage`, add-ons.
- Patterns: Decorator for add-ons (extra shot, milk), Factory for drinks, State for the machine.
- What interviewers probe: checking and deducting every ingredient as one atomic step, and refill alerts (Observer). → *Full design:* Ch 14 — 14.3 Design a Coffee Vending Machine

Q5. ⋆ Design an ATM

- Entities: `ATM`, `Card`, `Account`, `BankService`, `CashDispenser`, `Transaction` (withdraw, deposit, balance enquiry).
- Patterns: State (no card, card inserted, PIN verified), Chain of Responsibility for dispensing notes by denomination.
- What interviewers probe: what to do if cash is debited but not dispensed (reversal), PIN retry limits, and two ATMs withdrawing from one account at the same time. → *Full design:* Ch 14 — 14.4 Design an ATM

Q6. Design a Traffic Signal Control System

- Entities: `Intersection`, `TrafficLight`, `SignalPhase`, `SignalController`, `Timer`.
- Patterns: State for the light cycle, Strategy for timing plans (fixed, rush hour, sensor-driven), Observer for emergency vehicles.
- What interviewers probe: never letting conflicting directions be green together, and handling emergency override safely. → *Full design:* Ch 14 — 14.6 Design a Traffic Signal Control System

Q7. Design an Amazon Locker / shipping locker

- Entities: `LockerLocation`, `Locker` (by size), `Package`, `PickupCode`, `LockerService`.
- Patterns: Strategy for choosing a locker size, State for the locker (free, reserved, occupied).
- What interviewers probe: code expiry and returning the package to the sender, and two deliveries claiming the same locker. → *Full design:* Ch 14 — 14.7 Design a Shipping Locker System

Q8. Design a toll booth / parking-fee calculator (pricing strategies)

- This is the pricing part of the parking lot pulled out on its own: a `FeeCalculator` interface with hourly, flat, slab and weekend implementations, chosen by vehicle type and time.
- Patterns: Strategy (one class per pricing rule), optionally Decorator for surcharges and discounts layered on top.
- What interviewers probe: time zones and midnight boundaries, rounding (charge per started hour?), and keeping money as integer paise or `BigDecimal`, never `double`. → *Full design:* Ch 14 — 14.1 Design a Parking Lot (pricing section)

Q9. Design a home automation system

Requirements. A user owns devices (lights, fans, thermostats, door locks, sensors) made by different vendors. The user can switch a device on or off, set values (brightness, temperature), group actions into a "scene" ("Good night" turns off lights and locks doors), undo the last action, and create rules ("if motion after 23:00, turn on the hall light"). Schedules ("at 07:00") are nice to have.

Main classes.

- `Device` (interface): `id()`, `status()`. Subtypes `Light`, `Thermostat`, `Lock`, `MotionSensor`.
- `VendorAdapter`: wraps each vendor's protocol (Zigbee, Wi-Fi API) behind the `Device` interface.
- `DeviceCommand`: one action with `execute()` and `undo()`, for example `SetBrightness(light, 40)` remembers the old brightness.
- `Scene`: a list of commands run as one; it is itself a command.
- `Hub`: registry of devices, runs commands, keeps an undo stack.
- `Rule`: a trigger (sensor event, time) plus a condition plus a command.
- `EventBus`: sensors publish events; rules subscribe.

Patterns. Adapter (vendor protocols), Command (actions with undo), Composite (a scene is a command made of commands), Observer (sensor events reach rules), Strategy for the condition check.

Concurrency and edge cases. Send commands to one device in order: give each device its own single-threaded executor (or a lock), while different devices run in parallel on virtual threads. A device may be offline: queue the command with a timeout and report failure instead of hanging. Avoid rule loops (rule A turns on a light, which fires rule B, which turns it off) by tracking which rule caused an event and limiting the depth.

Extension question. "Support user roles: a guest can unlock the front door only between 10:00 and 18:00." Add an `AccessPolicy` checked by the `Hub` before running any command (a Proxy around devices works well).

Q10. Design a Jackpot / slot machine

Requirements. A player inserts credit, places a bet, and pulls the lever. Three (or five) reels spin and stop on symbols. A pay table decides the win. The machine tracks credit, pays out, and refuses to play when it is out of service or cannot cover the payout. Optional: a progressive jackpot shared by many machines.

Main classes.

- `SlotMachine`: holds the credit balance and the current state, and exposes `insert`, `bet`, `spin` and `cashOut`.
- `MachineState` (sealed interface): `Idle`, `HasCredit`, `Spinning`, `OutOfService`. Each state allows only its own actions.
- `Reel`: a list of symbols with weights, so a cherry can be more common than a seven.
- `RandomSource`: an interface wrapping `SecureRandom`. Tests swap in a fixed sequence, so results are repeatable.
- `PayTable`: maps a symbol combination to a multiplier. A Strategy, so different machines can pay differently.
- `JackpotPool`: one shared progressive pool. Each bet adds a small cut.

Patterns. State (the machine lifecycle), Strategy (pay rules, random source), Singleton-like shared `JackpotPool` (inject one instance rather than a static).

Concurrency and edge cases. Two machines can hit the jackpot at the same moment. Take the pool with one atomic step, `AtomicLong.getAndSet(seed)`, so only one machine wins the full amount. Use whole numbers for money. Reject a bet bigger than the credit. Go out of service if the machine's cash cannot cover the largest possible win. Log every spin (seed, result) for audits.

Extension question. "Add a bonus round after three scatter symbols." Add a `BonusRound` state that the machine moves into after the normal payout. No existing state class needs to change.

Q11. Design a music playlist

- Entities: `Song`, `Playlist` (ordered list), `Player` (current position, play mode), `User`.
- Patterns: Iterator (sequential, shuffle, repeat as different iterators) or Strategy for play order.
- What interviewers probe: shuffle without repeats until all songs played, and editing the playlist while it is playing. → *Full design:* Ch 1 — 1.3 A complete answer in five minutes

Booking and reservations

Q12. ⋆ Design BookMyShow / a movie ticket booking system

Asked at: Flipkart

- Entities: `Cinema` → `Screen` → `Seat`, `Movie`, `Show`, `ShowSeat` (seat state per show), `Booking`, `Payment`.
- Patterns: State for a seat (available, held, booked), Strategy for pricing, Observer for confirmations.
- What interviewers probe: two users picking the same seat (hold with expiry, then confirm), releasing holds when payment fails or times out. → *Full design:* Ch 15 — 15.2 Design a Movie Ticket Booking System

Q13. ⋆ Design a Library Management System

- Entities: `Book` (title info) vs `BookItem` (physical copy), `Member`, `Loan`, `Reservation`, `Fine`, `Catalog` (search).
- Patterns: Strategy for fine rules, Observer for "your reserved book is back".
- What interviewers probe: the book vs copy split, borrowing limits, and the reservation queue. → *Full design:* Ch 15 — 15.1 Design a Library Management System

Q14. ⋆ Design a Hotel Management System

- Entities: `Hotel`, `Room` (type, rate), `Guest`, `Reservation` (date range), `Invoice`, `HousekeepingTask`.
- Patterns: State for reservation and room status, Strategy for room pricing.
- What interviewers probe: date-range overlap checks, overbooking rules, and late check-out charges. → *Full design:* Ch 15 — 15.4 Design a Hotel Management System

Q15. Design a concert ticket booking system

- Like BookMyShow, but one big event with a sudden rush: seat sections, general admission, and a waiting room.
- Patterns: State for seats, plus a queue in front of booking.
- What interviewers probe: thousands of buyers at the same second, hold timeouts, and per-user ticket limits. → *Full design:* Ch 15 — 15.3 Design a Concert Ticket Booking System

Q16. ⋆ Design a car / vehicle rental system (Zoomcar)

Asked at: Flipkart (SDE2, 2024)

- Entities: `Branch`, `Vehicle`, `Reservation` (pickup and drop time and branch), `Customer`, `PricingPolicy`, `Invoice`.
- Patterns: Strategy for pricing and for vehicle selection, State for reservation lifecycle.
- What interviewers probe: availability across overlapping time windows, one-way rentals, and damage or late-return charges. → *Full design:* Ch 16 — Design a Car Rental System

Q17. Design a restaurant management system (table booking / waiting queue)

- Entities: `Table`, `Reservation`, `WaitlistEntry`, `Order`, `MenuItem`, `Bill`, `Kitchen`.
- Patterns: Command for orders sent to the kitchen, Observer for "your table is ready", State for orders.
- What interviewers probe: fitting party sizes to tables, walk-ins vs reservations, and splitting the bill. → *Full design:* Ch 16 — Design a Restaurant Management System

Q18. Design an airline management / flight booking system

- Entities: `Flight` (route) vs `FlightInstance` (a dated departure), `Aircraft`, `Seat`, `Itinerary`, `Passenger`, `Booking`.
- Patterns: State for booking, Strategy for fare rules.
- What interviewers probe: the route vs instance split, multi-leg trips booked all-or-nothing, and seat holds. → *Full design:* Ch 16 — Design an Airline Management System

Q19. Design a course registration system

- Entities: `Course`, `Section` (time slot, capacity), `Student`, `Enrollment`, `Prerequisite`, `Waitlist`.
- Patterns: Specification or Chain of Responsibility for eligibility checks (prerequisites, clashes, credit limit).
- What interviewers probe: the last seat taken by two students at once, and automatic waitlist promotion. → *Full design:* Ch 16 — Design a Course Registration System

Q20. Design a learning management system / online exam portal

- Start from course registration and add `Exam`, `Question` (single choice, multiple choice, text), `Attempt`, `Grader`.
- Patterns: Strategy for grading each question type, State for an attempt (started, submitted, graded), and a timer that auto-submits.
- What interviewers probe: saving answers often so a dropped connection loses nothing, and randomising question order per student. → *Related design:* Ch 16 — Design a Course Registration System

Q21. ⋆ Design a Meeting Scheduler / meeting room booking

- Entities: `Interval` (half-open), `Room` (owns a sorted calendar), `Meeting`, `Invitation` (RSVP), `MeetingScheduler`, `RoomSelectionStrategy`, `MeetingObserver`.
- Patterns: Strategy for room ranking, Observer for invitations and RSVPs.
- What interviewers probe: overlapping bookings of one room (check and insert under the room lock), room clash blocks while a busy person only warns, reschedule without losing the slot, and "first time everyone is free" (merge busy intervals, scan the gaps). → *Full design:* Ch 16 — Design a Meeting Scheduler

Q22. ⋆ Design a doctor appointment booking system

Asked at: Flipkart (SDE2, 2021)

Requirements. Doctors register with a speciality (cardiologist, dermatologist). Each doctor declares available one-hour slots for the day. Patients search by speciality and see free slots, sorted by start time by default. A patient books, cancels, and views their bookings. A patient may not hold two appointments at the same time. When a slot is full, the patient can join its waitlist, and a cancellation promotes the first person waiting.

Main classes.

- `Doctor` (id, name, speciality, slots) and `Patient` (id, name, appointments).
- `Slot`: doctor, start time, status (`FREE`, `BOOKED`), booked patient, and a waitlist (a `Deque<Patient>`).
- `Appointment`: a record (id, patient, slot).
- `BookingService`: `book`, `cancel` and `joinWaitlist`. It enforces "no double booking for the patient".
- `SearchService` with a `SlotRankingStrategy`. Order by start time by default, or by doctor rating later.
- Repositories: in-memory maps, as the machine-coding round expects.

Patterns. Strategy (ranking), Observer (tell a waitlisted patient they were promoted), a simple State on `Slot`.

Concurrency and edge cases. Two patients may book the last slot together. Guard the slot status with a per-slot lock, or a compare-and-set on an `AtomicReference<SlotStatus>`. Check the patient's own clash under the patient's lock. Lock patient then slot, always in that order, so no deadlock can form. Validate that slots are whole hours inside working hours. Cancelling an already-cancelled booking is a no-op.

Extension question. "Show the trending doctor: the one with the most bookings today." Keep a counter per doctor that `book` updates, and read it through a max-heap or a sorted view.

Q192. Meeting scheduler: why can't a concurrent map keyed by start time stop double-booking, and how do you find a clash fast? Overlapping meetings usually start at different times (10:00–11:00 and 10:30–11:15), so their keys differ and `putIfAbsent` lets both in. Keep each room's bookings in a `TreeMap` keyed by start; because stored bookings never overlap, only the one returned by `lowerEntry(newEnd)` can clash, an O(log n) check. Do that check and the insert under the room's lock. → *More:* Ch 16 — Design a Meeting Scheduler

Games

Q23. ⋆ Design Snake and Ladder

- Entities: `Board` (size, snakes and ladders as jumps), `Player`, `Dice`, `Game` (turn order queue).
- Patterns: Strategy for dice (one die, two dice, a loaded die in tests), a simple game loop.
- What interviewers probe: input-driven setup in machine coding, exact-roll-to-win rules, and chained jumps. → *Full design:* Ch 17 — Design Snake and Ladder

Q24. ⋆ Design Chess

- Entities: `Board`, `Square`, `Piece` subclasses, `Move`, `Player`, `Game` (status, history).
- Patterns: Strategy (or polymorphism) for each piece's moves, Command for moves with undo, State for game status.
- What interviewers probe: check / checkmate detection (simulate the move, then test the king), castling and en passant, and why moves should be objects. → *Full design:* Ch 17 — Design Chess

Q25. ⋆ Design a 2048 game

Asked at: Flipkart (2021)

- Entities: `Board` (grid), `Tile`, `Game` (score, state), `MoveDirection`, `TileSpawner`.
- Pattern: write one "slide and merge a row left" routine and rotate the grid for the other three directions.
- What interviewers probe: each tile merging only once per move, spawning only after a real change, and detecting game over. → *Full design:* Ch 17 — Design 2048

Q26. ⋆ Design Tic-Tac-Toe

- Entities: `Board`, `Player`, `Game`, `WinChecker`.
- Patterns: Strategy for players (human, AI), and per-row / column / diagonal counters so a win check is O(1).
- What interviewers probe: N×N boards, the O(1) win check, and keeping input handling out of the game rules. → *Full design:* Ch 17 — Design Tic-Tac-Toe

Q27. Design a deck of cards / Blackjack

- Entities: `Card` (rank, suit enums), `Deck`, `Shoe`, `Hand`, `Player`, `Dealer`, `Game`.
- Patterns: Template Method or Strategy for the dealer's rule, and a generic deck reused by other card games.

- What interviewers probe: ace counted as 1 or 11, and making a deck that works for any card game. → *Full design:* Ch 17 — Design Blackjack and a Deck of Cards

Q28. Design Minesweeper

Requirements. A rows × cols grid hides a number of mines. The player reveals a cell or toggles a flag. Revealing a mine loses the game. Revealing a number shows how many of the eight neighbours are mines. Revealing a zero opens all connected empty cells and their numbered borders. The player wins when every non-mine cell is revealed. The first click must never be a mine.

Main classes.

- `Cell`: `mine` (boolean), `adjacentMines`, and `state` (`HIDDEN`, `REVEALED`, `FLAGGED`).
- `Board`: the grid, `neighbours(r, c)`, `revealedCount`, `mineCount`.
- `MinePlacer` (interface): random placement with a seed. Tests use a fixed layout. Mines are placed *after* the first click, avoiding that cell (and its neighbours, if you like).
- `Game`: `reveal(r, c)`, `toggleFlag(r, c)`, and `status` (`PLAYING`, `WON`, `LOST`).
- `Renderer`: prints the board. Keep it separate from the rules.

Patterns. Strategy (mine placement), State for game status (no moves after a win or loss), and Observer if a UI must redraw.

Algorithm. For reveal, use a breadth-first search with a queue, not recursion. A recursive flood fill on a 1000×1000 empty board can overflow the stack. Mark each cell revealed when you add it to the queue, so no cell is processed twice. Flagged cells are not opened by the flood. Win check: `revealedCount == rows * cols - mineCount`, an O(1) test.

Edge cases. Ignore clicks outside the board, on revealed cells, or after the game ended. Reject more mines than cells minus one.

Extension question. "Add chording: clicking a number whose flags are all placed opens its other neighbours." Write it as one more method that reuses `reveal`, so no existing rule changes.

Q29. Design a Bowling Alley scoring system

- Entities: `Game`, `Frame` (rolls), `TenthFrame` (up to three rolls), `Player`, `Lane`.
- Key logic: a strike adds the next two rolls, a spare adds the next one. Score from the stored list of rolls instead of keeping running totals, so bonuses resolve naturally.
- What interviewers probe: the tenth frame, partial scores while bonuses are pending, and validating pin counts per frame.
- The scoring core for a finished game (validate rolls first: 0–10 each, at most 10 per frame outside the tenth):

```
static int score(List<Integer> rolls) {
    int total = 0, i = 0;
    for (int frame = 1; frame <= 10; frame++) {
        if (rolls.get(i) == 10) {                                  // strike
            total += 10 + rolls.get(i + 1) + rolls.get(i + 2); i += 1;
        } else if (rolls.get(i) + rolls.get(i + 1) == 10) {        // spare
            total += 10 + rolls.get(i + 2); i += 2;
        } else { total += rolls.get(i) + rolls.get(i + 1); i += 2; }
    }
    return total;   // 12 strikes -> 300; 21 rolls of 5 -> 150
}
```

This book has no Bowling Alley case study, so this outline is the whole answer here.

Q30. Design an online game engine

Requirements. Provide the core that any 2D game can build on: a loop that updates and draws the world at a steady rate, game objects with position and behaviour, keyboard and mouse input, collision detection, events between objects, loading images and sounds once, and screens (menu, playing, paused).

Main classes.

- `GameLoop`: runs a fixed-timestep update (say 60 times a second) and renders as often as possible. It catches up with several updates if a frame was slow.
- `Entity`: an id that holds `Component`s (`Position`, `Velocity`, `Sprite`, `Collider`). This is the Entity-Component-System (ECS) idea: behaviour comes from parts, not from a deep inheritance tree.

- `System` : `MovementSystem` , `CollisionSystem` and `RenderSystem` . Each one walks the entities that have the components it needs.
- `InputHandler` : maps keys to `Command` objects (`Jump` , `MoveLeft`), so key bindings can change.
- `EventBus` : lets objects publish "player died" or "coin collected" without knowing each other.
- `AssetManager` : loads each texture once and shares it.
- `SceneManager` : holds a stack of `Scene` s. Each scene is a state, such as `Menu` , `Level` or `Pause` .

Patterns. Component (composition over inheritance), Command (input), Observer (events), Flyweight (shared assets), State (scenes), Game Loop.

Concurrency and edge cases. Keep the world single-threaded. Input arrives from other threads into a `ConcurrentLinkedQueue` , and the loop drains it at the start of each tick. Asset loading can run on virtual threads, publishing "loaded" events. Use a fixed time step, so physics gives the same result on fast and slow machines.

Extension question. "Make it multiplayer." The server owns the true state and runs ticks. Clients send commands, predict movement locally, and correct themselves when the server state arrives.

Q193. Design a Jigsaw Puzzle

- Entities: `Side` (FLAT, TAB or BLANK plus a shape code), immutable `Piece` (four sides clockwise), `Board` with `Placement(piece, rotation)` , `PuzzleValidator` , `SolverStrategy` .
- Patterns: Strategy for solvers (brute-force backtracking vs edge-first with a `Side → pieces` index); all placement rules in one `Board.check` used by players and solvers.
- What interviewers probe: where rotation lives (in the placement, so pieces stay immutable), the facing side after `r` turns (`sides[(d − r) mod 4]`), why corners need no special rule, and the cost of solving (about n attempts with unique shapes, exponential with look-alikes, hence an attempt budget). → *Full design:* Ch 17 — Design a Jigsaw Puzzle

Platforms and marketplaces

Q31. ⋆ Design a cab booking / ride-sharing service (Uber; Flipkart "offer and select a ride")

Asked at: Flipkart

- Entities: `Rider` , `Driver` , `Vehicle` , `Location` , `RideRequest` , `Trip` , `FareCalculator` , `MatchingService` .
- Patterns: Strategy for matching (nearest, best rated) and for fares (surge), State for the trip, Observer for live updates.
- What interviewers probe: two riders matched to one driver, cancellation rules, and the carpool variant (offer a ride with seats, select the best ride). → *Full design:* Ch 20 — 20.4 Design a Ride Sharing Service (Uber)

Q32. ⋆ Design a food delivery system (Swiggy / Zomato; Flipkart "FoodKart")

Asked at: Flipkart

- Entities: `Restaurant` , `Menu` , `Item` , `Cart` , `Order` , `DeliveryAgent` , `RestaurantSelectionStrategy` .
- Patterns: State for the order, Strategy for choosing the restaurant (lowest cost, highest rating) and assigning agents, Observer for tracking.
- What interviewers probe: restaurant capacity (max orders at once) freed when an order completes, and one-restaurant-per-cart rules. → *Full design:* Ch 19 — Design an Online Food Delivery Service (Swiggy) + FoodKart

Q33. ⋆ Design CricInfo / a cricket scoreboard (Udaan "Cricket Match Dashboard")

Asked at: Udaan

- Entities: `Match` , `Team` , `Player` , `Innings` , `Over` , `Ball` (the event that everything derives from), `Scorecard` .
- Patterns: Template Method for match formats, Observer for live updates, "derive stats from balls, don't store totals".
- What interviewers probe: extras (wide, no-ball), strike rotation at the end of an over, and the machine-coding version that prints the scorecard after every over. → *Full design:* Ch 18 — 18.2 Design Cricinfo

Q34. ⋆ Design a stock exchange / order matching engine (LLD)

- Entities: `Order` (buy or sell, limit or market, quantity, price, time), `OrderBook` per symbol, `Trade`, `MatchingEngine`.
- Data structure: for each side, a `TreeMap<price, Deque<Order>>`. The best bid is the highest key and the best ask is the lowest. Match by price first, then by time.
- What interviewers probe: partial fills, cancelling an order sitting in the book, and a single writer thread per symbol so no locks are needed. → *Related design:* Ch 19 — Design an Online Stock Brokerage System

Q35. Design an online stock brokerage system

- Entities: `Account`, `Portfolio`, `Position`, `Order` (market, limit, stop), `Exchange` adapter, `Watchlist`.
- Patterns: Strategy/Factory for order types, State for order lifecycle, Observer for price alerts, Adapter for the exchange.
- What interviewers probe: blocking funds when an order is placed, partial fills, and order status kept in step with the exchange. → *Full design:* Ch 19 — Design an Online Stock Brokerage System

Q36. Design a cryptocurrency exchange

- Same core as Q34 (order book and matching) plus wallets per currency, deposits that wait for blockchain confirmations, and withdrawals that need approval.
- Patterns: State for deposits and withdrawals, Strategy for fees.
- What interviewers probe: exact decimal amounts (`BigDecimal` with each coin's scale), and never letting a balance go negative under concurrent trades. → *Related design:* Ch 19 — Design an Online Stock Brokerage System

Q37. Design Stack Overflow

- Entities: `User` (reputation, badges), `Question`, `Answer`, `Comment`, `Vote`, `Tag`, `Bounty`.
- Patterns: Observer for reputation changes and notifications, Strategy for search and ranking.
- What interviewers probe: one vote per user per post, accepting an answer, and permissions that depend on reputation. → *Full design:* Ch 18 — 18.1 Design Stack Overflow

Q38. Design LinkedIn

- Entities: `Member`, `Profile`, `Connection`, `Invitation`, `Post`, `Job`, `Message`.
- Patterns: State for invitations, BFS for 1st/2nd/3rd degree connections, Strategy for feed ranking.
- What interviewers probe: two-sided operations (accepting adds both sides), and degree computation limits. → *Full design:* Ch 18 — 18.3 Design LinkedIn

Q39. Design Facebook / a social network

- Entities: `User`, `Friendship`, `Follow`, `Post`, `Audience` (privacy), `Page`, `Group`, `NewsFeed`.
- Patterns: one visibility policy used by feed and search, Strategy for feed ranking, pull-model feed.
- What interviewers probe: friends vs followers, privacy checks on every read, and friend suggestions. → *Full design:* Ch 18 — 18.4 Design Facebook

Q40. ⋆ Design Twitter (post / follow / feed; LeetCode 355 "Design Twitter")

- The LeetCode version: `postTweet`, `getNewsFeed` (the 10 newest tweets from self and followees), `follow`, and `unfollow`.
- Data: a map from user to followees, and per user a list of tweets with a global increasing timestamp. Build the feed by merging the users' newest-first lists with a max-heap, k-way merge style, and stop after 10.
- What interviewers probe: why you pull (merge on read) rather than push here, self-follow, and unfollowing yourself. → *Related design:* Ch 18 — 18.4 Design Facebook (feed and follow)

Q41. ⋆ Design Amazon (online shopping system / cart / checkout)

- Entities: `Product`, `Catalog`, `Cart`, `Order`, `OrderItem`, `Payment`, `Shipment`, `Inventory`.
- Patterns: State for orders, Strategy for payment and shipping, Observer for status emails.
- What interviewers probe: reserving stock at checkout vs at payment, price changing while the item sits in the cart, and partial cancellations. → *Full design:* Ch 19 — Design Amazon — Online Shopping System

Q42. Design an online auction system (eBay; live IPL auction)

- Entities: `Auction` (item, reserve price, end time), `Bid`, `Bidder`, `Seller`, `AuctionHouse`.

- Patterns: State (scheduled, live, closed), Observer for outbid alerts, Strategy for auction type (English, sealed bid).
- What interviewers probe: two bids arriving at the same moment (serialise per auction), last-second sniping and extending the end time. → *Full design:* Ch 19 — Design an Online Auction System

Q43. Design Spotify / a music streaming service (LLD)

- Entities: `Track`, `Album`, `Artist`, `Playlist`, `User`, `Subscription`, `Player`, `PlayQueue`.
- Patterns: Strategy for recommendations and shuffle, State for the player, Iterator for queues.
- What interviewers probe: free vs premium limits, collaborative playlists, and the play-next queue vs the playlist order. → *Full design:* Ch 19 — Design a Music Streaming Service (Spotify)

Q44. Design a grocery store / POS / checkout (BigBasket)

- Entities: `Product` (priced by unit or by weight), `Inventory`, `Basket`, `Receipt`, `Discount`, `Payment`.
- Patterns: Strategy and Composite for discounts (buy-2-get-1, percent off a category), Decorator for tax.
- What interviewers probe: the order in which discounts apply, and returns. → *Full design:* Ch 19 — Design a Grocery Store (POS & Checkout)

Q45. Design a subscription-based sports score website

- Reuse the Cricinfo model for matches and add `Subscription` (plan, expiry), `Entitlement` checks and live score pushes.
- Patterns: Observer (subscribers to a match), Proxy for checking entitlements before content is served.
- What interviewers probe: free vs paid content per request, and fan-out of score updates to many subscribers. → *Related design:* Ch 18 — 18.2 Design Cricinfo

Q46. Design a logistics / shipment tracking system

- Entities: `Shipment`, `Parcel`, `Hub`, `Route`, `TrackingEvent`, `Courier`.
- Patterns: State for shipment status, Observer for customer updates, Strategy for route choice.
- What interviewers probe: events that arrive out of order (sort by event time, not arrival), and returns. → *Related design:* Ch 14 — 14.7 Design a Shipping Locker System

Q47. Design a newsletter service

- Entities: `Publication`, `Issue`, `Subscriber`, `Subscription` (topics, frequency), `Delivery`.
- Patterns: Observer / Pub-Sub (subscribers per topic), Strategy for channel (email, in-app), Template Method for rendering.
- What interviewers probe: unsubscribe taking effect at once, retries without double delivery, and daily digests. → *Related design:* Ch 20 — 20.6 Design a Pub-Sub System

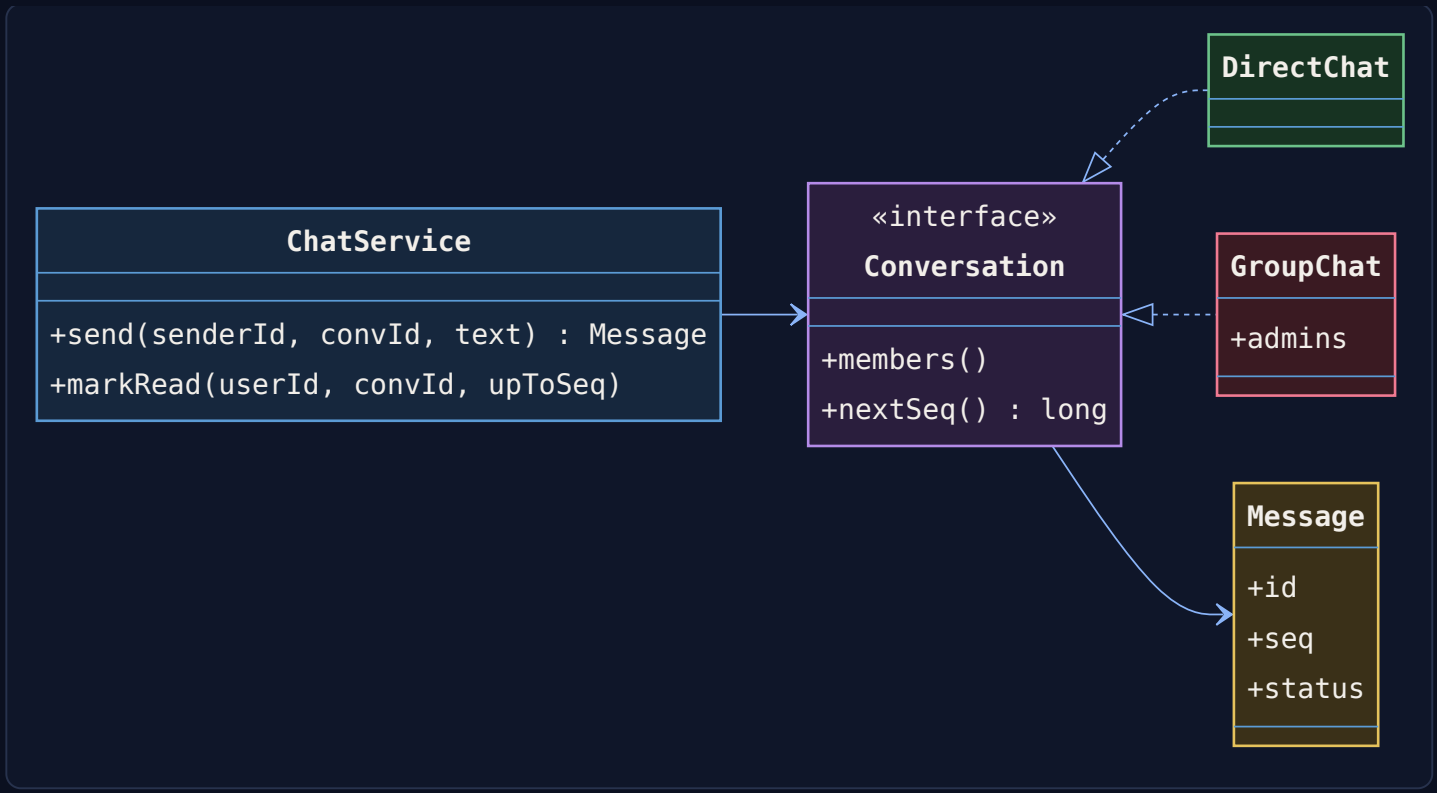
Q48. Design a maps navigator client (walk / car / bus / bike)

- Entities: `Navigator`, `Route`, `RouteStrategy` (one per travel mode), `Location`.
- Pattern: Strategy is the whole question. The navigator swaps strategies at runtime without `if` chains.
- What interviewers probe: adding a new travel mode without touching old code, and combining modes (walk + bus). → *Related design:* Ch 11 — 11.3 Strategy

Q49. Design a chat server / WhatsApp (LLD)

Requirements. Users send text messages one-to-one and in groups. Messages show sent, delivered and read ticks. Offline users receive their messages when they come back. Each conversation keeps its order. Groups have admins who add and remove members. Online and last-seen status is a nice extra.

Main classes.



- `Conversation` (sealed interface): `DirectChat` (two members) and `GroupChat` (members, admins).
- `Message` : a record (id, conversationId, senderId, seq, text, sentAt). Its delivery status is kept per recipient, because in a group each person reads at a different time.
- `Inbox`: per user, a queue of messages not yet delivered.
- `PresenceService`: online or offline, and last seen.
- `ConnectionRegistry` : the user's live sessions. Delivery listeners hang off it (Observer).

Patterns. Observer (deliver to connected sessions), Strategy (delivery channel: socket or push notification), and Mediator (`ChatService` coordinates instead of users talking to each other).

Concurrency and edge cases. Give each conversation an `AtomicLong` sequence. Sort messages by sequence, not by clock time, so senders on two phones cannot reorder history. Make delivery idempotent: a client acknowledges the highest sequence it has seen, and re-sent messages are ignored. A removed group member must not receive later messages, so check membership when the message is sent.

Extension question. "Add message editing and deletion." Store edits as new events that point at the original id, and let clients apply them, so the order stays clean.

Q50. Design Truecaller

Requirements. Users sign up with a phone number and may upload their contacts. Anyone can look up a number to see the most likely name and a spam score. Users report numbers as spam and block numbers. Users can search by name. A user can mark their own number as private.

Main classes.

- `PhoneNumber`: a value object. Its constructor normalises the input to E.164 format (for example `+919876543210`), so the same number can never be stored two ways.
- `User`: an id, an own number, contacts and a block list.
- `Contact`: a record (number, savedName).
- `GlobalDirectory`: maps each number to a name-count table (name → how many users saved it under that name). A lookup returns the most frequent name, and a registered user's own name always wins.
- `SpamService`: counts reports per number and gives a score. A number is "likely spam" past a threshold, weighted by how trustworthy each reporter is.
- `CallerIdService`: answers `lookup(number)` with `CallerInfo(name, spamScore, blocked)` .
- `NameSearchIndex`: a trie (prefix tree) from name prefixes to numbers.

Patterns. Value object, Strategy (spam scoring, name ranking), Observer (a new contact upload updates the directory and the index), Facade (`CallerIdService`).

Concurrency and edge cases. Uploads and reports arrive at the same time from many users. Use `ConcurrentHashMap<PhoneNumber, ConcurrentHashMap<String, LongAdder>>` for name counts and `LongAdder` for spam counts, so updates never block reads. Count one report per reporter per number, so repeat reports don't inflate the score. A private user's name should not be served from the crowd data.

Extension question. "Mark verified business numbers." Add a `VerifiedBusiness` record checked before the crowd data, so a verified name always beats crowd guesses.

Q51. Design a mentorship / mock-interview platform (Pramp)

Requirements. Users sign up with the topics they practise (data structures, system design) and their available time slots. The platform pairs two users who share a topic and a free slot. Each person interviews the other for half the session. Afterwards both leave feedback. Users can cancel, and they are reminded before a session.

Main classes.

- `User`: id, topics, a calendar of `AvailabilitySlot`s, and a rating.
- `AvailabilitySlot`: a record (start, end).
- `MatchRequest`: user, topic, preferred slots.
- `MatchingService`: keeps the open requests per topic and slot, and pairs requests through a `MatchingStrategy` (first come first served, similar rating).
- `Session`: two participants, topic, slot, question set, and a state (`SCHEDULED`, `IN_PROGRESS`, `COMPLETED`, `CANCELLED`, `NO_SHOW`).
- `Feedback`: a record (from, to, scores, notes).
- `ReminderScheduler`: uses a `ScheduledExecutorService` to fire reminders.

Patterns. Strategy (matching), State (session), Observer (send notifications on state change), Builder for sessions.

Concurrency and edge cases. One user must never be matched twice for the same slot. Run the matching in one step: lock both users' calendars in a fixed order (lowest id first) to avoid deadlock, check both are still free, then book both. A cancellation within an hour counts against the user's reliability, and the partner goes back into the pool. A no-show is detected when a session never moves to `IN_PROGRESS`.

Extension question. "Add paid sessions with expert mentors." Add a `Mentor` role with a price. Payment must succeed before the session moves to `SCHEDULED`. The rest of the flow stays the same.

Q52. Design an Application Tracking System (Greenhouse)

Requirements. A company posts jobs. Each job has a pipeline of stages (Applied → Phone screen → Onsite → Offer → Hired). Candidates apply, and recruiters move applications forward or reject them. Interviewers submit scorecards. Hiring managers see a dashboard of each job's pipeline. Every action is audited.

Main classes.

- `Job`: id, title, department, status (`OPEN`, `ON_HOLD`, `CLOSED`), and its own `Pipeline` (an ordered list of `Stage`s).
- `Candidate`: a person, who may have several `Application`s.
- `Application`: candidate + job, the current stage, a status (`ACTIVE`, `REJECTED`, `WITHDRAWN`, `HIRED`), and a history of stage changes.
- `Interview`: application, interviewers and time. `Scorecard`: interviewer, ratings and a recommendation.
- `ApplicationService`: `advance`, `reject` and `hire`. It checks the rules (no skipping to offer without scorecards, unless overridden).
- `PermissionPolicy`: decides what a recruiter, an interviewer or a hiring manager may see or do.
- `AuditLog` and `NotificationService`.

Patterns. State (application status), Template Method or configurable stage lists (each job has its own pipeline), Chain of Responsibility (the checks before advancing), Observer (emails and audit entries on each change).

Concurrency and edge cases. Two recruiters may move the same application at the same time. Use optimistic locking: a version number on `Application`, where an update fails if the version changed. Duplicate applications (same candidate, same job) are merged or rejected. Closing a job should not delete applications, only stop new ones.

Extension question. "Auto-reject candidates who fail a knock-out question." Add a rule to the advance chain that runs on `Applied`.

Q53. Design a call center

Requirements. The center has three levels of employee: respondents, managers and directors. An incoming call goes to a free respondent. If the respondent can't handle it, or none is free, it escalates to a manager, then to a director. If nobody is free, the call waits in a queue. When an employee finishes a call, they take the next waiting call at their level.

Main classes.

- `Employee` (abstract or sealed): id, `Rank`, and a current call. Subtypes `Respondent`, `Manager`, `Director`. Each can `canHandle(call)`.
- `Call`: caller, required minimum rank (it rises when escalated), state (`WAITING`, `IN_PROGRESS`, `ESCALATED`, `ENDED`).
- `CallDispatcher`: holds a queue of free employees per rank (`BlockingQueue<Employee>`) and a queue of waiting calls per rank. `dispatch(call)` tries the call's rank and above.
- `EscalationHandler`: a chain from respondent to manager to director.

Patterns. Chain of Responsibility (escalation), Strategy (routing: least busy, skills-based), State (call lifecycle).

Concurrency and edge cases. Calls arrive and end on many threads. Use `poll()` on the free-employee queues, so taking an employee is atomic. When an employee becomes free, first check the waiting-call queue for their rank (and lower ranks), and only then return to the free pool. Do these two steps under one lock, or a call could be missed. If a call is dropped mid-escalation, end it cleanly and free the employee. A call that waits too long can be escalated or offered a callback.

Extension question. "Add priority customers." Change the waiting queue to a `PriorityBlockingQueue` ordered by priority, then by arrival time.

Q54. Design a Survey / Google Forms system

Requirements. An owner builds a form with a title and questions of several types: short text, paragraph, single choice, multiple choice, rating, and date. Questions can be required and can have validation (email format, number range). The form can show or skip a section based on an earlier answer. Respondents submit responses. The owner sees a summary per question and can close the form.

Main classes.

- `Form`: id, owner, settings (one response per user, deadline), an ordered list of `Section`s, a version, and a status.
- `Question` (sealed interface): `TextQuestion`, `ChoiceQuestion` (options, multiSelect), `RatingQuestion` (scale), `DateQuestion`. Each has `validate(Answer)`.
- `ValidationRule`: composable rules (required, regex, range).
- `BranchRule`: "if the answer to question 3 is 'No', go to section 5".
- `Response`: form id, form version, respondent, and `Map<questionId, Answer>`.
- `FormBuilder`: builds a form step by step and checks it is valid.
- `Summary`: per question type, a different aggregation (counts, average, word list).

Patterns. Builder (forms), Composite (form → sections → questions), Strategy/Specification (validation), Visitor or a pattern-matching `switch` over the sealed `Question` types for summaries.

Edge cases. Editing a form after responses exist: create a new version, so old responses still point at the questions they answered. Enforce one response per user with a set of respondent ids, updated atomically with `ConcurrentHashMap.newKeySet().add(...)`. Validate every answer on the server, never trusting the client.

Extension question. "Add a quiz mode with scores." Add a correct answer and points to each question, plus a `Grader` that walks the response. Forms that are not quizzes are unaffected.

Q55. Design an employee management / payroll system

Requirements. Store employees with department, role and manager. Show the org chart and everyone who reports (directly or indirectly) to a manager. Track leave balances and requests. Run monthly payroll for salaried, hourly and commission employees, with deductions (tax, provident fund) and payslips.

Main classes.

- `Employee`: id, name, department, role, managerId, join date, and a `PayPolicy`.
- `OrgChart`: manager id → list of direct reports. `allReports(id)` walks the tree breadth-first.

- `PayPolicy` (interface): `SalariedPolicy` (monthly amount, prorated for part months), `HourlyPolicy` (rate × timesheet hours, overtime), `CommissionPolicy` (base + percent of sales).
- `Deduction` (interface): `IncomeTax`, `ProvidentFund`, `LeaveWithoutPay`, applied in order.
- `PayrollRun`: a month and the list of `Payslip`s it produced. Its status is `DRAFT`, `APPROVED` or `PAID`.
- `LeaveService`: balances per leave type, and requests that a manager approves.

Patterns. Strategy (pay policies), Chain of Responsibility or a list of deductions, Composite (org tree), State (payroll run and leave request).

Edge cases. Hold money as `long` paise or `BigDecimal`, and round only once, at the end. Mid-month joiners and leavers get prorated pay. A manager cycle (A reports to B, B to A) must be rejected when a manager is assigned. Once a payroll run is `PAID`, it is never edited. Corrections go into next month as adjustments.

Extension question. "Change an employee's salary from the middle of the month." Store pay policies with an effective date, and prorate between the two policies.

Components and infrastructure

Q56. ⋆ Design a Rate Limiter (LLD)

Asked at: Atlassian (P4 LLD round)

- Entities: `RateLimiter` interface, `TokenBucket`, `SlidingWindowLog`, `FixedWindowCounter`, a `RateLimiterRegistry` (one limiter per client key), and a config.
- Patterns: Strategy for the algorithm, Factory to build from config.
- What interviewers probe: thread safety (per-key lock or CAS), lazy refill based on elapsed time instead of a timer thread, and the burst behaviour of each algorithm. → *Full design:* Ch 20 — 20.7 Design a Rate Limiter; also Ch 13 — 13.11 Rate limiter — token bucket

Q57. ⋆ Design an LRU Cache

- Data structure: a `HashMap<key, Node>` plus a doubly linked list. Every `get` or `put` moves the node to the front, and eviction removes the tail. Both operations are O(1).
- In Java you can also show `LinkedHashMap` with `accessOrder = true` and an overridden `removeEldestEntry`.
- What interviewers probe: writing the linked list by hand with dummy head and tail nodes, and making it thread-safe. → *Full design:* Ch 20 — 20.3 Design an In-Memory Cache (and an LRU Cache from scratch)

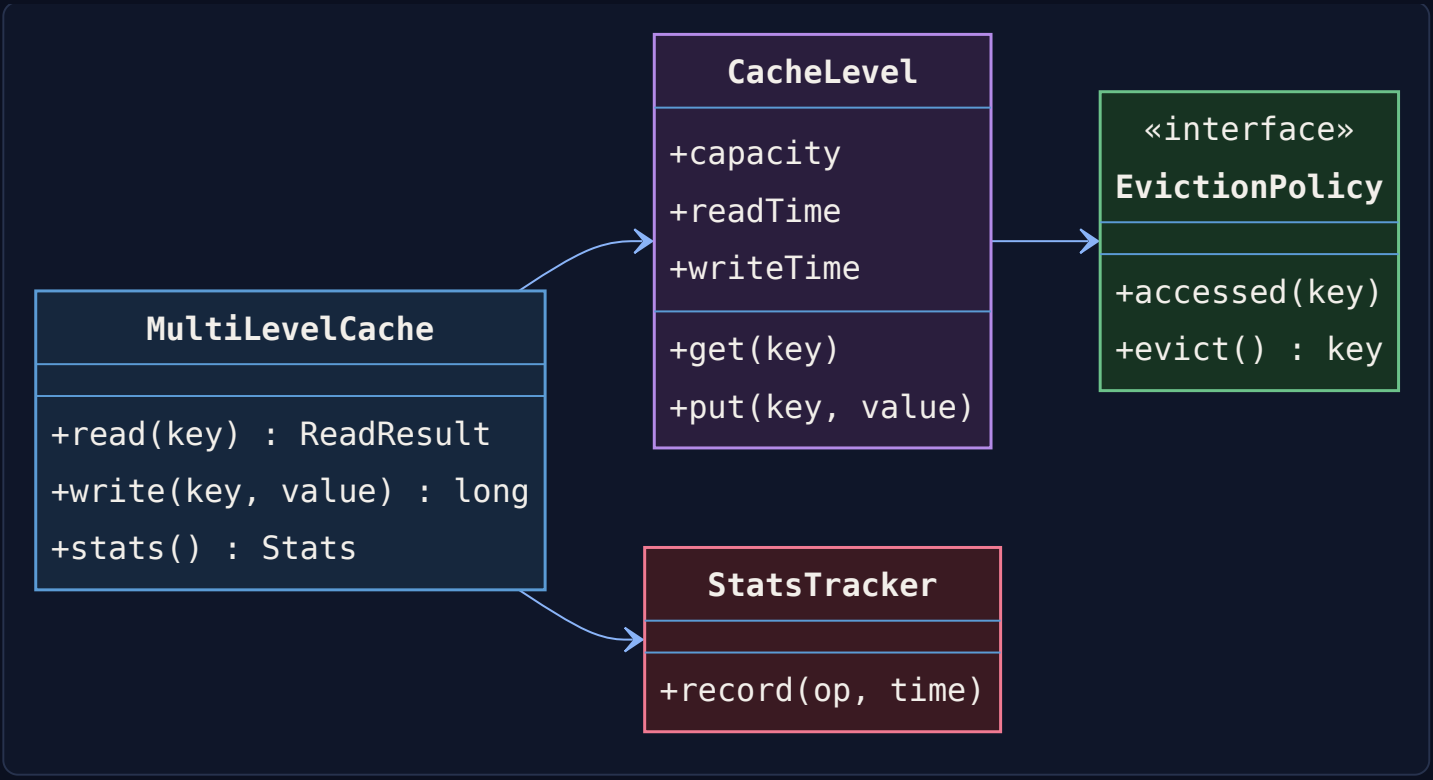
Q58. ⋆ Design a cache with pluggable eviction policies (LRU / LFU)

- Entities: `Cache<K, V>`, `Storage`, `EvictionPolicy` (`keyAccessed`, `evictKey`) with LRU, LFU and FIFO implementations.
- Pattern: Strategy for eviction. LFU in O(1) uses a frequency → linked set map plus a `minFreq` pointer (LeetCode 460).
- What interviewers probe: tie-breaking in LFU, and keeping the policy and the storage in step. → *Full design:* Ch 20 — 20.3 Design an In-Memory Cache

Q59. ⋆ Design a multi-level cache

Requirements. A cache has N levels (L1, L2, … Ln). Each level has a capacity, a read time and a write time. Reading a key checks L1 first, then each lower level, until the key is found. When found at level i, write the value into every level above i. Writing a key writes it to every level, skipping a level that already holds the same value. Report the time taken by every read and write, plus the average read and write time over the last N operations. Show usage per level.

Main classes.



- `CacheLevel`: one level, with its own storage and `EvictionPolicy` (LRU by default).
- `MultiLevelCache`: holds the ordered list of levels and runs the read and write paths. Time is the sum of each level's read time visited, plus the write times added.
- `StatsTracker`: a ring buffer of the last N timings for reads and for writes, with running sums so averages are O(1).

Patterns. Chain of Responsibility (a level passes a miss to the next one), Strategy (eviction), Builder (configure levels).

Concurrency and edge cases. In a machine-coding round, guard the whole cache with one lock (`synchronized` or a `ReentrantLock`). A read-write lock does not help here, because a read is not read-only: it updates each level's LRU order and copies the value into the levels above, so every read would need the write lock anyway. For more throughput, lock levels one at a time, always top-down, so no deadlock can form. An eviction at L1 does not remove the key from L2. Reject capacity 0 or an empty list of levels.

Extension question. "Make it write-back instead of write-through." Writes go to L1 only and are marked dirty. They reach lower levels when evicted, and a flush is forced on shutdown.

Q60. ⋆ Design an in-memory key-value store (with transactions / TTL / versioning)

- Base: `ConcurrentHashMap<String, Entry>` where `Entry` holds value and expiry; expired entries are removed lazily on read plus by a periodic sweeper.
- Transactions (`begin`, `set`, `get`, `commit`, `rollback`, nestable): keep a stack of change maps. Reads check the top map, then the maps below it, then the base store. `commit` merges the top map down, and `rollback` pops it.
- Versioning: store a list of `(version, value)` per key and answer `get(key, version)` by binary search.
- What interviewers probe: isolation between clients (one transaction stack per session), and deletes inside a transaction (store a tombstone marker). → *Related design:* Ch 20 — 20.3 Design an In-Memory Cache; Ch 13 — 13.6 Thread-safe cache with TTL

Q61. ⋆ Design a Logging Framework / Logger library

- Entities: `Logger`, `LogLevel`, `LogMessage`, `Appender` (console, file, database), `Formatter`, `LoggerConfig`.
- Patterns: Singleton or factory for loggers, Strategy for appenders and formatters, Chain of Responsibility for level filtering.
- What interviewers probe: async logging through a queue without losing messages on shutdown, and per-package log levels. → *Full design:* Ch 20 — 20.5 Design a Logging Framework

Q62. ⋆ Design a Pub-Sub system / message queue (LLD)

- Entities: `Broker`, `Topic`, `Message`, `Publisher`, `Subscriber`, `Subscription` (with offset).

- Patterns: Observer, and a per-subscriber offset so slow consumers don't block fast ones.
- What interviewers probe: at-least-once delivery with acks and retries, ordering within a topic, and replay from an offset. → *Full design:* Ch 20 — 20.6 Design a Pub-Sub System

Q63. ⋆ Design a Notification System (LLD)

Asked at: Flipkart

- Entities: `Notification`, `Channel` (email, SMS, push), `Template`, `UserPreference`, `NotificationService`, `RetryPolicy`.
- Patterns: Strategy per channel, Template Method or Builder for messages, Observer for events, Decorator for retries.
- What interviewers probe: user opt-outs, retries with backoff, falling back to another channel, and not sending duplicates. → *Full design:* Ch 20 — 20.2 Design a Notification Service

Q64. Design a HashMap

- Array of buckets; each bucket is a linked list (Java turns long chains into balanced trees); index = `hash(key) & (capacity - 1)`; resize and rehash when size exceeds capacity × load factor (0.75).
- For the concurrent version: lock striping (a lock per group of buckets) or CAS on empty buckets.
- What interviewers probe: `equals` / `hashCode` contract, why capacity is a power of two, and resizing under concurrency. → *Full design:* Ch 13 — 13.7 Design a (simplified) Concurrent HashMap

Q65. ⋆ Design a thread-safe blocking queue / producer-consumer

- A bounded buffer: a lock with two conditions, `notFull` and `notEmpty`. `put` waits while the queue is full, and `take` waits while it is empty. Always wait inside a `while` loop, never an `if`.
- Alternative: two `Semaphore`s (free slots, filled slots) plus a mutex.
- What interviewers probe: spurious wakeups, `signal` vs `signalAll`, and shutdown with a poison pill. → *Full design:* Ch 13 — 13.8 Thread-safe bounded Blocking Queue; Ch 12 — 12.22 Producer-consumer pattern

Q66. Design a thread-safe cache with TTL

- `ConcurrentHashMap` of entries with expiry time; lazy expiry on `get`; a `ScheduledExecutorService` or a `DelayQueue` for active cleanup.
- What interviewers probe: get-or-compute without two threads loading the same key (`computeIfAbsent`), and clock choice (`System.nanoTime` for durations). → *Full design:* Ch 13 — 13.6 Thread-safe cache with TTL

Q67. Design a concurrent Bloom filter

- A bit array with k hash functions. "Might contain" can give false positives but never false negatives.
- Concurrent version: an `AtomicLongArray` with CAS on each word, so adds are lock-free.
- What interviewers probe: choosing the size and k for a target false-positive rate, and why you cannot delete (use a counting Bloom filter for that). → *Full design:* Ch 13 — 13.9 Concurrent Bloom Filter

Q68. Design a thread pool / executor

- Entities: a work queue (`BlockingQueue<Runnable>`), worker threads looping on `take()`, and `submit` returning a `Future`. Shutdown has a graceful mode and an immediate mode.
- What interviewers probe: what happens when the queue is full (rejection policy), worker death on exceptions, and when virtual threads make a pool unnecessary. → *Full design:* Ch 12 — 12.21 Thread pool pattern

Q69. Dining philosophers / print FooBar alternately / building H2O (LeetCode concurrency)

- Short puzzles about making threads take turns or meet at a point. The tools are semaphores, `CyclicBarrier`, lock ordering, and conditions.
- Dining philosophers: break the deadlock by always taking the lower-numbered fork first, or by letting at most four philosophers sit at once.
- What interviewers probe: proving there is no deadlock and no starvation. → *Full design:* Ch 13 — 13.1–13.5

Q70. Design a connection pool / object pool

- Entities: `Pool` with idle and in-use sets, `acquire(timeout)` and `release`, max size, validation on borrow, and idle eviction.
- Pattern: Object Pool, with a `Semaphore` capping how many objects can be borrowed at once.

- What interviewers probe: leaked connections (borrow timeouts, try-with-resources wrapper), and broken connections returned to the pool. → *Full design:* Ch 8 — 8.14 Bonus — Object Pool

Q71. Design a payment gateway (Razorpay) — LLD

- Entities: `PaymentRequest`, `Payment` (state: created, authorised, captured, failed, refunded), `PaymentMethod` (card, UPI, net banking), `ProcessorAdapter` per bank, `Refund`, `WebhookNotifier`.
- Patterns: State (payment lifecycle), Strategy/Adapter (methods and processors), Observer (webhooks to merchants).
- What interviewers probe: an idempotency key so retries never double-charge, timeouts with an unknown result (check status before retrying), and refunds that are partial. → *Related design:* Ch 21 — 21.2 Design a Digital Wallet Service

Q72. Design a text editor with undo / redo (Notion / Google Docs LLD)

Asked at: Notion

- Entities: `Document` (a buffer, such as a list of lines or a gap buffer), `Cursor`, `EditCommand` (insert, delete), and `History` (an undo stack and a redo stack).
- Patterns: Command (each edit knows how to undo itself) or Memento (save snapshots). Any new edit clears the redo stack.
- What interviewers probe: Command vs Memento memory cost, grouping keystrokes into one undo step, and collaborative editing (operational transforms) as a stretch. → *Full design:* Ch 10 — 10.2 Command and 10.5 Memento

Q73. Design a business rule engine

- Entities: `Rule` (condition + action, priority), `Condition` (composable with and / or / not), `Fact` (input data), `RuleEngine` (evaluate in priority order, first-match or all-match).
- Patterns: Specification (composable conditions), Chain of Responsibility (rule order), Interpreter if rules come from a text language.
- What interviewers probe: loading rules from config without redeploying, and conflicts when two rules match. → *Related design:* Ch 11 — 11.7 Beyond GoF (Specification); Ch 10 — 10.1 Chain of Responsibility

Q74. Design a circuit breaker (LLD)

- States: `Closed` (count failures), `Open` (fail fast until a cool-down ends), `HalfOpen` (let a few trial calls through; success closes, failure re-opens).
- Patterns: State, plus Proxy/Decorator wrapping the protected call.
- What interviewers probe: a sliding window of failure rate vs a plain count, and thread-safe state changes (`AtomicReference` with CAS). → *Related design:* Ch 20 — 20.2 Design a Notification Service (provider circuit breaker)

Q75. Design a Hit Counter (LeetCode 362)

Requirements. `hit(timestamp)` records a hit at a time given in seconds. `getHits(timestamp)` returns how many hits happened in the past 300 seconds, including the current second. Timestamps arrive in non-decreasing order. Follow-ups: what if there are millions of hits per second, and what if it must be thread-safe?

Main classes.

- `HitCounter` (interface) with `hit` and `getHits`.
- `QueueHitCounter`: a deque of timestamps. `getHits` drops entries older than `t - 299` and returns the size. It is simple, but memory grows with the number of hits.
- `BucketHitCounter`: two arrays of size 300, `times[]` and `counts[]`. The index is `t % 300`. If `times[i] != t`, the bucket is stale, so reset it to `(t, 1)`; otherwise increment. `getHits` sums the buckets whose `times[i] > t - 300`. Memory is fixed, and `getHits` costs O(300).
- `HitCounterRegistry`: one counter per key (per URL or per user), in a `ConcurrentHashMap`.

Patterns. Strategy (swap implementations), with the sliding window as the core idea. This is the same window as a rate limiter, and interviewers often link the two.

Concurrency and edge cases. For thread safety, guard each bucket's reset-or-increment as one step. Lock per bucket (stripes), or store `(time, count)` packed into one `AtomicLong` updated with CAS. For very high hit rates, use a `LongAdder` per bucket. Hits from the future should not happen, but clamp them anyway. If time jumps forward by more than 300 seconds, the stale check already clears old buckets.

Extension question. "Make the window configurable and also report per minute and per hour." Keep several bucket arrays at different resolutions (seconds, minutes), like a time-series database's roll-ups.

Q76. Design a URL shortener (LLD / class design)

Requirements. `shorten(longUrl)` returns a short code such as `aZ3k9Q`. `resolve(code)` returns the long URL. Users may request a custom alias. Links can expire. The service counts clicks. The same long URL may get the same code or a new one; ask which.

Main classes.

- `ShortUrl` : a record (code, longUrl, ownerId, createdAt, expiresAt).
- `CodeGenerator` (interface): `CounterBase62Generator` (an `AtomicLong` counter encoded in base 62, short with no collisions, but guessable), `HashGenerator` (the first 7 characters of a hash of the URL, retrying with a salt on collision), and `RandomGenerator`.
- `UrlRepository` : `saveIfAbsent(ShortUrl)`, `find(code)`. The in-memory version is a `ConcurrentHashMap`.
- `UrlShortenerService` : validates the URL, applies the alias or generates a code, and handles expiry.
- `ClickTracker` : a `LongAdder` per code, with events sent on asynchronously for analytics.

Patterns. Strategy (code generation), Repository, Facade (the service), Observer (click events).

Concurrency and edge cases. Two users may ask for the same custom alias. `putIfAbsent` decides the winner atomically, and the loser gets "alias taken". The counter generator never collides, while hash and random generators must retry. Reject non-HTTP(S) URLs and links to the shortener itself (loops). Expired codes return "not found". Remove them lazily, plus a sweeper. Keep codes case-sensitive and avoid look-alike characters if humans type them.

Extension question. "Hide the order of codes, so people cannot enumerate links." Scramble the counter with a reversible mixing function before encoding. It stays collision-free but no longer looks sequential.

Q77. Design a JSON / CSV parser

Requirements. Parse JSON text into an in-memory tree: objects, arrays, strings (with escapes and `\uXXXX`), numbers (sign, fraction, exponent), `true`, `false` and `null`. Give clear errors with line and column. Serialise the tree back to text. For CSV: split rows into fields, where quoted fields may contain commas, newlines and doubled quotes (`""`).

Main classes.

- `JsonValue` (sealed interface) with records `JsonObject(Map)`, `JsonArray(List)`, `JsonString`, `JsonNumber(BigDecimal)`, `JsonBool`, `JsonNull`.
- `Lexer`: turns characters into `Token`s (`LBRACE`, `STRING`, `NUMBER`, …) and tracks position.
- `Parser` : recursive descent. `parseValue` looks at the next token and calls `parseObject`, `parseArray` and so on. One method per grammar rule.
- `JsonWriter`: walks the tree with a pattern-matching `switch` over the sealed types. The compiler (Java 21 and later) checks that every case is covered.
- `ParseException` : message, line, column.
- `CsvReader` : a small state machine with the states `FIELD_START`, `IN_UNQUOTED`, `IN_QUOTED`, `QUOTE_IN_QUOTED`.

Patterns. Interpreter/Composite (the value tree), Visitor-style traversal (here a sealed `switch`), State (the CSV reader), Builder for writing.

Edge cases. Duplicate keys (last wins, or an error). Very deep nesting, so add a depth limit against stack overflow. Leading zeros (`01` is invalid in JSON). Trailing commas (invalid). Big numbers (use `BigDecimal`). Text left over after the root value (an error). For huge inputs, offer a streaming, event-based API (`startObject`, `key`, `value`) instead of building the whole tree.

Extension question. "Map JSON onto Java records." Use reflection on record components (`getRecordComponents()`) to match keys to constructor parameters.

Q78. Design a configuration management system (LLD)

Requirements. Services read typed config values (`int`, `boolean`, `duration`) by key. Values can be set at several layers: defaults, then environment (dev, staging, prod), then service, then instance. The most specific layer wins. Every change creates a new version and records who changed what. Users can roll back to an earlier version. Running services are told about changes without restarting.

Main classes.

- `ConfigKey`: a record (namespace, name), plus a `ConfigSchema` that holds the type, default, and validation per key.
- `ConfigSnapshot`: an immutable map for one version: version number, author, time, values.
- `ConfigStore`: keeps the list of snapshots per scope, and an `AtomicReference<ConfigSnapshot>` for the current one. `set` builds a new snapshot (copy-on-write) and swaps it in with CAS.
- `ConfigResolver`: looks a key up through the layers in order (instance → service → environment → default).
- `ConfigWatcher` (Observer): a client registers a callback on keys or prefixes.
- `AuditLog`: who, what, when, old value and new value.

Patterns. Chain of Responsibility (layered lookup), Observer (change listeners), Memento/immutable snapshots (versions and rollback), Strategy for value parsers.

Concurrency and edge cases. Readers never lock: they read the current snapshot reference. Writers race through CAS and retry on conflict, or use compare-by-version so a stale edit is rejected. Validate against the schema before publishing, so a bad value never goes live. A listener that throws must not stop the other listeners.

Extension question. "Add feature flags rolled out to 10% of users." Add a `FlagRule` with a percentage. Hash the user id into 0–99 and compare it, so each user gets a stable result.

Q79. Design AWS Lambda (serverless) — LLD

Requirements. Users register a function: name, runtime (Java, Python), memory, timeout, code location. `invoke(name, payload)` runs it synchronously and returns the result, or runs it asynchronously with retries. Warm environments are reused, and new ones are started (cold start) when none is free. A per-function limit caps how many run at once. Idle environments shut down after a while.

Main classes.

- `FunctionConfig`: a record (name, runtime, memoryMb, timeout, maxConcurrency, codeRef).
- `FunctionRegistry`: create, update and delete functions, with versions.
- `ExecutionEnvironment`: one sandbox. Its state goes `INITIALISING` → `READY` → `BUSY` → `READY` …, and finally `TERMINATED`.
- `RuntimeFactory`: builds an environment for each runtime type.
- `EnvironmentPool`: per function, a deque of warm environments (newest first, so old ones age out).
- `Invoker`: takes a concurrency permit, borrows or creates an environment, runs with a timeout, and returns the environment to the pool.
- `AsyncQueue` + `RetryPolicy` + `DeadLetterQueue` for asynchronous calls.
- `IdleReaper`: a scheduled task that terminates environments idle longer than a limit.

Patterns. Object Pool (warm environments), Factory (runtimes), State (environment lifecycle), Proxy (the `Invoker` hides pooling and limits), Strategy (retry backoff).

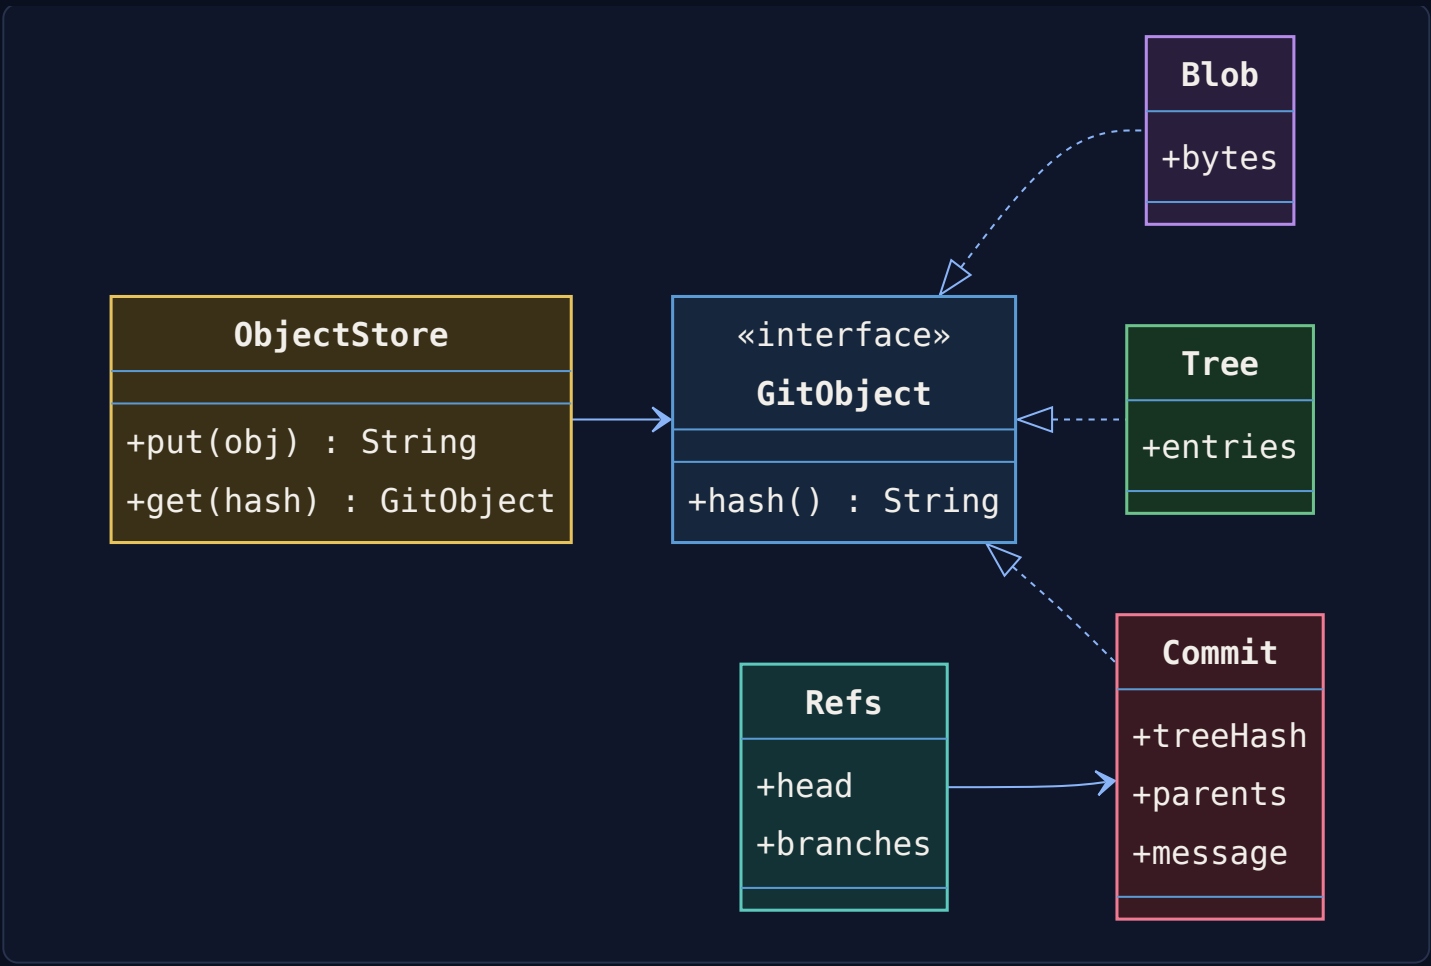
Concurrency and edge cases. A per-function `Semaphore` enforces `maxConcurrency`. Throttle by failing fast, or queue async calls. Run the handler on a virtual thread and wait with `Future.get(timeout)`. On timeout, kill and discard that environment, never reuse it. A crashed environment is removed, not returned. An update to the code marks old warm environments as stale.

Extension question. "Add provisioned concurrency." Pre-create N environments at deploy time, and keep the reaper from removing them.

Q80. Design a version-control / Git-like system

Requirements. Support `init`, `add` (stage files), `commit`, `log`, `branch`, `checkout`, `diff`, and `merge` for a single local repository. Identical content is stored once. History is never changed by normal commands.

Main classes.



- `Blob`: file content. `Tree`: a directory with name → (blob or tree) hash entries. `Commit`: tree hash, parent hashes, author, time, message. All of them are immutable records.
- `ObjectStore`: content-addressed storage, a map from hash (SHA-256 of the content) to object. The same content always gets the same hash, so duplicates cost nothing.
- `Index` (staging area): path → blob hash for the next commit.
- `Refs`: branch name → commit hash, and `HEAD` → a branch name (or a commit, the "detached" state).
- `Repository`: a facade exposing the commands, with each command as a `Command` object for the CLI.

Patterns. Composite (trees of trees), Command (CLI operations), immutable value objects, Facade.

Algorithms and edge cases. `log` follows parent links. `merge` first finds the merge base, the lowest common ancestor of the two commits (BFS from both sides). If one side is the base, just move the branch pointer (a fast-forward). Otherwise do a three-way merge per file and mark conflicts. `checkout` with uncommitted changes must refuse, or warn. Moving a branch is one atomic reference update.

Extension question. "Add `revert`." Create a new commit that undoes the changes of an earlier one. Never rewrite history.

Machine-coding round favourites

Q81. ⋆ Design Splitwise (expense sharing)

Asked at: Swiggy (SDE1; SDE2, twice)

- Entities: `User`, `Group`, `Expense` with a `Split` type (equal, exact, percent), `BalanceSheet` (who owes whom), `SettlementService`.
- Patterns: Strategy/Factory for split types, plus validation that the splits add up to the total.
- What interviewers probe: rounding (give the leftover paise to one person), simplifying debts (greedy matching of the biggest creditor with the biggest debtor), and the command-line input format. → *Full design:* Ch 20 — 20.1 Design Splitwise

Q82. ⋆ Design a task scheduler / distributed task scheduler (LLD)

Asked at: Flipkart (SDE2, Jan 2026)

- Entities: `Task` (id, payload, run-at time, recurrence, retries), `Scheduler` (`schedule`, `cancel`), a `DelayQueue` or `PriorityQueue` ordered by run time, a `Worker` pool, and an `ExecutionRecord`.
- Design: one dispatcher thread waits for the next due task (a `Condition` with a timed wait, woken early when a sooner task arrives) and hands it to a worker pool. A recurring task schedules its next run after finishing.
- What interviewers probe: cancelling a queued task, retries with backoff, and the distributed version (lease a task in storage with a timeout, so a crashed worker's task is picked up again). → *Related design:* Ch 12 — 12.21 Thread pool pattern

Q83. ⋆ Design an in-memory file system (mkdir / ls / find)

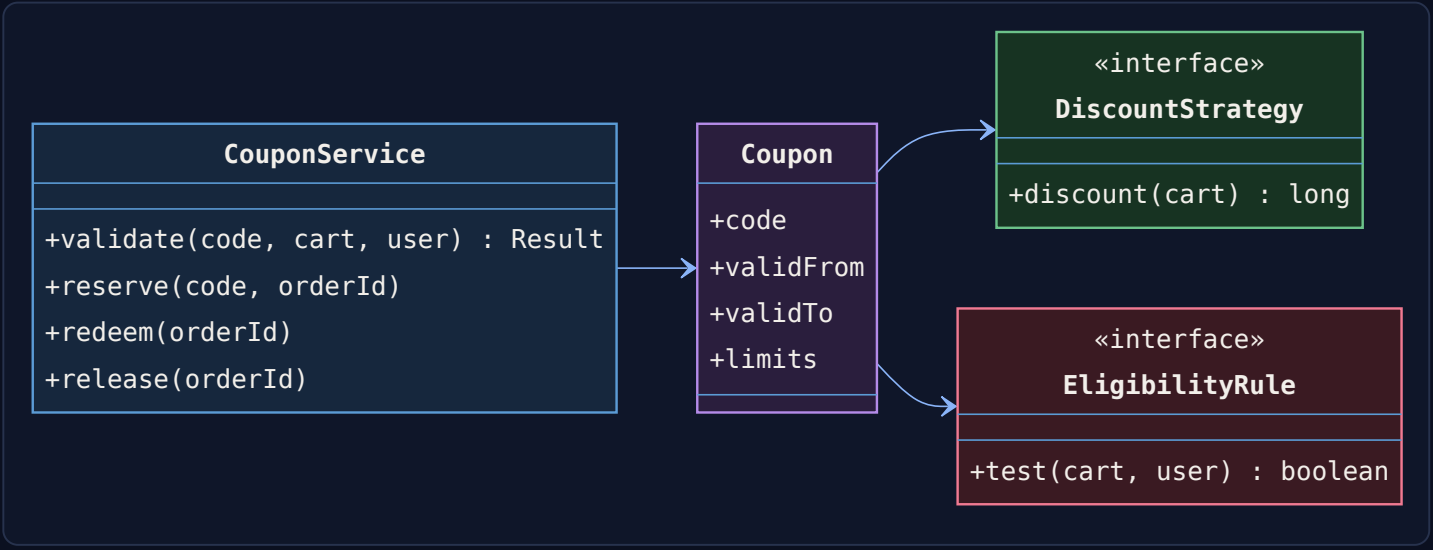
Asked at: ClearTrip and others (Unix `find`)

- Entities: `FileSystemNode` (sealed: `File`, `Directory`), `Path` parsing, and `FileSystem` (`mkdir -p`, `ls`, `addContent`, `readContent`).
- Patterns: Composite for the tree, Specification/Strategy for `find` filters (name, extension, size) combined with and / or.
- What interviewers probe: adding a new filter without changing existing ones, and walking a deep tree without recursion limits. → *Full design:* Ch 21 — 21.8 Design a Unix File Search (`find`); Ch 9 — 9.3 Composite Pattern

Q84. ⋆ Design a coupon / promo-code system

Requirements. Admins create coupons: a code, a discount type (flat ₹100 off, 20% off up to ₹150, buy-one-get-one), a validity window, a total usage limit, and a per-user limit. Eligibility rules include a minimum cart value, certain categories, first order only, and certain user segments. At checkout the user applies a code. The system validates it, computes the discount, and redeems it when payment succeeds. Coupons can or cannot be combined (stacking).

Main classes.



- `DiscountStrategy`: `FlatDiscount`, `PercentDiscount(cap)`, `BuyXGetY`. The result never exceeds the eligible amount.
- `EligibilityRule`: small rules joined with `and` / `or` (Specification). Each failure gives a reason the user can read.
- `Redemption`: a record (code, userId, orderId, status: `RESERVED`, `REDEEMED`, `RELEASED`).
- `UsageCounter`: total and per-user counts.

Patterns. Strategy (discounts), Specification / Chain of Responsibility (eligibility), Factory (build a coupon from an admin request).

Concurrency and edge cases. A "first 1,000 users" coupon must not be used 1,001 times. At checkout, *reserve* a use: increment the counter atomically only while it is below the limit (a CAS loop on an `AtomicInteger`). Move it to `REDEEMED` when payment succeeds, and *release* it on failure or timeout. Make redeeming idempotent per order id. Codes are case-insensitive. Check validity windows with a time zone in mind.

Extension question. "Auto-apply the best coupon." Evaluate every eligible coupon's discount against the cart and pick the biggest, and allow stacking only within the stacking rules.

Q85. ⋆ Design an Event Calendar (Flipkart; Google Calendar)

Asked at: Flipkart

- Entities: `User`, `Event` (time range, organiser, participants), `Team`, `Invite` (accept / reject), `Calendar` view.
- Patterns: Observer for invites, and an interval-overlap check for conflicts.
- What interviewers probe: finding a common free slot for a team, recurring events, and the per-user agenda for a day. → *Full design:* Ch 21 — 21.3 Design an Event Calendar (Flipkart)

Q86. ⋆ Design a LeetCode-like coding platform ("Coding Blox", Flipkart)

Asked at: Flipkart

- Entities: `User` (score), `Contest`, `Question` (difficulty level, score), `Submission`, `Leaderboard`.
- Patterns: Strategy for scoring and question selection, Observer for leaderboard updates.
- What interviewers probe: contest states (created, running, ended), only the creator can start a contest, and ordering the leaderboard. → *Full design:* Ch 21 — 21.4 Design a LeetCode-like Coding Platform

Q87. ⋆ Design a property listing site ("PropertyHunt", ClearTrip)

Asked at: ClearTrip

- Entities: `User`, `Property` (location, price, size, type, status), `Shortlist`, `SearchQuery`.
- Patterns: Specification/Builder for search filters, Strategy for sorting.
- What interviewers probe: combining many optional filters cleanly, and marking a property sold so it leaves search results. → *Full design:* Ch 21 — 21.5 Design a Property Listing Site

Q88. ⋆ Design a loan ledger ("Ledger Co", Navi)

Asked at: Navi

- Entities: `Loan` (principal, rate, years, EMI), `Payment` (lump sum after EMI n), `Bank`, `LedgerService`, and commands `LOAN`, `PAYMENT` and `BALANCE`.
- Patterns: Command for each input line, and pure functions for the EMI and balance maths.
- What interviewers probe: rounding EMIs up, how a lump sum shortens the remaining EMIs, and parsing the input file robustly. → *Full design:* Ch 21 — 21.6 Design a Loan Ledger

Q89. ⋆ Design a Task Management System / Jira / task planner

- Entities: `Task` (title, assignee, status, due date, priority), `Story`, `Bug`, `Sprint`, `User`, `TaskService`.
- Patterns: State for status transitions, Factory for task types, Observer for assignment notifications.
- What interviewers probe: valid transitions per task type, subtasks, and filtering or sorting by many fields. → *Full design:* Ch 21 — 21.1 Design a Task Management System

Q90. ⋆ Design a To-Do list app

Asked at: PhonePe

- A smaller task manager: `TodoItem` (text, done, due date, tags), `TodoList`, add / edit / complete / delete, filters, and an activity log.
- Patterns: Command for actions (gives undo and the activity log for free), Strategy for sorting.
- What interviewers probe: soft delete, undo, and recurring tasks. → *Related design:* Ch 21 — 21.1 Design a Task Management System

Q91. ⋆ Design a Digital Wallet (Paytm / PhonePe)

- Entities: `User`, `Wallet` (balance), `Transaction` (credit, debit, transfer), `Offer`, `Statement`.
- Patterns: Strategy for offers and cashback, Command or a ledger of immutable transactions.
- What interviewers probe: transfers between two wallets without deadlock (lock in id order), money as whole paise, and idempotent transfers. → *Full design:* Ch 21 — 21.2 Design a Digital Wallet Service

Q92. ⋆ Design an Inventory / Warehouse Management System

- Entities: `Product` / `SKU`, `Warehouse`, `StockLevel`, `StockMovement` (in, out, transfer), `Reservation`, `ReorderRule`.

- Patterns: Observer for low-stock alerts, Strategy for picking a warehouse, and reserve-then-commit for orders.
- What interviewers probe: preventing overselling under concurrent orders, and keeping a movement log so stock can be audited. → *Full design:* Ch 21 — 21.7 Design an Inventory Management System

Part B — Concept questions

These are the "explain X" and "X vs Y" questions that come up in a low-level design (LLD) round, usually in the first 10 minutes or while you are coding. Answer in two to four sentences, then give a tiny example from your design.

OOP basics

Q93. ⋆ What are the four pillars of OOP? Explain each with an example. **Encapsulation** keeps an object's data private and changes it only through methods that check the rules (a `BankAccount` whose `withdraw` refuses overdrafts). **Abstraction** shows *what* something does and hides *how* (a `PaymentGateway` interface with `pay()`). **Inheritance** lets a subclass reuse and specialise a parent (`SavingsAccount extends Account`). **Polymorphism** means one call runs different code depending on the real object type (`account.monthlyFee()`). → *More:* Ch 02 — 2.3–2.6 The four pillars

Q94. ⋆ Encapsulation vs abstraction: what is the difference? Encapsulation is about *protecting state*: fields are private and every change goes through a method that keeps the object valid. Abstraction is about *hiding complexity*: callers depend on a simple contract (an interface or a small public API) and never see the details. Encapsulation is done inside one class, while abstraction is a design decision about what to expose at all. → *More:* Ch 02 — 2.3 Encapsulation, 2.4 Abstraction

Q95. ⋆ Compile-time vs runtime polymorphism; overloading vs overriding. **Overloading** means several methods share one name but have different parameter lists. The compiler picks one from the declared argument types, so this is compile-time (static) polymorphism. **Overriding** means a subclass redefines an inherited instance method with the same signature. The JVM picks the version from the object's actual class while the program runs (dynamic dispatch), so this is runtime polymorphism. Return types may be covariant (a subtype) when overriding, and the override cannot reduce visibility or throw broader checked exceptions. → *More:* Ch 02 — 2.7 Overloading vs overriding

Q96. Can you override a static or a private method? No. A `static` method with the same signature in a subclass *hides* the parent's method, and the call is chosen by the declared type, not the object. A `private` method is invisible to subclasses, so a method with the same name there is a brand-new method. `final` methods cannot be overridden either. → *More:* Ch 02 — 2.7 Overloading vs overriding

Q97. ⋆ Interface vs abstract class: when do you use which? Use an **interface** for a capability that unrelated classes can share (`Comparable`, `PaymentMethod`). A class can implement many interfaces, and an interface has no instance state. Use an **abstract class** when closely related classes share state or a fixed algorithm skeleton (constructor, protected fields, Template Method). A class can extend only one class. In LLD rounds, default to interfaces and add an abstract base class only when real shared code appears. → *More:* Ch 02 — 2.9 Abstract classes vs interfaces

Q98. What are default, static and private methods in interfaces? A **default** method (Java 8) has a body that implementing classes inherit, so you can add methods to an interface without breaking existing implementers. A **static** interface method is a helper called as `Shape.unit()` and is not inherited. **Private** interface methods (Java 9) share code between default methods. Interfaces still cannot hold instance fields, only constants. → *More:* Ch 02 — 2.8 Interfaces

Q99. Can an abstract class have a constructor? Can it be instantiated? It can have constructors, and they run when a subclass is created (via `super(...)`) to initialise the shared fields. You cannot write `new` on it directly, although an anonymous subclass works. Since Java 25 (flexible constructor bodies, JEP 513), a constructor may run statements *before* calling `super(...)`: it can validate or prepare arguments and assign its own fields, but it cannot read fields, call instance methods or otherwise use `this`. → *More:* Ch 02 — 2.2 Classes, objects, constructors

Q100. ⋆ Composition vs inheritance: why "favor composition"? Inheritance ties a subclass to its parent's internals at compile time, so a parent change can break every child (the "fragile base class"), and Java allows only one parent. Composition means an object *holds* collaborators behind interfaces and delegates to them. You can swap them at runtime and combine behaviours without an explosion of subclasses. Use

inheritance only for a true, stable IS-A relationship where the subclass honours the parent's contract. → *More:* Ch 02 — 2.11 Composition over inheritance

Q101. IS-A vs HAS-A. IS-A is inheritance or interface implementation: a `Car` *is a* `Vehicle`, so it can be used wherever a `Vehicle` is expected. HAS-A is a field reference (association, aggregation or composition): a `Car` *has an* `Engine`. If you can't say "every X is a Y and can stand in for it" truthfully, use HAS-A. → *More:* Ch 03 — 3.8 Decision guide

Q102. ★ What is the diamond problem? How does Java handle multiple inheritance? The diamond problem appears when class D inherits the same method from two parents B and C that both override A's version, so it is unclear which one D gets. Java avoids it for state by allowing only one superclass. With interfaces, if two default methods clash, the class *must* override the method and may pick one with `B.super.method()`. A class method always beats an interface default. → *More:* Ch 02 — 2.8 Interfaces

Q103. ★ Explain the `equals()` and `hashCode()` contract. Two objects that are `equals` must return the same `hashCode`. The reverse is not required, because collisions are allowed. Hash collections find the bucket with `hashCode` and then compare with `equals`. If you override only `equals`, a `HashSet` will hold "duplicates" and `map.get` will miss keys. Build both from the same fields, never from mutable fields of a key while it sits in a map, and use records for value objects. → *More:* Ch 02 — 2.12 equals/hashCode contract

Q104. ★ How do you make a class immutable? What do records give you? Make the class `final`, all fields `private final`, set them once in the constructor, provide no setters, and make defensive copies of mutable inputs and outputs (`List.copyOf`). A **record** gives final fields, a canonical constructor, accessors, `equals`, `hashCode` and `toString` in one line. You still copy mutable components in a compact constructor.

```
record Order(String id, List<String> items) {
    Order { items = List.copyOf(items); }   // defensive copy, also rejects nulls
}
```

→ *More:* Ch 02 — 2.13 Immutability and records

Q105. Explain the access modifiers, and `static` vs `final`. `private` means this class only. *No modifier* (package-private) means the same package. `protected` means the package plus subclasses. `public` means everyone. `static` means "belongs to the class, not to each object" (one shared counter, a factory method). `final` means "cannot change": a final variable is assigned once, a final method cannot be overridden, a final class cannot be extended. They are independent, so `static final` is a class-level constant. → *More:* Ch 02 — 2.2 and 2.14

Q106. ★ Coupling vs cohesion. Cohesion measures how closely the parts of one class belong together. Aim high: one clear job. **Coupling** measures how much one class depends on others' details. Aim low: depend on small interfaces. High cohesion plus low coupling means a change stays in one place and is easy to test. These two yardsticks sit behind almost every principle and pattern. → *More:* Ch 02 — 2.16 Coupling and cohesion

Relationships and UML

Q107. ★ Association vs aggregation vs composition. Association means two classes know about each other (`Doctor` – `Patient`). **Aggregation** is a weak whole–part link where the part can live on its own and be shared (`Team` has `Player` s). **Composition** is a strong whole–part link where the part is created and destroyed with the whole and belongs to only one owner (`Order` owns its `OrderLine` s). In UML, aggregation is a hollow diamond and composition is a filled diamond on the whole's side. → *More:* Ch 03 — 3.2–3.4

Q108. Dependency vs association. A **dependency** is a short-lived use: the class receives another as a method parameter, local variable or return value (a dashed arrow in UML). An **association** is a lasting structural link stored in a field (a solid line). If deleting the field removes the link, it was an association. → *More:* Ch 03 — 3.5 Dependency

Q109. What is UML? Class vs sequence vs activity vs state diagrams. UML (Unified Modeling Language) is a standard notation for drawing software designs. **Class diagrams** show static structure: classes, fields, methods, relationships. **Sequence diagrams** show which objects call which, in time order, for one scenario. **Activity diagrams** are flowcharts of a process. **State machine diagrams** show an object's states and the events that move it between them, as for a vending machine or an order. → *More:* Ch 03 — 3.9–3.14

Q110. Draw a class diagram (or use-case diagram) for X. Start with a **use-case diagram** only if actors matter: actors as stick figures, use cases as ovals ("Park vehicle", "Pay fee"). For the class diagram, list the nouns from the requirements as classes, put two or three key fields and methods in each, and draw only the relationships that carry meaning, with multiplicities (`1` , `0..*`). Mark interfaces and enums. Keep it to what you will actually code. → *More:* Ch 03 — 3.10 Use-case, 3.11 Class diagrams

Q111. ⋆ How do you approach an LLD interview, and how much UML should you draw? First clarify the requirements and agree on the scope (core use cases, what is out of scope). Then find the entities, then draw a lean class diagram with the key interfaces. Walk through one main flow, and only then write code with the patterns that fit. Add concurrency and extensibility last. Draw just enough UML to agree on the shape, usually one class diagram and maybe one sequence or state sketch. Interviewers judge the reasoning and the code, not notation perfection. → *More:* Ch 01 — 1.4 The step-by-step answer framework

Q112. ⋆ LLD vs HLD: what is the difference? HLD (high-level design) chooses the services, databases, caches, queues and how they scale across machines, with boxes and arrows. LLD (low-level design) works inside one service: classes, interfaces, relationships, patterns, method signatures and thread safety, often ending in working code. HLD asks "which components and how do they scale?", LLD asks "how is this component built cleanly?" → *More:* Ch 01 — 1.1 What this round is

Q113. ⋆ Machine-coding round: how do you structure code in 90 minutes, and what do evaluators look for? They want working, runnable code that shows the core flows end to end with a demo `main` . They also look for clean separation into models, services and strategies, sensible names, input validation with clear errors, and an easy path to the extension they will ask for. Keep the code in memory (maps for storage), get a vertical slice working early, then add features one at a time. Part C below gives a minute-by-minute plan. → *More:* Ch 21 — 21.0 How the machine-coding round works

SOLID

Q114. ⋆ Explain SOLID with examples. **S**ingle Responsibility: one reason to change. **O**pen/Closed: add behaviour with new classes, not by editing tested ones. **L**iskov Substitution: subtypes must work wherever the parent is expected. **I**nterface Segregation: small role-based interfaces. **D**ependency Inversion: high-level code depends on abstractions, not on concrete details. Together they keep changes local and code testable. Apply them where change is likely, not everywhere. → *More:* Ch 04 — SOLID principles

Q115. ⋆ Single Responsibility Principle with an example. A class should have one reason to change, meaning one actor or concern drives its edits. An `Invoice` that calculates totals, prints PDFs and saves to the database changes for finance, design and DBA reasons. Split it into `Invoice` (rules), `InvoicePrinter` and `InvoiceRepository` . The test: describe the class without the word "and". → *More:* Ch 04 — 4.2 SRP

Q116. ⋆ Open/Closed Principle with an example. Code should be open for extension but closed for modification. A `FareCalculator` with a growing `switch (vehicleType)` must be edited for every new type. Instead, define a `PricingStrategy` interface and add `EvPricing` as a new class while the calculator stays untouched. Strategy, Decorator and Factory registries are the usual tools. → *More:* Ch 04 — 4.7 OCP

Q117. ⋆ Liskov Substitution: explain the rectangle/square problem. LSP says a subtype must keep the parent's promises: no stronger preconditions, no weaker postconditions, no surprise exceptions. If `Square extends Rectangle` , then `setWidth(5)` silently changes the height too, so code that expects area = width × height breaks. The fix is not to make Square a subtype of a mutable Rectangle. Use a common `Shape` interface or immutable shapes instead. → *More:* Ch 04 — 4.8 LSP

Q118. Interface Segregation with an example. Clients should not depend on methods they don't use. A fat `Machine` interface with `print` , `scan` and `fax` forces a basic printer to throw `UnsupportedOperationException` . Split it into `Printer` , `Scanner` and `Fax` , and let a multifunction device implement all three. → *More:* Ch 04 — 4.9 ISP

Q119. ⋆ Dependency Inversion vs Dependency Injection vs Inversion of Control. **DIP** is a principle: high-level policy and low-level details both depend on an abstraction that the high-level side owns (`OrderService` → `PaymentGateway` interface ← `StripeGateway`). **Dependency Injection** is a technique: dependencies are handed in from outside, usually through the constructor, instead of being created with `new` . **IoC** is the broader idea that a framework or container calls your code and wires objects ("don't call us, we'll call you"). DI is one way to do IoC, and both help you apply DIP. → *More:* Ch 04 — 4.10 DIP

Q120. Can SOLID be overdone? How do you use it in an interview? Yes. Interfaces with a single implementation that will never vary, factories for trivial objects and ten-class hierarchies for a two-case `if` are over-engineering. In an interview, apply SOLID at the points the interviewer is likely to

extend (pricing, notification channels, strategies) and say why. Leave stable parts simple. → *More:* Ch 04 — 4.15 Over-engineering

Design principles and clean code

Q121. ⋆ What are DRY, KISS and YAGNI? DRY (Don't Repeat Yourself): each piece of knowledge, such as a rule or formula, lives in one place. **KISS** (Keep It Simple): choose the simplest design that meets the need. **YAGNI** (You Aren't Gonna Need It): don't build features or hooks for imagined future needs. They balance each other, because too much DRY or flexibility hurts simplicity. → *More:* Ch 05 — §1–3

Q122. Law of Demeter and Tell, Don't Ask. The **Law of Demeter** says a method should talk only to its direct friends (its fields, its parameters, objects it creates), not to strangers through chains like `order.getCustomer().getWallet().debit(x)`. **Tell, Don't Ask** says tell an object to do the work (`order.pay(x)`) instead of pulling its data out and deciding for it. Both keep logic next to its data and cut coupling. → *More:* Ch 05 — §4–5

Q123. What does "program to an interface, not an implementation" mean? Declare fields, parameters and return types with the abstraction (`List` , `NotificationSender`), and create the concrete class in one place (a factory or the composition root). Callers then work with any implementation, including test fakes, without changes. → *More:* Ch 05 — §7

Q124. What does "encapsulate what varies" mean? Find the parts of the design that are likely to change, such as pricing rules, payment providers or eviction policies. Hide each one behind an interface so that the stable code around it doesn't change when they do. Most behavioural patterns (Strategy, State, Observer) are this idea applied to one kind of change. → *More:* Ch 05 — §8

Q125. What are the GRASP principles? GRASP (General Responsibility Assignment Software Patterns) are nine guidelines for deciding which class gets which job. **Information Expert**: give the job to the class that has the data. **Creator**: the class that holds or aggregates B creates B. **Controller**: one object receives system events. **Low Coupling** and **High Cohesion**. **Polymorphism**. **Pure Fabrication**: invent a service class when no domain class fits. **Indirection**. **Protected Variations**. → *More:* Ch 05 — §16 GRASP

Q126. DDD basics: entity vs value object, and what is an aggregate? An **entity** has an identity that lasts while its attributes change (a `Customer` with an ID). A **value object** is defined only by its values, is immutable and is compared by value (`Money` , `Address`). An **aggregate** is a cluster of objects changed as one unit through a single **root** (`Order` with its lines), which protects the cluster's rules. Outsiders hold references only to the root. → *More:* Ch 05 — §17 DDD essentials

Q127. What does "fail fast" mean? Detect invalid input or state as early as possible and stop with a clear exception, instead of carrying bad data forward until it fails somewhere confusing. Examples are `Objects.requireNonNull` in constructors, validating in a record's compact constructor and rejecting an invalid state transition right away. → *More:* Ch 05 — §13 Fail Fast

Q128. ⋆ What are anti-patterns? What is a God class? What is over-engineering? An **anti-pattern** is a common "solution" that causes more problems than it solves. A **God class** knows and does too much (a `ParkingLotManager` with 40 methods), so every change touches it and it can't be tested in isolation. **Over-engineering** adds layers, patterns and configurability nobody asked for. Other anti-patterns are Singleton abuse, anemic models, magic numbers and copy-paste code. → *More:* Ch 08 — 8.6 Anti-patterns

Q129. ⋆ What are code smells? Name a few and their refactorings. A **code smell** is a surface sign of a deeper design problem. **Long Method** → Extract Method. **Large Class** → Extract Class. **Long Parameter List** → Introduce Parameter Object. **Switch on type** → Replace Conditional with Polymorphism. **Feature Envy** (a method that uses another class's data more than its own) → Move Method. **Duplicate Code** → Extract and reuse. **Refactoring** means changing structure in small, behaviour-preserving steps, with tests as the safety net. → *More:* Ch 07 — 7.2 Code smells, 7.4 Lookup table

Q130. What makes code "clean" (naming, functions, error handling)? Names reveal intent (`activeBookings` , not `list2`). Functions are small, do one thing at one level of abstraction and take few arguments. Errors are reported with exceptions that carry context, not with error codes or `null` . Comments explain *why*, not *what*. Clean code reads like the problem description. → *More:* Ch 06 — Clean code

Q131. Checked vs unchecked exceptions: which do you use? Checked exceptions extend `Exception` (but not `RuntimeException`). The compiler forces callers to catch or declare them, which suits recoverable conditions from outside the program, such as `IOException` . **Unchecked** exceptions extend `RuntimeException` and usually signal programming errors or broken rules (`IllegalArgumentException` , a custom `SeatAlreadyBookedException`).

Modern Java design prefers unchecked domain exceptions with clear names, plus checked ones only where the caller can really recover. → *More:* Ch 06 — §6 Error Handling

Q132. Unit testing: mocks vs stubs, and how do you make a design testable? A **stub** returns canned answers so the code under test can run. A **mock** also *verifies* interactions (for example, "`send` was called once"). A **fake** is a working lightweight version, such as an in-memory repository. Designs are testable when dependencies are injected through interfaces, time and randomness are injected (`Clock`), and logic sits apart from I/O. → *More:* Ch 06 — §17 Testing strategies

Q194. ⋆ Layered vs hexagonal architecture: what is the difference? Both separate concerns. In **layered** architecture (controller → service → repository) the domain depends *down* on persistence, so storage tends to shape the model. In **hexagonal** (ports and adapters) the core owns interfaces such as `OrderRepository`, and adapters (JDBC, in-memory, REST) depend *in* on them, so the core imports no framework and tests run with in-memory adapters. Layered suits CRUD; hexagonal pays off when the rules are valuable. → *More:* Ch 11 — 11.9 Architectural patterns

Q195. Event-driven architecture: broker vs mediator topology? Broker (choreography): no coordinator; each consumer reacts to events and may publish new ones. Very decoupled, but no single place shows the whole flow. **Mediator** (orchestration): a central coordinator sends ordered commands and handles failures and compensation. Use a broker for independent reactions (email, analytics) and a mediator for an ordered process with undo (order fulfilment). Either way, delivery is usually at least once, so consumers must be idempotent. → *More:* Ch 11 — 11.9.5 Event-driven architecture

Q196. Pipe-and-filter vs Chain of Responsibility? In a **pipeline** every filter runs and *transforms* the data, possibly into a new type (`String → fields → command`); a rejection is just one possible output. In a **chain of responsibility** each handler may handle or *stop* the request, and usually one handler deals with it. Decorator differs from both: it wraps one object behind the *same* interface. → *More:* Ch 11 — 11.9.6 Pipe-and-filter

Q197. ⋆ Modular monolith or microservices for a new system? Start with a **modular monolith**: one deployable split into business modules with small public interfaces, which gives clear boundaries without network failures. Extract a module into a **microservice** only for a concrete reason (independent team releases, very different scaling, fault isolation), and accept the costs: timeouts and retries, idempotency, sagas instead of cross-service transactions, eventual consistency. Services that must be deployed together are a "distributed monolith", the worst of both. → *More:* Ch 11 — 11.9.4 Monolith, modular monolith and microservices

Creational patterns

Q133. ⋆ What is a design pattern? What are the three GoF categories? A design pattern is a named, reusable solution to a recurring design problem. It is a template you adapt, not a library you import. The "Gang of Four" (GoF) book lists 23 patterns in three groups. **Creational** patterns deal with how objects are made (Singleton, Factory Method, Abstract Factory, Builder, Prototype). **Structural** patterns deal with how objects are composed (Adapter, Bridge, Composite, Decorator, Facade, Flyweight, Proxy). **Behavioral** patterns deal with how objects communicate and share work (Strategy, Observer, State, Command and others). → *More:* Ch 08 — 8.1–8.3

Q134. ⋆ Singleton: how do you make it thread-safe? Use an **enum** (simplest, and safe against reflection and serialization), or the **holder idiom**: a static nested class that the JVM loads, thread-safely, on first use. **Double-checked locking** also works but needs a `volatile` field. Without `volatile`, another thread can see a half-constructed object.

```
final class Config {
    private Config() {}
    private static class Holder { static final Config INSTANCE = new Config(); }
    static Config get() { return Holder.INSTANCE; }
}
```

→ *More:* Ch 08 — 8.8 Singleton

Q135. How can a Singleton be broken? Reflection can call the private constructor. Deserialization creates a new instance unless you add `readResolve`. Cloning breaks it if the class is `Cloneable`. Each class loader can also hold its own copy. An enum singleton blocks the first three. → *More:* Ch 08 — 8.8 Singleton

Q136. ⋆ Why is Singleton often called an anti-pattern? It is global mutable state in disguise. It hides dependencies (callers reach out with `getInstance()` instead of receiving them), it makes unit tests hard because you can't swap or reset it, and it creates concurrency hot spots. Prefer creating one instance at the composition root (or a DI container's singleton scope) and injecting it. → *More:* Ch 08 — 8.8 Singleton

Q137. ⋆ Simple Factory vs Factory Method vs Abstract Factory. A **Simple Factory** is one method with a `switch` that returns the right subtype (`VehicleFactory.create(type)`). It is an idiom, not a GoF pattern. **Factory Method** lets subclasses decide which product to create by overriding a creation method. **Abstract Factory** creates *families* of related products that must match (a `WindowsUiFactory` makes `WinButton` + `WinCheckbox`). In interviews, a simple factory or a `Map<Type, Supplier<T>>` registry is usually enough. → *More:* Ch 08 — 8.9–8.11, 8.15

Q138. ⋆ When would you use Builder? Use it when an object has many parameters, several of them optional (the "telescoping constructor" problem), or when you want to validate everything before building an immutable object. The builder collects values with readable chained calls and `build()` checks them once. Records with a few fields don't need one.

```
var pizza = Pizza.builder().size(LARGE).cheese(true).topping("olive").build();
```

→ *More:* Ch 08 — 8.12 Builder

Q139. Prototype pattern; shallow vs deep copy. Prototype creates new objects by copying a configured existing one, which helps when construction is expensive or the configuration is complex. A **shallow copy** copies fields as they are, so the copy shares nested mutable objects with the original. A **deep copy** also copies those nested objects, so changes don't leak between them. In Java, prefer copy constructors or factory methods over `clone()` . → *More:* Ch 08 — 8.13 Prototype

Q140. What is an object pool, and when is it worth it? A pool keeps a fixed set of expensive, reusable objects (database connections, threads). It lends one out and takes it back, so you avoid repeated creation and cap the total. It needs a thread-safe borrow/return (often a `BlockingQueue` or a `Semaphore`) and must reset each object's state on return. For cheap objects it only adds bugs. → *More:* Ch 08 — 8.14 Object Pool

Structural patterns

Q141. ⋆ Adapter: when would you use it? Adapter vs Facade vs Bridge. An **Adapter** converts one interface into another that the client expects. A typical use is wrapping a third-party `LegacyPaymentApi` so it implements your `PaymentGateway` . A **Facade** gives a *simpler* front to a whole subsystem (`HomeTheater.watchMovie()`) without changing the interfaces. A **Bridge** splits an abstraction from its implementation up front, so the two can vary independently (`Shape` × `Renderer`). Adapter fixes a mismatch after the fact, Facade simplifies, and Bridge is planned in from the start. → *More:* Ch 09 — 9.1, 9.2, 9.5, 9.8

Q142. ⋆ Decorator vs Proxy (and vs inheritance). Both wrap an object behind the same interface. A **Decorator** *adds* behaviour, and decorators can be stacked at runtime (`new Logging(new Compressing(stream))`), which avoids a subclass for every combination. A **Proxy** *controls access* to the same behaviour: lazy loading, caching, security checks, remote calls. Usually the proxy, not the client, decides whether to create the real object. → *More:* Ch 09 — 9.4 Decorator, 9.7 Proxy

Q143. ⋆ Composite pattern (file system / org chart). Composite treats single objects and groups of objects through one interface, so the client doesn't care which it has. A `File` and a `Directory` both implement `FileSystemNode.size()` , and the directory sums its children recursively. Use it for any tree: menus, org charts, expression trees. → *More:* Ch 09 — 9.3 Composite

Q144. What is the Flyweight pattern? Flyweight saves memory when there are huge numbers of similar objects. It shares the *intrinsic* (unchanging, shareable) state and passes the *extrinsic* (per-use) state in from outside. Examples: one shared `Glyph` per character in a text editor, with the position stored separately, or tree types shared across a million trees in a game. `Integer.valueOf` caching small values is a JDK example. → *More:* Ch 09 — 9.6 Flyweight

Q145. How does Java support Proxy at runtime (dynamic proxies, Spring AOP)? `java.lang.reflect.Proxy` creates, at runtime, a class that implements given interfaces and sends every call to an `InvocationHandler` . That is how logging, transactions and retries get added around calls. Spring AOP builds either a JDK proxy (interfaces only) or a generated subclass (CGLIB). Plain Spring uses a JDK proxy when the bean has an

interface, while Spring Boot defaults to CGLIB subclasses for everything. Either way the advice lives in the proxy, so a `@Transactional` method called from inside the same class (through `this`) bypasses it. → *More:* Ch 09 — 9.7 Proxy

Behavioral patterns

Q146. ⋆ Strategy pattern: where did you use it? Strategy puts each interchangeable algorithm behind an interface and lets the context pick one at runtime. Examples are a parking fee per vehicle type, Splitwise's equal/exact/percent splits, cache eviction (LRU/LFU) and payment methods. It removes `switch` statements and follows Open/Closed. In modern Java a strategy is often just a lambda.

```
interface SplitStrategy { Map<User, Money> split(Money total, List<User> users); }
```

→ *More:* Ch 11 — 11.3 Strategy

Q147. ⋆ Strategy vs State. They have the same class shape but a different intent. With **Strategy**, the *client* picks the algorithm, and strategies don't know about each other. With **State**, the *object itself* changes behaviour as its internal state changes, and the states usually trigger the transitions (`IdleState` → `HasCoinState`). Ask "who chooses?": the caller means Strategy, the object's own lifecycle means State. → *More:* Ch 11 — 11.8 Choose the right behavioral pattern

Q148. Strategy vs Template Method. Template Method fixes the steps of an algorithm in a base class and lets subclasses fill in some of them (inheritance, chosen at compile time). **Strategy** swaps the whole algorithm through composition, at runtime. Prefer Strategy when variants change independently or need to be combined. Template Method suits a stable skeleton with small hooks. → *More:* Ch 11 — 11.4 Template Method

Q149. ⋆ Observer pattern; Observer vs Pub-Sub. In **Observer**, a subject keeps a list of observers and calls them directly when its state changes (a stock price notifying displays). The subject knows its observers' interface, and delivery is usually synchronous and in-process. **Pub-Sub** puts a broker or event bus between them: publishers send to a topic without knowing the subscribers, and delivery can be asynchronous and cross-process. Pub-Sub decouples more, at the cost of extra infrastructure. → *More:* Ch 11 — 11.1 Observer

Q150. ⋆ Command pattern: how does it give you undo/redo? Command turns a request into an object with `execute()` and `undo()`, so it can be queued, logged, retried or reversed. For undo/redo, push each executed command on an undo stack. Undo pops it, calls `undo()` and pushes it on the redo stack. Any new command clears the redo stack. Text editors, remote controls and job queues use it. → *More:* Ch 10 — 10.2 Command

Q151. Chain of Responsibility (middleware, approval flows). A request passes along a chain of handlers. Each one either handles it or passes it to the next, so the sender doesn't know who will deal with it. Examples are expense approval (manager → director → CFO by amount), ATM cash dispensing by note value, and servlet filters or HTTP middleware (auth → rate limit → logging). Handlers can be added or reordered without changing the client. → *More:* Ch 10 — 10.1 Chain of Responsibility

Q152. Iterator, Mediator, Memento, Interpreter, Visitor: when would you use each? Iterator: walk a collection without exposing how it is stored. **Mediator**: objects talk through a hub instead of to each other (an air-traffic control tower, a chat room). **Memento**: save and restore an object's state without breaking encapsulation (snapshots for undo). **Interpreter**: evaluate sentences of a small grammar (rule expressions). **Visitor**: add new operations to a stable class hierarchy without editing it (export or tax calculation over node types). In modern Java, sealed types with pattern-matching `switch` often replace Visitor. → *More:* Ch 10 — 10.3–10.6, Ch 11 — 11.5 Visitor

Q153. ⋆ State pattern (vending machine / ATM / elevator). Each state becomes a class that implements the same interface (`insertCoin`, `selectItem`, `dispense`). The machine delegates to its current state object, and states switch the machine to the next state. This replaces large `if (state == ...)` blocks, makes illegal actions explicit per state, and lets you add a state without touching the others. A small enum-based state machine is fine when the states have little behaviour. → *More:* Ch 11 — 11.2 State

Q154. What are Null Object, Repository, Specification and MVC? Null Object: a do-nothing implementation (`NoDiscount`) used instead of `null`, which removes null checks. **Repository**: a collection-like interface for storing and finding aggregates (`BookingRepository.findById`) that hides the storage. **Specification**: a business rule as a combinable object (`isAvailable.and(inCity)`). **MVC** (Model–View–Controller): the model holds data and rules, the view displays it, and the controller turns input into model updates. → *More:* Ch 11 — 11.7 Beyond GoF

Q155. ⋆ Which patterns would you use for a parking lot, an elevator or a vending machine? Parking lot: Strategy for pricing and spot allocation, Factory for vehicles or spots, Singleton or injected for the lot, Observer for display boards. **Elevator**: State for each car (moving/idle/doors open), Strategy for the dispatch algorithm, Command for requests. **Vending machine**: State for idle/has-money/dispensing, and Chain of Responsibility for giving change. Name a pattern only when you can say what change it absorbs. → *More:* Ch 22 — 22.9 Case-study map

Q156. How have lambdas and modern Java changed classic patterns? Patterns with a single method, such as Strategy, Command, Observer callbacks and Template Method hooks, are often just functional interfaces with lambdas (`Comparator`, `Runnable`, `Function`). Records cover simple value objects, and sealed interfaces with pattern-matching `switch` give exhaustive case handling in place of many Visitor uses. The *ideas* stay the same, with less boilerplate. → *More:* Ch 11 — 11.6 Quick-reference summary

Concurrency

Q157. ⋆ Process vs thread. A **process** is a running program with its own memory space. Processes are isolated and costly to create, and they communicate through IPC (inter-process communication: pipes, sockets). A **thread** is a path of execution *inside* a process. Threads share the heap with their siblings but each has its own stack, so they are cheap to switch and communicate through shared memory, which is why they need synchronisation. → *More:* Ch 12 — 12.2

Q158. ⋆ Concurrency vs parallelism. Concurrency means dealing with many tasks over the same period by interleaving them. It is about structure, and it works on one core. **Parallelism** means literally running tasks at the same instant on several cores. It is about speed. A web server handling 10,000 waiting connections is concurrent, while a multi-core merge sort is parallel. → *More:* Ch 12 — 12.1

Q159. What are the states in a thread's lifecycle? `Thread.State` has six values. **NEW** means created but not started. **RUNNABLE** means running or ready to run. **BLOCKED** means waiting to enter a `synchronized` monitor. **WAITING** means waiting without a time limit (`wait()`, `join()`, `LockSupport.park`). **TIMED_WAITING** means waiting with a timeout (`sleep(ms)`, `wait(ms)`). **TERMINATED** means finished. → *More:* Ch 12 — 12.3

Q160. Why call `start()` and not `run()`? `start()` asks the JVM to create a new thread, which then calls `run()`. Calling `run()` directly just executes the method on the *current* thread, so nothing runs concurrently. Calling `start()` twice on the same thread throws `IllegalThreadStateException`. → *More:* Ch 12 — 12.4

Q161. ⋆ Runnable vs Callable; ExecutorService; CompletableFuture. A `Runnable` returns nothing and can't throw checked exceptions. A `Callable<V>` returns a value and may throw. An `ExecutorService` runs tasks on a managed pool and hands back a `Future` for the result. `CompletableFuture` adds non-blocking chaining and combining (`thenApply`, `thenCombine`, `allOf`) and error handling. Always shut executors down. Since Java 19, `ExecutorService` is `AutoCloseable`, so try-with-resources works: `close()` shuts the executor down and waits for submitted tasks to finish. → *More:* Ch 12 — 12.4

Q162. ⋆ What are virtual threads, and when should you not use them? Virtual threads (final in Java 21) are cheap JVM-managed threads. When one blocks on I/O it unmounts from its OS carrier thread, so you can run one per request or task, even millions. They help blocking, I/O-heavy work, not CPU-bound work. Don't pool them, and limit access to scarce resources with a `Semaphore`. Since JDK 24 (JEP 491), blocking inside `synchronized` or in `Object.wait()` no longer pins the carrier. Pinning remains in rare cases: blocking while a native frame is on the stack, and blocking in class loading or a class initializer. → *More:* Ch 12 — 12.4

Q163. ⋆ What is a race condition? What is a critical section? A **race condition** happens when the result depends on the timing of threads touching shared mutable data. A classic case is `count++`, which is read–modify–write: two threads read 5 and both write 6. Another is "check-then-act", where two users both see seat A1 free. A **critical section** is the code that touches the shared state and must run by one thread at a time (or be made atomic). → *More:* Ch 12 — 12.5

Q164. ⋆ What does `volatile` do? What is happens-before (the Java Memory Model)? The Java Memory Model (JMM) defines when one thread's writes become visible to another. **Happens-before** is its guarantee: if A happens-before B, then B sees A's effects. Releasing a lock happens-before the next acquire of it, `Thread.start` and `join` create the edge too, and so do volatile writes and reads. `volatile` gives visibility and ordering for one variable but **not** atomicity, so `volatile int count; count++` is still a race. It is good for flags like `running = false`. → *More:* Ch 12 — 12.6

Q165. ⋆ `synchronized` vs `ReentrantLock`. Both give mutual exclusion and are **reentrant**, meaning the thread holding the lock can enter it again. `synchronized` is simpler and releases automatically. `ReentrantLock` adds

`tryLock` with a timeout, interruptible locking, optional fairness and several `Condition`s, but you must `unlock()` in `finally`. Start with `synchronized` and switch when you need those features.

```
lock.lock();
try { /* critical section */ } finally { lock.unlock(); }
```

→ *More:* Ch 12 — 12.7, 12.8

Q166. ReadWriteLock and StampedLock: when do you use them? A `ReentrantReadWriteLock` lets many readers in at once but a writer gets exclusive access, which helps when reads far outnumber writes (a config cache). `StampedLock` adds **optimistic reads**: read without locking, then `validate(stamp)`, and retry under a real read lock if a write happened. It is faster for read-heavy data but not reentrant and harder to use correctly. → *More:* Ch 12 — 12.10

Q167. ★ Semaphore vs mutex. A **mutex** (lock) lets exactly one thread in and has an owner: only the locker should unlock. A **semaphore** holds N permits, so up to N threads proceed at once (limiting a pool to 10 connections). It has no owner, so one thread can release a permit another acquired, which is useful for signalling. A binary semaphore looks like a mutex but lacks ownership and reentrancy. → *More:* Ch 12 — 12.12 Semaphores

Q168. `wait()` / `notify()` vs `Condition`; `notify()` vs `notifyAll()`. `wait()` releases the object's monitor and sleeps until notified. It must be called while holding the monitor (`synchronized`), and always inside a `while` loop that re-checks the condition, because of spurious wakeups. `notify()` wakes one arbitrary waiter, `notifyAll()` wakes all. `notifyAll` is safer when waiters wait for different conditions. A `Condition` from a `Lock` does the same job, but one lock can have several conditions (`notFull`, `notEmpty`), so you wake only the right group. → *More:* Ch 12 — 12.13 Condition variables

Q169. ★ `wait()` vs `sleep()`. `Object.wait()` **releases** the monitor, must be called inside `synchronized` (otherwise it throws `IllegalMonitorStateException`), and returns on `notify` / `notifyAll`, a timeout, an interrupt or a spurious wakeup. Use it to wait for a condition. `Thread.sleep()` is a static method that pauses the current thread for a time and **keeps** any locks it holds. Use it for delays, not coordination. Both throw `InterruptedException`. → *More:* Ch 12 — 12.3, 12.13

Q170. ★ `CountDownLatch` vs `CyclicBarrier`. A `CountDownLatch` lets threads wait until a count reaches zero. It is one-shot, and the threads counting down are usually different from the waiters ("start when all 5 services are ready"). A `CyclicBarrier` makes a fixed group of threads wait for *each other* at a point and can be reused round after round, with an optional barrier action ("each simulation step"). `Phaser` is the flexible, dynamic version. → *More:* Ch 12 — 12.14

Q171. ★ What is CAS? How do atomics like `AtomicInteger` work? CAS (compare-and-swap) is a CPU instruction: "set the value to new **only if** it still equals expected", which is atomic and lock-free. Atomic updates are a CAS loop: read, compute, CAS, and retry if another thread changed the value (`updateAndGet` works exactly like this, and on x86 the JIT turns `incrementAndGet` into a single atomic add instruction). That beats locks under low contention, and `LongAdder` spreads hot counters across cells under high contention. Watch for the **ABA problem** (a value changes A→B→A unnoticed). `AtomicStampedReference` solves it with a version stamp. → *More:* Ch 12 — 12.15

Q172. ★ How does `ConcurrentHashMap` work? How is it different from `Hashtable` and `synchronizedMap`? `Hashtable` and `Collections.synchronizedMap` lock the whole map for every call. Since Java 8, `ConcurrentHashMap` uses CAS to place the first node in an empty bucket and otherwise locks only that bucket's head node. Reads take no lock at all, and writes to different buckets run in parallel. It forbids `null` keys and values. Its iterators are weakly consistent and never throw `ConcurrentModificationException`. Use `compute` / `merge` / `putIfAbsent` for atomic read-modify-write. → *More:* Ch 12 — 12.16, Ch 13 — 13.7

Q173. ★ What is a deadlock? What are its conditions, and how do you prevent and detect it? A deadlock happens when threads wait on each other forever, for example T1 holds lock A and wants B while T2 holds B and wants A. The four Coffman conditions must all hold: mutual exclusion, hold-and-wait, no preemption and circular wait. To prevent it, break one condition. The most common way is to always acquire locks in a global order (for example by account ID). Others are `tryLock` with a timeout and backoff, or holding fewer locks. To detect it, take a thread dump (`jstack`), use `ThreadMXBean.findDeadlockedThreads()`, or use a wait-for graph. → *More:* Ch 12 — 12.17 Deadlock

Q174. Livelock vs starvation. In a **livelock**, threads are active but make no progress because they keep reacting to each other, like two people stepping aside the same way in a corridor. The fix is random backoff. In **starvation**, a thread never gets the CPU or the lock because others keep winning. The fix is

fair locks (`new ReentrantLock(true)`), bounded queues and avoiding priority tricks. → *More:* Ch 12 — 12.18, 12.19

Q175. ⋆ Explain the producer-consumer problem and how to solve it. Producers put items into a bounded buffer and consumers take them out. Producers must wait when the buffer is full, and consumers must wait when it is empty, without busy-spinning or racing. The easy solution is a `BlockingQueue` (`ArrayBlockingQueue`), whose `put` and `take` block. Hand-rolled versions use a lock with `notFull` / `notEmpty` conditions, or `wait` / `notifyAll` in `while` loops. → *More:* Ch 12 — 12.22, Ch 13 — 13.8

Q176. The reader-writer problem. Many readers may read shared data at once, but a writer needs exclusive access. A naive solution can starve writers while readers keep arriving. In Java, use `ReentrantReadWriteLock` (with the fair option if writer starvation matters), `StampedLock`, or copy-on-write structures when writes are rare. → *More:* Ch 12 — 12.23

Q177. ⋆ How do you prevent double booking of a seat or parking spot? The danger is check-then-act: two requests both see the seat free and both book it. In memory, make the check and the claim one atomic step, for example `seatHolds.putIfAbsent(seatId, userId)` on a `ConcurrentHashMap`, or lock per show or seat (fine-grained), acquiring several seats in sorted order. With a database, use **optimistic locking** (a version column, and the update fails if the version changed, so you retry or report "taken"), **pessimistic locking** (`SELECT ... FOR UPDATE`) or a unique constraint. Add temporary holds with a timeout during payment. → *More:* Ch 12 — 12.25

Q178. Optimistic vs pessimistic locking. Pessimistic locking takes the lock *before* working, assuming conflicts are likely. It is safe but blocks others. **Optimistic** locking works without a lock and checks at commit that nobody changed the data (a version number or CAS), retrying if they did. It is fast when conflicts are rare. Use pessimistic locking for hot, contended rows (the last seats of a flash sale) and optimistic locking for most other updates. → *More:* Ch 12 — 12.25

Q179. Coarse-grained vs fine-grained locking. Coarse-grained means one lock guards a lot (the whole parking lot). It is simple and deadlock-free but serialises everything. **Fine-grained** means many locks guard small parts (per floor, per spot, per bucket). Throughput goes up, but so do complexity and deadlock risk. Start coarse, measure, and split only the hot path. Lock striping is the middle ground. → *More:* Ch 12 — 12.11

Q180. Name the ways to make a class thread-safe. In rough order of preference: make it **immutable** (records, final fields). **Confine** the state to one thread (local variables, `ThreadLocal`, or a single-threaded owner reached via a queue). Use **thread-safe building blocks** (`ConcurrentHashMap`, atomics, `BlockingQueue`). Last, **lock** carefully around every invariant that spans several fields. Document which one you chose. → *More:* Ch 12 — 12.24 Immutability & thread confinement

Java specifics for LLD rounds

Q181. ⋆ Why is `String` immutable in Java? Strings are shared everywhere, so immutability makes that safe. Literals live in the **string pool** and are reused, which would be dangerous if one holder could change them. Immutable strings are automatically **thread-safe**. Their `hashCode` can be **cached**, which makes them fast, stable `HashMap` keys. They are also **secure** for values like file paths, URLs and class names, which cannot be changed after they are checked. "Modifying" methods such as `concat` and `replace` return new strings. → *More:* Ch 02 — 2.13 Immutability

Q182. ⋆ Is Java pass-by-value or pass-by-reference? Always **pass-by-value**. For objects, the value copied is the *reference*. So a method can change the object's fields through its copy (`list.add(x)` is visible to the caller), but reassigning the parameter (`list = new ArrayList<>()`) does not affect the caller's variable. That is why a `swap(a, b)` method can't swap the caller's objects.

```
void reset(StringBuilder sb) { sb.setLength(0); sb = new StringBuilder("x"); }
// the caller's sb is now empty, not "x"
```

→ *More:* Ch 02 — 2.2 Classes and objects

Q183. ⋆ `==` vs `equals()` for strings (and the string pool). `==` compares references (is this the same object?), while `equals` compares contents. `"ab" == "ab"` is true because string literals (and compile-time constants such as `"a" + "b"`) are interned into the pool. `new String("ab") == "ab"` is false, and so is a string built at runtime unless you call `intern()`. Always compare strings with `equals` (or `Objects.equals` for null-safety). → *More:* Ch 02 — 2.12 Identity vs equality

Q184. **`String` vs `StringBuilder` vs `StringBuffer`.** `String` is immutable, so concatenating in a loop creates many temporary objects. `StringBuilder` is mutable and fast but not thread-safe, which makes it the default for building text. `StringBuffer` is the old synchronised version and is rarely needed, since builders are almost always used inside one thread. → *More:* Ch 02 — 2.13 Immutability

Q185. **`final` vs `finally` vs `finalize`.** `final` is a modifier for a constant variable, a non-overridable method or a non-extendable class. `finally` is the block after `try` that always runs, used for cleanup, though try-with-resources is better. `finalize()` was a GC cleanup hook. It has been deprecated for removal since JDK 18 (JEP 421) because it is unpredictable and slow. Use `AutoCloseable` or `java.lang.ref.Cleaner` instead. → *More:* Ch 02 — 2.2 static and final

Q186. ⋆ **How does `HashMap` work internally?** `HashMap` keeps an array of buckets. It spreads the key's `hashCode` bits and masks the result to a bucket index, then compares with `equals` inside the bucket. Collisions form a linked list. When a bucket already holding 8 entries (`TREEIFY_THRESHOLD`) gets another, the list becomes a red-black tree, but only if the table has at least 64 slots (`MIN_TREEIFY_CAPACITY`); a smaller table is resized instead. A tree shrinks back to a list at 6 entries or fewer. When the size exceeds capacity × load factor (0.75 by default), the table doubles and entries are redistributed. It allows one `null` key and is not thread-safe. → *More:* Ch 13 — 13.7 Concurrent HashMap

Q187. **`Comparable` vs `Comparator`.** `Comparable` gives a class its single natural order through `compareTo`, for example `Money` by amount. A `Comparator` is an external, swappable ordering, which makes it a Strategy (`Comparator.comparing(Ride::fare).thenComparing(Ride::eta)`). Keep `compareTo` consistent with `equals`, or sorted sets and maps will behave oddly. → *More:* Ch 11 — 11.3 Strategy

Q188. **Fail-fast vs fail-safe iterators.** **Fail-fast** iterators (`ArrayList`, `HashMap`) throw `ConcurrentModificationException` if the collection is structurally changed during iteration, other than through the iterator's own `remove`. **Fail-safe** (weakly consistent or snapshot) iterators (`ConcurrentHashMap`, `CopyOnWriteArrayList`) never throw, but may not show the latest changes. To remove while looping, use `iterator.remove()` or `removeIf`. → *More:* Ch 12 — 12.16 Concurrent collections

Q189. ⋆ **Sealed classes, records and pattern matching: how do they help in LLD?** A `sealed` interface lists its only allowed implementations (Java 17). Combined with records and pattern-matching `switch` (Java 21), the compiler checks that every case is handled, with no `default` needed. This is a clean way to model closed sets like payment results or vehicle types.

```
sealed interface Payment permits Card, Upi {}
record Card(String last4) implements Payment {}
record Upi(String vpa) implements Payment {}
String describe(Payment p) { return switch (p) {
    case Card c -> "Card " + c.last4();
    case Upi u  -> "UPI " + u.vpa(); }; }
```

→ *More:* Ch 02 — 2.13 Immutability and records

Q190. **Generics: what are wildcards and PECS?** `List<? extends Shape>` accepts lists of any Shape subtype and lets you *read* Shapes, but you can't add to it. `List<? super Circle>` lets you *add* Circles. **PECS** ("Producer Extends, Consumer Super") is the rule: use `extends` for sources you read from and `super` for sinks you write to, as in `Collections.copy(List<? super T> dst, List<? extends T> src)`. Generics are erased at runtime, so `new T()` is impossible. → *More:* Ch 02 — 2.15 Generics basics

Q191. **When should you use `Optional`, and how do you avoid `null` in a design?** Return `Optional<T>` from lookups that may find nothing (`findSpot()`), so callers must handle absence. Don't use it for fields, parameters or collections; return an empty collection instead. Elsewhere, avoid `null` with validated constructors, the Null Object pattern and `Objects.requireNonNull`. → *More:* Ch 11 — 11.7.1 Null Object

Part C — How to run a 60–90 minute machine-coding round

A machine-coding round asks you to build a small working system, such as Splitwise, a parking lot or a ride-sharing service, in one sitting, usually in an IDE. It is judged on whether the code runs, how it is structured, and how easily it absorbs a change. The plan below assumes 90 minutes. For a 60-minute round, roughly halve each step.

Minutes 0–10: clarify and cut scope. Read the problem twice. Write down the must-have features (usually 3–5 commands), the input/output format, and what you will leave out (persistence, auth, UI). Ask about the

tricky rules, such as rounding, ties, invalid input and concurrent users. Confirm that in-memory storage and a console demo are fine.

Minutes 10–20: model. List the entities (nouns) and the services (verbs). Decide the few interfaces where change is likely, such as a pricing, split or allocation strategy. Sketch the packages: `model`, `service`, `strategy`, `repository` (in-memory maps), `exception`, and `Main` for the demo. Don't draw full UML. A short list of classes is enough.

Minutes 20–60: build a vertical slice first. Get the main flow working end to end (create → act → show) before adding anything else. Use records and enums for the models, constructor injection for the services, and clear custom exceptions for invalid operations. Commit or save each working state. Write code that compiles at every step, and don't leave a half-built feature that breaks the demo.

Minutes 60–75: complete the features and handle edge cases. Add the remaining commands one by one. Validate input, handle "not found" and duplicates, and make any shared state thread-safe if concurrency was asked (`ConcurrentHashMap`, per-entity locks).

Minutes 75–85: demo and test. Run `Main` on the sample input plus two or three edge cases. If time allows, add a few unit tests for the core rule, such as the split math or the fare calculation.

Minutes 85–90: walk through and extend. Explain the structure in one minute, then show how the likely extension fits: "a new split type is one new class". Be honest about anything you skipped.

Common ways to lose marks: one giant class, no working demo, over-engineering before the basics run, ignoring input validation, and staying silent. Talk while you code.

→ *More:* Ch 01 — The LLD interview approach, Ch 21 — 21.0 How the machine-coding round works

Sources of the question lists

Question wording and frequency only; every answer above is written for this book.

- LeetCode Discuss interview-experience posts (Flipkart SDE2, Swiggy SDE1/SDE2, Atlassian P4, PhonePe, Meta E5) and the "How do you prepare for machine coding rounds" thread
- prasadgujar/low-level-design-primer — `questions.md` (LeetCode-Discuss-sourced list)
- ashishps1/awesome-low-level-design — "LLD Interview Problems" and concurrency problems
- kumaransg/LLD — machine-coding questions tagged Flipkart / Udaan / Navi / ClearTrip
- donnemartin/system-design-primer — "Object-oriented design interview questions"
- workat.tech machine-coding practice problems
- Course tables of contents (titles only): Grokking the Object-Oriented Design Interview

Chapter 24 — Mock Interview Drills (Study Edition)

> **Goal.** Turn each of the 23 classic object-oriented design problems into a timed, spoken rehearsal. A drill is the interview script only: the prompt, the questions to ask back, a plan, a checklist, the usual follow-ups and a link. The designs themselves are taught in the case-study chapters; this chapter trains you to present them.

An **object-oriented design (OOD) round**, also called a low-level design (LLD) round, is a 45–60 minute interview in which you turn a one-line product idea into classes, relationships and a little code while explaining your choices. It grades the **model** and the **explanation**. These drills train the explanation.

24.0 How to use these drills

Running a drill

1. Read only the prompt and the clarifying questions. Cover the rest.
2. Start a 45-minute timer and talk out loud, sketching on paper and typing the core code.
3. When the timer stops, score yourself with the rubric below **before** looking at anything.
4. Uncover the checklist and follow-ups, then open the linked case study. Re-score, and write down the one dimension you will fix next time.
5. Repeat a drill the next day if you scored under 24.

The timebox

Every drill's "Minute-by-minute plan" uses these phases. Only the problem-specific content is listed in each drill.

45-min round	60-min round	Phase	What exists at the end
0–5	0–6	**Clarify**	6–10 numbered requirements, an out-of-scope line, 2–4 actors (people or systems that use the design)
5–8	6–10	**Use cases**	4–8 actions, one line each
8–18	10–22	**Entities and relationships**	Classes with one sentence of responsibility each; a class diagram with kinds and multiplicities (how many on each side)
18–23	22–28	**Walk one flow**	The main use case traced through the classes
23–35	28–45	**Code**	The 2–3 classes that carry the main flow, written properly; the rest only named
35–45	45–60	**Follow-ups**	Concurrency, one extension, trade-offs

Two rules for every problem. **Say the plan in the first minute**, so the interviewer can steer early. **Have a class diagram by minute 18**; if you don't, cut scope out loud ("refunds later") and move on. Extra time in a 60-minute round goes to code and follow-ups, never to longer clarifying.

How to explain clearly

Narrate **decisions, not actions**. "I'm drawing Ticket" is an action. "Ticket is its own class because it outlives the parking session and payment needs something to point at" is a decision, and decisions are what get graded.

- **Each class:** name, one sentence of responsibility, why it is separate.
- **Each relationship:** kind (composition, aggregation, association; §3.2–3.8) and multiplicity.
- **Each design pattern:** name the **change it absorbs**. "Strategy for pricing, because pricing changes most often and a new rule should be a new class" scores; "I'll use Strategy" alone does not.
- **At each phase boundary**, check in: "That's my class list. Shall I walk the booking flow, or would you rather I start on concurrency?"

Terms the drills use

Pattern and concurrency terms (Strategy, Command, compare-and-set, idempotent, half-open interval …) are defined in the glossary, §22.11. In the plans, "`A` owns many `B`" means composition, one to many, and "walk" means trace one use case through the classes.

Scoring rubric

Score each dimension 1–4 (total out of 28).

Dimension	1 — Weak	2 — Partial	3 — Solid	4 — Strong
Requirements	Drew classes at once	Asked questions, wrote nothing	Numbered functional list plus out of scope	Also non-functional (concurrency, failure), confirmed with the interviewer
Entities	One "Manager" does everything	Right nouns, blurred jobs	One-sentence responsibility per class	Also the non-obvious split (type vs instance, book vs copy)
Relationships	Unlabelled lines	Arrows without kind or multiplicity	Correct kinds and multiplicities	Also inheritance only where substitution holds, with reasons
Patterns	None, or named for show	Named, change not stated	One or two, each tied to the change it absorbs	Also one considered and rejected, with the reason
Concurrency	Not mentioned	"Add synchronized" with no target	Names the shared resource; locks at the right grain	Also the failure mode: hold expiry, lock order, idempotent retry
Extensibility	A new type edits many classes	Claims extensibility without showing it	Shows one extension as "one new class"	Also says what would *not* extend cheaply and why that is acceptable
Communication	Silent or rambling	Narrates actions	Narrates decisions, checks in	Also keeps to time and invites steering

Reading the total: treat 24–28 as interview-ready. 18–23 is borderline, usually dragged down by Requirements or Communication. Below 18, study the case and repeat the drill tomorrow.

Suggested order

Small state machines first, platforms last: Parking Lot, Vending Machine, ATM, Elevator, Amazon Locker → Library, Movie Tickets, Hotel, Car Rental, Airline, Restaurant, Meeting Scheduler → Blackjack, Chess, Jigsaw → Amazon Shopping, Stack Overflow, Facebook, LinkedIn, ESPNcricinfo, Stock Brokerage → Ride Sharing and Food Delivery with no notes at all. The sections below follow the original course order so you can find a problem quickly.

24.1 Drill — Parking Lot

The prompt as the interviewer says it

> "Design a parking lot. Cars come in, get a ticket, park, and pay when they leave. Take it wherever you like."

The open ending tests whether you bound the problem.

Clarifying questions to ask

Ask	Assume
Levels?	Several, each with a fixed set of spots
Vehicle and spot types?	Motorbike/car/truck; spots small/medium/large; a smaller vehicle may take a bigger spot
Gates?	Several entries and exits running at once
Fee?	Hourly, rate per vehicle type, easy to change
Display?	Free spots per level and size
Out of scope?	Advance reservations, valet, plate cameras

Minute-by-minute plan

- **0–8:** Actors: driver, attendant, admin. Use cases: take ticket, park, pay, exit, view free spots.
- **8–18:** `ParkingLot` owns many `Level` owns many `ParkingSpot`; `Vehicle`, `Ticket`, `EntryGate`, `ExitGate`, `SpotAllocationStrategy`, `FeeStrategy`, `Payment`.
- **18–23:** Walk "car enters": gate → lot → strategy picks spot → spot claimed → ticket issued.
- **23–35:** Code `ParkingSpot.tryOccupy`, one allocation strategy, `FeeStrategy.fee(ticket, exitTime)`.

What a strong answer must include

- [ ] Size enums for vehicle and spot, with one fit rule.
- [ ] `Ticket` as its own class (spot, vehicle, entry time).
- [ ] Allocation and fee each behind an interface; fee separate from payment method.
- [ ] Atomic spot claim across gates.
- [ ] Free counts maintained in one place, not by scanning.

Likely follow-ups and how to answer

- **Two gates, last spot:** compare-and-set on the spot (or a per-size free queue); the loser retries, then shows "full". A lot-wide lock is correct but serialises every gate.
- **EV charging spots:** a new spot kind plus a charging fee class; no gate code changes.
- **Weekend rates / first hour free:** a new `FeeStrategy`, or a decorator around an existing one.
- **Lost ticket:** flat daily maximum; look up by plate if stored.

Where to study it

Ch 14 — 14.1 Design a Parking Lot

24.2 Drill — Elevator System

The prompt as the interviewer says it

> "Design the elevators for a building. People press buttons on floors and inside cars. How do you decide which elevator goes where?"

Clarifying questions to ask

Ask	Assume
Size?	3–4 cars, ~20 floors
Buttons?	Up/down per floor (hall calls), floor panel per car (car calls), door and emergency buttons
Capacity?	Weight limit; an overloaded car holds its doors open
Dispatch goal?	Short waits, few direction changes, swappable
Special modes?	Maintenance, fire mode (all cars to ground)
Out of scope?	Motor physics, access cards

Minute-by-minute plan

- **0–8:** Actors: passenger, operator. Use cases: call car, choose floor, open/close door, maintenance, fire mode.
- **8–18:** `ElevatorSystem` → many `ElevatorCar` (state, current floor, up-stops and down-stops as sorted sets), `Request` (floor, direction, kind), `Dispatcher` interface, `Door`.
- **18–23:** Walk "floor 7 presses up": dispatcher scores cars, adds a stop to one.
- **23–35:** Code the dispatcher's cost function and `ElevatorCar.step()` using the LOOK rule (keep going while stops lie ahead, then reverse).

What a strong answer must include

- ☐ Hall calls assigned to one car; car calls belong to their car.
- ☐ Sorted stop sets per direction, not a FIFO queue; you can say why FIFO zig-zags.
- ☐ Dispatch as Strategy (nearest, least busy, zoned).
- ☐ Car states with legal transitions; doors never open while moving.
- ☐ Button presses from many threads handled safely.

Likely follow-ups and how to answer

- **Why not nearest idle car?** A car already heading towards the caller in the same direction is often better. Cost = distance + direction-change penalty + load penalty.
- **Many presses at once:** one controller thread drains a request queue, so dispatch logic stays single-threaded.
- **Fire mode:** a system-wide mode that clears stops and sends every car down; cars consult it each step.
- **A car breaks mid-trip:** maintenance state, re-dispatch its hall calls to other cars.

Where to study it

Ch 14 — 14.5 Design an Elevator System

24.3 Drill — Library Management System

The prompt as the interviewer says it

> "Design a library system. Members borrow and return books, and late returns are fined."

Clarifying questions to ask

Ask	Assume
Multiple copies of a title?	Yes; each copy has a barcode and rack
Loan rules?	14 days, max 5 books, renew once
Reservations?	Per-title queue; next in line is notified on return
Fines?	Per day late; may differ by membership type
Search?	Title, author, subject, date
Out of scope?	E-books, inter-library loans, payment processing

Minute-by-minute plan

- **0–8:** Actors: member, librarian, system. Use cases: search, borrow, return, renew, reserve, pay fine.
- **8–18:** `Book` (title data) owns many `BookItem` (physical copy), `Member`, `Loan`, `Reservation`, `Fine`, `Catalog`, `FinePolicy`, `Notifier`.
- **18–23:** Walk "return a late book": close loan → compute fine → check reservation queue → notify.
- **23–35:** Code `borrow` (limit, availability, reservation honoured) and `returnItem`.

What a strong answer must include

- [ ] The **book vs copy** split; this is what the interviewer looks for first.
- [ ] `Loan` as a class with its own dates, not fields on `Member`.
- [ ] Borrow limit enforced in one place, per membership type.
- [ ] Fine rule behind an interface; notifications decoupled.
- [ ] Search separated from lending.

Likely follow-ups and how to answer

- **Two members, last copy:** the copy's available → loaned change is atomic; the loser is offered a reservation.
- **Reserved copy never collected:** hold N days, then a scheduled sweep passes it on.
- **Add DVDs and magazines:** generalise to a catalog item only where lending rules match; a non-lendable magazine must not inherit `borrow`.

Where to study it

Ch 15 — 15.1 Design a Library Management System

24.4 Drill — Amazon Locker Service

The prompt as the interviewer says it

> "Customers can have parcels delivered to a self-service locker. Design delivery, pickup, and what happens when nobody collects."

Clarifying questions to ask

Ask	Assume
Sites and sizes?	Many sites; lockers small/medium/large
Opening a locker?	One-time code sent to the customer
Pickup window?	3 days, then return to sender
Returns?	Customer drops an item with a code; courier collects
Site full?	Courier redirected to the nearest site with space
Out of scope?	Ordering, payments, hardware drivers

Minute-by-minute plan

- **0–8:** Actors: customer, delivery agent, ops admin. Use cases: assign locker, drop parcel, pick up, return item, expire.
- **8–18:** `LockerSite` owns many `Locker` (size, state), `Parcel`, `LockerAssignment` (parcel, locker, code, deadline), `LockerSelectionStrategy`, `CodeGenerator`, `Notifier`, `ExpiryJob`.
- **18–23:** Walk "pick up with code": find assignment by code → check deadline → open → free locker.
- **23–35:** Code `LockerSite.assign(parcel)` and `pickUp(code)`.

What a strong answer must include

- [ ] Ordered size enums; "smallest that fits" behind an interface.
- [ ] Locker states: free → reserved → occupied → free, plus out of service.
- [ ] Code on the assignment, single-use, with a deadline.
- [ ] Scheduled expiry that frees lockers and starts return-to-sender.
- [ ] Customer returns reuse the same assignment mechanism.

Likely follow-ups and how to answer

- **Two couriers, last medium locker:** reserve under a per-site lock or poll a per-size free queue; fall back to a larger size, then another site.
- **Repeated wrong codes:** count failures per keypad (a wrong code matches no assignment, so a per-assignment count cannot work) and lock the keypad for a few minutes after five in a row.
- **Jammed locker on arrival:** mark out of service and reassign; the code belongs to the assignment, so it survives.

Where to study it

Ch 14 — 14.7 Design a Shipping Locker System (Amazon Locker)

24.5 Drill — Vending Machine

The prompt as the interviewer says it

> "Design a vending machine. A user inserts money, picks a product, and gets the product and change."

This is a State-pattern question; reach the states quickly.

Clarifying questions to ask

Ask	Assume
Payment?	Coins and notes; cards later
Order?	Pay first, then pick
Cancel?	Full refund of what was inserted
No change available?	Refuse and refund before dispensing
Products?	Slots (A1, B3) with product and count
Out of scope?	Telemetry, cooling

Minute-by-minute plan

- **0–8:** Actors: customer, operator. Use cases: insert, select, cancel, dispense, restock.
- **8–18:** `VendingMachine` (context) → `VendingState` with `Idle`, `HasMoney`, `Dispensing`, `SoldOut`, `Service`; `Inventory`, `Slot`, `Coin`, `ChangeMaker`.
- **18–23:** Walk the state diagram, including cancel and sold-out.
- **23–35:** Code the state interface, two states, and `ChangeMaker.makeChange`.

What a strong answer must include

- [ ] Each state lists the events it accepts; illegal events fail clearly.
- [ ] State pattern justified: each mode's rules in one class, no growing `switch`.
- [ ] Money as integer minor units, never `double`.
- [ ] Change checked against real coin stock **before** dispensing.
- [ ] Restock as a state, so sales stop during it.

Likely follow-ups and how to answer

- **Greedy change?** Works for normal coin sets, fails for odd ones or depleted stock; a small search is the fallback behind the same interface.
- **Card payment:** a payment-method abstraction; card authorises the exact price, so no change.
- **Product jams:** back to `HasMoney`, offer refund or another pick.
- **Enum vs State classes:** an enum with a switch is fine for four stateless modes; classes pay off once states carry data or multiply.

Where to study it

Ch 14 — 14.2 Design a Vending Machine

24.6 Drill — Online Blackjack Game

The prompt as the interviewer says it

> "Design an online Blackjack game: several players at a table against a dealer, with betting."

Clarifying questions to ask

Ask	Assume
Players per table?	1–5 plus the dealer
Decks?	Six-deck shoe, reshuffled at a cut card
Actions?	Hit, stand, double; split later
Dealer rule?	Hits below 17, stands on all 17s
Payouts?	1:1 win, 3:2 blackjack, push returns bet
Out of scope?	Side bets, insurance, networking details

Minute-by-minute plan

- **0–8:** Actors: player, dealer (system). Use cases: join, bet, deal, hit, stand, double, settle.
- **8–18:** `Card` (`Rank`, `Suit`), `Deck`, `Shoe`, `Hand`, `Player`, `Dealer`, `Table`, `Round`, `Bet`, `DealerPolicy`.
- **18–23:** Walk one round: bets → deal → players act in turn → dealer plays → settle.
- **23–35:** Code `Hand.bestValue()` and `Round.settle()`.

What a strong answer must include

- [ ] A game-agnostic `Card` / `Deck`, reusable for other card games.
- [ ] The ace rule inside `Hand` only.
- [ ] `Round` with phases, separate from the long-lived `Table`.
- [ ] Dealer rule as a policy, so "hit soft 17" is a new class.
- [ ] Bets debited once at start, settled once; an injectable seeded shuffle for tests.

Likely follow-ups and how to answer

- **Ace value:** count aces as 1; if total + 10 ≤ 21, add 10 once. That also yields the "soft" flag.
- **Splitting:** a player holds a list of `Hand`s, each with its own `Bet`; design for the list from the start.
- **Online multiplayer:** each table processes player commands from a queue on one thread, so game state needs no locks.
- **Reuse for Poker:** deck, card and shoe carry over; hand evaluation and round phases do not.

Where to study it

Ch 17 — Design Blackjack and a Deck of Cards

24.7 Drill — Ride Sharing (practice session)

The prompt as the interviewer says it

> "Design a ride-sharing service like Uber. Riders request rides, drivers accept them, and the rider pays at the end."

Clarifying questions to ask

Ask	Assume
Matching rule?	Nearest available driver within a radius; swappable
Ride types?	Standard and premium; carpool later
Fare?	Base + distance + time, with surge multiplier
Cancellation?	Either side may cancel; a fee after driver assignment
Ratings?	Both sides rate after the trip
Out of scope?	Maps/routing, real payment gateway

Minute-by-minute plan

- **0–8:** Actors: rider, driver, payment system. Use cases: request, match, accept, start, complete, pay, rate, cancel.
- **8–18:** `Rider`, `Driver` (status, `Vehicle`, `Location`), `RideRequest`, `Trip` (state), `MatchingStrategy`, `FareStrategy`, `DriverLocator`, `TripService`, `Payment`.
- **18–23:** Walk "request to completion": request → match → offer → accept → start → complete → fare → pay.
- **23–35:** Code `TripService.request` with matching and atomic driver claim, plus `Trip` state transitions.

What a strong answer must include

- [ ] `Trip` state machine: requested → driver assigned → in progress → completed, with cancelled from the early states.
- [ ] Matching and fare each as Strategy (nearest/best-rated; normal/surge).
- [ ] Driver availability claimed atomically.
- [ ] Location lookup behind its own class (a grid or geohash index), not a scan.
- [ ] Observer for live status updates to the rider.

Likely follow-ups and how to answer

- **Two requests match one driver:** compare-and-set the driver's status from available to assigned; the loser is re-matched to the next candidate.
- **Driver declines or times out:** offer to the next candidate; cap attempts, then tell the rider.
- **Carpool:** a trip holds several riders and seats; matching checks detour cost. A new strategy plus a seats field, not a new service.
- **Surge:** a fare decorator using a demand/supply ratio per area.

Where to study it

Ch 20 — 20.4 Design a Ride Sharing Service (Uber)

24.8 Drill — Meeting Scheduler

The prompt as the interviewer says it

> "Design a meeting scheduler for an office. People book rooms for meetings and invite colleagues."

Clarifying questions to ask

Ask	Assume
Room choice?	System picks a free room with enough capacity; user may name one
Time model?	Start and end on one day; half-open intervals, so 10–11 and 11–12 do not clash
Invitations?	Invitees accept or decline; organiser sees replies
Notifications?	On invite, change and cancel
Recurring meetings?	Later extension
Out of scope?	External calendars, time zones across offices

Minute-by-minute plan

- **0–8:** Actors: organiser, invitee, admin. Use cases: book, invite, respond, cancel, view room or person calendar.
- **8–18:** `Scheduler`, `Room` (capacity, own `Calendar` of bookings), `Meeting` (interval, organiser, room), `Invitation` (status), `Interval`, `RoomSelectionStrategy`, `Notifier`.
- **18–23:** Walk "book": filter rooms by capacity → strategy orders them → first room with no overlap is booked → invitations sent.
- **23–35:** Code `Interval.overlaps`, `Room.tryReserve` (sorted bookings, one `lowerEntry(end)` check) and `Scheduler.book`.

What a strong answer must include

- [ ] An `Interval` value type with a correct half-open overlap test.
- [ ] Bookings per room in a sorted map; clash check in O(log n).
- [ ] Check and insert under one per-room lock.
- [ ] Room selection as Strategy (smallest fitting, nearest).
- [ ] Invitations with RSVP status; notifications through Observer.

Likely follow-ups and how to answer

- **Two users, same room, overlapping times:** lock per room; when trying rooms in turn, release one lock before taking the next, so no deadlock.
- **First slot when all attendees are free:** gather busy intervals, sort, merge overlaps, scan the gaps for the required length.
- **Recurring meetings:** expand into instances; fail the whole booking on any clash, or offer to skip clashing dates.
- **Room goes out of service:** cancel or move its future meetings and notify organisers.

Where to study it

Ch 16 — Design a Meeting Scheduler

24.9 Drill — Movie Ticket Booking System

The prompt as the interviewer says it

> "Design BookMyShow: users find a movie, pick a show and seats, and pay."

Clarifying questions to ask

Ask	Assume
Scope?	Many cities, cinemas, screens and shows
Seat selection?	User picks exact seats on a seat map
Seat types and pricing?	Regular/premium; price per show and seat type
Hold time?	Seats held 10 minutes while the user pays
Cancellation?	Allowed up to a cutoff, with refund rules
Out of scope?	Reviews, food ordering, real payment gateway

Minute-by-minute plan

- **0–8:** Actors: customer, cinema admin, payment system. Use cases: search, view seats, hold, pay, cancel, add show.
- **8–18:** `Cinema` owns `Screen` owns `Seat`; `Movie`; `Show` (screen, movie, time); `ShowSeat` (seat state per show); `Booking`; `Payment`; `PricingStrategy`.
- **18–23:** Walk "book two seats": hold seats → create pending booking → pay → confirm, or release on timeout.
- **23–35:** Code `ShowSeat` hold/confirm/release and `BookingService.hold(show, seats, user)` as all-or-nothing.

What a strong answer must include

- [ ] `Seat` (physical) vs `ShowSeat` (state for one show); this is the key split.
- [ ] Seat states: available → held → booked, with hold expiry back to available.
- [ ] Holding several seats is all-or-nothing.
- [ ] Pricing as Strategy; booking confirmation via Observer.
- [ ] Payment failure releases the hold.

Likely follow-ups and how to answer

- **Two users click the same seat:** lock the show's seats in a fixed order (by seat id) and check-and-hold together; the second user sees "taken". Fixed order prevents deadlock between overlapping selections.
- **Payment succeeds after the hold expired:** check hold ownership at confirm; if lost, refund automatically.
- **Show cancelled by cinema:** cancel all bookings for the show and refund, notifying each customer.
- **Many servers:** the hold must live in shared storage with a conditional update, not in one JVM's memory.

Where to study it

Ch 15 — 15.2 Design a Movie Ticket Booking System

24.10 Drill — Car Rental System

The prompt as the interviewer says it

> "Design a car rental system. Customers reserve a car at a branch for some dates and return it later."

Clarifying questions to ask

Ask	Assume
Reserve a model or a specific car?	A vehicle type; the car is assigned at pickup
Branches?	Several; one-way rentals allowed
Pricing?	Daily rate per type, plus add-ons (insurance, GPS)
Late return / damage?	Extra charges at return
Cancellation?	Free until 24 hours before pickup
Out of scope?	Fleet maintenance scheduling, loyalty

Minute-by-minute plan

- **0–8:** Actors: customer, branch agent, admin. Use cases: search, reserve, pick up, return, cancel, pay.
- **8–18:** `Branch` holds many `Vehicle` (type, status); `Customer` ; `Reservation` (type, dates, pickup and drop branch, state); `RentalAgreement` ; `PricingPolicy` ; `Invoice` ; `AddOn` .
- **18–23:** Walk "return late": close rental → price base + add-ons + late fee → invoice → vehicle available at drop branch.
- **23–35:** Code availability for a date range and `Reservation` state transitions.

What a strong answer must include

- [ ] Availability as interval overlap, not a single "is free" flag.
- [ ] Reservation lifecycle: reserved → picked up → returned, plus cancelled and no-show.
- [ ] Pricing as Strategy; add-ons composable (Decorator or a list).
- [ ] One-way rental moves the vehicle to the drop branch.
- [ ] Atomic reserve against overlapping requests for the last car of a type.

Likely follow-ups and how to answer

- **Last SUV, two customers:** per-branch, per-type lock (or a count with compare-and-set) around check and reserve.
- **Overbooking:** allow a small configured margin per type; upgrade when short.
- **Dynamic pricing:** a new pricing class reading demand; nothing else changes.
- **Car breaks during rental:** swap vehicle on the same agreement; the agreement, not the car, carries the money.

Where to study it

Ch 16 — Design a Car Rental System

24.11 Drill — ATM

The prompt as the interviewer says it

> "Design an ATM. A customer inserts a card, enters a PIN, and withdraws cash or checks the balance."

Clarifying questions to ask

Ask	Assume
Transactions?	Withdraw, deposit, balance, transfer between own accounts
Who holds balances?	The bank, via a remote service
PIN rules?	Three wrong tries keep the card
Cash?	Cassettes per denomination; limited stock
Limits?	Daily withdrawal limit per card
Out of scope?	Bank's internal ledger, hardware drivers

Minute-by-minute plan

- **0–8:** Actors: customer, bank system, technician. Use cases: authenticate, withdraw, deposit, balance, refill cash.
- **8–18:** `ATM` (context) with states `Idle`, `CardInserted`, `Authenticated`, `Dispensing`; `Card`; `BankService` interface; `CashDispenser` with a chain of denomination handlers; `Transaction` subclasses.
- **18–23:** Walk "withdraw": check cash → debit via bank → dispense → print receipt; reverse on dispense failure.
- **23–35:** Code the state transitions and the dispenser chain.

What a strong answer must include

- [ ] ATM states with events allowed in each.
- [ ] Bank behind an interface; ATM never holds balances.
- [ ] Note dispensing as Chain of Responsibility (largest denomination first).
- [ ] Order: check cash → debit → dispense → on failure, reverse the debit.
- [ ] PIN retry limit and card retention.

Likely follow-ups and how to answer

- **Debited but no cash:** send a reversal with the same transaction id; the bank treats it idempotently.
- **Two ATMs, one account:** the bank serialises debits per account; the ATM just handles a "declined" reply.
- **Network drop mid-transaction:** retry with the same id; idempotency stops a double debit.
- **Add cardless (QR) withdrawal:** a new authentication state; the rest of the flow is unchanged.

Where to study it

Ch 14 — 14.4 Design an ATM

24.12 Drill — Chess Game

The prompt as the interviewer says it

> "Design a two-player chess game."

Clarifying questions to ask

Ask	Assume
Players?	Two humans; computer player later
Full rules?	All moves, including castling, en passant, promotion
End conditions?	Checkmate, stalemate, resignation, draw by agreement
Clocks?	Optional time control
Undo?	Yes, for practice games
Out of scope?	Ratings, online matchmaking

Minute-by-minute plan

- **0–8:** Actors: two players. Use cases: start, move, undo, resign, offer draw.
- **8–18:** `Game` (status, turn, history), `Board` owns 64 `Square`s, abstract `Piece` with subclasses, `Move` (from, to, captured, special kind), `Player`, `MoveValidator`.
- **18–23:** Walk "make a move": piece's pseudo-legal moves → apply to a copy → reject if own king in check → commit → update status.
- **23–35:** Code `Piece.candidateMoves` for two pieces and `Game.makeMove` with the check test.

What a strong answer must include

- [ ] Piece movement through polymorphism, one subclass each; sliding pieces share a helper.
- [ ] `Move` as an object (Command) storing enough to undo.
- [ ] Legality = pseudo-legal + "own king not in check after the move".
- [ ] Game status computed after every move (check, mate, stalemate).
- [ ] Turn enforcement in `Game`, not in pieces.

Likely follow-ups and how to answer

- **Detect checkmate:** in check and no legal move exists; generate all moves and test each.
- **Castling rules:** needs "has moved" flags, empty path, and no attacked squares on the king's path.
- **Add an AI player:** a `Player` strategy that picks a move; the game loop does not change.
- **Why not a giant switch on piece type?** Adding a fairy piece would edit it; a subclass is closed for modification.

Where to study it

Ch 17 — Design Chess

24.13 Drill — Hotel Management System

The prompt as the interviewer says it

> "Design a hotel management system: booking rooms, check-in, check-out, and billing."

Clarifying questions to ask

Ask	Assume
Book a room type or a room?	A type; the room is assigned at check-in
Date model?	Nights as half-open date ranges
Extra services?	Room service, laundry, added to the bill
Cancellation and no-show?	Free until a cutoff; one night charged for no-show
Housekeeping?	Room must be cleaned before the next guest
Out of scope?	Channel managers, loyalty

Minute-by-minute plan

- **0–8:** Actors: guest, receptionist, housekeeper. Use cases: search, reserve, check in, add charge, check out, cancel.
- **8–18:** `Hotel` owns `Room` (type, status); `Guest`; `Reservation` (type, date range, state); `Stay`; `Invoice` with `Charge` lines; `RoomPricing`; `HousekeepingTask`.
- **18–23:** Walk "check out": total charges → pay → room marked dirty → housekeeping task created.
- **23–35:** Code availability per room type over a date range and the reservation state machine.

What a strong answer must include

- [ ] Availability = count of rooms of a type minus overlapping reservations, per night.
- [ ] Reservation states: booked → checked in → checked out, plus cancelled and no-show.
- [ ] Room status separate from reservation status (occupied, dirty, clean, out of order).
- [ ] Pricing as Strategy (season, weekend).
- [ ] Charges as line items, so new services need no invoice change.

Likely follow-ups and how to answer

- **Last room, two bookings:** lock per room type and date range check-and-insert; or a per-night counter decremented atomically.
- **Early check-in / late check-out:** a charge line plus a room-status check.
- **Overbooking:** allow a configured margin and a "walk" procedure to a partner hotel.
- **Group bookings:** a parent reservation holding several room reservations, booked all-or-nothing.

Where to study it

Ch 15 — 15.4 Design a Hotel Management System

24.14 Drill — Food Delivery (practice session)

The prompt as the interviewer says it

> "Design a food delivery app like Swiggy: customers order from restaurants, and delivery agents bring the food."

Clarifying questions to ask

Ask	Assume
One restaurant per order?	Yes; the cart resets if the user switches restaurant
Restaurant capacity?	A kitchen cooks at most N items at once
Agent assignment?	Nearest free agent when the food is ready; swappable
Tracking?	Customer sees order status changes
Cancellation?	Until cooking starts; prepaid orders are refunded
Out of scope?	Routing maps, payment gateway internals

Minute-by-minute plan

- **0–8:** Actors: customer, restaurant, delivery agent, admin. Use cases: browse, add to cart, place order, accept, prepare, assign agent, deliver, cancel.
- **8–18:** `Restaurant` (menu, capacity, rating), `MenuItem`, `Cart`, `Order` (state), `DeliveryAgent` (status, location), `RestaurantSelectionStrategy`, `AgentAssignmentStrategy`, `OrderService`, `Notifier`.
- **18–23:** Walk "place order to delivered": cart → order placed → restaurant accepts (capacity taken) → ready → agent assigned → picked up → delivered (capacity freed).
- **23–35:** Code `Order` state transitions and capacity handling in `OrderService`.

What a strong answer must include

- ☐ Order state machine with who may trigger each transition.
- ☐ Kitchen capacity taken on accept, released when the food is ready or the order is cancelled; a refusal for a full kitchen refunds a prepaid order.
- ☐ Restaurant selection and agent assignment as separate Strategies.
- ☐ Atomic claim of a free agent.
- ☐ Status updates through Observer.

Likely follow-ups and how to answer

- **Two orders, one agent:** compare-and-set the agent's status; the loser goes to the next candidate.
- **Restaurant never accepts:** a timeout auto-cancels and refunds.
- **Batch two orders to one agent:** assignment strategy that allows a second pickup within a detour limit.
- **Pick the cheapest restaurant for a whole cart (FoodKart):** a selection strategy across restaurants serving every item.

Where to study it

Ch 19 — Design an Online Food Delivery Service (Swiggy) + FoodKart

24.15 Drill — Amazon Online Shopping System

The prompt as the interviewer says it

> "Design the core of Amazon: browse products, add to cart, check out, and track the order."

Clarifying questions to ask

Ask	Assume
Sellers?	Many; a product may be sold by several
Stock?	Per product per warehouse
Reserve stock when?	At checkout, released if payment fails
Payments?	Card, wallet, cash on delivery
Order changes?	Cancel before shipping; return after delivery
Out of scope?	Recommendations, search ranking

Minute-by-minute plan

- **0–8:** Actors: customer, seller, admin, payment system. Use cases: search, add to cart, checkout, pay, track, cancel, return, review.
- **8–18:** `Product`, `Catalog`, `Cart` with `CartItem`, `Order` with `OrderItem` (price copied), `Inventory`, `Payment` (Strategy), `Shipment`, `Customer`, `Notifier`.
- **18–23:** Walk "checkout": validate cart → reserve stock → create order → pay → confirm, or release on failure.
- **23–35:** Code `CheckoutService.checkout` and `Order` state transitions.

What a strong answer must include

- [ ] `OrderItem` copies price at purchase time; the cart reads live prices.
- [ ] Stock reserved atomically for all lines, released on failure.
- [ ] Order states: placed → paid → shipped → delivered, plus cancelled and returned.
- [ ] Payment and shipping methods as Strategy.
- [ ] Notifications via Observer on state change.

Likely follow-ups and how to answer

- **Price changes while in the cart:** recheck at checkout and show the new price before paying.
- **Last unit, two buyers:** conditional decrement on stock; the loser gets "out of stock" at checkout.
- **Partial cancellation:** cancel per `OrderItem`, refund that line, release its stock.
- **Payment timeout:** reservation has an expiry; a sweeper releases stock.

Where to study it

Ch 19 — Design Amazon — Online Shopping System

24.16 Drill — Stack Overflow

The prompt as the interviewer says it

> "Design Stack Overflow: users ask questions, answer them, vote, and earn reputation."

Clarifying questions to ask

Ask	Assume
Content types?	Question, answer, comment; tags on questions
Voting?	Up/down on questions and answers, one vote per user per post
Reputation?	Changes on votes and accepted answers; unlocks privileges
Bounties and badges?	Yes, as extensions
Moderation?	Close, delete, flag
Out of scope?	Full-text search engine internals, chat

Minute-by-minute plan

- **0–8:** Actors: guest, member, moderator, system. Use cases: ask, answer, comment, vote, accept, search, close.
- **8–18:** `User` (reputation), abstract `Post` → `Question`, `Answer`; `Comment`; `Vote`; `Tag`; `Badge`; `Bounty`; `ReputationService`; `SearchService`.
- **18–23:** Walk "upvote an answer": check privilege → record vote (unique per user and post) → update score → reputation event to the author.
- **23–35:** Code `VoteService.vote` and the reputation listener.

What a strong answer must include

- [ ] Shared `Post` base for voting and commenting; `Question` adds title, tags and accepted answer.
- [ ] `Vote` keyed by (user, post) so a second vote replaces, never adds.
- [ ] Reputation changed by events (Observer), not inline in the vote code.
- [ ] Privileges as reputation thresholds in one place.
- [ ] Search separate from posting.

Likely follow-ups and how to answer

- **Concurrent votes on a hot post:** atomic score counter; the unique (user, post) key stops double voting.
- **Changing a vote:** reverse the old reputation delta, then apply the new one.
- **Bounties:** reputation is escrowed when offered and paid on award or expiry.
- **Self-voting / vote rings:** a rule in the vote service; detection can run later as a job.

Where to study it

Ch 18 — 18.1 Design Stack Overflow

24.17 Drill — Restaurant Management System

The prompt as the interviewer says it

> "Design a restaurant system: table reservations, taking orders, the kitchen, and the bill."

Clarifying questions to ask

Ask	Assume
Reservations and walk-ins?	Both; walk-ins join a waitlist
Tables?	Sizes 2/4/6; tables may be joined
Ordering?	Waiter enters orders per table; items go to kitchen stations
Bill?	Per table, with tax and tip; may be split
Menu changes?	Items can be marked unavailable
Out of scope?	Inventory purchasing, payroll

Minute-by-minute plan

- **0–8:** Actors: guest, host, waiter, chef, cashier. Use cases: reserve, seat, order, prepare, serve, bill, pay.
- **8–18:** `Table` (size, status), `Reservation`, `WaitlistEntry`, `Menu` with `MenuItem`, `Order` with `OrderItem` (state), `Kitchen` with stations, `Bill`, `TableAllocator`.
- **18–23:** Walk "order to bill": waiter places order → items sent to stations → ready → served → bill.
- **23–35:** Code `TableAllocator` (smallest table that fits) and `OrderItem` state transitions.

What a strong answer must include

- [ ] Table allocation as Strategy, with waitlist fallback.
- [ ] Orders sent to the kitchen as Command objects (queued, can be cancelled).
- [ ] `OrderItem` state: ordered → cooking → ready → served.
- [ ] "Table ready" and "dish ready" via Observer.
- [ ] Bill built from order items plus tax/tip policy; split as a separate step.

Likely follow-ups and how to answer

- **Reservation vs walk-in for the same table:** hold reserved tables from a time before the slot; walk-ins only get tables free for a full sitting.
- **Split the bill:** by items or evenly; a split strategy producing several bills.
- **Item runs out mid-service:** mark unavailable; pending orders with it are flagged to the waiter.
- **Several waiters ordering for one table:** lock the order per table, or append-only items.

Where to study it

Ch 16 — Design a Restaurant Management System

24.18 Drill — Facebook

The prompt as the interviewer says it

> "Design the core of Facebook: profiles, friends, posts, a news feed and privacy."

Clarifying questions to ask

Ask	Assume
Friends or followers?	Both: mutual friendship and one-way follow
Post visibility?	Public, friends, only me, custom list
Feed?	Recent posts from friends, followed pages and groups, ranked
Interactions?	Like, comment, share
Groups and pages?	Yes, simple versions
Out of scope?	Ads, messaging, media storage

Minute-by-minute plan

- **0–8:** Actors: user, page admin, system. Use cases: send/accept friend request, post, comment, like, view feed, search.
- **8–18:** `User`, `Profile`, `FriendRequest` (state), `Friendship`, `Follow`, `Post` (audience), `Comment`, `Reaction`, `Group`, `Page`, `VisibilityPolicy`, `FeedService` with `FeedRanker`.
- **18–23:** Walk "load feed": collect candidate posts → filter by visibility → rank → page.
- **23–35:** Code `VisibilityPolicy.canSee(viewer, post)` and the feed assembly.

What a strong answer must include

- [ ] Friendship (two-sided, via request) distinct from follow (one-sided).
- [ ] One visibility policy reused by feed, profile and search.
- [ ] Feed ranking as Strategy.
- [ ] A stated choice of pull (merge at read) vs push (fan-out at write), with the trade-off.
- [ ] Friend request state: pending → accepted / rejected.

Likely follow-ups and how to answer

- **Push or pull?** Pull is simpler and right for an in-memory design; push makes reads fast but costs writes for users with many friends. Hybrid: push for normal users, pull for celebrities.
- **Privacy changed after posting:** check visibility at read time, never cache it into the feed.
- **Friend suggestions:** friends-of-friends counted by mutual friends (two-hop search).
- **Unfriend:** remove both sides of the friendship atomically.

Where to study it

Ch 18 — 18.4 Design Facebook

24.19 Drill — Online Stock Brokerage System

The prompt as the interviewer says it

> "Design an online stock brokerage: users fund an account, place buy and sell orders, and see their portfolio."

Clarifying questions to ask

Ask	Assume
Order types?	Market and limit; stop-loss later
Who matches orders?	An external exchange; we route and track
Partial fills?	Yes
Funds?	Blocked when a buy is placed, released on cancel
Alerts?	Price alerts and watchlists
Out of scope?	Margin trading, tax reports

Minute-by-minute plan

- **0–8:** Actors: investor, exchange, admin. Use cases: deposit, place order, cancel, receive fill, view portfolio, set alert.
- **8–18:** `Account` (cash, blocked cash), `Portfolio` with `Position`s, abstract `Order` → `MarketOrder`, `LimitOrder` (state), `Fill`, `ExchangeGateway` (Adapter), `OrderService`, `Watchlist`, `PriceAlert`.
- **18–23:** Walk "buy with partial fills": validate → block funds → send → fills update position and cash → filled or cancelled remainder releases funds.
- **23–35:** Code `OrderService.place` and `onFill`.

What a strong answer must include

- [ ] Funds or shares blocked at placement, not at fill.
- [ ] Order states: new → open → partially filled → filled / cancelled / rejected.
- [ ] Exchange behind an Adapter interface.
- [ ] Money as `BigDecimal` or integer units.
- [ ] Price alerts through Observer.

Likely follow-ups and how to answer

- **Two orders spend the same cash:** lock per account around check-and-block.
- **Fill arrives after cancel:** the exchange is the truth; apply the fill, cancel only the remainder.
- **Duplicate fill messages:** keep processed fill ids; ignore repeats (idempotency).
- **How would the exchange match?** Per symbol, a price-sorted book per side, matched by price then time.

Where to study it

Ch 19 — Design an Online Stock Brokerage System

24.20 Drill — Jigsaw Puzzle

The prompt as the interviewer says it

> "Design a jigsaw puzzle: pieces, the board, how to tell whether two pieces fit, and an algorithm to solve it."

Clarifying questions to ask

Ask	Assume
Piece shape?	Rectangular grid; each piece has four sides
Side kinds?	Flat (border), tab (sticks out), blank (goes in)
How do two sides fit?	A tab fits a blank with the same shape id
Rotation?	Pieces may be rotated
Solve or just play?	Both: a user places pieces; a solver can finish
Out of scope?	Image processing, 3-D puzzles

Minute-by-minute plan

- **0–8:** Actor: player (the solver is part of the system). Use cases: start (includes validating the box), place piece (includes choosing a rotation), remove, auto-solve.
- **8–18:** `Puzzle`, `Board` (grid of `Placement(piece, rotation)`), immutable `Piece` (four `Side`s), `Side` (kind, shape id), `Direction` enum (TOP, RIGHT, BOTTOM, LEFT), `SolverStrategy`.
- **18–23:** Walk the solver: corners (two flat sides) → borders (one flat side) → inside, each by matching neighbours.
- **23–35:** Code `Side.fits(other)`, `Piece.side(direction, rotation)`, `Board.check` and the solver loop with a side index.

What a strong answer must include

- [ ] Sides as their own type with a fit rule, not booleans on the piece.
- [ ] Rotation stored in the placement, not the piece; the side facing `d` after `r` turns is `sides[(d − r) mod 4]`.
- [ ] Piece category from flat-side count: corner, border, interior.
- [ ] Solver that fills cells in order and looks up candidates by shape id (hash map), not by scanning all pieces.
- [ ] Board validates a placement against already-placed neighbours.

Likely follow-ups and how to answer

- **Complexity:** with unique shapes and an index from side to pieces, each cell is O(1) average, so O(n) total; scanning every piece is $O(n^2)$. Look-alike pieces force backtracking, so budget the search.
- **Two pieces fit the same spot:** add a picture-similarity score as a tie-breaker, or backtrack.
- **Undo a placement:** placements as Command objects on a stack.
- **Non-rectangular puzzles:** sides become a list and `Edge` a direction count; most code is unchanged.

Where to study it

Ch 17 — Design a Jigsaw Puzzle

24.21 Drill — Airline Management System

The prompt as the interviewer says it

> "Design an airline reservation system: search flights, book seats, check in."

Clarifying questions to ask

Ask	Assume
Route vs dated departure?	A flight number recurs; each date is its own instance
Multi-leg trips?	Yes, booked all-or-nothing
Seat classes?	Economy, business; seat selection optional
Fares?	Per class, varies with time to departure
Changes?	Cancel with refund rules; crew assignment for admin
Out of scope?	Loyalty, baggage tracking

Minute-by-minute plan

- **0–8:** Actors: customer, agent, admin, crew. Use cases: search, book, pay, select seat, check in, cancel, assign crew.
- **8–18:** `Airport`, `Flight` (route, number), `FlightInstance` (date, aircraft, status), `Aircraft` with `Seat`s, `FlightSeat` (per instance), `Itinerary` with `Reservation`s, `Passenger`, `FarePolicy`, `Payment`.
- **18–23:** Walk "book a two-leg itinerary": hold seats on both instances → pay → confirm both, or release both.
- **23–35:** Code the itinerary hold as all-or-nothing and `FlightSeat` states.

What a strong answer must include

- ☐ `Flight` vs `FlightInstance` split, and seat state per instance.
- ☐ All-or-nothing booking across legs with release on failure.
- ☐ Fare rules as Strategy.
- ☐ Booking states: held → confirmed → checked in / cancelled.
- ☐ Flight status changes notify passengers (Observer).

Likely follow-ups and how to answer

- **Two people, same seat:** lock seats in a fixed order across instances to avoid deadlock; check-and-hold together.
- **Flight cancelled:** rebook passengers on the next instance or refund.
- **Overbooking:** allow a margin per class; deny boarding by a priority rule at check-in.
- **Multi-airline itineraries:** an adapter per partner airline behind one booking interface.

Where to study it

Ch 16 — Design an Airline Management System

24.22 Drill — ESPNcricinfo

The prompt as the interviewer says it

> "Design a cricket site like ESPNcricinfo: live scores, scorecards, commentary and player stats."

Clarifying questions to ask

Ask	Assume
Formats?	Test, ODI, T20
Data entry?	An admin enters each ball
Live view?	Scorecard and commentary update per ball
Stats?	Career stats per player, per format
Tournaments?	Series and tournaments with points tables
Out of scope?	Video, fantasy games

Minute-by-minute plan

- **0–8:** Actors: viewer, scorer/admin, system. Use cases: record ball, view scorecard, follow match, view stats, manage tournament.
- **8–18:** `Tournament`, `Match` (format), `Team`, `Player`, `Innings`, `Over`, `Ball` (runs, extras, wicket), `Wicket`, `Commentary`, `Scorecard`, `MatchFormatRules`.
- **18–23:** Walk "record a ball": validate → append ball → update innings totals and strike → notify subscribers.
- **23–35:** Code `Innings.record(ball)` including extras and strike rotation.

What a strong answer must include

- [ ] `Ball` as the source event; totals and stats derived from balls.
- [ ] Extras handled correctly: wide and no-ball do not count as legal deliveries.
- [ ] Strike rotation on odd runs and at the end of an over.
- [ ] Format rules (overs, powerplays) behind an interface or Template Method.
- [ ] Live updates via Observer.

Likely follow-ups and how to answer

- **Scorer corrects a ball:** replace that event and recompute derived totals.
- **Career stats fast to read:** maintain aggregates updated from ball events, rebuildable from the log.
- **Many viewers:** one writer per match; readers get immutable snapshots.
- **Rain-reduced match:** a format rule computes the revised target; the ball model is unchanged.

Where to study it

Ch 18 — 18.2 Design Cricinfo

24.23 Drill — LinkedIn

The prompt as the interviewer says it

> "Design LinkedIn: profiles, connections, a feed, jobs and messages."

Clarifying questions to ask

Ask	Assume
Connections?	Two-sided, by invitation; show 1st/2nd/3rd degree
Profile?	Experience, education, skills, endorsements
Jobs?	Companies post; members apply
Feed?	Posts from connections and followed companies
Messaging?	One-to-one, between connections
Out of scope?	Ads, premium analytics

Minute-by-minute plan

- **0–8:** Actors: member, recruiter, company admin. Use cases: edit profile, connect, accept, post, view feed, post job, apply, message.
- **8–18:** `Member`, `Profile` (`Experience`, `Education`, `Skill`), `Invitation` (state), `Connection`, `Company`, `Job`, `Application` (state), `Post`, `Message`, `FeedService`, `ConnectionGraph`.
- **18–23:** Walk "accept invitation": invitation → accepted → connection stored for both sides → feed and suggestions update.
- **23–35:** Code `ConnectionGraph.degree(a, b)` with a bounded breadth-first search, and `Invitation` transitions.

What a strong answer must include

- [ ] Invitation states: pending → accepted / ignored / withdrawn.
- [ ] Connection symmetric and created for both members in one step.
- [ ] Degree by BFS limited to depth 3.
- [ ] Job applications with their own state (applied → reviewed → rejected / offered).
- [ ] Feed ranking as Strategy; notifications via Observer.

Likely follow-ups and how to answer

- **Degree on a huge graph:** bidirectional BFS from both ends, stop at depth 3.
- **"People you may know":** second-degree members ranked by mutual connections.
- **Messages to non-connections:** a permission check (premium or recruiter) before sending.
- **Endorsements:** a (skill, endorser) pair stored once per endorser.

Where to study it

Ch 18 — 18.3 Design LinkedIn

Chapter summary

- The patterns repeat across problems: an atomic claim of a scarce resource (spot, seat, room, driver, locker), a state machine for the main entity, a Strategy for the rule most likely to change, and Observer for notifications. Recognise which one the interviewer is probing.