Chapter 15 — ML System Design: The Interview Framework

"Design YouTube's recommendations." "Design a fraud detector." "Add AI to our support desk." These now sit beside "Design Twitter" in system design loops, and are asked of ordinary **backend engineers**, not only ML specialists. Everything about load balancers, caches, queues and databases still applies (System Design Fundamentals). The new part: one box **learns from data** — its output can be wrong while the request succeeds, its quality can decay while the code is unchanged, and each call may cost real money.

This chapter gives a repeatable 45-minute method: the five question forms and role expectations; four classic assumptions that break; three question families (predictive, generative, agentic); a seven-step framework (frame → success metrics → data → baseline → serving → keep-it-working → failures and cost); turning a **business goal** into an **ML objective**; scoping a vague question (capstone); the mistakes that sink answers. Chapter 16 goes deep on each production block (training pipelines, feature stores, serving, evaluation, monitoring, A/B tests, **MLOps** — practices and tooling for running ML in production).

Code is in `code/`; outputs are real.

15.1 ML design vs classic system design

What it is

An **ML (or AI) system design question** asks for a product in which at least one component is a **model** — a function learned from data rather than written as rules — mapping an input (watch history, a transaction, a question) to an output (score, label, number, text).

Term	Meaning	Example
Training	Fitting or updating the model from data; usually a background job separate from user traffic	Nightly job learning from yesterday's orders
Inference	Running the trained model on an input	Scoring one transaction at checkout
Online inference	Inference inside a live request; must meet its latency limit	Fraud score before a payment is approved
Batch inference	Inference over many inputs in a scheduled job; results stored	Tonight's recommendation lists for every active user

A model is like a new hire who learned from past examples: you still need an office (infrastructure), checks on their work (evaluation), a way to notice decline (monitoring), a budget (cost) and a plan for mistakes (fallbacks). Model internals help, but the round mostly grades the **whole system** — candidates who spent weeks on attention and loss derivations often freeze at "where does the training data come from?" or "what if the model times out?".

The five forms of the question

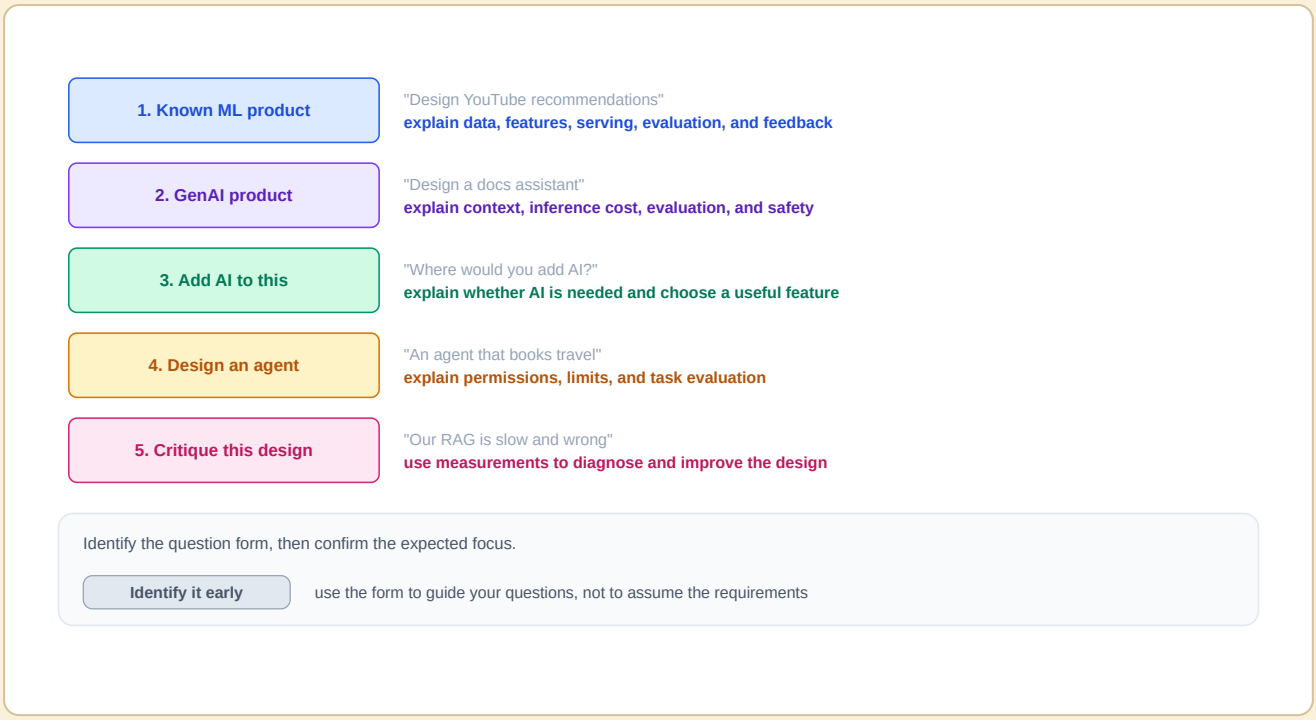


Figure — Five shapes of AI design prompts (course diagram)

#	Form	Example prompts	What to emphasise
1	**Known ML product**	"Design YouTube recommendations", "Design search autocomplete"	Data, features, training, serving, evaluation, feedback loop
2	**GenAI product**	"Design ChatGPT", "Design a docs Q&A assistant"	Retrieval, context assembly (building the text sent to the model), inference cost, output evaluation, safety
3	**Add AI to an existing system**	"Our store runs fine — which part should get AI?"	Whether AI is needed at all, one useful feature, simplest viable design
4	**Design an agent**	"An agent that books travel"	Tool permissions, approvals, step/cost limits, stop conditions, task-success evaluation
5	**Critique / scale an AI design**	"Our RAG system is slow and wrong — fix it"	Diagnose with measurements, logs and failed examples *before* changing anything

Terms: a **feature** is one input to the model (watch history, hour); the **feature pipeline** computes and delivers them; the **feedback loop** collects outcomes to evaluate and improve. **GenAI** = generative AI. **LLM** = large language model. **RAG** = retrieval-augmented generation (put relevant fetched documents in the model's input — grounds answers, doesn't guarantee correctness). **Fine-tuning** = extra training of an existing model for a narrower task.

Who gets asked what

Depth varies by role and company — **ask**.

Role	Forms to prepare	Depth expected
Software engineer (senior+)	2, 3, sometimes 1	Architecture, serving, cost, failure handling; model inputs/outputs/limits
ML engineer	1, 2, 5	Above + model choice, features, training pipelines, evaluation metrics
AI / ML platform engineer	2, 4, 5	Inference infrastructure: GPUs, batching, caching, latency under load
Engineering manager / staff+	3, 5	Build vs buy, cost limits, ownership, risk, gradual rollout

First-minute question: *"Would you like me to go deep on model development, or on the system that serves and maintains the model?"*

What interviewers grade (and what they don't)

Graded	Usually secondary
Requirements discipline — users, latency, quality measure, out of scope, numeric assumptions	Deriving the transformer or backprop by hand
A complete loop, not just the request path — outcomes come back and improve the system	Naming this month's newest model
Cost and latency realism — assume, estimate, optimise	Framework trivia without saying why a component is needed
Failure honesty — a model can be confidently wrong with HTTP 200; detection, **blast-radius** limits (how much one failure can damage), fallbacks	

How it works — four assumptions that break when a model enters

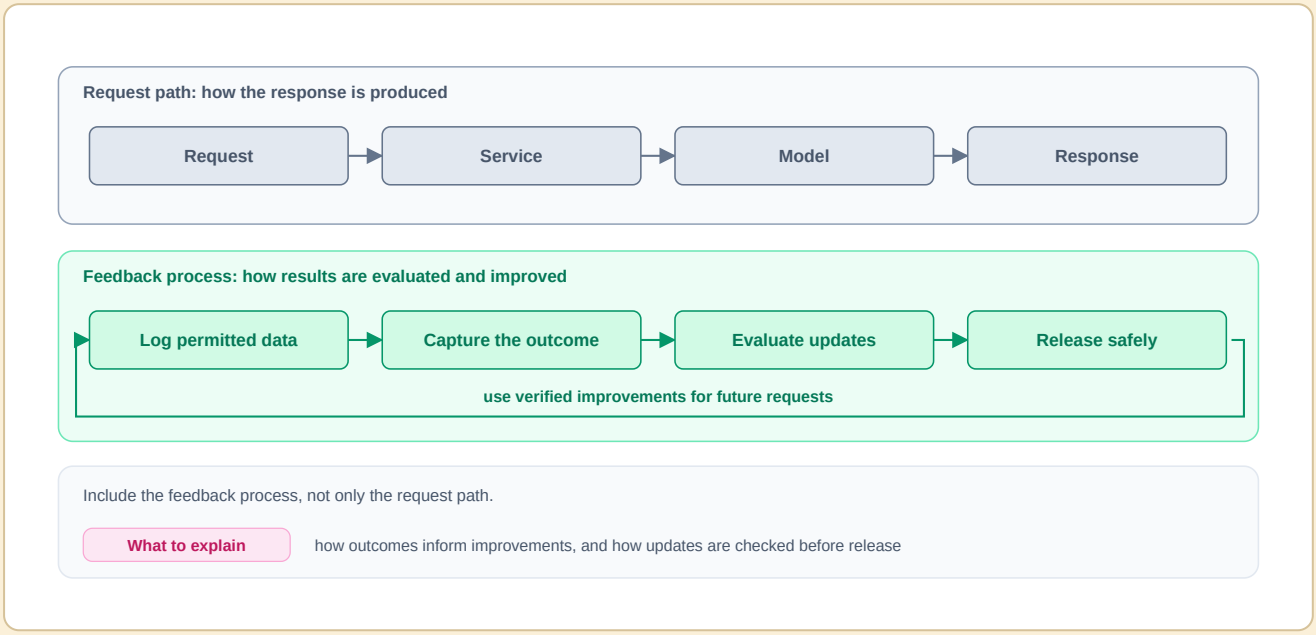


Figure — Serving path on top, improvement loop underneath (course diagram)

Contrasts, not absolutes — but where incomplete AI designs fall down.

#	Classic assumption	What changes with a model	Design response
1	**Correctness is binary** (a payment is right or a bug)	Quality is a spectrum: a recommendation can render fine and be irrelevant	**Quality metrics** (CTR — clicks ÷ impressions, resolution rate, suggestion-acceptance) *alongside* latency and availability **SLOs** (service-level objectives); log predictions and join them with later outcomes
2	**The code is the behaviour**	Behaviour depends on data. **Drift** — change in input patterns or the input→outcome relationship — degrades an unchanged model (new fraud tactics, changed tastes)	Feedback process: collect outcomes, investigate quality changes, evaluate candidate vs current, release gradually with rollback. Retrain on evidence, not on every alert
3	**Compute is cheap**	A model call can cost orders of magnitude more than a DB read, depending on model size, input/output length, hardware and utilisation	Put a number on it: 10M requests/day × $0.01 = $100,000/day ≈ $3M/month (illustrative). Then caching (when permissions/freshness allow), routing easy requests to a smaller model (if it meets the quality bar), batch precompute
4	**Failures are loud** (500s, timeouts)	A wrong answer arrives as `200 OK`: invented citation, missed fraud, harmful text	Treat user text and retrieved text as possibly hostile (**prompt injection** — untrusted text overriding instructions — can't be fully filtered; enforce permissions outside the model), check outputs, human approval for high-impact actions, fallbacks (unpersonalised results, hand-off to a human)

For every design ask: **"Who is hurt by a wrong output, and what limits the damage?"**

Build, host, or call an API?

Option	Pros	Cons / what you still own
Hosted API (pay per use)	Fastest launch, no GPU fleet	Data-handling terms, rate limits, provider availability, per-call cost at scale; you still build integration, evaluation, monitoring. Poor fit under strict **data-residency**/sensitivity rules, very high steady volume, or latency the provider can't meet
Self-host open-weights model (weights downloadable under a licence)	Control over data, deployment, latency; can be cheaper at high, steady utilisation	Capacity, security, upgrades, on-call, often GPUs; savings not guaranteed
Train / fine-tune your own	Needed when your interaction data is the product (ranking, fraud, recommendations); fine-tuning can help a narrow generative task	Data preparation, training infra, ongoing maintenance

Options overlap (fine-tune, then managed hosting). Compare on quality, data sensitivity, volume, latency, total cost and team skills; start simplest and say what evidence would make you switch.

Trade-offs

Going ML at all is a trade-off. Google's *Rules of ML*: **Rule #1: don't be afraid to launch without ML** — a **heuristic** (hand-written rule of thumb) may suffice to start, gives a baseline, and produces the data you'll learn from. **Rule #3: choose ML over a complex heuristic** — once rules become a tangle, a model is easier to maintain.

Common mistakes

- Answering an SWE design question like a research talk (15 minutes on fine-tuning, nothing on serving or evaluation).

- Drawing only the request path — no logging, outcome join, retraining or evaluation loop.
- Assuming model calls are free or cache-fast.
- Treating thumbs up/down as trustworthy training labels.
- Claiming a fixed traffic share can go to a small model without evidence.

Real-world examples

Netflix, YouTube, TikTok feeds (form 1); **GitHub Copilot**-style assistants measured by suggestion acceptance and **retention** (form 2); **Stripe Radar**, a fraud model inside an existing payments flow (form 3 in spirit); support agents issuing refunds through limited tools (form 4).

In the interview

> **30-second answer** — *"What changes compared with an ordinary backend service?" → "Distributed-systems fundamentals still hold. I add quality metrics beside latency and uptime SLOs; a loop that notices when quality slips and tests any new version before a staged rollout; a budget for per-call latency and spend; and safeguards for bad outputs — validation, fallbacks, human review where stakes warrant."*

- **"What do you mean by quality?"** → Feed: CTR plus a longer-term signal (watch time, return visits). Fraud: **precision** (share of flagged that is fraud) and **recall** (share of fraud flagged) at the chosen threshold. Support: correct resolution rate plus **CSAT** (customer-satisfaction score) and repeat-contact rate.
- **"Would you fine-tune?"** → Not first. Check whether an existing model with good instructions (plus retrieval for private data) meets quality and cost on an evaluation set; fine-tune only for a measured gap simpler changes don't close.
- **"Assume the model exists."** → Spend the time on integration, serving, cost, evaluation and feedback.

TL;DR

Classic design still applies; a model adds **quality measurement, drift, cost, and silent failure**. Five forms; confirm depth. Graded: requirements, full loop, cost/latency realism, failure honesty. Build vs host vs API is requirements-driven — start simplest.

15.2 The three question families

What it is

Classify what the system mainly **produces**:

Family	The model produces…	Typical questions	Focus of your answer	Classic failure
1. Predictive	A score, number or label the system acts on	Recommendations, feed/search ranking, ads CTR, fraud/spam, ETA, churn	Data & features, candidate generation + ranking (cheap shortlist, then careful ordering), offline & online metrics, feedback loop	Picking a model ("an LLM ranks the feed") without data, latency or cost reasoning
2. Generative	Content: answer, summary, code, image	ChatGPT-style assistant, docs bot, code assistant, meeting summaries	Retrieval & context assembly, inference latency/cost (streaming, caching, routing), output evaluation, safety	Training a foundation model from scratch; ignoring cost; no quality measurement
3. Agentic	Which **action** to take next (call a tool, read result, decide again)	Refund-issuing support agent, travel booking, coding agent opening PRs	Tool interface & permissions, step/time/cost limits, stop conditions, approval for risky actions, task-level evaluation	Unlimited tools and no stop rule — or human approval for every harmless step

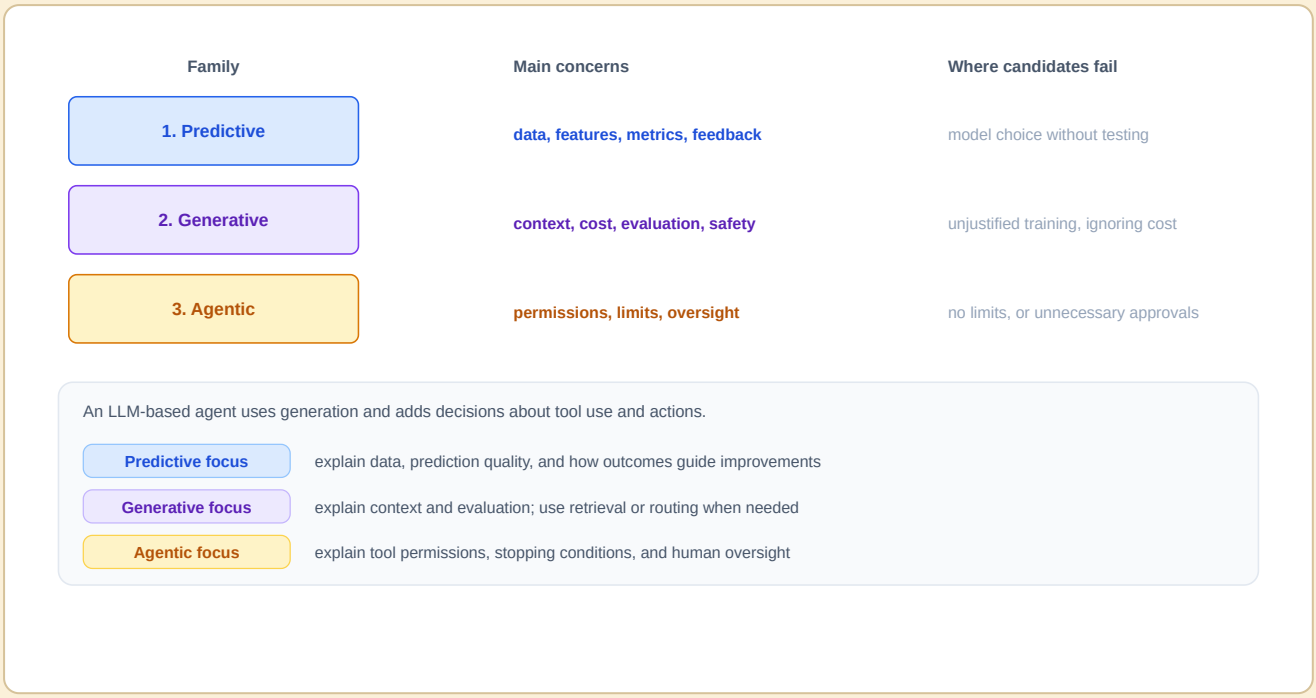


Figure — Predictive, generative and agentic compared (course diagram)

Terms: **label** = the target to predict. **Pretrained** = trained before your task. **Foundation model** = large model trained on broad data, adaptable to many tasks. **Token** = the text chunk an LLM reads/writes (word, fragment, punctuation). **Hallucination** = false or unsupported generated content. **Streaming** = showing output as it's generated. **Model routing** = picking which model handles a request by difficulty/cost.

"Design an AI shopping experience" and "Design fraud detection" need different designs, metrics and risks — classifying early avoids 30 minutes on the wrong system.

How it works — the output test

1. Does the system mainly output a **prediction, content, or an action**?
2. If vague, state your read and confirm.
3. Mixed products → split into subsystems (a streaming service ranks the homepage — predictive — and writes "why you might like this" blurbs — generative; a support platform routes tickets, drafts replies, issues refunds via an authorised tool — all three).
4. Name each subsystem and its interfaces, and ask which to go deep on — usually the riskiest.

This book focuses mostly on **family 1**; the AI Handbook covers LLM serving, RAG and agents. The framework works for all three.

Common mistakes

- An LLM for everything, e.g. one call per item to rank a 50M-item catalog (ignores personalisation data, evaluation, latency, cost). Sensible LLM roles are narrower — offline item enrichment, reordering a short list of finalists (~top 50) — and even then benchmark against a learned ranker on quality, speed and spend.
- Multi-agent designs when a fixed workflow would do (known steps → hard-code them).

In the interview

> **30-second answer** to an unclear prompt: *"First, what does the system output? A ranking or score puts my time on data, features and evaluation. Generated text shifts it to context, per-call spend and answer quality. Actions put tools, permissions and limits at the centre. My read is this is mostly ranking items per user — shall I go that way?"*

- **"It's all three."** → List subsystems and interfaces, go deep on the riskiest ("wrong refunds cost money"), let the interviewer redirect.
- *"Summarise tickets and auto-close duplicates" — family?* → Generative plus a small agentic part (closing) needing permission and review rules.

TL;DR

Predictions → data/features/metrics. Content → context/cost/quality/safety. Actions → permissions/limits/oversight. Mixed: split, connect, pick a focus.

15.3 The seven-step framework

What it is

A fixed order of decisions so requirements and data drive the architecture. It's the classic structure (requirements → estimates → API → high-level design → deep dive → wrap-up) with two ML concerns made explicit: **define quality early** and **show how the system keeps working after launch**.

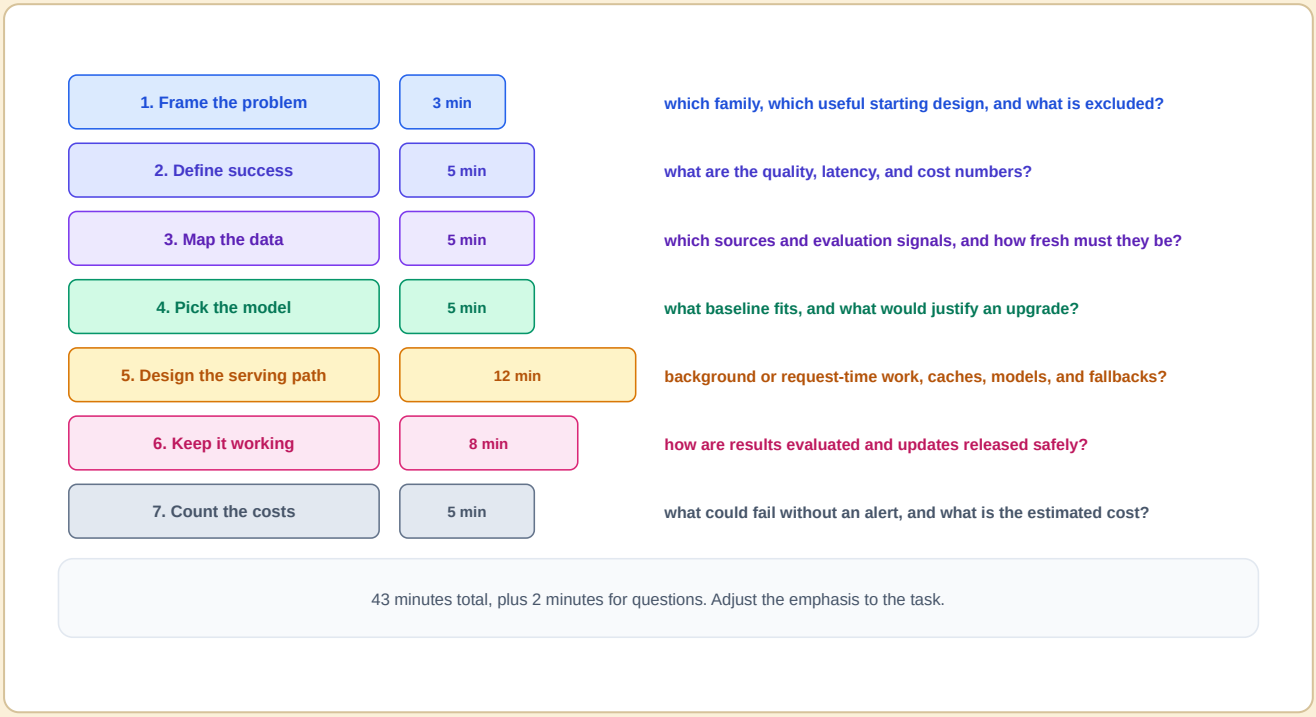
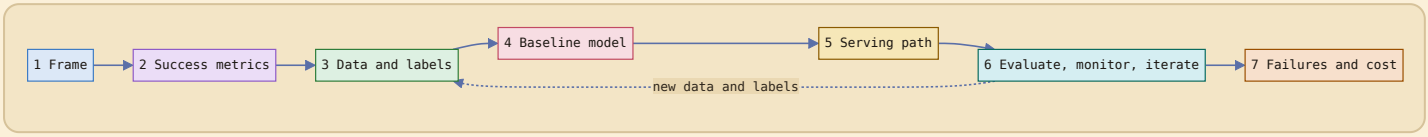


Figure — Seven steps with their guiding questions and minutes (course diagram)

Step	Question it answers	Minutes (45-min round)
1. Frame the problem	What family, what known system does it resemble, what's in/out of scope?	3
2. Define success	Quality, latency, scale and cost targets?	5
3. Map the data	Sources, label origin, required freshness?	5
4. Pick the model	Simplest baseline, and what would justify upgrading?	5
5. Design the serving path	Offline vs per request, caches, cheaper paths, failure handling?	12
6. Keep it working	How are outcomes logged, quality measured, changes evaluated and rolled out?	8
7. Count failures and cost	What breaks, what controls exist, cost vs budget?	5
Buffer	Questions	2

The same structure is in Aminian & Xu's *ML System Design Interview* (clarify requirements → frame as ML task → data preparation → model development → evaluation → deployment/serving → monitoring) and Chip Huyen's *Designing ML Systems* lifecycle (scoping → data engineering → model development → deployment → monitoring and continual learning → business analysis). Any wording; cover every step. Shift time toward data/models for predictive, inference/context for generative, permissions for agents.



Step 1 — Frame the problem (and frame it as an ML task)

1. **Name the family** (output test).
2. **Name the known system** it resembles ("a recommender where items are podcasts") and say which requirements differ — don't paste a memorised design.
3. **State scope**: "English text support first; voice and other languages later. OK?"
4. For predictive problems, **frame as an ML task** (Aminian & Xu's step 2):
    - **Input and output** (input = user, video, context; output = P(user watches ≥ 30 s)).
    - **ML category**: binary/multiclass classification, regression, ranking (learning-to-rank), retrieval, clustering, or generation.
    - **ML objective**: what the model optimises, chosen because it moves the business metric.

Business metric → ML objective (worked examples)

Business goal / metric	ML task	ML objective (what the model predicts)	Why not the obvious one
Increase time spent on video platform	Ranking	P(watch ≥ X seconds) or expected watch time	Raw clicks reward clickbait (watch-time objectives risk short-clip bias — guard with surveys, retention)
Reduce fraud losses without annoying customers	Binary classification	P(fraud) → thresholds set by expected cost	Accuracy is meaningless at 0.2% fraud
Accurate delivery promise at checkout	Regression	Minutes from order to delivery	Business also needs a *promise-kept rate*, handled by a display buffer, not the model
Improve marketplace search revenue	Ranking	P(purchase | query, item), possibly × price	Relevance-only ranking may ignore **conversion** (view → purchase)
Reduce subscription churn	Binary classification	P(cancel within 30 days)	Label arrives 30 days late — shapes monitoring

Rules of ML: **Rule #12 — don't overthink which objective you directly optimise** at first; **Rule #13 — pick a simple, observable and attributable metric** for the first objective (a click is attributable to the item shown; "user happiness" is not).

Step 2 — Define success

Numbers; state assumptions when not given.

1. **Quality** — a *business/online metric* plus a correlated *offline ML metric*:
   - "Users find videos they enjoy" → online: watch time, CTR, with **guardrails** (metrics that must not worsen: satisfaction surveys, policy-violating content); offline: NDCG / recall@k (Chapter 13) on **held-out** interactions.
   - "Helpful support bot" → correct resolution rate, CSAT, repeat-contact rate.
2. **Latency** — a feed in a few hundred ms; a chatbot has a **time-to-first-token** target *and* a full-answer target; a batch job has a completion deadline and freshness target. Latency decides what must be precomputed.
3. **Scale and cost** — users, requests/day, **peak QPS** (System Design Fundamentals), cost per request or per *completed task*. A free feature needs tiny per-request cost; a support bot can cost more if it really resolves tickets. A concrete budget, never "near zero".

Rules of ML **#2: design and implement metrics first** — log what you'll need before building.

Step 3 — Map the data

Sources: interaction logs, catalog/content, user profiles (where permitted), documents for generation. Then:

1. **Label source** — often late or imperfect:
   - Fraud: confirmed investigations and **chargebacks** (disputed card payments reversed) — weeks late, and not every chargeback is fraud.
   - Feeds: clicks — instant but only a **proxy** for satisfaction.
   - Support: reviewed conversations, verified resolutions; a thumbs-down is feedback, not a correctness label.
   - **Self-supervised learning** (labels made from the data itself, e.g. next word) and pretrained models reduce hand-labelling, but you still need labelled *evaluation* data.
2. **Freshness** — daily for long-term tastes; minutes or seconds for trending content and fraud. Decides batch vs **streaming** (continuous event processing) vs request-time computation.

Name **feature groups** (entity, recent-behaviour aggregates over time windows, request context, cross features such as user × category — §16.4) and flag **leakage** (features unavailable at prediction time) and **privacy** limits.

Step 4 — Choose a model, beginning with a baseline

A **baseline** is the least elaborate approach meeting basic requirements: it tests the pipeline end to end and is the yardstick for everything later (*Rules of ML #4: keep the first model simple and get the infrastructure right*) — and sometimes it's good enough.

Family	Baseline	Upgrade path	Trigger to upgrade
Predictive	Popularity / rules / per-entity averages; logistic regression	GBDT on tabular features → deep models (two-tower retrieval, multitask rankers — Chapter 13)	Baseline misses the offline target and the gain is worth the complexity
Generative	Existing model + clear instructions; add retrieval for private/fresh data	Better retrieval, model routing, fine-tuning	Evaluation shows a failure prompts/retrieval don't fix
Agentic	Fixed workflow if steps are known; else one agent with a few narrow tools	More tools, then multiple agents	Separation clearly helps and justifies coordination cost

Always state the condition: *"If the baseline misses the target, I'll compare a GBDT on the same held-out weeks."*

Code — baseline first (`code/baseline_first.py`)

Synthetic delivery-time task: naive baseline → heuristic → GBDT, then model error → the business metric "orders delivered within the promised time" via a display buffer.

```
# 1) naive baseline: everyone gets the global average
rows.append(report("global mean", yte, np.full(len(yte), ytr.mean())))

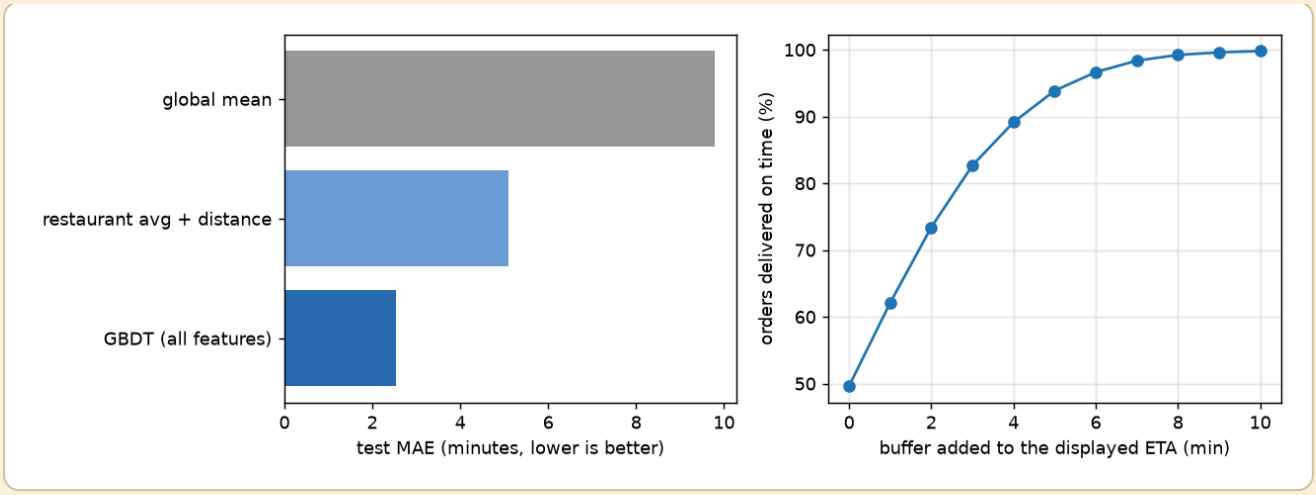
# 2) heuristic: per-restaurant average + a distance term (shippable in a week)
rest_avg = {r: ytr[Xtr[:, 0] == r].mean() for r in np.unique(Xtr[:, 0])}
resid = ytr - np.array([rest_avg[r] for r in Xtr[:, 0]])
dist_fit = LinearRegression().fit(Xtr[:, [1]], resid)
heur = np.array([rest_avg.get(r, ytr.mean()) for r in Xte[:, 0]]) + dist_fit.predict(Xte[:, [1]])

# 3) learned model: gradient-boosted trees on all features (restaurant id as categorical)
gbdt = HistGradientBoostingRegressor(random_state=0, categorical_features=[0]).fit(Xtr, ytr)
pred = gbdt.predict(Xte)

# business metric: % of orders that arrive within (prediction + buffer)
kept = [(yte <= pred + b).mean() * 100 for b in buffers]
```

```
model                          MAE  p90 abs err
global mean                   9.80        19.92
restaurant avg + distance     5.11        10.45
GBDT (all features)           2.54         5.31

buffer (min) -> promise kept % / avg promised minutes   [GBDT]
  + 0 min ->  49.6%  /  37.8
  + 3 min ->  82.7%  /  40.8
  + 5 min ->  93.9%  /  42.8
  + 8 min ->  99.2%  /  45.8
```



- (MAE in minutes; **p90 abs err** = error size 90% of predictions stay within.) Every number is judged **against a baseline**: heuristic halves naive error, GBDT halves it again (synthetic data — real gains are usually smaller).
- An unbiased regressor's errors are roughly symmetric, so it's late **half the time** (49.6% kept). The business metric depends on a **product decision** (the buffer): +5 min keeps ~94% but makes the ETA look slower. Keep that knob outside the model so product owns it.

Step 5 — Design the serving path

Draw the architecture and trace one request. Whatever the shape, answer four questions:

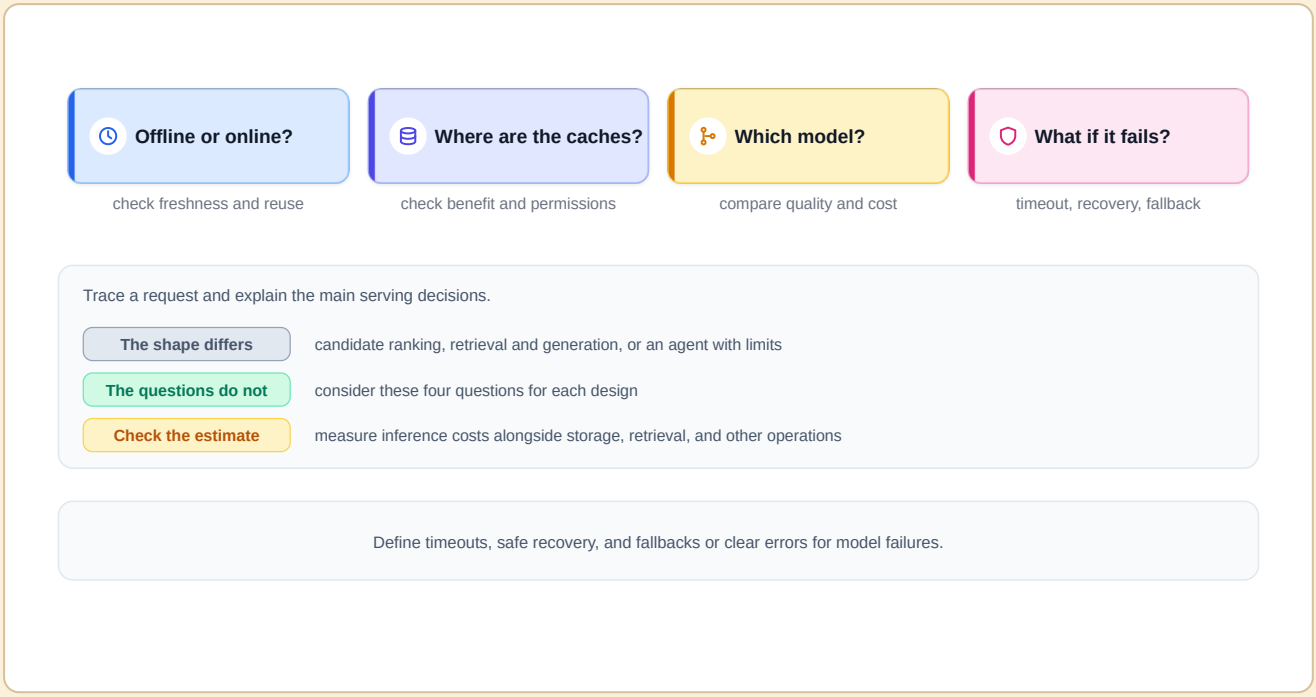
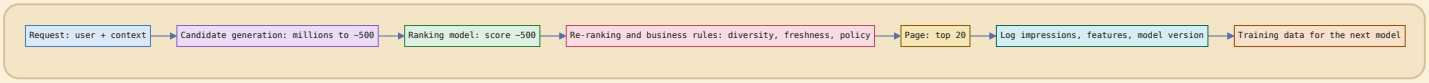


Figure — Step 5 checklist: offline/online, caches, model choice, failure (course diagram)

1. **Offline or online?** What can be precomputed vs needs request-time data (§16.5)?
2. **Where are the caches?** Latency/cost saved, freshness, **invalidation**, and **permission** rules so no user sees another's cached result.
3. **Which requests can take a cheaper path** (smaller model, cached answer, rule)?
4. **What if the model is slow, down or rate-limited?** Timeouts, bounded retries (only when safe), a fallback (previous model, cached score, heuristic, unpersonalised list) or a clear failure; for batch, deadlines and recovery.

A typical predictive design for a large catalog is a **funnel**:



Step 6 — Keep it working (evaluation, monitoring, iteration)

Show the feedback process **in the diagram**:

1. **Log** inputs, outputs, features used and model version (privacy controls, retention limits).
2. **Capture outcomes**: clicks, ratings, escalations, confirmed fraud, actual delivery time.
3. **Measure** step-2 metrics on live outcomes and human-reviewed samples.
4. **Improve** data, features, retrieval, prompts or model.
5. **Release safely**: offline evaluation → **shadow deployment** (new version scores copied traffic, outputs hidden, side effects disabled) and/or **A/B test** → gradual ramp with explicit stop/**rollback** conditions.

	Offline evaluation	**Online evaluation**
Data	Held-out historical data (time-based split)	Live users
Metrics	ML metrics: PR-AUC, NDCG, recall@k, MAE, **calibration** (predicted probabilities match real frequencies)	Business metrics: CTR, watch time, revenue, resolution rate + guardrails
Cost / risk	Cheap, no user risk	Real traffic; needs time and **statistical power**
Blind spots	Proxy labels, feedback loops, serving-time degradation	Slow, noisy, needs good experiment design
Role	Filter candidates	Decide the launch

(Details §16.7, §16.9.) Monitoring in one line: system health → input feature distributions → prediction distributions → outcome quality (as labels arrive), sliced per segment. Retrain on a schedule justified by measured decay or after an investigation — never auto-promote because an alert fired.

Step 7 — Failure modes and the bill

Failure list specific to your system — unsupported answers, drifting data, unsafe actions, feedback loops reinforcing bad behaviour (clicks-only → clickbait) — each with a control (thresholds, quality monitors, content checks, permission limits, a **kill switch** to turn the feature off) and what the control *cannot* guarantee.

Cost ≈ requests × share hitting each model × calls per request × cost per call, plus retrieval, storage, retries and human escalation.

Code — napkin math (`code/napkin_math.py`)

```python
def qps(per_day, peak_factor=3.0):
    avg = per_day / 86_400
    return avg, avg * peak_factor

def model_cost_per_day(requests_per_day, share, calls_per_request, cost_per_call):
    return requests_per_day * share * calls_per_request * cost_per_call
```

```
10M req/day -> avg 116 QPS, peak ~347 QPS (3x peak factor assumed)
all requests x $0.01        -> $100,000/day, $3,000,000/30-day month
20% to large model x $0.01  -> $20,000/day (excludes small model, retrieval, retries)
support: 40,000 conversations/day x $0.10 = $4,000/day model cost
  compare per resolved issue, not per call: human ticket assumed $4.00
  resolve 30%: expected cost/ticket $2.90 vs $4.00 human-only
  resolve 50%: expected cost/ticket $2.10 vs $4.00 human-only
  resolve 70%: expected cost/ticket $1.30 vs $4.00 human-only
precompute 25M users x 500 recs = 12.5B rows; active-only (5M) = 2.5B rows
```

The habit is the **unit**: cost *per resolved issue* (AI cost + human cost for the unresolved share), not one model call vs one human ticket. All prices are illustrative assumptions.

> Short on time? Steps 6–7 in two sentences: *"Every prediction is matched to its eventual outcome; when quality moves I dig in, and any new version passes offline tests, then shadow or **canary** (a small slice of real traffic), before a staged ramp — with cost per request on the dashboard. The two risks I'd watch hardest are X and Y."*

Trade-offs of the framework

Benefit	Cost
Decisions follow requirements and data	Feels slow if the interviewer wants the diagram — state assumptions in one sentence each and jump to step 5
You don't forget evaluation/monitoring/cost	Following it rigidly when steered elsewhere looks robotic — adapt order and depth

When not to use it rigidly

- "Skip to the architecture": state quality and latency as assumptions, draw, revisit requirements when they change a choice.
- "You have 10 minutes": scope, main architecture, biggest risk, cost — one sentence each on quality target, data, evaluation.

In the interview

30-second opener: *"Plan: agree the task and boundaries; put numbers on quality, latency and spend; find the data and label source. Then a baseline and the serving path. I'll close with evaluation, safe rollout, failures, and cost. To start — this reads like a ranking problem…"*

- **"Why a baseline if GBDT will win?"** → Validates pipeline and logging, gives a number to beat, sometimes good enough (Rules #1, #4).
- **"What's your ML objective?"** → Name the target, how it moves the business metric, the proxy risk and the guardrail that catches it.
- **"How will you know it works after launch?"** → Four monitoring layers, outcome join, per-segment quality, drift alerts → investigation, gated retraining.

TL;DR

Frame (3) → Success (5) → Data (5) → Baseline (5) → Serving (12) → Keep it working (8) → Failures & cost (5). Observable ML objective, beat a baseline, show the loop, price it.

15.4 Capstone — scoping a vague AI question

The exercise

Hypothetical prompt (numbers are assumptions):

> "We're a food-delivery company: ~25 million monthly users, 500,000 orders a day. Leadership wants AI to improve customer experience. Design something."

Feature, metric and family are unspecified. The first ten minutes turn this into an agreed, measurable problem — not an architecture.

Step A — Generate options by family, then choose

Family	Options
Predictive	ETA prediction; restaurant/dish recommendations; fraud on refund claims; predictions feeding courier assignment (assignment itself also needs optimisation/rules)
Generative	Support assistant for order status and refund policy; AI-written dish descriptions; review summaries
Agentic	"Dinner for four under $60" assistant that places an order after checks and confirmation

Choose on **impact × feasibility × risk** (feasibility = can the team build and run it with available data, time, people):

- **ETA** — touches every order; delivered orders give labels (check completeness).
- **Support assistant** — savings show directly in support spend; order records and policy documents likely exist; risk limited by restricted actions and escalation.
- **Ordering agent** — spends users' money; needs payment authorisation, confirmation, cancellation. Feasible, more controls.

No universally right pick — state your preference under these assumptions and what would change your mind.

Step B — Framework opening (support assistant chosen)

1. **Quality**: primary = *correct resolution without escalation*; plus CSAT and repeat contacts (closed ≠ fixed). Human ticket assumed $4. Target: resolve ≥ 50% of eligible **tier-1** tickets (order status, coupons, simple cancellations) with CSAT no worse than humans. Explaining the cancellation policy differs from *cancelling* — actions need their own authorisation and rule checks.
2. **Latency**: stream; p95 (95% of requests meet it) time-to-first-token ≤ 1 s, p95 completion ≤ 5 s for short answers; tool calls may get separate targets.
3. **Scale and cost**: 500,000 orders × 8% contact rate = **40,000 conversations/day** (average; peaks and repeats separately). 10 calls × $0.01 → $0.10/conversation → **$4,000/day**. Compare *total cost per correctly resolved issue* (retrieval, infrastructure, retries, escalations included).
4. **Out of scope**: voice, non-English. Refunds above an assumed $10 go to a human; smaller ones still need identity, eligibility and duplicate-payment checks enforced **outside** the model.

Step C — Classify and preview the shape

"At its core this generates answers grounded in order data the user may see and retrieved policy text. A refund tool adds an agentic slice, so reading data and moving money get separate safeguards. Deep dives: retrieval accuracy, answer grading, cost per conversation, and backend checks guarding refunds ≤ $10."

Sample opening (paraphrased dialogue)

- **Interviewer:** "Use AI to improve customer experience."
- **You:** "Three candidates: ETA prediction, a support assistant, or an ordering agent — the agent needs payment and confirmation controls. I'd start with support: savings are measurable against current support costs and we likely have order data and policies. Support, or ETA?"
- **Interviewer:** "Support works."
- **You:** "Target: correctly resolve at least half of eligible routine tickets, CSAT at least human-level, no rise in repeat contacts. 500k orders × 8% ≈ 40k conversations a day. p95 first output within 1 s, completion within 5 s. Ten one-cent calls is ten cents a conversation before retrieval and infrastructure, so we track total spend per issue actually resolved. V1 skips voice and non-English. Refunds above $10 go to a person; smaller ones still pass authorisation and policy checks. So: answer generation over orders and policies plus a refund action; deep dive on retrieval, grading, spend and guarding the action — agreed?"

Now you (10 minutes each, out loud)

Cover options by family, one choice with reasons, quality/latency/scale/cost targets with assumptions flagged, scope, family and key decisions.

1. "Our music-streaming app wants listeners to discover more. Where does AI help?"
2. "Bank leadership wants something AI-powered in the mobile app." (Weigh the damage of a wrong output: personal money advice is far riskier than explaining a purchase's category; forecasts and fraud warnings still need watching. No family is automatically safe — say what it may do, must not do, and when a human reviews.)
3. "Candidates and jobs on our job board are poorly matched. Fix it with AI."

> **TL;DR**
>
> Vague prompt → options by family → choose on impact, data, feasibility, risk → quality/latency/scale/cost numbers in comparable units → confirm scope, family, key decisions.

15.5 Common mistakes (the whole interview)

Mistake	Better
"I'd use a transformer" in minute 2	Frame, success and data first; baseline before deep models
Optimising a proxy with no guardrail (clicks → clickbait)	Pair the objective with guardrails (satisfaction, retention, policy violations)
"We'll have labels"	Where from, how late, how noisy, who pays
Random train/test split on time-dependent data	Time-based split; **point-in-time-correct** features (only data available at that moment) (Rules of ML #33)
Accuracy on imbalanced data	Precision/recall at a threshold, PR-AUC; threshold owned by product/risk
No fallback	Timeouts + fallback prediction agreed with product and risk owners
Only the request path	Logging → outcome join → evaluation → gated retraining → staged rollout
Training features computed differently from serving	One feature definition; log serving features and train on them (Rules of ML #29, #32)
"GPU cluster" for a small tabular model	Benchmark on CPU first
Auto-retraining on every drift alert	Alert → investigate (often a broken pipeline) → validated retrain → gates
Hand-waving cost	Requests × share × calls × price, per completed task
Ignoring the interviewer's redirect	The framework is a checklist, not a script

Chapter TL;DR

- ML design = classic system design **plus** quality metrics, a data/label story, baseline-first modelling, a feedback loop, cost realism.
- Identify the **form** and the **family** (prediction, content, action) in the first minute; confirm depth.
- Seven steps in ~43 minutes: frame (3), success (5), data (5), model (5), serving (12), keep-it-working (8), failures and cost (5).
- Business metric → observable ML objective; product knobs (buffers, thresholds) outside the model.
- Numbers mean something only next to a baseline; every design needs a fallback; cost goes in the same unit as benefit.
- Next: Chapter 16 builds each production piece and assembles an ETA platform.