Chapter 7 — Functions

A **function** is a named, reusable block of code that takes some inputs, does one job, and (usually) hands back a result. Almost everything you write in Python beyond a ten-line script is built out of functions, and many "advanced" features — decorators, callbacks, `sorted(key=...)`, caching — are just functions being passed around like ordinary values.

This chapter goes from `def` to closures, decorators and recursion. Every script lives in `code/` and was run with Python 3.14 (all code is 3.12-compatible); the output shown is real.

What you'll learn

#	Topic	Script
7.1	Defining and calling functions, `return`, `None`, `pass`	`01_defining_and_arguments.py`
7.2	Every kind of argument: positional, keyword, default, `*args`, `**kwargs`, positional-only `/`, keyword-only `*`, unpacking, the mutable-default trap, returning several values	`01_defining_and_arguments.py`
7.3	Scope and namespaces: LEGB, `global`, `nonlocal`	`02_scope_and_closures.py`
7.4	Closures and the late-binding trap	`02_scope_and_closures.py`
7.5	First-class and higher-order functions, `lambda`, `map`/`filter`, `sorted`/`min`/`max` with `key`	`03_first_class_and_functools.py`
7.6	Decorators and `functools`: `wraps`, `partial`, `reduce`, `lru_cache`, `cache`	`03_first_class_and_functools.py`
7.7	Recursion, the call stack, the recursion limit, memoization	`04_recursion.py`
7.8	Docstrings and type annotations	`05_docstrings_annotations.py`

7.1 Defining and calling functions

What it is

You **define** a function with the `def` keyword, a name, a pair of parentheses holding the **parameters** (the input names), and a colon. The indented lines below form the **body**. Defining a function does not run it — it only creates a function object and binds it to the name. You **call** (run) it by writing its name followed by parentheses, passing **arguments** (the actual values) inside.

Analogy: a function definition is a recipe card; a call is actually cooking the dish. You can cook from the same card as many times as you like, with different ingredients each time.

> **Parameter vs argument.** The *parameter* is the variable name in the definition (`def cube(x)` → `x`). The *argument* is the value supplied at the call (`cube(3)` → `3`). People mix them up constantly; interviewers notice.

Why it matters

Functions give you four things at once:

1. **Modularity** — a big program is split into small pieces, each with one clear job, which can be written, read and debugged on its own.

2. **Reuse** — write the logic once, call it from anywhere (or from other programs via modules, Chapter 8). Fixing a bug fixes it everywhere.
3. **Scope** — names created inside a function are private to it, so different parts of a program don't trample each other's variables (7.3).
4. **Readability and teamwork** — a well-named function (`send_invoice(order)`) documents intent, and different people can own different functions.

How it works

1. Python executes the `def` statement: it compiles the body, creates a *function object*, evaluates any default values (important later), and binds the object to the name.
2. At a call, Python creates a fresh **frame** (a workspace holding the local variables), binds each parameter to its argument, and runs the body.
3. `return value` ends the call immediately and gives `value` back to the caller. If the body finishes without `return` — or with a bare `return` — the call gives back `None`.
4. The frame is thrown away (unless a closure keeps some of it alive, 7.4).

`pass` is a statement that does nothing. Python forbids an empty block, so `pass` is the placeholder when you want to define a function (or class, or `if` branch) now and fill it in later. `...` (Ellipsis) is often used the same way in stubs.

```
def say_hi():
    """A function with no parameters and no return statement."""
    print("hi from say_hi()")

def cube(x):
    """Return x cubed."""
    return x ** 3

def todo():
    pass  # placeholder body - keeps the syntax valid

say_hi()
print("say_hi() returned:", say_hi())
print("cube(3) =", cube(3))
print("todo() returned:", todo())
```

Output:

```
hi from say_hi()
hi from say_hi()
say_hi() returned: None
cube(3) = 27
todo() returned: None
```

Notice that `say_hi()` printed twice: once for the plain call and once inside the `print(...)` — and then `None` appeared, because a function that only *prints* still *returns* `None`.

Trade-offs: print vs return

	`print(...)` inside the function	`return ...`
Result usable by caller	No — it goes to the screen	Yes — can be stored, tested, combined
Testable	Only by capturing stdout	Directly (`assert cube(3) == 27`)
Good for	Debugging, CLI output at the edges	Almost all real logic

Common mistakes

- **Forgetting the parentheses**: `result = cube` stores the *function*, not a result. `result = cube(3)` calls it.
- **Printing instead of returning**, then wondering why `x = f()` is `None`.
- **Code after `return`** in the same block never runs.
- **Calling before defining** at module level: the `def` must have executed before the call runs (inside other functions this is fine, because the lookup happens at call time).

Real-world examples

Every library you'll touch is a collection of functions: `len`, `print`, `open`, `json.dumps`, `requests.get`, `numpy.mean`. Web frameworks such as Flask and FastAPI map each URL to a plain Python function ("view" or "endpoint") that returns the response.

In the interview

> **30-second answer:** "A function is defined with `def`; defining creates a function object bound to a name, and calling it creates a new frame with its own local variables. `return` sends a value back; with no `return` the result is `None`. Functions give modularity, reuse and a private scope."

- *What does a function without `return` return?* → `None`.
- *Are functions objects?* → Yes; they have attributes (`__name__`, `__doc__`, `__defaults__`) and can be passed around (7.5).
- *What is `pass` for?* → A no-op placeholder where syntax needs a statement.

Self-check

1. What does `x = print("a")` store in `x`? → `None`.
2. Parameter or argument: the `3` in `cube(3)`? → Argument.
3. When does the body of a function run? → Only when it's called, not when `def` runs.
4. Why prefer `return` over `print` in logic? → Results can be reused and tested.

TL;DR

`def name(params):` + indented body. Defining ≠ running. `return` gives a value back; otherwise you get `None`. `pass` fills an empty body.

7.2 Function arguments — every kind

What it is

Python gives you several ways to feed values into a function. At the **call site** you can pass arguments **positionally** (matched by order) or **by keyword** (matched by name). In the **definition** you can add **defaults**, collect extras with `*args` / `**kwargs`, and force a style with `/` (positional-only) and `*` (keyword-only).

The full shape of a parameter list, in the only order Python allows:

```
def f(pos_only, /, pos_or_kw, *args, kw_only, **kwargs): ...
#     ^^^^^^^^     ^^^^^^^^^^  ^^^^^  ^^^^^^^  ^^^^^^^^
#     before /     normal      extra  after *  extra keywords
#                              positionals     (dict)
#                              (tuple)
```

Kind	Declared as	Called as	Notes
Positional-or-keyword	`def f(a, b)`	`f(1, 2)` or `f(b=2, a=1)`	The default kind
Default	`def f(a, b=10)`	`f(1)` or `f(1, 5)`	Params with defaults must come after those without (in the positional part)
Variable positional	`def f(*args)`	`f(1, 2, 3)`	`args` is a **tuple**
Variable keyword	`def f(**kwargs)`	`f(x=1, y=2)`	`kwargs` is a **dict** (insertion-ordered)
Positional-only	`def f(a, b, /)`	`f(1, 2)` only	Python 3.8+ (PEP 570)
Keyword-only	`def f(*, a, b)` or after `*args`	`f(a=1, b=2)` only	Can be required (no default)

Why it matters

Good argument design makes calls **self-documenting and hard to misuse**. `connect("db", 5432, True, False)` is a riddle; `connect("db", port=5432, ssl=True, retry=False)` explains itself. Defaults keep simple calls simple, `*args`/`**kwargs` let you write flexible wrappers, and `/`/`*` let a library author keep freedom to rename parameters later.

How it works — binding arguments to parameters

When you call a function, Python:

1. Fills parameters left-to-right from the **positional arguments**. Extra positionals go into `*args` (or raise `TypeError` if there is no `*args`).
2. Matches each **keyword argument** to the parameter with that name. Unknown names go into `**kwargs` (or raise `TypeError`).
3. A parameter filled twice (by position *and* by keyword) → `TypeError: got multiple values`.
4. Any still-empty parameter takes its **default**; if it has none → `TypeError: missing ... required argument`.

At the call site, positional arguments must come before keyword arguments: `f(1, b=2)` is fine, `f(a=1, 2)` is a `SyntaxError`.

Positional and keyword arguments

Positional arguments are matched purely by order, so every required one must be supplied and in the right order — swap two and you silently get wrong results (the function can't tell a first name from a last name). **Keyword** arguments name the target parameter, so order no longer matters and the call reads clearly; they shine when a function has many parameters or several optional ones.

Default arguments

A **default** makes a parameter optional: if the caller leaves it out, the default is used; if the caller supplies it, it overrides the default.

```python
def book_flight(origin, destination, seat="economy", *, meal="standard"):
    """origin/destination: positional-or-keyword; seat: default; meal: keyword-only."""
    return f"{origin}->{destination} [{seat}, meal={meal}]"

print(book_flight("SFO", "JFK"))                                    # defaults used
print(book_flight(destination="JFK", origin="SFO", seat="aisle"))  # any order by name
print(book_flight("SFO", "JFK", meal="vegan"))                      # keyword-only
book_flight("SFO")                          # missing arg
book_flight("SFO", "JFK", "aisle", "vegan")  # meal given positionally
book_flight("SFO", origin="LAX")             # same arg twice
```

Output (errors caught by a small helper in the script):

```
SFO->JFK [economy, meal=standard]
SFO->JFK [aisle, meal=standard]
SFO->JFK [economy, meal=vegan]
  missing arg: TypeError: book_flight() missing 1 required positional argument: 'destination'
  meal given positionally: TypeError: book_flight() takes from 2 to 3 positional arguments but 4 were
given
  same arg twice: TypeError: book_flight() got multiple values for argument 'origin'
```

Positional-only (`/`) and keyword-only (`*`)

- Everything **before** `/` must be passed by position. Their names become an internal detail, so the author may rename them later without breaking anyone, and a name like `value` can't collide with a `**kwargs` key. Many built-ins work this way: `len(obj, /)`, `abs(x, /)`, `divmod(x, y, /)` — try `len(obj=[1])` and you get a `TypeError`.
- Everything **after** `*` (or after `*args`) must be passed by name. Use it for flags and options whose meaning isn't obvious from position: `sorted(iterable, /, *, key=None, reverse=False)` — you can't write `sorted(xs, None, True)`. Keyword-only parameters may still be *required* (no default).

```python
def clamp(value, low, high, /):
    return max(low, min(value, high))

print("clamp(15, 0, 10) =", clamp(15, 0, 10))
clamp(value=15, low=0, high=10)
```

```
clamp(15, 0, 10) = 10
  clamp(value=15, ...): TypeError: clamp() got some positional-only arguments passed as keyword
arguments: 'value, low, high'
```

Variable-length arguments: `*args` and `**kwargs`

`*args` gathers any **extra positional** arguments into a tuple; `**kwargs` gathers any **extra keyword** arguments into a dict. The names `args`/`kwargs` are only convention — the `*` and `**` do the work.

```
def total(*numbers, start=0):
    print("  numbers is a", type(numbers).__name__, numbers)
    return sum(numbers, start)

def make_tag(tag, /, text, **attrs):
    attr_str = "".join(f' {k}="{v}"' for k, v in attrs.items())
    return f"<{tag}{attr_str}>{text}</{tag}>"

print("total(1, 2, 3) =", total(1, 2, 3))
print("total() =", total())
print(make_tag("a", "docs", href="https://docs.python.org", target="_blank"))
```

```
  numbers is a tuple (1, 2, 3)
total(1, 2, 3) = 6
  numbers is a tuple ()
total() = 0
<a href="https://docs.python.org" target="_blank">docs</a>
```

A typical `**kwargs` use is a settings function: `configure(host="db1", port=5432, timeout="30s")` receives `{"host": "db1", "port": 5432, "timeout": "30s"}` and can loop over `.items()`.

Unpacking at the call site

The same stars work in reverse when *calling*: `*iterable` spreads a sequence into positional arguments, `**mapping` spreads a dict into keyword arguments.

```
nums = [4, 5, 6]
opts = {"seat": "window", "meal": "kosher"}
print("total(*nums) =", total(*nums))
print(book_flight("SFO", "JFK", **opts))
```

```
  numbers is a tuple (4, 5, 6)
total(*nums) = 15
SFO->JFK [window, meal=kosher]
```

Returning several values

A function returns exactly **one** object — but that object can be a tuple, and `return a, b, c` builds one. The caller then *unpacks* it.

```
def min_max_mean(values):
    return min(values), max(values), sum(values) / len(values)

result = min_max_mean([3, 9, 6])
print("raw result:", result, type(result).__name__)
lo, hi, avg = result
```

```
raw result: (3, 9, 6.0) tuple
lo=3 hi=9 avg=6.0
```

For more than ~3 values, return a `dataclass` or `typing.NamedTuple` instead, so callers use names (`stats.mean`) not positions.

The mutable-default pitfall

Default values are evaluated once — when `def` runs — not on each call. If the default is a mutable object (list, dict, set), every call that uses the default shares *the same object*, and changes pile up.

```
def append_bad(item, bucket=[]):          # BUG: default evaluated once
    bucket.append(item)
    return bucket

def append_good(item, bucket=None):       # FIX: None sentinel
    if bucket is None:
        bucket = []
    bucket.append(item)
    return bucket

print(append_bad("a"), append_bad("b"))
print("default is now:", append_bad.__defaults__)
print(append_good("a"), append_good("b"))
```

```
['a', 'b'] ['a', 'b']
default is now: (['a', 'b'],)
['a'] ['b']
```

Both `append_bad` calls returned the *same* list, and `__defaults__` shows the stored default has been mutated. The fix is the **`None` sentinel** idiom: default to `None` and create a new object inside. (The same "evaluated once" rule means `def log(msg, when=datetime.now())` freezes the time at definition.) `dataclasses` refuses a bare mutable default for exactly this reason and asks for `field(default_factory=list)`.

Trade-offs

Choice	Benefit	Cost
Positional	Short calls	Order-dependent, easy to swap by mistake
Keyword	Readable, order-free	Longer calls; parameter names become public API
Keyword-only (`*`)	Forces clarity for flags/options	Slightly more typing for callers
Positional-only (`/`)	Freedom to rename; no clash with `**kwargs`	Callers can't use names for clarity
`*args` / `**kwargs`	Flexible wrappers, forwarding	Hides the real signature from readers, IDEs and type checkers

When not to use it

- Don't reach for `**kwargs` when you know the options — explicit keyword parameters give errors on typos (`timout=5`) and IDE help; `**kwargs` silently swallows the typo.
- Don't pile on boolean positional flags (`f(x, True, False)`) — make them keyword-only.

Common mistakes

- Mutable default (`=[]`, `={}`) — use `None`.
- Putting a non-default parameter after a default one: `def f(a=1, b):` → `SyntaxError`.
- Positional after keyword at the call: `f(a=1, 2)` → `SyntaxError`.
- Forgetting that `*args` is a **tuple** (immutable) — convert with `list(args)` if you need to modify it.

- Mixing up `return a, b` (one tuple) with "two return values".

Real-world examples

- `print(*objects, sep=" ", end="\n", file=None, flush=False)` — `*objects` plus keyword-only options.
- `requests.get(url, params=None, **kwargs)` forwards options such as `timeout=` and `headers=` through `**kwargs`.
- Matplotlib plotting calls accept `**kwargs` and pass them on to the underlying line/patch objects (`color=`, `linewidth=`).
- `dataclasses.field(default_factory=list)` exists because of the mutable-default trap.

In the interview

> **30-second answer:** "Arguments bind positionally by order or by keyword by name. Defaults make parameters optional but are evaluated once at definition time, so never use a mutable default — use `None`. `*args` collects extra positionals into a tuple, `**kwargs` extra keywords into a dict. Parameters before `/` are positional-only, after `*` keyword-only."

- *Why is `def f(x=[])` dangerous?* → The list is created once and shared by every call that uses the default; mutations accumulate.
- *Why would a library use positional-only parameters?* → To keep parameter names private (can be renamed without breaking callers) and to allow a `**kwargs` key with the same name.
- *How do you pass a list as separate arguments?* → `f(*my_list)`; a dict as keywords → `f(**my_dict)`.

Self-check

1. Inside `def f(*a, **k)`, what are the types of `a` and `k`? → `tuple` and `dict`.
2. `def f(a, /, b, *, c)` — which calls work: `f(1, 2, 3)`, `f(1, b=2, c=3)`, `f(a=1, b=2, c=3)`? → Only `f(1, b=2, c=3)` (and `f(1, 2, c=3)`).
3. What does `return 1, 2` return? → The tuple `(1, 2)`.
4. When is a default value computed? → Once, when the `def` statement runs.

TL;DR

Positional = by order, keyword = by name. Defaults are evaluated **once** → use `None` for mutables. `*args` → tuple, `**kwargs` → dict; `*`/`**` at the call site unpack. `/` = positional-only before it, `*` = keyword-only after it. Multiple return values = one tuple.

7.3 Scope and namespaces — LEGB, `global`, `nonlocal`

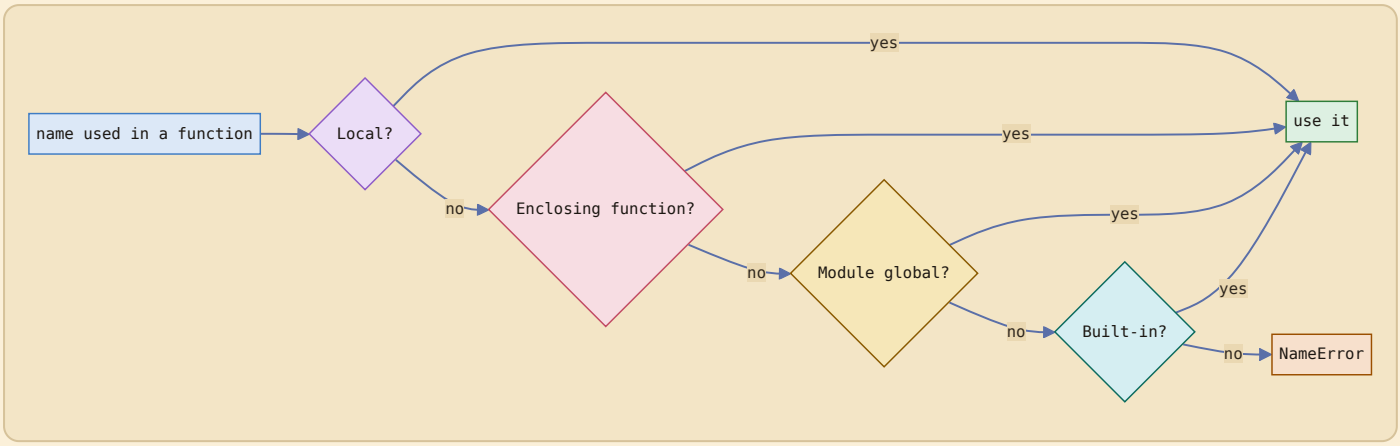
What it is

A **namespace** is a mapping from names to objects — literally a dictionary in most cases (`globals()` shows the module's one). Each name inside one namespace is unique, but the same name can exist in different namespaces without clashing. The **scope** of a name is the region of code where that name can be used directly.

Python resolves a bare name using the **LEGB rule** — it searches four scopes in order and stops at the first hit:

Letter	Scope	Where it comes from	Lifetime
L	Local	Names assigned inside the current function (including parameters)	Created when the call starts, gone when it ends
E	Enclosing	Locals of any *outer* function(s), for nested functions	As long as something references them (see closures)
G	Global	Names assigned at the top level of the **module** (file)	Until the program ends (or module is unloaded)
B	Built-in	`len`, `print`, `ValueError`… from the `builtins` module	Whole program

Not found anywhere → `NameError`.



Analogy: looking for scissors — first your own desk (local), then the room you're in (enclosing), then the whole house (global), then the hardware store (built-ins).

Why it matters

Scope is what makes functions safe to combine: a `total` inside one function can't clobber a `total` inside another. Misunderstanding it causes one of Python's most confusing errors, `UnboundLocalError`.

How it works — the key rule

Whether a name is local is decided at compile time, for the whole function body: if the function *assigns* to a name anywhere (`=`, `+=`, `for x in`, `import x`, `def x`), that name is local everywhere in the function — unless declared `global` or `nonlocal`. Merely *reading* a name doesn't make it local, so reading a global works without any declaration.

1. `global name` → assignments to `name` in this function rebind the **module-level** variable.
2. `nonlocal name` → assignments rebind the variable in the **nearest enclosing function** scope (it must already exist there; module globals don't count). Added in Python 3.
3. Only functions (and classes, comprehensions, lambdas) create scopes. **`if`, `for`, `while` and `with` blocks do not** — a variable assigned inside an `if` is visible after it.

```python
level = "global"                    # G

def outer():
    level = "enclosing"             # E for inner()
    def inner():
        level = "local"             # L
        print("  inner sees:", level)
    inner()
    print("  outer sees:", level)

def reads_global():
    print("  reads_global sees:", level)   # no assignment here -> global

counter = 0

def broken_increment():
    counter += 1                    # assignment makes counter LOCAL -> error

def global_increment():
    global counter
    counter += 1
```

```
== LEGB lookup ==
  inner sees: local
  outer sees: enclosing
  reads_global sees: global
  B (built-in) len is builtins.len: True

== assigning without global ==
  UnboundLocalError: cannot access local variable 'counter' where it is not associated with a value

== global ==
  counter = 2
```

`counter += 1` means `counter = counter + 1`. Because of the assignment, `counter` is local; the right-hand side then reads the local before it has a value → `UnboundLocalError` (a subclass of `NameError`).

`nonlocal` does the same job one level in — see `make_counter` in 7.4.

Trade-offs

Approach to shared state	Good	Bad
Pass in / return values	Explicit, testable, thread-friendly	More parameters
`global`	Quick for scripts	Hidden coupling; any function can change it; hard to test
`nonlocal` / closure	State private to one function family	Can be less obvious than a class
A class with attributes	Clear home for state + behaviour	More code

When not to use it

Avoid `global` for anything except true module-level configuration or caches in small scripts. Mutating a global *object* (e.g. `CACHE[key] = v`) doesn't even need `global` — only rebinding the *name*

does — but it is still shared mutable state.

Common mistakes

- `UnboundLocalError` from `x += 1` on a global — add `global x`, or better, return the new value.
- **Shadowing built-ins**: naming a variable `list`, `dict`, `id`, `input`, `sum` or `type` hides the built-in in that scope. The script shows `list = [1, 2]` followed by `list('xy')` → `'list' object is not callable`.
- Expecting a class body to be an enclosing scope for its methods — it isn't; methods must use `self.x` or `ClassName.x`.
- Expecting the loop variable of a **comprehension** to leak — it doesn't (comprehensions have their own scope), while a plain `for` loop's variable does remain after the loop.

Real-world examples

- `logging.getLogger(__name__)` is conventionally stored as a module-level (global) `logger` and only *read* inside functions.
- Configuration constants (`MAX_RETRIES = 3`) live at module level in UPPER_CASE by PEP 8 convention: global, but treated as read-only.

In the interview

> **30-second answer:** "Python looks names up via LEGB: local, enclosing, global, built-in. Any assignment in a function makes that name local for the whole function, which is why `x += 1` on a global raises `UnboundLocalError`. `global` rebinds module-level names, `nonlocal` rebinds the nearest enclosing function's name. Blocks like `if` and `for` don't create scopes."

- *Do you need `global` to append to a global list?* → No — `.append` mutates, it doesn't rebind. You need it only for `name = ...`.
- *What's the difference between `global` and `nonlocal`?* → `global` targets module scope; `nonlocal` targets the nearest enclosing *function* scope and requires the name to exist there.
- *Is a variable defined inside an `if` visible after it?* → Yes, `if`/`for` don't create scopes.

Self-check

1. What does LEGB stand for? → Local, Enclosing, Global, Built-in.
2. Why does `def f(): print(x); x = 1` fail even if a global `x` exists? → The later assignment makes `x` local for the whole body, so the `print` reads an unassigned local.
3. Can `nonlocal` refer to a module-level variable? → No — `SyntaxError: no binding for nonlocal ... found`.
4. What error do you get when a name isn't in any scope? → `NameError`.

TL;DR

Lookup order is **L → E → G → B**. Assignment anywhere in a function ⇒ local everywhere in it. `global` for module names, `nonlocal` for outer-function names. Prefer passing and returning values over either.

7.4 Closures

What it is

A **closure** is an inner function that *remembers* variables from the enclosing function's scope, even after the outer function has returned. Those remembered variables are called **free variables**; Python

stores them in "cells" attached to the function (`fn.__closure__`).

Analogy: a backpack. When the inner function leaves home (is returned), it takes a backpack with the variables it needs.

Why it matters

Closures let you build **functions from other functions** — configured functions (`make_multiplier(3)`), counters with private state, callbacks that carry context — without writing a class. Decorators (7.6) are built on them.

How it works

1. `make_multiplier(3)` runs; `factor = 3` lives in its local scope.
2. It defines `multiply`, which reads `factor` → the compiler marks `factor` as a free variable and stores it in a cell.
3. `make_multiplier` returns `multiply`. Its frame ends, but the cell survives because `multiply` references it.
4. Each call to `make_multiplier` creates a **new** cell, so `double` and `triple` are independent.
5. To *rebind* a captured variable (not just read it), the inner function needs `nonlocal`.

```
def make_counter():
    count = 0
    def step():
        nonlocal count              # rebind the enclosing variable
        count += 1
        return count
    return step

def make_multiplier(factor):
    def multiply(x):
        return x * factor           # 'factor' is captured (a free variable)
    return multiply

c1, c2 = make_counter(), make_counter()
print("  c1:", c1(), c1(), c1(), " c2:", c2())
double, triple = make_multiplier(2), make_multiplier(3)
print("  double(5) =", double(5), " triple(5) =", triple(5))
print("  free vars of double:", double.__code__.co_freevars,
      "cell value:", double.__closure__[0].cell_contents)
```

```
  c1: 1 2 3  c2: 1
  double(5) = 10  triple(5) = 15
  free vars of double: ('factor',) cell value: 2
```

The late-binding trap

Closures capture **variables, not values**. The value is looked up when the inner function *runs*.

```
funcs = [lambda: i for i in range(3)]
print([f() for f in funcs])          # [2, 2, 2]  - all see the final i
funcs = [lambda i=i: i for i in range(3)]
print([f() for f in funcs])          # [0, 1, 2]  - default captures the current value
```

The `i=i` trick works *because* defaults are evaluated at definition time (7.2). `functools.partial` is a cleaner alternative.

Trade-offs

	Closure	Class
Code size	Tiny	Larger
Number of operations	Best for one (a callable)	Many methods
Inspectability	State hidden in cells	Attributes visible, easy to debug

When not to use it

When the "function" needs several operations (increment, reset, read), or other people must inspect/serialise its state, write a small class instead. Closures also can't be pickled by the standard `pickle` module.

Common mistakes

- Late binding in loops (above).
- Forgetting `nonlocal` and getting `UnboundLocalError` on `count += 1`.

Real-world examples

- Every decorator (`@functools.lru_cache`, Flask's `@app.route("/")`) returns a closure wrapping your function.
- Event callbacks in GUI toolkits (Tkinter `command=`) often use closures or `partial` to carry the context of which button was pressed.

In the interview

> **30-second answer:** "A closure is a nested function that keeps references to variables of its enclosing scope after that scope has returned. It captures variables, not values, so lambdas in a loop all see the last value unless you bind with a default argument or `partial`. `nonlocal` lets it rebind the captured variable."

- *Where are captured variables stored?* → In cell objects in `fn.__closure__`; names in `fn.__code__.co_freevars`.
- *How to fix the loop-lambda bug?* → `lambda i=i: ...` or `functools.partial(f, i)`.

Self-check

1. Do `c1` and `c2` from `make_counter` share state? → No, each call creates a new cell.
2. Why does `[lambda: i for i in range(3)]` give `[2, 2, 2]`? → `i` is looked up at call time, when it's 2.
3. What keyword lets a closure change a captured variable? → `nonlocal`.

TL;DR

Inner function + remembered outer variables = closure. Captures variables (late binding), use `nonlocal` to rebind. Foundation of decorators.

7.5 First-class and higher-order functions, `lambda`, `map`/`filter`, `key=`

What it is

Python functions are **first-class objects**: you can assign them to variables, store them in lists and dicts, pass them as arguments and return them from other functions. A **higher-order function** is one

that takes a function as an argument or returns one (`map`, `sorted`, decorators, `make_multiplier`).

A `lambda` is an anonymous, single-expression function: `lambda x: x * 2` is the same as a `def` whose body is `return x * 2`, but it has no statements, no annotations and its name is just `<lambda>`.

Why it matters

It lets you pass **behaviour** as data: "sort by this", "retry this", "call this when done". It removes whole families of near-duplicate functions.

How it works

```
yell = shout                                  # assign to a name
table = {"up": str.upper, "low": str.lower}  # store in a dict
print(apply_twice(shout, "hey"))              # pass as argument  -> HEY!!

nums = [5, -2, 8, -7, 3]
list(map(abs, nums))                  # [5, 2, 8, 7, 3]
list(filter(lambda n: n > 0, nums))   # [5, 8, 3]
[n * n for n in nums if n > 0]        # [25, 64, 9]
```

- `map(fn, iterable)` applies `fn` to each item; `filter(pred, iterable)` keeps items where `pred(item)` is truthy. Both return **lazy iterators** in Python 3 (the script prints `map`) — wrap in `list()` to see values, and they can be consumed only once.
- `sorted(iterable, key=fn, reverse=False)` returns a **new list**; `list.sort()` sorts in place and returns `None`. `min`, `max`, `heapq.nlargest` and `itertools.groupby` accept `key=` too.
- `key` is called **once per element**; the elements are then compared by their key. Python's sort is **stable**: equal keys keep their original order — which is how multi-level sorting works.
- `operator.itemgetter(1)`, `operator.attrgetter("age")` and `str.lower` are ready-made, faster alternatives to many lambdas.

```
people = [("Ana", 31), ("bob", 25), ("Cy", 31), ("dee", 19)]
sorted(people, key=lambda p: p[1])            # by age
sorted(people, key=lambda p: (-p[1], p[0]))   # age desc, then name
sorted(["bob", "Ana", "dee", "Cy"], key=str.lower)
max(people, key=operator.itemgetter(1))
```

```
by age:             [('dee', 19), ('bob', 25), ('Ana', 31), ('Cy', 31)]
by age desc, name:  [('Ana', 31), ('Cy', 31), ('bob', 25), ('dee', 19)]
case-insensitive:   ['Ana', 'bob', 'Cy', 'dee']
itemgetter(1):      ('dee', 19)
oldest:             ('Ana', 31)
```

Returning a **tuple** from `key` sorts by the first element, then the second, and so on; negating a number flips that level's direction.

Trade-offs

Tool	Best for	Watch out
Comprehension	Transform + filter in one readable line	Builds a whole list (use a generator expression `( … )` for laziness)
`map` / `filter` with a named function	`map(str.strip, lines)`	With a lambda it's usually less readable than a comprehension
`lambda`	Tiny one-off `key=` functions	No statements, name `<lambda>` in tracebacks, PEP 8 says don't assign one to a name — use `def`

When not to use it

Don't write `f = lambda x: x + 1` — use `def f(x): return x + 1` (better name in tracebacks, can have a docstring). Don't nest `map(lambda..., filter(lambda..., ...))` when a comprehension reads plainly.

Common mistakes

- `sorted_list = my_list.sort()` → `None`. Use `sorted()` or call `.sort()` alone.
- Consuming a `map` / `filter` iterator twice — the second pass is empty.
- `key=len()` (calling it) instead of `key=len` (passing it).

Real-world examples

- pandas: `df["name"].map(str.title)`, `df.sort_values(key=...)`, `df.apply(fn)`.
- `sorted(files, key=os.path.getmtime)` — "newest file" scripts.
- Callbacks: `threading.Thread(target=worker)`, `concurrent.futures` `executor.map(fetch, urls)`.

In the interview

30-second answer: "Functions are objects, so they can be passed, stored and returned; a higher-order function takes or returns one. `lambda` makes a one-expression anonymous function, mostly for `key=`. `map` / `filter` return lazy iterators; comprehensions are usually preferred. `sorted` with `key` computes the key once per element and is stable, so tuple keys give multi-level sorts."

- *`sorted` vs `.sort()`?* → `sorted` takes any iterable and returns a new list; `.sort()` is a list method, in place, returns `None`.
- *Is Python's sort stable?* → Yes, guaranteed.
- *Sort by two fields, one descending?* → Tuple key with a negated number, or two stable sorts (secondary key first).

Self-check

1. What does `map(abs, [-1])` return? → A lazy `map` iterator, not a list.
2. How many times is `key` called when sorting 1,000 items? → 1,000 (once each).
3. Can a lambda contain an `if` statement? → No statements; a conditional *expression* (`a if c else b`) is fine.

TL;DR

Functions are values. Higher-order functions take/return functions. `lambda` for tiny keys. `map` / `filter` are lazy; comprehensions are clearer. `sorted(key=…)` — once per item, stable, tuple keys for multi-level.

7.6 Decorators and `functools`

What it is

A **decorator** is a higher-order function that takes a function and returns a replacement (usually a closure that calls the original plus extra behaviour). The `@` syntax is shorthand:

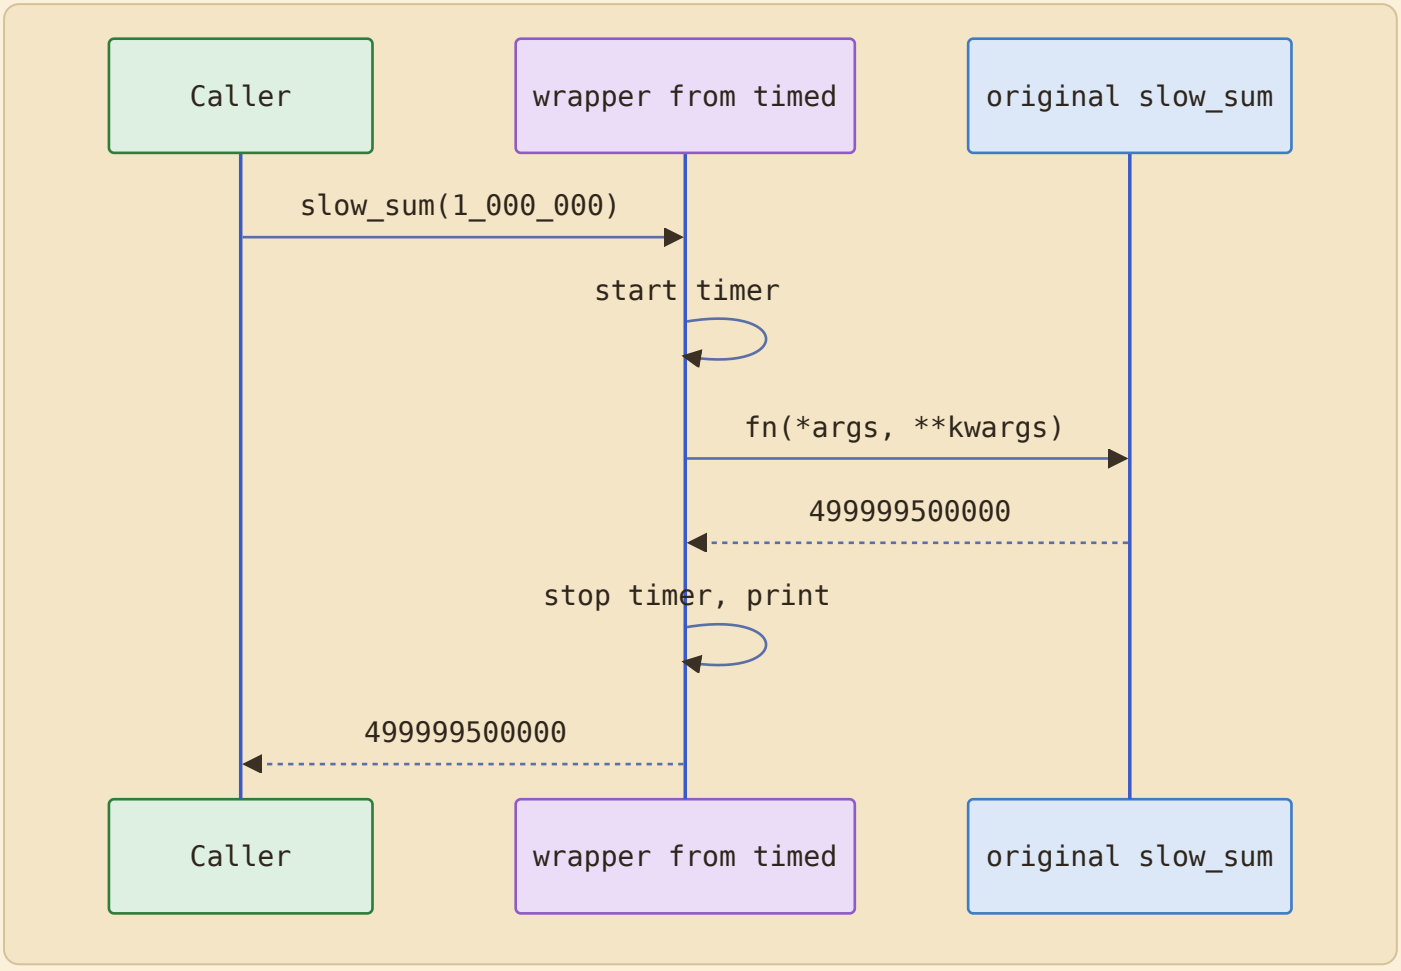
```
@timed
def slow_sum(n): ...
# is exactly
def slow_sum(n): ...
slow_sum = timed(slow_sum)
```

The standard-library module `functools` collects tools for working with functions: `wraps`, `partial`, `reduce`, `lru_cache`, `cache`, `cached_property`, `singledispatch`, `total_ordering`.

Why it matters

Decorators add cross-cutting behaviour — timing, logging, caching, auth checks, retries, route registration — without editing each function. `functools` gives you battle-tested versions so you don't write them yourself.

How it works



```python
def timed(fn):
    @functools.wraps(fn)                # copy __name__, __doc__, etc.
    def wrapper(*args, **kwargs):       # accept any signature
        t0 = time.perf_counter()
        result = fn(*args, **kwargs)    # call the original
        ms = (time.perf_counter() - t0) * 1000
        print(f"  [{fn.__name__} ran in {'under' if ms < 500 else 'over'} 500 ms]")
        return result                   # don't forget to return!
    return wrapper
```

```
  [slow_sum ran in under 500 ms]
  slow_sum(1_000_000) = 499999500000
  with wraps:    name = slow_sum | doc = Add 0..n-1 the slow way.
  without wraps: name = wrapper | doc = None
```

`functools.wraps` copies the original's `__name__`, `__doc__`, `__module__`, `__qualname__`, `__dict__` and sets `__wrapped__`. Without it, every decorated function would appear as `wrapper` in logs, `help()` and tracebacks.

`functools.partial(fn, *args, **kwargs)` "pre-fills" some arguments and returns a new callable:

```python
int_from_binary = functools.partial(int, base=2)
int_from_binary("1011")   # 11
```

`functools.reduce(fn, iterable[, initial])` folds a sequence into one value: `reduce(mul, [1,2,3,4,5])` computes ((((1·2)·3)·4)·5) = 120. Give an `initial` value or an empty input raises `TypeError`. For sums, products, max/min, `any`/`all` and joins, the dedicated built-ins (`sum`, `math.prod`, `max`, `"".join`) are clearer.

`functools.lru_cache(maxsize=128)` memoizes: it stores results keyed by the arguments and returns the stored result for repeat calls. When more than `maxsize` distinct keys are stored, the **least recently used** entry is evicted. `functools.cache` (3.9+) is `lru_cache(maxsize=None)` — unbounded, slightly faster.

```python
@functools.lru_cache(maxsize=128)
def square_slowly(n):
    time.sleep(0.05)
    return n * n
for n in (4, 4, 5, 4):
    square_slowly(n)
print(square_slowly.cache_info())
```

```
  CacheInfo(hits=2, misses=2, maxsize=128, currsize=2)
  after cache_clear: CacheInfo(hits=0, misses=0, maxsize=128, currsize=0)
  unhashable argument -> unhashable type: 'list'
```

Calls with 4, 4, 5, 4 → misses on the first 4 and on 5, hits on the two repeats. Arguments must be **hashable** (lists fail).

Trade-offs

Tool	Win	Cost / caveat
`lru_cache(maxsize=N)`	Bounded memory	Bookkeeping overhead; stale results if data changes
`cache`	Fastest, simplest	Grows forever — a memory leak for unbounded inputs
`partial`	Picklable, introspectable (`.func`, `.args`)	Only pre-fills; can't reorder arguments
`reduce`	General fold	Often less readable than a loop or built-in

When not to use it

- Don't cache functions with side effects (sending email) or whose result depends on time, randomness or mutable global state.
- `lru_cache` on an **instance method** keys on `self` and keeps every instance alive for the cache's lifetime; prefer `functools.cached_property` for per-instance values, or a module-level cached function.

Common mistakes

- Forgetting `return result` in the wrapper → the decorated function returns `None`.
- Forgetting `@functools.wraps`.
- Writing `@timed()` for a decorator that takes no arguments (or vice-versa). A decorator *with* arguments is a function that returns a decorator — three levels of nesting.

Real-world examples

- Flask/FastAPI: `@app.get("/items")` registers the function as a route handler.
- pytest: `@pytest.fixture`, `@pytest.mark.parametrize(...)`.
- Standard library: `@property`, `@staticmethod`, `@classmethod`, `@dataclass`, `@contextlib.contextmanager`.
- The `re` module keeps an internal cache of compiled patterns so repeated `re.match(pattern, s)` doesn't recompile.

In the interview

30-second answer: "A decorator is a callable that takes a function and returns a replacement; `@d` above `def f` means `f = d(f)`. Inside you define a `wrapper(*args, **kwargs)` closure and use `functools.wraps` to keep the original's name and docstring. `functools` also gives `partial` to pre-fill arguments, `reduce` to fold, and `lru_cache` / `cache` to memoize pure functions with hashable arguments."

- *What does `wraps` do?* → Copies metadata (`__name__`, `__doc__`, …) and sets `__wrapped__`.
- *How do you write a decorator that takes arguments, like `@retry(3)`?* → An outer function taking `3` returns the actual decorator, which returns the wrapper.
- *When is `lru_cache` wrong?* → Impure functions, unhashable args, unbounded key spaces with `maxsize=None`, instance methods.

Self-check

1. Rewrite `@d` / `def f` without the `@`. → `f = d(f)`.
2. What's the default `maxsize` of `lru_cache`? → 128.
3. `reduce(operator.add, [])` → ? → `TypeError` (no initial value).

4. Why `*args, **kwargs` in a wrapper? → So it can wrap any signature and forward everything.

TL;DR

Decorator = function in, function out; always `@wraps` and `return result`. `partial` pre-fills args, `reduce` folds, `lru_cache`/`cache` memoize pure functions with hashable arguments.

7.7 Recursion

What it is

A **recursive function** calls itself to solve a smaller version of the same problem. It's the programming form of *divide and conquer*: split the problem into smaller copies, solve each, combine the answers. Every recursive function needs:

1. A **base case** — an input small enough to answer directly, with no further call.
2. A **recursive case** — reduce the input and call yourself, making guaranteed progress toward the base case.

Analogy: Russian nesting dolls — open one, find a smaller one, repeat until the tiny solid doll (base case).

Why it matters

Recursion is the natural fit for **self-similar data and problems**: nested lists/JSON, folder trees, parse trees, tree and graph traversal, merge sort and quicksort, backtracking (permutations, sudoku). Interviews use it constantly.

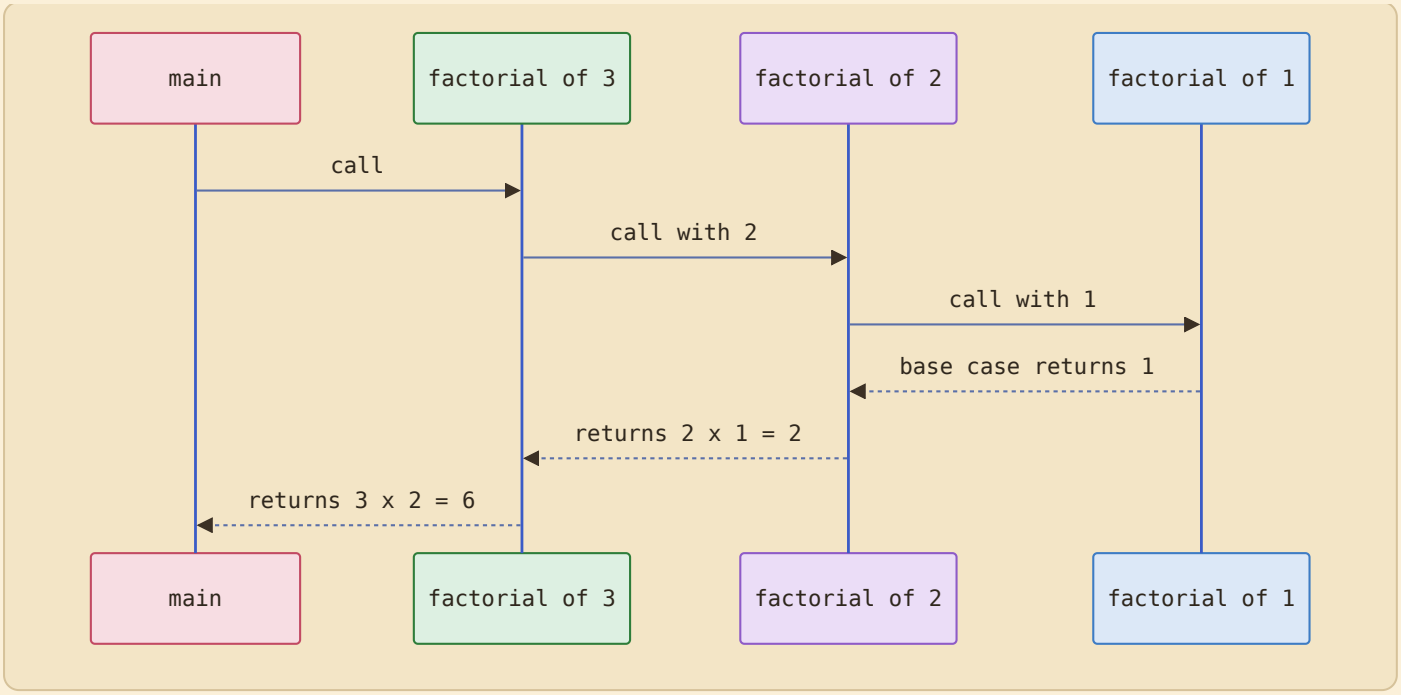
How it works — the call stack

Each call gets its own frame on the **call stack** (a last-in-first-out pile of active calls). Calls pile up until the base case, then results flow back as frames are popped. Factorial: $n! = n \times (n-1)!$, with $0! = 1! = 1$.

```
def factorial(n):
    if n <= 1:                # base case (also handles 0! = 1)
        return 1
    return n * factorial(n - 1)
```

Trace of `factorial(3)` from the script, and the arithmetic for `factorial(5)`: $5 \times 4 \times 3 \times 2 \times 1 = 120$.

```
factorial(5) = 120 | factorial(0) = 1
factorial(3)
  factorial(2)
    factorial(1)
    -> 1
  -> 2
-> 6
```



The recursion limit

CPython has **no tail-call optimisation**, so every level costs a frame. To protect the interpreter from running out of stack, it caps the depth: `sys.getrecursionlimit()` is **1000** by default. Exceed it and you get `RecursionError` (a subclass of `RuntimeError`).

```
sys.getrecursionlimit() = 1000
depth_of(500) = 500
RecursionError: maximum recursion depth exceeded
```

`sys.setrecursionlimit(n)` can raise it, but deep recursion still costs memory and can crash the process if pushed too far — for depths in the thousands or more, **rewrite it as a loop** (with an explicit stack list for tree walks).

Memoization — fixing exponential recursion

Fibonacci: F(0)=0, F(1)=1, F(n)=F(n−1)+F(n−2), so F(7) = 13. The naive version recomputes the same values over and over: F(5) calls F(4) and F(3), F(4) calls F(3) *again*, and so on. The number of calls grows roughly like 1.6^n. **Memoization** — remembering results of earlier calls — makes each F(k) computed once, so the work becomes linear.

```
@functools.cache
def fib_memo(n):
    if n < 2:
        return n
    return fib_memo(n - 1) + fib_memo(n - 2)

def fib_iter(n):              # bottom-up: no recursion at all
    a, b = 0, 1
    for _ in range(n):
        a, b = b, a + b
    return a
```

```
fib(7) = 13
naive calls for n=10, 20, 25: 177 21891 242785
memoized: CacheInfo(hits=23, misses=26, maxsize=None, currsize=26)
fib_iter(90) = 2880067194370816120
```

For n = 25 the naive version makes 242,785 calls; the memoized one computes 26 distinct values (misses) and answers 23 repeats from the cache.

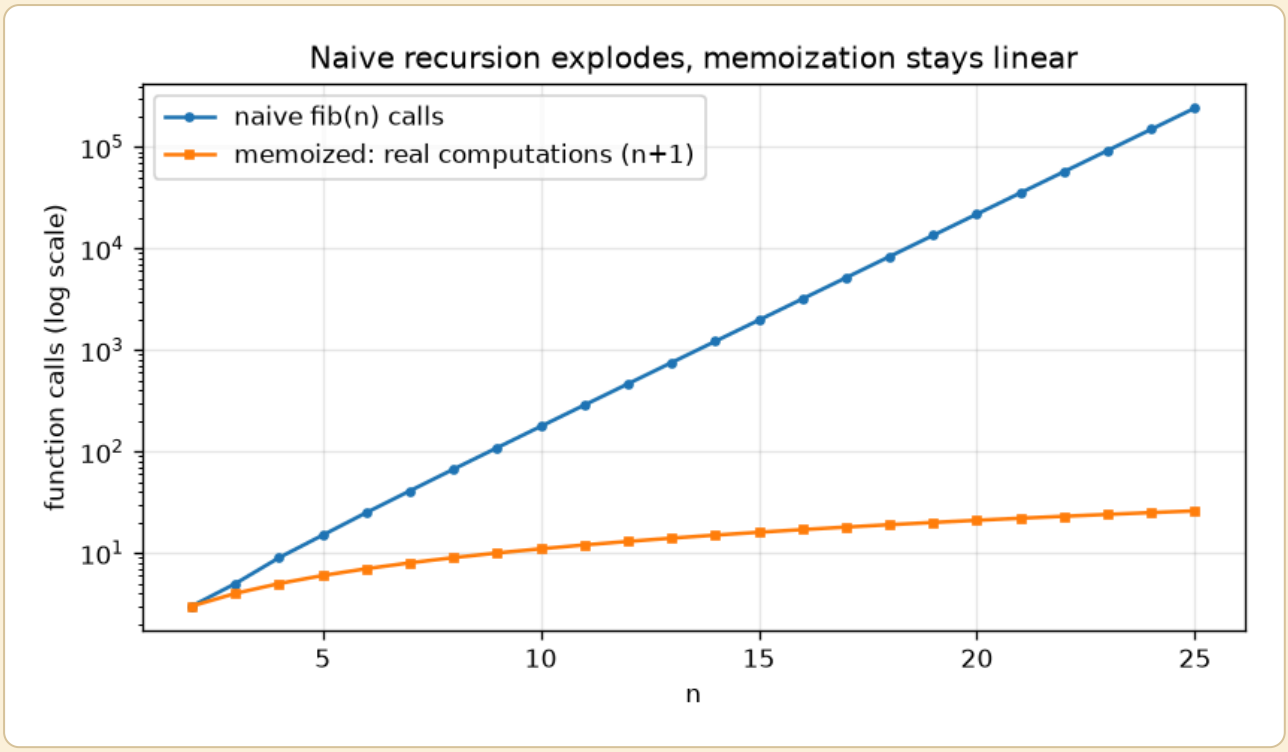


Figure — Function calls needed for fib(n): naive recursion grows exponentially (a straight line on a log axis), memoization grows linearly.

Recursion over nested data

```python
def flatten(nested):
    out = []
    for item in nested:
        if isinstance(item, list):
            out.extend(flatten(item))   # recursive case: a smaller list
        else:
            out.append(item)            # base case: a plain value
    return out

flatten([1, [2, [3, [4]], 5], [], 6])   # [1, 2, 3, 4, 5, 6]
```

The same shape walks directories (`os.walk` / `Path.iterdir`), JSON documents and trees.

Trade-offs

	Recursion	Iteration
Fits	Trees, nested data, divide and conquer, backtracking	Linear sequences, counting, very deep processes
Clarity	Mirrors the math/definition	Can need an explicit stack
Memory	One frame per level	Constant (or your own stack)
Depth limit	~1000 by default	None
Speed in CPython	Function-call overhead per level	Usually faster

When not to use it

Linear problems (sum a list, count down, walk a linked list of 100k nodes) — use a loop. Anything whose depth depends on untrusted input size.

Common mistakes

- **Missing or unreachable base case.** A factorial with base case `n == 1` recurses forever on `factorial(0)` (0 → −1 → −2 …) until `RecursionError`; `n <= 1` fixes it.
- Not shrinking the input (`f(n)` calling `f(n)`).
- Forgetting to `return` the recursive result (`factorial(n - 1)` without `return` gives `None`).
- Using naive recursion for overlapping sub-problems (Fibonacci, grid paths) — memoize.
- Mutable default accumulators in recursive helpers (`def walk(node, seen=set())`).

Real-world examples

- `os.walk`, `shutil.copytree`, `json` encoding of nested structures all process nested data (CPython's `json` encoder is recursive and can raise `RecursionError` for extremely deep input).
- `copy.deepcopy` recursively copies nested objects.
- Recursive-descent parsers (Python's own grammar tooling, many config/expression parsers).

In the interview

> **30-second answer:** "Recursion is a function calling itself on a smaller input until a base case. Each call uses a stack frame; CPython doesn't optimise tail calls and limits depth to about 1000, raising `RecursionError`. Overlapping sub-problems like Fibonacci are exponential unless memoized with `functools.cache`, which makes them linear. For very deep or linear problems, convert to a loop."

- *Complexity of naive vs memoized Fibonacci?* → $O(\varphi^n) \approx O(1.6^n)$ time vs $O(n)$ time and $O(n)$ space.
- *Does Python do tail-call optimisation?* → No, by design (keeps tracebacks complete).
- *How do you convert recursive DFS to iterative?* → Keep an explicit stack list: pop a node, process, push its children.

Self-check

1. What two parts does every recursive function need? → Base case and recursive case that progresses toward it.
2. What happens past the recursion limit? → `RecursionError`.
3. How many distinct values does memoized `fib(25)` compute? → 26 (F(0)…F(25)).
4. Why might `factorial(0)` crash with base case `n == 1`? → It never reaches 1.

TL;DR

Base case + shrinking recursive case. Stack frames pile up; limit ≈ 1000; no TCO. Memoize overlapping sub-problems (`@cache`), loop when deep.

7.8 Docstrings and annotations

What it is

- A **docstring** is a string literal placed as the **first statement** of a function, class or module. Python stores it in `__doc__`, and `help()`, IDEs and documentation generators (Sphinx) show it. Conventions are in PEP 257: a one-line summary, a blank line, then details.
- **Annotations** (type hints) attach expected types to parameters and the return value: `def f(x: int) -> str`. They are **not enforced** at runtime — Python stores them and moves on. Tools such as **mypy**, **pyright** and IDEs read them to catch bugs before you run the code; libraries such as **FastAPI** and **Pydantic** read them at runtime to validate data.

Why it matters

Docstrings explain *what and why*; annotations make the contract checkable. Together they make functions usable without reading their bodies — and let a type checker catch `average("abc")`-style mistakes in CI.

How it works

```
from collections.abc import Callable, Iterable

def average(values: Iterable[float], *, default: float | None = None) -> float | None:
    """Return the arithmetic mean of *values*.

    Args:
        values: numbers to average.
        default: returned when *values* is empty.

    Returns:
        The mean, or *default* for empty input.

    >>> average([1, 2, 3])
    2.0
    """
    items = list(values)
    if not items:
        return default
    return sum(items) / len(items)

def apply(fn: Callable[[int], int], x: int) -> int:
    return fn(x)

type Pair = tuple[int, int]          # 3.12+ type alias statement
def swap(p: Pair) -> Pair:
    return p[1], p[0]
```

```
first docstring line: Return the arithmetic mean of *values*.
signature: (values: collections.abc.Iterable[float], *, default: float | None = None) -> float | None
return annotation: float | None
average([1, 2, 3]) = 2.0
average([], default=0.0) = 0.0
swap(('x', 'y')) = ('y', 'x') <- str, not int, and no error
apply(lambda n: n + 1, 41) = 42
swap((1, 2)) = (2, 1)
doctest: 1 test(s), 0 failure(s)
```

Key syntax, by version:

Syntax	Meaning	Since
`x: int`, `-> str`	Parameter / return annotation	3.0 (PEP 3107); typing semantics 3.5 (PEP 484)
`list[int]`, `dict[str, float]`	Built-in generics	3.9
`int | None`	Union (instead of `Optional[int]`)	3.10
`type Alias = ...`, `def f[T](x: T) -> T`	Type alias / generic syntax	3.12 (PEP 695)
Lazy evaluation of annotations	Forward references work without quotes	3.14 (PEP 649)

`inspect.signature(fn)` reads parameters, defaults and annotations at runtime. The `>>>` lines in the docstring are a **doctest**: `doctest.testmod()` runs them and compares the output.

Trade-offs

Annotations add some noise and need a type checker to pay off, but on shared code they catch whole classes of bugs and power autocompletion. Docstrings drift if not maintained — keep them short and about behaviour, not implementation.

When not to use it

Tiny throwaway scripts don't need full annotations. Don't restate the obvious in a docstring (`"""Returns x."""`).

Common mistakes

- Believing annotations validate input (they don't — `swap(("x", "y"))` ran fine).
- Putting a comment `# does X` instead of a docstring → invisible to `help()`.
- Docstring not the *first* statement → `__doc__` is `None`.

Real-world examples

- FastAPI builds request validation and OpenAPI docs from function annotations.
- NumPy/Google docstring styles are rendered by Sphinx on docs sites such as numpy.org.
- Large companies with big Python codebases (Dropbox is a well-known example) run mypy in CI to type-check millions of lines.

In the interview

30-second answer: "A docstring is the first string in a function, stored in `__doc__` and shown by `help()`. Annotations declare expected types but aren't enforced at runtime; static checkers

like mypy use them, and frameworks like FastAPI/Pydantic read them to validate."

- *Are type hints checked when the code runs?* → No, only by tools or libraries that choose to read them.
- *How do you read a function's annotations at runtime?* → `inspect.signature(f)` or `inspect.get_annotations(f)` / `f.__annotations__`.

Self-check

1. Where must a docstring go? → First statement in the body.
2. What does `float | None` mean? → Either a float or `None`.
3. Does `def f(x: int)` reject `f("a")`? → No.

TL;DR

Docstring = first-line string, shown by `help()`. Annotations = documentation checkable by mypy/pyright, not enforced at runtime.

Chapter TL;DR

- `def` creates a function object; calling runs it in a fresh frame; no `return` ⇒ `None`; `pass` for empty bodies.
- Arguments: positional (order) and keyword (name); defaults evaluated **once** ⇒ never a mutable default, use `None`; `*args` tuple, `**kwargs` dict; `/` positional-only, `*` keyword-only; `*`/`**` unpack at call sites; `return a, b` returns a tuple.
- Scope: **LEGB**. Assignment makes a name local for the whole function (`UnboundLocalError`). `global` → module, `nonlocal` → enclosing function. `if`/`for` don't create scopes.
- Closures remember enclosing variables (by reference ⇒ late binding). Decorators are closures: `@d` means `f = d(f)`; use `functools.wraps`.
- Functions are first-class: pass them, return them, `key=` them. `map`/`filter` are lazy; `sorted` is stable.
- `functools`: `partial`, `reduce`, `lru_cache`/`cache` (pure functions, hashable args), `wraps`.
- Recursion: base case + progress; limit 1000, no tail-call optimisation; memoize overlapping sub-problems; loop when deep.
- Docstrings for `help()`, annotations for tools — not runtime checks.