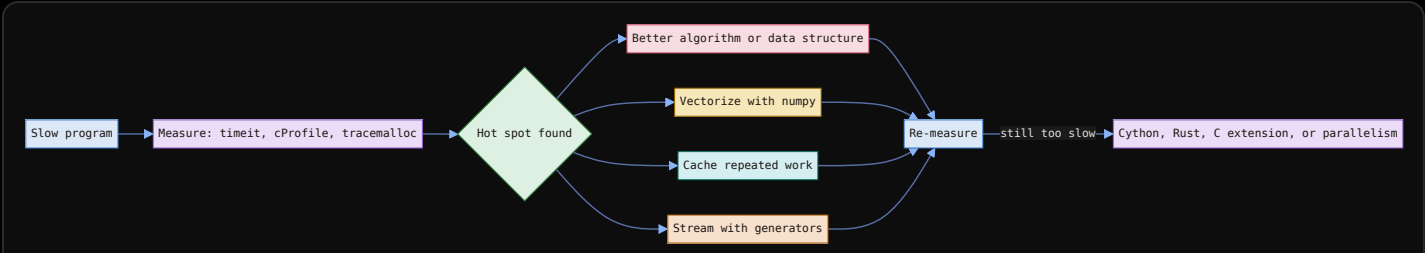
Chapter 17 — Performance, Memory and CPython Internals

What you will learn. How CPython represents every value as an object (`PyObject` : type + reference count + value); how memory is freed (reference counting plus the cyclic garbage collector); weak references; why `256 is 256` but maybe not `257 is 257` (small-int cache, string interning — one shared copy of equal values); mutability and copying recap; the **big-O cost of built-in operations** (list, dict, set, deque); how to **profile** (measure where time and memory go: `timeit` , `cProfile` , `tracemalloc` , line-level tools); and the optimisations that pay off — data structures, numpy vectorisation (one C-speed operation over a whole array instead of a Python loop), caching, generators, local lookups, `str.join` , `__slots__` — plus when to drop to C/Cython/Rust and what the "Faster CPython" releases (3.11+) changed.

Prerequisites. Ch 2 (names and objects), Ch 5–6 (containers), Ch 10 (classes), Ch 11 (generators), Ch 16 (the GIL).

OS fundamentals. What sits under every Python process — processes vs threads, CPU scheduling (which program runs on which core, when), virtual memory (each process sees a private address space the OS maps onto RAM) and page faults (the trap when touching memory not currently mapped), system calls (requests to the kernel, e.g. "read this file"), file descriptors (small integers identifying open files and sockets), blocking vs `epoll` / `kqueue` I/O (the mechanism under `asyncio` and `selectors`), and the page cache (file data kept in RAM for fast re-reads) — is explained once, language-neutrally, in the Java handbook's §16.13 Operating-system fundamentals for engineers. Python equivalents: `os.getpid()` / `multiprocessing` (processes), `threading` (threads under the GIL, Ch 16), `mmap` , `os.fsync` , `resource.getrlimit(resource.RLIMIT_NOFILE)` and `selectors` .

Code. `chapters/17-performance-memory-internals/code/` — `object_model.py` , `gc_weakref.py` , `complexity_bench.py` (plot), `profiling_demo.py` , `optimizations.py` (plot). Output from CPython 3.14.7, 64-bit macOS. Byte sizes and timings are implementation details and vary by version and machine.



17.1 The CPython object model

What it is

Everything is an object — numbers, strings, functions, classes, modules, `None` , types. Every object has:

Part	Meaning	How to see it
Identity	Which object it is; never changes during its life. In CPython it is the memory address.	`id(x)` , `x is y`
Type	What it can do; also never changes.	`type(x)`
Value	Its data. Changeable for *mutable* objects (list, dict, set), fixed for *immutable* ones (int, str, tuple, frozenset).	`==` compares values

In C, every object starts with a header struct `PyObject` holding `ob_refcnt` (the **reference count** — how many references point at it) and `ob_type` (a pointer to its type object). A variable is a **name** holding a reference to a heap object; when no name refers to it any more, it is freed.

How it works — demo

```
data = [1, 2, 3]
sys.getrefcount(data)             # 1
alias = data; holder = {"k": data}
sys.getrefcount(data)             # 3  (two more references)
del alias, holder                 # back to 1

x = 10; before = id(x); x += 1    # int is immutable -> x now names a NEW object
lst = [1]; before = id(lst); lst += [2]   # list is mutable -> same object, changed in place
```

Output (`object_model.py`):

```
refcount of new list: 1  (exact numbers are an implementation detail)
after alias + dict entry: 3
after del: 1
refcount of None: 3221225472  (immortal object, 3.12+)
a is b: True, a == c: True, a is c: False
type(a): list, type(type(a)): type, isinstance(len, object): True
int += changes id: True
list += keeps id: True
```

- `sys.getrefcount` traditionally reports one extra (its own argument), yet a fresh list shows 1: newer interpreters skip some refcount updates, so values vary by version. The docs say don't rely on it beyond 0 or 1 — watch it go **up and down**.
- `None`'s huge count: since 3.12 (PEP 683), `None`, `True`, small ints and some strings are **immortal** — refcount never changed, avoiding cache-line traffic (cores repeatedly invalidating each other's copies of the same memory) and helping the free-threaded build (Ch 16).
- `type(type(a))` is `type`: classes are objects whose type is `type`.

Object sizes (64-bit CPython 3.14)

```
getsizeof(int 0       ) =   28 bytes
getsizeof(int 1       ) =   28 bytes
getsizeof(int 2**30   ) =   32 bytes
getsizeof(int 2**100  ) =   40 bytes
getsizeof(float       ) =   24 bytes
getsizeof(str ''      ) =   41 bytes
getsizeof(str 'abc'   ) =   44 bytes
getsizeof(str 'é'     ) =   61 bytes
getsizeof(tuple ()    ) =   48 bytes
getsizeof(list []     ) =   56 bytes
getsizeof(list 3 items) =   88 bytes
getsizeof(dict {}     ) =   64 bytes
getsizeof(set()       ) =  216 bytes
getsizeof(None        ) =   16 bytes
```

A Python `int` costs ~28 bytes vs 8 in a numpy `int64` array; ints grow as needed (arbitrary precision); an ASCII string costs ~41 bytes + 1 per char, and one non-ASCII char switches the whole string to wider storage (PEP 393). `sys.getsizeof` is **shallow**: a list counts its pointer array (8 bytes per slot), not the items.

Common mistakes

- `is` for values (`if x is 1000`). `is` checks identity — only for singletons (`None`, `True`/`False`) and sentinels.
- Treating `id()` as unique forever — a dead object's address can be reused.

In the interview

> **30-second answer:** "In CPython every value is a heap object starting with a `PyObject` header — a reference count and a type pointer. Variables are names bound to objects. Identity is `id` / `is`,

> type is fixed, and the value may or may not be mutable. Memory is freed when the refcount hits zero, with a cycle collector for reference cycles."

- `==` *vs* `is`*?* — Value equality (`__eq__`) vs same object.
- *Why is a list of a million ints ~36 MB but a numpy array ~8 MB?* — The list stores 8-byte pointers to separate ~28-byte ints; numpy stores raw 8-byte values contiguously.

17.2 Memory management and garbage collection

What it is

1. **Reference counting** — the main mechanism. Each new reference increments, each dropped one decrements; at 0 the object is freed **immediately**.
2. **Cyclic garbage collector (`gc`)** — a backup that finds **reference cycles** (objects referencing each other) unreachable from anything else; their counts never reach 0 on their own.

Underneath, **pymalloc** allocates small objects (≤ 512 bytes) from pools and arenas (big chunks requested from the OS once, then split) so tiny allocations are fast. Freed memory usually returns to pymalloc's pools, not the OS — why process memory can stay high after you delete data.

How it works

- **Refcounting is deterministic**: files close when the last reference goes, `__del__` runs at a predictable point (in CPython — other implementations differ, so still use `with`).
- **Cyclic GC** sorts container objects (lists, dicts, instances…) into **generations**: new objects start young, survivors are promoted; young ones are checked often, old rarely ("most objects die young"). A collection starts when allocations minus deallocations pass a threshold (`gc.get_threshold()`). Only containers are tracked; ints and strings can't form cycles.
- Version note: 3.14.0–3.14.4 shipped an *incremental* collector (old-generation work in slices to shorten pauses); 3.14.5 (May 2026) reverted to the 3.13-style generational collector after memory-pressure problems in production. The ideas hold for both.

```
class Node:
    def __init__(self, name): self.name, self.other = name, None

n = Node("solo")
weakref.finalize(n, print, "  solo freed as soon as last reference went away")
del n                                   # refcount -> 0 -> freed right here

gc.collect(); gc.disable()
a, b = Node("a"), Node("b")
a.other, b.other = b, a                 # reference cycle
weakref.finalize(a, print, "  cycle freed by gc.collect()")
del a, b                                # refcounts stay at 1 -> not freed yet
found = gc.collect()
```

Output (`gc_weakref.py`):

```
gc thresholds: (2000, 10, 10), gc enabled: True
  solo freed as soon as last reference went away
  (after del solo)
  (after del a, b: cycle still alive)
  cycle freed by gc.collect()
  gc.collect() found 3 unreachable objects
```

Weak references

A **weak reference** points to an object *without* increasing its refcount, so it doesn't keep it alive. `weakref.ref(obj)()` returns the object, or `None` once it died.

```
r = weakref.ref(obj);  del obj;  r()          # -> None
cache = weakref.WeakValueDictionary()          # entries disappear with their values
```

```
weakref alive: True
weakref after del: None
cache keys while x alive: ['x']
cache keys after del x: []
```

Uses: caches that shouldn't pin big objects, child → parent pointers in trees (avoids cycles), observer lists. `weakref.finalize(obj, fn)` is the safe way to run clean-up at death (better than `__del__`). `int`, `list`, `tuple` can't be weakly referenced; ordinary class instances can.

Trade-offs

	Reference counting	Cyclic GC
When memory is freed	Immediately	Periodically
Handles cycles	No	Yes
Cost	Small, on every reference change	Pauses proportional to tracked objects
Control	—	`gc.collect()`, `gc.disable()`, `gc.freeze()`, `gc.set_threshold()`

Common mistakes

- `__del__` for important clean-up (files, locks) — use `with` or `weakref.finalize`.
- Accidental cycles in long-lived caches, closures or tracebacks (a stored exception keeps its frames alive).
- `gc.collect()` in a loop "to be safe" — costs time, doesn't fix leaks from references you still hold.

Real-world examples

Instagram disabled, then **froze**, the cyclic GC in its Django workers: after `fork`, parent and child share memory pages until one *writes* (copy-on-write), and GC bookkeeping writes into every object, forcing huge copies; that work led to `gc.freeze()` (3.7). ORMs like SQLAlchemy keep an identity map (table of already-loaded rows) with weak references so unused rows can be freed.

In the interview

> **30-second answer:** "CPython primarily uses reference counting: at zero an object is freed immediately. Cycles never reach zero, so a generational cyclic garbage collector periodically finds unreachable groups of container objects. Weak references refer without keeping objects alive — useful for caches and back-pointers."

- *Can Python leak memory?* — Yes: objects still referenced by a global cache, an ever-growing list, or uncollectable cycles.
- *Finding a leak?* — Compare `tracemalloc` snapshots over time; `gc.get_referrers`; `objgraph`.

17.3 Small-int caching and string interning

Two optimisations that make *different-looking* values share one object:

- **Small-int cache**: CPython pre-creates ints **−5 to 256**; every result in that range reuses them.
- **String interning**: one shared copy per equal string. CPython auto-interns identifiers and many short identifier-like literals; force it with `sys.intern(s)`.

```
a, b = int("256"), int("256")           # built at run time, not constants
c, d = int("257"), int("257")
s1 = "".join(["hello", "_", "world"])     # built at run time: not interned automatically
s2 = "hello_world"
```

```
256 is 256: True   257 is 257: False   (small ints -5..256 are cached)
runtime str is literal: False, after sys.intern: True
```

(`int("...")` avoids the compiler merging equal constants in one code block, which would hide the effect.)

Why it matters: interned dict keys can match by pointer comparison (speed); millions of repeated strings (a "country" column) can be deduplicated with `sys.intern` (memory); and the interview trap — `is` on ints/strings is implementation-dependent. **Never use `is` for value comparison.**

17.4 Mutability and copying (recap)

Chapters 2 and 5 cover this in detail; the memory-level view:

```
grid = [[0, 0], [0, 0]]
alias = grid                      # same object
shallow = grid.copy()             # new outer list, SAME inner lists
deep = copy.deepcopy(grid)        # new everything
grid[0][0] = 99
bad = [[0] * 2] * 2               # two references to ONE inner list
def append_to(item, bucket=[]):  # default evaluated once, shared across calls
```

```
alias [99, 0], shallow [99, 0], deep [0, 0]
[[0]*2]*2 after bad[0][0]=1 -> [[1, 0], [1, 0]]
mutable default after two calls: [1, 2]
```

Operation	New outer object?	New inner objects?
`b = a`	No	No
`a.copy()`, `list(a)`, `a[:]`, `dict(a)`	Yes	No (shallow)
`copy.deepcopy(a)`	Yes	Yes (recursive, handles cycles)

Copying is O(n) — shallow copies copy pointers, deep copies everything (much slower). Immutables are safe to share, hashable, and thread-safe to read. Use `[[0] * cols for _ in range(rows)]` for grids and `None` as the default for mutable parameters.

17.5 Big-O of built-in operations

What it is

Big-O describes how cost grows with container size *n*: O(1) constant, O(n) proportional, O(log n) very slow growth. "Average" for dict/set assumes a good hash spread; "amortised" = occasional expensive resizes averaged over many cheap calls.

The table (CPython)

Operation	list	dict	set	deque
Index `x[i]` / key lookup `d[k]`	O(1)	O(1) avg	—	O(1) at ends, O(n) in middle
`x in c`	**O(n)**	O(1) avg	O(1) avg	O(n)
Append / add at right end	O(1) amortised	O(1) avg	O(1) avg	O(1)
Pop from right end	O(1)	`popitem` O(1)	—	O(1)
Insert / pop at left end	**O(n)** (`insert(0,x)`, `pop(0)`)	—	—	**O(1)** (`appendleft`, `popleft`)
Insert / delete in middle	O(n)	O(1) avg (`del d[k]`)	O(1) avg (`remove`)	O(n)
`len`	O(1)	O(1)	O(1)	O(1)
Iterate	O(n)	O(n)	O(n)	O(n)
Copy	O(n)	O(n)	O(n)	O(n)
Sort	O(n log n) (Timsort, Python's hybrid sort)	—	—	—
Slice `x[a:b]`	O(k) (k = slice length)	—	—	—
`min` / `max` / `sum`	O(n)	O(n)	O(n)	O(n)
Union / intersection	—	`d1 | d2` O(len d1 + len d2)	`s | t` O(len s + len t); `s & t` O(min) avg	—

Others: `heapq.heappush/heappop` O(log n), `heapify` O(n); `bisect` search O(log n) (`insort` O(n) — the list shift); `a + b` on strings O(len a + len b); `str.join` O(total length). Dict/set worst case O(n) when many keys collide (same hash slot) — rare with built-in hashes (string hashing is randomised per process to resist deliberate collisions).

Measured

```
setup = f"n={n}; L=list(range(n)); S=set(L); D=deque(L)"
per_op_us("(n-1) in L", setup, 200)               # worst case: last item
per_op_us("(n-1) in S", setup, 200_000)
per_op_us("L.append(L.pop(0))", setup, 2_000)
per_op_us("D.append(D.popleft())", setup, 200_000)
```

Output (`complexity_bench.py`, µs per operation):

```
         n            x in list            x in set   list.pop(0)+append  deque.popleft+append
     1,000                5.346               0.022                0.098                 0.018
    10,000               54.597               0.023                0.788                 0.017
   100,000              541.614               0.023               13.024                 0.018
 1,000,000             5346.586               0.022              132.574                 0.017
```

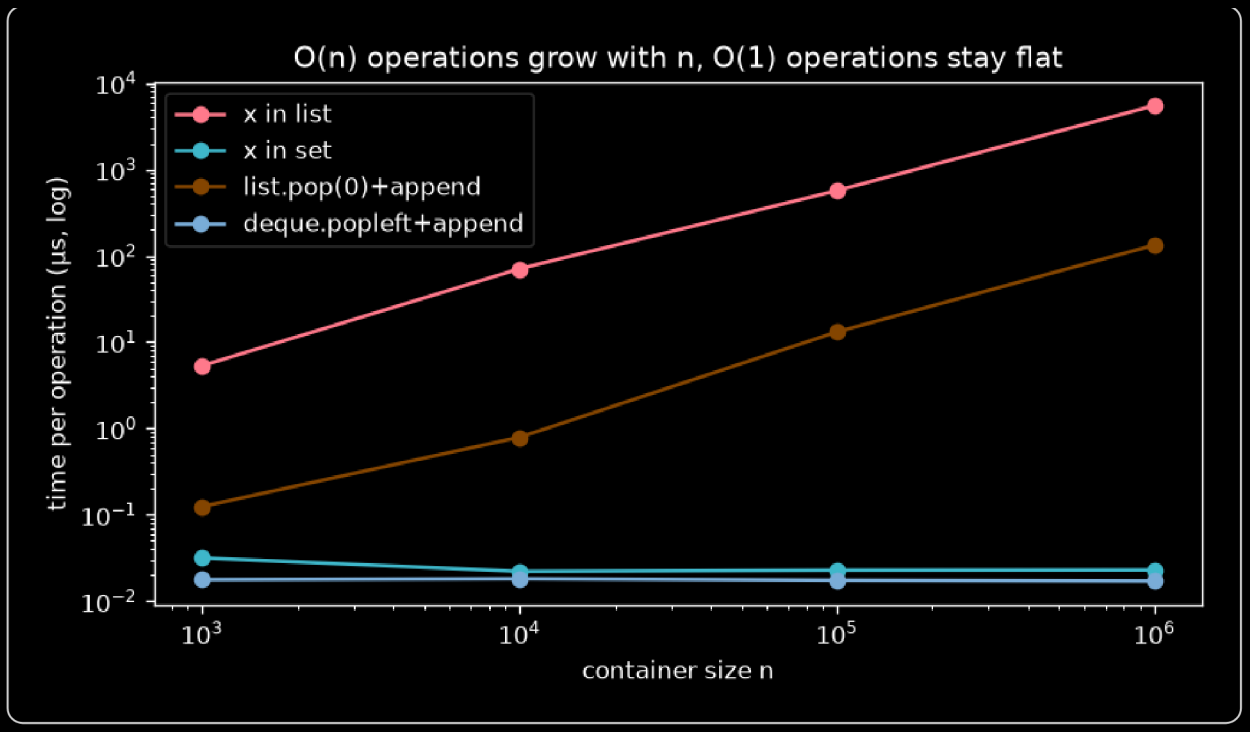


Figure — each 10× larger list makes `in` ~10× slower; set and deque don't care about n. At a million items, set lookup is ~250,000× faster.

Common mistakes

- `if x in some_list` in a loop → O(n²) (10× data, 100× slower). Build a `set` once.
- List as a queue (`pop(0)`) → `collections.deque`.
- Repeated `sorted()` in a loop for a running minimum → `heapq`.
- Building a set (O(n)) for a *single* lookup — the list is fine.

In the interview

30-second answer: "Lists are dynamic arrays (they grow by reallocating a bigger block and copying): O(1) index and append, O(n) insert/pop at the front and O(n) membership. Dicts and sets are hash tables with O(1) average lookup, insert and delete. Deque gives O(1) at both ends. The classic fix for slow code is turning repeated list membership into a set lookup."

Key term — hash table: items stored in an array of slots. A **hash function** turns each key into a number that picks the slot, so a lookup is one hash plus one probe instead of a scan — O(1) on average. When two keys pick the same slot (a *collision*) the table probes for another; with many collisions lookups degrade toward O(n). `dict` and `set` are hash tables, which is why keys must be hashable.

- *Why is `list.append` "amortised" O(1)?* — The array over-allocates; occasional resizes copy everything, but averaged per append it's O(1).
- *When is dict lookup O(n)?* — Pathological hash collisions.

17.6 Profiling — measure before you optimise

Profiling measures where time or memory goes. Guessing is usually wrong; the hot spot is often one innocent-looking line.

Tool	Answers	How
`timeit`	"How fast is this small snippet?"	`python -m timeit "sum(range(1000))"` or `timeit.repeat(...)`; in Jupyter `%timeit`
`cProfile` + `pstats`	"Which *functions* take the time?" (deterministic: every call counted)	`python -m cProfile -s cumtime script.py`
`line_profiler` (3rd-party)	"Which *lines* inside a function?"	decorate with `@profile`, run `kernprof -l -v script.py`; Jupyter `%lprun`
`py-spy`, `scalene`, `pyinstrument` (3rd-party)	Sampling profilers (peek at the program many times a second instead of recording every call): low overhead, attach to running processes; scalene also profiles memory and CPU vs native time	`py-spy top --pid 1234`
`tracemalloc`	"Which lines allocate memory?"	snapshots and `statistics("lineno")`
`time.perf_counter()`	Quick wall-clock timing of a block	subtract two readings

Python 3.15 (release candidates since August 2026; final due October 2026) adds a statistical sampling profiler to the stdlib, `profiling.sampling` (PEP 799); on 3.14 and earlier use the tools above.

How it works — a real hunt

```
def is_vip(customer_id: int, vip_ids: list[int]) -> bool:
    return customer_id in vip_ids             # O(len(vip_ids)) per call!

def build_report(n: int = 20_000) -> dict:
    records = load_records(n)
    vip_ids = list(range(0, n, 7))            # a list, not a set
    ...
prof = cProfile.Profile(); prof.enable(); build_report(); prof.disable()
pstats.Stats(prof).strip_dirs().sort_stats("cumulative").print_stats(4)
```

Output (`profiling_demo.py`):

```
timeit sum(range(1000)): 3.79 µs per call
         22861 function calls in 0.287 seconds
   ncalls  tottime  percall  cumtime  percall filename:lineno(function)
        1    0.003    0.003    0.287    0.287 profiling_demo.py:19(build_report)
    20000    0.281    0.000    0.281    0.000 profiling_demo.py:15(is_vip)
        1    0.003    0.003    0.003    0.003 profiling_demo.py:10(load_records)
     2858    0.000    0.000    0.000    0.000 {method 'get' of 'dict' objects}
list lookup: 278 ms  set lookup: 2.8 ms  speed-up x101  same result: True
records: 50,000, current 11.2 MB, peak 11.2 MB
top allocation site: line 11, 9.6 MB
```

Columns: **ncalls** = calls; **tottime** = time in the function itself; **cumtime** = including callees; **percall** = averages. `is_vip` has 98 % of the time → make `vip_ids` a `set` → ~100× faster. `tracemalloc` flags line 11 (the list of dicts) as the memory hog.

Steps that work every time:

1. Reproduce with realistic input size.
2. Profile the whole run (`cProfile`, or `py-spy` on a live server).
3. Sort by `cumtime` to find the expensive branch, then `tottime` for the leaf doing the work.
4. Fix the **algorithm/data structure** first; micro-optimise last.
5. Re-measure with `timeit`; check the result is unchanged.

Common mistakes

- A single timing run (noise) — `timeit.repeat`, take the minimum.
- Profiling tiny inputs where start-up dominates.
- Forgetting `cProfile`'s per-call overhead makes many tiny functions look worse; sampling profilers avoid this.

17.7 Optimisation techniques

Measured on 3.14

```
def concat_plus():              s = ""; for i in range(N): s += str(i)
def concat_join():              return "".join(str(i) for i in range(N))
def global_lookup(): ... total += math.sqrt(i)          # attribute lookup each time
def local_lookup():  sqrt = math.sqrt; ... total += sqrt(i)
def numpy_sum_squares():        a = np.arange(N, dtype=np.int64); return int((a * a).sum())
@lru_cache(maxsize=None)
def fib_cached(n): return n if n < 2 else fib_cached(n - 1) + fib_cached(n - 2)
```

Output (`optimizations.py`, N = 100,000):

```
string += vs join             6.33 ms ->    5.11 ms   x  1.2
for+append vs listcomp        2.45 ms ->    2.15 ms   x  1.1
math.sqrt vs local sqrt       3.27 ms ->    2.83 ms   x  1.2
python sum vs numpy           2.70 ms ->    0.04 ms   x 61.1
fib(25) vs lru_cache          5.85 ms ->   0.002 ms   x 2419   CacheInfo(hits=23, misses=26, maxsize=None,
currsize=26)
sum over 1M: list peak 40.4 MB vs generator peak 0.4 KB
100k objects: regular class 12.8 MB vs __slots__ 8.8 MB
__slots__ blocks new attributes: 'PointSlots' object has no attribute 'z' and no __dict__ for setting new
attributes
```

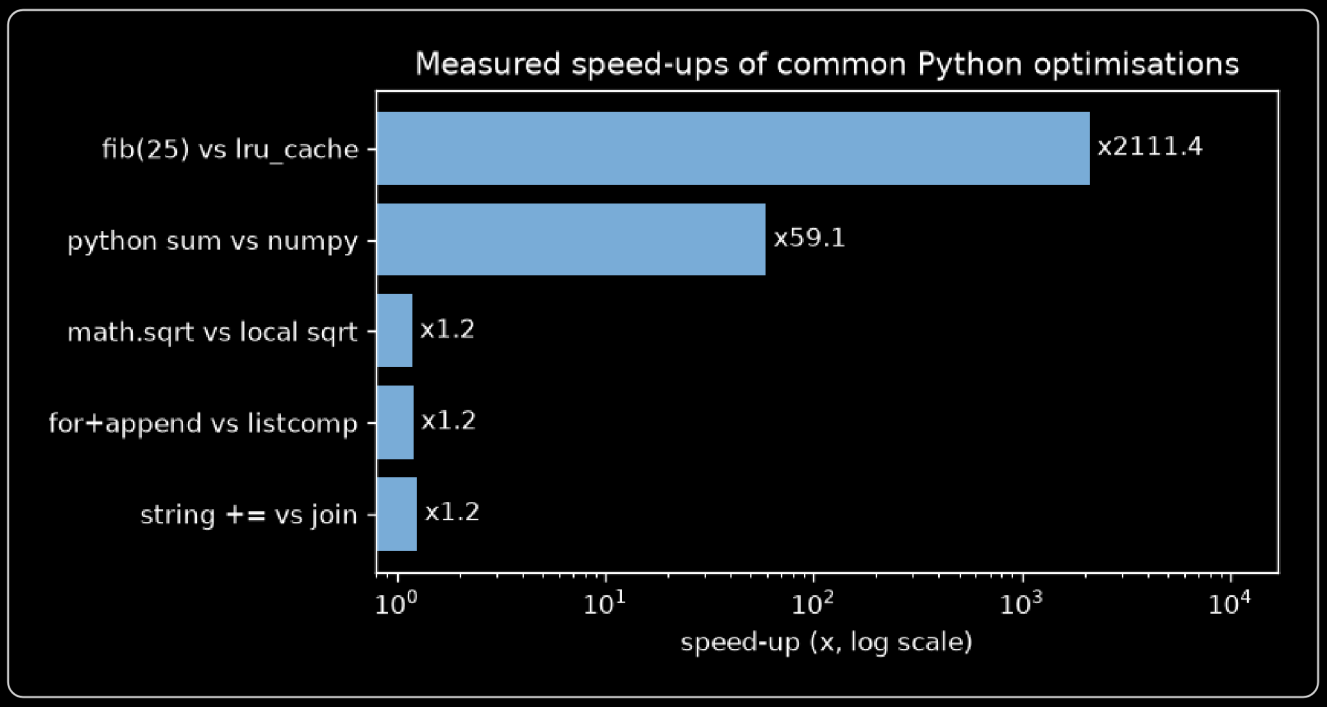


Figure — algorithmic and vectorisation changes win by orders of magnitude; micro-tweaks give 10–20 %.

The techniques, ranked by typical pay-off

#	Technique	Why it helps	Typical gain
1	**Better algorithm / data structure** (set for membership, dict for lookup, deque for queues, heapq for top-k)	Changes the big-O	10×–100,000×
2	**Caching** (`functools.lru_cache` / `cache` , memoisation, dicts)	Don't recompute the same answer	Up to exponential → linear
3	**Vectorise with numpy/pandas**	One C loop over raw numbers instead of a Python loop over objects	10×–100×
4	**Generators / streaming**	Constant memory, earlier results	Memory: GB → KB
5	**Built-ins and comprehensions** (`sum` , `max` , `sorted` , `"".join` , list comprehensions)	Loops run in C / specialised bytecode	1.1×–3×
6	**Local variable lookups** (alias `math.sqrt` before a hot loop)	Local access is an array index; globals/attributes are dict lookups	~1.1× on 3.11+
7	`__slots__` on classes with millions of instances	No per-instance `__dict__`	Less memory per object (31 % less total here)
8	**Parallelism** (Ch 16)	Use more cores	≤ number of cores

Details worth knowing:

- **`str.join`**: repeated `s += piece` can be O(n²), since strings are immutable and each `+=` may copy. CPython optimises in place when the string has one reference — hence the small gap above — but that's an implementation detail (not on PyPy; breaks once another reference exists). `"".join(parts)` is always linear. For bytes: `bytearray` or `io.BytesIO`.
- **`lru_cache`**: hashable arguments; `maxsize=None` (or `@functools.cache`) is unbounded — leak risk with many distinct inputs. Caches per function, not per instance; on methods it keeps `self` alive (use `functools.cached_property` for per-instance values).
- **Local lookups**: the 3.11+ specialising interpreter already caches global and attribute lookups, so the gain is small — only in measured hot loops.
- **Avoid needless work in loops**: hoist invariant computations, compute `len()` once, avoid exceptions for normal control flow on very hot paths (though `try` itself is zero-cost in 3.11+ when nothing is raised).

`__slots__`

Instances normally keep attributes in a per-instance dict; `__slots__ = ("x", "y")` replaces it with fixed slots:

```
class PointSlots:
    __slots__ = ("x", "y")              # fixed attributes, no per-instance __dict__
    def __init__(self, x, y):
        self.x, self.y = x, y
```

100,000 points: 12.8 MB vs 8.8 MB (the total includes the list and the ints, so the per-object saving is bigger than the ratio). Slightly faster attribute access too. Costs: no new attributes, no `__dict__` (some tools expect it), no weak references unless `"__weakref__"` is listed, and every class in the inheritance chain needs its own `__slots__`. `@dataclass(slots=True)` (3.10+) generates them. Since 3.11–3.13 ordinary instances store attributes more compactly ("inline values"), so the gap is smaller than old blog posts claim — measure.

When to drop to C, Cython or Rust

When a profiled hot loop can't be expressed in numpy/pandas and is still too slow:

Option	What it is	Good for
numba	JIT-compiles (compiles while the program runs) numeric Python functions (`@njit`) to machine code	Loops over numpy arrays
Cython	Python-like language with C types, compiled to a C extension	Gradual speed-up of existing Python; wrapping C libraries
Rust + PyO3/maturin	Write the extension in Rust	Safe, fast new libraries (pydantic-core, polars, ruff, tokenizers are written this way)
C/C++ + pybind11 / C API	Classic native extension	Existing C/C++ code
PyPy	Alternative Python with a tracing JIT (compiles the loops it sees running hot)	Pure-Python long-running programs (check C-extension support)
mypyc	Compiles type-annotated Python to C extensions	Typed codebases (mypy and the black formatter use it)

Exhaust algorithmic fixes and vectorisation first; native code adds a build step and a second language.

Python version performance (Faster CPython)

Version	What changed (from the official "What's New" notes)
3.11 (2022)	Specialising adaptive interpreter (PEP 659): bytecode rewrites itself for the types it sees. Cheaper frames, inlined Python-to-Python calls, "zero-cost" exceptions. Official figure: **10–60 % faster than 3.10, ~1.25× on average** on pyperformance (the standard benchmark suite).
3.12 (2023)	Inlined comprehensions (PEP 709, up to 2× faster for comprehension-heavy code), immortal objects (PEP 683), per-interpreter GIL (PEP 684).
3.13 (2024)	Experimental JIT compiler (PEP 744, off by default), experimental free-threaded build (PEP 703), improved interactive REPL.
3.14 (2025)	Free-threaded build officially supported (PEP 779); optional tail-calling interpreter (new internal dispatch; ~3–5 % faster when built with recent Clang); experimental JIT in the Windows and macOS installers, still off by default.

Upgrading from 3.10 to 3.11+ is often the cheapest speed-up you'll ever get.

When not to optimise

Code that runs once, I/O-bound code (1 ms CPU saved next to a 200 ms network call), anything a profiler doesn't show hot. Readability usually wins.

Common mistakes

- Micro-optimising before fixing an $O(n^2)$ algorithm.
- "Vectorising" by calling numpy element-by-element in a Python loop (slower than plain Python!).
- Unbounded caches growing forever in long-running services.
- Premature `__slots__` on classes with 50 instances.

Real-world examples

pandas and scikit-learn are fast because inner loops are C/Cython; polars and pydantic v2 moved their cores to Rust; Django uses `lru_cache` / `cached_property` for settings and URL resolution; mypyc was built by the mypy team (originally at Dropbox) to compile mypy itself, making it several times faster.

In the interview

> **30-second answer:** "First I profile — cProfile or py-spy — to find the hot spot. The biggest wins are algorithmic: the right data structure, caching repeated work, vectorising numeric loops with numpy. Then smaller things: built-ins, comprehensions, `join` instead of `+=`, generators for memory, `__slots__` for many small objects. If that's not enough, I move the hot loop to numba, Cython or Rust, or parallelise."

- *Why is numpy faster than a Python loop?* — Contiguous typed memory, one C loop, no per-element object creation or dispatch, often SIMD (single instruction, multiple data — one CPU instruction processing several numbers).
- *Downside of `lru_cache`?* — Memory growth and stale results; arguments must be hashable.
- *How much faster was 3.11 than 3.10?* — ~25 % on average (10–60 %).

17.8 Chapter self-check

1. Three properties of every object, and what starts every `PyObject`? → Identity, type, value; refcount and type pointer.
2. Does `x += 1` on an int change the object? → No, it rebinds `x` to a new int.
3. Is `sys.getsizeof([big objects])` the total memory? → No, it's shallow.
4. When is an object freed in CPython? → Refcount reaches zero, or the cyclic GC finds it unreachable in a cycle.
5. Why do GC generations exist? → Most objects die young; scan young often, old rarely.
6. Does a weak reference keep an object alive? What does `WeakValueDictionary` do when a value dies? → No; removes the entry.
7. Which ints are cached? → −5 to 256.
8. Cost of `x in list` vs `x in set`? → O(n) vs O(1) average.
9. Best structure for a FIFO queue? → `collections.deque`.
10. `tottime` vs `cumtime`? → Own time vs including callees.
11. Why does `sum(i*i for i in range(10**6))` use almost no memory? → The generator streams values.
12. Why `''.join` over `+=`? → Guaranteed linear; `+=` can be quadratic.
13. Downside of `__slots__`? → No new attributes, no `__dict__`, no weakrefs unless listed.
14. Two ways to speed up a numeric loop numpy can't express? → numba, Cython (or Rust/PyO3).
15. First thing to try before rewriting in Rust? → Profile and fix the algorithm / vectorise.

Chapter TL;DR

- **Objects**: identity (`id`, `is`), type, value; `PyObject` = refcount + type pointer; names are references; immortal objects since 3.12.
- **Memory**: refcounting frees immediately; generational cyclic GC handles cycles; `weakref` refers without owning; `gc.freeze()` helps pre-fork servers.
- **Caching quirks**: small ints −5…256 and interned strings are shared — always compare with `==`.
- **Copying**: assignment shares, `.copy()` is shallow, `deepcopy` is deep; immutables are safe to share.
- **Big-O**: list O(1) end / O(n) front and `in`; dict/set O(1) average; deque O(1) both ends; heapq O(log n).
- **Profile**: `timeit`, `cProfile` / `pstats`, `line_profiler`, `py-spy` / `scalene`, `tracemalloc`.
- **Optimise**: data structure → caching → numpy → generators → built-ins/join/locals → `__slots__` → numba/Cython/Rust → parallelism. Newer CPython (3.11+) is substantially faster for free.
- **OS layer** (processes, threads, paging, syscalls, fds, epoll, page cache): see Java handbook §16.13.