Chapter 10 — Data Modeling and ER Diagrams

In real work nobody hands you tables: someone describes a business ("we rent bikes to people") and you decide what to store, how the pieces connect and which rules the database enforces. That is **data modeling**, and it matters more than any query, because a fact left out of the model is a question no SQL can ever answer.

The path: a plain-language brief → a drawing business people can check (an **ER diagram**, in **Chen notation** and **crow's-foot notation** — the one `mermaid erDiagram` draws; Mermaid turns short text into diagrams) → precise tables with keys and constraints (the **relational model**, §10.11) → `CREATE TABLE` on PostgreSQL. Covered: the three model levels, a seven-step process and best practices, entities, the five kinds of attribute, relationships (degree, cardinality, participation, attributes), weak entities, a worked employee example, relations, domains, `NULL`, keys, integrity constraints, and ER-to-table mapping. Chapter 11 then tests tables for repeated data; Chapter 17 applies everything to nine case studies.

Code in `code/`: `00_schema.sql` builds a small bike-rental database; `01`–`04` each run on a fresh database (Chapter 1 server). Outputs are real psql output.

10.1 What data modeling is

What it is

Data modeling is deciding three things before you write any SQL:

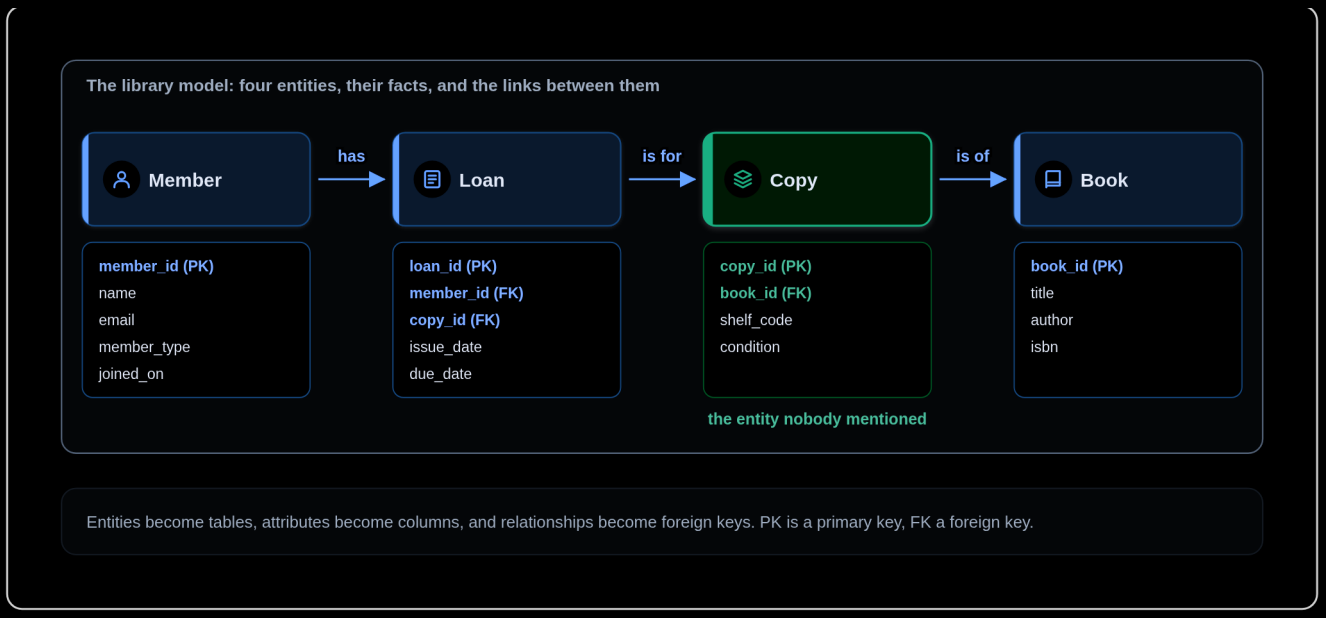
Part	Plain meaning	Usually becomes
Entity	a *thing* you keep data about (Customer, Bike, Rental) — almost always a noun	a table
Attribute	one *fact* about an entity (a customer's email, a bike's frame number)	a column
Relationship	a *link* between entities (a customer *rents* a bike)	a foreign key (a column holding another table's key), or a table of its own

A **data model** pictures those three things and nothing else — no indexes (lookup structures that speed up searches), `VARCHAR` lengths or product choice: the architect's floor plan before anyone picks bricks.

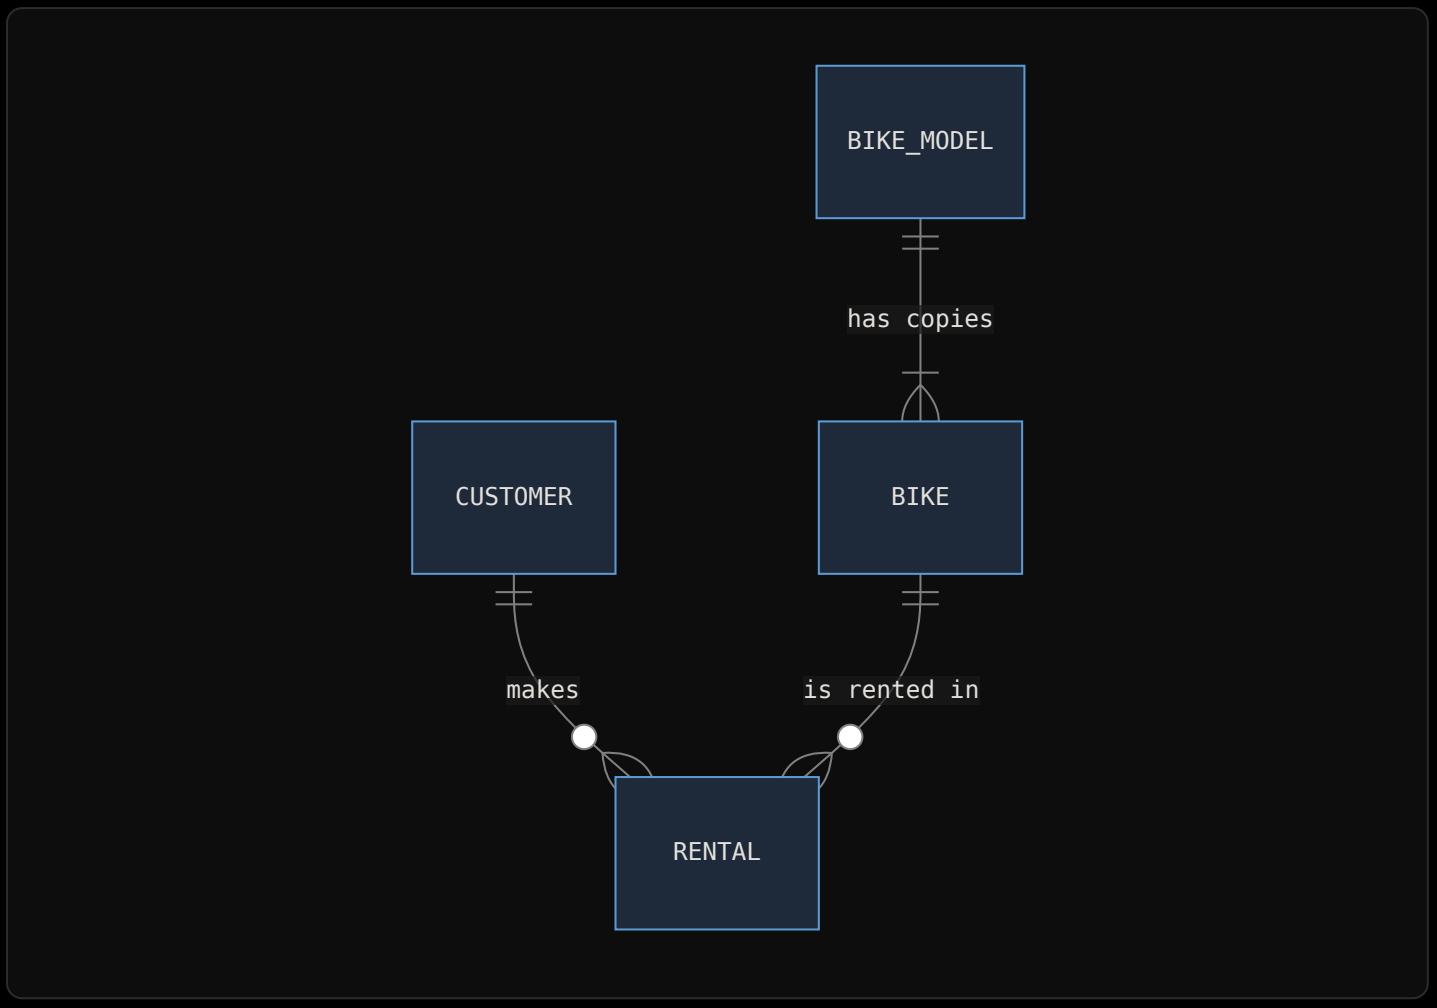
The problem it solves — the question you cannot answer

"We rent bikes to customers" sounds like `customer` and `bike`. Then the owner asks: *"We own three 'Urban 3' bikes. One came back with a bent wheel — which one, and who had it?"* If `bike` describes a **model** (name, type, hourly rate), it cannot answer: the physical object with a frame number differs from the product it exemplifies. Add an entity for the physical bike and the question is a short join — the database simply records one more fact.

Figure — the same idea in the course's library example: a Book (a title) is not a Copy (the object on the shelf) (course diagram)



The bike shop's model therefore has four entities: `customer`, `bike_model` (the "title"), `bike` (the physical object) and `rental` (who had which bike, when).



Who reads a data model

Reader	What they check
Business people (shop owner, registrar, librarian)	"Is anything missing? Are the rules right?" — they spot a forgotten entity in seconds
Developers	which tables exist and what each one means
Database administrators	storage, indexes, backups (mostly at the physical level)
Analysts	where each fact lives, so reports are correct

Non-technical readers must follow it, so describe it in everyday language first, detail later — hence §10.2's three levels.

What data modeling is *not*

- **Not writing SQL** — that comes last and is nearly mechanical (§10.15).
- **Not performance tuning** — indexes and query plans (the database's chosen execution strategy) come later (Chapter 12); quickly answering the wrong question helps nobody.
- **Not application code** — very different apps can share one model.

Why it matters

- **Fewer wrong answers** — rules in the model become constraints that every program must obey.
- **Less repeated data** — each fact stored once cannot contradict itself.
- **Cheaper change** — a model that matches reality absorbs new requirements; one that doesn't fights each of them.
- **Shared vocabulary** — when the team says "rental", everyone means the same thing.

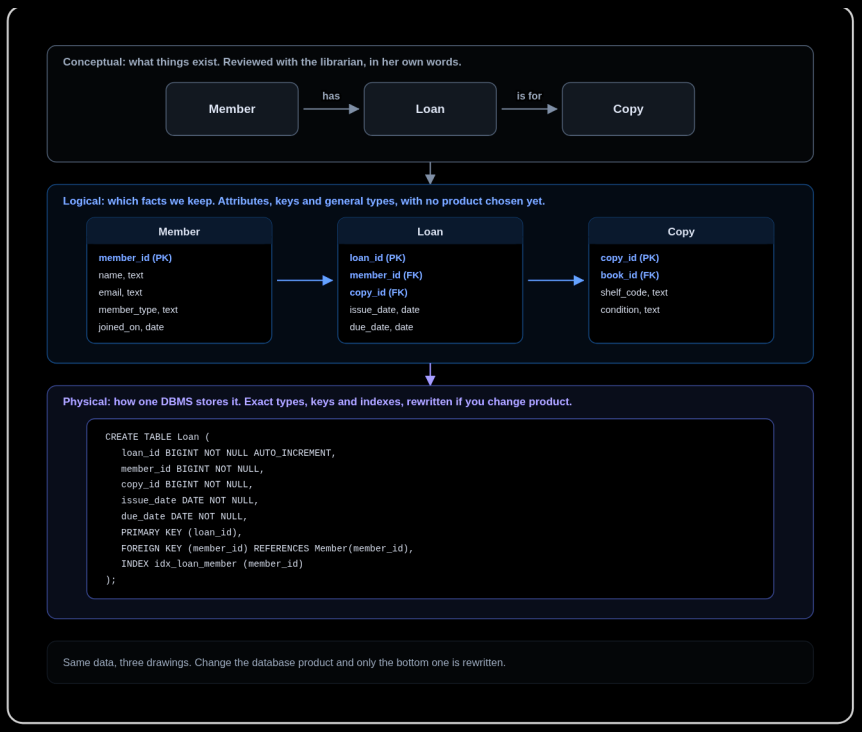
10.2 The three levels of a data model

What it is

The same model drawn three times, with rising detail, for different audiences:

	Conceptual	Logical	Physical
Shows	entities and relationships only	+ attributes, keys, general types (text/number/date)	+ exact types, indexes, partitioning (splitting a big table into pieces), storage options
Detail	low	medium	high
Main audience	business people	designers, developers	DBAs (database administrators), developers
Tied to a DBMS (database management system, e.g. PostgreSQL) product?	no	no	**yes**
Question it answers	"Is every thing present, and is each link correct?"	"Which facts are recorded, and which keys join the tables?"	"How do we store this well on *this* database?"
Bike-shop example	Customer *rents* Bike; Bike *is of* BikeModel	`rental(rental_id, customer_id → customer, bike_id → bike, out_at, due_at, back_at)`	`CREATE TABLE rental (... timestamp ..., INDEX on customer_id)`

Figure — the same data at the three levels (course diagram)



How it works

1. **Conceptual** — boxes and lines in the users' words; valid for a relational database, a document store (JSON-like documents instead of tables) or a paper ledger.
2. **Logical** — attributes, primary/foreign keys, general types; **normalization** (removing repeated data, Chapter 11) happens here. Still product-neutral.
3. **Physical** — for one DBMS: `bigint` vs `int`, `numeric(6,2)`, `timestamp`, indexes for known queries, identity columns (auto-numbered keys), partitioning.

The physical version of the rental table from `code/00_schema.sql`:

```
CREATE TABLE rental (
  rental_id   bigint GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
  customer_id bigint      NOT NULL REFERENCES customer(customer_id),
  bike_id     bigint      NOT NULL REFERENCES bike(bike_id),
  out_at      timestamp   NOT NULL,
  due_at      timestamp   NOT NULL,
  back_at     timestamp,                          -- NULL = bike still out (meaningful NULL)
  CHECK (due_at > out_at),
  CHECK (back_at IS NULL OR back_at >= out_at)
);
CREATE INDEX rental_customer_idx ON rental(customer_id);   -- "rentals of this customer"
CREATE INDEX rental_bike_idx     ON rental(bike_id);       -- "who has this bike"
```

Entity from the conceptual model, columns and keys from the logical; only exact types, the identity column and the indexes are physical — and identity syntax differs by product:

Auto-numbered key	Syntax
PostgreSQL 10+	`bigint GENERATED ALWAYS AS IDENTITY` (older: `bigserial`)
MySQL	`BIGINT AUTO_INCREMENT`
SQLite	`INTEGER PRIMARY KEY` (becomes the rowid, SQLite's built-in row number)
SQL Server	`BIGINT IDENTITY(1,1)`

Why bother with three levels

- **Each reader gets a suitable version** — an owner reviewing three boxes spots a missing entity; shown `CREATE TABLE`, nothing.

- **Mistakes caught cheaply** — a missing entity costs a conversation at the conceptual level, a migration (script changing an existing schema and data) at two million rows.
- **Technology isolated** — MySQL → PostgreSQL rewrites only the physical model: **data independence** (Chapter 1) at design time.

Working out the logical model often shows the conceptual one was wrong; going back is cheap because the drawing is small.

Common mistakes

- Showing column types to business reviewers (they stop reading).
- Calling a diagram with attributes and keys "conceptual" — once keys appear it is logical.
- Treating logical and physical as the same: *logical = what the business needs; physical = how one product stores it.*

Self-check

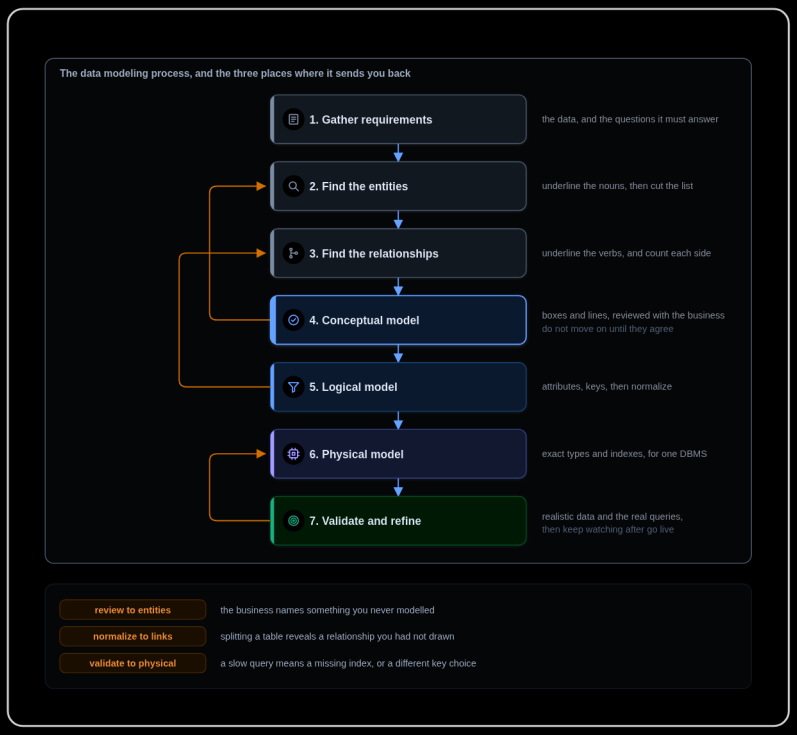
1. Which level first shows primary keys, and hosts normalization? → **Logical.**
2. Migrating SQL Server → PostgreSQL changes which levels? → **Only the physical.**

10.3 The data modeling process

What it is

A fixed order of work, each step producing something reviewable. Starting with tables is the classic mistake — it fixes the answer before you understand the question.

Figure — seven steps and the three places the work loops back (course diagram)



How it works — the bike shop, step by step

1. **Gather requirements — including the questions** the system must answer:
    - Customers are casual or members; members may keep a bike longer.
    - The shop owns several physical bikes of each model.
    - A customer takes a bike at a time and must return it by a due time.
    - Late returns cost 2.00 per started hour.
    - Staff must see which bikes are out and who has them.

An unanswerable question is a design bug, free to fix now.

2. **Find the entities: underline the nouns** (customer, member, model, bike, time, fee, staff). Keep a noun only if it has *several facts* and *many instances*:
   - customer, bike_model, bike, rental → entities.
   - time → one value → an attribute.
   - casual/member → two kinds of customer → start with a `customer_type` attribute.
   - late fee → calculable from times → a **derived** value (§10.6), not stored.
   - staff → users of the system, not necessarily data in it.
3. **Find the relationships: underline the verbs** (customer *rents* bike; bike *is of* bike_model); ask "how many on each side?" (cardinality, §10.7) and write it in words.
4. **Review the conceptual model** with the owner, read back as sentences. She will correct you ("customers can also *reserve* a bike") — a new entity for the price of one box. Don't continue until the business agrees.
5. **Logical model** — attributes, keys, general types; normalize.
6. **Physical model** — exact types, an index for every step-1 question (`rental.customer_id`, `rental.bike_id`).
7. **Validate, implement, refine** — realistic volumes (enough for slowness to show), run the step-1 questions. After go-live keep watching: right for 10 thousand rentals may need a new index at 10 million.

The loops back are part of the process

Trigger	Sends you back to
Stakeholder (anyone with a say in the system) review reveals something new (reservations)	step 2 (entities)
Normalization splits a table and exposes a hidden link	step 3 (relationships)
Validation finds a slow query	step 6 (index or key choice)

Looping back is the method working: every loop costs less than releasing the error.

Trade-offs to keep in mind throughout

- **Normalize vs denormalize** — normalize first (each fact once); add redundancy (deliberate copies) only for a measured slow query (Chapter 11).
- **Growth** — plan for rows in three years; it changes key and index choices.
- **Access control** — mark sensitive data early (emails, payments).
- **Integrity** — each step-1 rule belongs in a constraint, not an application `if`.

The model answering the step-1 questions (`code/01_bike_questions.sql`)

```sql
-- Q1: which bikes are out right now, and who has them?
SELECT b.frame_no, m.model_name, c.full_name, r.due_at
FROM rental r
JOIN bike b       ON b.bike_id = r.bike_id
JOIN bike_model m ON m.model_id = b.model_id
JOIN customer c   ON c.customer_id = r.customer_id
WHERE r.back_at IS NULL
ORDER BY b.frame_no;
-- Q2: bikes never rented (NOT EXISTS anti-join)   Q4: late fee, derived at read time
SELECT r.rental_id, c.full_name,
       ceil(extract(epoch FROM (r.back_at - r.due_at)) / 3600) * 2.00 AS late_fee
FROM rental r JOIN customer c USING (customer_id)
WHERE r.back_at > r.due_at;
```

Output:

```
--- Q1: which bikes are out right now, and who has them?
 frame_no | model_name |  full_name   |       due_at
----------+------------+--------------+---------------------
 FR-2001  | Ridge 29   | Meera Nair   | 2026-09-27 18:00:00
 FR-3001  | Volt City  | Lena Fischer | 2026-09-27 10:30:00
(2 rows)

--- Q2: which physical bikes have never been rented?
 frame_no | model_name
----------+------------
 FR-1002  | Urban 3
 FR-1003  | Urban 3
(2 rows)

--- Q3: how many bikes of each model exist, and how many are out?
 model_name | bikes | out_now
------------+-------+---------
 Ridge 29   |     1 |       1
 Urban 3    |     3 |       0
 Volt City  |     1 |       1
(3 rows)

--- Q4: late returns, with a late fee of 2.00 per started hour (fee is derived, not stored)
 rental_id | full_name  | late_fee
-----------+------------+----------
         2 | Tom Okafor |     2.00
(1 row)
```

Q2 and Q3 work only because `bike` (the object) is separate from `bike_model`.

In the interview

30-second answer: "I start from requirements *and the questions the system must answer*. Nouns give candidate entities, verbs give relationships. I draw a conceptual model and review it with the business, then add attributes and keys and normalize (logical), then choose types and indexes for the target database (physical), and validate with realistic data and the real queries. I expect to loop back."

- *How do you decide whether a noun is an entity or an attribute?* → Entity if it carries several facts and has many instances; a single value (a date, a colour) is an attribute.
- *When do you add indexes?* → At the physical level, for the queries written down in requirements — not "on every column".

10.4 Best practices in data modeling

Two models that both work on day one can differ enormously a year later because of small early decisions:

Practice	Rule	Why
Questions first	Write the list of questions before drawing.	A model is only "correct" relative to the questions it must answer.
Rules in the database	Every stated rule → a `PRIMARY KEY`, `FOREIGN KEY`, `NOT NULL`, `UNIQUE` or `CHECK`.	A constraint protects every program that ever connects, including tomorrow's ad-hoc (one-off, written for a single task) script; an `if` protects only the program it is in.
One naming convention	Decide once (`snake_case` = lower case with underscores, singular *or* plural tables) and never mix. Same thing = same name everywhere (`customer_id`, never `cust_no` elsewhere). Name FKs after their target.	Joins read naturally; nobody has to guess.
Types on purpose	Dates as `date` / `timestamp`, never text; money as `numeric`, never `float` / `double`; text sized for real data. Prefer the narrowest type that fits.	Text dates accept "31 February" and sort wrongly; floating point cannot represent most decimals exactly, so totals drift; narrower rows = more rows per page (the fixed-size block the database reads from disk) = fewer disk reads.
Decide what NULL means	`NOT NULL` whenever the fact must exist; allow `NULL` only when "missing" is meaningful (`back_at IS NULL` = bike still out).	`NULL` means unknown/not applicable, is not 0 or '' and is not equal to anything (§10.12).
Keys deliberately	Surrogate primary key (a generated id with no business meaning) + `UNIQUE` on the natural key (a real-world identifier: `email`, `frame_no`). Keep keys narrow.	Natural values change; a PK is copied into every FK and index.
Normalize, then denormalize with evidence	Default to 3NF (third normal form, the usual target of normalization); denormalize only after a measured slow query that an index could not fix — and document why.	Redundant copies can disagree (Chapter 11).
Index for real queries	Index what the requirement questions filter/join on; nothing else.	Every index slows every write and costs disk.
Design for predictable growth	Estimate rows in 3 years; don't build tables for imagined requirements.	Speculative tables become the parts nobody understands.
Document meaning	For each entity/column: what it means, allowed values, units.	Anyone can see `status` is text; only docs say which 4 values it holds.

The bike schema's `CHECK (due_at > out_at)` and `CHECK (back_at IS NULL OR back_at >= out_at)` (§10.2) apply "rules in the database" for every client.

Common mistakes

- **A list in one column** (`phones = '555-1, 555-2'`) — can't be searched, counted or constrained → separate table.
- **Numbered columns** (`phone1, phone2, phone3`) — a list across columns; breaks at the fourth phone.
- **One huge table** — 60 mostly-empty columns are several unseparated entities.
- **Over-normalizing** — seven joins for an everyday question is a real cost; 3NF is the usual stop.
- **Rules only in code** — another program will break them.

Real-world examples

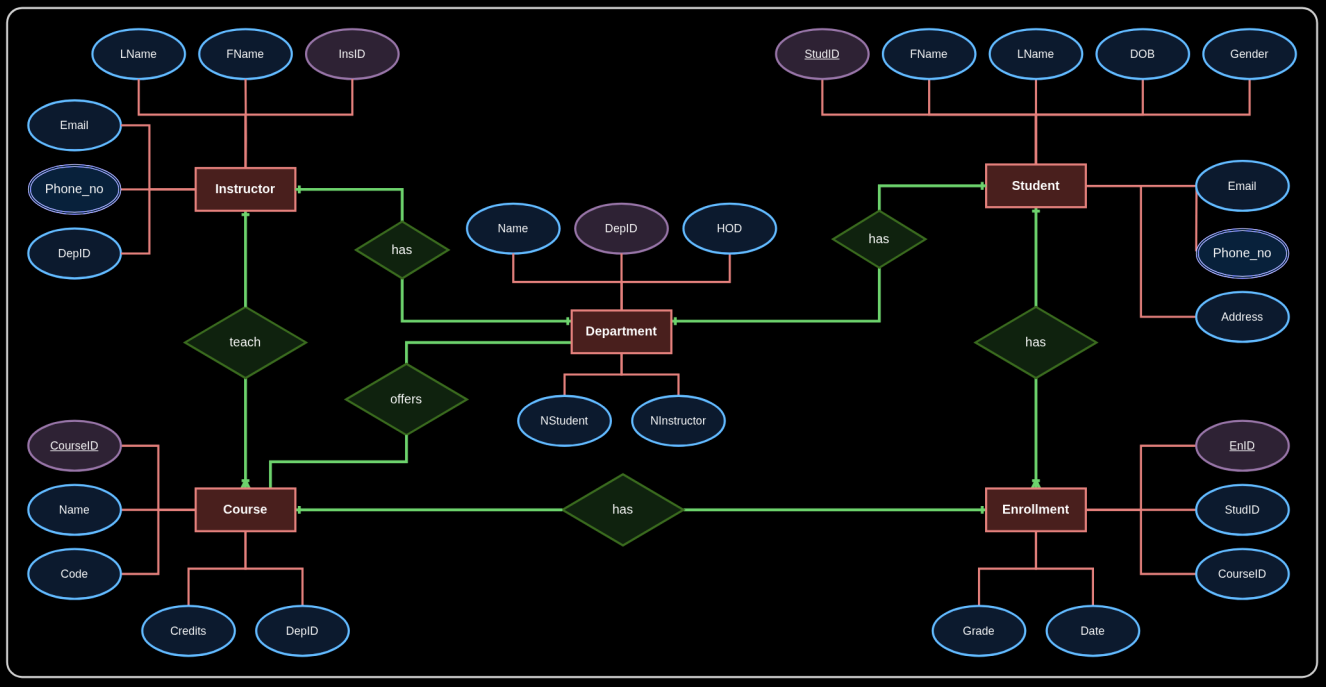
- **Stripe's API** (the payments company's programming interface) stores amounts as integers in the smallest currency unit (cents) — "never float for money".
- **ISO country codes** (two-letter standard codes such as `IN`, `GB`) are `char(2)` — "size text for the data".

10.5 The Entity-Relationship (ER) model and its notations

What it is

The **Entity-Relationship model** (ER model), proposed by Peter Chen in 1976, sketches entities, attributes and relationships before any table exists. An **ER diagram (ERD)** draws the conceptual model, then the logical one (plus attributes, keys, cardinalities); it never names a DBMS.

Figure — a complete Chen-notation ER diagram for a student management system; by §10.10 you can read every symbol (course diagram)



Two notations for the same model

	Chen notation	Crow's-foot notation (IE / Martin)
Entity	rectangle	rectangle with a title bar
Attribute	ellipse attached by a line	a row listed *inside* the rectangle
Relationship	diamond labelled with a verb	the line itself, labelled with a verb
Cardinality	`1`, `N`, `M` written near the diamond (or `(min,max)` pairs)	the *shape* of each line end (see table below)
Participation	double line = total, single = partial	inner symbol: bar = at least one, circle = may be zero
Used by	textbooks, exams, teaching	nearly every design tool (dbdiagram, MySQL Workbench, ERwin, Lucidchart), and `mermaid erDiagram`
Strength	shows the concepts explicitly	compact; fits 40 tables on a screen

Neither is "more correct": learn Chen for the concepts, then crow's foot reads itself. This book shows concepts in the course's Chen figures and every schema in mermaid crow's foot.

Reading crow's-foot line ends — *outer* symbol = maximum, *inner* = minimum:

Mermaid syntax (left end / right end)	Picture	Meaning (min, max)
`|o` / `o|`	circle + bar	zero or one `(0,1)`
`||` / `||`	bar + bar	exactly one `(1,1)`
`}o` / `o{`	circle + crow's foot	zero or more `(0,N)`
`}|` / `|{`	bar + crow's foot	one or more `(1,N)`

`CUSTOMER ||--o{ RENTAL : makes` reads: *each rental belongs to exactly one customer; a customer makes zero or more rentals*. Strictly, mermaid's solid `--` marks an **identifying** relationship (the child's key includes the parent's — a weak entity, §10.9) and dashed `..` a non-identifying one; this book draws all lines solid, so spot a weak entity by a PK containing the owner's key.

Strengths and limits of an ERD

- **Good at** showing things, links and counts so even non-technical reviewers notice gaps.
- **Not** a table diagram (one M:N diamond becomes *three* tables, §10.7), a performance document (no indexes or storage), or a process diagram (a student enrols in a course — not which screens or approvals).

10.6 Entities and the five kinds of attribute

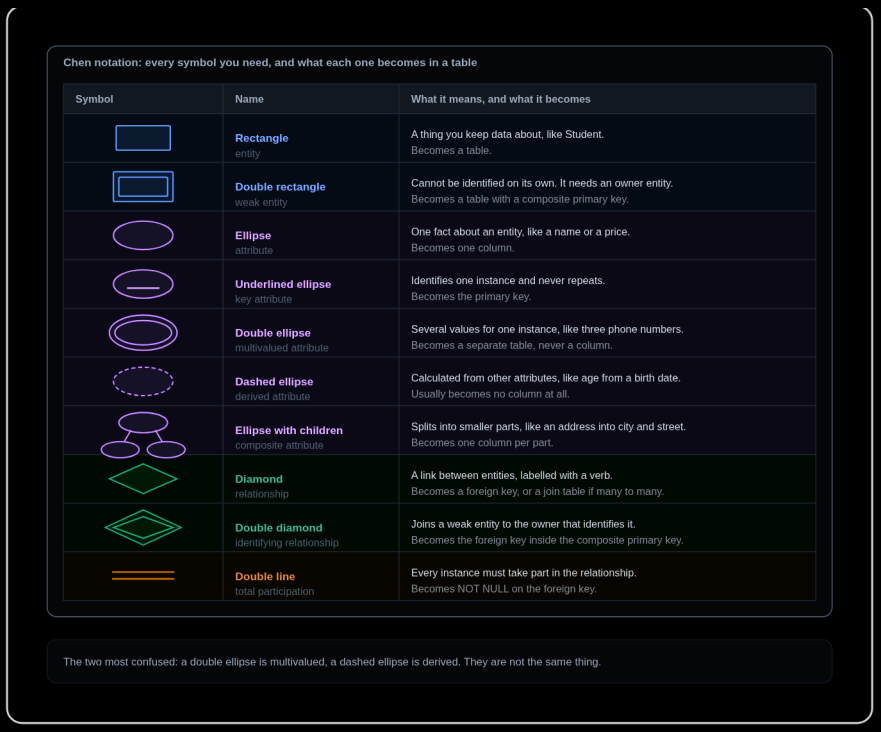
Entities: type, instance, set

An **entity** is an object or concept you record data about and can distinguish from others like it. Three loosely used words:

Term	Meaning	Example
Entity type	the general kind; what the rectangle is labelled	`Employee`
Entity instance	one particular member	employee #2, Ravi Menon
Entity set	all instances that exist right now	every row in the `employee` table today

You *draw* types; the database *stores* sets.

Figure — every Chen symbol and its meaning (course diagram)



The five kinds of attribute

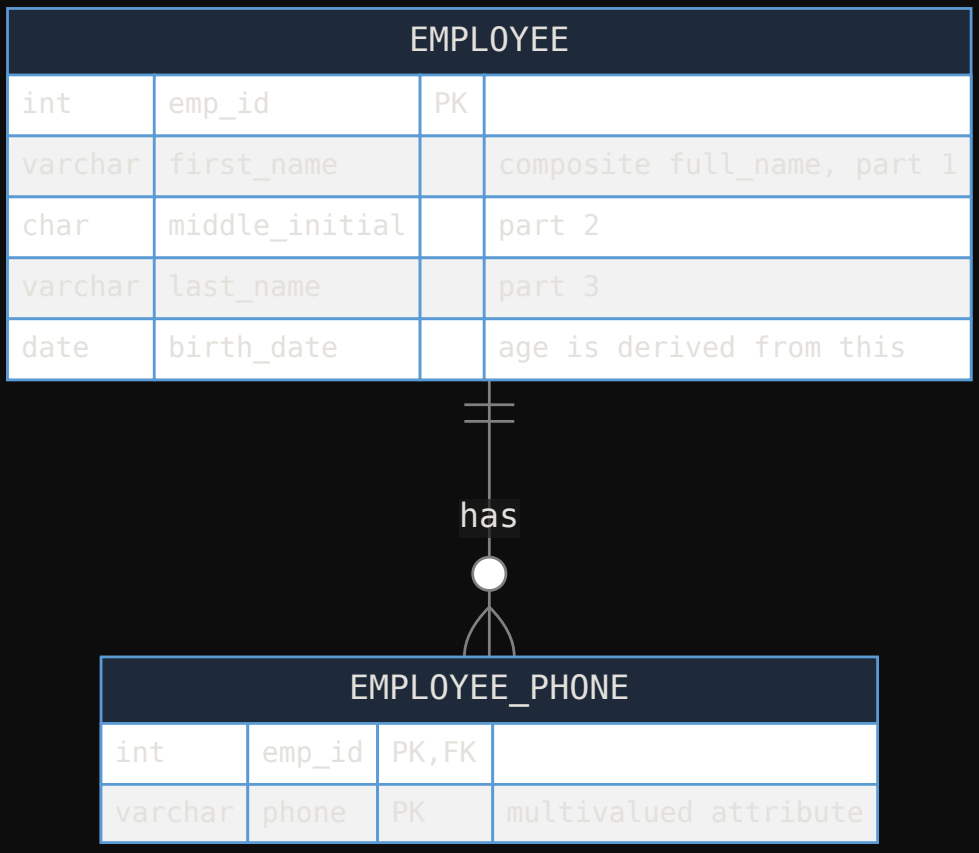
Kind	Chen symbol	Example	Becomes in tables
Key	ellipse, name **underlined**	`emp_id`	the primary key
Simple (atomic)	plain ellipse	`birth_date`	one column
Composite	ellipse with child ellipses	`full_name` = first + middle initial + last	one column **per part**
Multivalued	**double** ellipse	`phone` (an employee may have several)	a **separate table** (owner key + value)
Derived	**dashed** ellipse	`age` (from `birth_date`)	usually **no column** — computed on read

The two most confused: **double ellipse = multivalued**, **dashed ellipse = derived**.

(All three special shapes appear together in the employee diagram of §10.10.) Why each rule holds:

- **Composite → split** — an address in one column can't be searched by city (if you *never* query the parts, one column is fine).
- **Multivalued → own table** — a cell holds one value (§10.12); `phone1, phone2, phone3` is the mistake the double ellipse prevents.
- **Derived → compute** — a stored `age` is wrong a year later. Store only if expensive *and* kept fresh, or **frozen in time** (an order total at purchase-day prices isn't derivable from *today's* data).

Crow's foot has no special shapes: composites are already columns, multivalued attributes separate boxes, derived ones omitted or commented:



Self-check

1. `age` next to `birth_date` — which kind, what column? → **Derived**; no column, compute it.

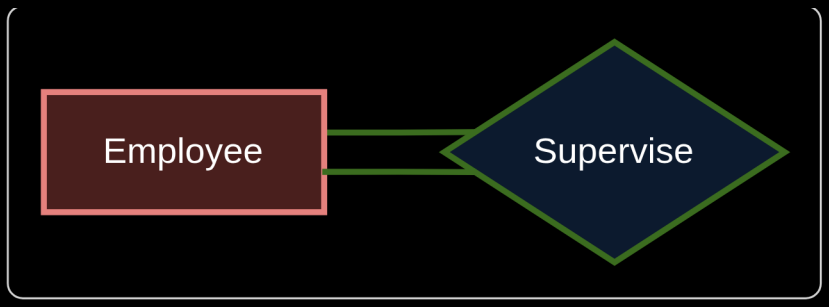
10.7 Relationships: degree, cardinality and participation

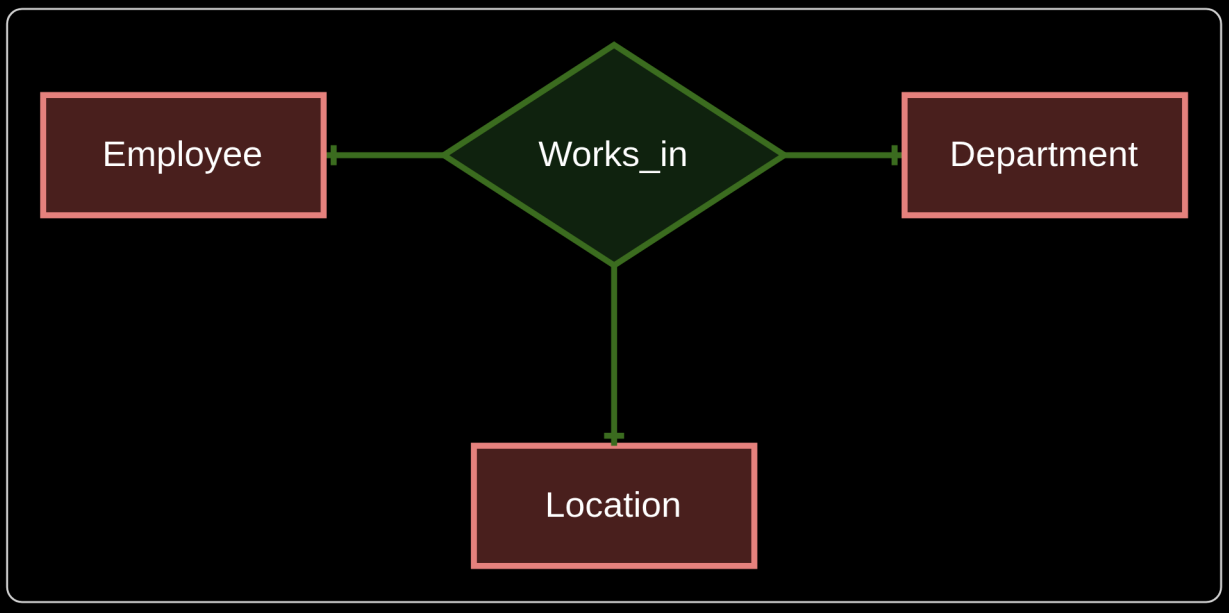
A **relationship** is an association between entities — a Chen diamond with a verb that reads well aloud ("Employee *works in* Department", not "*has*"). The **relationship type** is what you draw; the **relationship set** the actual stored pairs. Three separate questions describe it: **degree** (how many entity *types*), **cardinality** (how many *instances* per side), **participation** (is taking part *required*).

Degree

Degree	Name	Example	Tables
1	**Unary** (recursive)	an employee *supervises* employees	self-referencing FK: `employee.supervisor_id → employee.emp_id`
2	**Binary**	employee *works in* department	the usual case
3	**Ternary**	supplier *supplies* part *to* site	one table with three FKs, usually all three in the PK

Figure — unary and ternary relationships (course diagrams)





A ternary relationship records what three binary ones cannot: *which* combination is true. `code/04_ternary_and_weak_entity.sql` stores three supplier-part-site facts, keeps only binary pairs and joins them back:

```
WITH sp AS (SELECT DISTINCT supplier_id, part_id FROM supplies),
     ss AS (SELECT DISTINCT supplier_id, site_id FROM supplies)
SELECT ... FROM sp JOIN ss USING (supplier_id) ...   -- rebuild from pairs
```

Output:

```
--- the ternary fact: who supplies what, where
 supplier | part  | city
----------+-------+-------
 Acme     | bolt  | Leeds
 Acme     | panel | Pune
 Bolt Co  | bolt  | Pune
(3 rows)

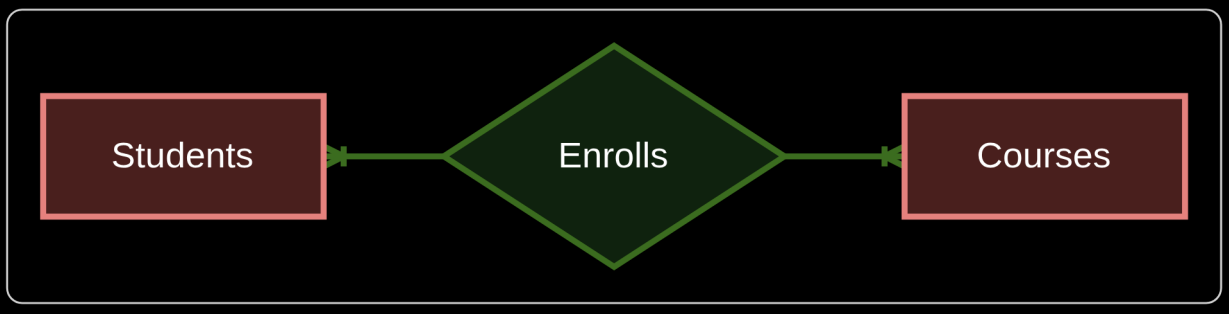
--- keep only binary pairs (supplier-part, supplier-site) and join them back:
 supplier | part  | city  | really_true
----------+-------+-------+-------------
 Acme     | bolt  | Leeds | t
 Acme     | bolt  | Pune  | f
 Acme     | panel | Leeds | f
 Acme     | panel | Pune  | t
 Bolt Co  | bolt  | Pune  | t
(5 rows)
```

From the pairs you can't tell whether Acme sends bolts to Pune. Ternaries are rare — check the fact truly needs all three (Chapter 11's 5NF, §11.13, formalises this).

Cardinality (mapping cardinality)

Cardinality	Meaning	Example
1:1	each side links to at most one of the other	employee ↔ payroll record
1:N	one on the left, many on the right	a department employs many employees
N:1	1:N viewed from the opposite side	many employees work in one department
M:N	many on both sides	employees work on many projects; a project has many employees

Figure — many-to-many in Chen notation (course diagram)



How each cardinality becomes tables

Cardinality	Rule	Why
1:N / N:1	FK on the **many** side (`employee.dept_id`)	the "one" side would need a list in a cell
M:N	a new **junction table** (also *join*, *bridge* or *associative* table) holding both FKs, which together form its PK (`works_on(emp_id, project_id)`)	there is no single side to put it on
1:1	FK on **either** side **+** `UNIQUE`; prefer the side with total participation so it can be `NOT NULL`	without `UNIQUE` it silently becomes 1:N
Unary	FK to the same table	

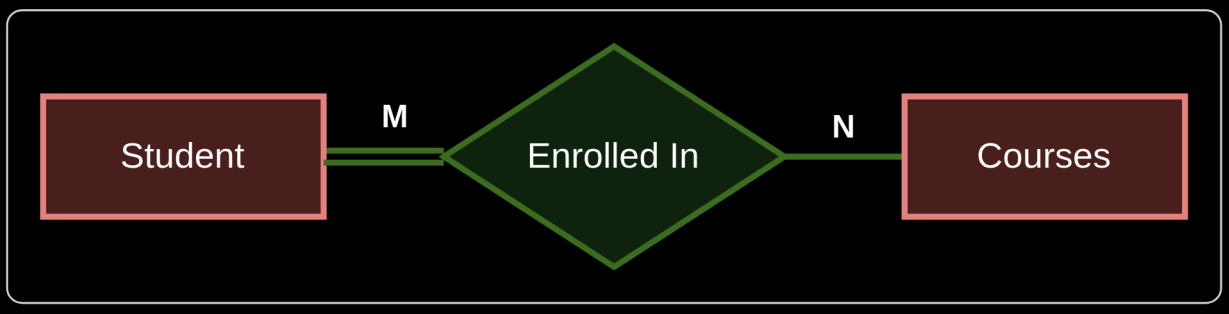
Memorise: **every M:N link costs three tables in SQL.**

Participation

Participation says whether every instance *must* take part.

	Chen drawing	Meaning	Becomes
Total (mandatory)	double line	every employee must work in a department	FK `NOT NULL`
Partial (optional)	single line	an employee may have no supervisor	FK allows `NULL`

Figure — total participation (double line) (course diagram)



Skipping participation is easy and expensive: it decides whether a column can be empty, affecting every query (outer vs inner joins, `IS NULL` checks).

The (min, max) notation

Tools often label each end **(min, max)** — fewest and most times one instance takes part:

Pair	Meaning	Participation	Crow's foot
`(0,1)`	at most once, maybe never	partial	`o|`
`(1,1)`	exactly once	total	`||`
`(0,N)`	any number, including none	partial	`o{`
`(1,N)`	at least once	total	`|{`

Minimum = participation; maximum = cardinality — one label says both. Careful: Chen's `(min,max)` is usually written next to the entity it describes ("each employee works in (1,1) departments") — the *opposite* end from crow's foot. Check the legend.

In the interview

30-second answer: "Degree counts the entity types involved — one (unary), two (binary) or three (ternary). Cardinality is how many instances link: 1:1, 1:N or M:N. Participation is whether taking part is mandatory. 1:N puts the FK on the many side; M:N needs a junction table; 1:1 is an FK plus UNIQUE; total participation makes the FK NOT NULL."

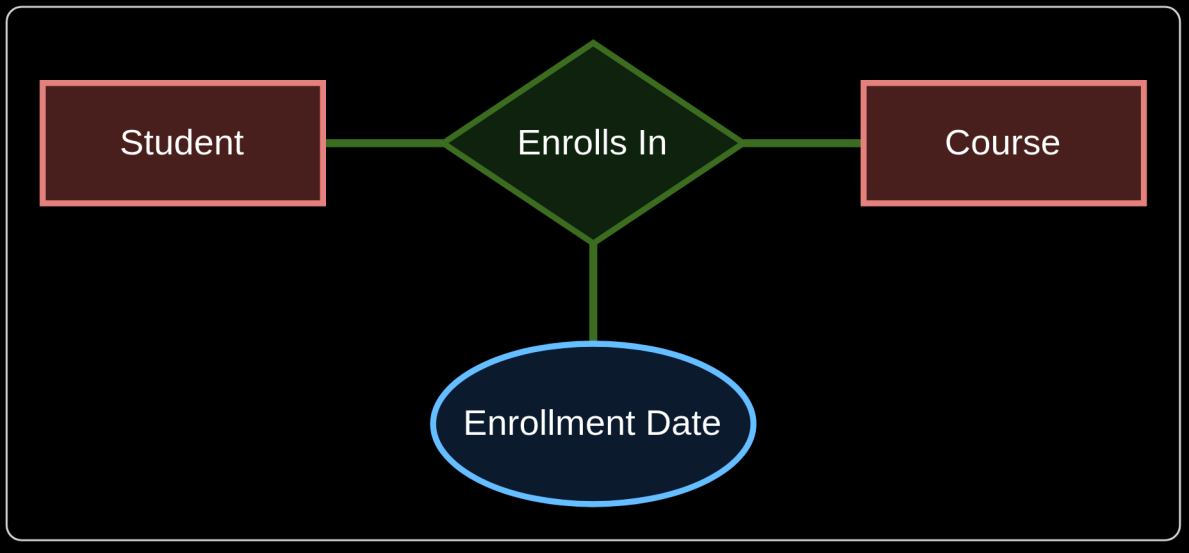
- *Where does the FK go in a 1:N?* → Many side. *Why not the one side?* → It would need a list in one cell, which violates 1NF (first normal form: one value per cell).
- *How do you enforce 1:1 in SQL?* → FK + `UNIQUE` on it (or share the same PK in both tables).
- *Enforce "every department has at least one employee" with a plain FK?* → No; an FK constrains the child side. Use a deferred constraint trigger (code run automatically on a change, checked at commit) or application logic.

10.8 Attributes of relationships

What it is

Some facts describe the **link**: an employee's hours on a project depend on *both*. In Chen such an attribute hangs off the diamond.

Figure — an attribute on a relationship (course diagram)



The test: *would you need to know both entities to know the value?*

Attribute	Needs	Belongs on
`weekly_hours`	employee **and** project	the relationship *works on*
`start_date` of an assignment	employee **and** project	the relationship
`title`	project only	Project
`last_name`	employee only	Employee

Where it lands in tables

Put the attribute wherever the relationship itself ended up:

- **M:N** → junction table column: `works_on(emp_id, project_id, weekly_hours)`.
- **1:N** → the many side beside the FK (`role_in_dept` on `employee` next to `dept_id`).
- **1:1** → wherever the FK went.

On an entity it forces a false rule: `weekly_hours` on `employee` allows one number however many projects.

When a relationship should become an entity

A link with many facts (rental: out, due, return times, fee) becomes an entity — same tables, one more box for a noun people already use. **Test:** if the business names the link ("rental", "loan", "enrolment", "booking"), make it an entity — why `rental` is an entity, not a diamond.

10.9 Weak entities

What it is

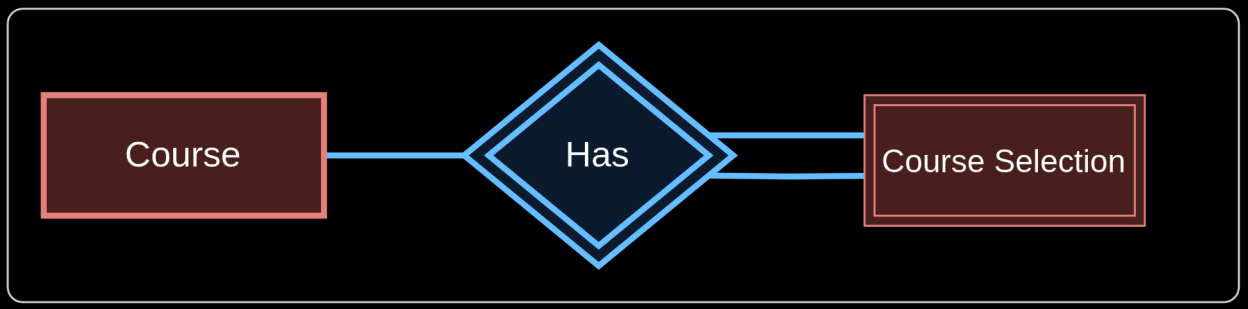
A **weak entity** has nothing of its own that singles it out: "Room 101" is ambiguous; only *(Main building, room 101)* identifies it. It depends on its **owner** (identifying / strong entity). Four properties:

1. **No key of its own** — nothing inside it is unique across all instances.
2. **A partial key** (also called a **discriminator**) — unique *within one owner* (`room_no`).
3. **Existence dependency** — delete the building and its rooms mean nothing.
4. **Total participation in the identifying relationship** — no room belongs to no building.

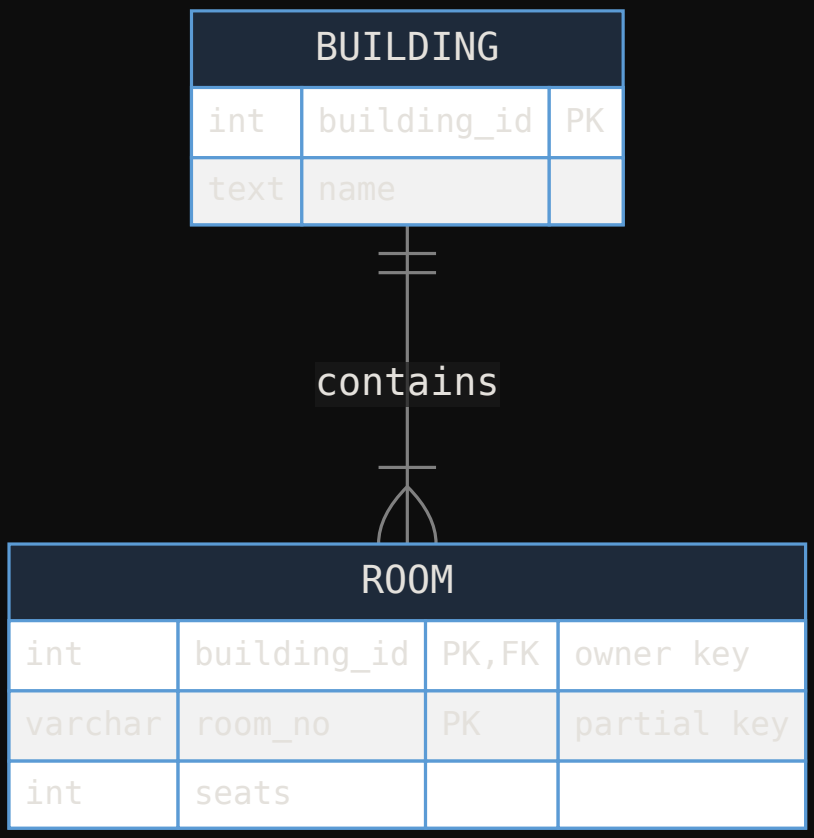
How it is drawn

Chen symbol	Meaning
double rectangle	the weak entity
double diamond	the **identifying relationship**
dashed underline	the partial key
double line (weak side)	total participation

Figure — a weak entity in Chen notation (course diagram)



In crow's foot (mermaid), a weak entity is shown by an **identifying (solid) line** and a PK that includes the owner's key:



What it becomes in SQL — two valid options

```sql
-- Option A: composite key = owner key + partial key
CREATE TABLE room_a (
  building_id int NOT NULL REFERENCES building ON DELETE CASCADE,
  room_no     varchar(6) NOT NULL,                -- partial key (discriminator)
  seats       int,
  PRIMARY KEY (building_id, room_no)
);
-- Option B: surrogate key + UNIQUE on (owner key, partial key)
CREATE TABLE room_b (
  room_id     int GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
  building_id int NOT NULL REFERENCES building ON DELETE CASCADE,
  room_no     varchar(6) NOT NULL,
  seats       int,
  UNIQUE (building_id, room_no)                   -- without this, duplicates slip in
);
```

Output (both options refuse a duplicate room, and cascade when the owner goes):

```
--- both options refuse a second room 101 in building 1
NOTICE:  room_a: duplicate key value violates unique constraint "room_a_pkey"
NOTICE:  room_b: duplicate key value violates unique constraint "room_b_building_id_room_no_key"
--- existence dependency: delete the Annex, its rooms go too
 room_a_rows | room_b_rows
-------------+-------------
           1 |           1
(1 row)
```

	Composite PK (A)	Surrogate PK (B)
Enforces "no duplicate partial key per owner"	automatically	only with the extra `UNIQUE` (easy to forget)
Referencing it from other tables	two columns	one column
Key width copied into child FKs/indexes	wider	narrow

Examples: order → order line (`(order_id, line_no)`), building → room, invoice → invoice line. **Weak or not is a design decision**: per-account transaction numbers are weak, a global transaction id strong. Default: option B *with* the unique constraint.

Self-check

1. Which SQL clause expresses existence dependency? → `ON DELETE CASCADE` on the owner FK.
2. Why is the owner FK `NOT NULL`? → Participation in the identifying relationship is total.

10.10 Worked example — an employee-management ER diagram

The brief

A company wants to:

1. Record each employee: name, date of birth, one or more phone numbers, and the department they belong to (every employee has exactly one department).
2. Keep one payroll record per employee (monthly salary, cash advance, tax); no payroll record is shared by two employees. Net pay is salary − advance − tax.
3. Record who supervises whom (the CEO has no supervisor).
4. Record which projects employees work on and for how many hours a week (many-to-many).
5. Record each employee's dependants (children, spouse) for insurance; a dependant is identified by name *within* an employee.

Step 1 — entities (nouns)

Employee, Department, Payroll, Project, Dependant ("record", "information", "number" are words about data, not things).

Step 2 — attributes, classified

Entity	Attribute	Kind	Reason
Employee	`emp_id`	key	issued once, never repeats
	`full_name`	composite	first name, middle initial, last name
	`birth_date`	simple	one value
	`phone`	**multivalued**	"one or more"
	`age`	**derived**	from `birth_date`
Payroll	`payroll_id`	key	
	`monthly_salary`, `cash_advance`, `tax`	simple, **stored**	typed in by a person — nothing computes them
	`net_pay`	**derived**	salary − advance − tax
Department	`dept_id` (key), `dept_name`		
Project	`project_id` (key), `title`		
Dependant	`dep_name`	**partial key**	unique only within one employee
	`relationship`	simple	

The trap is salary vs net pay: **typed in → stored; computable from stored attributes → derived.**

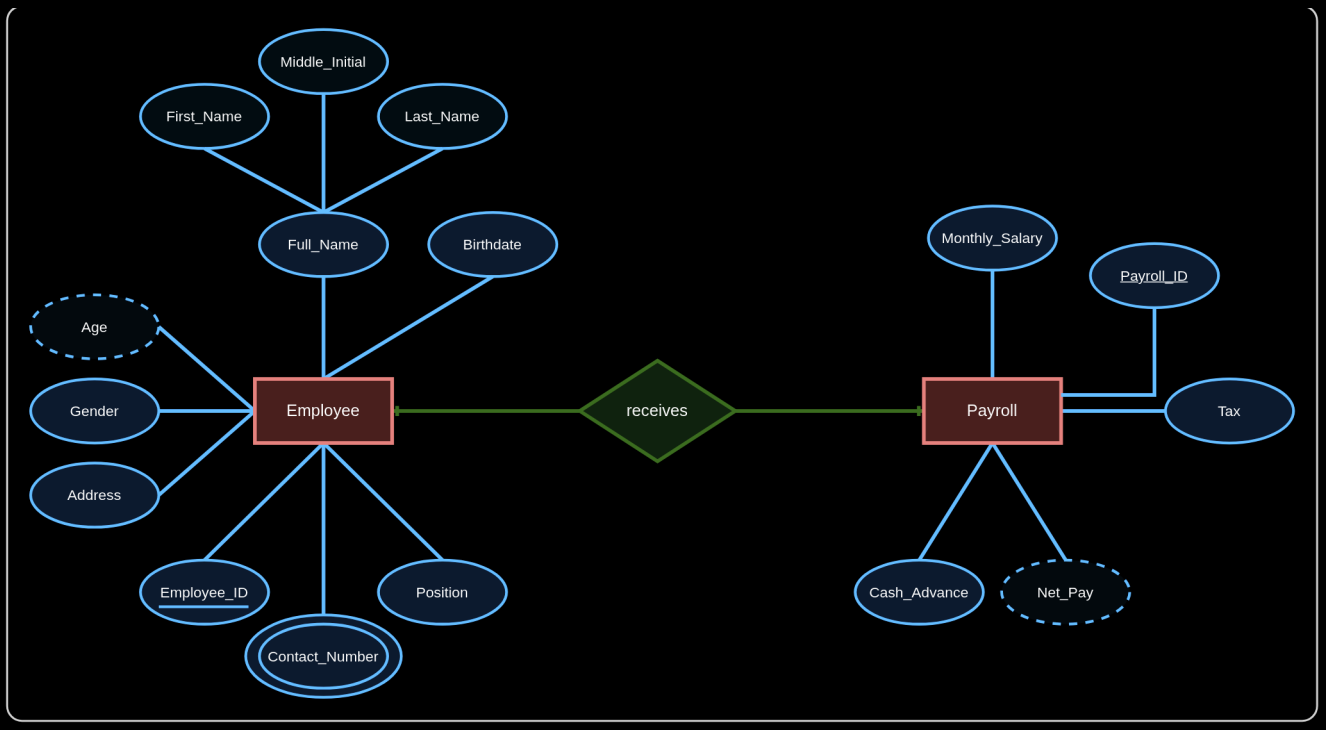
Steps 3–4 — relationships with degree, cardinality and participation

Relationship	Degree	Cardinality	Participation
Employee *works in* Department	binary	N:1	employee total (brief says "exactly one")
Employee *is paid through* Payroll	binary	1:1	both total
Employee *supervises* Employee	unary	1:N	partial (CEO has none)
Employee *works on* Project, attribute `weekly_hours`	binary	M:N	partial
Employee *has* Dependant (identifying)	binary	1:N	dependant total

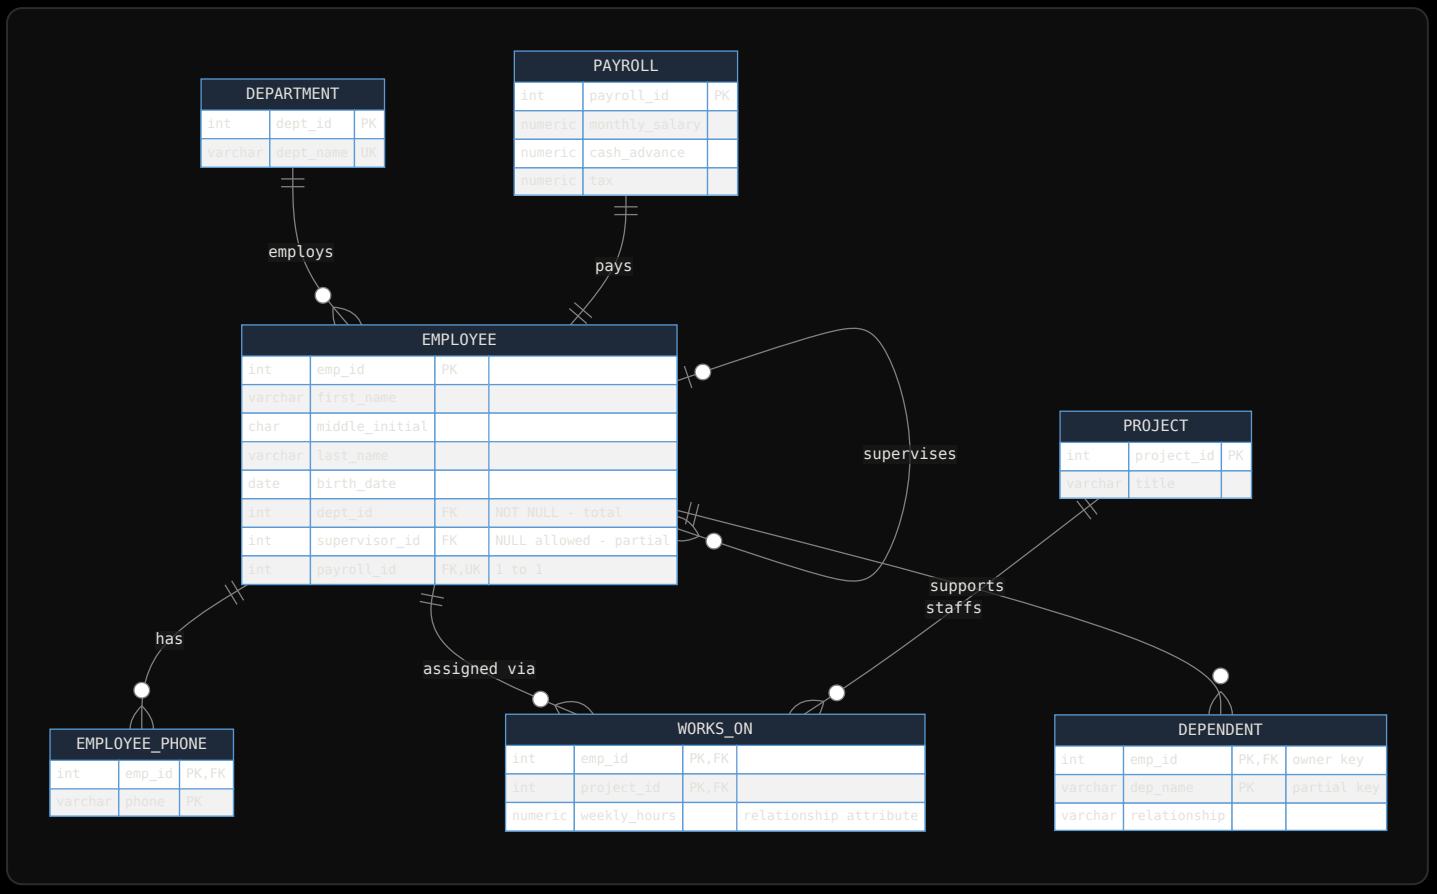
Step 5 — the diagram

The course's Chen diagram of the core (Employee + Payroll) shows every special symbol at once:

Figure — Employee and Payroll in Chen notation: underlined keys, composite Full_Name, double-ellipse Contact_Number, dashed Age and Net_Pay (course diagram)



The full design in crow's-foot notation:



Read it back as sentences — if one sounds wrong, the diagram is: "Each employee works in exactly one department; a department employs zero or more (the brief doesn't require staff). Each employee has exactly one payroll record and vice versa. An employee may supervise several and has at most one supervisor…"

Design check: does Payroll deserve its own table?

A 1:1 total on **both** sides usually is one thing in two places. Keep two tables for **access control** (`GRANT SELECT ON payroll TO hr` — `GRANT` gives a role a permission, Chapter 14), **row width** (keep the hot `employee` row narrow) or **optional data** (no row instead of several `NULL`s). Merge to avoid the join and make

"employee without payroll" impossible. Either is fine in an interview **if you notice the question and justify it**; here, access control wins. §10.15 converts and runs it.

10.11 The relational model

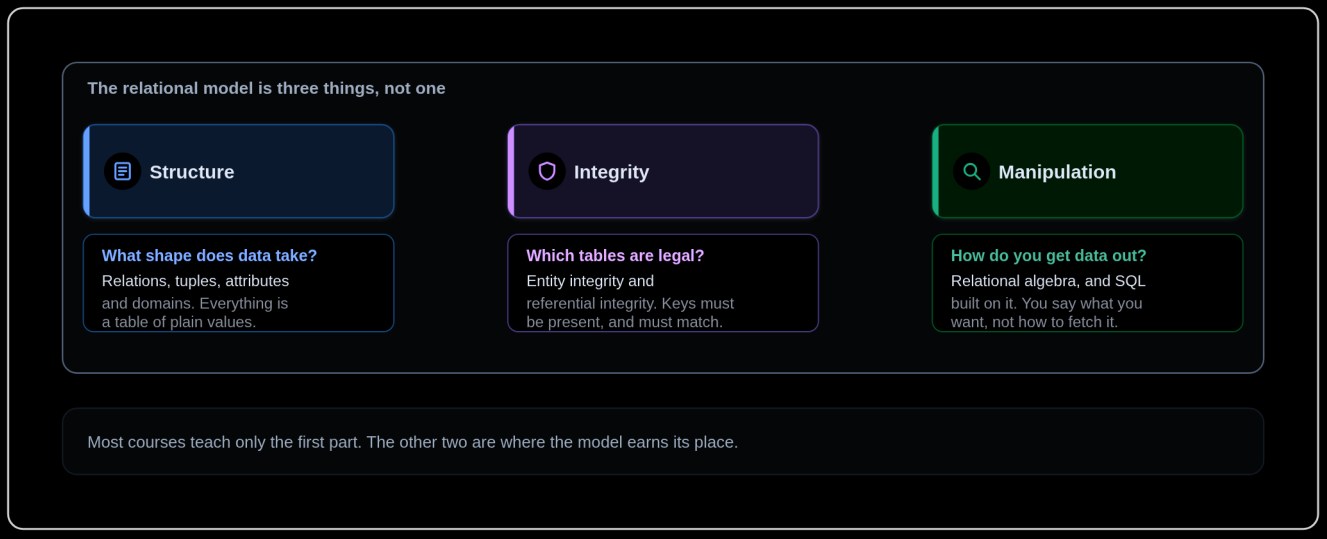
What it is

A database stores **relations** (tables). The **relational model** (**Edgar F. Codd, 1970**): store all data as tables of values; programs *describe* what they want and the system works out how. It replaced hierarchical and network databases (records linked in trees or webs), whose programs walked stored pointers (disk addresses) and broke on every reorganization (Chapter 1).

The three parts of the model

Part	Question it answers	Content
Structure	What shape is the data?	Only relations: a name, named columns, rows of single values. No nested records, arrays or pointers.
Integrity	Which tables are legal?	Keys and constraints declared once on the table; the DBMS refuses any change that breaks them.
Manipulation	How do I get data out?	Declarative queries (SQL: you state *what* you want, not how to fetch it), built on **relational algebra**: operations (select, project, join, union, difference …) that take relations and return a relation.

Figure — structure, integrity, manipulation (course diagram)



Links are values, not pointers

`employee.dept_id = 10` links to department 10 — the shared *value* is the whole link. It survives copying to another server, reloading in another order or rebuilding indexes, and the DBMS can **check** it (an FK refuses `dept_id = 99` if no department 99 exists), which a raw pointer can't.

What you get

- **One rule reaches every program** — constraints live with the data.
- **Physical data independence** — queries name *what*, not *how*, so storage and indexes change freely.
- **Some logical data independence** — adding a column or hiding a split behind a view (a saved query acting as a table) breaks nothing; *not* total — renaming a column or shrinking its type still breaks users.

What it does not promise

- **Nothing about storage** — files, pages, indexes belong to the storage engine (the DBMS part reading and writing files; Chapter 12).
- **Joins aren't free** — work grows with table size.
- **A table isn't always an entity** — also M:N relationships (`works_on`) and multivalued attributes (`employee_phone`).

10.12 Relations, tuples, attributes — the precise vocabulary

The core words

Formal term	SQL term	Older file term	Meaning
relation	table	file	named set of rows with fixed named columns
tuple	row	record	one value for each attribute
attribute	column	field	a named column
domain	data type + constraints	field type	the set of allowed values for an attribute
degree (arity)	number of columns	number of fields	a property of the schema
cardinality	number of rows	number of records	a property of the current data

Figure — anatomy of a relation (course diagram)



- **Domain** is what you *mean*, more than a type: `int` is a type; "integers 1–500" is room capacity's domain, enforced with `CHECK`.
- **Schema vs instance.** The **relation schema** is the design (`room(room_id, code, capacity)`), rarely changing; the **relation instance** is the current rows. A key belongs to the *schema* — true for any legal data; current rows prove nothing.
- **Two meanings of "cardinality"**: instances per link (1:N, §10.7) vs row count here.

Five properties of every relation

1. **One value per cell** (atomic) — why multivalued attributes need a table; the heart of 1NF (§11.7).
2. **Rows have no inherent order** — ask with `ORDER BY`.

3. **Columns have no inherent order** — referred to by name (SQL tables remember declaration order for `SELECT *`; still list columns explicitly).
4. **Attribute names are unique within a relation.**
5. **No duplicate tuples** — *SQL differs*: identical rows are accepted unless a `PRIMARY KEY` or `UNIQUE` forbids them.

NULL

`NULL` marks a **missing** value — *unknown* (grade not computed) or *not applicable* (no middle name); not zero or ''. Under **three-valued logic** comparing with `NULL` gives *unknown* and `WHERE` keeps only *true* (Chapter 2). From `code/03_keys_and_integrity.sql`:

```
SELECT (SELECT count(*) FROM lecturer WHERE room_id = NULL)  AS eq_null,
       (SELECT count(*) FROM lecturer WHERE room_id IS NULL) AS is_null,
       count(room_id) AS count_col, count(*) AS count_star
FROM lecturer;
```

Output:

```
 eq_null | is_null | count_col | count_star
---------+---------+-----------+------------
       0 |       2 |         1 |          3
(1 row)
```

`= NULL` matches nothing, `IS NULL` works, `count(col)` skips nulls, `count(*)` counts rows — every query on a nullable column must decide how to treat missing values.

Self-check

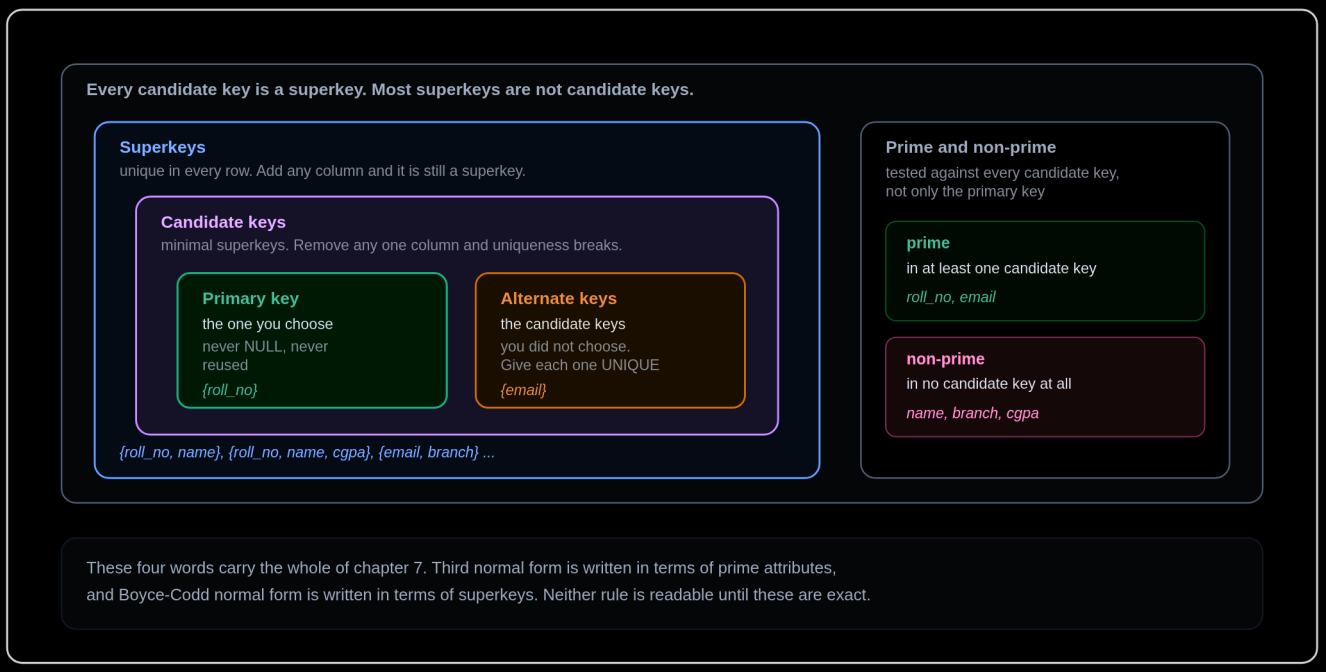
1. `student(roll_no, name, cgpa, dept)` with 8,000 rows: degree, cardinality? → **4**, **8,000**.
2. Why can't rows prove `{name}` is a key? → Keys are schema properties; rows can only *disprove* one (two equal names).

10.13 Keys in relational databases

Two identical-looking rows and a ticket to change "that student" have no answer without **keys**, which make identity exist. Chapter 11's rules use exactly these definitions (§9.3 shows them as DDL). Example: `lecturer(lecturer_id, email, name, room_id)`, ids issued once, emails unique.

Key	Definition	Example
Superkey	a set of attributes that can **never** repeat in two rows (in any legal instance)	`{lecturer_id}`, `{email}`, `{lecturer_id, name}`, all four columns …
Candidate key	a **minimal** superkey — remove any attribute and it stops being unique	`{lecturer_id}`, `{email}`
Primary key	the one candidate key you *choose*; unique and never `NULL`	`lecturer_id`
Alternate key	any candidate key other than the chosen primary one — declare it `UNIQUE`	`email`
Composite key	any key with 2+ attributes (a shape, not a separate kind)	`booking(room_id, day, slot)`
Foreign key	attributes whose values must match a key in a (possibly the same) table, or be `NULL`	`lecturer.room_id → room.room_id`
Natural key	made of real-world data	email, ISBN, frame number
Surrogate key	system-generated, no business meaning	identity integer, UUID

Figure — superkeys ⊇ candidate keys ⊇ primary key, with prime and non-prime attributes (course diagram)



Key facts to be exact about:

- **Adding columns to a superkey keeps it a superkey**, so there are many; candidate keys have nothing spare. Minimality is *removability*, not size — a 3-column candidate key is fine if none can go. All candidate keys are superkeys, rarely the reverse.
- **Prime attribute** = in *some* candidate key (any, not only the primary); **non-prime** = in none. In `lecturer`: `lecturer_id`, `email` prime; `name`, `room_id` non-prime. 3NF uses these words (§11.9).
- **Composite**: `enrolment(student_id, course_id, term, grade)` — retakes mean only `{student_id, course_id, term}` identifies a row.
- **Foreign keys** may be `NULL` (no link yet — not a violation), may reference any `PRIMARY KEY` **or** `UNIQUE` set, and may reference their own table (`employee.supervisor_id` — a unary relationship as a column).
- **Surrogate vs natural**: surrogates are narrow and never change; natural keys are readable and may save a join but change with the business (a new roll-number format updates every FK). **Default: surrogate PK + `UNIQUE` natural key** (as `customer` in `00_schema.sql`).

In the interview

> **30-second answer:** "A superkey uniquely identifies rows; a candidate key is a minimal superkey; the primary key is the candidate I pick — unique and not null; the other candidates are alternate keys and get UNIQUE constraints. A foreign key references a primary or unique key and may be null. My default is a generated surrogate PK, with the natural key kept unique."

- *Is `{id, name}` a candidate key?* → A superkey, not a candidate key: `name` can be dropped.
- *UUID vs bigint surrogate?* → See §9.4: coordination-free but 16 bytes vs 8, and random UUIDs scatter B-tree inserts; v7 reduces that.
- *Is an attribute of an alternate key prime?* → Yes — prime means "in any candidate key".

10.14 Relational integrity constraints

What it is

An **integrity constraint** is a condition stored data may never break; the DBMS tests it on every `INSERT`, `UPDATE`, `DELETE` from any program and refuses violations.

Figure — the integrity constraints and the SQL that declares each (course diagram)

Five rules the database enforces for every application that connects

Constraint	What it refuses	How you declare it
Domain the allowed set of values	Marks of 210 out of 100. The text "twenty" in a number column.	`marks INT` `CHECK (marks BETWEEN 0 AND 100)`
Entity integrity every row is identifiable	A row whose primary key is empty. Half of a composite key left blank.	`PRIMARY KEY (loan_id)` NOT NULL comes with it automatically

Referential integrity
a foreign key points at a real row, or at nothing
	A loan for a member id that does not exist.
	Deleting a member who still has loans, unless you chose CASCADE or SET NULL.
		`FOREIGN KEY (member_id)`
		`REFERENCES Member(member_id)`
		`ON DELETE RESTRICT`

Key constraint no duplicate values	Two members sharing one email address. Several NULLs are usually still allowed.	`UNIQUE KEY uq_member_email (email)` add NOT NULL if the value is required
Check constraint any condition you can state	A due date earlier than the issue date. A negative price.	`CHECK (due_date > issue_date)` one rule, held for every application

The same rule written as an if statement in one application holds only for that application. In the schema it holds for all of them.

Constraint	Rule	SQL
Domain	a value must be in the attribute's domain	the data type, plus `CHECK (capacity BETWEEN 1 AND 500)`
Entity integrity	no part of a primary key may be `NULL`	`PRIMARY KEY` (implies `NOT NULL`)
Key	candidate keys are unique	`PRIMARY KEY`, `UNIQUE`
Referential integrity	an FK value matches an existing referenced row, or is `NULL` — no **orphan rows**	`FOREIGN KEY … REFERENCES …`
(also declared the same way)	must always have a value / default when omitted	`NOT NULL`, `DEFAULT` (DEFAULT never refuses anything, so it is not strictly a constraint)

Entity integrity: a PK identifies a row, and a row of unknown (`NULL`) identity can't be distinguished; in composite keys *every* column must be present.

What happens when a referenced (parent) row is deleted or its key updated

The `ON DELETE` / `ON UPDATE` actions are tabulated in §9.3: `NO ACTION` (default; error at statement end, can be `DEFERRABLE` to transaction end), `RESTRICT` (immediate error, not deferrable), `CASCADE` (for weak entities and pure child data — phones, order lines — dangerous elsewhere), `SET NULL` (optional relationships; column must be nullable), `SET DEFAULT` (rare). Dialects: MySQL/InnoDB treats `NO ACTION` as `RESTRICT` and rejects `SET DEFAULT`; SQL Server has no `RESTRICT` keyword; SQLite enforces FKs only after `PRAGMA foreign_keys = ON`; MySQL enforces `CHECK` only from 8.0.16.

Running it (`code/03_keys_and_integrity.sql`)

Each failing statement runs in a `DO` block (one-off PostgreSQL procedural code) that prints the error as a `NOTICE`, so the script continues.

```
CREATE TABLE room (
  room_id  int PRIMARY KEY,                                  -- entity integrity: unique + NOT NULL
  code     varchar(10) NOT NULL UNIQUE,                      -- alternate (candidate) key
  capacity int NOT NULL CHECK (capacity BETWEEN 1 AND 500)   -- domain constraint
);
CREATE TABLE lecturer (
  lecturer_id int PRIMARY KEY,
  name        varchar(60) NOT NULL,
  email       varchar(120) UNIQUE,                           -- UNIQUE but nullable
  room_id     int REFERENCES room(room_id) ON DELETE SET NULL  -- optional link
);
CREATE TABLE booking (
  room_id  int  NOT NULL REFERENCES room(room_id) ON DELETE RESTRICT,
  day      date NOT NULL,
  slot     int  NOT NULL CHECK (slot BETWEEN 1 AND 8),
  PRIMARY KEY (room_id, day, slot)                           -- composite key
);
-- each violation is attempted inside:  DO $$ BEGIN ... EXCEPTION WHEN ... THEN RAISE NOTICE ... END $$;
```

Output:

```
--- 1. domain: capacity 900 is outside the allowed values
NOTICE:  new row for relation "room" violates check constraint "room_capacity_check"
--- 2. entity integrity: a primary key may not be NULL
NOTICE:  null value in column "room_id" of relation "room" violates not-null constraint
--- 3. key constraint: the alternate key code must be unique
NOTICE:  duplicate key value violates unique constraint "room_code_key"
--- 4. referential integrity: no orphan rows
NOTICE:  insert or update on table "booking" violates foreign key constraint "booking_room_id_fkey"
--- 5. UNIQUE lets several rows have NULL (Dr Silva and Dr Wu both have no email)
   name   | email
----------+-------
 Dr Silva |
 Dr Wu    |
(2 rows)

--- 6. ON DELETE RESTRICT: room 1 still has bookings
NOTICE:  update or delete on table "room" violates foreign key constraint "booking_room_id_fkey" on table
"booking"
--- 7. ON DELETE SET NULL: deleting room 2 empties Dr Silva's room_id
   name   | room_id
----------+---------
 Dr Park  |       1
 Dr Silva |
 Dr Wu    |
(3 rows)
```

UNIQUE and NULL — a classic surprise

`NULL` ≠ `NULL`, so most databases let a `UNIQUE` column hold **many** nulls (output 5) — add `NOT NULL` if "unique when present" isn't wanted. PostgreSQL, MySQL and SQLite allow many (PostgreSQL 15+ `UNIQUE NULLS NOT DISTINCT` allows one); SQL Server allows **only one** (use a filtered unique index `WHERE col IS NOT NULL`) — see §9.3.

Why rules belong in the schema

They **hold for everybody** (web app, mobile backend, nightly import, a midnight hand fix), **can't drift** (one copy), and **document meaning**. Costs — a little work per write, and a thoughtless `CASCADE` wiping far more than meant — are small next to one central enforcement point.

Common mistakes

- Undeclared alternate keys → two customers with one email.
- `ON DELETE CASCADE` on a "real" relationship (deleting a customer silently deletes invoices).
- Assuming a nullable `UNIQUE` means "at most one missing".
- FKs left off "for speed" — orphans appear, joins silently lose rows.

10.15 Converting an ER diagram to tables

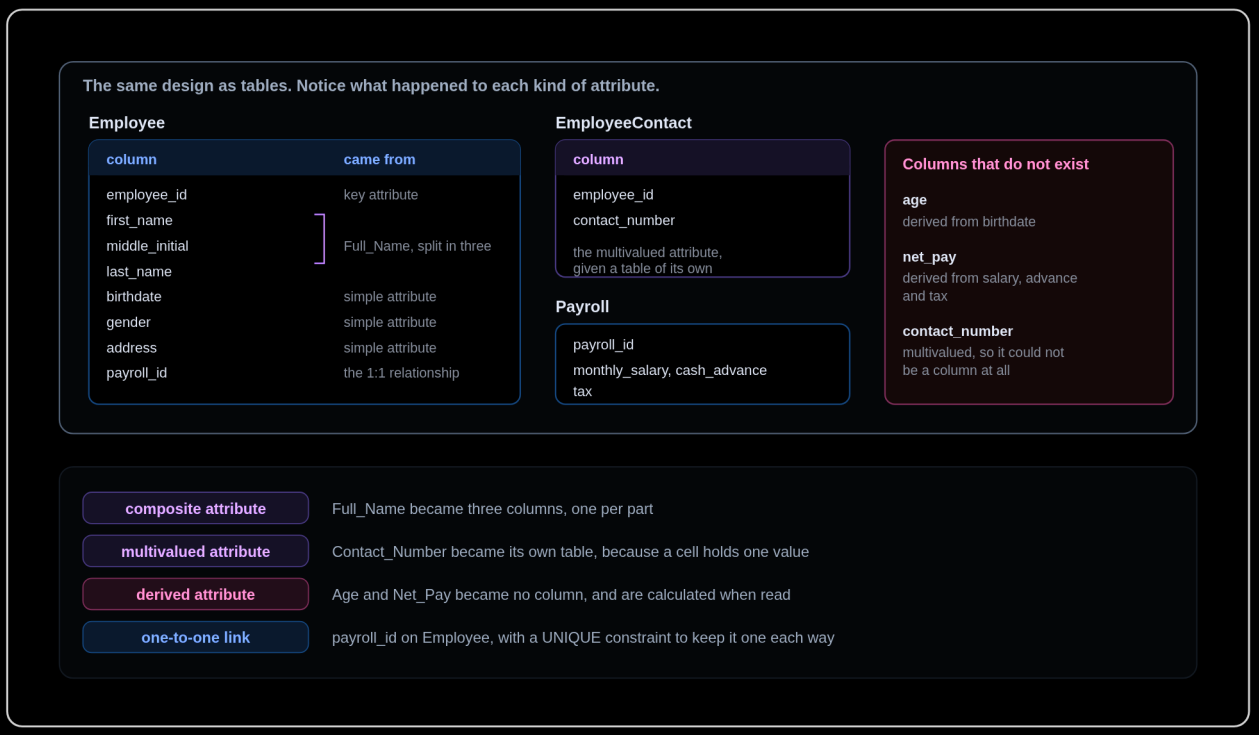
The procedure

Once the diagram is right, conversion is mechanical:

#	ER construct	Relational result
1	Strong entity	a table; key attribute → primary key
2	Simple attribute	a column
3	Composite attribute	one column per part (no column for the whole)
4	Derived attribute	no column (compute in the query or a view)
5	Multivalued attribute	own table: owner key + value, both in the PK, `ON DELETE CASCADE`
6a	1:N / N:1 relationship	FK on the many side
6b	M:N relationship	junction table: both FKs, together the PK
6c	1:1 relationship	FK + `UNIQUE` on either side, prefer the total-participation side
6d	Unary relationship	FK to the same table
6e	Ternary relationship	table with three FKs (PK usually all three)
7	Relationship attribute	follows the relationship: junction table (M:N) or many side (1:N)
8	Weak entity	table with PK = owner key + partial key; owner FK `NOT NULL`, usually `ON DELETE CASCADE`
9	Participation	total → FK `NOT NULL`; partial → FK nullable

Create tables in **dependency order** — parents first, since an FK can only point at an existing table.

Figure — the course's Employee/Payroll example as tables: composite split, multivalued attribute a table, derived attributes gone (course diagram)



Running the procedure on §10.10 (`code/02_employee_er_to_tables.sql`)

```
CREATE TABLE department (                       -- strong entity -> table
  dept_id   int PRIMARY KEY,
  dept_name varchar(60) NOT NULL UNIQUE
);
CREATE TABLE payroll (                          -- strong entity -> table
  payroll_id     int PRIMARY KEY,
  monthly_salary numeric(10,2) NOT NULL CHECK (monthly_salary >= 0),  -- stored (typed in)
  cash_advance   numeric(10,2) NOT NULL DEFAULT 0,
  tax            numeric(10,2) NOT NULL DEFAULT 0
  -- net_pay is DERIVED -> no column (computed in the view below)
);
CREATE TABLE employee (
  emp_id         int PRIMARY KEY,                              -- key attribute
  first_name     varchar(40) NOT NULL,                         -- composite attribute full_name
  middle_initial char(1),                                      --   split into
  last_name      varchar(40) NOT NULL,                         --   one column per part
  birth_date     date NOT NULL,                                -- simple; age is DERIVED
  dept_id        int  NOT NULL REFERENCES department(dept_id),  -- 1:N, total participation -> NOT NULL
  supervisor_id  int  REFERENCES employee(emp_id),             -- unary 1:N, partial -> NULL allowed
  payroll_id     int  NOT NULL UNIQUE REFERENCES payroll(payroll_id)  -- 1:1 -> FK + UNIQUE
);
CREATE TABLE employee_phone (                   -- multivalued attribute -> own table
  emp_id int         NOT NULL REFERENCES employee(emp_id) ON DELETE CASCADE,
  phone  varchar(20) NOT NULL,
  PRIMARY KEY (emp_id, phone)
);
CREATE TABLE project (project_id int PRIMARY KEY, title varchar(80) NOT NULL);
CREATE TABLE works_on (                         -- M:N -> junction table
  emp_id     int NOT NULL REFERENCES employee(emp_id),
  project_id int NOT NULL REFERENCES project(project_id),
  weekly_hours numeric(4,1) NOT NULL CHECK (weekly_hours > 0),  -- relationship attribute
  PRIMARY KEY (emp_id, project_id)
);
CREATE TABLE dependent (                        -- weak entity
  emp_id       int NOT NULL REFERENCES employee(emp_id) ON DELETE CASCADE,  -- owner key
  dep_name     varchar(40) NOT NULL,                                       -- partial key
  relationship varchar(20) NOT NULL,
  PRIMARY KEY (emp_id, dep_name)                                           -- owner key + partial key
);
-- Derived attributes live in a view (fixed "as of" date so the output is repeatable):
CREATE VIEW employee_card AS
SELECT e.emp_id,
       concat_ws(' ', e.first_name, e.middle_initial, e.last_name)   AS full_name,
       date_part('year', age(date '2026-09-28', e.birth_date))::int  AS age,
       p.monthly_salary - p.cash_advance - p.tax                    AS net_pay,
       (SELECT string_agg(phone, ', ' ORDER BY phone) FROM employee_phone ph
         WHERE ph.emp_id = e.emp_id)                                 AS phones
FROM employee e JOIN payroll p USING (payroll_id);
```

Output:

```
--- employee_card (derived: full_name, age, net_pay; multivalued: phones)
 emp_id |  full_name   | age | net_pay |              phones
--------+--------------+-----+---------+---------------------------------
      1 | Aisha Khan   |  41 | 4900.00 | +44 20 7946 0001
      2 | Ravi K Menon |  33 | 4250.00 | +91 80 4000 0003, +91 98450 00002
      3 | Sofia Rossi  |  28 | 3100.00 |
(3 rows)

--- the unary relationship: who reports to whom
 employee | supervisor
----------+------------
 Aisha    |
 Ravi     | Aisha
 Sofia    | Aisha
(3 rows)

--- 1:1 is enforced by UNIQUE: a second employee on payroll 501 is refused
NOTICE:  refused: duplicate key value violates unique constraint "employee_payroll_id_key"
--- weak entity: deleting the owner removes its dependents (ON DELETE CASCADE)
 dependents_left | phones_left
-----------------+-------------
               0 |           0
(1 row)
```

(Before deleting employee 2 the script removes their `works_on` rows and clears `supervisor_id` pointing at them — those FKs use default `NO ACTION`, so the database would otherwise refuse: referential integrity at work.)

Checking a conversion

1. **Count:** 5 entities + 1 multivalued attribute + 1 M:N = 7 tables. ✓
2. **Every relationship produced something** (FK, UNIQUE FK or table) — otherwise it was missed.
3. **No derived attribute survived** as a column (or a comment says why).
4. **No multivalued attribute left in its parent** — a list column violates 1NF (Chapter 11).

10.16 Best practices for ER → relational mapping

Large diagrams need judgement in a few places:

Situation	Recommendation	Why / exception
Composite attribute	Store the parts; do **not** also store the whole. Build the whole in a query or view.	Two copies disagree after the first partial update. Exception: never query the parts → one column is fine.
Multivalued attribute	Always a separate table with a composite PK (owner key + value). Never numbered columns.	The composite PK stops the same phone being stored twice for one person.
Derived attribute	Compute on read. Store only if expensive *or* it must be frozen in time (order total at purchase prices) — and comment the column.	Otherwise the next engineer deletes it as a "mistake".
1:1 with total participation both sides	Probably one table. Keep two only for access control, row width, or mostly-absent optional data.	Two tables always joined 1:1 are cost without benefit.
M:N	The junction table is mandatory; relationship attributes live there; give it a business name (`enrolment`, `booking`) not `student_course`.	Add a surrogate key only if other tables must reference a single row — and keep `UNIQUE` on the pair so nobody enrols twice.
Weak entity with a surrogate key	Keep `UNIQUE (owner_key, partial_key)`.	Otherwise the rule the weak entity expressed is enforced nowhere.
Naming	One style for tables and columns; FKs named after their target; **role names** when two FKs hit the same table (`origin_airport_id`, `destination_airport_id`, not `airport_id_1/2`).	Readable joins.
Normalization	Run the normal forms over the result (Chapter 11); expect one or two further splits.	A careful ERD lands *near* 3NF, rarely exactly on it.
Large diagrams	Create parent tables before children and finish one area (say orders) before starting the next (inventory); record each choice you made that the rules left open (merged 1:1s, stored derived columns, surrogate keys on weak entities).	Those are the choices someone will question later.

Common mistakes

- Converting an M:N with an FK on one side (only works if it was really 1:N).
- Keeping `full_name` *and* `first_name` / `last_name`.
- Surrogate key on a junction table without the pair `UNIQUE`.
- Creating child tables before parents in a migration script.

In the interview

30-second answer (for "turn this ERD into tables"): "Each strong entity becomes a table with its key as PK. Composite attributes split into columns, derived ones are computed, multivalued ones get their own table. 1:N puts the FK on the many side, M:N gets a junction table holding both keys plus any relationship attributes, 1:1 is an FK with UNIQUE. A weak entity's PK combines its owner's key with its partial key, with cascade delete, and total participation means NOT NULL. Then I check the result against the normal forms."

- *Where does `grade` go in student–course?* → In the enrolment junction table.
- *Would you merge a 1:1?* → Yes unless access control, row width or sparse optional data justify splitting.

Self-check

1. A junction table has a surrogate PK. What else must it have? → `UNIQUE` on the two FK columns.
2. Is the converted schema automatically in 3NF? → No; check the normal forms.

Key terms

Defined where first used: data model (conceptual / logical / physical), data independence, entity type / instance / set, key / simple / composite / multivalued / derived attribute, relationship type / set, degree, cardinality (ER: 1:1, 1:N, M:N; relation: row count), participation (total / partial), (min, max), junction table, weak entity / owner / partial key (discriminator) / identifying relationship, Chen vs crow's foot, relation / tuple / attribute / domain, schema vs instance, superkey / candidate / primary / alternate / composite / foreign / natural / surrogate key, prime / non-prime attribute, domain / entity / key / referential integrity, orphan row (child whose FK points at nothing), RESTRICT / NO ACTION / CASCADE / SET NULL / SET DEFAULT. Glossary: chapter 20.

Chapter 10 TL;DR

- **Model before tables.** Entities (nouns) → tables, attributes → columns, relationships (verbs) → FKs or junction tables. A fact you do not model is a question you cannot answer.
- **Three levels:** conceptual (for the business), logical (attributes, keys, normalization), physical (types, indexes, one DBMS).
- **Process:** requirements + questions → entities → relationships → conceptual review → logical → physical → validate; loop back freely.
- **Best practices:** constraints in the schema, one naming style, dates as dates, money as `numeric`, explicit `NOT NULL`, surrogate PK + unique natural key, 3NF by default, index real queries.
- **ER notation:** Chen (rectangle, ellipse, diamond; double ellipse = multivalued, dashed = derived, double rectangle/diamond = weak) vs crow's foot (`||`, `o|`, `o{`, `|{`), which mermaid draws.
- **Relationships:** degree, cardinality, participation are three separate questions; (min,max) = (participation, cardinality).
- **Relational model (Codd 1970):** relations of atomic values, links by value, declarative queries; SQL tables allow duplicates unless keyed; `NULL` needs `IS NULL`.
- **Keys:** superkey ⊇ candidate (minimal) ⊇ primary; alternates get `UNIQUE`; prime = in any candidate key.
- **Integrity:** domain, entity, key, referential (+ ON DELETE choices); UNIQUE usually allows many NULLs.
- **Mapping:** composite → parts, multivalued → table, derived → view, 1:N → FK on many, M:N → junction, 1:1 → FK + UNIQUE, weak → owner key + partial key + CASCADE, total → NOT NULL. Then normalize — Chapter 11.

Next: Chapter 11 — Normalization, which checks a design for redundancy and splits tables safely.