Module 21 — Case Studies: Developer and AI Platforms

Module Introduction

Four more end-to-end designs in the format of Modules 13–17 and 20, for platform systems common in modern interviews. The algorithms they rely on are explained step by step in Module 19.

Case	Centre of gravity
1 — Online Judge (LeetCode-like)	Running untrusted code safely and fairly: sandboxes, resource limits, judge queues, contest leaderboards
2 — Global Deployment System	Getting one build to many servers safely: artifact distribution, rollout waves, health gates, rollback
3 — Conversational AI Service (ChatGPT-like)	Token streaming, GPU batching, KV-cache-aware routing, quotas and cost
4 — Distributed Cache Cluster	Hash slots, replication and failover, hot keys, routing, eviction and resharding

Case 1 — Online Judge (LeetCode-like)

The prompt (in short). A LeetCode/Codeforces/HackerRank-style site: users write code in one of ~20 languages, press Submit, and within seconds see "Accepted", "Wrong Answer", "Time Limit Exceeded" and so on. Several times a week a timed **contest** (e.g. 90 minutes, same problems for everyone) runs with a live leaderboard.

What makes it hard. (1) It **runs code written by strangers**: any submission may try to read others' code, open connections, fork a million processes or fill the disk, so every submission is assumed hostile. (2) Load is **extremely spiky**: tens of thousands crowd into the same 90 minutes, bunching at the start and end. (3) Results must be **fair and repeatable**: same code, same verdict and roughly the same time on any machine, or the leaderboard is meaningless.

1. Requirements and numbers

Functional: browse problems; submit in a chosen language; compile, run against hidden tests and return a **verdict** (Accepted, Wrong Answer, Time Limit Exceeded, Memory Limit Exceeded, Runtime Error, Compile Error) with runtime and memory; a "Run" button on the user's own sample input; submission history; contests with registration, start/end, live leaderboard and final ratings; post-contest plagiarism checks.

Non-functional: security first (no sandbox escape, no other users' data, no network); verdict **p95 < 5 s** outside contests and **< 15 s** at contest peak; **determinism** (same verdict, time within a few percent); **no lost submissions** (after "Pending", a verdict always appears); fairness (contest submissions judged in arrival order, never starved by practice traffic).

Numbers (assumptions, rounded):

- **20M registered, 2M daily active**; **1M practice submissions/day ≈ 12/s**, ~40/s at the daily peak.
- **Contest:** 50K participants × 4 problems × ~2.5 attempts = **500K submissions in 90 min ≈ 93/s**; first and last 10 minutes ~5× → **~500/s peak**.
- **One submission:** ~1 s CPU compile + ~50 tests × ~60 ms ≈ **~4 CPU-seconds**. Peak 500/s × 4 = **2,000 cores busy**; one dedicated core per sandbox (Deep dive 2) and 30% headroom → **~2,600 cores ≈ 80 32-core workers**. Off-contest 40/s × 4 = 160 cores, so the fleet **scales down ~10×** between contests.
- **Code:** 3 KB average → **3–5 GB/day ≈ 1.5 TB/year**: tiny.
- **Test data:** 3,000 problems × 50 tests × ~1 MB (some 100 MB) ≈ **150 GB**; any worker may need any problem, but a contest touches 4 problems (~200 MB), which fits in RAM.
- **Leaderboard:** participants plus spectators refreshing every 10–30 s → **~5K reads/s** of a ranking changing ~100 times/s.

What the numbers say: storage is trivial; **CPU for untrusted execution**, the **10× contest spike** and the **isolation boundary** decide the architecture.

2. The five-questions pass

- **Flow: asynchronous**: the API stores the code, enqueues a job and returns a `submission_id` at once; the browser waits via Server-Sent Events or polling (Message Queue, Server-Sent Events). No HTTP connection is held open while a compiler runs.
- **Storage:** the **submission record** (code + final verdict) in a relational database is the truth; test cases are immutable versioned files in object storage; the leaderboard is **derived** and rebuildable (CQRS).
- **Speed:** execution dominates, so keep compilers warm in pre-built images, cache test data on local disk, and exit at the first failing test.
- **Failure:** workers crash, hang or are killed by user code. Jobs are **leased**, not consumed (unreported jobs return to the queue), and judging is **idempotent** (Idempotency).
- **Growth:** with contest size, not users; pre-scale before every scheduled contest (Auto-Scaling).
- *Skeleton:* stateless submit API → durable priority queue → judge workers running each submission in a locked-down sandbox → verdict store → result push + leaderboard service.

3. API

Method and path	Purpose
`POST /submissions {problem_id, language, code, contest_id?}`	Returns `202 Accepted {submission_id}` at once. Requires an `Idempotency-Key` header so a double-click does not judge twice
`GET /submissions/{id}`	Status (`queued`, `compiling`, `running 17/50`, final verdict), time and memory
`GET /submissions/{id}/events` (SSE)	Server pushes status changes until the final verdict
`POST /run {problem_id, language, code, stdin}`	"Run code" on custom input. Same sandbox, lower priority queue, no verdict stored
`GET /contests/{id}/leaderboard?cursor=…&limit=50`	Paged ranking (Cursor-Based Pagination)
`GET /contests/{id}/leaderboard/me`	My rank and the rows around me
Internal: `lease(job)`, `report(submission_id, attempt, verdict)`	Worker ↔ queue and worker ↔ verdict store

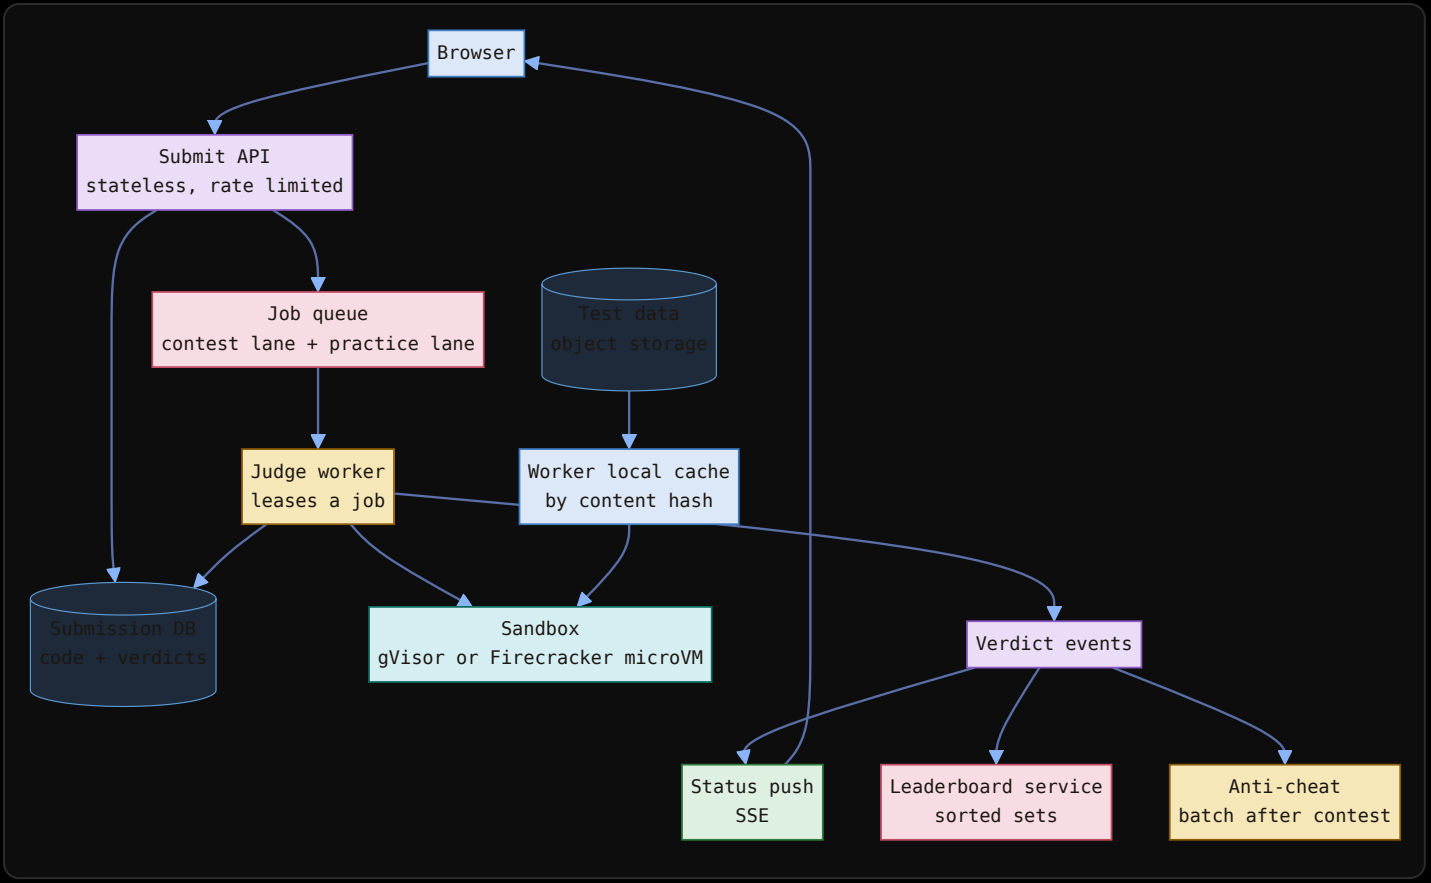
Rate limits, e.g. 1 submission per 5 s per user per problem plus a `/run` cap, stop scripts flooding the judge.

4. Data model

- **`problems`** — `problem_id, title, statement_ref, time_limit_ms, memory_limit_mb, checker_type (exact / token / float tolerance / custom), testset_version`. The **checker** decides whether output is correct: "exact" compares bytes, "float tolerance" accepts 1e-6 differences, "custom" handles problems with many valid answers.
- **`testsets`** — `(problem_id, version) → manifest` of each test file's name, size and **SHA-256**. Files sit at content-addressed paths (the path *is* the hash), so a cached file can never be stale: a new version gets new paths.
- **`submissions`** — `submission_id (time-ordered 64-bit), user_id, problem_id, contest_id?, language, code_ref, created_at, status, verdict, time_ms, memory_kb, failed_test_index, judge_attempt, testset_version`; partitioned by `user_id` (Sharding), secondary index `(contest_id, created_at)`.
- **`contest_scores`** (in-memory sorted set + durable copy) — per contest `user_id → (solved_count, penalty_minutes, last_ac_time)` plus per-problem `(attempts, first_ac_time)`.
- **`language_images`** — `language → container image digest, compile command, run command, time multiplier`. The image is a frozen file system with compiler and runtime; its **digest** (content hash) guarantees every

worker runs byte-for-byte the same toolchain.

5. The architecture, annotated



1. The **Submit API** checks contest registration and timing, writes the code to the **submission DB** as `queued`, enqueues `{submission_id}` on the **contest lane**, and returns `202` in ~50 ms.
2. **Separate lanes** (priority classes) keep 50K contest users from waiting behind practice traffic without starving practice: workers take, say, 9 contest jobs per practice job (Bulkhead); FIFO within a lane.
3. A **judge worker** leases the job (invisible to others for, say, 120 s), loads the code, picks the language image, and fetches any missing test files by hash into its **local cache**.
4. It compiles in one sandbox with a compile-time limit, then runs each test in a fresh locked-down sandbox with CPU, memory, process and output limits (Deep dives 1 and 2); the **checker** compares output, stopping at the first failure.
5. A **conditional update** (`only if status != final`) stores the verdict, then a **verdict event** goes to the status-push service (to the browser) and the **leaderboard service** (Event-Driven Architecture).
6. After the contest, **anti-cheat** jobs compare all accepted submissions (Deep dive 4); ratings publish only after that pass.

6. Deep dives

Deep dive 1 — the sandbox: containers, gVisor or Firecracker. The sandbox is the product; a plain process is unsafe, so the question is how thick the wall is.

- **Plain Linux container** (Docker-style): **namespaces** (a private view of process IDs, mounts, network and users), **cgroups** (control groups: kernel caps on CPU, memory and process count), **seccomp** (a filter on which **system calls**, a program's requests to the kernel such as "open file" or "create socket", are allowed) and all privileges dropped. The IOI's open-source **isolate** sandbox does this, with no measurable overhead per its authors, and handles multi-threaded programs. *Weakness:* the code still talks to the **real host kernel**; one reachable kernel bug is a full escape to a worker holding other people's code.

- **gVisor** (Google): a user-space "application kernel" whose **Sentry** intercepts system calls and implements them itself, with file access via a separate **Gofer** process; plugs in as an OCI runtime (`runsc`) behind Docker or Kubernetes; used under App Engine, Cloud Run and Cloud Functions. *Cost:* syscall-heavy programs slow down, skewing time limits on I/O-heavy problems.
- **Firecracker microVM** (AWS, built for Lambda): a tiny **virtual machine** (guest OS on emulated hardware, isolated by CPU virtualisation) booting in under 125 ms with under 5 MiB overhead, up to 150 created per second per host. The guest has its own kernel, so a kernel bug inside doesn't reach the host.
- **Choice:** a **Firecracker microVM (or gVisor) as the outer wall**, **isolate-style cgroups + seccomp inside** for precise per-test limits: two independent layers must both fail. Keep a **pool of pre-booted microVMs** per language so boot time is off the critical path, and destroy each VM after one submission (never reuse a sandbox between users).
- **Network:** none, no interface at all. **File system:** read-only image plus a size-capped `tmpfs` (memory-backed scratch). **Workers** sit in an isolated segment reaching only the queue, the DB writer and object storage.

Deep dive 2 — resource limits and fair timing.

- **Time: CPU time** (actually running on a core) from cgroup counters decides TLE (Time Limit Exceeded); a **wall-clock limit** (real elapsed time, e.g. 3× the CPU limit) kills programs that sleep or block on input.
- **Memory:** a cgroup cap; the out-of-memory killer firing inside the cgroup means MLE.
- **Processes:** `pids.max` of, say, 64 stops a **fork bomb** (a program copying itself endlessly).
- **Output:** cap stdout at e.g. 64 MB; endless printing gets "Output Limit Exceeded".
- **Fairness:** pin each sandbox to **one dedicated physical core** (no hyper-threading sibling), disable CPU frequency boost, never run more sandboxes than cores, calibrate each worker class with a fixed benchmark and reject drifting machines, and automatically **re-run** borderline TLEs (within ~5% of the limit) on another worker, keeping the better time.
- **Language fairness:** a published **time multiplier** per language in `language_images` (e.g. Python 3× the C++ limit).

Deep dive 3 — the queue at contest time.

- If workers run short at 500/s the queue and the "Pending" spinner grow; acceptable if bounded, and showing queue position ("your submission is #1,240") turns anxiety into patience.
- **Pre-scale** 30 minutes before start from registrations (registrations × expected submissions × 4 CPU-s) and hold 30 minutes after the end for rejudges: the spike arrives in seconds, faster than reactive scaling.
- **Lease and retry:** a dead worker's job is redelivered at lease expiry; each judging has an `attempt` number and the conditional final write means two workers can't both win (visibility timeouts: Module 13 Case 4).
- **Poison submissions** that crash the *worker* (not just the sandbox) three times go to a dead letter queue with "System Error — under review".
- **Rejudge:** wrong tests mid-contest → admins publish a new `testset_version` and enqueue affected submissions on a **rejudge lane** below live submissions.

Deep dive 4 — leaderboard and anti-cheat.

- **Ranking rule (ICPC-style):** problems solved (more is better), then **penalty time** (less is better): per solved problem, minutes from start to first accept plus a fixed penalty per earlier rejected attempt (ICPC: 20 minutes; LeetCode and others publish their own). Encode as one **sorted-set** score (e.g. Redis ZSET): `score = solved × 10^7 − penalty_minutes × 10 − tiebreak`, so one `ZADD` per accepted verdict keeps order and "rank of X" or "rows 1–50" are O(log N) (Skip lists; sorted sets in depth: Module 15 Design 4).

- **Only verdict events change scores**, idempotent per `(user, problem)` (a second Accepted does nothing); `contest_scores` rebuilds the sorted set after cache loss.
- **Freeze:** many contests hide the last hour's changes; the service keeps a public (frozen) view and the true view, serving the public one until reveal.
- **Anti-cheat signals:** (1) **code similarity** with MOSS-style **winnowing**: strip comments, whitespace and identifiers, cut the token stream into overlapping k-grams, hash them, keep a sample of fingerprints, flag pairs sharing many; an inverted index (fingerprint → submissions) avoids comparing all 25K² pairs on a popular problem. (2) **Behaviour:** many accounts per device or IP, a hard problem solved 40 s after a friend, sudden rating jumps, large pastes (the editor can record paste events). (3) **Leaked-solution matching** against code posted publicly during the contest.
- Flags go to **human review**, never straight to a ban: short standard algorithms look alike and fingerprint tools over-flag them. Ratings finalise after review and roll back for confirmed cheaters.

7. The failure walk

1. **Worker dies mid-judging:** leases expire, jobs are redelivered, the conditional write keeps one result; users wait ~2 minutes longer.
2. **Sandbox escape discovered** (kernel or gVisor CVE): the worst case, but one failed layer isn't yet a breach. Workers are **cattle** (identical, disposable): rebuild every image with the patch, drain and replace the fleet, rotate credentials. Workers hold only the current job and can write only verdicts, limiting blast radius.
3. **Test-data store slow or down:** workers use their hash-addressed cache, and contest test sets are **pre-warmed** to every worker before start.
4. **Queue backlog at peak:** show queue position, **pause the practice lane** and `/run` (Graceful Degradation), extend the contest past a published delay threshold; never drop a queued contest submission.
5. **Leaderboard sorted set lost:** rebuild from `contest_scores` (50K rows) in seconds; verdicts live in the DB.
6. **Bad compiler image** (every C++ submission gets Compile Error): pin images by digest per contest, canary new images on a worker slice outside contest hours, watch the verdict mix per language (Canary Deployment), rejudge affected submissions.
7. **A program hogs I/O or memory bandwidth**, slowing neighbours: cgroup I/O limits, one core per sandbox and borderline re-runs keep verdicts fair.

8. Trade-offs

Decision	Chosen	Alternative	Why
Isolation	MicroVM or gVisor outside + cgroups/seccomp inside	Plain containers only	Code is hostile by definition; one shared-kernel bug should not be game over
Submission flow	Async queue + SSE push	Synchronous HTTP	Judging takes seconds and spikes 10×; a queue absorbs the burst
Queue	Separate contest, practice, rejudge lanes	One FIFO	Contest fairness and protection from bulk rejudges
Timing	CPU time on a dedicated core, re-run if borderline	Wall-clock on shared cores	Repeatable verdicts are the product
Test data	Content-addressed files + local cache, pre-warmed	Read from shared storage per run	Hundreds of reads/s of the same files; no stale caches
Leaderboard	Sorted set derived from verdict events	SQL `ORDER BY` per request	5K reads/s on a constantly changing ranking
Anti-cheat	Offline fingerprinting + human review	Real-time automatic bans	False positives on short standard code are common
Fleet	Scheduled pre-scaling around contests	Reactive autoscaling only	The spike arrives faster than machines boot

9. The follow-up questions

- **"Why not just Docker?"** Docker shares the host kernel, so isolation depends on every allowed system call being bug-free; hostile code needs a second boundary with its own kernel (microVM) or syscall implementation (gVisor). *Trap:* "containers are secure because of namespaces" without mentioning the shared kernel.
- **"Interactive problems** (the program converses with a judge program)?" Two sandboxes joined by pipes, the interactor deciding the verdict; time both, and a deadlock becomes TLE via the wall-clock limit.
- **"New language?"** Build an image, run every problem's reference solutions to set its time multiplier, publish the digest, enable it; workers are image-driven, so nothing else changes.
- **"Serverless judging?"** Fine for practice (each Lambda invocation already runs in a Firecracker microVM), but timing noise and cold starts make contest timing less fair; dedicated calibrated workers are safer for rated contests.
- **"AI-cheating detection?"** One more advisory signal: compare against popular code models' outputs for the problem plus behaviour (one large paste, no edits), with human review; prefer resistant contest design (proctored or in-person finals).

Pattern manifest

Pattern	Where used	Taught in
Message Queue (with leases)	Submission jobs, contest/practice/rejudge lanes	Module 02
Server-Sent Events	Pushing status and verdict to the browser	Module 02
Idempotency	`Idempotency-Key` on submit; conditional verdict write; one AC per problem	Module 05
Bulkhead	Separate lanes and worker shares for contest vs practice	Module 05
Dead Letter Queue	Submissions that crash workers repeatedly	Module 05
Graceful Degradation	Pause practice and `/run` at contest peak	Module 05
Auto-Scaling (scheduled)	Worker fleet before each contest	Module 06
Rate Limiting	Per-user submit and run caps	Module 08
Containers and sandboxing	Language images, microVM/gVisor sandboxes	Module 09
CQRS	Leaderboard as a derived read model	Module 03
Sorted-set leaderboard	Contest ranking	Module 15

Interview one-liner: "Submit is asynchronous: store the code, enqueue a job, return an ID. Calibrated workers lease jobs and run each submission in a throwaway microVM or gVisor sandbox with no network, cgroup limits on CPU, memory, processes and output, and content-hashed test data cached locally. Verdict events feed the browser and a sorted-set leaderboard; contests get their own queue lane and pre-scaled capacity; plagiarism is a fingerprinting batch job plus human review."

TL;DR

- **Problem:** run hostile code from millions of users fairly and repeatably, with a 10× contest spike.
- **Mechanism:** async submit → leased priority queue → judge workers → per-submission microVM/gVisor sandbox with cgroup and seccomp limits → verdict events → SSE + sorted-set leaderboard.
- **Hard parts:** isolation (two layers, no network, never reuse), fair timing (dedicated cores, CPU time, borderline re-runs), contest bursts (lanes, pre-scaling, pre-warmed tests), anti-cheat (winnowing + humans).
- **Price:** a large mostly idle fleet between contests and constant sandbox patching.

Sources

- Mareš, M. and Blackham, B., "A New Contest Sandbox", *Olympiads in Informatics* 6 (2012): https://ioinformatics.org/journal/INFOL094.pdf and the isolate project: https://github.com/ioi/isolate

- gVisor documentation (Sentry, Gofer, runsc): https://gvisor.dev/docs/
- Google Cloud blog, "How gVisor protects Google Cloud services from CVE-2020-14386": https://cloud.google.com/blog/products/containers-kubernetes/how-gvisor-protects-google-cloud-services-from-cve-2020-14386
- Firecracker project site (boot time, overhead, creation rate): https://firecracker-microvm.github.io/
- ICPC scoring rules (20-minute penalty per rejected run on a solved problem), e.g. ECNA regional rules: https://icpc-ecna.ysu.edu/rules.html
- Schleimer, Wilkerson and Aiken, "Winnowing: Local Algorithms for Document Fingerprinting", SIGMOD 2003 (the algorithm behind MOSS): https://theory.stanford.edu/~aiken/publications/papers/sigmod03.pdf

---

Case 2 — Global Deployment System

The prompt (in short). The internal system a large company ships software with: an engineer merges, and within hours the build runs safely on hundreds of thousands of servers in dozens of regions (think Facebook's continuous web-tier push, Amazon's pipelines to every AWS region, Google's rollout tooling).

What makes it hard. Google's SRE book observes that roughly 70% of outages come from changes to a live system, so deployment is both the biggest risk and the main defence. Hard parts: (1) **moving gigabytes to 200K machines** without melting the network or artifact store; (2) **gradual exposure** so a bad build hurts 0.1% of users, not 100%; (3) **automatically judging** stage health and rolling back without a human; (4) treating **configuration changes**, which spread in seconds and cause many outages, as carefully as code.

1. Requirements and numbers

Functional: build a versioned immutable artifact per commit; store and distribute it; roll out in staged waves across regions; check health per stage and promote or roll back automatically; a big red **stop button**; deploy **config** separately from code; audit who shipped what, where, when; deployment windows and blockers (no deploys during an active incident).

Non-functional: safety beats speed (catch bad changes while exposure is small); **rollback in minutes** anywhere; usable during an outage (it ships the fix); fully **automated and repeatable**; thousands of services, tens of deploys per service per week.

Numbers (assumptions):

- **200K servers** in **20 regions** (~10K each over 3 **availability zones**, AZ = one or more data centres with independent power and network inside a region).
- **2,000 services** deploying about daily → **~2,000 rollouts/day** plus **~5,000 config changes/day**.
- **Artifacts:** a big monolith binary or image **~1 GB**; a typical microservice image ~200 MB, with only changed layers moving (often ~50 MB).
- **Naive fan-out** of 1 GB to 200K hosts = **200 TB**; a central store of 20 servers × 10 Gbps = 25 GB/s → 8,000 s ≈ **2.2 hours** with the store saturated and other deploys blocked. **Peer-to-peer** (Deep dive 1) makes every host that holds a piece upload it, so capacity grows with hosts; at ~50% of a 10 Gbps NIC, time is about one host's download (~2 s per GB) plus a few relay hops: **well under a minute** (Uber reports 20K blobs of 100 MB–1 GB in under 30 s).
- **Rollout duration** is dominated by **bake time** (running under real traffic before promotion): five waves with 1–12 h bakes → **multi-day** global rollouts for risky backends (Amazon's published zonal-and-regional wave pipeline takes days), a few hours for front-ends with fast signals (Facebook's web tier pushes several times a day).

What the numbers say: P2P solves the bytes; the design is about **exposure control and fast, automatic judgement**.

2. The five-questions pass

- **Flow:** asynchronous and long-running: a rollout is a durable **workflow** run by a pipeline orchestrator, not a laptop script. Hosts **pull** their target from a desired-state store; the orchestrator only changes desired state (the reconcile loop of Containers and Kubernetes).
- **Storage:** the **artifact** (immutable, content-addressed: named by its bytes' hash) in blob storage; **desired state** per (service, cell), e.g. "cell 7 of region X runs build b-4812 with config v-991", in a strongly consistent etcd/ZooKeeper-style store; history and audit in a relational DB.
- **Speed:** P2P plus pre-staging (download before switching) make the switch a seconds-long restart; the long part is deliberate waiting.
- **Failure:** the main failure is **a bad change**; health gates, automatic rollback and small blast radii defend. The deploy system's own outage must never block an emergency fix (a minimal "break glass" path).
- **Growth:** services × hosts × change rate; distribution must be sub-linear in hosts (P2P) and control hierarchical (region → cell → host).
- *Skeleton:* CI build → artifact registry → P2P distribution → per-host agent → rollout orchestrator with waves and gates → health/canary analysis → automatic rollback; a separate fast lane for config.

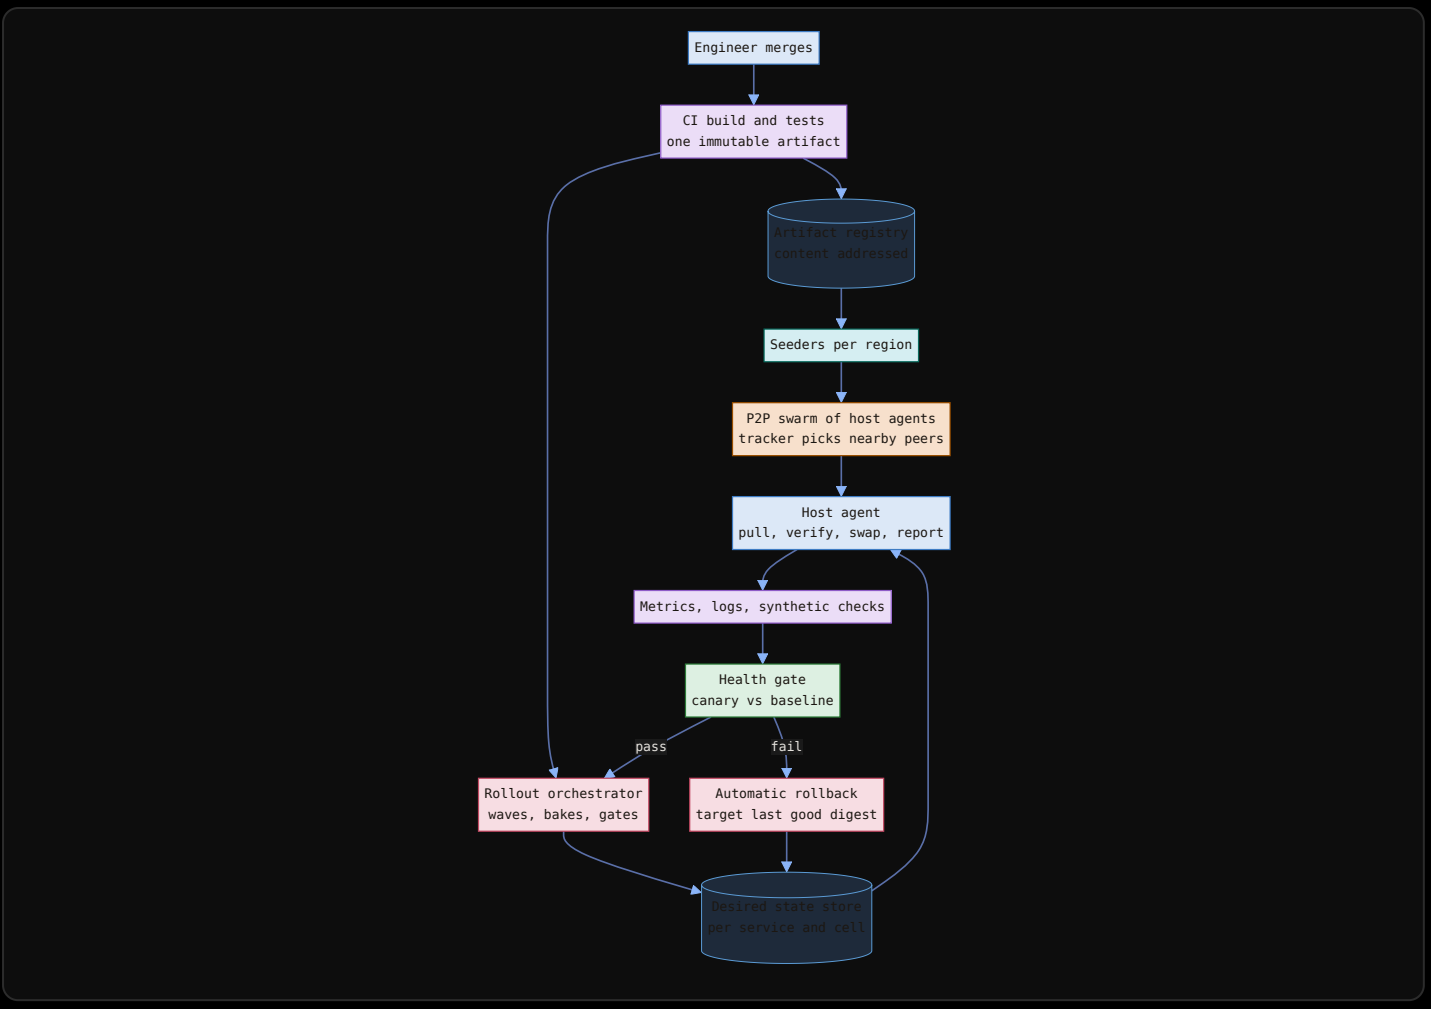
3. API

Method and path	Purpose
`POST /artifacts {service, commit, digest, sbom_ref, signature}`	Register a build (CI calls this). **Immutable**: same digest, same bytes, forever
`POST /rollouts {service, artifact_digest, plan_id}`	Start a rollout using a named plan (waves, bakes, gates). Returns `rollout_id`
`GET /rollouts/{id}`	Current wave, per-cell state, gate results, ETA
`POST /rollouts/{id}:pause` / `:resume` / `:rollback`	Human controls; rollback re-targets the last good digest
`POST /config-changes {key, new_value, canary_spec}`	Propose a config change; goes through review, validation and its own canary
`POST /freeze {scope, reason, until}`	Org-wide or service-level deploy blocker during incidents and holidays
Agent ↔ control plane: `GET /desired?host=…` (long poll), `POST /report {host, running_digest, health}`	Pull model; the agent reports what is actually running

4. Data model

- **`artifacts`** — `digest (SHA-256) → service, commit, build_time, size, layers[], signature, provenance` (who built it, from what source, on which builder); bytes stored as layers/chunks keyed by their own hashes, so identical layers are stored and sent once.
- **`rollout_plans`** — `plan_id → waves[]`, each `{targets (cells/regions), batch_size %, bake_time, gates[]}`; plans are reviewed like code.
- **`rollouts`** — `rollout_id, service, from_digest, to_digest, plan_id, state (running / paused / rolling_back / done), current_wave, started_by, started_at`.
- **`desired_state`** — `(service, cell) → {artifact_digest, config_version, generation}`; the increasing **generation** stops an agent applying an older instruction after a newer one (ordered durable changes, Write-Ahead Log thinking).
- **`actual_state`** — `(host) → running_digest, config_version, healthy, last_report_at`, expiring if the host stops reporting.
- **`gate_results`** — `(rollout_id, wave, gate) → pass / fail / inconclusive, metrics snapshot, decided_at`.
- **`config_store`** — versioned key → value with schema, owner and history (in Git, mirrored to a distribution tree).

5. The architecture, annotated



1. The merge triggers **CI**, which produces **one** signed artifact with a digest and recorded provenance. **Build once, promote everywhere**: the bytes that passed tests go to every stage (CI/CD Pipelines and Deployment Strategies).
2. The artifact goes to the **registry**; the **orchestrator** starts the service's plan: pre-production stages, then production **wave 1**.
3. Per target cell, agents first **pre-stage** (download and verify); regional **seeders** hold the first copy and the **P2P swarm** spreads it (Deep dive 1).
4. The orchestrator bumps the wave's **desired state**. Each **host agent** sees the new generation, verifies signature and hash, **drains** the local service (no new requests, finish in-flight ones), swaps versions, runs local health checks and reports `actual_state` (Health Check Endpoint); hosts switch in batches (e.g. ≤ 10%) so capacity holds.
5. The **health gate** compares new hosts with old over the **bake time** (Deep dive 3): pass → next wave; fail → **automatic rollback** to the last good digest, still cached on every host, so rollback is a restart, not a download.

6. Deep dives

Deep dive 1 — moving the bytes: registry, P2P and pre-staging.

- **Content addressing:** chunks or layers named by hash; hosts skip chunks they hold, so a build changing 3% of a 1 GB binary moves ~30 MB, and the recomputed hash rejects corrupt or tampered chunks.
- **Why P2P:** a central server's fixed upload capacity makes time linear in hosts. In a **BitTorrent-style swarm** the file is cut into **pieces** (BitTorrent's spec: most commonly 256 KB, each with its own SHA-1 hash), each host downloads from several peers and **uploads pieces it has**, so capacity grows with the swarm and 200K hosts take about as long as 2,000.

- **History:** Facebook pushed its front-end binary with internal BitTorrent around 2010–2012, with a custom tracker preferring same-rack peers; Twitter's "Murder" did the same; Uber's open-source **Kraken** is a P2P Docker registry (a **tracker** says which peers have which pieces; **origin** servers seed from backing storage such as S3), reporting 20K blobs of 100 MB–1 GB in under 30 s at peak at > 50% of each host's max download speed.
- **Topology awareness:** the tracker hands out peers in the **same rack, then cluster, then region**, so most bytes stay off expensive links; per-region seeders mean one copy crosses an ocean.
- **Pre-staging:** separates the slow, failure-prone network part from the risky restart, so a 10K-host switch takes seconds; keep the **previous N versions** on disk so rollback needs no network.
- **Guard rails:** cap per-host P2P upload (e.g. 30% of NIC) so production traffic isn't starved; heavily loaded hosts opt out of seeding.

Deep dive 2 — staged rollout waves and blast radius.

- **Blast radius** = share of users hurt if this step is bad; every wave caps it.
- **Typical plan** (the shape Amazon publishes):
  1. Pre-production: integration tests, then a **gamma** stage mirroring production config.
  2. **One-box** in one AZ of a low-traffic region (one host or a small traffic slice), bake ≥ 1 h.
  3. Wave 1: the rest of that AZ, then that region's other AZs **one at a time**; bake ~12 h after the first full region to catch daily-cycle bugs (overnight batch jobs, memory leaks, certificate rotations).
  4. Wave 2: one high-traffic region, AZ by AZ.
  5. Waves 3–5: 3, then ~12, then all remaining regions, with 2–4 h bakes.
- **Never two AZs of a region at once**, never two regions at once early: a failure stays inside a **cell** (a self-contained stack serving a slice of users) (Multi-Region and Cell-Based Architecture).
- **Front-end variant:** Facebook's web tier goes to **employees**, then **~2% of production**, then 100%, over a few hours, several times a day; enormous traffic gives strong signals fast.
- **Windows and blockers:** production waves start only in the owning team's working hours (not nights, weekends, late Fridays); **blockers** halt new waves automatically during related high-severity alarms or org-wide incidents.
- **Within a cell: rolling** batches (≤ 10–33% of hosts) or **blue-green** where instant switch-back matters (Blue-Green Deployment).

Deep dive 3 — health gates and automatic rollback.

- **Watched signals:** 5xx rate, latency p50/p90/p99, CPU, memory, crash/restart count, log error rate, failed health checks, **business metrics** (checkouts per minute, messages sent), and **synthetic canaries** (scripted fake users calling the API every minute from outside).
- **Compare like with like:** absolute thresholds miss subtle regressions and misfire at traffic peaks, so compare **canary vs baseline hosts** (old version, same cell, window and host count). Netflix and Google's open-source **Kayenta** runs a **Mann-Whitney U test** per metric (do two sets of numbers come from different distributions, without assuming a bell curve), flagging only at ~98% confidence of a meaningful difference, and rolls results into a pass/fail/ask-a-human score (Canary Deployment).
- **Outcomes:** `pass` → promote after the bake; `fail` → rollback; `inconclusive` (too little traffic) → extend the bake, never promote silently.
- **Fast tripwires** roll back immediately: request count to zero, fault-rate spike, health checks failing on > N% of new hosts.
- **Automatic rollback:** desired state back to the **last digest that passed all gates** in the cell; agents restart into the cached version. On a team-level aggregate alarm, roll back **every in-flight deployment** of that team in the region (the culprit is unknown). Rollback needs neither the failing service healthy nor CI to rebuild.

- **Rollback needs backward compatibility:** schema changes go **expand → migrate → contract** (add column, deploy code writing both, backfill, drop the old column later) so N−1 runs against data N wrote (Capstone - The Zero-Downtime Migration); irreversible changes (data deletions, protocol bumps) go behind **feature flags**.

Deep dive 4 — config deploys vs code deploys.

- **Why separate:** code needs a new binary everywhere (hours of waves); many changes are just *values* (a timeout, a rate limit, which 5% see a feature) that as config reach the fleet and undo in **seconds**.
- **The danger is that speed:** a bad value hits every server at once with no rollout to slow it; many well-known global outages were config pushes.
- **Facebook's model** (SOSP 2015, "Holistic Configuration Management at Facebook"): configs in Git, often **generated by code** with validators, every change reviewed; an automated **canary spec** applies it to ~20 servers, then a cluster, comparing health (~10 minutes for a full canary); distribution via **Zeus** (a ZooKeeper fork) through a high fan-out tree of observers and per-server proxies, reaching hundreds of thousands of servers in seconds; large blobs via a separate P2P path. Each proxy's **on-disk cache** lets apps read the last good config if the config system is down.
- **Design rules:** schema-validate every value (types, ranges); **stage config like code** (one-box → cell → region, shorter bakes); full history with **one-click, seconds-long revert**; apps **reject unparseable config and keep the last good one**; separate **feature flags** (per-user exposure, Feature Flags) from **operational config** (timeouts, pool sizes) so each gets the right review rules.

7. The failure walk

1. **Bad build passes tests** (a leak showing after 6 h): one-box and wave-1 bakes catch daily-cycle problems; the memory alarm trips on the one-box and rollback restores 1 host: ~0.01% of traffic.
2. **Bad build fails fast** (crash on start): hosts fail local health checks before joining the load balancer and the agent stops after the first batch: zero user impact.
3. **Gates are blind** (an unwatched feature breaks): it reaches more waves. Add business-metric gates and synthetic tests per critical journey; every escaped bug becomes a new gate.
4. **Orchestrator or desired-state store down:** hosts keep running their last known state (pull model) and no rollouts start, which is safe; a documented **break-glass** tool lets on-call push a pinned digest to a cell.
5. **Registry or seeders down:** pre-staged versions and peer copies continue; new builds wait; the registry is replicated per region (Primary-Replica).
6. **Rollback fails** because the new version wrote data the old can't read: hence backward-compatible schema rules and flags; otherwise **roll forward** a hotfix through an expedited plan (shorter bakes, still staged, never all at once).
7. **Bad config push:** the config canary catches most; escaped ones revert in seconds; apps keep the previous value on unparseable config.
8. **Compromised build machine** shipping malware: hosts accept only artifacts signed by the build system with provenance from reviewed source, verifying signature and digest before running.

8. Trade-offs

Decision	Chosen	Alternative	Why
Distribution	P2P swarm with regional seeders, rack-aware peers	Central servers / CDN only	Total capacity grows with hosts; 200 TB per release is otherwise hours
Control	Desired state + host agents pull	Orchestrator pushes over SSH	Survives orchestrator outages, scales, self-heals drift
Exposure	One-box → AZ → region waves with bakes	Deploy everywhere at once	Caps blast radius; 70% of outages are changes
Gate logic	Canary vs baseline statistics + hard tripwires	Static thresholds	Adapts to traffic patterns, catches subtle regressions
Rollback	Automatic to last good, cached digest	Human decides	Machines notice in minutes, humans at 3 a.m. take longer
Speed	Days for backends, hours for web tier	Same plan for all	Match bake time to how fast signals arrive and how bad failure is
Config	Separate fast lane, but still canaried and staged	Treat config as code deploy	Seconds-level changes and undo, without global instant blast

9. The follow-up questions

- **"Why not a CDN for artifacts?"** It helps seed the first copy per region, but inside a data centre all hosts pulling from one edge concentrate load; P2P uses every host's idle upload and keeps bytes in racks. Many systems combine them: CDN/object storage to seed, P2P to spread (CDN).
- **"Deploying a service 300 others depend on?"** Same waves, with **callers'** error rates in the gates and cell-aligned rollouts so callers and dependency change together; CI contract tests keep the API compatible (API Versioning).
- **"Mobile apps?"** No rollback on phones: staged store rollouts (a percentage of users), server-driven feature flags to switch behaviour off, server compatibility with old app versions for months.
- **"Did a deploy cause this incident?"** Dashboards overlay deploy and config events, the deploy system publishes to the incident tool, and triage starts with "what changed in the last hour, where?".
- **"2,000 rollouts a day without collisions?"** Per-cell locking (one rollout per service per cell), a per-service rollout queue that **batches** commits into one artifact (Facebook's push-from-master ships many diffs together), and blockers pausing a region when it is unhealthy.

Pattern manifest

Pattern	Where used	Taught in
CI/CD, build once, promote	One signed artifact through every stage	Module 09
Canary Deployment	One-box and wave gates, canary vs baseline	Module 09
Blue-Green Deployment	Instant switch-back within a cell	Module 09
Feature Flags	Decouple risky behaviour from deploys; non-reversible changes	Module 09
Health Check Endpoint	Agent's local check before joining the load balancer	Module 09
Containers and reconcile loops	Desired vs actual state, agents pull	Module 09
Multi-Region and Cell-Based Architecture	Waves follow AZ and cell boundaries	Module 06
Circuit-breaker-style blockers	Halt all deploys during alarms or incidents	Module 05
Publish-Subscribe / fan-out tree	Config distribution through observers and proxies	Module 02
Zero-downtime migration	Expand-migrate-contract so rollback stays safe	Module 09
Distributed Job Scheduler	Long-running rollout workflows	Module 16

Interview one-liner: "Build one signed, content-addressed artifact; spread it with a rack-aware P2P swarm and pre-stage it; change only the desired state and let host agents pull; roll out in waves — one-box, one AZ, one region, then wider — with bake times and canary-vs-baseline gates; roll back automatically to the cached last-good digest; and push config through its own fast but still canaried lane."

TL;DR

- **Problem:** ship thousands of changes a day to 200K servers without any one change taking everything down.
- **Mechanism:** CI → immutable artifact → P2P distribution with regional seeders → desired-state store + pulling host agents → staged waves with bakes → statistical health gates → automatic rollback; a separate canaried config pipeline.
- **Hard parts:** bytes at scale (P2P, content addressing, pre-staging), blast radius (cells, waves), automatic health judgement (canary vs baseline), keeping rollback possible (compatible data changes), config's dangerous speed.
- **Price:** slow global rollouts (days for backends), heavy telemetry plumbing, compatibility discipline.

Sources

- Liguori, C., "Automating safe, hands-off deployments", Amazon Builders' Library: https://aws.amazon.com/builders-library/automating-safe-hands-off-deployments/

- Facebook Engineering, "Rapid release at massive scale" (2017): https://engineering.fb.com/2017/08/31/web/rapid-release-at-massive-scale/
- Tang, C. et al., "Holistic Configuration Management at Facebook", SOSP 2015: https://doi.org/10.1145/2815400.2815401
- Uber Engineering, "Introducing Kraken, an Open Source Peer-to-Peer Docker Registry": https://www.uber.com/blog/introducing-kraken/ and https://github.com/uber/kraken
- BitTorrent protocol specification (BEP 3): https://www.bittorrent.org/beps/bep_0003.html
- Spinnaker docs, "How canary judgment works" (Kayenta, Mann-Whitney U): https://spinnaker.io/docs/guides/user/canary/judge
- Google Cloud blog, "Introducing Kayenta": https://cloud.google.com/blog/products/gcp/introducing-kayenta-an-open-automated-canary-analysis-tool-from-google-and-netflix
- Beyer, B. et al., *Site Reliability Engineering* (O'Reilly, 2016), on ~70% of outages coming from changes: https://sre.google/sre-book/introduction/
- TorrentFreak, "Facebook Uses BitTorrent, and They Love It" (2010, reporting a Velocity talk): https://torrentfreak.com/facebook-uses-bittorrent-and-they-love-it-100625/

Case 3 — Conversational AI Service (ChatGPT-like)

The prompt (in short). A ChatGPT/Claude/Gemini-style assistant: users sign in, type a message and watch the answer appear word by word; conversations are saved and reopen on any device; free and paid tiers get different limits; harmful requests and answers are filtered.

What makes it hard. Sign-in and history are ordinary CRUD. The hard parts: (1) every answer comes from a **large language model (LLM)** generating one **token** (a word piece, roughly ¾ of an English word) at a time on **GPUs**, by far the largest cost; (2) answers take **seconds to tens of seconds**, so they must be **streamed**; (3) the GPU memory each in-flight conversation needs (its **KV cache**, Deep dive 2) caps users per GPU, so scheduling and routing are memory problems.

1. Requirements and numbers

Functional: create a conversation; send a message and receive a streamed answer; stop mid-way; regenerate or edit (creating a branch); list, rename, search, delete conversations; choose a model; per-tier quotas (free, paid, team). Attachments only "stored in blob storage".

Non-functional:

- **Time to first token (TTFT)** p95 < 1.5 s.
- **Time per output token (TPOT)** ≈ 20–30 ms (~35–50 tokens/s, faster than reading).
- **No lost answers:** finished answers are saved; a dropped connection can resume the stream.
- **Availability** 99.9% for chat; degrade (smaller model, queue) rather than hard-fail.

Numbers (assumptions, rounded):

- **100M DAU × 10 messages/day = 1B/day ≈ 11.6K/s, ~35K/s peak** (3×).
- Answers average **500 output tokens** at **50 tokens/s** → **~10 s** of GPU time each.
- Concurrent streams at peak by Little's law: 35K/s × 10 s = **~350K open streams** → **17.5M output tokens/s**.
- Input per message (system prompt + history + new message) ≈ **3,000 tokens**; if 60% is a prefix already cached on the GPU (Deep dive 3), 1,200 new tokens each → **42M input tokens/s** of **prefill** (Deep dive 2).
- **KV cache per token**, 70B model shaped like Llama 3.1 70B (80 layers, 8 key/value heads, head size 128, 16-bit): 2 (key and value) × 80 × 8 × 128 × 2 bytes = **327,680 bytes ≈ 320 KB**; a 3,500-token conversation holds **~1.1 GB** while answered.
- **One replica** = 8 H100s × 80 GB = 640 GB; weights 70B × 2 bytes = 140 GB; after activations ~400 GB for KV → **~1.25M tokens** ≈ 350 conversations of 3,500 tokens at once.
- **Replica throughput (assumption, measure it):** ~6,000 output tokens/s at target TPOT → 17.5M ÷ 6,000 ≈ **2,900 replicas ≈ 23,000 GPUs**, +~30% for prefill and headroom → **~30,000 GPUs**.

- **Cost (assumption):** $2.50 per GPU-hour → **$1.8M/day ≈ $0.0018 per message**; a free user at 10 messages/day costs ~$0.55/month in GPU time, which drives quota design.
- **History:** ~2.5 KB per message → 2.5 TB/day ≈ **~0.9 PB/year** before replication, tiny next to the GPU bill.

What the numbers say: database and web tier are ordinary; **GPU memory and time** are the bottleneck and the bill. Judge every choice by "does it raise tokens per GPU-second, or cut tokens we must compute?"

2. The five-questions pass

- **Flow:** synchronous to a watching user but **streamed** token by token over one long HTTP response with Server-Sent Events (SSE: the server writes small `data:` events into an open response); history, billing and analytics are **asynchronous** after the stream (Five Questions).
- **Storage:** the **conversation store** (messages as a tree, partitioned by user) is truth; a short-lived **generation buffer** holds in-flight tokens so a dropped client can resume; the GPU KV cache is a disposable performance cache.
- **Speed:** TTFT = queueing + prefill; TPOT = decode batching. The hot read, "this conversation's messages for the prompt", is one partition read.
- **Failure:** GPU servers fail more than web servers: retry a failed generation **on another replica**; a client disconnect must **not** waste the generation.
- **Growth:** with users and **context length** (KV memory grows linearly); hot spots are launches and each region's working-day start.
- *Skeleton:* API/edge tier (auth, rate limits), a chat orchestrator (prompt building, safety), a **KV-cache-aware inference router**, GPU pools with **continuous batching**, a stream relay to SSE gateways, a conversation store, an async usage pipeline.

3. API

Method and path	Purpose
`POST /v1/conversations` → `{conversation_id}`	Start a conversation
`POST /v1/conversations/{id}/messages` body `{client_msg_id, parent_message_id, content, model}` → `200 text/event-stream`	Send a message and stream the answer. `client_msg_id` is an idempotency key (a client-made ID reused on retries so the server can spot repeats) (Idempotency)
SSE events: `message_start{message_id, generation_id}`, `delta{seq, text}`, `message_end{finish_reason, usage}`, `error{code}`	Each event carries an `id:` line equal to its `seq` (the token chunk's position)
`GET /v1/generations/{generation_id}/stream` with header `Last-Event-ID: 812`	Resume a dropped stream from chunk 813 (SSE clients send `Last-Event-ID` automatically on reconnect)
`POST /v1/generations/{generation_id}/cancel`	The "stop" button; frees the GPU slot
`GET /v1/conversations?cursor=…`	List conversations, newest first (Cursor-Based Pagination)
`GET /v1/conversations/{id}`	All messages (the tree) for rendering
`GET /v1/usage`	Remaining quota for the current window

Rate limits return `429 Too Many Requests` with limit, remaining and reset headers (OpenAI's API returns `x-ratelimit-remaining-tokens` and similar). SSE beats WebSockets here: one-way traffic, plain HTTP that proxies understand, built-in resume via `Last-Event-ID`.

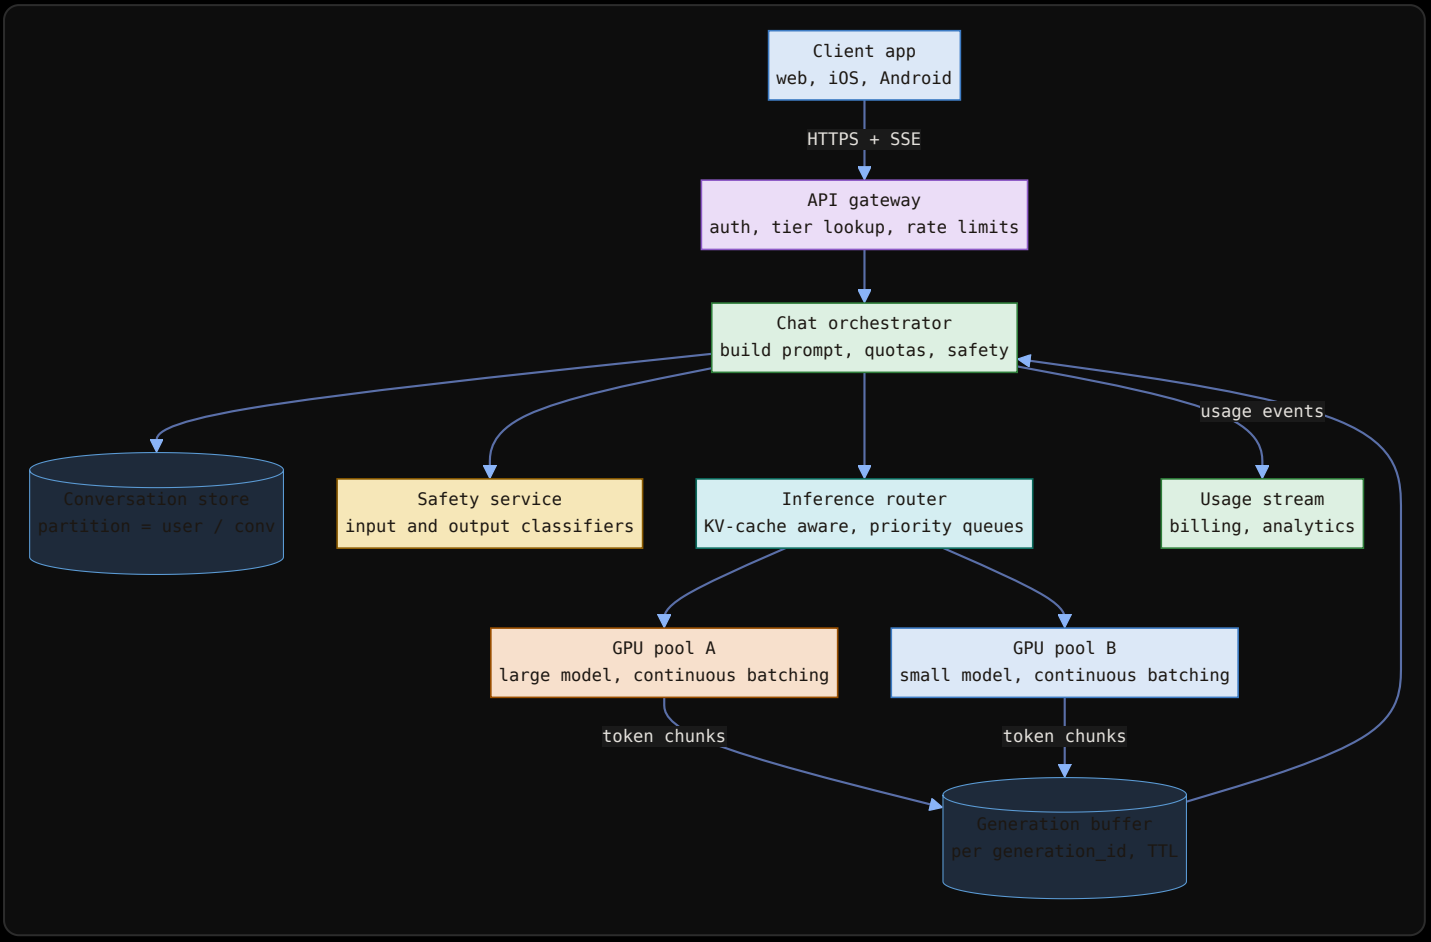
4. Data model

- **`conversations`** — partition `user_id`, clustering `updated_at` descending: `conv_id, title, model, created_at, archived`; listing chats is one range read.

- **`messages`** — partition `conv_id`, clustering `message_id` (time-ordered 64-bit): `parent_id, role (system/user/assistant/tool), content, content_ref (blob for large text or files), model, token_counts {input, cached_input, output}, status (streaming/complete/stopped/failed/blocked), safety_labels, created_at`. **Messages form a tree**: editing or regenerating adds a sibling under the same parent; the UI shows the selected branch, and a new turn's prompt walks from the new message up to the root.
- **`generation_buffer`** (in-memory with TTL, e.g. Redis Streams) — `generation_id` → ordered `(seq, text)` chunks, `status`, `replica_id`, TTL 15 minutes; decouples the GPU from the client connection.
- **`usage_ledger`** — `(user_id, window_start) → input_tokens, cached_tokens, output_tokens, cost_units`, hot counters in memory, settled totals durable. **`tiers`** — `tier → models, tokens per minute, messages per window, max context, priority`.
- **Router state** (in memory, rebuildable) — `prefix_block_hash → set of replicas` plus per-replica load (queue depth, free KV blocks).

A wide-column or key-value store (Cassandra/ScyllaDB/DynamoDB-style) fits both main tables: every query is by partition key, writes are append-mostly, data grows without bound (Sharding).

5. The architecture, annotated



1. `POST /messages` with `client_msg_id`: the **API gateway** checks the session, looks up the **tier** and applies per-user and per-organisation token-bucket request limits (Rate Limiting); the SSE stream opens immediately.
2. The **orchestrator** dedupes on `client_msg_id`, saves the user message, **builds the prompt** (system prompt + branch history, truncated or summarised to fit + new message), counts tokens and checks the **token quota** (Deep dive 4).
3. The **safety service** classifies the input; blocked requests get a refusal and never reach a GPU (Deep dive 5).
4. The **inference router** picks a **pool** by model and tier, then a **replica** by prefix-cache hits and load (Deep dive 3); if all are full the request waits in a **priority queue** (paid before free).
5. The **inference server** joins the request to its running batch (**continuous batching**, Deep dive 2), runs prefill, and appends `(seq, text)` chunks to the **generation buffer**.

6. The orchestrator reads the buffer, runs the **streaming safety check** and writes `delta` events.
7. At the end it saves the assistant message with token counts, settles the quota, and emits a **usage event** (Message Queue).

6. Deep dives

Deep dive 1 — Token streaming: SSE, resume and cancellation.

- *Wire format:* each event is a few lines (`id: 812` , `event: delta` , `data: {"text":" world"}` , blank line), parsed with a `fetch` stream reader (`EventSource` can't send a POST body). One event per token would be 17.5M events/s; ~4 tokens per event cuts that to ~4.4M/s.
- *Proxies:* no hop may **buffer**: disable buffering at load balancers and reverse proxies, send `Cache-Control: no-cache` , use HTTP/2 to multiplex streams over one TCP connection, and send a `: ping` comment every ~15 s while queued so idle timeouts don't fire.
- *Resume:* read the buffer from `Last-Event-ID + 1` , then follow the live tail ("cursor + replay + live" for one answer). Without the buffer, a phone switching from Wi-Fi to 5G kills an 80%-done generation.
- *Back-pressure:* the GPU never waits on a client; the gateway reads the buffer at the client's pace (Backpressure).
- *Cancellation:* "stop" sets a flag checked between decode steps, freeing the sequence's KV blocks at once. A silent disconnect does *not* cancel (the user may reconnect), but a generation unread for 60 s is cancelled.

Deep dive 2 — GPU inference: prefill, decode, continuous batching and KV-cache memory.

- *Two phases.* **Prefill** processes all prompt tokens in parallel, compute-bound, producing the first token. **Decode** produces one token per step per sequence, reading all weights and every sequence's KV cache each step: **memory-bandwidth-bound** (3.35 TB/s on an H100 SXM).
- *KV cache.* In a transformer every token yields a **key** and **value** vector per layer; keeping them (320 KB/token here) avoids recomputing the prompt for each new token.
- *Why batching is everything.* A decode step reads 140 GB ÷ 8 = 17.5 GB per GPU → ≈ **5.2 ms** whether the batch is 1 or 256 sequences; 1 sequence is ~190 tokens/s per replica. With 256 sequences of 3,500 tokens the step also reads 256 × 3,500 × 320 KB ≈ 287 GB of KV (≈36 GB per GPU, ≈10.7 ms), so ~16 ms per step → ~60 tokens/s per sequence × 256 ≈ **15K tokens/s** theoretical ceiling; real systems reach perhaps half, hence the 6,000 planning figure.
- *Static vs continuous batching.* Static runs a fixed batch until its **longest** answer ends, so short answers and arrivals wait. **Continuous batching** (iteration-level scheduling, Orca, OSDI 2022) re-forms the batch **every decode step**: finished sequences leave, waiting ones join, utilisation stays high (Model Serving (and Batching)).
- *PagedAttention.* Reserving contiguous KV for each sequence's *maximum* length wastes most of it. **PagedAttention** (vLLM, SOSP 2023) stores KV in fixed-size **blocks** (e.g. 16 tokens) allocated on demand via a block table, like OS virtual-memory pages; waste drops to the last partial block, and the paper reports 2–4× throughput at the same latency.
- *Prefill hurts decode.* A 30,000-token prefill can stall every decoding sequence for hundreds of milliseconds (a visible "hiccup"). Fixes: **chunked prefill** (split long prefills to share steps with decodes; the simple default) or **prefill/decode disaggregation** (separate pools per phase, shipping KV between them; for large fleets with long prompts).
- *Admission control.* Admit a sequence only if free KV blocks cover its prompt plus an output reserve; if memory runs out mid-generation, **preempt** the lowest-priority sequence (drop its KV, recompute later) rather than crash.

Deep dive 3 — KV-cache-aware routing and prefix caching.

- *The opportunity.* Turn 12 repeats turns 1–11 and every request shares the system prompt. If that **prefix**'s KV is still on a replica, prefill runs only on new tokens: a 60% hit rate turns 3,000

prefill tokens into 1,200, **2.5× less prefill** and much lower TTFT.

- *Finding a prefix.* vLLM's **automatic prefix caching** hashes each full KV block together with all tokens before it (so a block hash identifies the whole prefix); a new request is hashed block by block, matches reused, the rest computed; finished requests' blocks stay in an LRU pool. SGLang's **RadixAttention** uses a radix (prefix) tree.
- *The catch: the cache is per replica.* Round-robin over 3,000 replicas almost never sends turn 12 to turn 11's replica, so the **router** must be cache-aware:
  1. Compute the request's block hashes (or the hash of the first N tokens).
  2. Find replicas holding the longest matching prefix, **approximately** (remember where recent prefixes went) or **precisely** (servers publish "block stored/evicted" events into a `block_hash → replicas` index; the open-source llm-d scheduler offers both).
  3. Score `w1 × cached_fraction − w2 × queue_depth − w3 × KV_memory_used` and pick the best; load terms stop hot prefixes piling onto one replica.
- *Cheap version, session affinity:* consistent hashing on `conversation_id` picks a "home" replica, spilling to the least-loaded when home is busy; captures most multi-turn benefit.
- *Pricing follows:* public APIs charge less for cached input; OpenAI's prompt-caching guide describes routing on a hash of the prompt's first tokens. Report **cache hit rate** per pool beside GPU utilisation.
- *Not semantic caching:* prefix caching reuses **computation** and yields exactly the model's answer; Semantic Caching reuses a **previous answer** to a similar question, risky for open-ended chat.

Deep dive 4 — Quotas, rate limits per tier, priority and cost.

- *What to limit:* requests per minute protect API servers; **tokens per minute** protect GPUs (one 100K-token request costs as much as hundreds of short ones); product quotas ("N messages on the top model per few hours" for free users) protect the budget. Public LLM APIs use the same split per organisation and project.
- *Reserve, then settle:* output length is unknown, so **reserve** `input_tokens + max_output_tokens` from the user's token bucket and **refund** the unused part at the end; otherwise 50 parallel streams could each pass and overspend together.
- *Counters:* in memory, sharded by `user_id`, atomic check-and-decrement script. A few percent of over-use during failover is acceptable, so the limiter **fails open** for paid tiers and **closed** for anonymous traffic.
- *Priority and isolation:* paid traffic gets its own queues and large customers **reserved replicas**; free traffic uses the rest: a bulkhead so a free-tier surge can't starve paying users.
- *Cost per message* = uncached prefill tokens × prefill cost + output tokens × decode cost. Levers, biggest first: route easy requests to a **smaller model**; raise **batch size** and **cache hit rate**; cap `max_output_tokens` per tier; **auto-scale** pools over the day and fill idle night GPUs with batch jobs (GPU Auto-Scaling). Loading 140 GB of weights takes minutes, so scaling follows a forecast.

Deep dive 5 — Safety filters without killing latency.

- *Input:* a small classifier (e.g. Meta's Llama Guard, a fine-tuned model labelling prompts and responses against a risk taxonomy) runs **in parallel** with prompt building and quota checks, adding little TTFT; blocked prompts never reach the large model.
- *Output, while streaming:* hold back a short window (say 20–50 tokens) and classify the text so far each time it fills; if flagged, end with `error{code: "content_filtered"}` and the client replaces the shown text. A slightly later first token buys never showing a full harmful answer; provider docs warn partial output is harder to moderate, hence the window.
- *Failure policy:* safety down → **fail closed** for free and anonymous traffic and alert; on the critical path, it gets its own capacity and a circuit breaker. Account-level abuse detection runs offline on the usage stream.

7. The failure walk

1. **Replica crashes mid-generation:** its KV is gone; the orchestrator **re-submits prompt plus tokens so far** to another replica, which prefills and continues. A 1–2 s pause, `seq` numbers continue (no duplicates); after three failures the partial answer is saved as `failed`.
2. **Client connection drops:** the GPU keeps writing to the buffer; the client reconnects with `Last-Event-ID` to any gateway and resumes.
3. **Router restarts, losing its prefix index:** correctness unaffected (a cache of a cache); hit rate and TTFT suffer for minutes while it refills.
4. **Launch-day 5× spike:** admission control caps per-tier queue length; beyond it free users get "at capacity" or the smaller model (Graceful Degradation); paid traffic keeps reserved replicas; pre-scale before announced launches.
5. **Conversation store slow or down:** generation proceeds from the history the client already shows; writes go to a durable retry queue; listing old chats shows an error banner.
6. **Pathological prompt** (a 200K-token paste, or one built to make the model loop): per-tier context limits, `max_output_tokens` and a per-request time budget bound it; tokens-per-minute limits stop repeats.
7. **A region loses its GPUs:** stateless orchestrators send new and retried generations to another region within data-residency rules (Multi-Region and Cell-Based Architecture).

8. Trade-offs

Decision	Chosen	Alternative	Why
Streaming transport	SSE over HTTP/2	WebSockets	One-way traffic, built-in resume, proxy-friendly
GPU to client coupling	Generation buffer with `seq`	Direct socket write	Survives disconnects, enables resume and remote stop
Batching	Continuous batching + paged KV	Static batches	Higher GPU use, short answers return early
Replica choice	Prefix-cache score + load	Round robin	Round robin destroys multi-turn cache hits
Long prompts	Chunked prefill	Disaggregated prefill/decode pools	Simpler; disaggregate later at larger scale
Quota accounting	Reserve max then settle	Charge after completion	Prevents parallel overspend
Output safety	Incremental check with small hold-back	Check after full answer	Harmful text never fully shown, small latency cost

9. The follow-up questions

- **"Measuring and setting latency targets?"** Track **TTFT** (queueing + prefill) and **TPOT** separately per model and tier, plus tokens/s per GPU. Bigger batches raise throughput *and* TPOT, so cap batch size to keep TPOT under reading speed, then scale replicas for TTFT. *Trap:* reporting only average latency.
- **"1M-token contexts?"** At 320 KB/token that's ~320 GB, nearly a replica's free memory for one user: KV quantisation (8-bit or lower), offloading cold KV blocks to CPU memory or SSD, and a separate long-context pool with its own quota and price.
- **"Tools and web search?"** The model emits a tool call; the orchestrator pauses the stream (`event: tool_call`), runs the tool, appends the result and resumes; the prefix is unchanged, so the request returns to the same replica to reuse its KV.
- **"Rolling out a new model?"** A new pool, a canary traffic share, compare quality, safety and cost per message, then shift gradually (Canary Deployment).

Pattern manifest

Pattern	Where it served	Module
Server-Sent Events	Token streaming, `Last-Event-ID` resume	Module 02
Message Queue	Usage events to billing and analytics	Module 02
Sharding	Conversations by user, messages by conversation	Module 03
Consistent Hashing	Home replica per conversation for cache affinity	Module 03
Idempotency	`client_msg_id` on send, `seq` on stream chunks	Module 05
Bulkhead	Separate queues and reserved capacity per tier	Module 05
Circuit Breaker	Safety service and conversation store calls	Module 05
Graceful Degradation	Smaller model or "at capacity" for free tier under load	Module 05
Rate Limiting	Requests and tokens per minute per tier	Module 08
Cursor-Based Pagination	Conversation list, stream resume cursor	Module 08
Backpressure	Slow clients read from the buffer, never block GPUs	Module 10
Model Serving (and Batching)	Continuous batching, paged KV cache	Module 11
GPU Auto-Scaling	Forecast-driven pool sizing over the day	Module 11

Interview one-liner: "The web tier is ordinary; the GPUs are the system. Stream tokens over SSE from a resumable generation buffer, batch continuously with a paged KV cache, route each turn to the replica that already holds its prefix, meter tokens with reserve-then-settle quotas per tier, and check safety incrementally on the stream."

TL;DR

- **Problem:** 350K concurrent streamed answers, each holding about 1 GB of GPU memory, at a cost that dominates the business.
- **Mechanism:** API gateway with tier limits → orchestrator (prompt, quota, safety) → KV-cache-aware router → GPU pools with continuous batching and PagedAttention → generation buffer → SSE.
- **Hard parts:** GPU memory (KV sizing, admission, preemption), cache-aware routing without hot spots, resumable streams, token quotas with unknown output length, safety on partial text.
- **Price:** a huge, slow-to-scale GPU fleet; routing tightly coupled to engine internals; per-tier policy complexity.

Sources

- Kwon et al., "Efficient Memory Management for Large Language Model Serving with PagedAttention", SOSP 2023 — https://arxiv.org/abs/2309.06180
- Yu et al., "Orca: A Distributed Serving System for Transformer-Based Generative Models", OSDI 2022 — https://www.usenix.org/conference/osdi22/presentation/yu
- Zheng et al., "SGLang: Efficient Execution of Structured Language Model Programs" (RadixAttention) — https://arxiv.org/abs/2312.07104
- vLLM docs, Automatic Prefix Caching — https://docs.vllm.ai/design/prefix_caching.html
- llm-d docs, precise prefix-cache-aware routing — https://llm-d.ai/docs/guides/precise-prefix-cache-aware
- OpenAI API docs: rate limits, streaming responses, prompt caching — https://developers.openai.com/api/docs/guides/rate-limits , https://developers.openai.com/api/docs/guides/streaming-responses , https://developers.openai.com/api/docs/guides/prompt-caching

- Meta Llama 3.1 70B model configuration (layers, KV heads, hidden size) — https://huggingface.co/unsloth/Meta-Llama-3.1-70B-Instruct/raw/main/config.json
- NVIDIA H100 product specifications (80 GB, 3.35 TB/s) — https://www.nvidia.com/en-us/data-center/h100/
- Inan et al., "Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations" — https://arxiv.org/abs/2312.06674
- WHATWG HTML Standard, Server-sent events (`Last-Event-ID`) — https://html.spec.whatwg.org/multipage/server-sent-events.html

---

Case 4 — Distributed Cache Cluster (Redis Cluster / Memcached-like)

The prompt (in short). The shared in-memory key-value cache a large company's services use: sub-millisecond `GET` / `SET`, terabytes across many machines, surviving machine failures, resizing without downtime (Redis Cluster, Memcached behind Facebook's mcrouter, ElastiCache, Memorystore).

What makes it hard. (1) **Spreading keys** over 100+ nodes so every client finds the right one without a central lookup; (2) **moving keys while traffic flows**; (3) **failover** with asynchronous replication, leaving a window where acknowledged writes are lost; (4) **skew**: one celebrity key can exceed any node's capacity, and one dead node can stampede misses onto the database.

1. Requirements and numbers

Functional: `GET`, `SET` with TTL, `DEL`, multi-key `MGET`, atomic `INCR`, compare-and-set; hashes and sorted sets for the Redis flavour; online node add and remove; per-tenant namespaces.

Non-functional: single-key p99 < 1 ms in a data centre; ≥ 99.95% availability; linear scaling; a node failure loses at most a few seconds of writes and recovers in seconds; a cache outage must not take down the database behind it.

Numbers (assumptions):

- **10M ops/s at peak**, 90% reads.
- **2B keys** × (1 KB value + ~100 bytes overhead: key string, pointers, TTL, eviction metadata) ≈ **2.2 TB**.
- **Planning throughput:** one single-threaded **shard** (process owning part of the key space) handles **~100K ops/s** with headroom (an assumption; measure with `redis-benchmark` or `memtier_benchmark`).
- Shards for throughput: **100 primaries**; for memory: 2.2 TB ÷ 48 GB ≈ 46. **Throughput decides**: 100 primaries of **22 GB** each.
- One **replica** per primary → **200 shard processes**, 4 per machine (32 cores, 256 GB RAM, leaving room for fork copy-on-write and buffers) → **50 machines**, each replica on a different machine and rack.
- Network: 10M × 1.1 KB ≈ 11 GB/s ≈ 88 Gbit/s ÷ 50 ≈ **1.8 Gbit/s per machine**, comfortable on 25 Gbit/s NICs, but one hot key can saturate a machine's link (Deep dive 4).
- Redis Cluster's **16,384 hash slots** ÷ 100 ≈ **164 slots per primary**.

What the numbers say: per-shard CPU and per-machine network limit us, not memory: spread *load*, not just *data*.

2. The five-questions pass

- **Flow:** synchronous, sub-millisecond, over long-lived pooled TCP connections with **pipelining** (several commands sent before reading replies) (Five Questions).
- **Storage:** RAM is the store; for cache-aside data the **database behind** is truth and the cache can always refill. When Redis is a primary store (sessions, counters, leaderboards), persistence and replication matter more (Deep dive 5).
- **Speed:** one round trip to the right node, never two, so the **routing table lives in the client or a proxy**.
- **Failure:** detect in seconds, promote a replica, accept a small loss window, protect the database from the miss storm.

- **Growth:** add nodes by **moving key slices**, never rehashing everything; hot spots are viral keys, big keys (a 50 MB list) and keys expiring together.
- *Skeleton:* primary shards with replicas, a key → slot → shard map in every client, a failure detector, a migration protocol, per-shard eviction and persistence.

3. API

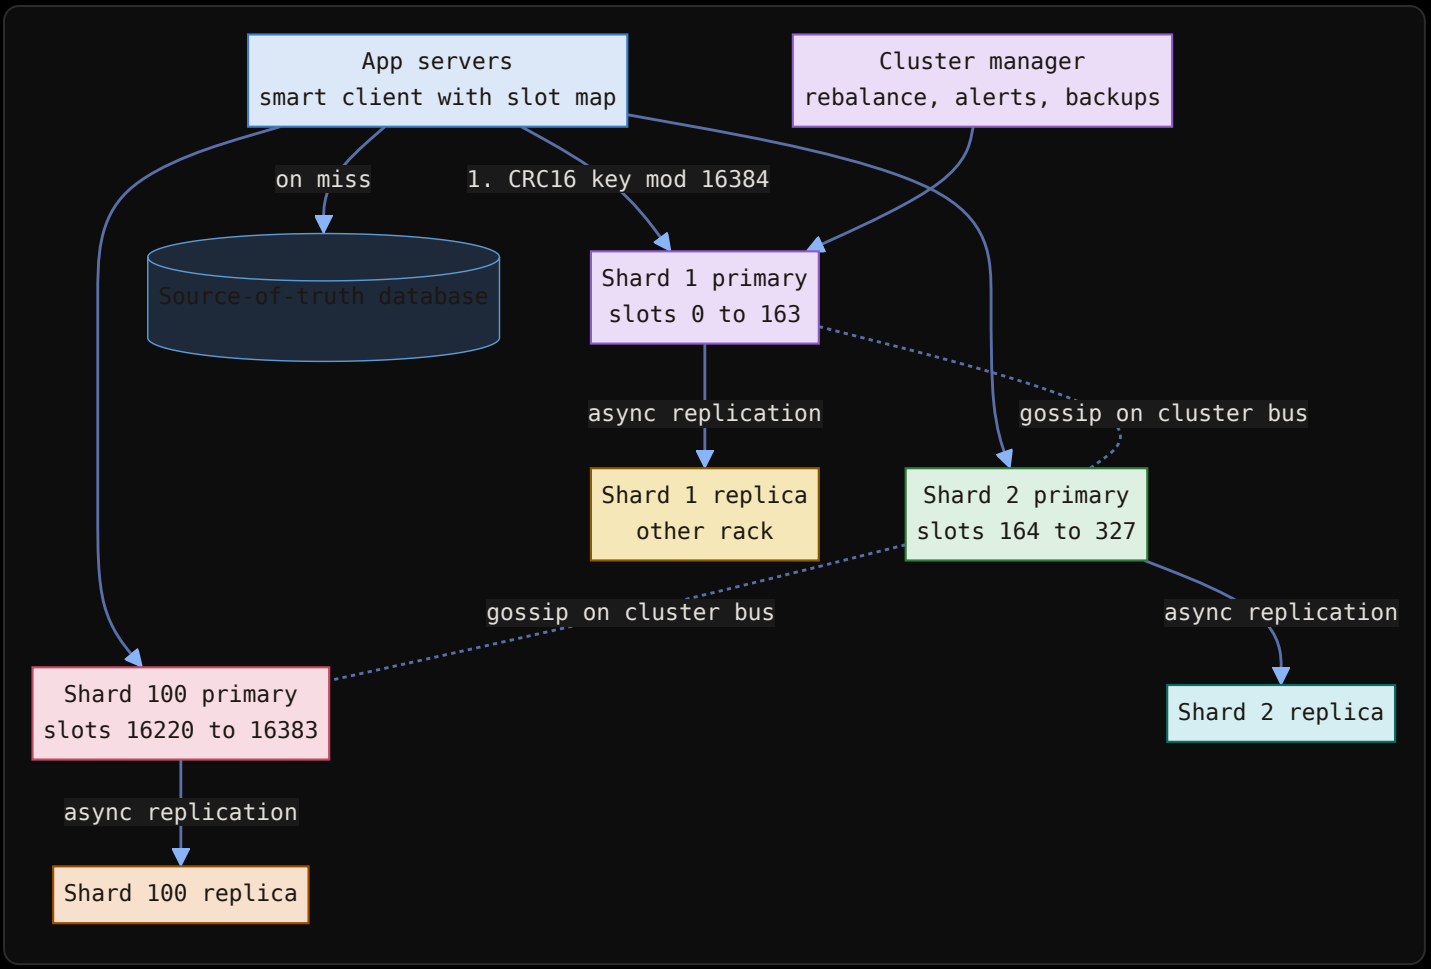
Command	Meaning
`GET k` / `MGET k1 k2 …`	Read; `MGET` across shards is split by the client and fanned out in parallel
`SET k v [EX seconds] [NX]`	Write with optional TTL; `NX` = only if absent (used for locks and leases)
`DEL k`, `EXPIRE k s`	Delete, set TTL
`INCR k`	Atomic counter on the owning shard
`gets k` → `(v, cas_token)` / `cas k v cas_token`	Memcached compare-and-set: the write succeeds only if nobody changed the value since the read
`WAIT numreplicas timeout`	Redis: block until this connection's earlier writes reach N replicas (narrows, not removes, the loss window)
`CLUSTER SLOTS` / `CLUSTER SHARDS`	Redis: return the slot → node map so a client can route directly
Admin: `add-node`, `reshard --from A --to B --slots 164`, `failover`	Operator or automation commands

Routing errors: `-MOVED 3999 10.0.0.7:6379` ("this slot lives there now, update your map") and `-ASK 3999 10.0.0.8:6379` ("slot migrating, try there for this one request only").

4. Data model

- **Inside a shard:** a hash table `key → object` with the value (string, hash, list, set, sorted set), optional expiry, and a few bits of eviction metadata (Redis: an approximate last-access clock or a small logarithmic frequency counter).
- **Memcached memory layout:** 1 MB **pages** assigned to **slab classes**, each cutting pages into fixed chunks (80 B, 104 B, 136 B, …); an item takes the smallest fitting chunk and each class has its own LRU. This avoids `malloc` fragmentation but wastes the item-to-chunk gap.
- **Slot map:** `slot (0..16383) → {primary_node_id, replica_node_ids, config_epoch}`, held by every node and cached by clients. The **config epoch** increases on every failover or slot move, so the newest claim wins disagreements.
- **Node table:** `node_id → address, role, flags (ok / PFAIL / FAIL), last pong, slots`.
- **Key naming:** `tenant:entity:id` (`shop:user:42`), with a **hash tag** in braces (`shop:{user:42}:cart`) when keys must share a shard (Deep dive 1).

5. The architecture, annotated



1. At startup the **smart client** (a library that knows the cluster layout) calls `CLUSTER SHARDS` on any node and caches the 16,384-entry **slot map**.
2. For `GET user:42`: `slot = CRC16(key) mod 16384` (**CRC16**, a cheap 16-bit checksum used as a hash; Redis Cluster uses 14 of its bits), then the slot's primary on a pooled connection: **one hop**.
3. The shard answers from RAM; on a miss the app reads the database and writes back with a TTL (Cache-Aside).
4. Writes go to the primary, which replies immediately and **streams the write asynchronously** to its replica (Primary-Replica).
5. Nodes **gossip** on a separate **cluster bus** (by default data port + 10,000) about health and slot ownership; no coordinator on the request path.
6. The **cluster manager** (automation off the data path) watches load, moves slots, runs backups and pages humans.

Memcached variant: servers know nothing of each other; distribution lives in the client (consistent hashing) or a proxy such as Facebook's **mcrouter** or Twitter's **twemproxy**, which can also replicate by writing to several pools.

6. Deep dives

Deep dive 1 — Sharding: hash slots vs consistent hashing, and hash tags.

- *Plain modulo* (`hash(key) mod N`) moves almost every key when N goes from 100 to 101: a near-total cache wipe.
- **Consistent hashing** (the Memcached tradition, often **ketama**): servers at many **virtual node** points on a ring, a key owned by the first server clockwise; adding a server moves about 1/(N+1) of keys and needs no metadata beyond the server list, suiting stateless clients and proxies (Consistent Hashing; other schemes in Module 19).

- **Hash slots** (Redis Cluster) add indirection: keys hash to one of a **fixed** 16,384 slots and an **explicit table** maps slots to nodes. `CRC16("user:42")` = 0x3E08 = 15,880 → **slot 15,880**, on shard [15,880 ÷ 163.84] + 1 = **shard 97** (implementation check from the spec: CRC16("123456789") = 0x31C3). Adding a node **moves whole slots**, and operators can move slots off a hot node; the cost is a table every client must keep current.
- **Which:** slots for precise control and a versionable ownership record for failover; a ring for no shared metadata. Both move few keys on resize (Sharding).
- **Hash tags.** Multi-key operations (`MGET` , transactions, Lua) need all keys on one shard; with `{…}` only the first braces' content is hashed, so `{user:42}:profile` and `{user:42}:cart` both land in slot 15,880. Trap: over-tagging (`{tenant}` on a big tenant's keys) recreates a hot shard.
- **Resharding without downtime (Redis).** Moving slot 8 from A to B: mark B `IMPORTING 8` and A `MIGRATING 8` ; repeatedly `MIGRATE` batches of slot-8 keys from A to B (each key atomically). A still answers keys it holds and replies `-ASK B` for moved ones, which the client sends to B prefixed with `ASKING` . When A's slot is empty both get `SETSLOT 8 NODE B` , the new owner bumps its config epoch, and A answers `-MOVED` , updating client maps. Moving 164 slots (22 GB) at a throttled 100 MB/s takes ~4 minutes per added node, with a small share of requests paying an extra hop.

Deep dive 2 — Replication and failover.

- *Asynchronous replication:* the primary acks before the replica has the write, so **an acknowledged write can be lost** if the primary dies first; the Redis Cluster spec states this window and calls its merge rule "last failover wins". `WAIT 1 100` blocks until one replica has the write, shrinking the window for important keys, but it is not consensus.
- *Failure detection:* nodes ping over the cluster bus; no answer within `cluster-node-timeout` (e.g. 5 s) → **PFAIL** (a local opinion); a **majority of primaries** reporting PFAIL/FAIL within twice the timeout → **FAIL**, broadcast cluster-wide (gossip-based detection: Module 19 — Failure detection).
- *Election:* the failed primary's replicas wait a delay ranked by how current their copy is (most current first), then request votes; a primary votes for at most one replica of a given failed primary within twice the timeout, so two can't both win. A replica with a **majority of primaries'** votes promotes itself with a new, higher config epoch and announces its slots; clients get `-MOVED` and refresh. Outage for those slots ≈ node timeout + 1–2 s.
- *Partitions:* a primary cut off with a minority keeps accepting writes until the node timeout, then **stops**; those are the writes that can be lost when the majority promotes a replica. A smaller timeout means less loss but more false failovers on GC-like pauses or network blips.
- *Placement:* never a primary and its replica on one machine or rack; **replica migration** moves a spare replica to a primary that lost its own, so a second failure doesn't leave slots uncovered.
- *Memcached tier:* no replication, so a dead node is a block of misses. Facebook's small **Gutter** pool (~1% of servers) takes retried failed gets with quickly expiring entries, shielding the database.

Deep dive 3 — Routing: smart client vs proxy.

- **Smart client:** holds the slot map or ring and connects to every shard. *Pros:* one hop, lowest latency, no extra servers. *Cons:* every language needs a correct client (MOVED/ASK, map refresh, pooling), and 5,000 app servers × 100 shards × 4 connections = **2M TCP connections**, each costing shard memory.
- **Proxy:** clients speak the plain protocol to a nearby proxy (mcrouter, twemproxy, Envoy's Redis proxy) that knows the topology. *Pros:* thin clients, **connection pooling and pipelining**, one place for prefix routing, replicated pools, shadow traffic and failover policy. *Cons:* an extra hop (~0.1–0.3 ms) and more machines. mcrouter's README cites almost 5 billion requests/s at peak across Facebook and Instagram.
- **Sidecar proxy** (one per app host on localhost) is the common middle: the hop is a local socket and connections drop to proxies × shards (Sidecar).

- Facebook's NSDI 2013 hybrid: **gets go directly to memcached over UDP** (connectionless, lower latency; a lost packet counts as a miss), **sets and deletes via a local mcrouter over TCP** for reliability.

Deep dive 4 — Hot keys, big keys and thundering herds.

- *The problem:* a viral key at 500K reads/s vs one shard's ~100K ops/s budget, and 500K × 1 KB = 500 MB/s ≈ 4 Gbit/s on one link. Even keys don't mean even *traffic*.
- *Find them:* top-K counters on sampled commands at client or proxy (`redis-cli --hotkeys` under an LFU policy); alert on per-shard skew.
- *Fix 1, near cache:* clients keep the hottest few thousand keys in process for 1–5 s; 500K reads/s over 2,000 app servers is 250 each, fully absorbed. Staleness is bounded by the TTL or removed with Redis 6+ **client tracking** (invalidation messages when a read key changes).
- *Fix 2, key splitting:* N copies `post:9:#0 … post:9:#7` in different slots; readers pick one at random, writers update all eight; good for read-mostly data.
- *Fix 3, replica reads:* send hot-slot reads to replicas (`READONLY`), accepting slight staleness.
- *Thundering herd:* when a hot key expires thousands miss at once. **Leases** (Facebook's memcache paper): on a miss the server gives a lease token to **one** client, at most once every 10 seconds per key; others briefly wait and retry, and a set is accepted only with a valid token (also blocking stale sets after a delete). The paper reports a database peak falling from 17K/s to 1.3K/s for herd-prone keys (Cache Stampede Prevention).
- *Synchronized expiry:* **jitter** TTLs (e.g. 3,600 s ± 10%).
- *Big keys:* a 50 MB list blocks a single-threaded shard while read or deleted; cap value size, bucket collections, delete with `UNLINK` (background free).

Deep dive 5 — Eviction, memory and persistence.

- *Eviction* at `maxmemory` , by `maxmemory-policy` : `allkeys-lru` (usual for a pure cache), `allkeys-lfu` (better when a stable key set is popular), `volatile-*` (only TTL keys), `allkeys-random` , `volatile-ttl` (soonest-expiring first), or `noeviction` (reject writes; right for a primary store). Redis approximates LRU by **sampling** a few keys (default 5) and evicting the oldest, far cheaper than a true LRU list; LFU uses a small decaying logarithmic counter per key (policy families: Module 19 — Cache eviction).
- *Headroom:* keep 20–30% free: replication and AOF buffers don't count toward `maxmemory` , and the persistence fork's **copy-on-write** (pages copied only when the parent modifies them) can nearly double memory under heavy writes.
- *Persistence (Redis):*
  - **None:** pure cache; restart = empty shard refilled from the database. Fastest, simplest.
  - **RDB snapshots:** a forked child writes a compact point-in-time file every few minutes: fast restarts, good backups, minutes of writes lost on crash.
  - **AOF (append-only file):** every write appended and fsynced per `appendfsync` : `always` (safest, slowest), `everysec` (default, lose at most about 1 s), `no` (OS decides); rewritten in the background to stay compact (Write-Ahead Log).
  - **RDB + AOF:** the Redis docs' recommendation for PostgreSQL-like safety; AOF is used on restart as the more complete.
- *Choice:* the shared cache tier runs **no persistence** on primaries and RDB on replicas for warm restarts; primary-store namespaces (sessions, rate-limit counters) get a separate cluster with AOF `everysec` and `noeviction` . Mixing "may evict" and "must keep" data on one cluster is the classic mistake.
- *Warm-up:* route reads to a new or restarted node gradually, or pre-fill it from a replica or snapshot, so its misses don't spike the database (Graceful Degradation).

7. The failure walk

1. **Primary crashes:** its 164 slots time out; after the node timeout (5 s) plus election (~1–2 s) the replica is promoted and clients follow `-MOVED`. Writes acked in the last instant may be lost. Apps fall back to the database for those keys behind a circuit breaker and request coalescing.
2. **A machine (4 shards) dies:** replicas elsewhere fail over in parallel; capacity drops 4% until a replacement resyncs.
3. **Partition isolates 10 primaries:** the minority side refuses writes after the node timeout and the majority promotes replicas; on healing, old primaries see a higher config epoch and become replicas, discarding divergent writes.
4. **Cache-wide cold start** (a deploy bug flushed everything): hit rate zero and 9M reads/s would hit the database. Leases or single-flight per key, per-app database concurrency caps (load shedding), stale or default data for non-critical features, gradual refill; never assume the database absorbs a full miss storm.
5. **Hot key overloads a shard** (p99 up, CPU 100%): client near-caches absorb reads; ops can move that slot alone to a less-loaded node.
6. **Resharding stalls mid-slot** (target dies): the slot stays `MIGRATING` and the source serves keys it still has; fix the target or migrate back. Each `MIGRATE` is atomic, so no key is in two places.
7. **Fork pause** during a snapshot: page-table copying stalls the process briefly, longer for bigger shards; keep shards small and snapshot on replicas.

8. Trade-offs

Decision	Chosen	Alternative	Why
Key placement	16,384 hash slots + slot map	Consistent hashing ring	Explicit ownership, targeted moves, works with failover epochs
Routing	Smart client (sidecar proxy where needed)	Central proxy tier	One hop; proxies only where clients are weak
Replication	Async, one replica per primary	Synchronous or consensus per write	Sub-ms writes; small loss window acceptable for a cache
Failure detection	Gossip + majority-of-primaries vote	External coordinator (Sentinel, ZooKeeper)	No extra system; slots and health in one protocol
Shard size	Many small shards (22 GB)	Few large shards	Faster failover and resync, shorter fork pauses, spreads CPU
Eviction	`allkeys-lru` for cache, `noeviction` for stores	One policy everywhere	Different data, different guarantees
Persistence	None on cache primaries, RDB on replicas	AOF everywhere	Cache is rebuildable; disk I/O costs latency
Hot keys	Near cache + key splitting + leases	Bigger machines	Skew beats any single machine

9. The follow-up questions

- **"Consistent with the database?"** Not strictly: cache-aside with **delete on write**, short TTLs, and invalidation from the database's change log (Change Data Capture; Facebook's daemons tail the commit log to issue deletes); leases stop a slow reader writing an old value after a delete. *Trap:* claiming atomic cache-and-database updates without a distributed transaction.
- **"Why not just Memcached?"** Multithreaded and lean for plain strings, but no replication, persistence or data structures: right for a pure look-aside cache behind a router; Redis for structures, atomic ops or failover. (Since 2024 Redis has changed licences, and the Linux Foundation's **Valkey** fork continues the BSD-licensed codebase; managed caches offer either.)
- **"Across regions?"** One cluster per region caching its local database replica, invalidations following database replication (Multi-Region and Cell-Based Architecture).

- **"Add 20% capacity at peak?"** Add 20 primaries with replicas; at 120 primaries each owns 16,384 ÷ 120 ≈ 137 slots, so every old primary gives up ~27 and every new one receives ~137, moved a few at a time watching p99, with ASK/MOVED keeping traffic flowing. Better: scale before the peak.
- **"Multi-tenancy?"** Separate clusters for noisy or critical tenants (a bulkhead); prefixes plus per-tenant proxy quotas for small ones.

Pattern manifest

Pattern	Where it served	Module
Sharding	16,384 slots spread over 100 primaries	Module 03
Consistent Hashing	Memcached clients and proxies (ketama ring)	Module 03
Primary-Replica	Async replica per shard, promotion on failure	Module 03
Write-Ahead Log	AOF persistence for store-mode clusters	Module 03
Cache-Aside	How applications fill and invalidate the cache	Module 04
Cache Stampede Prevention	Leases, single-flight, TTL jitter	Module 04
Circuit Breaker	Database fallback during shard failover	Module 05
Bulkhead	Separate clusters for cache vs store, noisy tenants	Module 05
Graceful Degradation	Gradual warm-up, stale data on cold start	Module 05
Sidecar	Local proxy for routing and connection pooling	Module 08
Change Data Capture	Database-driven invalidation	Module 10

Interview one-liner: "Hash keys to 16,384 slots on many small primaries with async replicas on other racks; smart clients route in one hop and follow MOVED/ASK during migration; gossip detects failures and a majority vote promotes replicas; skew is the real enemy, fought with near caches, key splitting, leases and TTL jitter."

TL;DR

- **Problem:** 10M ops/s at sub-millisecond latency over 2.2 TB, with online resizing and failover.
- **Mechanism:** hash slots (or a consistent-hash ring) → primary shards with async replicas → routing table in smart clients or proxies → gossip failure detection and replica promotion → per-shard eviction and optional persistence.
- **Hard parts:** moving keys under load (MIGRATING/IMPORTING with ASK), the async-replication loss window, hot and big keys, protecting the database from miss storms.
- **Price:** small windows of lost writes, client complexity or an extra proxy hop, capacity planned around CPU and network rather than memory.

Sources

- Redis Cluster specification (16,384 slots, CRC16, hash tags, MOVED/ASK, gossip, PFAIL/FAIL, elections, write-loss windows) — https://redis.io/docs/latest/operate/oss_and_stack/reference/cluster-spec/
- Redis docs, Key eviction (policies, sampled LRU, LFU) — https://redis.io/docs/latest/develop/reference/eviction/
- Redis docs, Persistence (RDB, AOF, `appendfsync`) — https://redis.io/docs/latest/operate/oss_and_stack/management/persistence/
- Redis docs, Client-side caching reference (tracking, invalidation) — https://redis.io/docs/latest/develop/reference/client-side-caching/

- Nishtala et al., "Scaling Memcache at Facebook", NSDI 2013 (leases, Gutter, UDP gets, mcrouter, mcsqueal) — https://www.usenix.org/system/files/conference/nsdi13/nsdi13-final170_update.pdf
- mcrouter wiki (features, scale) — https://github.com/facebook/mcrouter/wiki
- twemproxy README (memcached/Redis proxy, ketama, pipelining) — https://github.com/twitter/twemproxy
- memcached wiki, User Internals (pages, slab classes, per-class LRU) — https://github.com/memcached/memcached/wiki/UserInternals
- Linux Foundation launches Valkey (Phoronix report, March 2024) — https://www.phoronix.com/news/Linux-Foundation-Valkey